
Sage Reference Manual: Matrices and Spaces of Matrices

Release 6.3

The Sage Development Team

August 11, 2014

CONTENTS

| | | |
|-----------|--|------------|
| 1 | Matrix Spaces | 3 |
| 2 | Matrix Constructor | 15 |
| 3 | Matrices over an arbitrary ring | 59 |
| 3.1 | Indexing | 60 |
| 4 | Abstract base class for matrices | 67 |
| 5 | Base class for matrices, part 0 | 69 |
| 6 | Base class for matrices, part 1 | 101 |
| 7 | Base class for matrices, part 2 | 123 |
| 8 | Generic Asymptotically Fast Strassen Algorithms | 275 |
| 9 | Minimal Polynomials of Linear Recurrence Sequences | 279 |
| 10 | Base class for dense matrices | 281 |
| 11 | Base class for sparse matrices | 285 |
| 12 | Dense Matrices over a general ring | 293 |
| 13 | Sparse Matrices over a general ring | 295 |
| 14 | Dense matrices over $\mathbb{Z}/n\mathbb{Z}$ for n small | 297 |
| 15 | Sparse matrices over $\mathbb{Z}/n\mathbb{Z}$ for n small | 303 |
| 16 | Symbolic matrices | 307 |
| 17 | Dense matrices over the integer ring | 317 |
| 18 | Dense matrices over the rational field | 343 |
| 19 | Dense matrices using a NumPy backend. | 353 |
| 20 | Dense matrices over the Real Double Field using NumPy | 395 |
| 21 | Dense matrices over the Complex Double Field using NumPy | 397 |

| | |
|--|------------|
| 22 Dense matrices over multivariate polynomials over fields | 399 |
| 23 Operation Tables | 403 |
| 24 Benchmarks for matrices | 413 |
| 25 Indices and Tables | 423 |
| Bibliography | 425 |

Sage provides native support for working with matrices over any commutative or noncommutative ring. The parent object for a matrix is a matrix space `MatrixSpace(R, n, m)` of all $n \times m$ matrices over a ring R .

To create a matrix, either use the `matrix(...)` function or create a matrix space using the `MatrixSpace` command and coerce an object into it.

Matrices also act on row vectors, which you create using the `vector(...)` command or by making a `VectorSpace` and coercing lists into it. The natural action of matrices on row vectors is from the right. Sage currently does not have a column vector class (on which matrices would act from the left), but this is planned.

In addition to native Sage matrices, Sage also includes the following additional ways to compute with matrices:

- Several math software systems included with Sage have their own native matrix support, which can be used from Sage. E.g., PARI, GAP, Maxima, and Singular all have a notion of matrices.
- The GSL C-library is included with Sage, and can be used via Cython.
- The `scipy` module provides support for *sparse* numerical linear algebra, among many other things.
- The `numpy` module, which you load by typing `import numpy` is included standard with Sage. It contains a very sophisticated and well developed array class, plus optimized support for *numerical linear algebra*. Sage's matrices over RDF and CDF (native floating-point real and complex numbers) use `numpy`.

Finally, this module contains some data-structures for matrix-like objects like operation tables (e.g. the multiplication table of a group).

MATRIX SPACES

You can create any space $\text{Mat}_{n \times m}(R)$ of either dense or sparse matrices with given number of rows and columns over any commutative or noncommutative ring.

EXAMPLES:

```
sage: MS = MatrixSpace(QQ, 6, 6, sparse=True); MS
Full MatrixSpace of 6 by 6 sparse matrices over Rational Field
sage: MS.base_ring()
Rational Field
sage: MS = MatrixSpace(ZZ, 3, 5, sparse=False); MS
Full MatrixSpace of 3 by 5 dense matrices over Integer Ring
```

TESTS:

```
sage: matrix(RR, 2, 2, sparse=True)
[0.0000000000000000 0.0000000000000000]
[0.0000000000000000 0.0000000000000000]
sage: matrix(GF(11), 2, 2, sparse=True)
[0 0]
[0 0]
```

```
class sage.matrix.matrix_space.MatrixSpace(base_ring, nrows, ncols=None, sparse=False)
  Bases: sage.structure.unique_representation.UniqueRepresentation,
         sage.structure.parent_gens.ParentWithGens
```

The space of all $nrows \times ncols$ matrices over $base_ring$.

INPUT:

- `base_ring` - a ring
- `nrows` - int, the number of rows
- `ncols` - (default `nrows`) int, the number of columns
- `sparse` - (default `false`) whether or not matrices are given a sparse representation

EXAMPLES:

```
sage: MatrixSpace(ZZ, 10, 5)
Full MatrixSpace of 10 by 5 dense matrices over Integer Ring
sage: MatrixSpace(ZZ, 10, 5).category()
Category of modules over euclidean domains
```

```
sage: MatrixSpace(ZZ, 10, 2^31)
```

```
...
```

```
# 32-bit
```

```
ValueError: number of rows and columns must be less than 2^31 (on a 32-bit computer -- use a 64-
```

```
Full MatrixSpace of 10 by 2147483648 dense matrices over Integer Ring # 64-bit
sage: MatrixSpace(ZZ, 2^31, 10)
... # 32-bit
ValueError: number of rows and columns must be less than 2^31 (on a 32-bit computer -- use a 64-
Full MatrixSpace of 2147483648 by 10 dense matrices over Integer Ring # 64-bit
sage: MatrixSpace(ZZ, 10, 10).category()
Category of algebras over euclidean domains
sage: MatrixSpace(QQ, 10).category()
Category of algebras over quotient fields
```

base_extend(R)

Return base extension of this matrix space to R.

INPUT:

•R - ring

OUTPUT: a matrix space

EXAMPLES:

```
sage: Mat(ZZ, 3, 5).base_extend(QQ)
Full MatrixSpace of 3 by 5 dense matrices over Rational Field
sage: Mat(QQ, 3, 5).base_extend(GF(7))
Traceback (most recent call last):
...
TypeError: no base extension defined
```

basis()

Returns a basis for this matrix space.

Warning: This will of course compute every generator of this matrix space. So for large matrices, this could take a long time, waste a massive amount of memory (for dense matrices), and is likely not very useful. Don't use this on large matrix spaces.

EXAMPLES:

```
sage: Mat(ZZ, 2, 2).basis()
[
[1 0] [0 1] [0 0] [0 0]
[0 0], [0 0], [1 0], [0 1]
]
```

cached_method(f, name=None, key=None)

A decorator for cached methods.

EXAMPLES:

In the following examples, one can see how a cached method works in application. Below, we demonstrate what is done behind the scenes:

```
sage: class C:
....:     @cached_method
....:     def __hash__(self):
....:         print "compute hash"
....:         return int(5)
....:     @cached_method
....:     def f(self, x):
....:         print "computing cached method"
....:         return x*2
```



```
sage: c = C()
sage: type(C.__hash__)
<type 'sage.misc.cachefunc.CachedMethodCallerNoArgs'>
sage: hash(c)
compute hash
5
```

When calling a cached method for the second time with the same arguments, the value is gotten from the cache, so that a new computation is not needed:

```
sage: hash(c)
5
sage: c.f(4)
computing cached method
8
sage: c.f(4) is c.f(4)
True
```

Different instances have distinct caches:

```
sage: d = C()
sage: d.f(4) is c.f(4)
computing cached method
False
sage: d.f.clear_cache()
sage: c.f(4)
8
sage: d.f(4)
computing cached method
8
```

Using cached methods for the hash and other special methods was implemented in [trac ticket #12601](#), by means of `CachedSpecialMethod`. We show that it is used behind the scenes:

```
sage: cached_method(c.__hash__)
<sage.misc.cachefunc.CachedSpecialMethod object at ...>
sage: cached_method(c.f)
<sage.misc.cachefunc.CachedMethod object at ...>
```

change_ring(*R*)

Return matrix space over *R* with otherwise same parameters as self.

INPUT:

- *R* - ring

OUTPUT: a matrix space

EXAMPLES:

```
sage: Mat(QQ, 3, 5).change_ring(GF(7))
Full MatrixSpace of 3 by 5 dense matrices over Finite Field of size 7
```

column_space()

Return the module spanned by all columns of matrices in this matrix space. This is a free module of rank the number of columns. It will be sparse or dense as this matrix space is sparse or dense.

EXAMPLES:

```
sage: M = Mat(GF(9, 'a'), 20, 5, sparse=True); M.column_space()
Sparse vector space of dimension 20 over Finite Field in a of size 3^2
```

construction()

EXAMPLES:

```
sage: A = matrix(ZZ, 2, [1..4], sparse=True)
sage: A.parent().construction()
(MatrixFunctor, Integer Ring)
sage: A.parent().construction()[0](QQ['x'])
Full MatrixSpace of 2 by 2 sparse matrices over Univariate Polynomial Ring in x over Rational Field
sage: parent(A/2)
Full MatrixSpace of 2 by 2 sparse matrices over Rational Field
```

dimension()

Returns (m rows) * (n cols) of self as Integer

EXAMPLES:

```
sage: MS = MatrixSpace(ZZ, 4, 6)
sage: u = MS.dimension()
sage: u - 24 == 0
True
```

dims()

Returns (m row, n col) representation of self dimension

EXAMPLES:

```
sage: MS = MatrixSpace(ZZ, 4, 6)
sage: MS.dims()
(4, 6)
```

full_category_initialisation()

Make full use of the category framework.

Note: It turns out that it causes a massive speed regression in computations with elliptic curves, if a full initialisation of the category framework of matrix spaces happens at initialisation: The elliptic curves code treats matrix spaces as containers, not as objects of a category. Therefore, making full use of the category framework is now provided by a separate method (see [trac ticket #11900](#)).

EXAMPLES:

```
sage: MS = MatrixSpace(QQ, 8)
sage: TestSuite(MS).run()
sage: type(MS)
<class 'sage.matrix.matrix_space.MatrixSpace_with_category'>
sage: MS.full_category_initialisation()
doctest:...: DeprecationWarning: the full_category_initialization
method does nothing, as a matrix space now has its category
systematically fully initialized
See http://trac.sagemath.org/15801 for details.
```

gen(n)

Return the n-th generator of this matrix space.

This doesn't compute all basis matrices, so it is reasonably intelligent.

EXAMPLES:

```
sage: M = Mat(GF(7), 10000, 5); M.ngens()
50000
sage: a = M.10
```

```
sage: a[:4]
[0 0 0 0 0]
[0 0 0 0 0]
[1 0 0 0 0]
[0 0 0 0 0]
```

get_action_impl(*S, op, self_on_left*)

x.__init__(...) initializes *x*; see `help(type(x))` for signature

identity_matrix()

Returns the identity matrix in *self*.

self must be a space of square matrices. The returned matrix is immutable. Please use `copy` if you want a modified copy.

EXAMPLES:

```
sage: MS1 = MatrixSpace(ZZ, 4)
sage: MS2 = MatrixSpace(QQ, 3, 4)
sage: I = MS1.identity_matrix()
sage: I
[1 0 0 0]
[0 1 0 0]
[0 0 1 0]
[0 0 0 1]
sage: Er = MS2.identity_matrix()
Traceback (most recent call last):
...
TypeError: self must be a space of square matrices
```

TESTS:

```
sage: MS1.one()[1,2] = 3
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead (i.e., use copy(M) to change a
```

is_dense()

Returns True if matrices in *self* are dense and False otherwise.

EXAMPLES:

```
sage: Mat(RDF, 2, 3).is_sparse()
False
sage: Mat(RR, 123456, 22, sparse=True).is_sparse()
True
```

is_finite()

EXAMPLES:

```
sage: MatrixSpace(GF(101), 10000).is_finite()
True
sage: MatrixSpace(QQ, 2).is_finite()
False
```

is_sparse()

Returns True if matrices in *self* are sparse and False otherwise.

EXAMPLES:

```
sage: Mat(GF(2011), 10000).is_sparse()
False
sage: Mat(GF(2011), 10000, sparse=True).is_sparse()
True
```

matrix (*x=0, coerce=True, copy=True*)

Create a matrix in `self`.

INPUT:

- `x` – (default: 0) data to construct a new matrix from. Can be one of the following:
 - 0, corresponding to the zero matrix;
 - 1, corresponding to the identity_matrix;
 - a matrix, whose dimensions must match `self` and whose base ring must be convertible to the base ring of `self`;
 - a list of entries corresponding to all elements of the new matrix;
 - a list of rows with each row given as an iterable;
- `coerce` – (default: True) whether to coerce `x` into `self`;
- `copy` – (default: True) whether to copy `x` during construction (makes a difference only if `x` is a matrix in `self`).

OUTPUT:

- a matrix in `self`.

EXAMPLES:

```
sage: M = MatrixSpace(ZZ, 2)
sage: M.matrix([[1, 0], [0, -1]])
[ 1  0]
[ 0 -1]
sage: M.matrix([1, 0, 0, -1])
[ 1  0]
[ 0 -1]
sage: M.matrix([1, 2, 3, 4])
[1 2]
[3 4]
```

Note that the last “flip” cannot be performed if `x` is a matrix, no matter what is `rows` (it used to be possible but was fixed by Trac 10793):

```
sage: projection = matrix(ZZ, [[1, 0, 0], [0, 1, 0]])
sage: projection
[1 0 0]
[0 1 0]
sage: projection.parent()
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
sage: M = MatrixSpace(ZZ, 3, 2)
sage: M
Full MatrixSpace of 3 by 2 dense matrices over Integer Ring
sage: M(projection)
Traceback (most recent call last):
...
ValueError: a matrix from
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
```

cannot be converted to a matrix in
Full MatrixSpace of 3 by 2 dense matrices over Integer Ring!

If you really want to make from a matrix another matrix of different dimensions, use either transpose method or explicit conversion to a list:

```
sage: M(projection.list())
[1 0]
[0 0]
[1 0]
```

TESTS:

The following corner cases were problematic while working on #10628:

```
sage: MS = MatrixSpace(ZZ, 2, 1)
sage: MS([[1], [2]])
[1]
[2]
sage: MS = MatrixSpace(CC, 2, 1)
sage: F = NumberField(x^2+1, name='x')
sage: MS([F(1), F(0)])
[ 1.000000000000000]
[0.000000000000000]
```

Trac ticket #10628 allowed to provide the data be lists of matrices, but [trac ticket #13012](#) prohibited it:

```
sage: MS = MatrixSpace(ZZ, 4, 2)
sage: MS0 = MatrixSpace(ZZ, 2)
sage: MS.matrix([MS0([1, 2, 3, 4]), MS0([5, 6, 7, 8])])
Traceback (most recent call last):
...
TypeError: cannot construct an element of
Full MatrixSpace of 4 by 2 dense matrices over Integer Ring
from [[1 2]
[3 4], [5 6]
[7 8]]!
```

A mixed list of matrices and vectors is prohibited as well:

```
sage: MS.matrix([MS0([1, 2, 3, 4])] + list(MS0([5, 6, 7, 8])))
Traceback (most recent call last):
...
TypeError: cannot construct an element of
Full MatrixSpace of 4 by 2 dense matrices over Integer Ring
from [[1 2]
[3 4], (5, 6), (7, 8)]!
```

Check that [trac ticket #13302](#) is fixed:

```
sage: MatrixSpace(Qp(3), 1, 1)([Qp(3).zero()])
[0]
sage: MatrixSpace(Qp(3), 1, 1)([Qp(3)(4/3)])
[3^-1 + 1 + O(3^19)]
```

matrix_space (*nrows=None, ncols=None, sparse=False*)

Return the matrix space with given number of rows, columns and sparsity over the same base ring as self, and defaults the same as self.

EXAMPLES:

```
sage: M = Mat(GF(7), 100, 200)
sage: M.matrix_space(5000)
Full MatrixSpace of 5000 by 200 dense matrices over Finite Field of size 7
sage: M.matrix_space(ncols=5000)
Full MatrixSpace of 100 by 5000 dense matrices over Finite Field of size 7
sage: M.matrix_space(sparse=True)
Full MatrixSpace of 100 by 200 sparse matrices over Finite Field of size 7
```

ncols()

Return the number of columns of matrices in this space.

EXAMPLES:

```
sage: M = Mat(ZZ['x'], 200000, 500000, sparse=True)
sage: M.ncols()
500000
```

ngens()

Return the number of generators of this matrix space, which is the number of entries in the matrices in this space.

EXAMPLES:

```
sage: M = Mat(GF(7), 100, 200); M.ngens()
20000
```

nrows()

Return the number of rows of matrices in this space.

EXAMPLES:

```
sage: M = Mat(ZZ, 200000, 500000)
sage: M.nrows()
200000
```

one()

Returns the identity matrix in self.

self must be a space of square matrices. The returned matrix is immutable. Please use `copy` if you want a modified copy.

EXAMPLES:

```
sage: MS1 = MatrixSpace(ZZ, 4)
sage: MS2 = MatrixSpace(QQ, 3, 4)
sage: I = MS1.identity_matrix()
sage: I
[1 0 0 0]
[0 1 0 0]
[0 0 1 0]
[0 0 0 1]
sage: Er = MS2.identity_matrix()
Traceback (most recent call last):
...
TypeError: self must be a space of square matrices
```

TESTS:

```
sage: MS1.one()[1,2] = 3
Traceback (most recent call last):
...
```

ValueError: matrix is immutable; please change a copy instead (i.e., use `copy(M)` to change a

random_element (*density=None*, *args, **kws)

Returns a random element from this matrix space.

INPUT:

- *density* - float or None (default: None); rough measure of the proportion of nonzero entries in the random matrix; if set to None, all entries of the matrix are randomized, allowing for any element of the underlying ring, but if set to a float, a proportion of entries is selected and randomized to non-zero elements of the ring
- *args, **kws - remaining parameters, which may be passed to the `random_element` function of the base ring. (“may be”, since this function calls the `randomize` function on the zero matrix, which need not call the `random_element` function of the base ring at all in general.)

OUTPUT:

- Matrix

Note: This method will randomize a proportion of roughly *density* entries in a newly allocated zero matrix.

By default, if the user sets the value of *density* explicitly, this method will enforce that these entries are set to non-zero values. However, if the test for equality with zero in the base ring is too expensive, the user can override this behaviour by passing the argument `nonzero=False` to this method.

Otherwise, if the user does not set the value of *density*, the default value is taken to be 1, and the option `nonzero=False` is passed to the `randomize` method.

EXAMPLES:

```
sage: Mat(ZZ,2,5).random_element()
[ -8   2   0   0   1]
[ -1   2   1 -95  -1]
sage: Mat(QQ,2,5).random_element(density=0.5)
[ 2  0  0  0  1]
[ 0  0  0 -1  0]
sage: Mat(QQ,3,sparse=True).random_element()
[ -1  -1  -1]
[ -3 -1/3 -1]
[  0  -1   1]
sage: Mat(GF(9,'a'),3,sparse=True).random_element()
[  a  2*a  1]
[  2    1 a + 2]
[ 2*a    2    2]
```

row_space ()

Return the module spanned by all rows of matrices in this matrix space. This is a free module of rank the number of rows. It will be sparse or dense as this matrix space is sparse or dense.

EXAMPLES:

```
sage: M = Mat(ZZ,20,5,sparse=False); M.row_space()
Ambient free module of rank 5 over the principal ideal domain Integer Ring
```

zero ()

Returns the zero matrix in `self`.

`self` must be a space of square matrices. The returned matrix is immutable. Please use `copy` if you want a modified copy.

EXAMPLES:

```
sage: z = MatrixSpace(GF(7), 2, 4).zero_matrix(); z
[0 0 0 0]
[0 0 0 0]
sage: z.is_mutable()
False
```

TESTS:

```
sage: MM = MatrixSpace(RDF, 1, 1, sparse=False); mat = MM.zero_matrix()
sage: copy(mat)
[0.0]
sage: MM = MatrixSpace(RDF, 0, 0, sparse=False); mat = MM.zero_matrix()
sage: copy(mat)
[]
sage: mat.is_mutable()
False
sage: MM.zero().is_mutable()
False
```

zero_matrix()

Returns the zero matrix in `self`.

`self` must be a space of square matrices. The returned matrix is immutable. Please use `copy` if you want a modified copy.

EXAMPLES:

```
sage: z = MatrixSpace(GF(7), 2, 4).zero_matrix(); z
[0 0 0 0]
[0 0 0 0]
sage: z.is_mutable()
False
```

TESTS:

```
sage: MM = MatrixSpace(RDF, 1, 1, sparse=False); mat = MM.zero_matrix()
sage: copy(mat)
[0.0]
sage: MM = MatrixSpace(RDF, 0, 0, sparse=False); mat = MM.zero_matrix()
sage: copy(mat)
[]
sage: mat.is_mutable()
False
sage: MM.zero().is_mutable()
False
```

`sage.matrix.matrix_space.dict_to_list` (*entries*, *nrows*, *ncols*)

Given a dictionary of coordinate tuples, return the list given by reading off the `nrows*ncols` matrix in row order.

EXAMPLES:

```
sage: from sage.matrix.matrix_space import dict_to_list
sage: d = {}
sage: d[(0,0)] = 1
sage: d[(1,1)] = 2
sage: dict_to_list(d, 2, 2)
[1, 0, 0, 2]
```



```
sage: dict_to_list(d, 2, 3)
[1, 0, 0, 0, 2, 0]
```

```
sage.matrix.matrix_space.is_MatrixSpace(x)
```

Returns True if self is an instance of MatrixSpace returns false if self is not an instance of MatrixSpace

EXAMPLES:

```
sage: from sage.matrix.matrix_space import is_MatrixSpace
sage: MS = MatrixSpace(QQ,2)
sage: A = MS.random_element()
sage: is_MatrixSpace(MS)
True
sage: is_MatrixSpace(A)
False
sage: is_MatrixSpace(5)
False
```

```
sage.matrix.matrix_space.list_to_dict(entries, nrows, ncols, rows=True)
```

Given a list of entries, create a dictionary whose keys are coordinate tuples and values are the entries.

EXAMPLES:

```
sage: from sage.matrix.matrix_space import list_to_dict
sage: d = list_to_dict([1,2,3,4],2,2)
sage: d[(0,1)]
2
sage: d = list_to_dict([1,2,3,4],2,2,rows=False)
sage: d[(0,1)]
3
```

```
sage.matrix.matrix_space.test_trivial_matrices_inverse(ring, sparse=True, check-
rank=True)
```

Tests inversion, determinant and is_invertible for trivial matrices.

This function is a helper to check that the inversion of trivial matrices (of size 0×0 , $n \times 0$, $0 \times n$ or 1×1) is handled consistently by the various implementation of matrices. The coherency is checked through a bunch of assertions. If an inconsistency is found, an AssertionError is raised which should make clear what is the problem.

INPUT:

- ring - a ring
- sparse - a boolean
- checkrank - a boolean

OUTPUT:

- nothing if everything is correct, otherwise raise an AssertionError

The methods determinant, is_invertible, rank and inverse are checked for

- the 0×0 empty identity matrix
- the 0×3 and 3×0 matrices
- the 1×1 null matrix [0]
- the 1×1 identity matrix [1]

If checkrank is False then the rank is not checked. This is used the check matrix over ring where echelon form is not implemented.

TODO: must be adapted to category check framework when ready (see trac #5274).

TESTS:

```
sage: from sage.matrix.matrix_space import test_trivial_matrices_inverse as tinv
sage: tinv(ZZ, sparse=True)
sage: tinv(ZZ, sparse=False)
sage: tinv(QQ, sparse=True)
sage: tinv(QQ, sparse=False)
sage: tinv(GF(11), sparse=True)
sage: tinv(GF(11), sparse=False)
sage: tinv(GF(2), sparse=True)
sage: tinv(GF(2), sparse=False)
sage: tinv(SR, sparse=True)
sage: tinv(SR, sparse=False)
sage: tinv(RDF, sparse=True)
sage: tinv(RDF, sparse=False)
sage: tinv(CDF, sparse=True)
sage: tinv(CDF, sparse=False)
sage: tinv(CyclotomicField(7), sparse=True)
sage: tinv(CyclotomicField(7), sparse=False)
sage: tinv(QQ['x,y'], sparse=True)
sage: tinv(QQ['x,y'], sparse=False)
```

MATRIX CONSTRUCTOR

```
class sage.matrix.constructor.MatrixFactory
```

```
    Bases: object
```

Create a matrix.

This implements the `matrix` constructor:

```
sage: matrix([[1, 2], [3, 4]])  
[1 2]  
[3 4]
```

It also contains methods to create special types of matrices, see `matrix.[tab]` for more options. For example:

```
sage: matrix.identity(2)  
[1 0]  
[0 1]
```

INPUT:

The `matrix` command takes the entries of a matrix, optionally preceded by a ring and the dimensions of the matrix, and returns a matrix.

The entries of a matrix can be specified as a flat list of elements, a list of lists (i.e., a list of rows), a list of Sage vectors, a callable object, or a dictionary having positions as keys and matrix entries as values (see the examples). If you pass in a callable object, then you must specify the number of rows and columns. You can create a matrix of zeros by passing an empty list or the integer zero for the entries. To construct a multiple of the identity (cI), you can specify square dimensions and pass in c . Calling `matrix()` with a Sage object may return something that makes sense. Calling `matrix()` with a NumPy array will convert the array to a matrix.

The ring, number of rows, and number of columns of the matrix can be specified by setting the `ring`, `nrows`, or `ncols` parameters or by passing them as the first arguments to the function in the order `ring`, `nrows`, `ncols`. The ring defaults to \mathbb{Z} if it is not specified or cannot be determined from the entries. If the numbers of rows and columns are not specified and cannot be determined, then an empty 0×0 matrix is returned.

- `ring` - the base ring for the entries of the matrix.
- `nrows` - the number of rows in the matrix.
- `ncols` - the number of columns in the matrix.
- `sparse` - create a sparse matrix. This defaults to `True` when the entries are given as a dictionary, otherwise defaults to `False`.

OUTPUT:

a matrix

EXAMPLES:

```
sage: m=matrix(2); m; m.parent()
```

```
[0 0]
```

```
[0 0]
```

Full MatrixSpace of 2 by 2 dense matrices over Integer Ring

```
sage: m=matrix(2,3); m; m.parent()
```

```
[0 0 0]
```

```
[0 0 0]
```

Full MatrixSpace of 2 by 3 dense matrices over Integer Ring

```
sage: m=matrix(QQ,[[1,2,3],[4,5,6]]); m; m.parent()
```

```
[1 2 3]
```

```
[4 5 6]
```

Full MatrixSpace of 2 by 3 dense matrices over Rational Field

```
sage: m = matrix(QQ, 3, 3, lambda i, j: i+j); m
```

```
[0 1 2]
```

```
[1 2 3]
```

```
[2 3 4]
```

```
sage: m = matrix(3, lambda i, j: i-j); m
```

```
[ 0 -1 -2]
```

```
[ 1  0 -1]
```

```
[ 2  1  0]
```

```
sage: matrix(QQ,2,3,lambda x, y: x+y)
```

```
[0 1 2]
```

```
[1 2 3]
```

```
sage: matrix(QQ,3,2,lambda x, y: x+y)
```

```
[0 1]
```

```
[1 2]
```

```
[2 3]
```

```
sage: v1=vector((1,2,3))
```

```
sage: v2=vector((4,5,6))
```

```
sage: m=matrix([v1,v2]); m; m.parent()
```

```
[1 2 3]
```

```
[4 5 6]
```

Full MatrixSpace of 2 by 3 dense matrices over Integer Ring

```
sage: m=matrix(QQ,2,[1,2,3,4,5,6]); m; m.parent()
```

```
[1 2 3]
```

```
[4 5 6]
```

Full MatrixSpace of 2 by 3 dense matrices over Rational Field

```
sage: m=matrix(QQ,2,3,[1,2,3,4,5,6]); m; m.parent()
```

```
[1 2 3]
```

```
[4 5 6]
```

Full MatrixSpace of 2 by 3 dense matrices over Rational Field

```
sage: m=matrix(({(0,1): 2, (1,1):2/5})); m; m.parent()
```

```
[ 0  2]
```

```
[ 0 2/5]
```

Full MatrixSpace of 2 by 2 sparse matrices over Rational Field

```
sage: m=matrix(QQ,2,3,{(1,1): 2}); m; m.parent()
```

```
[0 0 0]
```

```
[0 2 0]
```

Full MatrixSpace of 2 by 3 sparse matrices over Rational Field

```
sage: import numpy
sage: n=numpy.array([[1,2],[3,4]],float)
sage: m=matrix(n); m; m.parent()
[1.0 2.0]
[3.0 4.0]
```

Full MatrixSpace of 2 by 2 dense matrices over Real Double Field

```
sage: v = vector(ZZ, [1, 10, 100])
sage: m=matrix(v); m; m.parent()
[ 1 10 100]
```

Full MatrixSpace of 1 by 3 dense matrices over Integer Ring

```
sage: m=matrix(GF(7), v); m; m.parent()
[1 3 2]
```

Full MatrixSpace of 1 by 3 dense matrices over Finite Field of size 7

```
sage: g = graphs.PetersenGraph()
sage: m = matrix(g); m; m.parent()
```

```
[0 1 0 0 1 1 0 0 0 0]
[1 0 1 0 0 0 1 0 0 0]
[0 1 0 1 0 0 0 1 0 0]
[0 0 1 0 1 0 0 0 1 0]
[1 0 0 1 0 0 0 0 0 1]
[1 0 0 0 0 0 0 1 1 0]
[0 1 0 0 0 0 0 0 1 1]
[0 0 1 0 0 1 0 0 0 1]
[0 0 0 1 0 1 1 0 0 0]
[0 0 0 0 1 0 1 1 0 0]
```

Full MatrixSpace of 10 by 10 dense matrices over Integer Ring

```
sage: matrix(ZZ, 10, 10, range(100), sparse=True).parent()
Full MatrixSpace of 10 by 10 sparse matrices over Integer Ring
```

```
sage: R = PolynomialRing(QQ, 9, 'x')
sage: A = matrix(R, 3, 3, R.gens()); A
[x0 x1 x2]
[x3 x4 x5]
[x6 x7 x8]
sage: det(A)
-x2*x4*x6 + x1*x5*x6 + x2*x3*x7 - x0*x5*x7 - x1*x3*x8 + x0*x4*x8
```

TESTS:

```
sage: m=matrix(); m; m.parent()
[]
Full MatrixSpace of 0 by 0 dense matrices over Integer Ring
sage: m=matrix(QQ); m; m.parent()
[]
Full MatrixSpace of 0 by 0 dense matrices over Rational Field
sage: m=matrix(QQ,2); m; m.parent()
[0 0]
[0 0]
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: m=matrix(QQ,2,3); m; m.parent()
[0 0 0]
[0 0 0]
Full MatrixSpace of 2 by 3 dense matrices over Rational Field
sage: m=matrix([]); m; m.parent()
```

```
[ ]
Full MatrixSpace of 0 by 0 dense matrices over Integer Ring
sage: m=matrix(QQ,[]); m; m.parent()
[ ]
Full MatrixSpace of 0 by 0 dense matrices over Rational Field
sage: m=matrix(2,2,1); m; m.parent()
[1 0]
[0 1]
Full MatrixSpace of 2 by 2 dense matrices over Integer Ring
sage: m=matrix(QQ,2,2,1); m; m.parent()
[1 0]
[0 1]
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: m=matrix(2,3,0); m; m.parent()
[0 0 0]
[0 0 0]
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
sage: m=matrix(QQ,2,3,0); m; m.parent()
[0 0 0]
[0 0 0]
Full MatrixSpace of 2 by 3 dense matrices over Rational Field
sage: m=matrix([[1,2,3],[4,5,6]]); m; m.parent()
[1 2 3]
[4 5 6]
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
sage: m=matrix(QQ,2,[[1,2,3],[4,5,6]]); m; m.parent()
[1 2 3]
[4 5 6]
Full MatrixSpace of 2 by 3 dense matrices over Rational Field
sage: m=matrix(QQ,3,[[1,2,3],[4,5,6]]); m; m.parent()
Traceback (most recent call last):
...
ValueError: Number of rows does not match up with specified number.
sage: m=matrix(QQ,2,3,[[1,2,3],[4,5,6]]); m; m.parent()
[1 2 3]
[4 5 6]
Full MatrixSpace of 2 by 3 dense matrices over Rational Field
sage: m=matrix(QQ,2,4,[[1,2,3],[4,5,6]]); m; m.parent()
Traceback (most recent call last):
...
ValueError: Number of columns does not match up with specified number.
sage: m=matrix([(1,2,3),(4,5,6)]); m; m.parent()
[1 2 3]
[4 5 6]
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
sage: m=matrix([1,2,3,4,5,6]); m; m.parent()
[1 2 3 4 5 6]
Full MatrixSpace of 1 by 6 dense matrices over Integer Ring
sage: m=matrix((1,2,3,4,5,6)); m; m.parent()
[1 2 3 4 5 6]
Full MatrixSpace of 1 by 6 dense matrices over Integer Ring
sage: m=matrix(QQ,[1,2,3,4,5,6]); m; m.parent()
[1 2 3 4 5 6]
Full MatrixSpace of 1 by 6 dense matrices over Rational Field
sage: m=matrix(QQ,3,2,[1,2,3,4,5,6]); m; m.parent()
[1 2]
[3 4]
[5 6]
```

```

Full MatrixSpace of 3 by 2 dense matrices over Rational Field
sage: m=matrix(QQ,2,4,[1,2,3,4,5,6]); m; m.parent()
Traceback (most recent call last):
...
ValueError: entries has the wrong length
sage: m=matrix(QQ,5,[1,2,3,4,5,6]); m; m.parent()
Traceback (most recent call last):
...
TypeError: cannot construct an element of
Full MatrixSpace of 5 by 1 dense matrices over Rational Field
from [1, 2, 3, 4, 5, 6]!
sage: m=matrix(({(1,1): 2})); m; m.parent()
[0 0]
[0 2]
Full MatrixSpace of 2 by 2 sparse matrices over Integer Ring
sage: m=matrix(QQ,{(1,1): 2}); m; m.parent()
[0 0]
[0 2]
Full MatrixSpace of 2 by 2 sparse matrices over Rational Field
sage: m=matrix(QQ,3,{(1,1): 2}); m; m.parent()
[0 0 0]
[0 2 0]
[0 0 0]
Full MatrixSpace of 3 by 3 sparse matrices over Rational Field
sage: m=matrix(QQ,3,4,{(1,1): 2}); m; m.parent()
[0 0 0 0]
[0 2 0 0]
[0 0 0 0]
Full MatrixSpace of 3 by 4 sparse matrices over Rational Field
sage: m=matrix(QQ,2,{(1,1): 2}); m; m.parent()
[0 0]
[0 2]
Full MatrixSpace of 2 by 2 sparse matrices over Rational Field
sage: m=matrix(QQ,1,{(1,1): 2}); m; m.parent()
Traceback (most recent call last):
...
IndexError: invalid entries list
sage: m=matrix({}); m; m.parent()
[]
Full MatrixSpace of 0 by 0 sparse matrices over Integer Ring
sage: m=matrix(QQ,{}); m; m.parent()
[]
Full MatrixSpace of 0 by 0 sparse matrices over Rational Field
sage: m=matrix(QQ,2,{}); m; m.parent()
[0 0]
[0 0]
Full MatrixSpace of 2 by 2 sparse matrices over Rational Field
sage: m=matrix(QQ,2,3,{}); m; m.parent()
[0 0 0]
[0 0 0]
Full MatrixSpace of 2 by 3 sparse matrices over Rational Field
sage: m=matrix(2,{}); m; m.parent()
[0 0]
[0 0]
Full MatrixSpace of 2 by 2 sparse matrices over Integer Ring
sage: m=matrix(2,3,{}); m; m.parent()
[0 0 0]
[0 0 0]

```

```
Full MatrixSpace of 2 by 3 sparse matrices over Integer Ring
sage: m=matrix(0); m; m.parent()
[]
Full MatrixSpace of 0 by 0 dense matrices over Integer Ring
sage: m=matrix(0,2); m; m.parent()
[]
Full MatrixSpace of 0 by 2 dense matrices over Integer Ring
sage: m=matrix(2,0); m; m.parent()
[]
Full MatrixSpace of 2 by 0 dense matrices over Integer Ring
sage: m=matrix(0,[1]); m; m.parent()
Traceback (most recent call last):
...
ValueError: entries has the wrong length
sage: m=matrix(1,0,[]); m; m.parent()
[]
Full MatrixSpace of 1 by 0 dense matrices over Integer Ring
sage: m=matrix(0,1,[]); m; m.parent()
[]
Full MatrixSpace of 0 by 1 dense matrices over Integer Ring
sage: m=matrix(0,[]); m; m.parent()
[]
Full MatrixSpace of 0 by 0 dense matrices over Integer Ring
sage: m=matrix(0,{}); m; m.parent()
[]
Full MatrixSpace of 0 by 0 sparse matrices over Integer Ring
sage: m=matrix(0,{(1,1):2}); m; m.parent()
Traceback (most recent call last):
...
IndexError: invalid entries list
sage: m=matrix(2,0,{(1,1):2}); m; m.parent()
Traceback (most recent call last):
...
IndexError: invalid entries list
sage: import numpy
sage: n=numpy.array([[numpy.complex(0,1),numpy.complex(0,2)],[3,4]],complex)
sage: m=matrix(n); m; m.parent()
[1.0*I 2.0*I]
[ 3.0  4.0]
Full MatrixSpace of 2 by 2 dense matrices over Complex Double Field
sage: n=numpy.array([[1,2],[3,4]], 'int32')
sage: m=matrix(n); m; m.parent()
[1 2]
[3 4]
Full MatrixSpace of 2 by 2 dense matrices over Integer Ring
sage: n = numpy.array([[1,2,3],[4,5,6],[7,8,9]], 'float32')
sage: m=matrix(n); m; m.parent()
[1.0 2.0 3.0]
[4.0 5.0 6.0]
[7.0 8.0 9.0]
Full MatrixSpace of 3 by 3 dense matrices over Real Double Field
sage: n=numpy.array([[1,2,3],[4,5,6],[7,8,9]], 'float64')
sage: m=matrix(n); m; m.parent()
[1.0 2.0 3.0]
[4.0 5.0 6.0]
[7.0 8.0 9.0]
Full MatrixSpace of 3 by 3 dense matrices over Real Double Field
sage: n=numpy.array([[1,2,3],[4,5,6],[7,8,9]], 'complex64')
```



```

sage: m=matrix(n); m; m.parent()
[1.0 2.0 3.0]
[4.0 5.0 6.0]
[7.0 8.0 9.0]
Full MatrixSpace of 3 by 3 dense matrices over Complex Double Field
sage: n=numpy.array([[1,2,3],[4,5,6],[7,8,9]], 'complex128')
sage: m=matrix(n); m; m.parent()
[1.0 2.0 3.0]
[4.0 5.0 6.0]
[7.0 8.0 9.0]
Full MatrixSpace of 3 by 3 dense matrices over Complex Double Field
sage: a = matrix([[1,2],[3,4]])
sage: b = matrix(a.numpy()); b
[1 2]
[3 4]
sage: a == b
True
sage: c = matrix(a.numpy('float32')); c
[1.0 2.0]
[3.0 4.0]
sage: matrix(numpy.array([[5]]))
[5]
sage: v = vector(ZZ, [1, 10, 100])
sage: m=matrix(ZZ['x'], v); m; m.parent()
[ 1 10 100]
Full MatrixSpace of 1 by 3 dense matrices over Univariate Polynomial Ring in x over Integer Ring
sage: matrix(ZZ, 10, 10, range(100)).parent()
Full MatrixSpace of 10 by 10 dense matrices over Integer Ring
sage: m = matrix(GF(7), [[1/3,2/3,1/2], [3/4,4/5,7]]); m; m.parent()
[5 3 4]
[6 5 0]
Full MatrixSpace of 2 by 3 dense matrices over Finite Field of size 7
sage: m = matrix([[1,2,3], [RDF(2), CDF(1,2), 3]]); m; m.parent()
[      1.0      2.0      3.0]
[      2.0 1.0 + 2.0*I      3.0]
Full MatrixSpace of 2 by 3 dense matrices over Complex Double Field
sage: m=matrix(3,3,1/2); m; m.parent()
[1/2  0  0]
[ 0 1/2  0]
[ 0  0 1/2]
Full MatrixSpace of 3 by 3 dense matrices over Rational Field
sage: matrix([[1],[2,3]])
Traceback (most recent call last):
...
ValueError: List of rows is not valid (rows are wrong types or lengths)
sage: matrix([[1],2])
Traceback (most recent call last):
...
ValueError: List of rows is not valid (rows are wrong types or lengths)
sage: matrix(vector(RR, [1,2,3])).parent()
Full MatrixSpace of 1 by 3 dense matrices over Real Field with 53 bits of precision
sage: matrix(ZZ, [[0] for i in range(10^5)]).is_zero() # see #10158
True

```

AUTHORS:

- ?: Initial implementation
- Jason Grout (2008-03): almost a complete rewrite, with bits and pieces from the original implementation

`sage.matrix.constructor.block_diagonal_matrix(*sub_matrices, **kws)`

This function is available as `block_diagonal_matrix(...)` and `matrix.block_diagonal(...)`.

Create a block matrix whose diagonal block entries are given by `sub_matrices`, with zero elsewhere.

See also `block_matrix()`.

EXAMPLES:

```
sage: A = matrix(ZZ, 2, [1,2,3,4])
```

```
sage: block_diagonal_matrix(A, A)
```

```
[1 2|0 0]
[3 4|0 0]
[---+---]
[0 0|1 2]
[0 0|3 4]
```

The sub-matrices need not be square:

```
sage: B = matrix(QQ, 2, 3, range(6))
```

```
sage: block_diagonal_matrix(~A, B)
```

```
[ -2      1|  0      0      0]
[ 3/2 -1/2|  0      0      0]
[-----+-----]
[  0      0|  0      1      2]
[  0      0|  3      4      5]
```

`sage.matrix.constructor.block_matrix(*args, **kws)`

This function is available as `block_matrix(...)` and `matrix.block(...)`.

Returns a larger matrix made by concatenating submatrices (rows first, then columns). For example, the matrix

```
[ A B ]
[ C D ]
```

is made up of submatrices A, B, C, and D.

INPUT:

The `block_matrix` command takes a list of submatrices to add as blocks, optionally preceded by a ring and the number of block rows and block columns, and returns a matrix.

The submatrices can be specified as a list of matrices (using `nrows` and `ncols` to determine their layout), or a list of lists of matrices, where each list forms a row.

- `ring` - the base ring
- `nrows` - the number of block rows
- `ncols` - the number of block cols
- `sub_matrices` - matrices (see below for syntax)
- `subdivide` - boolean, whether or not to add subdivision information to the matrix
- `sparse` - boolean, whether to make the resulting matrix sparse

EXAMPLES:

```
sage: A = matrix(QQ, 2, 2, [3,9,6,10])
```

```
sage: block_matrix([ [A, -A], [~A, 100*A] ])
```

```
[  3      9| -3     -9]
[  6     10| -6    -10]
[-----+-----]
```

```
[-5/12  3/8|  300  900]
[  1/4 -1/8|  600 1000]
```

If the number of submatrices in each row is the same, you can specify the submatrices as a single list too:

```
sage: block_matrix(2, 2, [ A, A, A, A ])
[ 3  9| 3  9]
[ 6 10| 6 10]
[-----+-----]
[ 3  9| 3  9]
[ 6 10| 6 10]
```

One can use constant entries:

```
sage: block_matrix([ [1, A], [0, 1] ])
[ 1  0| 3  9]
[ 0  1| 6 10]
[-----+-----]
[ 0  0| 1  0]
[ 0  0| 0  1]
```

A zero entry may represent any square or non-square zero matrix:

```
sage: B = matrix(QQ, 1, 1, [ 1 ] )
sage: C = matrix(QQ, 2, 2, [ 2, 3, 4, 5 ] )
sage: block_matrix([ [B, 0], [0, C] ])
[1|0 0]
[-+----]
[0|2 3]
[0|4 5]
```

One can specify the number of rows or columns as keywords too:

```
sage: block_matrix([A, -A, ~A, 100*A], ncols=4)
[   3   9|   -3   -9|-5/12  3/8|  300  900]
[   6  10|   -6  -10|  1/4  -1/8|  600 1000]

sage: block_matrix([A, -A, ~A, 100*A], nrows=1)
[   3   9|   -3   -9|-5/12  3/8|  300  900]
[   6  10|   -6  -10|  1/4  -1/8|  600 1000]
```

It handles base rings nicely too:

```
sage: R.<x> = ZZ['x']
sage: block_matrix(2, 2, [1/2, A, 0, x-1])
[ 1/2  0|  3  9]
[  0  1/2|  6 10]
[-----+-----]
[  0  0|x - 1  0]
[  0  0|  0 x - 1]
```

```
sage: block_matrix(2, 2, [1/2, A, 0, x-1]).parent()
```

Full MatrixSpace of 4 by 4 dense matrices over Univariate Polynomial Ring in x over Rational Field

Subdivisions are optional. If they are disabled, the columns need not line up:

```
sage: B = matrix(QQ, 2, 3, range(6))
sage: block_matrix([ [~A, B], [B, ~A] ], subdivide=False)
[-5/12  3/8  0  1  2]
[  1/4 -1/8  3  4  5]
```

```
[ 0 1 2 -5/12 3/8]
[ 3 4 5 1/4 -1/8]
```

Without subdivisions it also deduces dimensions for scalars if possible:

```
sage: C = matrix(ZZ, 1, 2, range(2))
sage: block_matrix([ [ C, 0 ], [ 3, 4 ], [ 5, 6, C ] ], subdivide=False )
[0 1 0 0]
[3 0 4 0]
[0 3 0 4]
[5 6 0 1]
```

If all submatrices are sparse (unless there are none at all), the result will be a sparse matrix. Otherwise it will be dense by default. The `sparse` keyword can be used to override this:

```
sage: A = Matrix(ZZ, 2, 2, [0, 1, 0, 0], sparse=True)
sage: block_matrix([ [ A ], [ A ] ]).parent()
Full MatrixSpace of 4 by 2 sparse matrices over Integer Ring
sage: block_matrix([ [ A ], [ A ] ], sparse=False).parent()
Full MatrixSpace of 4 by 2 dense matrices over Integer Ring
```

Consecutive zero submatrices are consolidated.

```
sage: B = matrix(2, range(4))
sage: C = matrix(2, 8, range(16))
sage: block_matrix(2, [[B,0,0,B],[C]], subdivide=False)
[ 0 1 0 0 0 0 0 0 1]
[ 2 3 0 0 0 0 2 3]
[ 0 1 2 3 4 5 6 7]
[ 8 9 10 11 12 13 14 15]
```

Ambiguity is not tolerated.

```
sage: B = matrix(2, range(4))
sage: C = matrix(2, 8, range(16))
sage: block_matrix(2, [[B,0,B,0],[C]], subdivide=False)
Traceback (most recent call last):
...
ValueError: insufficient information to determine submatrix widths
```

Historically, giving only a flat list of submatrices, whose number was a perfect square, would create a new matrix by laying out the submatrices in a square grid. This behavior is now deprecated.

```
sage: A = matrix(2, 3, range(6))
sage: B = matrix(3, 3, range(9))
sage: block_matrix([A, A, B, B])
doctest:...: DeprecationWarning: invocation of block_matrix with just a list whose length is a p
[0 1 2|0 1 2]
[3 4 5|3 4 5]
[-----+-----]
[0 1 2|0 1 2]
[3 4 5|3 4 5]
[6 7 8|6 7 8]
```

Historically, a flat list of matrices whose number is not a perfect square, with no specification of the number of rows or columns, would raise an error. This behavior continues, but could be removed when the deprecation above is completed.

```
sage: A = matrix(2, 3, range(6))
sage: B = matrix(3, 3, range(9))
```

```
sage: block_matrix([A, A, A, B, B, B])
Traceback (most recent call last):
...
ValueError: must specify nrows or ncols for non-square block matrix.
```

`sage.matrix.constructor.column_matrix(*args, **kws)`
 This function is available as `column_matrix(...)` and `matrix.column(...)`.

Constructs a matrix, and then swaps rows for columns and columns for rows.

Note: Linear algebra in Sage favors rows over columns. So, generally, when creating a matrix, input vectors and lists are treated as rows. This function is a convenience that turns around this convention when creating a matrix. If you are not familiar with the usual `matrix` constructor, you might want to consider it first.

INPUT:

Inputs are almost exactly the same as for the `matrix` constructor, which are documented there. But see examples below for how dimensions are handled.

OUTPUT:

Output is exactly the transpose of what the `matrix` constructor would return. In other words, the `matrix` constructor builds a matrix and then this function exchanges rows for columns, and columns for rows.

EXAMPLES:

The most compelling use of this function is when you have a collection of lists or vectors that you would like to become the columns of a matrix. In almost any other situation, the `matrix` constructor can probably do the job just as easily, or easier.

```
sage: col_1 = [1, 2, 3]
sage: col_2 = [4, 5, 6]
sage: column_matrix([col_1, col_2])
[1 4]
[2 5]
[3 6]
```

```
sage: v1 = vector(QQ, [10, 20])
sage: v2 = vector(QQ, [30, 40])
sage: column_matrix(QQ, [v1, v2])
[10 30]
[20 40]
```

If you only specify one dimension along with a flat list of entries, then it will be the number of columns in the result (which is different from the behavior of the `matrix` constructor).

```
sage: column_matrix(ZZ, 8, range(24))
[ 0  3  6  9 12 15 18 21]
[ 1  4  7 10 13 16 19 22]
[ 2  5  8 11 14 17 20 23]
```

And when you specify two dimensions, then they should be number of columns first, then the number of rows, which is the reverse of how they would be specified for the `matrix` constructor.

```
sage: column_matrix(QQ, 5, 3, range(15))
[ 0  3  6  9 12]
[ 1  4  7 10 13]
[ 2  5  8 11 14]
```

And a few unproductive, but illustrative, examples.

```
sage: A = matrix(ZZ, 3, 4, range(12))
sage: B = column_matrix(ZZ, 3, 4, range(12))
sage: A == B.transpose()
True

sage: A = matrix(QQ, 7, 12, range(84))
sage: A == column_matrix(A.columns())
True

sage: A=column_matrix(QQ, matrix(ZZ, 3, 2, range(6)) )
sage: A
[0 2 4]
[1 3 5]
sage: A.parent()
Full MatrixSpace of 2 by 3 dense matrices over Rational Field
```

`sage.matrix.constructor.companion_matrix(poly, format='right')`

This function is available as `companion_matrix(...)` and `matrix.companion(...)`.

Create a companion matrix from a monic polynomial.

INPUT:

- `poly` - a univariate polynomial, or an iterable containing the coefficients of a polynomial, with low-degree coefficients first. The polynomial (or the polynomial implied by the coefficients) must be monic. In other words, the leading coefficient must be one. A symbolic expression that might also be a polynomial is not proper input, see examples below.
- `format` - default: 'right' - specifies one of four variations of a companion matrix. Allowable values are 'right', 'left', 'top' and 'bottom', which indicates which border of the matrix contains the negatives of the coefficients.

OUTPUT:

A square matrix with a size equal to the degree of the polynomial. The returned matrix has ones above, or below the diagonal, and the negatives of the coefficients along the indicated border of the matrix (excepting the leading one coefficient). See the first examples below for precise illustrations.

EXAMPLES:

Each of the four possibilities. Notice that the coefficients are specified and their negatives become the entries of the matrix. The leading one must be given, but is not used. The permutation matrix P is the identity matrix, with the columns reversed. The last three statements test the general relationships between the four variants.

```
sage: poly = [-2, -3, -4, -5, -6, 1]
sage: R = companion_matrix(poly, format='right'); R
[0 0 0 0 2]
[1 0 0 0 3]
[0 1 0 0 4]
[0 0 1 0 5]
[0 0 0 1 6]
sage: L = companion_matrix(poly, format='left'); L
[6 1 0 0 0]
[5 0 1 0 0]
[4 0 0 1 0]
[3 0 0 0 1]
[2 0 0 0 0]
sage: B = companion_matrix(poly, format='bottom'); B
[0 1 0 0 0]
[0 0 1 0 0]
```

```

[0 0 0 1 0]
[0 0 0 0 1]
[2 3 4 5 6]
sage: T = companion_matrix(poly, format='top'); T
[6 5 4 3 2]
[1 0 0 0 0]
[0 1 0 0 0]
[0 0 1 0 0]
[0 0 0 1 0]

```

```

sage: perm = Permutation([5, 4, 3, 2, 1])
sage: P = perm.to_matrix()
sage: L == P*R*P
True
sage: B == R.transpose()
True
sage: T == P*R.transpose()*P
True

```

A polynomial may be used as input, however a symbolic expression, even if it looks like a polynomial, is not regarded as such when used as input to this routine. Obtaining the list of coefficients from a symbolic polynomial is one route to the companion matrix.

```

sage: x = polygen(QQ, 'x')
sage: p = x^3 - 4*x^2 + 8*x - 12
sage: companion_matrix(p)
[ 0  0 12]
[ 1  0 -8]
[ 0  1  4]

```

```

sage: y = var('y')
sage: q = y^3 - 2*y + 1
sage: companion_matrix(q)
Traceback (most recent call last):
...

```

TypeError: input must be a polynomial (not a symbolic expression, see docstring), or other iterable

```

sage: coeff_list = [q(y=0)] + [q.coeff(y^k) for k in range(1, q.degree(y)+1)]
sage: coeff_list
[1, -2, 0, 1]
sage: companion_matrix(coeff_list)
[ 0  0 -1]
[ 1  0  2]
[ 0  1  0]

```

The minimal polynomial of a companion matrix is equal to the polynomial used to create it. Used in a block diagonal construction, they can be used to create matrices with any desired minimal polynomial, or characteristic polynomial.

```

sage: t = polygen(QQ, 't')
sage: p = t^12 - 7*t^4 + 28*t^2 - 456
sage: C = companion_matrix(p, format='top')
sage: q = C.minpoly(var=t); q
t^12 - 7*t^4 + 28*t^2 - 456
sage: p == q
True

sage: p = t^3 + 3*t - 8
sage: q = t^5 + t - 17

```

```
sage: A = block_diagonal_matrix( companion_matrix(p),
...                             companion_matrix(p^2),
...                             companion_matrix(q),
...                             companion_matrix(q) )
sage: A.charpoly(var=t).factor()
(t^3 + 3*t - 8)^3 * (t^5 + t - 17)^2
sage: A.minpoly(var=t).factor()
(t^3 + 3*t - 8)^2 * (t^5 + t - 17)
```

TESTS:

```
sage: companion_matrix([4, 5, 1], format='junk')
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: format must be 'right', 'left', 'top' or 'bottom', not junk
```

```
sage: companion_matrix(sin(x))
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: input must be a polynomial (not a symbolic expression, see docstring), or other iterable
```

```
sage: companion_matrix([2, 3, 896])
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: polynomial (or the polynomial implied by coefficients) must be monic, not a leading
```

```
sage: F.<a> = GF(2^2)
```

```
sage: companion_matrix([4/3, a+1, 1])
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: unable to find common ring for coefficients from polynomial
```

```
sage: A = companion_matrix([1])
```

```
sage: A.nrows(); A.ncols()
```

```
0
```

```
0
```

```
sage: A = companion_matrix([])
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: polynomial cannot be specified by an empty list
```

AUTHOR:

•Rob Beezer (2011-05-19)

sage.matrix.constructor.**diagonal_matrix**(arg0=None, arg1=None, arg2=None,
sparse=True)

This function is available as `diagonal_matrix(...)` and `matrix.diagonal(...)`.

Return a square matrix with specified diagonal entries, and zeros elsewhere.

FORMATS:

- 1.`diagonal_matrix(entries)`
- 2.`diagonal_matrix(nrows, entries)`
- 3.`diagonal_matrix(ring, entries)`
- 4.`diagonal_matrix(ring, nrows, entries)`

INPUT:

- `entries` - the values to place along the diagonal of the returned matrix. This may be a flat list, a flat tuple, a vector or free module element, or a one-dimensional NumPy array.
- `nrows` - the size of the returned matrix, which will have an equal number of columns
- `ring` - the ring containing the entries of the diagonal entries. This may not be specified in combination with a NumPy array.
- `sparse` - default: `True` - whether or not the result has a sparse implementation.

OUTPUT:

A square matrix over the given `ring` with a size given by `nrows`. If the ring is not given it is inferred from the given entries. The values on the diagonal of the returned matrix come from `entries`. If the number of entries is not enough to fill the whole diagonal, it is padded with zeros.

EXAMPLES:

We first demonstrate each of the input formats with various different ways to specify the entries.

Format 1: a flat list of entries.

```
sage: A = diagonal_matrix([2, 1.3, 5]); A
[ 2.000000000000000 0.000000000000000 0.000000000000000]
[0.000000000000000 1.300000000000000 0.000000000000000]
[0.000000000000000 0.000000000000000 5.000000000000000]
sage: A.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Real Field with 53 bits of precision
```

Format 2: size specified, a tuple with initial entries. Note that a short list of entries is effectively padded with zeros.

```
sage: A = diagonal_matrix(3, (4, 5)); A
[4 0 0]
[0 5 0]
[0 0 0]
sage: A.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Integer Ring
```

Format 3: ring specified, a vector of entries.

```
sage: A = diagonal_matrix(QQ, vector(ZZ, [1,2,3])); A
[1 0 0]
[0 2 0]
[0 0 3]
sage: A.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Rational Field
```

Format 4: ring, size and list of entries.

```
sage: A = diagonal_matrix(FiniteField(3), 3, [2, 16]); A
[2 0 0]
[0 1 0]
[0 0 0]
sage: A.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Finite Field of size 3
```

NumPy arrays may be used as input.

```
sage: import numpy
sage: entries = numpy.array([1.2, 5.6]); entries
```

```
array([ 1.2,  5.6])
sage: A = diagonal_matrix(3, entries); A
[1.2 0.0 0.0]
[0.0 5.6 0.0]
[0.0 0.0 0.0]
sage: A.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Real Double Field

sage: j = numpy.complex(0,1)
sage: entries = numpy.array([2.0+j, 8.1, 3.4+2.6*j]); entries
array([ 2.0+1.j ,  8.1+0.j ,  3.4+2.6j])
sage: A = diagonal_matrix(entries); A
[2.0 + 1.0*I      0.0      0.0]
[      0.0      8.1      0.0]
[      0.0      0.0 3.4 + 2.6*I]
sage: A.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Complex Double Field

sage: entries = numpy.array([4, 5, 6])
sage: A = diagonal_matrix(entries); A
[4 0 0]
[0 5 0]
[0 0 6]
sage: A.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Integer Ring

sage: entries = numpy.array([4.1, 5.2, 6.3])
sage: A = diagonal_matrix(ZZ, entries); A
Traceback (most recent call last):
...
TypeError: Cannot convert non-integral float to integer
```

By default returned matrices have a sparse implementation. This can be changed when using any of the formats.

```
sage: A = diagonal_matrix([1,2,3], sparse=False)
sage: A.parent()
Full MatrixSpace of 3 by 3 dense matrices over Integer Ring
```

An empty list and no ring specified defaults to the integers.

```
sage: A = diagonal_matrix([])
sage: A.parent()
Full MatrixSpace of 0 by 0 sparse matrices over Integer Ring
```

Giving the entries improperly may first complain about not having a length.

```
sage: diagonal_matrix(QQ, 5, 10)
Traceback (most recent call last):
...
TypeError: unable to determine number of entries for diagonal matrix construction
```

Giving too many entries will raise an error.

```
sage: diagonal_matrix(QQ, 3, [1,2,3,4])
Traceback (most recent call last):
...
ValueError: number of diagonal matrix entries (4) exceeds the requested matrix size (3)
```

A negative size sometimes causes the error that there are too many elements.

```
sage: diagonal_matrix(-2, [2])
Traceback (most recent call last):
...
ValueError: number of diagonal matrix entries (1) exceeds the requested matrix size (-2)
```

Types for the entries are limited, even though they may have a length.

```
sage: diagonal_matrix(x^2)
Traceback (most recent call last):
...
TypeError: diagonal matrix entries are not a supported type (list, tuple, vector, or NumPy array)
```

AUTHOR:

- Rob Beezer (2011-01-11): total rewrite

`sage.matrix.constructor.elementary_matrix` (*arg0*, *arg1*=None, ***kws*)

This function is available as `elementary_matrix(...)` and `matrix.elementary(...)`.

Creates a square matrix that corresponds to a row operation or a column operation.

FORMATS:

In each case, *R* is the base ring, and is optional. *n* is the size of the square matrix created. Any call may include the `sparse` keyword to determine the representation used. The default is `False` which leads to a dense representation. We describe the matrices by their associated row operation, see the output description for more.

- `elementary_matrix(R, n, row1=i, row2=j)`

The matrix which swaps rows *i* and *j*.

- `elementary_matrix(R, n, row1=i, scale=s)`

The matrix which multiplies row *i* by *s*.

- `elementary_matrix(R, n, row1=i, row2=j, scale=s)`

The matrix which multiplies row *j* by *s* and adds it to row *i*.

Elementary matrices representing column operations are created in an entirely analogous way, replacing `row1` by `col1` and replacing `row2` by `col2`.

Specifying the ring for entries of the matrix is optional. If it is not given, and a scale parameter is provided, then a ring containing the value of `scale` will be used. Otherwise, the ring defaults to the integers.

OUTPUT:

An elementary matrix is a square matrix that is very close to being an identity matrix. If *E* is an elementary matrix and *A* is any matrix with the same number of rows, then *E***A* is the result of applying a row operation to *A*. This is how the three types created by this function are described. Similarly, an elementary matrix can be associated with a column operation, so if *E* has the same number of columns as *A* then *A***E* is the result of performing a column operation on *A*.

An elementary matrix representing a row operation is created if `row1` is specified, while an elementary matrix representing a column operation is created if `col1` is specified.

EXAMPLES:

Over the integers, creating row operations. Recall that row and column numbering begins at zero.

```
sage: A = matrix(ZZ, 4, 10, range(40)); A
[ 0  1  2  3  4  5  6  7  8  9]
[10 11 12 13 14 15 16 17 18 19]
```

```
[20 21 22 23 24 25 26 27 28 29]
[30 31 32 33 34 35 36 37 38 39]

sage: E = elementary_matrix(4, row1=1, row2=3); E
[1 0 0 0]
[0 0 0 1]
[0 0 1 0]
[0 1 0 0]
sage: E*A
[ 0  1  2  3  4  5  6  7  8  9]
[30 31 32 33 34 35 36 37 38 39]
[20 21 22 23 24 25 26 27 28 29]
[10 11 12 13 14 15 16 17 18 19]

sage: E = elementary_matrix(4, row1=2, scale=10); E
[ 1  0  0  0]
[ 0  1  0  0]
[ 0  0 10  0]
[ 0  0  0  1]
sage: E*A
[  0  1  2  3  4  5  6  7  8  9]
[10 11 12 13 14 15 16 17 18 19]
[200 210 220 230 240 250 260 270 280 290]
[ 30  31  32  33  34  35  36  37  38  39]

sage: E = elementary_matrix(4, row1=2, row2=1, scale=10); E
[ 1  0  0  0]
[ 0  1  0  0]
[ 0 10  1  0]
[ 0  0  0  1]
sage: E*A
[  0  1  2  3  4  5  6  7  8  9]
[10 11 12 13 14 15 16 17 18 19]
[120 131 142 153 164 175 186 197 208 219]
[ 30  31  32  33  34  35  36  37  38  39]
```

Over the rationals, now as column operations. Recall that row and column numbering begins at zero. Checks now have the elementary matrix on the right.

```
sage: A = matrix(QQ, 5, 4, range(20)); A
[ 0  1  2  3]
[ 4  5  6  7]
[ 8  9 10 11]
[12 13 14 15]
[16 17 18 19]

sage: E = elementary_matrix(QQ, 4, col1=1, col2=3); E
[1 0 0 0]
[0 0 0 1]
[0 0 1 0]
[0 1 0 0]
sage: A*E
[ 0  3  2  1]
[ 4  7  6  5]
[ 8 11 10  9]
[12 15 14 13]
[16 19 18 17]
```

```

sage: E = elementary_matrix(QQ, 4, coll=2, scale=1/2); E
[ 1  0  0  0]
[ 0  1  0  0]
[ 0  0 1/2  0]
[ 0  0  0  1]
sage: A*E
[ 0  1  1  3]
[ 4  5  3  7]
[ 8  9  5 11]
[12 13  7 15]
[16 17  9 19]

sage: E = elementary_matrix(QQ, 4, coll=2, col2=1, scale=10); E
[ 1  0  0  0]
[ 0  1 10  0]
[ 0  0  1  0]
[ 0  0  0  1]
sage: A*E
[ 0  1 12  3]
[ 4  5 56  7]
[ 8  9 100 11]
[12 13 144 15]
[16 17 188 19]

```

An elementary matrix is always nonsingular. Then repeated row operations can be represented by products of elementary matrices, and this product is again nonsingular. If row operations are to preserve fundamental properties of a matrix (like rank), we do not allow scaling a row by zero. Similarly, the corresponding elementary matrix is not constructed. Also, we do not allow adding a multiple of a row to itself, since this could also lead to a new zero row.

```

sage: A = matrix(QQ, 4, 10, range(40)); A
[ 0  1  2  3  4  5  6  7  8  9]
[10 11 12 13 14 15 16 17 18 19]
[20 21 22 23 24 25 26 27 28 29]
[30 31 32 33 34 35 36 37 38 39]

sage: E1 = elementary_matrix(QQ, 4, row1=0, row2=1)
sage: E2 = elementary_matrix(QQ, 4, row1=3, row2=0, scale=100)
sage: E = E2*E1
sage: E.is_singular()
False
sage: E*A
[ 10  11  12  13  14  15  16  17  18  19]
[  0   1   2   3   4   5   6   7   8   9]
[ 20  21  22  23  24  25  26  27  28  29]
[1030 1131 1232 1333 1434 1535 1636 1737 1838 1939]

sage: E3 = elementary_matrix(QQ, 4, row1=3, scale=0)
Traceback (most recent call last):
...
ValueError: scale parameter of row of elementary matrix must be non-zero

sage: E4 = elementary_matrix(QQ, 4, row1=3, row2=3, scale=12)
Traceback (most recent call last):
...
ValueError: cannot add a multiple of a row to itself

```

If the ring is not specified, and a scale parameter is given, the base ring for the matrix is chosen to contain the

scale parameter. Otherwise, if no ring is given, the default is the integers.

```
sage: E = elementary_matrix(4, row1=1, row2=3)
```

```
sage: E.parent()
```

```
Full MatrixSpace of 4 by 4 dense matrices over Integer Ring
```

```
sage: E = elementary_matrix(4, row1=1, scale=4/3)
```

```
sage: E.parent()
```

```
Full MatrixSpace of 4 by 4 dense matrices over Rational Field
```

```
sage: E = elementary_matrix(4, row1=1, scale=I)
```

```
sage: E.parent()
```

```
Full MatrixSpace of 4 by 4 dense matrices over Symbolic Ring
```

```
sage: E = elementary_matrix(4, row1=1, scale=CDF(I))
```

```
sage: E.parent()
```

```
Full MatrixSpace of 4 by 4 dense matrices over Complex Double Field
```

```
sage: E = elementary_matrix(4, row1=1, scale=QQbar(I))
```

```
sage: E.parent()
```

```
Full MatrixSpace of 4 by 4 dense matrices over Algebraic Field
```

Returned matrices have a dense implementation by default, but a sparse implementation may be requested.

```
sage: E = elementary_matrix(4, row1=0, row2=1)
```

```
sage: E.is_dense()
```

```
True
```

```
sage: E = elementary_matrix(4, row1=0, row2=1, sparse=True)
```

```
sage: E.is_sparse()
```

```
True
```

And the ridiculously small cases. The zero-row matrix cannot be built since then there are no rows to manipulate.

```
sage: elementary_matrix(QQ, 1, row1=0, row2=0)
```

```
[1]
```

```
sage: elementary_matrix(QQ, 0, row1=0, row2=0)
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: size of elementary matrix must be 1 or greater, not 0
```

TESTS:

```
sage: E = elementary_matrix('junk', 5, row1=3, row2=1, scale=12)
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: optional first parameter must be a ring, not junk
```

```
sage: E = elementary_matrix(5, row1=3, scale='junk')
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: scale must be an element of some ring, not junk
```

```
sage: E = elementary_matrix(ZZ, 5, row1=3, col2=3, scale=12)
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: received an unexpected keyword: col2=3
```

```
sage: E = elementary_matrix(QQ, row1=3, scale=12)
```

```
Traceback (most recent call last):
```

```
...
```

```

ValueError: size of elementary matrix must be given

sage: E = elementary_matrix(ZZ, 4/3, row1=3, row2=1, scale=12)
Traceback (most recent call last):
...
TypeError: size of elementary matrix must be an integer, not 4/3

sage: E = elementary_matrix(ZZ, -3, row1=3, row2=1, scale=12)
Traceback (most recent call last):
...
ValueError: size of elementary matrix must be 1 or greater, not -3

sage: E = elementary_matrix(ZZ, 5, row2=1, scale=12)
Traceback (most recent call last):
...
ValueError: row1 or col1 must be specified

sage: E = elementary_matrix(ZZ, 5, row1=3, col1=3, scale=12)
Traceback (most recent call last):
...
ValueError: cannot specify both row1 and col1

sage: E = elementary_matrix(ZZ, 5, row1=4/3, row2=1, scale=12)
Traceback (most recent call last):
...
TypeError: row of elementary matrix must be an integer, not 4/3

sage: E = elementary_matrix(ZZ, 5, col1=5, col2=1, scale=12)
Traceback (most recent call last):
...
ValueError: column of elementary matrix must be positive and smaller than 5, not 5

sage: E = elementary_matrix(ZZ, 5, col1=3, col2=4/3, scale=12)
Traceback (most recent call last):
...
TypeError: column of elementary matrix must be an integer, not 4/3

sage: E = elementary_matrix(ZZ, 5, row1=3, row2=-1, scale=12)
Traceback (most recent call last):
...
ValueError: row of elementary matrix must be positive and smaller than 5, not -1

sage: E = elementary_matrix(ZZ, 5, row1=3, row2=1, scale=4/3)
Traceback (most recent call last):
...
TypeError: scale parameter of elementary matrix must an element of Integer Ring, not 4/3

sage: E = elementary_matrix(ZZ, 5, row1=3)
Traceback (most recent call last):
...
ValueError: insufficient parameters provided to construct elementary matrix

sage: E = elementary_matrix(ZZ, 5, row1=3, row2=3, scale=12)
Traceback (most recent call last):
...
ValueError: cannot add a multiple of a row to itself

sage: E = elementary_matrix(ZZ, 5, col1=3, scale=0)

```

```
Traceback (most recent call last):
...
ValueError: scale parameter of column of elementary matrix must be non-zero
```

AUTHOR:

•Rob Beezer (2011-03-04)

`sage.matrix.constructor.identity_matrix(ring, n=0, sparse=False)`

This function is available as `identity_matrix(...)` and `matrix.identity(...)`.

Return the $n \times n$ identity matrix over the given ring.

The default ring is the integers.

EXAMPLES:

```
sage: M = identity_matrix(QQ, 2); M
[1 0]
[0 1]
sage: M.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: M = identity_matrix(2); M
[1 0]
[0 1]
sage: M.parent()
Full MatrixSpace of 2 by 2 dense matrices over Integer Ring
sage: M.is_mutable()
True
sage: M = identity_matrix(3, sparse=True); M
[1 0 0]
[0 1 0]
[0 0 1]
sage: M.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Integer Ring
sage: M.is_mutable()
True
```

`sage.matrix.constructor.ith_to_zero_rotation_matrix(v, i, ring=None)`

This function is available as `ith_to_zero_rotation_matrix(...)` and `matrix.ith_to_zero_rotation(...)`.

Return a rotation matrix that sends the i -th coordinates of the vector v to zero by doing a rotation with the $(i-1)$ -th coordinate.

INPUT:

- v - vector
- i - integer
- $ring$ - ring (optional, default: None) of the resulting matrix

OUTPUT:

A matrix

EXAMPLES:

```
sage: from sage.matrix.constructor import ith_to_zero_rotation_matrix
sage: v = vector((1, 2, 3))
sage: ith_to_zero_rotation_matrix(v, 2)
[          1          0          0]
[          0 2/13*sqrt(13) 3/13*sqrt(13)]
```



```

[
      0 -3/13*sqrt(13)  2/13*sqrt(13)]
sage: ith_to_zero_rotation_matrix(v, 2) * v
(1, sqrt(13), 0)

sage: ith_to_zero_rotation_matrix(v, 0)
[ 3/10*sqrt(10)      0 -1/10*sqrt(10)]
[      0      1      0]
[ 1/10*sqrt(10)      0  3/10*sqrt(10)]
sage: ith_to_zero_rotation_matrix(v, 1)
[ 1/5*sqrt(5)  2/5*sqrt(5)      0]
[-2/5*sqrt(5)  1/5*sqrt(5)      0]
[      0      0      1]
sage: ith_to_zero_rotation_matrix(v, 2)
[      1      0      0]
[      0  2/13*sqrt(13)  3/13*sqrt(13)]
[      0 -3/13*sqrt(13)  2/13*sqrt(13)]

sage: ith_to_zero_rotation_matrix(v, 0) * v
(0, 2, sqrt(10))
sage: ith_to_zero_rotation_matrix(v, 1) * v
(sqrt(5), 0, 3)
sage: ith_to_zero_rotation_matrix(v, 2) * v
(1, sqrt(13), 0)

```

Other ring:

```

sage: ith_to_zero_rotation_matrix(v, 2, ring=RR)
[ 1.000000000000000 0.000000000000000 0.000000000000000]
[ 0.000000000000000 0.554700196225229 0.832050294337844]
[ 0.000000000000000 -0.832050294337844 0.554700196225229]
sage: ith_to_zero_rotation_matrix(v, 2, ring=RDF)
[      1.0      0.0      0.0]
[      0.0 0.554700196225 0.832050294338]
[      0.0 -0.832050294338 0.554700196225]

```

On the symbolic ring:

```

sage: x,y,z = var('x,y,z')
sage: v = vector((x,y,z))
sage: ith_to_zero_rotation_matrix(v, 2)
[      1      0      0]
[      0  y/sqrt(y^2 + z^2)  z/sqrt(y^2 + z^2)]
[      0 -z/sqrt(y^2 + z^2)  y/sqrt(y^2 + z^2)]
sage: ith_to_zero_rotation_matrix(v, 2) * v
(x, y^2/sqrt(y^2 + z^2) + z^2/sqrt(y^2 + z^2), 0)

```

TESTS:

```

sage: ith_to_zero_rotation_matrix((1,0,0), 0)
[ 0  0 -1]
[ 0  1  0]
[ 1  0  0]
sage: ith_to_zero_rotation_matrix((1,0,0), 1)
[1 0 0]
[0 1 0]
[0 0 1]
sage: ith_to_zero_rotation_matrix((1,0,0), 2)
[1 0 0]
[0 1 0]

```

```
[0 0 1]
```

AUTHORS:

Sebastien Labbe (April 2010)

`sage.matrix.constructor.jordan_block` (*eigenvalue*, *size*, *sparse=False*)

This function is available as `jordan_block(...)` and `matrix.jordan_block(...)`.

Returns the Jordan block for the given eigenvalue with given size.

INPUT:

- *eigenvalue* - eigenvalue for the diagonal entries of the block
- *size* - size of the square matrix
- *sparse* - (default: False) - if True, return a sparse matrix

EXAMPLE:

```
sage: jordan_block(5, 3)
[5 1 0]
[0 5 1]
[0 0 5]
```

TESTS:

```
sage: jordan_block(6.2, 'junk')
Traceback (most recent call last):
...
TypeError: size of Jordan block needs to be an integer, not junk
sage: jordan_block(6.2, -1)
Traceback (most recent call last):
...
ValueError: size of Jordan block must be non-negative, not -1
```

`sage.matrix.constructor.matrix_method` (*func=None*, *name=None*)

Allows a function to be tab-completed on the global matrix constructor object.

INPUT:

- **function* - a single argument. The function that is being decorated.
- ***kwds* - a single optional keyword argument *name=<string>*. The name of the corresponding method in the global matrix constructor object. If not given, it is derived from the function name.

EXAMPLES:

```
sage: from sage.matrix.constructor import matrix_method
sage: def foo_matrix(n): return matrix.diagonal(range(n))
sage: matrix_method(foo_matrix)
<function foo_matrix at ...>
sage: matrix.foo(5)
[0 0 0 0 0]
[0 1 0 0 0]
[0 0 2 0 0]
[0 0 0 3 0]
[0 0 0 0 4]
sage: matrix_method(foo_matrix, name='bar')
<function foo_matrix at ...>
sage: matrix.bar(3)
[0 0 0]
```

```
[0 1 0]
[0 0 2]
```

`sage.matrix.constructor.ncols_from_dict(d)`

Given a dictionary that defines a sparse matrix, return the number of columns that matrix should have.

This is for internal use by the matrix function.

INPUT:

- *d* - dict

OUTPUT:

integer

EXAMPLES:

```
sage: sage.matrix.constructor.ncols_from_dict({})
0
```

Here the answer is 301 not 300, since there is a 0-th row.

```
sage: sage.matrix.constructor.ncols_from_dict({(4,300):10})
301
```

`sage.matrix.constructor.nrows_from_dict(d)`

Given a dictionary that defines a sparse matrix, return the number of rows that matrix should have.

This is for internal use by the matrix function.

INPUT:

- *d* - dict

OUTPUT:

integer

EXAMPLES:

```
sage: sage.matrix.constructor.nrows_from_dict({})
0
```

Here the answer is 301 not 300, since there is a 0-th row.

```
:: sage: sage.matrix.constructor.nrows_from_dict({(300,4):10}) 301
```

`sage.matrix.constructor.ones_matrix(ring, nrows=None, ncols=None, sparse=False)`

This function is available as `ones_matrix(...)` and `matrix.ones(...)`.

Return a matrix with all entries equal to 1.

CALL FORMATS:

In each case, the optional keyword `sparse` can be used.

1. `ones_matrix(ring, nrows, ncols)`

2. `ones_matrix(ring, nrows)`

3. `ones_matrix(nrows, ncols)`

4. `ones_matrix(nrows)`

INPUT:

- ring - default: ZZ - base ring for the matrix.
- nrows - number of rows in the matrix.
- ncols - number of columns in the matrix. If omitted, defaults to the number of rows, producing a square matrix.
- sparse - default: False - if True creates a sparse representation.

OUTPUT:

A matrix of size `nrows` by `ncols` over the `ring` with every entry equal to 1. While the result is far from sparse, you may wish to choose a sparse representation when mixing this matrix with other sparse matrices.

EXAMPLES:

A call specifying the ring and the size.

```
sage: M= ones_matrix(QQ, 2, 5); M
[1 1 1 1 1]
[1 1 1 1 1]
sage: M.parent()
Full MatrixSpace of 2 by 5 dense matrices over Rational Field
```

Without specifying the number of columns, the result is square.

```
sage: M = ones_matrix(RR, 2); M
[1.000000000000000 1.000000000000000]
[1.000000000000000 1.000000000000000]
sage: M.parent()
Full MatrixSpace of 2 by 2 dense matrices over Real Field with 53 bits of precision
```

The ring defaults to the integers if not given.

```
sage: M = ones_matrix(2, 3); M
[1 1 1]
[1 1 1]
sage: M.parent()
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
```

A lone integer input produces a square matrix over the integers.

```
sage: M = ones_matrix(3); M
[1 1 1]
[1 1 1]
[1 1 1]
sage: M.parent()
Full MatrixSpace of 3 by 3 dense matrices over Integer Ring
```

The result can have a sparse implementation.

```
sage: M = ones_matrix(3, 1, sparse=True); M
[1]
[1]
[1]
sage: M.parent()
Full MatrixSpace of 3 by 1 sparse matrices over Integer Ring
```

Giving just a ring will yield an error.

```
sage: ones_matrix(CC)
Traceback (most recent call last):
```

```
...
ValueError: constructing an all ones matrix requires at least one dimension
```

`sage.matrix.constructor.prepare(w)`

Given a list `w` of numbers, find a common ring that they all canonically map to, and return the list of images of the elements of `w` in that ring along with the ring.

This is for internal use by the matrix function.

INPUT:

- `w` - list

OUTPUT:

list, ring

EXAMPLES:

```
sage: sage.matrix.constructor.prepare([-2, Mod(1,7)])
([5, 1], Ring of integers modulo 7)
```

Notice that the elements must all canonically coerce to a common ring (since `Sequence` is called):

```
sage: sage.matrix.constructor.prepare([2/1, Mod(1,7)])
Traceback (most recent call last):
...
TypeError: unable to find a common ring for all elements
```

`sage.matrix.constructor.prepare_dict(w)`

Given a dictionary `w` of numbers, find a common ring that they all canonically map to, and return the dictionary of images of the elements of `w` in that ring along with the ring.

This is for internal use by the matrix function.

INPUT:

- `w` - dict

OUTPUT:

dict, ring

EXAMPLES:

```
sage: sage.matrix.constructor.prepare_dict({(0,1):2, (4,10):Mod(1,7)})
({(0, 1): 2, (4, 10): 1}, Ring of integers modulo 7)
```

`sage.matrix.constructor.random_diagonalizable_matrix(parent, eigenvalues=None, dimensions=None)`

This function is available as `random_diagonalizable_matrix(...)` and `matrix.random_diagonalizable(...)`.

Create a random matrix that diagonalizes nicely.

To be used as a teaching tool. Return matrices have only real eigenvalues.

INPUT:

If eigenvalues and dimensions are not specified in a list, they will be assigned randomly.

- `parent` - the desired size of the square matrix.
- `eigenvalues` - the list of desired eigenvalues (default=None).
- `dimensions` - the list of dimensions corresponding to each eigenspace (default=None).

OUTPUT:

A square, diagonalizable, matrix with only integer entries. The eigenspaces of this matrix, if computed by hand, give basis vectors with only integer entries.

Note: It is easiest to use this function via a call to the `random_matrix()` function with the `algorithm='diagonalizable'` keyword. We provide one example accessing this function directly, while the remainder will use this more general function.

EXAMPLES:

A diagonalizable matrix, size 5.

```
sage: from sage.matrix.constructor import random_diagonalizable_matrix
sage: matrix_space = sage.matrix.matrix_space.MatrixSpace(QQ, 5)
sage: A=random_diagonalizable_matrix(matrix_space); A # random
[ 10  18   8   4 -18]
[ 20  10   8   4 -16]
[-60 -54 -22 -12  18]
[-60 -54 -24  -6   6]
[-20 -18  -8  -4   8]
sage: A.eigenvalues() # random
[10, 6, 2, -8, -10]
sage: S=A.right_eigenmatrix()[1]; S # random
[ 1  1  1  1  0]
[ 1  1  1  0  1]
[-3 -3 -4 -3 -3]
[-3 -4 -3 -3 -3]
[-1 -1 -1 -1 -1]
sage: S_inverse=S.inverse(); S_inverse # random
[ 1  1  1  1 -5]
[ 0  0  0 -1  3]
[ 0  0 -1  0  3]
[ 0 -1  0  0 -1]
[-1  0  0  0 -1]
sage: S_inverse*A*S # random
[ 10  0  0  0  0]
[  0  6  0  0  0]
[  0  0  2  0  0]
[  0  0  0 -8  0]
[  0  0  0  0 -10]
```

A diagonalizable matrix with eigenvalues and dimensions designated, with a check that if eigenvectors were calculated by hand entries would all be integers.

```
sage: B=random_matrix(QQ, 6, algorithm='diagonalizable', eigenvalues=[-12,4,6],dimensions=[2,3,1])
[ -52   32  240 -464  -96 -520]
[   6    4  -48   72   36   90]
[  46  -32 -108  296  -12  274]
[  24  -16  -64  164    0  152]
[  18  -16    0   72  -48   30]
[   2    0 -16   24   12   34]
sage: all([x in ZZ for x in (B-(-12*identity_matrix(6))).rref().list()])
True
sage: all([x in ZZ for x in (B-(4*identity_matrix(6))).rref().list()])
True
sage: all([x in ZZ for x in (B-(6*identity_matrix(6))).rref().list()])
True
sage: S=B.right_eigenmatrix()[1]; S_inverse=S.inverse(); S_inverse*B*S # random
```

```
[ 6  0  0  0  0  0]
[ 0 -12 0  0  0  0]
[ 0  0 -12 0  0  0]
[ 0  0  0  4  0  0]
[ 0  0  0  0  4  0]
[ 0  0  0  0  0  4]
```

TESTS:

Eigenvalues must all be integers.

```
sage: random_matrix(QQ, 3, algorithm='diagonalizable', eigenvalues=[2+I, 2-I, 2], dimensions=[1, 1, 1])
Traceback (most recent call last):
...
TypeError: eigenvalues must be integers.
```

Diagonal matrices must be square.

```
sage: random_matrix(QQ, 5, 7, algorithm='diagonalizable', eigenvalues=[-5, 2, -3], dimensions=[1, 1, 1, 1, 1])
Traceback (most recent call last):
...
TypeError: a diagonalizable matrix must be square.
```

A list of eigenvalues must be accompanied with a list of dimensions.

```
sage: random_matrix(QQ, 10, algorithm='diagonalizable', eigenvalues=[4, 8])
Traceback (most recent call last):
...
ValueError: the list of eigenvalues must have a list of dimensions corresponding to each eigenvalue.
```

A list of dimensions must be accompanied with a list of eigenvalues.

```
sage: random_matrix(QQ, 10, algorithm='diagonalizable', dimensions=[2, 2, 4, 2])
Traceback (most recent call last):
...
ValueError: the list of dimensions must have a list of corresponding eigenvalues.
```

The sum of the eigenvalue dimensions must equal the size of the matrix.

```
sage: random_matrix(QQ, 12, algorithm='diagonalizable', eigenvalues=[4, 2, 6, -1], dimensions=[2, 3, 5, 1])
Traceback (most recent call last):
...
ValueError: the size of the matrix must equal the sum of the dimensions.
```

Each eigenspace dimension must be at least 1.

```
sage: random_matrix(QQ, 9, algorithm='diagonalizable', eigenvalues=[-15, 22, 8, -4, 90, 12], dimensions=[1, 1, 1, 1, 1, 1])
Traceback (most recent call last):
...
ValueError: eigenspaces must have a dimension of at least 1.
```

Each eigenvalue must have a corresponding eigenspace dimension.

```
sage: random_matrix(QQ, 12, algorithm='diagonalizable', eigenvalues=[4, 2, 6, -1], dimensions=[4, 3, 5])
Traceback (most recent call last):
...
ValueError: each eigenvalue must have a corresponding dimension and each dimension a corresponding eigenvalue.
```

Each dimension must have an eigenvalue paired to it.

```
sage: random_matrix(QQ, 12, algorithm='diagonalizable', eigenvalues=[4, 2, 6], dimensions=[2, 3, 5, 2])
Traceback (most recent call last):
...
ValueError: each eigenvalue must have a corresponding dimension and each dimension a corresponding
```

TODO:

Modify the routine to allow for complex eigenvalues.

AUTHOR:

Billy Wonderly (2010-07)

```
sage.matrix.constructor.random_echelonizable_matrix(parent, rank, upper_bound=None,
                                                    max_tries=100)
```

This function is available as `random_echelonizable_matrix(...)` and `matrix.random_echelonizable(...)`.

Generate a matrix of a desired size and rank, over a desired ring, whose reduced row-echelon form has only integral values.

INPUT:

- `parent` - A matrix space specifying the base ring, dimensions and representation (dense/sparse) for the result. The base ring must be exact.
- `rank` - Rank of result, i.e the number of non-zero rows in the reduced row echelon form.
- `upper_bound` - If designated, size control of the matrix entries is desired. Set `upper_bound` to 1 more than the maximum value entries can achieve. If `None`, no size control occurs. But see the warning below. (default: `None`)
- `max_tries` - If designated, number of tries used to generate each new random row; only matters when `upper_bound!=None`. Used to prevent endless looping. (default: 100)

OUTPUT:

A matrix not in reduced row-echelon form with the desired dimensions and properties.

Warning: When `upper_bound` is set, it is possible for this constructor to fail with a `ValueError`. This may happen when the `upper_bound`, `rank` and/or matrix dimensions are all so small that it becomes infeasible or unlikely to create the requested matrix. If you *must* have this routine return successfully, do not set `upper_bound`.

Note: It is easiest to use this function via a call to the `random_matrix()` function with the `algorithm='echelonizable'` keyword. We provide one example accessing this function directly, while the remainder will use this more general function.

EXAMPLES:

Generated matrices have the desired dimensions, rank and entry size. The matrix in reduced row-echelon form has only integer entries.

```
sage: from sage.matrix.constructor import random_echelonizable_matrix
sage: matrix_space = sage.matrix.matrix_space.MatrixSpace(QQ, 5, 6)
sage: A=random_echelonizable_matrix(matrix_space, rank=4, upper_bound=40); A # random
[ 1 -1  1 -3 -4  6]
[ 5 -4  0  8  4 19]
[ -3  3 -2  4  7 -16]
[ -4  5 -7 26 31 -31]
```



```
[ 2 -3 4 -11 -14 17]
sage: A.rank()
4
sage: max(map(abs,A.list()))<40
True
sage: A.rref()==A.rref().change_ring(ZZ)
True
```

An example with default settings (i.e. no entry size control).

```
sage: C=random_matrix(QQ, 6, 7, algorithm='echelonizable', rank=5); C # random
[ 1 0 5 -2 -26 -16 0]
[ -3 1 -19 6 97 61 1]
[ 0 4 -15 -1 71 50 3]
[ 2 4 -9 0 39 25 8]
[ 2 2 3 -3 -18 -9 3]
[ -3 -4 -2 14 14 -6 4]
sage: C.rank()
5
sage: C.rref()==C.rref().change_ring(ZZ)
True
```

A matrix without size control may have very large entry sizes.

```
sage: D=random_matrix(ZZ, 7, 8, algorithm='echelonizable', rank=6); D # random
[ 9 -53 -255 45 -1519 4043 9819 3324]
[ 3 -14 -64 8 -369 972 2350 810]
[ 2 -14 -65 9 -377 1000 2420 829]
[ 4 -24 -116 21 -693 1846 4485 1516]
[ -3 14 68 -16 426 -1134 -2767 -919]
[ -5 21 92 -13 548 -1432 -3466 -1183]
[ 1 -9 -42 7 -254 670 1624 547]
```

Matrices can be generated over any exact ring.

```
sage: F.<a>=GF(2^3)
sage: B=random_matrix(F, 4, 5, algorithm='echelonizable', rank=4, upper_bound=None); B # random
[ a + 1 a^2 + a + 1 a^2 0 1]
[ 1 a a + 1 a^2 a^2 + a]
[ a a^2 a^2 + a + 1 a + 1 1]
[ 1 0 a^2 + a + 1 0 a^2 + 1]
sage: B.rank()
4
```

Square matrices over ZZ or QQ with full rank are always unimodular.

```
sage: E=random_matrix(QQ, 7, 7, algorithm='echelonizable', rank=7); E # random
[ 1 -1 5 12 -24 -41 47]
[ 0 1 -1 3 0 -11 40]
[ 1 -1 6 6 -19 -20 -11]
[ -2 1 -10 -12 35 44 -4]
[ 3 -1 9 7 -35 -40 -18]
[ 0 0 0 -4 4 13 -32]
[ 3 -3 11 6 -33 -31 -35]
sage: det(E)
1
```

TESTS:

Matrices must have a rank zero or greater, and less than both the number of rows and the number of columns.

```
sage: random_matrix(QQ, 3, 4, algorithm='echelonizable', rank=-1)
Traceback (most recent call last):
...
ValueError: matrices must have rank zero or greater.
sage: random_matrix(QQ, 3, 8, algorithm='echelonizable', rank=4)
Traceback (most recent call last):
...
ValueError: matrices cannot have rank greater than min(ncols,nrows).
sage: random_matrix(QQ, 8, 3, algorithm='echelonizable', rank=4)
Traceback (most recent call last):
...
ValueError: matrices cannot have rank greater than min(ncols,nrows).
```

The base ring must be exact.

```
sage: random_matrix(RR, 3, 3, algorithm='echelonizable', rank=2)
Traceback (most recent call last):
...
TypeError: the base ring must be exact.
```

Works for rank==1, too.

```
sage: random_matrix(QQ, 3, 3, algorithm='echelonizable', rank=1).ncols()
3
```

AUTHOR:

Billy Wonderly (2010-07)

```
sage.matrix.constructor.random_matrix(ring, nrows, ncols=None, algorithm='randomize',
                                         *args, **kws)
```

This function is available as `random_matrix(...)` and `matrix.random(...)`.

Return a random matrix with entries in a specified ring, and possibly with additional properties.

INPUT:

- `ring` - base ring for entries of the matrix
- `nrows` - Integer; number of rows
- `ncols` - (default: None); number of columns; if None defaults to `nrows`
- `algorithm` - (default: `randomize`); determines what properties the matrix will have. See examples below for possible additional arguments.
 - `randomize` - randomize the elements of the matrix, possibly controlling the density of non-zero entries.
 - `echelon_form` - creates a matrix in echelon form
 - `echelonizable` - creates a matrix that has a predictable echelon form
 - `subspaces` - creates a matrix whose four subspaces, when explored, have reasonably sized, integral valued, entries.
 - `unimodular` - creates a matrix of determinant 1.
 - `diagonalizable` - creates a diagonalizable matrix whose eigenvectors, if computed by hand, will have only integer entries.
- `*args, **kws` - arguments and keywords to describe additional properties. See more detailed documentation below.

Warning: An upper bound on the absolute value of the entries may be set when the algorithm is echelonizable or unimodular. In these cases it is possible for this constructor to fail with a `ValueError`. If you *must* have this routine return successfully, do not set `upper_bound`. This behavior can be partially controlled by a `max_tries` keyword.

Note: When constructing matrices with random entries and no additional properties (i.e. when `algorithm='randomize'`), most of the randomness is controlled by the `random_element` method for elements of the base ring of the matrix, so the documentation of that method may be relevant or useful. Also, the default is to not create zero entries, unless the `density` keyword is set to something strictly less than one.

EXAMPLES:

Random integer matrices. With no arguments, the majority of the entries are -1 and 1, never zero, and rarely “large.”

```
sage: random_matrix(ZZ, 5, 5)
[ -8   2   0   0   1]
[ -1   2   1 -95  -1]
[ -2 -12   0   0   1]
[ -1   1  -1  -2  -1]
[  4  -4  -6   5   0]
```

The `distribution` keyword set to `uniform` will limit values between -2 and 2, and never zero.

```
sage: random_matrix(ZZ, 5, 5, distribution='uniform')
[ 1  0 -2  1  1]
[ 1  0  0  0  2]
[-1 -2  0  2 -2]
[-1 -1  1  1  2]
[ 0 -2 -1  0  0]
```

The `x` and `y` keywords can be used to distribute entries uniformly. When both are used `x` is the minimum and `y` is one greater than the the maximum. But still entries are never zero, even if the range contains zero.

```
sage: random_matrix(ZZ, 4, 8, x=70, y=100)
[81 82 70 81 78 71 79 94]
[80 98 89 87 91 94 94 77]
[86 89 85 92 95 94 72 89]
[78 80 89 82 94 72 90 92]
```

```
sage: random_matrix(ZZ, 3, 7, x=-5, y=5)
[-3  3  1 -5  3  1  2]
[ 3  3  0  3 -5 -2  1]
[ 0 -2 -2  2 -3 -4 -2]
```

If only `x` is given, then it is used as the upper bound of a range starting at 0.

```
sage: random_matrix(ZZ, 5, 5, x=25)
[20 16  8  3  8]
[ 8  2  2 14  5]
[18 18 10 20 11]
[19 16 17 15  7]
[ 0 24  3 17 24]
```

To allow, and control, zero entries use the `density` keyword at a value strictly below the default of 1.0, even if distributing entries across an interval that does not contain zero already. Note that for a square matrix it is only necessary to set a single dimension.

```
sage: random_matrix(ZZ, 5, x=-10, y=10, density=0.75)
[-6  1  0  0  0]
[ 9  0  0  4  1]
[-6  0  0 -8  0]
[ 0  4  0  6  0]
[ 1 -9  0  0 -8]
```

```
sage: random_matrix(ZZ, 5, x=20, y=30, density=0.75)
[ 0 28  0 27  0]
[25 28 20  0  0]
[ 0 21  0 21  0]
[ 0 28 22  0  0]
[ 0  0  0 26 24]
```

It is possible to construct sparse matrices, where it may now be advisable (but not required) to control the density of nonzero entries.

```
sage: A=random_matrix(ZZ, 5, 5)
sage: A.is_sparse()
False
sage: A=random_matrix(ZZ, 5, 5, sparse=True)
sage: A.is_sparse()
True

sage: random_matrix(ZZ, 5, 5, density=0.3, sparse=True)
[ 4  0  0  0 -1]
[ 0  0  0  0 -7]
[ 0  0  2  0  0]
[ 0  0  1  0 -4]
[ 0  0  0  0  0]
```

For algorithm testing you might want to control the number of bits, say 10,000 entries, each limited to 16 bits.

```
sage: A = random_matrix(ZZ, 100, 100, x=2^16); A
100 x 100 dense matrix over Integer Ring (use the '.str()' method to see the entries)
```

Random rational matrices. Now `num_bound` and `den_bound` control the generation of random elements, by specifying limits on the absolute value of numerators and denominators (respectively). Entries will be positive and negative (map the absolute value function through the entries to get all positive values), and zeros are avoided unless the density is set. If either the numerator or denominator bound (or both) is not used, then the values default to the distribution for `ZZ` described above that is most frequently positive or negative one.

```
sage: random_matrix(QQ, 2, 8, num_bound=20, den_bound=4)
[-1/2   6   13  -12  -2/3  -1/4   5   5]
[-9/2   5/3   19 15/2 19/2 20/3 -13/4   0]

sage: random_matrix(QQ, 4, density = 0.5, sparse=True)
[  0   71   0 -1/2]
[  0   0   0   0]
[31/85   0 -31/2   0]
[  1 -1/4   0   0]

sage: A = random_matrix(QQ, 3, 10, num_bound = 99, den_bound = 99)
sage: positives = map(abs, A.list())
sage: matrix(QQ, 3, 10, positives)
[61/18 47/41 1/22  1/2 75/68  6/7   1   1/2 72/41  7/3]
[33/13  9/2 40/21 45/46 17/22   1 70/79 97/71  7/24 12/5]
[13/8  8/25  1/3 61/14 92/45  4/85  3/38 95/16 82/71  1/5]
```

```
sage: random_matrix(QQ, 4, 10, den_bound = 10)
[ -1      0 1/8 1/6 2/9 -1/6 1/5 -1/8 1/5 -1/5]
[ 1/9 1/5  -1 2/9 1/4 -1/7 1/8 -1/9  0    2]
[ 2/3  2 1/8  -2  0    0  -2    2    0 -1/2]
[  0    2  1 -2/3  0    0 1/6    0 -1/3 -2/9]
```

Random matrices over other rings. Several classes of matrices have specialized `randomize()` methods. You can locate these with the Sage command:

```
search_def('randomize')
```

The default implementation of `randomize()` relies on the `random_element()` method for the base ring. The density and sparse keywords behave as described above.

```
sage: K.<a>=FiniteField(3^2)
sage: random_matrix(K, 2, 5)
[      1      a      1 2*a + 1      2]
[ 2*a  a + 2      0      2      1]

sage: random_matrix(RR, 3, 4, density=0.66)
[ 0.0000000000000000 -0.806696574554030 -0.693915509972359  0.0000000000000000]
[ 0.629781664418083  0.0000000000000000 -0.833709843116637  0.0000000000000000]
[ 0.922346867410064  0.0000000000000000  0.0000000000000000 -0.940316454178921]

sage: A = random_matrix(ComplexField(32), 3, density=0.8, sparse=True); A
[      0.000000000  0.399739209 + 0.909948633*I      0.000000000]
[-0.361911424 - 0.455087671*I -0.687810605 + 0.460619713*I  0.625520058 - 0.360952012*I]
[      0.000000000      0.000000000 -0.162196416 - 0.193242896*I]
sage: A.is_sparse()
True
```

Random matrices in echelon form. The `algorithm='echelon_form'` keyword, along with a requested number of non-zero rows (`num_pivots`) will return a random matrix in echelon form. When the base ring is `QQ` the result has integer entries. Other exact rings may be also specified.

```
sage: A=random_matrix(QQ, 4, 8, algorithm='echelon_form', num_pivots=3); A # random
[ 1 -5  0 -2  0  1  1 -2]
[ 0  0  1 -5  0 -3 -1  0]
[ 0  0  0  0  1  2 -2  1]
[ 0  0  0  0  0  0  0  0]
sage: A.base_ring()
Rational Field
sage: (A.nrows(), A.ncols())
(4, 8)
sage: A in sage.matrix.matrix_space.MatrixSpace(ZZ, 4, 8)
True
sage: A.rank()
3
sage: A==A.rref()
True
```

For more, see the documentation of the `random_rref_matrix()` function. In the notebook or at the Sage command-line, first execute the following to make this further documentation available:

```
from sage.matrix.constructor import random_rref_matrix
```

Random matrices with predictable echelon forms. The `algorithm='echelonizable'` keyword, along with a requested rank (`rank`) and optional size control (`upper_bound`) will return a random matrix in echelon form. When the base ring is `ZZ` or `QQ` the result has integer entries, whose magnitudes can be limited by the

value of `upper_bound`, and the echelon form of the matrix also has integer entries. Other exact rings may be also specified, but there is no notion of controlling the size. Square matrices of full rank generated by this function always have determinant one, and can be constructed with the `unimodular` keyword.

```
sage: A=random_matrix(QQ, 4, 8, algorithm='echelonizable', rank=3, upper_bound=60); A # random
sage: A.base_ring()
Rational Field
sage: (A.nrows(), A.ncols())
(4, 8)
sage: A in sage.matrix.matrix_space.MatrixSpace(ZZ, 4, 8)
True
sage: A.rank()
3
sage: all([abs(x)<60 for x in A.list()])
True
sage: A.rref() in sage.matrix.matrix_space.MatrixSpace(ZZ, 4, 8)
True
```

For more, see the documentation of the `random_echelonizable_matrix()` function. In the notebook or at the Sage command-line, first execute the following to make this further documentation available:

```
from sage.matrix.constructor import random_echelonizable_matrix
```

Random diagonalizable matrices. The `algorithm='diagonalizable'` keyword, along with a requested matrix size (`size`) and optional lists of eigenvalues (`eigenvalues`) and the corresponding eigenspace dimensions (`dimensions`) will return a random diagonalizable matrix. When the eigenvalues and dimensions are not specified the result will have randomly generated values for both that fit with the designated size.

```
sage: A=random_matrix(QQ, 5, algorithm='diagonalizable', eigenvalues=[2,3,-1], dimensions=[1,2,2])
sage: all([x in ZZ for x in (A-(2*identity_matrix(5))).rref().list()])
True
sage: all([x in ZZ for x in (A-(3*identity_matrix(5))).rref().list()])
True
sage: all([x in ZZ for x in (A-(-1*identity_matrix(5))).rref().list()])
True
sage: A.jordan_form()
[ 2| 0| 0| 0| 0]
[---+---+---+---]
[ 0| 3| 0| 0| 0]
[---+---+---+---]
[ 0| 0| 3| 0| 0]
[---+---+---+---]
[ 0| 0| 0|-1| 0]
[---+---+---+---]
[ 0| 0| 0| 0|-1]
```

For more, see the documentation of the `random_diagonalizable_matrix()` function. In the notebook or at the Sage command-line, first execute the following to make this further documentation available:

```
from sage.matrix.constructor import random_diagonalizable_matrix
```

Random matrices with predictable subspaces. The `algorithm='subspaces'` keyword, along with an optional rank (`rank`) will return a matrix whose natural basis vectors for its four fundamental subspaces, if computed as described in the documentation of the `random_subspaces_matrix()` contain only integer entries. If `rank` is not set, the rank of the matrix will be generated randomly.

```
sage: B=random_matrix(QQ, 5, 6, algorithm='subspaces', rank=3); B #random
sage: B_expanded=B.augment(identity_matrix(5)).rref()
sage: (B.nrows(), B.ncols())
(5, 6)
```

```

sage: all([x in ZZ for x in B_expanded.list()])
True
sage: C=B_expanded.submatrix(0,0,B.nrows()-B.nullity(),B.ncols())
sage: L=B_expanded.submatrix(B.nrows()-B.nullity(),B.ncols())
sage: B.right_kernel()==C.right_kernel()
True
sage: B.row_space()==C.row_space()
True
sage: B.column_space()==L.right_kernel()
True
sage: B.left_kernel()==L.row_space()
True

```

For more, see the documentation of the `random_subspaces_matrix()` function. In the notebook or at the Sage command-line, first execute the following to make this further documentation available:

```
from sage.matrix.constructor import random_subspaces_matrix
```

Random unimodular matrices. The `algorithm='unimodular'` keyword, along with an optional entry size control (`upper_bound`) will return a matrix of determinant 1. When the base ring is `ZZ` or `QQ` the result has integer entries, whose magnitudes can be limited by the value of `upper_bound`.

```

sage: C=random_matrix(QQ, 5, algorithm='unimodular', upper_bound=70); C # random
sage: det(C)
1
sage: C.base_ring()
Rational Field
sage: (C.nrows(), C.ncols())
(5, 5)
sage: all([abs(x)<70 for x in C.list()])
True

```

For more, see the documentation of the `random_unimodular_matrix()` function. In the notebook or at the Sage command-line, first execute the following to make this further documentation available:

```
from sage.matrix.constructor import random_unimodular_matrix
```

TESTS:

We return an error for a bogus value of `algorithm`:

```

sage: random_matrix(ZZ, 5, algorithm = 'bogus')
Traceback (most recent call last):
...
ValueError: random matrix algorithm "bogus" is not recognized

```

AUTHOR:

- William Stein (2007-02-06)
- Rob Beezer (2010-08-25) Documentation, code to allow additional types of output

`sage.matrix.constructor.random_rref_matrix` (*parent*, *num_pivots*)

This function is available as `random_rref_matrix(...)` and `matrix.random_rref(...)`.

Generate a matrix in reduced row-echelon form with a specified number of non-zero rows.

INPUT:

- parent* - A matrix space specifying the base ring, dimensions and representation (dense/sparse) for the result. The base ring must be exact.

- `num_pivots` - The number of non-zero rows in the result, i.e. the rank.

OUTPUT:

A matrix in reduced row echelon form with `num_pivots` non-zero rows. If the base ring is $\mathbb{Z}\mathbb{Z}$ or $\mathbb{Q}\mathbb{Q}$ then the entries are all integers.

Note: It is easiest to use this function via a call to the `random_matrix()` function with the `algorithm='echelon_form'` keyword. We provide one example accessing this function directly, while the remainder will use this more general function.

EXAMPLES:

Matrices generated are in reduced row-echelon form with specified rank. If the base ring is $\mathbb{Q}\mathbb{Q}$ the result has only integer entries.

```
sage: from sage.matrix.constructor import random_rref_matrix
sage: matrix_space = sage.matrix.matrix_space.MatrixSpace(QQ, 5, 6)
sage: A=random_rref_matrix(matrix_space, num_pivots=4); A # random
[ 1  0  0 -6  0 -3]
[ 0  1  0  2  0  3]
[ 0  0  1 -4  0 -2]
[ 0  0  0  0  1  3]
[ 0  0  0  0  0  0]
sage: A.base_ring()
Rational Field
sage: (A.nrows(), A.ncols())
(5, 6)
sage: A in sage.matrix.matrix_space.MatrixSpace(ZZ, 5, 6)
True
sage: A.rank()
4
sage: A==A.rref()
True
```

Matrices can be generated over other exact rings.

```
sage: B=random_matrix(FiniteField(7), 4, 4, algorithm='echelon_form', num_pivots=3); B # random
[1 0 0 0]
[0 1 0 6]
[0 0 1 4]
[0 0 0 0]
sage: B.rank() == 3
True
sage: B.base_ring()
Finite Field of size 7
sage: B==B.rref()
True
```

TESTS:

Rank of a matrix must be an integer.

```
sage: random_matrix(QQ, 120, 56, algorithm='echelon_form', num_pivots=61/2)
Traceback (most recent call last):
...
TypeError: the number of pivots must be an integer.
```

Matrices must be generated over exact fields.


```
sage: random_matrix(RR, 40, 88, algorithm='echelon_form', num_pivots=39)
Traceback (most recent call last):
...
TypeError: the base ring must be exact.
```

Matrices must have the number of pivot columns be less than or equal to the number of rows.

```
sage: C=random_matrix(ZZ, 6,4, algorithm='echelon_form', num_pivots=7); C
Traceback (most recent call last):
...
ValueError: number of pivots cannot exceed the number of rows or columns.
```

Matrices must have the number of pivot columns be less than or equal to the number of columns.

```
sage: D=random_matrix(QQ, 1,3, algorithm='echelon_form', num_pivots=5); D
Traceback (most recent call last):
...
ValueError: number of pivots cannot exceed the number of rows or columns.
```

Matrices must have the number of pivot columns be greater than zero.

```
sage: random_matrix(QQ, 5, 4, algorithm='echelon_form', num_pivots=-1)
Traceback (most recent call last):
...
ValueError: the number of pivots must be zero or greater.
```

AUTHOR:

Billy Wonderly (2010-07)

sage.matrix.constructor.**random_subspaces_matrix**(parent, rank=None)

This function is available as random_subspaces_matrix(...) and matrix.random_subspaces(...).

Create a matrix of the designated size and rank whose right and left null spaces, column space, and row space have desirable properties that simplify the subspaces.

INPUT:

- parent - A matrix space specifying the base ring, dimensions, and representation (dense/sparse) for the result. The base ring must be exact.
- rank - The desired rank of the return matrix (default: None).

OUTPUT:

A matrix whose natural basis vectors for its four subspaces, when computed, have reasonably sized, integral valued, entries.

Note: It is easiest to use this function via a call to the `random_matrix()` function with the `algorithm='subspaces'` keyword. We provide one example accessing this function directly, while the remainder will use this more general function.

EXAMPLES:

A 6x8 matrix with designated rank of 3. The four subspaces are determined using one simple routine in which we augment the original matrix with the equal row dimension identity matrix. The resulting matrix is then put in reduced row-echelon form and the subspaces can then be determined by analyzing subdivisions of this matrix. See the four subspaces routine in [BEEZER] for more.

```
sage: from sage.matrix.constructor import random_subspaces_matrix
sage: matrix_space = sage.matrix.matrix_space.MatrixSpace(QQ, 6, 8)
```

```

sage: B=random_subspaces_matrix(matrix_space, rank=3); B # random
[ 113  339 -46  218 -243  641 -269 -306]
[ -33  -99  13  -63   69 -185   77   90]
[  35  105 -14   67  -74  197  -82  -95]
[ -18  -54   7  -34   37 -100   41   49]
[ -26  -78  10  -49   53 -144   59   71]
[  -8  -24   3  -15   16  -44   18   22]
sage: B.rank()
3
sage: B.nullity()
3
sage: (B.nrows(), B.ncols())
(6, 8)
sage: all([x in ZZ for x in B.list()])
True
sage: B_expanded=B.augment(identity_matrix(6)).rref()
sage: all([x in ZZ for x in B_expanded.list()])
True
sage: B_expanded # random
[ 1  3  0  0  1  1  3 -2  0  0 -3  0 -9 16]
[ 0  0  1  0  3 -2 -1 -3  0  0  2  0 11 -27]
[ 0  0  0  1 -1  2 -3 -1  0  0  2  0  7 -14]
[ 0  0  0  0  0  0  0  0  1  0 -5  0 -3  2]
[ 0  0  0  0  0  0  0  0  0  1  1  0  1 -3]
[ 0  0  0  0  0  0  0  0  0  0  0  1 -1  1]

```

Check that we fixed Trac #10543 (echelon forms should be immutable):

```

sage: B_expanded.is_immutable()
True

```

We want to modify `B_expanded`, so replace it with a copy:

```

sage: B_expanded = copy(B_expanded)
sage: B_expanded.subdivide(B.nrows()-B.nullity(), B.ncols()); B_expanded # random
[ 1  3  0  0  1  1  3 -2| 0  0 -3  0 -9 16]
[ 0  0  1  0  3 -2 -1 -3| 0  0  2  0 11 -27]
[ 0  0  0  1 -1  2 -3 -1| 0  0  2  0  7 -14]
[-----+-----]
[ 0  0  0  0  0  0  0  0| 1  0 -5  0 -3  2]
[ 0  0  0  0  0  0  0  0| 0  1  1  0  1 -3]
[ 0  0  0  0  0  0  0  0| 0  0  0  1 -1  1]
sage: C=B_expanded.subdivision(0,0)
sage: C # random
[ 1  3  0  0  1  1  3 -2]
[ 0  0  1  0  3 -2 -1 -3]
[ 0  0  0  1 -1  2 -3 -1]
sage: L=B_expanded.subdivision(1,1)
sage: L # random
[ 1  0 -5  0 -3  2]
[ 0  1  1  0  1 -3]
[ 0  0  0  1 -1  1]
sage: B.right_kernel()==C.right_kernel()
True
sage: B.row_space()==C.row_space()
True
sage: B.column_space()==L.right_kernel()
True
sage: B.left_kernel()==L.row_space()

```

True

A matrix to show that the null space of the L matrix is the column space of the starting matrix.

```
sage: A=random_matrix(QQ, 5, 7, algorithm='subspaces', rank=None); A # random
[-31 12 -9 -27 21 2 -15]
[105 -24 6 103 -30 -34 79]
[ 29 -9 5 26 -14 -5 17]
[233 -55 16 228 -71 -73 173]
[-42 10 -3 -41 13 13 -31]
sage: (A.nrows(), A.ncols())
(5, 7)
sage: all([x in ZZ for x in A.list()])
True
sage: A.nullity() # random
1
sage: A_expanded=A.augment(identity_matrix(5)).rref()
sage: A_expanded # random
[ 1 0 0 0 0 0 1 0 0 3 7 25 151]
[ 0 1 0 0 0 1 2 1 0 5 21 84 493]
[ 0 0 1 0 -1 2 0 0 2 13 53 308]
[ 0 0 0 1 0 -1 1 0 -2 -3 -9 -57]
[ 0 0 0 0 0 0 0 1 -3 1 1 -2]
sage: all([x in ZZ for x in A_expanded.list()])
True
sage: C=A_expanded.submatrix(0,0,A.nrows()-A.nullity(),A.ncols())
sage: L=A_expanded.submatrix(A.nrows()-A.nullity(),A.ncols())
sage: A.right_kernel()==C.right_kernel()
True
sage: A.row_space()==C.row_space()
True
sage: A.column_space()==L.right_kernel()
True
sage: A.left_kernel()==L.row_space()
True
```

TESTS:

The designated rank of the L matrix cannot be greater than the number of desired rows, nor can the rank be negative.

```
sage: random_matrix(QQ, 19, 20, algorithm='subspaces', rank=21)
Traceback (most recent call last):
...
ValueError: rank cannot exceed the number of rows or columns.
sage: random_matrix(QQ, 19, 20, algorithm='subspaces', rank=-1)
Traceback (most recent call last):
...
ValueError: matrices must have rank zero or greater.
```

REFERENCES:

AUTHOR:

Billy Wonderly (2010-07)

sage.matrix.constructor.**random_unimodular_matrix**(parent, upper_bound=None, max_tries=100)
 This function is available as random_unimodular_matrix(...) and matrix.random_unimodular(...).

Generate a random unimodular (determinant 1) matrix of a desired size over a desired ring.

INPUT:

- `parent` - A matrix space specifying the base ring, dimensions and representation (dense/sparse) for the result. The base ring must be exact.
- `upper_bound` - For large matrices over `QQ` or `ZZ`, `upper_bound` is the largest value matrix entries can achieve. But see the warning below.
- `max_tries` - If designated, number of tries used to generate each new random row; only matters when `upper_bound!=None`. Used to prevent endless looping. (default: 100)

A matrix not in reduced row-echelon form with the desired dimensions and properties.

OUTPUT:

An invertible matrix with the desired properties and determinant 1.

Warning: When `upper_bound` is set, it is possible for this constructor to fail with a `ValueError`. This may happen when the `upper_bound`, rank and/or matrix dimensions are all so small that it becomes infeasible or unlikely to create the requested matrix. If you *must* have this routine return successfully, do not set `upper_bound`.

Note: It is easiest to use this function via a call to the `random_matrix()` function with the `algorithm='unimodular'` keyword. We provide one example accessing this function directly, while the remainder will use this more general function.

EXAMPLES:

A matrix size 5 over `QQ`.

```
sage: from sage.matrix.constructor import random_unimodular_matrix
sage: matrix_space = sage.matrix.matrix_space.MatrixSpace(QQ, 5)
sage: A=random_unimodular_matrix(matrix_space); A # random
[ -8   31   85  148 -419]
[  2   -9  -25  -45  127]
[ -1   10   30   65 -176]
[ -3   12   33   58 -164]
[  5  -21  -59 -109  304]
sage: det(A)
1
```

A matrix size 6 with entries no larger than 50.

```
sage: B=random_matrix(ZZ, 7, algorithm='unimodular', upper_bound=50);B # random
[ -1   0   3   1  -2   2   9]
[  1   2  -5   0  14  19 -49]
[ -3  -2  12   5  -6  -4  24]
[  1   2  -9  -3   3   4  -7]
[ -2  -1   7   2  -8  -5  31]
[  2   2  -6  -3   8  16 -32]
[  1   2  -9  -2   5   6 -12]
sage: det(B)
1
```

A matrix over the number Field in y with defining polynomial $y^2 - 2y - 2$.

```
sage: y = var('y')
sage: K=NumberField(y^2-2*y-2,'y')
```

```

sage: C=random_matrix(K, 3, algorithm='unimodular');C # random
[ 2*y - 33 681*y - 787 31*y - 37]
[      y + 6 -155*y + 83  -7*y + 4]
[      -y  24*y + 51      y + 3]
sage: det(C)
1

```

TESTS:

Unimodular matrices are square.

```

sage: random_matrix(QQ, 5, 6, algorithm='unimodular')
Traceback (most recent call last):
...
TypeError: a unimodular matrix must be square.

```

Only matrices over ZZ and QQ can have size control.

```

sage: F.<a>=GF(5^7)
sage: random_matrix(F, 5, algorithm='unimodular', upper_bound=20)
Traceback (most recent call last):
...
TypeError: only matrices over ZZ or QQ can have size control.

```

AUTHOR:

Billy Wonderly (2010-07)

sage.matrix.constructor.**vector_on_axis_rotation_matrix**(*v*, *i*, *ring=None*)

This function is available as `vector_on_axis_rotation_matrix(...)` and `matrix.vector_on_axis_rotation(...)`.

Return a rotation matrix M such that $\det(M) = 1$ sending the vector v on the i -th axis so that all other coordinates of Mv are zero.

Note: Such a matrix is not uniquely determined. This function returns one such matrix.

INPUT:

- *v* - vector
- *i* - integer
- *ring* - ring (optional, default: None) of the resulting matrix

OUTPUT:

A matrix

EXAMPLES:

```

sage: from sage.matrix.constructor import vector_on_axis_rotation_matrix
sage: v = vector((1,2,3))
sage: vector_on_axis_rotation_matrix(v, 2) * v
(0, 0, sqrt(14))
sage: vector_on_axis_rotation_matrix(v, 1) * v
(0, sqrt(14), 0)
sage: vector_on_axis_rotation_matrix(v, 0) * v
(sqrt(14), 0, 0)

sage: x,y = var('x,y')
sage: v = vector((x,y))

```

```
sage: vector_on_axis_rotation_matrix(v, 1)
[ y/sqrt(x^2 + y^2) -x/sqrt(x^2 + y^2)]
[ x/sqrt(x^2 + y^2)  y/sqrt(x^2 + y^2)]
sage: vector_on_axis_rotation_matrix(v, 0)
[ x/sqrt(x^2 + y^2)  y/sqrt(x^2 + y^2)]
[-y/sqrt(x^2 + y^2)  x/sqrt(x^2 + y^2)]
sage: vector_on_axis_rotation_matrix(v, 0) * v
(x^2/sqrt(x^2 + y^2) + y^2/sqrt(x^2 + y^2), 0)
sage: vector_on_axis_rotation_matrix(v, 1) * v
(0, x^2/sqrt(x^2 + y^2) + y^2/sqrt(x^2 + y^2))

sage: v = vector((1,2,3,4))
sage: vector_on_axis_rotation_matrix(v, 0) * v
(sqrt(30), 0, 0, 0)
sage: vector_on_axis_rotation_matrix(v, 0, ring=RealField(10))
[ 0.18  0.37  0.55  0.73]
[-0.98 0.068  0.10  0.14]
[ 0.00 -0.93  0.22  0.30]
[ 0.00  0.00 -0.80  0.60]
sage: vector_on_axis_rotation_matrix(v, 0, ring=RealField(10)) * v
(5.5, 0.00098, 0.00098, 0.00)
```

AUTHORS:

Sebastien Labbe (April 2010)

`sage.matrix.constructor.zero_matrix(ring, nrows, ncols=None, sparse=False)`

This function is available as `zero_matrix(...)` and `matrix.zero(...)`.

Return the $nrows \times ncols$ zero matrix over the given ring.

The default ring is the integers.

EXAMPLES:

```
sage: M = zero_matrix(QQ, 2); M
[0 0]
[0 0]
sage: M.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: M = zero_matrix(2, 3); M
[0 0 0]
[0 0 0]
sage: M.parent()
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
sage: M.is_mutable()
True
sage: M = zero_matrix(3, 1, sparse=True); M
[0]
[0]
[0]
sage: M.parent()
Full MatrixSpace of 3 by 1 sparse matrices over Integer Ring
sage: M.is_mutable()
True
```

MATRICES OVER AN ARBITRARY RING

AUTHORS:

- William Stein
- Martin Albrecht: conversion to Pyrex
- Jaap Spies: various functions
- Gary Zablackis: fixed a sign bug in generic determinant.
- William Stein and Robert Bradshaw - complete restructuring.
- Rob Beezer - refactor kernel functions.

Elements of matrix spaces are of class `Matrix` (or a class derived from `Matrix`). They can be either sparse or dense, and can be defined over any base ring.

EXAMPLES:

We create the 2×3 matrix

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{pmatrix}$$

as an element of a matrix space over \mathbb{Q} :

```
sage: M = MatrixSpace(QQ, 2, 3)
sage: A = M([1, 2, 3, 4, 5, 6]); A
[1 2 3]
[4 5 6]
sage: A.parent()
Full MatrixSpace of 2 by 3 dense matrices over Rational Field
```

Alternatively, we could create `A` more directly as follows (which would completely avoid having to create the matrix space):

```
sage: A = matrix(QQ, 2, [1, 2, 3, 4, 5, 6]); A
[1 2 3]
[4 5 6]
```

We next change the top-right entry of `A`. Note that matrix indexing is 0-based in Sage, so the top right entry is $(0, 2)$, which should be thought of as “row number 0, column number 2”.

```
sage: A[0, 2] = 389
sage: A
[ 1  2 389]
[ 4  5  6]
```

Also notice how matrices print. All columns have the same width and entries in a given column are right justified. Next we compute the reduced row echelon form of A .

```
sage: A.rref()
[      1      0 -1933/3]
[      0      1 1550/3]
```

3.1 Indexing

Sage has quite flexible ways of extracting elements or submatrices from a matrix:

```
sage: m=[(1, -2, -1, -1,9), (1, 8, 6, 2,2), (1, 1, -1, 1,4), (-1, 2, -2, -1,4)];M= matrix(m)
sage: M
[ 1 -2 -1 -1  9]
[ 1  8  6  2  2]
[ 1  1 -1  1  4]
[-1  2 -2 -1  4]
```

Get the 2 x 2 submatrix of M , starting at row index and column index 1:

```
sage: M[1:3,1:3]
[ 8  6]
[ 1 -1]
```

Get the 2 x 3 submatrix of M starting at row index and column index 1:

```
sage: M[1:3, [1..3]]
[ 8  6  2]
[ 1 -1  1]
```

Get the second column of M :

```
sage: M[:,1]
[-2]
[ 8]
[ 1]
[ 2]
```

Get the first row of M :

```
sage: M[0,:]
[ 1 -2 -1 -1  9]
```

Get the last row of M (negative numbers count from the end):

```
sage: M[-1,:]
[-1  2 -2 -1  4]
```

More examples:

```
sage: M[range(2),:]
[ 1 -2 -1 -1  9]
[ 1  8  6  2  2]
sage: M[range(2),4]
[9]
[2]
```



```
sage: M[range(3), range(5)]
[ 1 -2 -1 -1  9]
[ 1  8  6  2  2]
[ 1  1 -1  1  4]
```

```
sage: M[3, range(5)]
[-1  2 -2 -1  4]
```

```
sage: M[3, :]
[-1  2 -2 -1  4]
```

```
sage: M[3, 4]
4
```

```
sage: M[-1, :]
[-1  2 -2 -1  4]
```

```
sage: A = matrix(ZZ, 3, 4, [3, 2, -5, 0, 1, -1, 1, -4, 1, 0, 1, -3]); A
[ 3  2 -5  0]
[ 1 -1  1 -4]
[ 1  0  1 -3]
```

A series of three numbers, separated by colons, like $n:m:s$, means numbers from n up to (but not including) m , in steps of s . So $0:5:2$ means the sequence $[0, 2, 4]$:

```
sage: A[:, 0:4:2]
[ 3 -5]
[ 1  1]
[ 1  1]
```

```
sage: A[1:, 0:4:2]
[1 1]
[1 1]
```

```
sage: A[2::-1, :]
[ 1  0  1 -3]
[ 1 -1  1 -4]
[ 3  2 -5  0]
```

```
sage: A[1:, 3::-1]
[-4  1 -1  1]
[-3  1  0  1]
```

```
sage: A[1:, 3::-2]
[-4 -1]
[-3  0]
```

```
sage: A[2::-1, 3:1:-1]
[-3  1]
[-4  1]
[ 0 -5]
```

We can also change submatrices using these indexing features:

```
sage: M=matrix([(1, -2, -1, -1, 9), (1, 8, 6, 2, 2), (1, 1, -1, 1, 4), (-1, 2, -2, -1, 4)]); M
[ 1 -2 -1 -1  9]
[ 1  8  6  2  2]
[ 1  1 -1  1  4]
[-1  2 -2 -1  4]
```

Set the 2 x 2 submatrix of M, starting at row index and column index 1:

```
sage: M[1:3, 1:3] = [[1, 0], [0, 1]]; M
[ 1 -2 -1 -1  9]
[ 1  1  0  2  2]
[ 1  0  1  1  4]
[-1  2 -2 -1  4]
```

Set the 2 x 3 submatrix of M starting at row index and column index 1:

```
sage: M[1:3, [1..3]] = M[2:4, 0:3]; M
[ 1 -2 -1 -1  9]
[ 1  1  0  1  2]
[ 1 -1  2 -2  4]
[-1  2 -2 -1  4]
```

Set part of the first column of M:

```
sage: M[1:, 0] = [[2], [3], [4]]; M
[ 1 -2 -1 -1  9]
[ 2  1  0  1  2]
[ 3 -1  2 -2  4]
[ 4  2 -2 -1  4]
```

Or do a similar thing with a vector:

```
sage: M[1:, 0] = vector([-2, -3, -4]); M
[ 1 -2 -1 -1  9]
[-2  1  0  1  2]
[-3 -1  2 -2  4]
[-4  2 -2 -1  4]
```

Or a constant:

```
sage: M[1:, 0] = 30; M
[ 1 -2 -1 -1  9]
[30  1  0  1  2]
[30 -1  2 -2  4]
[30  2 -2 -1  4]
```

Set the first row of M:

```
sage: M[0, :] = [20, 21, 22, 23, 24]; M
[20 21 22 23 24]
[30  1  0  1  2]
[30 -1  2 -2  4]
[30  2 -2 -1  4]
sage: M[0, :] = vector([0, 1, 2, 3, 4]); M
[ 0  1  2  3  4]
[30  1  0  1  2]
[30 -1  2 -2  4]
[30  2 -2 -1  4]
sage: M[0, :] = -3; M
[-3 -3 -3 -3 -3]
[30  1  0  1  2]
[30 -1  2 -2  4]
[30  2 -2 -1  4]
```

```
sage: A = matrix(ZZ, 3, 4, [3, 2, -5, 0, 1, -1, 1, -4, 1, 0, 1, -3]); A
[ 3  2 -5  0]
[ 1 -1  1 -4]
[ 1  0  1 -3]
```

We can use the step feature of slices to set every other column:

```
sage: A[:, 0:3:2] = 5; A
[ 5  2  5  0]
[ 5 -1  5 -4]
[ 5  0  5 -3]
```

```
sage: A[1:, 0:4:2] = [[100, 200], [300, 400]]; A
[  5   2   5   0]
[100 -1 200 -4]
[300  0 400 -3]
```

We can also count backwards to flip the matrix upside down:

```
sage: A[::-1, :] = A; A
[300  0 400 -3]
[100 -1 200 -4]
[  5   2   5   0]
```

```
sage: A[1:, 3::-1] = [[2, 3, 0, 1], [9, 8, 7, 6]]; A
[300  0 400 -3]
[  1   0   3   2]
[  6   7   8   9]
```

```
sage: A[1:, ::-2] = A[1:, ::2]; A
[300  0 400 -3]
[  1   3   3   1]
[  6   8   8   6]
```

```
sage: A[:, -1, 3:1:-1] = [[4, 3], [1, 2], [-1, -2]]; A
[300  0 -2 -1]
[  1   3   2   1]
[  6   8   3   4]
```

We save and load a matrix:

```
sage: A = matrix(Integers(8), 3, range(9))
sage: loads(dumps(A)) == A
True
```

MUTABILITY: Matrices are either immutable or not. When initially created, matrices are typically mutable, so one can change their entries. Once a matrix A is made immutable using `A.set_immutable()` the entries of A cannot be changed, and A can never be made mutable again. However, properties of A such as its rank, characteristic polynomial, etc., are all cached so computations involving A may be more efficient. Once A is made immutable it cannot be changed back. However, one can obtain a mutable copy of A using `copy(A)`.

EXAMPLES:

```
sage: A = matrix(RR, 2, [1, 10, 3.5, 2])
sage: A.set_immutable()
sage: copy(A) is A
False
```

The echelon form method always returns immutable matrices with known rank.

EXAMPLES:

```
sage: A = matrix(Integers(8), 3, range(9))
sage: A.determinant()
0
sage: A[0,0] = 5
sage: A.determinant()
1
sage: A.set_immutable()
sage: A[0,0] = 5
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead (i.e., use copy(M) to change a copy of
```

3.1.1 Implementation and Design

Class Diagram (an x means that class is currently supported):

```
x Matrix
x   Matrix_sparse
x   Matrix_generic_sparse
x   Matrix_integer_sparse
x   Matrix_rational_sparse
x   Matrix_cyclo_sparse
x   Matrix_modn_sparse
x   Matrix_RR_sparse
x   Matrix_CC_sparse
x   Matrix_RDF_sparse
x   Matrix_CDF_sparse

x Matrix_dense
x   Matrix_generic_dense
x   Matrix_integer_dense
x   Matrix_integer_2x2_dense
x   Matrix_rational_dense
x   Matrix_cyclo_dense    -- idea: restrict scalars to QQ, compute charpoly there, then factor
x   Matrix_modn_dense
x   Matrix_RR_dense
x   Matrix_CC_dense
x   Matrix_real_double_dense
x   Matrix_complex_double_dense
```

The corresponding files in the sage/matrix library code directory are named

[matrix] [base ring] [dense or sparse].

See the files `matrix_template.pxd` and `matrix_template.pyx`.

New matrices types can only be implemented in Cython.

```
***** LEVEL 1 *****
NON-OPTIONAL
For each base field it is absolutely essential to completely
implement the following functionality for that base ring:
```

```

* __cinit__      -- should use sage_malloc from ext/stdsage.pxi (only
                  needed if allocate memory)
* __init__       -- this signature: 'def __init__(self, parent, entries, copy, coerce)'
* __dealloc__    -- use sage_free (only needed if allocate memory)
* set_unsafe(self, size_t i, size_t j, x) -- doesn't do bounds or any other checks; assumes x is :
* get_unsafe(self, size_t i, size_t j) -- doesn't do checks
* __richcmp__    -- always the same (I don't know why its needed -- bug in PYREX).

```

Note that the `__init__` function must construct the all zero matrix if `entries == None`.

***** LEVEL 2 *****

IMPORTANT (and *highly* recommended):

After getting the special class with all level 1 functionality to work, implement all of the following (they should not change functionality, except speed (always faster!) in any way):

```

* def _pickle(self):
    return data, version
* def _unpickle(self, data, int version)
    reconstruct matrix from given data and version; may assume _parent, _nrows, and _ncols are
    Use version numbers >= 0 so if you change the pickle strategy then
    old objects still unpickle.
* cdef _list -- list of underlying elements (need not be a copy)
* cdef _dict -- sparse dictionary of underlying elements
* cdef _add -- add two matrices with identical parents
* _matrix_times_matrix_c_impl -- multiply two matrices with compatible dimensions and
    identical base rings (both sparse or both dense)
* cdef _cmp_c_impl -- compare two matrices with identical parents
* cdef _lmul_c_impl -- multiply this matrix on the right by a scalar, i.e., self * scalar
* cdef _rmul_c_impl -- multiply this matrix on the left by a scalar, i.e., scalar * self
* __copy__
* __neg__

```

The list and dict returned by `_list` and `_dict` will *not* be changed by any internal algorithms and are not accessible to the user.

***** LEVEL 3 *****

OPTIONAL:

```

* cdef _sub_
* __invert__
* _multiply_classical
* __deepcopy__

```

Further special support:

- * Matrix windows -- to support Strassen multiplication for a given base ring.
- * Other functions, e.g., transpose, for which knowing the specific representation can be helpful.

.. note::

- For caching, use `self.fetch` and `self.cache`.
- Any method that can change the matrix should call `__check_mutability()` first. There are also many fast cdef'd bounds checking methods.

- Kernels of matrices
Implement only a `left_kernel()` or `right_kernel()` method, whichever requires the least overhead (usually meaning little or no transposing). Let the methods in the `matrix2` class handle left, right, generic kernel distinctions.

ABSTRACT BASE CLASS FOR MATRICES

Abstract base class for matrices

For design documentation see `matrix/docs.py`.

```
class sage.matrix.matrix.Matrix
    Bases: sage.matrix.matrix2.Matrix
```

The initialization routine of the `Matrix` base class ensures that it sets the attributes `self._parent`, `self._base_ring`, `self._nrows`, `self._ncols`. It sets the latter ones by accessing the relevant information on `parent`, which is often slower than what a more specific subclass can do.

Subclasses of `Matrix` can safely skip calling `Matrix.__init__` provided they take care of initializing these attributes themselves.

The private attributes `self._is_immutable` and `self._cache` are implicitly initialized to valid values upon memory allocation.

EXAMPLES:

```
sage: import sage.matrix.matrix0
sage: A = sage.matrix.matrix0.Matrix(MatrixSpace(QQ, 2))
sage: type(A)
<type 'sage.matrix.matrix0.Matrix'>
```

```
sage.matrix.matrix.is_Matrix(x)
```

EXAMPLES:

```
sage: from sage.matrix.matrix import is_Matrix
sage: is_Matrix(0)
False
sage: is_Matrix(matrix([[1, 2], [3, 4]]))
True
```


BASE CLASS FOR MATRICES, PART 0

Base class for matrices, part 0

Note: For design documentation see `matrix/docs.py`.

EXAMPLES:

```
sage: matrix(2, [1, 2, 3, 4])
[1 2]
[3 4]
```

```
class sage.matrix.matrix0.Matrix
    Bases: sage.structure.element.Matrix
```

A generic matrix.

The `Matrix` class is the base class for all matrix classes. To create a `Matrix`, first create a `MatrixSpace`, then coerce a list of elements into the `MatrixSpace`. See the documentation of `MatrixSpace` for more details.

EXAMPLES:

We illustrate matrices and matrix spaces. Note that no actual matrix that you make should have class `Matrix`; the class should always be derived from `Matrix`.

```
sage: M = MatrixSpace(CDF, 2, 3); M
Full MatrixSpace of 2 by 3 dense matrices over Complex Double Field
sage: a = M([1, 2, 3, 4, 5, 6]); a
[1.0 2.0 3.0]
[4.0 5.0 6.0]
sage: type(a)
<type 'sage.matrix.matrix_complex_double_dense.Matrix_complex_double_dense'>
sage: parent(a)
Full MatrixSpace of 2 by 3 dense matrices over Complex Double Field

sage: matrix(CDF, 2, 3, [1, 2, 3, 4, 5, 6])
[1.0 2.0 3.0]
[4.0 5.0 6.0]
sage: Mat(CDF, 2, 3)(range(1, 7))
[1.0 2.0 3.0]
[4.0 5.0 6.0]

sage: Q.<i, j, k> = QuaternionAlgebra(QQ, -1, -1)
sage: matrix(Q, 2, 1, [1, 2])
```

```
[1]
[2]
```

act_on_polynomial(*f*)

Returns the polynomial $f(\text{self} \cdot x)$.

INPUT:

- *self* - an $n \times n$ matrix
- *f* - a polynomial in n variables $x=(x_1, \dots, x_n)$

OUTPUT: The polynomial $f(\text{self} \cdot x)$.

EXAMPLES:

```
sage: R.<x,y> = QQ[]
sage: x, y = R.gens()
sage: f = x**2 - y**2
sage: M = MatrixSpace(QQ, 2)
sage: A = M([1,2,3,4])
sage: A.act_on_polynomial(f)
-8*x^2 - 20*x*y - 12*y^2
```

add_multiple_of_column(*i, j, s, start_row=0*)

Add s times column j to column i .

EXAMPLES: We add -1 times the third column to the second column of an integer matrix, remembering to start numbering cols at zero:

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.add_multiple_of_column(1, 2, -1)
sage: a
[ 0 -1  2]
[ 3 -1  5]
```

To add a rational multiple, we first need to change the base ring:

```
sage: a = a.change_ring(QQ)
sage: a.add_multiple_of_column(1, 0, 1/3)
sage: a
[ 0 -1  2]
[ 3  0  5]
```

If not, we get an error message:

```
sage: a.add_multiple_of_column(1, 0, i)
Traceback (most recent call last):
```

```
...
```

```
TypeError: Multiplying column by Symbolic Ring element cannot be done over Rational Field, u
```

add_multiple_of_row(*i, j, s, start_col=0*)

Add s times row j to row i .

EXAMPLES: We add -3 times the first row to the second row of an integer matrix, remembering to start numbering rows at zero:

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
```

```
sage: a.add_multiple_of_row(1,0,-3)
sage: a
[ 0  1  2]
[ 3  1 -1]
```

To add a rational multiple, we first need to change the base ring:

```
sage: a = a.change_ring(QQ)
sage: a.add_multiple_of_row(1,0,1/3)
sage: a
[ 0  1  2]
[ 3  4/3 -1/3]
```

If not, we get an error message:

```
sage: a.add_multiple_of_row(1,0,i)
Traceback (most recent call last):
...
TypeError: Multiplying row by Symbolic Ring element cannot be done over Rational Field, use
```

anticommutator (*other*)

Return the anticommutator self and other.

The *anticommutator* of two $n \times n$ matrices A and B is defined as $\{A, B\} := AB + BA$ (sometimes this is written as $[A, B]_+$).

EXAMPLES:

```
sage: A = Matrix(ZZ, 2, 2, range(4))
sage: B = Matrix(ZZ, 2, 2, [0, 1, 0, 0])
sage: A.anticommutator(B)
[2 3]
[0 2]
sage: A.anticommutator(B) == B.anticommutator(A)
True
sage: A.commutator(B) + B.anticommutator(A) == 2*A*B
True
```

base_ring ()

Returns the base ring of the matrix.

EXAMPLES:

```
sage: m=matrix(QQ, 2, [1,2,3,4])
sage: m.base_ring()
Rational Field
```

change_ring (*ring*)

Return the matrix obtained by coercing the entries of this matrix into the given ring.

Always returns a copy (unless self is immutable, in which case returns self).

EXAMPLES:

```
sage: A = Matrix(QQ, 2, 2, [1/2, 1/3, 1/3, 1/4])
sage: A.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: A.change_ring(GF(25, 'a'))
[3 2]
[2 4]
sage: A.change_ring(GF(25, 'a')).parent()
```

```
Full MatrixSpace of 2 by 2 dense matrices over Finite Field in a of size 5^2
sage: A.change_ring(ZZ)
Traceback (most recent call last):
...
TypeError: matrix has denominators so can't change to ZZ.
```

Changing rings preserves subdivisions:

```
sage: A.subdivide([1], []); A
[1/2 1/3]
[-----]
[1/3 1/4]
sage: A.change_ring(GF(25, 'a'))
[3 2]
[---]
[2 4]
```

commutator(*other*)

Return the commutator $\text{self} * \text{other} - \text{other} * \text{self}$.

EXAMPLES:

```
sage: A = Matrix(ZZ, 2, 2, range(4))
sage: B = Matrix(ZZ, 2, 2, [0, 1, 0, 0])
sage: A.commutator(B)
[-2 -3]
[ 0  2]
sage: A.commutator(B) == -B.commutator(A)
True
```

dict()

Dictionary of the elements of self with keys pairs (i,j) and values the nonzero entries of self.

It is safe to change the returned dictionary.

EXAMPLES:

```
sage: R.<x,y> = QQ[]
sage: a = matrix(R, 2, [x,y,0, 0,0,2*x+y]); a
[      x      y      0]
[      0      0 2*x + y]
sage: d = a.dict(); d
{(0, 1): y, (1, 2): 2*x + y, (0, 0): x}
```

Notice that changing the returned list does not change a (the list is a copy):

```
sage: d[0,0] = 25
sage: a
[      x      y      0]
[      0      0 2*x + y]
```

dimensions()

Returns the dimensions of this matrix as the tuple (nrows, ncols).

EXAMPLES:

```
sage: M = matrix([[1,2,3],[4,5,6]])
sage: N = M.transpose()
sage: M.dimensions()
(2, 3)
```

```
sage: N.dimensions()
(3, 2)
```

AUTHORS:

•Benjamin Lundell (2012-02-09): examples

is_alternating()

Return True if self is an alternating matrix.

Here, “alternating matrix” means a square matrix A satisfying $A^T = -A$ and such that the diagonal entries of A are 0. Notice that the condition that the diagonal entries be 0 is not redundant for matrices over arbitrary ground rings (but it is redundant when 2 is invertible in the ground ring). A square matrix A only required to satisfy $A^T = -A$ is said to be “skew-symmetric”, and this property is checked by the `is_skew_symmetric()` method.

EXAMPLES:

```
sage: m = matrix(QQ, [[0,2], [-2,0]])
sage: m.is_alternating()
True
sage: m = matrix(QQ, [[1,2], [2,1]])
sage: m.is_alternating()
False
```

In contrast to the property of being skew-symmetric, the property of being alternating does not tolerate nonzero entries on the diagonal even if they are their own negatives:

```
sage: n = matrix(Zmod(4), [[0, 1], [-1, 2]])
sage: n.is_alternating()
False
```

is_dense()

Returns True if this is a dense matrix.

In Sage, being dense is a property of the underlying representation, not the number of nonzero entries.

EXAMPLES:

```
sage: matrix(QQ, 2, 2, range(4)).is_dense()
True
sage: matrix(QQ, 2, 2, range(4), sparse=True).is_dense()
False
```

is_hermitian()

Returns True if the matrix is equal to its conjugate-transpose.

OUTPUT:

True if the matrix is square and equal to the transpose with every entry conjugated, and False otherwise.

Note that if conjugation has no effect on elements of the base ring (such as for integers), then the `is_symmetric()` method is equivalent and faster.

This routine is for matrices over exact rings and so may not work properly for matrices over `RR` or `CC`. For matrices with approximate entries, the rings of double-precision floating-point numbers, `RDF` and `CDF`, are a better choice since the `sage.matrix.matrix_double_dense.Matrix_double_dense.is_hermitian()` method has a tolerance parameter. This provides control over allowing for minor discrepancies between entries when checking equality.

The result is cached.

EXAMPLES:

```
sage: A = matrix(QQbar, [[ 1 + I, 1 - 6*I, -1 - I],
...                      [-3 - I, -4*I, -2],
...                      [-1 + I, -2 - 8*I, 2 + I]])
sage: A.is_hermitian()
False
sage: B = A*A.conjugate_transpose()
sage: B.is_hermitian()
True
```

Sage has several fields besides the entire complex numbers where conjugation is non-trivial.

```
sage: F.<b> = QuadraticField(-7)
sage: C = matrix(F, [[-2*b - 3, 7*b - 6, -b + 3],
...                  [-2*b - 3, -3*b + 2, -2*b],
...                  [ b + 1, 0, -2]])
sage: C.is_hermitian()
False
sage: C = C*C.conjugate_transpose()
sage: C.is_hermitian()
True
```

A matrix that is nearly Hermitian, but for a non-real diagonal entry.

```
sage: A = matrix(QQbar, [[ 2, 2-I, 1+4*I],
...                      [ 2+I, 3+I, 2-6*I],
...                      [1-4*I, 2+6*I, 5]])
sage: A.is_hermitian()
False
sage: A[1,1] = 132
sage: A.is_hermitian()
True
```

Rectangular matrices are never Hermitian.

```
sage: A = matrix(QQbar, 3, 4)
sage: A.is_hermitian()
False
```

A square, empty matrix is trivially Hermitian.

```
sage: A = matrix(QQ, 0, 0)
sage: A.is_hermitian()
True
```

is_immutable()

Return True if this matrix is immutable.

See the documentation for `self.set_immutable` for more details about mutability.

EXAMPLES:

```
sage: A = Matrix(QQ['t','s'], 2, 2, range(4))
sage: A.is_immutable()
False
sage: A.set_immutable()
sage: A.is_immutable()
True
```

is_invertible()

Return True if this matrix is invertible.

EXAMPLES: The following matrix is invertible over \mathbb{Q} but not over \mathbb{Z} .

```
sage: A = MatrixSpace(ZZ, 2)(range(4))
sage: A.is_invertible()
False
sage: A.matrix_over_field().is_invertible()
True
```

The inverse function is a constructor for matrices over the fraction field, so it can work even if A is not invertible.

```
sage: ~A      # inverse of A
[-3/2  1/2]
[  1   0]
```

The next matrix is invertible over \mathbb{Z} .

```
sage: A = MatrixSpace(IntegerRing(), 2)([1, 10, 0, -1])
sage: A.is_invertible()
True
sage: ~A      # compute the inverse
[ 1 10]
[ 0 -1]
```

The following nontrivial matrix is invertible over $\mathbb{Z}[x]$.

```
sage: R.<x> = PolynomialRing(IntegerRing())
sage: A = MatrixSpace(R, 2)([1, x, 0, -1])
sage: A.is_invertible()
True
sage: ~A
[ 1  x]
[ 0 -1]
```

is_mutable()

Return True if this matrix is mutable.

See the documentation for `self.set_immutable` for more details about mutability.

EXAMPLES:

```
sage: A = Matrix(QQ['t', 's'], 2, 2, range(4))
sage: A.is_mutable()
True
sage: A.set_immutable()
sage: A.is_mutable()
False
```

is_singular()

Returns True if self is singular.

OUTPUT:

A square matrix is singular if it has a zero determinant and this method will return True in exactly this case. When the entries of the matrix come from a field, this is equivalent to having a nontrivial kernel, or lacking an inverse, or having linearly dependent rows, or having linearly dependent columns.

For square matrices over a field the methods `is_invertible()` and `is_singular()` are logical opposites. However, it is an error to apply `is_singular()` to a matrix that is not square, while `is_invertible()` will always return False for a matrix that is not square.

EXAMPLES:

A singular matrix over the field $\mathbb{Q}\mathbb{Q}$.

```
sage: A = matrix(QQ, 4, [-1, 2, -3, 6, 0, -1, -1, 0, -1, 1, -5, 7, -1, 6, 5, 2])
sage: A.is_singular()
True
sage: A.right_kernel().dimension()
1
```

A matrix that is not singular, i.e. nonsingular, over a field.

```
sage: B = matrix(QQ, 4, [1, -3, -1, -5, 2, -5, -2, -7, -2, 5, 3, 4, -1, 4, 2, 6])
sage: B.is_singular()
False
sage: B.left_kernel().dimension()
0
```

For *rectangular* matrices, invertibility is always False, but asking about singularity will give an error.

```
sage: C = matrix(QQ, 5, range(30))
sage: C.is_invertible()
False
sage: C.is_singular()
Traceback (most recent call last):
...
ValueError: self must be a square matrix
```

When the base ring is not a field, then a matrix may be both not invertible and not singular.

```
sage: D = matrix(ZZ, 4, [2, 0, -4, 8, 2, 1, -2, 7, 2, 5, 7, 0, 0, 1, 4, -6])
sage: D.is_invertible()
False
sage: D.is_singular()
False
sage: d = D.determinant(); d
2
sage: d.is_unit()
False
```

is_skew_symmetric()

Return True if self is a skew-symmetric matrix.

Here, “skew-symmetric matrix” means a square matrix A satisfying $A^T = -A$. It does not require that the diagonal entries of A are 0 (although this automatically follows from $A^T = -A$ when 2 is invertible in the ground ring over which the matrix is considered). Skew-symmetric matrices A whose diagonal entries are 0 are said to be “alternating”, and this property is checked by the `is_alternating()` method.

EXAMPLES:

```
sage: m = matrix(QQ, [[0, 2], [-2, 0]])
sage: m.is_skew_symmetric()
True
sage: m = matrix(QQ, [[1, 2], [2, 1]])
sage: m.is_skew_symmetric()
False
```

Skew-symmetric is not the same as alternating when 2 is a zero-divisor in the ground ring:

```
sage: n = matrix(Zmod(4), [[0, 1], [-1, 2]])
sage: n.is_skew_symmetric()
True
```

but yet the diagonal cannot be completely arbitrary in this case:


```
sage: n = matrix(Zmod(4), [[0, 1], [-1, 3]])
sage: n.is_skew_symmetric()
False
```

is_skew_symmetrizable (*return_diag=False, positive=True*)

This function takes a square matrix over an *ordered integral domain* and checks if it is skew-symmetrizable. A matrix B is skew-symmetrizable iff there exists an invertible diagonal matrix D such that DB is skew-symmetric.

Warning: Expects `self` to be a matrix over an *ordered integral domain*.

INPUT:

- `return_diag` – bool (default: False) if True and `self` is skew-symmetrizable the diagonal entries of the matrix D are returned.
- `positive` – bool (default: True) if True, the condition that D has positive entries is added.

OUTPUT:

- True – if `self` is skew-symmetrizable and `return_diag` is False
- the diagonal entries of a matrix D such that DB is skew-symmetric – iff `self` is skew-symmetrizable and `return_diag` is True
- False – iff `self` is not skew-symmetrizable

EXAMPLES:

```
sage: matrix([[0, 6], [3, 0]]).is_skew_symmetrizable(positive=False)
True
sage: matrix([[0, 6], [3, 0]]).is_skew_symmetrizable(positive=True)
False
```

```
sage: M = matrix(4, [0, 1, 0, 0, -1, 0, -1, 0, 0, 2, 0, 1, 0, 0, -1, 0]); M
[ 0  1  0  0]
[-1  0 -1  0]
[ 0  2  0  1]
[ 0  0 -1  0]
```

```
sage: M.is_skew_symmetrizable(return_diag=True)
[1, 1, 1/2, 1/2]
```

```
sage: M2 = diagonal_matrix([1, 1, 1/2, 1/2])*M; M2
[ 0  1  0  0]
[-1  0 -1  0]
[ 0  1  0  1/2]
[ 0  0 -1/2  0]
```

```
sage: M2.is_skew_symmetric()
True
```

REFERENCES:

- [FZ2001] S. Fomin, A. Zelevinsky. Cluster Algebras 1: Foundations, arXiv:math/0104151 (2001).

is_sparse ()

Return True if this is a sparse matrix.

In Sage, being sparse is a property of the underlying representation, not the number of nonzero entries.

EXAMPLES:

```
sage: matrix(QQ, 2, 2, range(4)).is_sparse()
False
sage: matrix(QQ, 2, 2, range(4), sparse=True).is_sparse()
True
```

is_square()

Return True precisely if this matrix is square, i.e., has the same number of rows and columns.

EXAMPLES:

```
sage: matrix(QQ, 2, 2, range(4)).is_square()
True
sage: matrix(QQ, 2, 3, range(6)).is_square()
False
```

is_symmetric()

Returns True if this is a symmetric matrix.

EXAMPLES:

```
sage: m=Matrix(QQ, 2, range(0, 4))
sage: m.is_symmetric()
False

sage: m=Matrix(QQ, 2, (1, 1, 1, 1, 1, 1))
sage: m.is_symmetric()
False

sage: m=Matrix(QQ, 1, (2,))
sage: m.is_symmetric()
True
```

is_symmetrizable (*return_diag=False, positive=True*)

This function takes a square matrix over an *ordered integral domain* and checks if it is symmetrizable. A matrix B is symmetrizable iff there exists an invertible diagonal matrix D such that DB is symmetric.

Warning: Expects `self` to be a matrix over an *ordered integral domain*.

INPUT:

- `return_diag` – bool(default:False) if True and `self` is symmetrizable the diagonal entries of the matrix D are returned.
- `positive` – bool(default:True) if True, the condition that D has positive entries is added.

OUTPUT:

- True – if `self` is symmetrizable and `return_diag` is False
- the diagonal entries of a matrix D such that DB is symmetric – iff `self` is symmetrizable and `return_diag` is True
- False – iff `self` is not symmetrizable

EXAMPLES:

```
sage: matrix([[0, 6], [3, 0]]).is_symmetrizable(positive=False)
True

sage: matrix([[0, 6], [3, 0]]).is_symmetrizable(positive=True)
```

```
True
```

```
sage: matrix([[0,6],[0,0]]).is_symmetrizable(return_diag=True)
False
```

```
sage: matrix([2]).is_symmetrizable(positive=True)
True
```

```
sage: matrix([[1,2],[3,4]]).is_symmetrizable(return_diag=true)
[1, 2/3]
```

REFERENCES:

- [FZ2001] S. Fomin, A. Zelevinsky. Cluster Algebras 1: Foundations, arXiv:math/0104151 (2001).

iterates (*v*, *n*, *rows=True*)

Let A be this matrix and v be a free module element. If *rows* is True, return a matrix whose rows are the entries of the following vectors:

$$v, vA, vA^2, \dots, vA^{n-1}.$$

If *rows* is False, return a matrix whose columns are the entries of the following vectors:

$$v, Av, A^2v, \dots, A^{n-1}v.$$

INPUT:

- v* - free module element
- n* - nonnegative integer

EXAMPLES:

```
sage: A = matrix(ZZ, 2, [1,1,3,5]); A
[1 1]
[3 5]
sage: v = vector([1,0])
sage: A.iterates(v,0)
[]
sage: A.iterates(v,5)
[ 1  0]
[ 1  1]
[ 4  6]
[22 34]
[124 192]
```

Another example:

```
sage: a = matrix(ZZ, 3, range(9)); a
[0 1 2]
[3 4 5]
[6 7 8]
sage: v = vector([1,0,0])
sage: a.iterates(v,4)
[ 1  0  0]
[ 0  1  2]
[15 18 21]
[180 234 288]
sage: a.iterates(v,4,rows=False)
[ 1  0 15 180]
[ 0  3 42 558]
[ 0  6 69 936]
```

linear_combination_of_columns(*v*)

Return the linear combination of the columns of `self` given by the coefficients in the list *v*.

INPUT:

- *v* - a list of scalars. The length can be less than the number of columns of `self` but not greater.

OUTPUT:

The vector (or free module element) that is a linear combination of the columns of `self`. If the list of scalars has fewer entries than the number of columns, additional zeros are appended to the list until it has as many entries as the number of columns.

EXAMPLES:

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.linear_combination_of_columns([1, 1, 1])
(3, 12)
```

```
sage: a.linear_combination_of_columns([0, 0, 0])
(0, 0)
```

```
sage: a.linear_combination_of_columns([1/2, 2/3, 3/4])
(13/6, 95/12)
```

The list *v* can be anything that is iterable. Perhaps most naturally, a vector may be used.

```
sage: v = vector(ZZ, [1, 2, 3])
sage: a.linear_combination_of_columns(v)
(8, 26)
```

We check that a matrix with no columns behaves properly.

```
sage: matrix(QQ, 2, 0).linear_combination_of_columns([])
(0, 0)
```

The object returned is a vector, or a free module element.

```
sage: B = matrix(ZZ, 4, 3, range(12))
sage: w = B.linear_combination_of_columns([-1, 2, -3])
sage: w
(-4, -10, -16, -22)
sage: w.parent()
Ambient free module of rank 4 over the principal ideal domain Integer Ring
sage: x = B.linear_combination_of_columns([1/2, 1/3, 1/4])
sage: x
(5/6, 49/12, 22/3, 127/12)
sage: x.parent()
Vector space of dimension 4 over Rational Field
```

The length of *v* can be less than the number of columns, but not greater.

```
sage: A = matrix(QQ, 3, 5, range(15))
sage: A.linear_combination_of_columns([1, -2, 3, -4])
(-8, -18, -28)
sage: A.linear_combination_of_columns([1, 2, 3, 4, 5, 6])
Traceback (most recent call last):
...
ValueError: length of v must be at most the number of columns of self
```

linear_combination_of_rows(*v*)

Return the linear combination of the rows of *self* given by the coefficients in the list *v*.

INPUT:

- *v* - a list of scalars. The length can be less than the number of rows of *self* but not greater.

OUTPUT:

The vector (or free module element) that is a linear combination of the rows of *self*. If the list of scalars has fewer entries than the number of rows, additional zeros are appended to the list until it has as many entries as the number of rows.

EXAMPLES:

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.linear_combination_of_rows([1, 2])
(6, 9, 12)

sage: a.linear_combination_of_rows([0, 0])
(0, 0, 0)

sage: a.linear_combination_of_rows([1/2, 2/3])
(2, 19/6, 13/3)
```

The list *v* can be anything that is iterable. Perhaps most naturally, a vector may be used.

```
sage: v = vector(ZZ, [1, 2])
sage: a.linear_combination_of_rows(v)
(6, 9, 12)
```

We check that a matrix with no rows behaves properly.

```
sage: matrix(QQ, 0, 2).linear_combination_of_rows([])
(0, 0)
```

The object returned is a vector, or a free module element.

```
sage: B = matrix(ZZ, 4, 3, range(12))
sage: w = B.linear_combination_of_rows([-1, 2, -3, 4])
sage: w
(24, 26, 28)
sage: w.parent()
Ambient free module of rank 3 over the principal ideal domain Integer Ring
sage: x = B.linear_combination_of_rows([1/2, 1/3, 1/4, 1/5])
sage: x
(43/10, 67/12, 103/15)
sage: x.parent()
Vector space of dimension 3 over Rational Field
```

The length of *v* can be less than the number of rows, but not greater.

```
sage: A = matrix(QQ, 3, 4, range(12))
sage: A.linear_combination_of_rows([2, 3])
(12, 17, 22, 27)
sage: A.linear_combination_of_rows([1, 2, 3, 4])
Traceback (most recent call last):
...
ValueError: length of v must be at most the number of rows of self
```

list()

List of the elements of `self` ordered by elements in each row. It is safe to change the returned list.

Warning: This function returns a list of the entries in the matrix `self`. It does not return a list of the rows of `self`, so it is different than the output of `list(self)`, which returns `[self[0], self[1], ...]`.

EXAMPLES:

```
sage: R.<x,y> = QQ[]
sage: a = matrix(R,2,[x,y,x*y, y,x,2*x+y]); a
[      x      y      x*y]
[      y      x 2*x + y]
sage: v = a.list(); v
[x, y, x*y, y, x, 2*x + y]
```

Note that `list(a)` is different than `a.list()`:

```
sage: a.list()
[x, y, x*y, y, x, 2*x + y]
sage: list(a)
[(x, y, x*y), (y, x, 2*x + y)]
```

Notice that changing the returned list does not change `a` (the list is a copy):

```
sage: v[0] = 25
sage: a
[      x      y      x*y]
[      y      x 2*x + y]
```

mod(p)

Return matrix mod p , over the reduced ring.

EXAMPLES:

```
sage: M = matrix(ZZ, 2, 2, [5, 9, 13, 15])
sage: M.mod(7)
[5 2]
[6 1]
sage: parent(M.mod(7))
Full MatrixSpace of 2 by 2 dense matrices over Ring of integers modulo 7
```

multiplicative_order()

Return the multiplicative order of this matrix, which must therefore be invertible.

EXAMPLES:

```
sage: A = matrix(GF(59), 3, [10, 56, 39, 53, 56, 33, 58, 24, 55])
sage: A.multiplicative_order()
580
sage: (A^580).is_one()
True

sage: B = matrix(GF(10007^3, 'b'), 0)
sage: B.multiplicative_order()
1

sage: C = matrix(GF(2^10, 'c'), 2, 3, [1]*6)
sage: C.multiplicative_order()
Traceback (most recent call last):
```

```

...
ArithmeticError: self must be invertible ...

sage: D = matrix(IntegerModRing(6), 3, [5, 5, 3, 0, 2, 5, 5, 4, 0])
sage: D.multiplicative_order()
Traceback (most recent call last):
...
NotImplementedError: ... only ... over finite fields

sage: E = MatrixSpace(GF(11^2, 'e'), 5).random_element()
sage: (E^E.multiplicative_order()).is_one()
True

```

REFERENCES:

- Frank Celler and C. R. Leedham-Green, “Calculating the Order of an Invertible Matrix”, 1997

mutate(*k*)

Mutates *self* at row and column index *k*.

Warning: Only makes sense if *self* is skew-symmetrizable.

INPUT:

- k* – integer at which row/column *self* is mutated.

EXAMPLES:

Mutation of the B-matrix of the quiver of type A_3 :

```

sage: M = matrix(ZZ, 3, [0, 1, 0, -1, 0, -1, 0, 1, 0]); M
[ 0  1  0]
[-1  0 -1]
[ 0  1  0]

sage: M.mutate(0); M
[ 0 -1  0]
[ 1  0 -1]
[ 0  1  0]

sage: M.mutate(1); M
[ 0  1 -1]
[-1  0  1]
[ 1 -1  0]

sage: M = matrix(ZZ, 6, [0, 1, 0, -1, 0, -1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1]); M
[ 0  1  0]
[-1  0 -1]
[ 0  1  0]
[ 1  0  0]
[ 0  1  0]
[ 0  0  1]

sage: M.mutate(0); M
[ 0 -1  0]
[ 1  0 -1]
[ 0  1  0]
[-1  1  0]
[ 0  1  0]
[ 0  0  1]

```

REFERENCES:

- [FZ2001] S. Fomin, A. Zelevinsky. Cluster Algebras 1: Foundations, arXiv:math/0104151 (2001).

ncols()

Return the number of columns of this matrix.

EXAMPLES:

```
sage: M = MatrixSpace(QQ, 2, 3)
sage: A = M([1, 2, 3, 4, 5, 6])
sage: A
[1 2 3]
[4 5 6]
sage: A.ncols()
3
sage: A.nrows()
2
```

AUTHORS:

- Naqi Jaffery (2006-01-24): examples

nonpivots()

Return the list of i such that the i -th column of self is NOT a pivot column of the reduced row echelon form of self.

OUTPUT: sorted tuple of (Python) integers

EXAMPLES:

```
sage: a = matrix(QQ, 3, 3, range(9)); a
[0 1 2]
[3 4 5]
[6 7 8]
sage: a.echelon_form()
[ 1  0 -1]
[ 0  1  2]
[ 0  0  0]
sage: a.nonpivots()
(2,)
```

nonzero_positions (*copy=True, column_order=False*)

Returns the sorted list of pairs (i, j) such that $\text{self}[i, j] \neq 0$.

INPUT:

- copy* - (default: True) It is safe to change the resulting list (unless you give the option *copy=False*).
- column_order* - (default: False) If true, returns the list of pairs (i, j) such that $\text{self}[i, j] \neq 0$, but sorted by columns, i.e., column $j=0$ entries occur first, then column $j=1$ entries, etc.

EXAMPLES:

```
sage: a = matrix(QQ, 2, 3, [1, 2, 0, 2, 0, 0]); a
[1 2 0]
[2 0 0]
sage: a.nonzero_positions()
[(0, 0), (0, 1), (1, 0)]
sage: a.nonzero_positions(copy=False)
[(0, 0), (0, 1), (1, 0)]
sage: a.nonzero_positions(column_order=True)
[(0, 0), (1, 0), (0, 1)]
sage: a = matrix(QQ, 2, 3, [1, 2, 0, 2, 0, 0], sparse=True); a
```



```

[1 2 0]
[2 0 0]
sage: a.nonzero_positions()
[(0, 0), (0, 1), (1, 0)]
sage: a.nonzero_positions(copy=False)
[(0, 0), (0, 1), (1, 0)]
sage: a.nonzero_positions(column_order=True)
[(0, 0), (1, 0), (0, 1)]

```

nonzero_positions_in_column(*i*)

Return a sorted list of the integers *j* such that `self[j,i]` is nonzero, i.e., such that the *j*-th position of the *i*-th column is nonzero.

INPUT:

- *i* - an integer

OUTPUT: list

EXAMPLES:

```

sage: a = matrix(QQ, 3, 2, [1, 2, 0, 2, 0, 0]); a
[1 2]
[0 2]
[0 0]
sage: a.nonzero_positions_in_column(0)
[0]
sage: a.nonzero_positions_in_column(1)
[0, 1]

```

You'll get an `IndexError`, if you select an invalid column:

```

sage: a.nonzero_positions_in_column(2)
Traceback (most recent call last):
...
IndexError: matrix column index out of range

```

nonzero_positions_in_row(*i*)

Return the integers *j* such that `self[i,j]` is nonzero, i.e., such that the *j*-th position of the *i*-th row is nonzero.

INPUT:

- *i* - an integer

OUTPUT: list

EXAMPLES:

```

sage: a = matrix(QQ, 3, 2, [1, 2, 0, 2, 0, 0]); a
[1 2]
[0 2]
[0 0]
sage: a.nonzero_positions_in_row(0)
[0, 1]
sage: a.nonzero_positions_in_row(1)
[1]
sage: a.nonzero_positions_in_row(2)
[]

```

nrows()

Return the number of rows of this matrix.

EXAMPLES:

```

sage: M = MatrixSpace(QQ, 6, 7)
sage: A = M([1, 2, 3, 4, 5, 6, 7, 22, 3/4, 34, 11, 7, 5, 3, 99, 65, 1/2, 2/3, 3/5, 4/5, 5/6, 9, 8/9, 9/8, 7/6, 6/7])
sage: A
[ 1  2  3  4  5  6  7]
[ 22 3/4 34 11 7 5 3]
[ 99 65 1/2 2/3 3/5 4/5 5/6]
[ 9 8/9 9/8 7/6 6/7 76 4]
[ 0 9 8 7 6 5 4]
[ 123 99 91 28 6 1024 1]
sage: A.ncols()
7
sage: A.nrows()
6

```

AUTHORS:

- Naqi Jaffery (2006-01-24): examples

permute_columns (*permutation*)

Permute the columns of `self` by applying the permutation group element `permutation`.

As a permutation group element acts on integers $\{1, \dots, n\}$ the columns are considered as being numbered from 1 for this operation.

INPUT:

- `permutation` – a `PermutationGroupElement`.

EXAMPLE: We create a matrix:

```

sage: M = matrix(ZZ, [[1, 0, 0, 0, 0], [0, 2, 0, 0, 0], [0, 0, 3, 0, 0], [0, 0, 0, 4, 0], [0, 0, 0, 0, 5]])
sage: M
[1 0 0 0 0]
[0 2 0 0 0]
[0 0 3 0 0]
[0 0 0 4 0]
[0 0 0 0 5]

```

Next of all, create a permutation group element and act on M with it:

```

sage: G = PermutationGroup(['(1,2,3)(4,5)', '(1,2,3,4,5)'])
sage: sigma, tau = G.gens()
sage: sigma
(1,2,3)(4,5)
sage: M.permute_columns(sigma)
sage: M
[0 0 1 0 0]
[2 0 0 0 0]
[0 3 0 0 0]
[0 0 0 0 4]
[0 0 0 5 0]

```

permute_rows (*permutation*)

Permute the rows of `self` by applying the permutation group element `permutation`.

As a permutation group element acts on integers $\{1, \dots, n\}$ the rows are considered as being numbered from 1 for this operation.

INPUT:

- `permutation` – a `PermutationGroupElement`

EXAMPLE: We create a matrix:

```
sage: M = matrix(ZZ, [[1,0,0,0,0],[0,2,0,0,0],[0,0,3,0,0],[0,0,0,4,0],[0,0,0,0,5]])
sage: M
[1 0 0 0 0]
[0 2 0 0 0]
[0 0 3 0 0]
[0 0 0 4 0]
[0 0 0 0 5]
```

Next of all, create a permutation group element and act on M:

```
sage: G = PermutationGroup(['(1,2,3)(4,5)', '(1,2,3,4,5)'])
sage: sigma, tau = G.gens()
sage: sigma
(1,2,3)(4,5)
sage: M.permute_rows(sigma)
sage: M
[0 2 0 0 0]
[0 0 3 0 0]
[1 0 0 0 0]
[0 0 0 0 5]
[0 0 0 4 0]
```

permute_rows_and_columns(*row_permutation*, *column_permutation*)

Permute the rows and columns of *self* by applying the permutation group elements *row_permutation* and *column_permutation* respectively.

As a permutation group element acts on integers $\{1, \dots, n\}$ the rows and columns are considered as being numbered from 1 for this operation.

INPUT:

- *row_permutation* – a `PermutationGroupElement`
- *column_permutation* – a `PermutationGroupElement`

OUTPUT:

- A matrix.

EXAMPLE: We create a matrix:

```
sage: M = matrix(ZZ, [[1,0,0,0,0],[0,2,0,0,0],[0,0,3,0,0],[0,0,0,4,0],[0,0,0,0,5]])
sage: M
[1 0 0 0 0]
[0 2 0 0 0]
[0 0 3 0 0]
[0 0 0 4 0]
[0 0 0 0 5]
```

Next of all, create a permutation group element and act on M:

```
sage: G = PermutationGroup(['(1,2,3)(4,5)', '(1,2,3,4,5)'])
sage: sigma, tau = G.gens()
sage: sigma
(1,2,3)(4,5)
sage: M.permute_rows_and_columns(sigma,tau)
sage: M
[2 0 0 0 0]
[0 3 0 0 0]
[0 0 0 0 1]
```

```
[0 0 0 5 0]
[0 0 4 0 0]
```

pivots()

Return the pivot column positions of this matrix.

OUTPUT: a tuple of Python integers: the position of the first nonzero entry in each row of the echelon form.

This returns a tuple so it is immutable; see #10752.

EXAMPLES:

```
sage: A = matrix(QQ, 2, 2, range(4))
sage: A.pivots()
(0, 1)
```

rank()

TESTS:

We should be able to compute the rank of a matrix whose entries are polynomials over a finite field (trac #5014):

```
sage: P.<x> = PolynomialRing(GF(17))
sage: m = matrix(P, [ [ 6*x^2 + 8*x + 12, 10*x^2 + 4*x + 11],
...                  [8*x^2 + 12*x + 15, 8*x^2 + 9*x + 16] ])
sage: m.rank()
2
```

rescale_col(i, s, start_row=0)

Replace i-th col of self by s times i-th col of self.

INPUT:

- i - ith column
- s - scalar
- start_row - only rescale entries at this row and lower

EXAMPLES: We rescale the last column of a matrix over the rational numbers:

```
sage: a = matrix(QQ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.rescale_col(2, 1/2); a
[ 0 1 1]
[ 3 4 5/2]
sage: R.<x> = QQ[]
```

We rescale the last column of a matrix over a polynomial ring:

```
sage: a = matrix(R, 2, 3, [1, x, x^2, x^3, x^4, x^5]); a
[ 1 x x^2]
[x^3 x^4 x^5]
sage: a.rescale_col(2, 1/2); a
[ 1 x 1/2*x^2]
[ x^3 x^4 1/2*x^5]
```

We try and fail to rescale a matrix over the integers by a non-integer:

```
sage: a = matrix(ZZ, 2, 3, [0, 1, 2, 3, 4, 4]); a
[0 1 2]
```

```
[3 4 4]
sage: a.rescale_col(2,1/2)
Traceback (most recent call last):
...
TypeError: Rescaling column by Rational Field element cannot be done over Integer Ring, use
```

To rescale the matrix by 1/2, you must change the base ring to the rationals:

```
sage: a = a.change_ring(QQ); a
[0 1 2]
[3 4 4]
sage: a.rescale_col(2,1/2); a
[0 1 1]
[3 4 2]
```

rescale_row (*i*, *s*, *start_col*=0)

Replace *i*-th row of self by *s* times *i*-th row of self.

INPUT:

- *i* - *i*th row
- *s* - scalar
- *start_col* - only rescale entries at this column and to the right

EXAMPLES: We rescale the second row of a matrix over the rational numbers:

```
sage: a = matrix(QQ, 3, range(6)); a
[0 1]
[2 3]
[4 5]
sage: a.rescale_row(1,1/2); a
[ 0 1]
[ 1 3/2]
[ 4 5]
```

We rescale the second row of a matrix over a polynomial ring:

```
sage: R.<x> = QQ[]
sage: a = matrix(R, 3, [1, x, x^2, x^3, x^4, x^5]); a
[ 1 x]
[x^2 x^3]
[x^4 x^5]
sage: a.rescale_row(1,1/2); a
[ 1 x]
[1/2*x^2 1/2*x^3]
[ x^4 x^5]
```

We try and fail to rescale a matrix over the integers by a non-integer:

```
sage: a = matrix(ZZ, 2, 3, [0, 1, 2, 3, 4, 4]); a
[0 1 2]
[3 4 4]
sage: a.rescale_row(1,1/2)
Traceback (most recent call last):
...
TypeError: Rescaling row by Rational Field element cannot be done over Integer Ring, use cha
```

To rescale the matrix by 1/2, you must change the base ring to the rationals:

```
sage: a = a.change_ring(QQ); a
[0 1 2]
[3 4 4]
sage: a.rescale_col(1, 1/2); a
[ 0 1/2 2]
[ 3 2 4]
```

set_col_to_multiple_of_col(*i, j, s*)

Set column *i* equal to *s* times column *j*.

EXAMPLES: We change the second column to -3 times the first column.

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.set_col_to_multiple_of_col(1, 0, -3)
sage: a
[ 0 0 2]
[ 3 -9 5]
```

If we try to multiply a column by a rational number, we get an error message:

```
sage: a.set_col_to_multiple_of_col(1, 0, 1/2)
Traceback (most recent call last):
...
TypeError: Multiplying column by Rational Field element cannot be done over Integer Ring, us
```

set_immutable()

Call this function to set the matrix as immutable.

Matrices are always mutable by default, i.e., you can change their entries using $A[i, j] = x$. However, mutable matrices aren't hashable, so can't be used as keys in dictionaries, etc. Also, often when implementing a class, you might compute a matrix associated to it, e.g., the matrix of a Hecke operator. If you return this matrix to the user you're really returning a reference and the user could then change an entry; this could be confusing. Thus you should set such a matrix immutable.

EXAMPLES:

```
sage: A = Matrix(QQ, 2, 2, range(4))
sage: A.is_mutable()
True
sage: A[0,0] = 10
sage: A
[10 1]
[ 2 3]
```

Mutable matrices are not hashable, so can't be used as keys for dictionaries:

```
sage: hash(A)
Traceback (most recent call last):
...
TypeError: mutable matrices are unhashable
sage: v = {A:1}
Traceback (most recent call last):
...
TypeError: mutable matrices are unhashable
```

If we make *A* immutable it suddenly is hashable.

```

sage: A.set_immutable()
sage: A.is_mutable()
False
sage: A[0,0] = 10
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead (i.e., use copy(M) to change a
sage: hash(A) #random
12
sage: v = {A:1}; v
{[10  1]
 [ 2  3]: 1}

```

set_row_to_multiple_of_row(*i,j,s*)

Set row *i* equal to *s* times row *j*.

EXAMPLES: We change the second row to -3 times the first row:

```

sage: a = matrix(ZZ,2,3,range(6)); a
[0 1 2]
[3 4 5]
sage: a.set_row_to_multiple_of_row(1,0,-3)
sage: a
[ 0  1  2]
[ 0 -3 -6]

```

If we try to multiply a row by a rational number, we get an error message:

```

sage: a.set_row_to_multiple_of_row(1,0,1/2)
Traceback (most recent call last):
...
TypeError: Multiplying row by Rational Field element cannot be done over Integer Ring, use c

```

str(*rep_mapping=None, zero=None, plus_one=None, minus_one=None*)

Return a nice string representation of the matrix.

INPUT:

- *rep_mapping* - a dictionary or callable used to override the usual representation of elements.

If *rep_mapping* is a dictionary then keys should be elements of the base ring and values the desired string representation. Values sent in via the other keyword arguments will override values in the dictionary. Use of a dictionary can potentially take a very long time due to the need to hash entries of the matrix. Matrices with entries from $\mathbb{Q}\overline{\mathbb{Q}}$ are one example.

If *rep_mapping* is callable then it will be called with elements of the matrix and must return a string. Simply call `repr()` on elements which should have the default representation.

- *zero* - string (default: None); if not None use the value of *zero* as the representation of the zero element.
- *plus_one* - string (default: None); if not None use the value of *plus_one* as the representation of the one element.
- *minus_one* - string (default: None); if not None use the value of *minus_one* as the representation of the negative of the one element.

EXAMPLES:

```

sage: R = PolynomialRing(QQ,6,'z')
sage: a = matrix(2,3, R.gens())
sage: a.__repr__()

```

```

'[z0 z1 z2]\n[z3 z4 z5]'

sage: M = matrix([[1,0],[2,-1]])
sage: M.str()
'[ 1  0]\n[ 2 -1]'
sage: M.str(plus_one='+',minus_one='-',zero='.')
'[+ .]\n[2 -]'
sage: M.str({1:"not this one",2:"II"},minus_one="*",plus_one="I")
'[ I  0]\n[II  *]'

sage: def print_entry(x):
...     if x>0:
...         return '+'
...     elif x<0:
...         return '-'
...     else: return '.'
...
sage: M.str(print_entry)
'[+ .]\n[+ -]'
sage: M.str(repr)
'[ 1  0]\n[ 2 -1]'

```

TESTS:

Prior to Trac #11544 this could take a full minute to run (2011).

```

sage: A = matrix(QQ, 4, 4, [1, 2, -2, 2, 1, 0, -1, -1, 0, -1, 1, 1, -1, 2, 1/2, 0])
sage: e = A.eigenvalues()[3]
sage: K = (A-e).kernel()
sage: P = K.basis_matrix()
sage: P.str()
'[
      1.0000000000000000? + 0.?e-17*I -2.116651487479748? + 0.0255565807096352?*I -0

```

swap_columns(c1, c2)

Swap columns c1 and c2 of self.

EXAMPLES: We create a rational matrix:

```

sage: M = MatrixSpace(QQ, 3, 3)
sage: A = M([1, 9, -7, 4/5, 4, 3, 6, 4, 3])
sage: A
[ 1  9 -7]
[4/5 4  3]
[ 6  4  3]

```

Since the first column is numbered zero, this swaps the second and third columns:

```

sage: A.swap_columns(1, 2); A
[ 1 -7  9]
[4/5 3  4]
[ 6  3  4]

```

swap_rows(r1, r2)

Swap rows r1 and r2 of self.

EXAMPLES: We create a rational matrix:

```

sage: M = MatrixSpace(QQ, 3, 3)
sage: A = M([1, 9, -7, 4/5, 4, 3, 6, 4, 3])
sage: A
[ 1  9 -7]

```



```
[4/5  4  3]
[  6  4  3]
```

Since the first row is numbered zero, this swaps the first and third rows:

```
sage: A.swap_rows(0,2); A
[  6  4  3]
[4/5  4  3]
[  1  9 -7]
```

with_added_multiple_of_column(*i, j, s, start_row=0*)

Add *s* times column *j* to column *i*, returning new matrix.

EXAMPLES: We add -1 times the third column to the second column of an integer matrix, remembering to start numbering cols at zero:

```
sage: a = matrix(ZZ,2,3,range(6)); a
[0 1 2]
[3 4 5]
sage: b = a.with_added_multiple_of_column(1,2,-1); b
[ 0 -1  2]
[ 3 -1  5]
```

The original matrix is unchanged:

```
sage: a
[0 1 2]
[3 4 5]
```

Adding a rational multiple is okay, and reassigning a variable is okay:

```
sage: a = a.with_added_multiple_of_column(0,1,1/3); a
[ 1/3  1  2]
[13/3  4  5]
```

with_added_multiple_of_row(*i, j, s, start_col=0*)

Add *s* times row *j* to row *i*, returning new matrix.

EXAMPLES: We add -3 times the first row to the second row of an integer matrix, remembering to start numbering rows at zero:

```
sage: a = matrix(ZZ,2,3,range(6)); a
[0 1 2]
[3 4 5]
sage: b = a.with_added_multiple_of_row(1,0,-3); b
[ 0  1  2]
[ 3  1 -1]
```

The original matrix is unchanged:

```
sage: a
[0 1 2]
[3 4 5]
```

Adding a rational multiple is okay, and reassigning a variable is okay:

```
sage: a = a.with_added_multiple_of_row(0,1,1/3); a
[  1  7/3 11/3]
[  3   4   5]
```

with_col_set_to_multiple_of_col(*i, j, s*)

Set column i equal to s times column j , returning a new matrix.

EXAMPLES: We change the second column to -3 times the first column.

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: b = a.with_col_set_to_multiple_of_col(1, 0, -3); b
[ 0 0 2]
[ 3 -9 5]
```

Note that the original matrix is unchanged:

```
sage: a
[0 1 2]
[3 4 5]
```

Adding a rational multiple is okay, and reassigning a variable is okay:

```
sage: a = a.with_col_set_to_multiple_of_col(1, 0, 1/2); a
[ 0 0 2]
[ 3 3/2 5]
```

with_permuted_columns (*permutation*)

Return the matrix obtained from permuting the columns of `self` by applying the permutation group element `permutation`.

As a permutation group element acts on integers $\{1, \dots, n\}$ the columns are considered as being numbered from 1 for this operation.

INPUT:

- `permutation`, a `PermutationGroupElement`

OUTPUT:

- A matrix.

EXAMPLE: We create some matrix:

```
sage: M = matrix(ZZ, [[1, 0, 0, 0, 0], [0, 2, 0, 0, 0], [0, 0, 3, 0, 0], [0, 0, 0, 4, 0], [0, 0, 0, 0, 5]])
sage: M
[1 0 0 0 0]
[0 2 0 0 0]
[0 0 3 0 0]
[0 0 0 4 0]
[0 0 0 0 5]
```

Next of all, create a permutation group element and act on `M`:

```
sage: G = PermutationGroup(['(1, 2, 3) (4, 5)', '(1, 2, 3, 4, 5)'])
sage: sigma, tau = G.gens()
sage: sigma
(1, 2, 3) (4, 5)
sage: M.with_permuted_columns(sigma)
[0 0 1 0 0]
[2 0 0 0 0]
[0 3 0 0 0]
[0 0 0 0 4]
[0 0 0 5 0]
```

with_permuted_rows (*permutation*)

Return the matrix obtained from permuting the rows of `self` by applying the permutation group element

permutation.

As a permutation group element acts on integers $\{1, \dots, n\}$ the rows are considered as being numbered from 1 for this operation.

INPUT:

- `permutation` – a `PermutationGroupElement`

OUTPUT:

- A matrix.

EXAMPLE: We create a matrix:

```
sage: M = matrix(ZZ, [[1,0,0,0,0],[0,2,0,0,0],[0,0,3,0,0],[0,0,0,4,0],[0,0,0,0,5]])
sage: M
[1 0 0 0 0]
[0 2 0 0 0]
[0 0 3 0 0]
[0 0 0 4 0]
[0 0 0 0 5]
```

Next of all, create a permutation group element and act on M:

```
sage: G = PermutationGroup(['(1,2,3)(4,5)', '(1,2,3,4,5)'])
sage: sigma, tau = G.gens()
sage: sigma
(1,2,3)(4,5)
sage: M.with_permuted_rows(sigma)
[0 2 0 0 0]
[0 0 3 0 0]
[1 0 0 0 0]
[0 0 0 0 5]
[0 0 0 4 0]
```

`with_permuted_rows_and_columns` (*row_permutation*, *column_permutation*)

Return the matrix obtained from permuting the rows and columns of `self` by applying the permutation group elements `row_permutation` and `column_permutation`.

As a permutation group element acts on integers $\{1, \dots, n\}$ the rows are considered as being numbered from 1 for this operation.

INPUT:

- `row_permutation` – a `PermutationGroupElement`
- `column_permutation` – a `PermutationGroupElement`

OUTPUT:

- A matrix.

EXAMPLE: We create a matrix:

```
sage: M = matrix(ZZ, [[1,0,0,0,0],[0,2,0,0,0],[0,0,3,0,0],[0,0,0,4,0],[0,0,0,0,5]])
sage: M
[1 0 0 0 0]
[0 2 0 0 0]
[0 0 3 0 0]
[0 0 0 4 0]
[0 0 0 0 5]
```

Next of all, create a permutation group element and act on M:

```
sage: G = PermutationGroup(['(1,2,3)(4,5)', '(1,2,3,4,5)'])
sage: sigma, tau = G.gens()
sage: sigma
(1,2,3)(4,5)
sage: M.with_permuted_rows_and_columns(sigma,tau)
[2 0 0 0 0]
[0 3 0 0 0]
[0 0 0 0 1]
[0 0 0 5 0]
[0 0 4 0 0]
```

with_rescaled_col (*i*, *s*, *start_row*=0)

Replaces *i*-th col of self by *s* times *i*-th col of self, returning new matrix.

EXAMPLES: We rescale the last column of a matrix over the integers:

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: b = a.with_rescaled_col(2, -2); b
[ 0  1 -4]
[ 3  4 -10]
```

The original matrix is unchanged:

```
sage: a
[0 1 2]
[3 4 5]
```

Adding a rational multiple is okay, and reassigning a variable is okay:

```
sage: a = a.with_rescaled_col(1, 1/3); a
[ 0 1/3  2]
[ 3 4/3  5]
```

with_rescaled_row (*i*, *s*, *start_col*=0)

Replaces *i*-th row of self by *s* times *i*-th row of self, returning new matrix.

EXAMPLES: We rescale the second row of a matrix over the integers:

```
sage: a = matrix(ZZ, 3, 2, range(6)); a
[0 1]
[2 3]
[4 5]
sage: b = a.with_rescaled_row(1, -2); b
[ 0  1]
[-4 -6]
[ 4  5]
```

The original matrix is unchanged:

```
sage: a
[0 1]
[2 3]
[4 5]
```

Adding a rational multiple is okay, and reassigning a variable is okay:

```
sage: a = a.with_rescaled_row(2, 1/3); a
[ 0  1]
[ 2  3]
```

```
[4/3 5/3]
```

with_row_set_to_multiple_of_row (*i*, *j*, *s*)

Set row *i* equal to *s* times row *j*, returning a new matrix.

EXAMPLES: We change the second row to -3 times the first row:

```
sage: a = matrix(ZZ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: b = a.with_row_set_to_multiple_of_row(1, 0, -3); b
[ 0 1 2]
[ 0 -3 -6]
```

Note that the original matrix is unchanged:

```
sage: a
[0 1 2]
[3 4 5]
```

Adding a rational multiple is okay, and reassigning a variable is okay:

```
sage: a = a.with_row_set_to_multiple_of_row(1, 0, 1/2); a
[ 0 1 2]
[ 0 1/2 1]
```

with_swapped_columns (*c1*, *c2*)

Swap columns *c1* and *c2* of *self* and return a new matrix.

INPUT:

- *c1*, *c2* - integers specifying columns of *self* to interchange

OUTPUT:

A new matrix, identical to *self* except that columns *c1* and *c2* are swapped.

EXAMPLES:

Remember that columns are numbered starting from zero.

```
sage: A = matrix(QQ, 4, range(20))
sage: A.with_swapped_columns(1, 2)
[ 0 2 1 3 4]
[ 5 7 6 8 9]
[10 12 11 13 14]
[15 17 16 18 19]
```

Trying to swap a column with itself will succeed, but still return a new matrix.

```
sage: A = matrix(QQ, 4, range(20))
sage: B = A.with_swapped_columns(2, 2)
sage: A == B
True
sage: A is B
False
```

The column specifications are checked.

```
sage: A = matrix(4, range(20))
sage: A.with_swapped_columns(-1, 2)
Traceback (most recent call last):
...
```

```
IndexError: matrix column index out of range
```

```
sage: A.with_swapped_columns(2, 5)
Traceback (most recent call last):
...
IndexError: matrix column index out of range
```

with_swapped_rows (*r1*, *r2*)

Swap rows *r1* and *r2* of *self* and return a new matrix.

INPUT:

- *r1*, *r2* - integers specifying rows of *self* to interchange

OUTPUT:

A new matrix, identical to *self* except that rows *r1* and *r2* are swapped.

EXAMPLES:

Remember that rows are numbered starting from zero.

```
sage: A = matrix(QQ, 4, range(20))
sage: A.with_swapped_rows(1, 2)
[ 0  1  2  3  4]
[10 11 12 13 14]
[ 5  6  7  8  9]
[15 16 17 18 19]
```

Trying to swap a row with itself will succeed, but still return a new matrix.

```
sage: A = matrix(QQ, 4, range(20))
sage: B = A.with_swapped_rows(2, 2)
sage: A == B
True
sage: A is B
False
```

The row specifications are checked.

```
sage: A = matrix(4, range(20))
sage: A.with_swapped_rows(-1, 2)
Traceback (most recent call last):
...
IndexError: matrix row index out of range

sage: A.with_swapped_rows(2, 5)
Traceback (most recent call last):
...
IndexError: matrix row index out of range
```

`sage.matrix.matrix0.set_max_cols(n)`

Sets the global variable `max_cols` (which is used in deciding how to output a matrix).

EXAMPLES:

```
sage: from sage.matrix.matrix0 import set_max_cols
sage: set_max_cols(50)
```

`sage.matrix.matrix0.set_max_rows(n)`

Sets the global variable `max_rows` (which is used in deciding how to output a matrix).

EXAMPLES:

```
sage: from sage.matrix.matrix0 import set_max_rows
sage: set_max_rows(20)
```

```
sage.matrix.matrix0.unpickle(cls, parent, immutability, cache, data, version)
```

Unpickle a matrix. This is only used internally by Sage. Users should never call this function directly.

EXAMPLES: We illustrating saving and loading several different types of matrices.

OVER \mathbb{Z} :

```
sage: A = matrix(ZZ, 2, range(4))
sage: loads(dumps(A)) # indirect doctest
[0 1]
[2 3]
```

Sparse OVER \mathbb{Q} :

Dense over $\mathbb{Q}[x, y]$:

Dense over finite field.

BASE CLASS FOR MATRICES, PART 1

Base class for matrices, part 1

For design documentation see [sage.matrix.docs](#).

TESTS:

```
sage: A = Matrix(GF(5), 3, 3, xrange(9))
sage: TestSuite(A).run()
```

```
class sage.matrix.matrix1.Matrix
    Bases: sage.matrix.matrix0.Matrix
```

The initialization routine of the `Matrix` base class ensures that it sets the attributes `self._parent`, `self._base_ring`, `self._nrows`, `self._ncols`. It sets the latter ones by accessing the relevant information on `parent`, which is often slower than what a more specific subclass can do.

Subclasses of `Matrix` can safely skip calling `Matrix.__init__` provided they take care of initializing these attributes themselves.

The private attributes `self._is_immutable` and `self._cache` are implicitly initialized to valid values upon memory allocation.

EXAMPLES:

```
sage: import sage.matrix.matrix0
sage: A = sage.matrix.matrix0.Matrix(MatrixSpace(QQ, 2))
sage: type(A)
<type 'sage.matrix.matrix0.Matrix'>
```

augment (*right*, *subdivide=False*)

Returns a new matrix formed by appending the matrix (or vector) *right* on the right side of *self*.

INPUT:

- *right* - a matrix, vector or free module element, whose dimensions are compatible with *self*.
- *subdivide* - default: `False` - request the resulting matrix to have a new subdivision, separating *self* from *right*.

OUTPUT:

A new matrix formed by appending *right* onto the right side of *self*. If *right* is a vector (or free module element) then in this context it is appropriate to consider it as a column vector. (The code first converts a vector to a 1-column matrix.)

If *subdivide* is `True` then any column subdivisions for the two matrices are preserved, and a new subdivision is added between *self* and *right*. If the row divisions are identical, then they are preserved,

otherwise they are discarded. When `subdivide` is `False` there is no subdivision information in the result.

Warning: If `subdivide` is `True` then unequal row subdivisions will be discarded, since it would be ambiguous how to interpret them. If the subdivision behavior is not what you need, you can manage subdivisions yourself with methods like `get_subdivisions()` and `subdivide()`. You might also find `block_matrix()` or `block_diagonal_matrix()` useful and simpler in some instances.

EXAMPLES:

Augmenting with a matrix.

```
sage: A = matrix(QQ, 3, range(12))
sage: B = matrix(QQ, 3, range(9))
sage: A.augment(B)
[ 0  1  2  3  0  1  2]
[ 4  5  6  7  3  4  5]
[ 8  9 10 11  6  7  8]
```

Augmenting with a vector.

```
sage: A = matrix(QQ, 2, [0, 2, 4, 6, 8, 10])
sage: v = vector(QQ, 2, [100, 200])
sage: A.augment(v)
[ 0  2  4 100]
[ 6  8 10 200]
```

Errors are raised if the sizes are incompatible.

```
sage: A = matrix(RR, [[1, 2], [3, 4]])
sage: B = matrix(RR, [[10, 20], [30, 40], [50, 60]])
sage: A.augment(B)
Traceback (most recent call last):
...
TypeError: number of rows must be the same, 2 != 3

sage: v = vector(RR, [100, 200, 300])
sage: A.augment(v)
Traceback (most recent call last):
...
TypeError: number of rows must be the same, 2 != 3
```

Setting `subdivide` to `True` will, in its simplest form, add a subdivision between `self` and `right`.

```
sage: A = matrix(QQ, 3, range(12))
sage: B = matrix(QQ, 3, range(15))
sage: A.augment(B, subdivide=True)
[ 0  1  2  3| 0  1  2  3  4]
[ 4  5  6  7| 5  6  7  8  9]
[ 8  9 10 11|10 11 12 13 14]
```

Column subdivisions are preserved by augmentation, and enriched, if subdivisions are requested. (So multiple augmentations can be recorded.)

```
sage: A = matrix(QQ, 3, range(6))
sage: A.subdivide(None, [1])
sage: B = matrix(QQ, 3, range(9))
sage: B.subdivide(None, [2])
sage: A.augment(B, subdivide=True)
```

```
[0|1|0 1|2]
[2|3|3 4|5]
[4|5|6 7|8]
```

Row subdivisions can be preserved, but only if they are identical. Otherwise, this information is discarded and must be managed separately.

```
sage: A = matrix(QQ, 3, range(6))
sage: A.subdivide([1,3], None)
sage: B = matrix(QQ, 3, range(9))
sage: B.subdivide([1,3], None)
sage: A.augment(B, subdivide=True)
[0 1|0 1 2]
[---+-----]
[2 3|3 4 5]
[4 5|6 7 8]
[---+-----]

sage: A.subdivide([1,2], None)
sage: A.augment(B, subdivide=True)
[0 1|0 1 2]
[2 3|3 4 5]
[4 5|6 7 8]
```

The result retains the base ring of `self` by coercing the elements of `right` into the base ring of `self`.

```
sage: A = matrix(QQ, 2, [1,2])
sage: B = matrix(RR, 2, [sin(1.1), sin(2.2)])
sage: C = A.augment(B); C
[          1 183017397/205358938]
[          2 106580492/131825561]
```

```
sage: C.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
```

```
sage: D = B.augment(A); D
[0.89120736006... 1.0000000000000000]
[0.80849640381... 2.0000000000000000]
sage: D.parent()
```

```
Full MatrixSpace of 2 by 2 dense matrices over Real Field with 53 bits of precision
```

Sometimes it is not possible to coerce into the base ring of `self`. A solution is to change the base ring of `self` to a more expansive ring. Here we mix the rationals with a ring of polynomials with rational coefficients.

```
sage: R = PolynomialRing(QQ, 'y')
sage: A = matrix(QQ, 1, [1,2])
sage: B = matrix(R, 1, ['y', 'y^2'])
```

```
sage: C = B.augment(A); C
[ y y^2 1 2]
sage: C.parent()
```

```
Full MatrixSpace of 1 by 4 dense matrices over Univariate Polynomial Ring in y over Rational Field
```

```
sage: D = A.augment(B)
Traceback (most recent call last):
...
TypeError: not a constant polynomial

sage: E = A.change_ring(R)
```

```
sage: F = E.augment(B); F
[ 1  2  y y^2]
sage: F.parent()
Full MatrixSpace of 1 by 4 dense matrices over Univariate Polynomial Ring in y over Rational
```

AUTHORS:

- Naqi Jaffery (2006-01-24): examples
- Rob Beezer (2010-12-07): vector argument, docstring, subdivisions

block_sum(*other*)

Return the block matrix that has self and other on the diagonal:

```
[ self  0 ]
[  0 other ]
```

EXAMPLES:

```
sage: A = matrix(QQ[['t']], 2, range(1, 5))
sage: A.block_sum(100*A)
[ 1  2  0  0]
[ 3  4  0  0]
[ 0  0 100 200]
[ 0  0 300 400]
```

column(*i*, *from_list=False*)

Return the *i*'th column of this matrix as a vector.

This column is a dense vector if and only if the matrix is a dense matrix.

INPUT:

- i* - integer
- from_list* - bool (default: False); if true, returns the *i*'th element of `self.columns()` (see `columns()`), which may be faster, but requires building a list of all columns the first time it is called after an entry of the matrix is changed.

EXAMPLES:

```
sage: a = matrix(2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.column(1)
(1, 4)
```

If the column is negative, it wraps around, just like with list indexing, e.g., -1 gives the right-most column:

```
sage: a.column(-1)
(2, 5)
```

TESTS:

```
sage: a = matrix(2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.column(3)
Traceback (most recent call last):
...
IndexError: column index out of range
sage: a.column(-4)
```

```
Traceback (most recent call last):
...
IndexError: column index out of range
```

columns (*copy=True*)

Return a list of the columns of self.

INPUT:

- *copy* - (default: True) if True, return a copy of the list of columns which is safe to change.

If self is a sparse matrix, columns are returned as sparse vectors, otherwise returned vectors are dense.

EXAMPLES:

```
sage: matrix(3, [1..9]).columns()
[(1, 4, 7), (2, 5, 8), (3, 6, 9)]
sage: matrix(RR, 2, [sqrt(2), pi, exp(1), 0]).columns()
[(1.41421356237310, 2.71828182845905), (3.14159265358979, 0.000000000000000)]
sage: matrix(RR, 0, 2, []).columns()
[(), ()]
sage: matrix(RR, 2, 0, []).columns()
[]
sage: m = matrix(RR, 3, 3, {(1,2): pi, (2, 2): -1, (0,1): sqrt(2)})
sage: parent(m.columns()[0])
Sparse vector space of dimension 3 over Real Field with 53 bits of precision
```

Sparse matrices produce sparse columns.

```
sage: A = matrix(QQ, 2, range(4), sparse=True)
sage: v = A.columns()[0]
sage: v.is_sparse()
True
```

TESTS:

```
sage: A = matrix(QQ, 4, range(16))
sage: A.columns('junk')
Traceback (most recent call last):
...
ValueError: 'copy' must be True or False, not junk
```

delete_columns (*dcols, check=True*)

Return the matrix constructed from deleting the columns with indices in the *dcols* list.

INPUT:

- *dcols* - list of indices of columns to be deleted from self.
- *check* - checks whether any index in *dcols* is out of range. Defaults to True.

SEE ALSO: The methods `delete_rows()` and `matrix_from_columns()` are related.

EXAMPLES:

```
sage: A = Matrix(3, 4, range(12)); A
[ 0  1  2  3]
[ 4  5  6  7]
[ 8  9 10 11]
sage: A.delete_columns([0, 2])
[ 1  3]
```

```
[ 5  7]
[ 9 11]
```

`dcols` can be a tuple. But only the underlying set of indices matters.

```
sage: A.delete_columns((2,0,2))
[ 1  3]
[ 5  7]
[ 9 11]
```

The default is to check whether any index in `dcols` is out of range.

```
sage: A.delete_columns([-1,2,4])
Traceback (most recent call last):
...
IndexError: [4, -1] contains invalid indices.
sage: A.delete_columns([-1,2,4], check=False)
[ 0  1  3]
[ 4  5  7]
[ 8  9 11]
```

TESTS:

The list of indices is checked.

```
sage: A.delete_columns('junk')
Traceback (most recent call last):
...
TypeError: The argument must be a list or a tuple, not junk
```

AUTHORS:

- Wai Yan Pong (2012-03-05)

`delete_rows` (*drows*, *check=True*)

Return the matrix constructed from deleting the rows with indices in the `drows` list.

INPUT:

- `drows` - list of indices of rows to be deleted from self.
- `check` - checks whether any index in `drows` is out of range. Defaults to `True`.

SEE ALSO: The methods `delete_columns()` and `matrix_from_rows()` are related.

EXAMPLES:

```
sage: A = Matrix(4,3,range(12)); A
[ 0  1  2]
[ 3  4  5]
[ 6  7  8]
[ 9 10 11]
sage: A.delete_rows([0,2])
[ 3  4  5]
[ 9 10 11]
```

`drows` can be a tuple. But only the underlying set of indices matters.

```
sage: A.delete_rows((2,0,2))
[ 3  4  5]
[ 9 10 11]
```

The default is to check whether the any index in `drows` is out of range.

```
sage: A.delete_rows([-1,2,4])
Traceback (most recent call last):
...
IndexError: [4, -1] contains invalid indices.
sage: A.delete_rows([-1,2,4], check=False)
[ 0  1  2]
[ 3  4  5]
[ 9 10 11]
```

TESTS:

The list of indices is checked.

```
sage: A.delete_rows('junk')
Traceback (most recent call last):
...
TypeError: The argument must be a list or a tuple, not junk
```

AUTHORS:

- Wai Yan Pong (2012-03-05)

dense_columns (*copy=True*)

Return list of the dense columns of self.

INPUT:

- `copy` - (default: `True`) if `True`, return a copy so you can modify it safely

EXAMPLES:

An example over the integers:

```
sage: a = matrix(3,3,range(9)); a
[0 1 2]
[3 4 5]
[6 7 8]
sage: a.dense_columns()
[(0, 3, 6), (1, 4, 7), (2, 5, 8)]
```

We do an example over a polynomial ring:

```
sage: R.<x> = QQ[ ]
sage: a = matrix(R, 2, [x,x^2, 2/3*x,1+x^5]); a
[      x      x^2]
[ 2/3*x x^5 + 1]
sage: a.dense_columns()
[(x, 2/3*x), (x^2, x^5 + 1)]
sage: a = matrix(R, 2, [x,x^2, 2/3*x,1+x^5], sparse=True)
sage: c = a.dense_columns(); c
[(x, 2/3*x), (x^2, x^5 + 1)]
sage: parent(c[1])
Ambient free module of rank 2 over the principal ideal domain Univariate Polynomial Ring in
```

TESTS:

Check that the returned rows are immutable as per [trac ticket #14874](#):

```
sage: m = Mat(ZZ,3,3)(range(9))
sage: v = m.dense_columns()
```

```
sage: map(lambda x: x.is_mutable(), v)
[False, False, False]
```

dense_matrix()

If this matrix is sparse, return a dense matrix with the same entries. If this matrix is dense, return this matrix (not a copy).

Note: The definition of “dense” and “sparse” in Sage have nothing to do with the number of nonzero entries. Sparse and dense are properties of the underlying representation of the matrix.

EXAMPLES:

```
sage: A = MatrixSpace(QQ, 2, sparse=True) ([1, 2, 0, 1])
sage: A.is_sparse()
True
sage: B = A.dense_matrix()
sage: B.is_sparse()
False
sage: A*B
[1 4]
[0 1]
sage: A.parent()
Full MatrixSpace of 2 by 2 sparse matrices over Rational Field
sage: B.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
```

In Sage, the product of a sparse and a dense matrix is always dense:

```
sage: (A*B).parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: (B*A).parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
```

TESTS:

Make sure that subdivisions are preserved when switching between dense and sparse matrices:

```
sage: a = matrix(ZZ, 3, range(9))
sage: a.subdivide([1, 2], 2)
sage: a.subdivisions()
([1, 2], [2])
sage: b = a.sparse_matrix().dense_matrix()
sage: b.subdivisions()
([1, 2], [2])
```

dense_rows (copy=True)

Return list of the dense rows of self.

INPUT:

- `copy` - (default: `True`) if `True`, return a copy so you can modify it safely (note that the individual vectors in the copy should not be modified since they are mutable!)

EXAMPLES:

```
sage: m = matrix(3, range(9)); m
[0 1 2]
[3 4 5]
[6 7 8]
```



```

sage: v = m.dense_rows(); v
[(0, 1, 2), (3, 4, 5), (6, 7, 8)]
sage: v is m.dense_rows()
False
sage: m.dense_rows(copy=False) is m.dense_rows(copy=False)
True
sage: m[0,0] = 10
sage: m.dense_rows()
[(10, 1, 2), (3, 4, 5), (6, 7, 8)]

```

TESTS:

Check that the returned rows are immutable as per [trac ticket #14874](#):

```

sage: m = Mat(ZZ, 3, 3)(range(9))
sage: v = m.dense_rows()
sage: map(lambda x: x.is_mutable(), v)
[False, False, False]

```

lift()

Return lift of self to the covering ring of the base ring R, which is by definition the ring returned by calling `cover_ring()` on R, or just R itself if the `cover_ring` method is not defined.

EXAMPLES:

```

sage: M = Matrix(Integers(7), 2, 2, [5, 9, 13, 15]) ; M
[5 2]
[6 1]
sage: M.lift()
[5 2]
[6 1]
sage: parent(M.lift())
Full MatrixSpace of 2 by 2 dense matrices over Integer Ring

```

The field QQ doesn't have a `cover_ring` method:

```

sage: hasattr(QQ, 'cover_ring')
False

```

So lifting a matrix over QQ gives back the same exact matrix.

```

sage: B = matrix(QQ, 2, [1..4])
sage: B.lift()
[1 2]
[3 4]
sage: B.lift() is B
True

```

matrix_from_columns(columns)

Return the matrix constructed from self using columns with indices in the columns list.

EXAMPLES:

```

sage: M = MatrixSpace(Integers(8), 3, 3)
sage: A = M(range(9)); A
[0 1 2]
[3 4 5]
[6 7 0]
sage: A.matrix_from_columns([2,1])
[2 1]

```

```
[5 4]
[0 7]
```

matrix_from_rows (*rows*)

Return the matrix constructed from self using rows with indices in the rows list.

EXAMPLES:

```
sage: M = MatrixSpace(Integers(8), 3, 3)
sage: A = M(range(9)); A
[0 1 2]
[3 4 5]
[6 7 0]
sage: A.matrix_from_rows([2, 1])
[6 7 0]
[3 4 5]
```

matrix_from_rows_and_columns (*rows*, *columns*)

Return the matrix constructed from self from the given rows and columns.

EXAMPLES:

```
sage: M = MatrixSpace(Integers(8), 3, 3)
sage: A = M(range(9)); A
[0 1 2]
[3 4 5]
[6 7 0]
sage: A.matrix_from_rows_and_columns([1], [0, 2])
[3 5]
sage: A.matrix_from_rows_and_columns([1, 2], [1, 2])
[4 5]
[7 0]
```

Note that row and column indices can be reordered or repeated:

```
sage: A.matrix_from_rows_and_columns([2, 1], [2, 1])
[0 7]
[5 4]
```

For example here we take from row 1 columns 2 then 0 twice, and do this 3 times.

```
sage: A.matrix_from_rows_and_columns([1, 1, 1], [2, 0, 0])
[5 3 3]
[5 3 3]
[5 3 3]
```

AUTHORS:

- Jaap Spies (2006-02-18)
- Didier Deshommes: some Pyrex speedups implemented

matrix_over_field ()

Return copy of this matrix, but with entries viewed as elements of the fraction field of the base ring (assuming it is defined).

EXAMPLES:

```
sage: A = MatrixSpace(IntegerRing(), 2) ([1, 2, 3, 4])
sage: B = A.matrix_over_field()
sage: B
[1 2]
```

```
[3 4]
sage: B.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
```

matrix_space (*nrows=None, ncols=None, sparse=None*)

Return the ambient matrix space of self.

INPUT:

- *nrows, ncols* - (optional) number of rows and columns in returned matrix space.
- *sparse* - whether the returned matrix space uses sparse or dense matrices.

EXAMPLES:

```
sage: m = matrix(3, [1..9])
sage: m.matrix_space()
Full MatrixSpace of 3 by 3 dense matrices over Integer Ring
sage: m.matrix_space(ncols=2)
Full MatrixSpace of 3 by 2 dense matrices over Integer Ring
sage: m.matrix_space(1)
Full MatrixSpace of 1 by 3 dense matrices over Integer Ring
sage: m.matrix_space(1, 2, True)
Full MatrixSpace of 1 by 2 sparse matrices over Integer Ring
```

new_matrix (*nrows=None, ncols=None, entries=None, coerce=True, copy=True, sparse=None*)

Create a matrix in the parent of this matrix with the given number of rows, columns, etc. The default parameters are the same as for self.

INPUT:

These three variables get sent to `matrix_space()`:

- *nrows, ncols* - number of rows and columns in returned matrix. If not specified, defaults to None and will give a matrix of the same size as self.
- *sparse* - whether returned matrix is sparse or not. Defaults to same value as self.

The remaining three variables (*coerce*, *entries*, and *copy*) are used by `sage.matrix.matrix_space.MatrixSpace()` to construct the new matrix.

Warning: This function called with no arguments returns the zero matrix of the same dimension and sparseness of self.

EXAMPLES:

```
sage: A = matrix(ZZ, 2, 2, [1, 2, 3, 4]); A
[1 2]
[3 4]
sage: A.new_matrix()
[0 0]
[0 0]
sage: A.new_matrix(1, 1)
[0]
sage: A.new_matrix(3, 3).parent()
Full MatrixSpace of 3 by 3 dense matrices over Integer Ring

sage: A = matrix(RR, 2, 3, [1.1, 2.2, 3.3, 4.4, 5.5, 6.6]); A
[1.10000000000000 2.20000000000000 3.30000000000000]
[4.40000000000000 5.50000000000000 6.60000000000000]
sage: A.new_matrix()
```

```
[0.0000000000000000 0.0000000000000000 0.0000000000000000]
[0.0000000000000000 0.0000000000000000 0.0000000000000000]
sage: A.new_matrix().parent()
Full MatrixSpace of 2 by 3 dense matrices over Real Field with 53 bits of precision
```

numpy (*dtype=None*)

Return the Numpy matrix associated to this matrix.

INPUT:

- *dtype* - The desired data-type for the array. If not given, then the type will be determined as the minimum type required to hold the objects in the sequence.

EXAMPLES:

```
sage: a = matrix(3, range(12))
sage: a.numpy()
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
sage: a.numpy('f')
array([[ 0.,  1.,  2.,  3.],
       [ 4.,  5.,  6.,  7.],
       [ 8.,  9., 10., 11.]], dtype=float32)
sage: a.numpy('d')
array([[ 0.,  1.,  2.,  3.],
       [ 4.,  5.,  6.,  7.],
       [ 8.,  9., 10., 11.]])
sage: a.numpy('B')
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]], dtype=uint8)
```

Type `numpy.typecodes` for a list of the possible typecodes:

```
sage: import numpy
sage: sorted(numpy.typecodes.items())
[('All', '?bhlqpBHILQPefdgFDGSUVOMm'), ('AllFloat', 'efdgFDG'), ('AllInteger', 'bBhHiIlLqQp
```

Alternatively, `numpy` automatically calls this function (via the magic `__array__()` method) to convert Sage matrices to numpy arrays:

```
sage: import numpy
sage: b=numpy.array(a); b
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
sage: b.dtype
dtype('int32') # 32-bit
dtype('int64') # 64-bit
sage: b.shape
(3, 4)
```

row (*i, from_list=False*)

Return the *i*'th row of this matrix as a vector.

This row is a dense vector if and only if the matrix is a dense matrix.

INPUT:

- *i* - integer

- `from_list` - bool (default: False); if true, returns the `i`'th element of `self.rows()` (see `rows()`), which may be faster, but requires building a list of all rows the first time it is called after an entry of the matrix is changed.

EXAMPLES:

```
sage: a = matrix(2,3,range(6)); a
[0 1 2]
[3 4 5]
sage: a.row(0)
(0, 1, 2)
sage: a.row(1)
(3, 4, 5)
sage: a.row(-1) # last row
(3, 4, 5)
```

TESTS:

```
sage: a = matrix(2,3,range(6)); a
[0 1 2]
[3 4 5]
sage: a.row(2)
Traceback (most recent call last):
...
IndexError: row index out of range
sage: a.row(-3)
Traceback (most recent call last):
...
IndexError: row index out of range
```

rows (*copy=True*)

Return a list of the rows of self.

INPUT:

- `copy` - (default: True) if True, return a copy of the list of rows which is safe to change.

If self is a sparse matrix, rows are returned as sparse vectors, otherwise returned vectors are dense.

EXAMPLES:

```
sage: matrix(3, [1..9]).rows()
[(1, 2, 3), (4, 5, 6), (7, 8, 9)]
sage: matrix(RR, 2, [sqrt(2), pi, exp(1), 0]).rows()
[(1.41421356237310, 3.14159265358979), (2.71828182845905, 0.000000000000000)]
sage: matrix(RR, 0, 2, []).rows()
[]
sage: matrix(RR, 2, 0, []).rows()
[(), ()]
sage: m = matrix(RR, 3, 3, {(1,2): pi, (2, 2): -1, (0,1): sqrt(2)})
sage: parent(m.rows()[0])
Sparse vector space of dimension 3 over Real Field with 53 bits of precision
```

Sparse matrices produce sparse rows.

```
sage: A = matrix(QQ, 2, range(4), sparse=True)
sage: v = A.rows()[0]
sage: v.is_sparse()
True
```

TESTS:

```
sage: A = matrix(QQ, 4, range(16))
sage: A.rows('junk')
Traceback (most recent call last):
...
ValueError: 'copy' must be True or False, not junk
```

set_column(col, v)

Sets the entries of column `col` to the entries of `v`.

INPUT:

- `col` - index of column to be set.
- `v` - a list or vector of the new entries.

OUTPUT:

Changes the matrix in-place, so there is no output.

EXAMPLES:

New entries may be contained in a vector.:

```
sage: A = matrix(QQ, 5, range(25))
sage: u = vector(QQ, [0, -1, -2, -3, -4])
sage: A.set_column(2, u)
sage: A
[ 0  1  0  3  4]
[ 5  6 -1  8  9]
[10 11 -2 13 14]
[15 16 -3 18 19]
[20 21 -4 23 24]
```

New entries may be in any sort of list.:

```
sage: A = matrix([[1, 2], [3, 4]]); A
[1 2]
[3 4]
sage: A.set_column(0, [0, 0]); A
[0 2]
[0 4]
sage: A.set_column(1, (0, 0)); A
[0 0]
[0 0]
```

TESTS:

```
sage: A = matrix([[1, 2], [3, 4]])
sage: A.set_column(2, [0, 0]); A
Traceback (most recent call last):
...
ValueError: column number must be between 0 and 1 (inclusive), not 2

sage: A.set_column(0, [0, 0, 0])
Traceback (most recent call last):
...
ValueError: list of new entries must be of length 2 (not 3)

sage: A = matrix(2, [1, 2, 3, 4])
sage: A.set_column(0, [1/4, 1]); A
Traceback (most recent call last):
```

```
...
TypeError: Cannot set column with Rational Field elements over Integer Ring, use change_ring
```

set_row(row, v)

Sets the entries of row `row` to the entries of `v`.

INPUT:

- `row` - index of row to be set.
- `v` - a list or vector of the new entries.

OUTPUT:

Changes the matrix in-place, so there is no output.

EXAMPLES:

New entries may be contained in a vector.:

```
sage: A = matrix(QQ, 5, range(25))
sage: u = vector(QQ, [0, -1, -2, -3, -4])
sage: A.set_row(2, u)
sage: A
[ 0  1  2  3  4]
[ 5  6  7  8  9]
[ 0 -1 -2 -3 -4]
[15 16 17 18 19]
[20 21 22 23 24]
```

New entries may be in any sort of list.:

```
sage: A = matrix([[1, 2], [3, 4]]); A
[1 2]
[3 4]
sage: A.set_row(0, [0, 0]); A
[0 0]
[3 4]
sage: A.set_row(1, (0, 0)); A
[0 0]
[0 0]
```

TESTS:

```
sage: A = matrix([[1, 2], [3, 4]])
sage: A.set_row(2, [0, 0]); A
Traceback (most recent call last):
...
ValueError: row number must be between 0 and 1 (inclusive), not 2
```

```
sage: A.set_row(0, [0, 0, 0])
Traceback (most recent call last):
...
ValueError: list of new entries must be of length 2 (not 3)
```

```
sage: A = matrix(2, [1, 2, 3, 4])
sage: A.set_row(0, [1/3, 1]); A
Traceback (most recent call last):
...
TypeError: Cannot set row with Rational Field elements over Integer Ring, use change_ring fi
```

sparse_columns (*copy=True*)Return a list of the columns of `self` as sparse vectors (or free module elements).

INPUT:

•copy - (default: True) if True, return a copy so you can modify it safely

EXAMPLES:

```
sage: a = matrix(2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: v = a.sparse_columns(); v
[(0, 3), (1, 4), (2, 5)]
sage: v[1].is_sparse()
True
```

TESTS:

Columns of sparse matrices having no columns were fixed on [trac ticket #10714](#):

```
sage: m = matrix(10, 0, sparse=True)
sage: m.ncols()
0
sage: m.columns()
[]
```

Check that the returned columns are immutable as per [trac ticket #14874](#):

```
sage: m = Mat(ZZ, 3, 3, sparse=True)(range(9))
sage: v = m.sparse_columns()
sage: map(lambda x: x.is_mutable(), v)
[False, False, False]
```

sparse_matrix ()

If this matrix is dense, return a sparse matrix with the same entries. If this matrix is sparse, return this matrix (not a copy).

Note: The definition of “dense” and “sparse” in Sage have nothing to do with the number of nonzero entries. Sparse and dense are properties of the underlying representation of the matrix.

EXAMPLES:

```
sage: A = MatrixSpace(QQ, 2, sparse=False)([1, 2, 0, 1])
sage: A.is_sparse()
False
sage: B = A.sparse_matrix()
sage: B.is_sparse()
True
sage: A
[1 2]
[0 1]
sage: B
[1 2]
[0 1]
sage: A*B
[1 4]
[0 1]
sage: A.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: B.parent()
```



```
Full MatrixSpace of 2 by 2 sparse matrices over Rational Field
sage: (A*B).parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
sage: (B*A).parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field
```

sparse_rows (*copy=True*)

Return a list of the rows of `self` as sparse vectors (or free module elements).

INPUT:

- **copy** - (default: **True**) if **True**, return a copy so you can modify it safely

EXAMPLES:

```
sage: m = Mat(ZZ, 3, 3, sparse=True)(range(9)); m
[0 1 2]
[3 4 5]
[6 7 8]
sage: v = m.sparse_rows(); v
[(0, 1, 2), (3, 4, 5), (6, 7, 8)]
sage: m.sparse_rows(copy=False) is m.sparse_rows(copy=False)
True
sage: v[1].is_sparse()
True
sage: m[0,0] = 10
sage: m.sparse_rows()
[(10, 1, 2), (3, 4, 5), (6, 7, 8)]
```

TESTS:

Rows of sparse matrices having no rows were fixed on [trac ticket #10714](#):

```
sage: m = matrix(0, 10, sparse=True)
sage: m.nrows()
0
sage: m.rows()
[]
```

Check that the returned rows are immutable as per [trac ticket #14874](#):

```
sage: m = Mat(ZZ, 3, 3, sparse=True)(range(9))
sage: v = m.sparse_rows()
sage: map(lambda x: x.is_mutable(), v)
[False, False, False]
```

stack (*bottom, subdivide=False*)

Returns a new matrix formed by appending the matrix (or vector) `bottom` beneath `self`.

INPUT:

- **bottom** - a matrix, vector or free module element, whose dimensions are compatible with `self`.
- **subdivide** - default: **False** - request the resulting matrix to have a new subdivision, separating `self` from `bottom`.

OUTPUT:

A new matrix formed by appending `bottom` beneath `self`. If `bottom` is a vector (or free module element) then in this context it is appropriate to consider it as a row vector. (The code first converts a vector to a 1-row matrix.)

If `subdivide` is `True` then any row subdivisions for the two matrices are preserved, and a new subdivision is added between `self` and `bottom`. If the column divisions are identical, then they are preserved, otherwise they are discarded. When `subdivide` is `False` there is no subdivision information in the result.

Warning: If `subdivide` is `True` then unequal column subdivisions will be discarded, since it would be ambiguous how to interpret them. If the subdivision behavior is not what you need, you can manage subdivisions yourself with methods like `subdivisions()` and `subdivide()`. You might also find `block_matrix()` or `block_diagonal_matrix()` useful and simpler in some instances.

EXAMPLES:

Stacking with a matrix.

```
sage: A = matrix(QQ, 4, 3, range(12))
sage: B = matrix(QQ, 3, 3, range(9))
sage: A.stack(B)
[ 0  1  2]
[ 3  4  5]
[ 6  7  8]
[ 9 10 11]
[ 0  1  2]
[ 3  4  5]
[ 6  7  8]
```

Stacking with a vector.

```
sage: A = matrix(QQ, 3, 2, [0, 2, 4, 6, 8, 10])
sage: v = vector(QQ, 2, [100, 200])
sage: A.stack(v)
[ 0  2]
[ 4  6]
[ 8 10]
[100 200]
```

Errors are raised if the sizes are incompatible.

```
sage: A = matrix(RR, [[1, 2], [3, 4]])
sage: B = matrix(RR, [[10, 20, 30], [40, 50, 60]])
sage: A.stack(B)
Traceback (most recent call last):
...
TypeError: number of columns must be the same, 2 != 3

sage: v = vector(RR, [100, 200, 300])
sage: A.stack(v)
Traceback (most recent call last):
...
TypeError: number of columns must be the same, 2 != 3
```

Setting `subdivide` to `True` will, in its simplest form, add a subdivision between `self` and `bottom`.

```
sage: A = matrix(QQ, 2, 5, range(10))
sage: B = matrix(QQ, 3, 5, range(15))
sage: A.stack(B, subdivide=True)
[ 0  1  2  3  4]
[ 5  6  7  8  9]
[-----]
[ 0  1  2  3  4]
```

```
[ 5  6  7  8  9]
[10 11 12 13 14]
```

Row subdivisions are preserved by stacking, and enriched, if subdivisions are requested. (So multiple stackings can be recorded.)

```
sage: A = matrix(QQ, 2, 4, range(8))
sage: A.subdivide([1], None)
sage: B = matrix(QQ, 3, 4, range(12))
sage: B.subdivide([2], None)
sage: A.stack(B, subdivide=True)
[ 0  1  2  3]
[-----]
[ 4  5  6  7]
[-----]
[ 0  1  2  3]
[ 4  5  6  7]
[-----]
[ 8  9 10 11]
```

Column subdivisions can be preserved, but only if they are identical. Otherwise, this information is discarded and must be managed separately.

```
sage: A = matrix(QQ, 2, 5, range(10))
sage: A.subdivide(None, [2,4])
sage: B = matrix(QQ, 3, 5, range(15))
sage: B.subdivide(None, [2,4])
sage: A.stack(B, subdivide=True)
[ 0  1| 2  3| 4]
[ 5  6| 7  8| 9]
[-----+-----+--]
[ 0  1| 2  3| 4]
[ 5  6| 7  8| 9]
[10 11|12 13|14]
```

```
sage: A.subdivide(None, [1,2])
sage: A.stack(B, subdivide=True)
[ 0  1  2  3  4]
[ 5  6  7  8  9]
[-----]
[ 0  1  2  3  4]
[ 5  6  7  8  9]
[10 11 12 13 14]
```

The result retains the base ring of `self` by coercing the elements of `bottom` into the base ring of `self`.

```
sage: A = matrix(QQ, 1, 2, [1,2])
sage: B = matrix(RR, 1, 2, [sin(1.1), sin(2.2)])
sage: C = A.stack(B); C
[
1                               2]
[183017397/205358938 106580492/131825561]
sage: C.parent()
Full MatrixSpace of 2 by 2 dense matrices over Rational Field

sage: D = B.stack(A); D
[0.891207360061435 0.808496403819590]
[ 1.000000000000000  2.000000000000000]
sage: D.parent()
Full MatrixSpace of 2 by 2 dense matrices over Real Field with 53 bits of precision
```

Sometimes it is not possible to coerce into the base ring of `self`. A solution is to change the base ring of `self` to a more expansive ring. Here we mix the rationals with a ring of polynomials with rational coefficients.

```
sage: R = PolynomialRing(QQ, 'y')
sage: A = matrix(QQ, 1, 2, [1,2])
sage: B = matrix(R, 1, 2, ['y', 'y^2'])

sage: C = B.stack(A); C
[  y y^2]
[  1  2]
sage: C.parent()
Full MatrixSpace of 2 by 2 dense matrices over Univariate Polynomial Ring in y over Rational

sage: D = A.stack(B)
Traceback (most recent call last):
...
TypeError: not a constant polynomial

sage: E = A.change_ring(R)
sage: F = E.stack(B); F
[  1  2]
[  y y^2]
sage: F.parent()
Full MatrixSpace of 2 by 2 dense matrices over Univariate Polynomial Ring in y over Rational
```

TESTS:

A legacy test from the original implementation.

```
sage: M = Matrix(QQ, 2, 3, range(6))
sage: N = Matrix(QQ, 1, 3, [10,11,12])
sage: M.stack(N)
[ 0  1  2]
[ 3  4  5]
[10 11 12]
```

AUTHOR:

- Rob Beezer (2011-03-19) - rewritten to mirror code for `augment()`

submatrix (*row=0, col=0, nrows=-1, ncols=-1*)

Return the matrix constructed from `self` using the specified range of rows and columns.

INPUT:

- `row, col` - index of the starting row and column. Indices start at zero.
- `nrows, ncols` - (optional) number of rows and columns to take. If not provided, take all rows below and all columns to the right of the starting entry.

SEE ALSO:

The functions `matrix_from_rows()`, `matrix_from_columns()`, and `matrix_from_rows_and_columns()` allow one to select arbitrary subsets of rows and/or columns.

EXAMPLES:

Take the 3×3 submatrix starting from entry (1,1) in a 4×4 matrix:

```
sage: m = matrix(4, [1..16])
sage: m.submatrix(1, 1)
```

```
[ 6  7  8]
[10 11 12]
[14 15 16]
```

Same thing, except take only two rows:

```
sage: m.submatrix(1, 1, 2)
[ 6  7  8]
[10 11 12]
```

And now take only one column:

```
sage: m.submatrix(1, 1, 2, 1)
[ 6]
[10]
```

You can take zero rows or columns if you want:

```
sage: m.submatrix(1, 1, 0)
[]
sage: parent(m.submatrix(1, 1, 0))
Full MatrixSpace of 0 by 3 dense matrices over Integer Ring
```


BASE CLASS FOR MATRICES, PART 2

Base class for matrices, part 2

For design documentation see `matrix/docs.py`.

AUTHORS:

- William Stein: initial version
- Miguel Marco (2010-06-19): modified eigenvalues and eigenvectors functions to allow the option `extend=False`
- Rob Beezer (2011-02-05): refactored all of the matrix kernel routines

TESTS:

```
sage: m = matrix(ZZ['x'], 2, 3, [1..6])
sage: TestSuite(m).run()
```

```
class sage.matrix.matrix2.Matrix
    Bases: sage.matrix.matrix1.Matrix
```

The initialization routine of the `Matrix` base class ensures that it sets the attributes `self._parent`, `self._base_ring`, `self._nrows`, `self._ncols`. It sets the latter ones by accessing the relevant information on `parent`, which is often slower than what a more specific subclass can do.

Subclasses of `Matrix` can safely skip calling `Matrix.__init__` provided they take care of initializing these attributes themselves.

The private attributes `self._is_immutable` and `self._cache` are implicitly initialized to valid values upon memory allocation.

EXAMPLES:

```
sage: import sage.matrix.matrix0
sage: A = sage.matrix.matrix0.Matrix(MatrixSpace(QQ, 2))
sage: type(A)
<type 'sage.matrix.matrix0.Matrix'>
```

C

Returns the conjugate matrix.

EXAMPLE:

```
sage: A = matrix(QQbar, [[ -3, 5 - 3*I, 7 - 4*I],
...                      [7 + 3*I, -1 + 6*I, 3 + 5*I],
...                      [3 + 3*I, -3 + 6*I, 5 + I]])
sage: A.C
[ -3 5 + 3*I 7 + 4*I]
```

```
[ 7 - 3*I -1 - 6*I  3 - 5*I]
[ 3 - 3*I -3 - 6*I  5 - 1*I]
```

H

Returns the conjugate-transpose (Hermitian) matrix.

EXAMPLE:

```
sage: A = matrix(QQbar, [[ -3,  5 - 3*I, 7 - 4*I],
...                       [7 + 3*I, -1 + 6*I, 3 + 5*I],
...                       [3 + 3*I, -3 + 6*I, 5 +   I]])
sage: A.H
[      -3  7 - 3*I  3 - 3*I]
[ 5 + 3*I -1 - 6*I -3 - 6*I]
[ 7 + 4*I  3 - 5*I  5 - 1*I]
```

I

Returns the inverse of the matrix, if it exists.

EXAMPLES:

```
sage: A = matrix(QQ, [[-5, -3, -1, -7],
...                  [ 1,  1,  1,  0],
...                  [-1, -2, -2,  0],
...                  [-2, -1,  0, -4]])
sage: A.I
[ 0  2  1  0]
[-4 -8 -2  7]
[ 4  7  1 -7]
[ 1  1  0 -2]

sage: B = matrix(QQ, [[-11, -5, 18, -6],
...                  [ 1,  2, -6,  8],
...                  [-4, -2,  7, -3],
...                  [ 1, -2,  5, -11]])
sage: B.I
Traceback (most recent call last):
...
ZeroDivisionError: input matrix must be nonsingular
```

LU (*pivot=None, format='plu'*)

Finds a decomposition into a lower-triangular matrix and an upper-triangular matrix.

INPUT:

- **pivot** - pivoting strategy
 - ‘auto’ (default) - see if the matrix entries are ordered (i.e. if they have an absolute value method), and if so, use a the partial pivoting strategy. Otherwise, fall back to the nonzero strategy. This is the best choice for general routines that may call this for matrix entries of a variety of types.
 - ‘partial’ - each column is examined for the element with the largest absolute value and the row containing this element is swapped into place.
 - ‘nonzero’ - the first nonzero element in a column is located and the row with this element is used.
- **format** - contents of output, see more discussion below about output.
 - ‘plu’ (default) - a triple; matrices P, L and U such that $A = P * L * U$.
 - ‘compact’ - a pair; row permutation as a tuple, and the matrices L and U combined into one matrix.

OUTPUT:

Suppose that A is an $m \times n$ matrix, then an LU decomposition is a lower-triangular $m \times m$ matrix L with every diagonal element equal to 1, and an upper-triangular $m \times n$ matrix, U such that the product LU , after a permutation of the rows, is then equal to A . For the ‘plu’ format the permutation is returned as an $m \times m$ permutation matrix P such that

$$A = PLU$$

It is more common to place the permutation matrix just to the left of A . If you desire this version, then use the inverse of P which is computed most efficiently as its transpose.

If the ‘partial’ pivoting strategy is used, then the non-diagonal entries of L will be less than or equal to 1 in absolute value. The ‘nonzero’ pivot strategy may be faster, but the growth of data structures for elements of the decomposition might counteract the advantage.

By necessity, returned matrices have a base ring equal to the fraction field of the base ring of the original matrix.

In the ‘compact’ format, the first returned value is a tuple that is a permutation of the rows of LU that yields A . See the doctest for how you might employ this permutation. Then the matrices L and U are merged into one matrix – remove the diagonal of ones in L and the remaining nonzero entries can replace the entries of U beneath the diagonal.

The results are cached, only in the compact format, separately for each pivot strategy called. Repeated requests for the ‘plu’ format will require just a small amount of overhead in each call to bust out the compact format to the three matrices. Since only the compact format is cached, the components of the compact format are immutable, while the components of the ‘plu’ format are regenerated, and hence are mutable.

Notice that while U is similar to row-echelon form and the rows of U span the row space of A , the rows of U are not generally linearly independent. Nor are the pivot columns (or rank) immediately obvious. However for rings without specialized echelon form routines, this method is about twice as fast as the generic echelon form routine since it only acts “below the diagonal”, as would be predicted from a theoretical analysis of the algorithms.

Note: This is an exact computation, so limited to exact rings. If you need numerical results, convert the base ring to the field of real double numbers, `RDF` or the field of complex double numbers, `CDF`, which will use a faster routine that is careful about numerical subtleties.

ALGORITHM:

“Gaussian Elimination with Partial Pivoting,” Algorithm 21.1 of [\[TREFETHEN-BAU\]](#).

EXAMPLES:

Notice the difference in the L matrix as a result of different pivoting strategies. With partial pivoting, every entry of L has absolute value 1 or less.

```
sage: A = matrix(QQ, [[1, -1, 0, 2, 4, 7, -1],
...                  [2, -1, 0, 6, 4, 8, -2],
...                  [2, 0, 1, 4, 2, 6, 0],
...                  [1, 0, -1, 8, -1, -1, -3],
...                  [1, 1, 2, -2, -1, 1, 3]])
sage: P, L, U = A.LU(pivot='partial')
sage: P
[0 0 0 0 1]
[1 0 0 0 0]
[0 0 0 1 0]
[0 0 1 0 0]
```

```

[0 1 0 0 0]
sage: L
[ 1 0 0 0 0]
[ 1/2 1 0 0 0]
[ 1/2 1/3 1 0 0]
[ 1 2/3 1/5 1 0]
[ 1/2 -1/3 -2/5 0 1]
sage: U
[ 2 -1 0 6 4 8 -2]
[ 0 3/2 2 -5 -3 -3 4]
[ 0 0 -5/3 20/3 -2 -4 -10/3]
[ 0 0 0 0 0 2/5 4/5 0]
[ 0 0 0 0 0 1/5 2/5 0]
sage: A == P*L*U
True
sage: P, L, U = A.LU(pivot='nonzero')
sage: P
[1 0 0 0 0]
[0 1 0 0 0]
[0 0 1 0 0]
[0 0 0 1 0]
[0 0 0 0 1]
sage: L
[ 1 0 0 0 0]
[ 2 1 0 0 0]
[ 2 2 1 0 0]
[ 1 1 -1 1 0]
[ 1 2 2 0 1]
sage: U
[ 1 -1 0 2 4 7 -1]
[ 0 1 0 2 -4 -6 0]
[ 0 0 1 -4 2 4 2]
[ 0 0 0 0 1 2 0]
[ 0 0 0 0 -1 -2 0]
sage: A == P*L*U
True

```

An example of the compact format.

```

sage: B = matrix(QQ, [[ 1, 3, 5, 5],
...                  [ 1, 4, 7, 8],
...                  [-1, -4, -6, -6],
...                  [ 0, -2, -5, -8],
...                  [-2, -6, -6, -2]])
sage: perm, M = B.LU(format='compact')
sage: perm
(4, 3, 0, 1, 2)
sage: M
[ -2 -6 -6 -2]
[ 0 -2 -5 -8]
[-1/2 0 2 4]
[-1/2 -1/2 3/4 0]
[ 1/2 1/2 -1/4 0]

```

We can easily illustrate the relationships between the two formats with a square matrix.

```

sage: C = matrix(QQ, [[-2, 3, -2, -5],
...                  [ 1, -2, 1, 3],
...                  [-4, 7, -3, -8],

```

```

...          [-3, 8, -1, -5]])
sage: P, L, U = C.LU(format='plu')
sage: perm, M = C.LU(format='compact')
sage: (L - identity_matrix(4)) + U == M
True
sage: p = [perm[i]+1 for i in range(len(perm))]
sage: PP = Permutation(p).to_matrix()
sage: PP == P
True

```

For a nonsingular matrix, and the ‘nonzero’ pivot strategy there is no need to permute rows, so the permutation matrix will be the identity. Furthermore, it can be shown that then the L and U matrices are uniquely determined by requiring L to have ones on the diagonal.

```

sage: D = matrix(QQ, [[ 1, 0, 2, 0, -2, -1],
...                  [ 3, -2, 3, -1, 0, 6],
...                  [-4, 2, -3, 1, -1, -8],
...                  [-2, 2, -3, 2, 1, 0],
...                  [ 0, -1, -1, 0, 2, 5],
...                  [-1, 2, -4, -1, 5, -3]])
sage: P, L, U = D.LU(pivot='nonzero')
sage: P
[1 0 0 0 0 0]
[0 1 0 0 0 0]
[0 0 1 0 0 0]
[0 0 0 1 0 0]
[0 0 0 0 1 0]
[0 0 0 0 0 1]
sage: L
[ 1 0 0 0 0 0]
[ 3 1 0 0 0 0]
[-4 -1 1 0 0 0]
[-2 -1 -1 1 0 0]
[ 0 1/2 1/4 1/2 1 0]
[-1 -1 -5/2 -2 -6 1]
sage: U
[ 1 0 2 0 -2 -1]
[ 0 -2 -3 -1 6 9]
[ 0 0 2 0 -3 -3]
[ 0 0 0 1 0 4]
[ 0 0 0 0 -1/4 -3/4]
[ 0 0 0 0 0 1]
sage: D == L*U
True

```

The base ring of the matrix may be any field, or a ring which has a fraction field implemented in Sage. The ring needs to be exact (there is a numerical LU decomposition for matrices over RDF and CDF). Matrices returned are over the original field, or the fraction field of the ring. If the field is not ordered (i.e. the absolute value function is not implemented), then the pivot strategy needs to be ‘nonzero’.

```

sage: A = matrix(RealField(100), 3, 3, range(9))
sage: P, L, U = A.LU()
Traceback (most recent call last):
...
TypeError: base ring of the matrix must be exact, not Real Field with 100 bits of precision

sage: A = matrix(Integers(6), 3, 2, range(6))
sage: A.LU()
Traceback (most recent call last):

```

```

...
TypeError: base ring of the matrix needs a field of fractions, not Ring of integers modulo 6

sage: R.<y> = PolynomialRing(QQ, 'y')
sage: B = matrix(R, [[y+1, y^2+y], [y^2, y^3]])
sage: P, L, U = B.LU(pivot='partial')
Traceback (most recent call last):
...
TypeError: cannot take absolute value of matrix entries, try 'pivot=nonzero'
sage: P, L, U = B.LU(pivot='nonzero')
sage: P
[1 0]
[0 1]
sage: L
[          1          0]
[y^2/(y + 1)      1]
sage: U
[ y + 1 y^2 + y]
[      0      0]
sage: L.base_ring()
Fraction Field of Univariate Polynomial Ring in y over Rational Field
sage: B == P*L*U
True

sage: F.<a> = FiniteField(5^2)
sage: C = matrix(F, [[a + 3, 4*a + 4, 2, 4*a + 2],
...                  [3, 2*a + 4, 2*a + 4, 2*a + 1],
...                  [3*a + 1, a + 3, 2*a + 4, 4*a + 3],
...                  [a, 3, 3*a + 1, a]])
sage: P, L, U = C.LU(pivot='nonzero')
sage: P
[1 0 0 0]
[0 1 0 0]
[0 0 1 0]
[0 0 0 1]
sage: L
[          1          0          0          0]
[3*a + 3          1          0          0]
[          2*a 4*a + 2          1          0]
[2*a + 3          2 2*a + 4          1]
sage: U
[ a + 3 4*a + 4          2 4*a + 2]
[          0 a + 1 a + 3 2*a + 4]
[          0          0          1 4*a + 2]
[          0          0          0          0]
sage: L.base_ring()
Finite Field in a of size 5^2
sage: C == P*L*U
True

```

With no pivoting strategy given (i.e. `pivot=None`) the routine will try to use partial pivoting, but then fall back to the nonzero strategy. For the nonsingular matrix below, we see evidence of pivoting when viewed over the rationals, and no pivoting over the integers mod 29.

```

sage: entries = [3, 20, 11, 7, 16, 28, 5, 15, 21, 23, 22, 18, 8, 23, 15, 2]
sage: A = matrix(Integers(29), 4, 4, entries)
sage: perm, _ = A.LU(format='compact'); perm
(0, 1, 2, 3)
sage: B = matrix(QQ, 4, 4, entries)

```

```
sage: perm, _ = B.LU(format='compact'); perm
(2, 0, 1, 3)
```

The U matrix is only guaranteed to be upper-triangular. The rows are not necessarily linearly independent, nor are the pivots columns or rank in evidence.

```
sage: A = matrix(QQ, [[ 1, -4, 1, 0, -2, 1, 3, 3, 2],
...                  [-1, 4, 0, -4, 0, -4, 5, -7, -7],
...                  [ 0, 0, 1, -4, -1, -3, 6, -5, -6],
...                  [-2, 8, -1, -4, 2, -4, 1, -8, -7],
...                  [ 1, -4, 2, -4, -3, 2, 5, 6, 4]])
sage: P, L, U = A.LU()
sage: U
[ -2      8     -1     -4      2     -4      1     -8     -7]
[  0      0    1/2     -2     -1     -2     9/2     -3    -7/2]
[  0      0    3/2     -6     -2      0    11/2      2     1/2]
[  0      0      0      0    -1/3     -1     5/3     -5/3    -5/3]
[  0      0      0      0     1/3     -3     7/3    -19/3   -19/3]
sage: A.rref()
[ 1 -4  0  4  0  0 -1 -1 -1]
[ 0  0  1 -4  0  0  1  0 -1]
[ 0  0  0  0  1  0 -2 -1 -1]
[ 0  0  0  0  0  1 -1  2  2]
[ 0  0  0  0  0  0  0  0  0]
sage: A.pivots()
(0, 2, 4, 5)
```

TESTS:

Unknown keywords are caught.

```
sage: A = matrix(ZZ, 2, range(4))
sage: A.LU(pivot='junk')
Traceback (most recent call last):
...
ValueError: pivot strategy must be None, 'partial' or 'nonzero', not junk
sage: A.LU(format='garbage')
Traceback (most recent call last):
...
ValueError: format must be 'plu' or 'compact', not garbage
```

Components of the 'compact' format are immutable, while components of the 'plu' format are not.

```
sage: A = matrix(ZZ, 2, range(4))
sage: perm, M = A.LU(format='compact')
sage: perm[0] = 25
Traceback (most recent call last):
...
TypeError: 'tuple' object does not support item assignment
sage: M.is_immutable()
True
sage: P, L, U = A.LU(format='plu')
sage: all([A.is_mutable() for A in [P, L, U]])
True
```

Partial pivoting is based on the absolute values of entries of a column. Trac #12208 shows that the return value of the absolute value must be handled carefully. This tests that situation in the case of cyclotomic fields.

```
sage: C = SymmetricGroup(5).character_table()
sage: C.base_ring()
Cyclotomic Field of order 1 and degree 1
sage: P, L, U = C.LU(pivot='partial')
sage: C == P*L*U
True
```

AUTHOR:

•Rob Beezer (2011-04-26)

N (*prec=None, digits=None, algorithm=None*)

Return a numerical approximation of `self` as either a real or complex number with at least the requested number of bits or digits of precision.

INPUT:

- `prec` - an integer: the number of bits of precision
- `digits` - an integer: digits of precision

OUTPUT: A matrix coerced to a real or complex field with `prec` bits of precision.

EXAMPLES:

```
sage: d = matrix([[3, 0], [0, sqrt(2)]]) ;
sage: b = matrix([[1, -1], [2, 2]]) ; e = b * d * b.inverse(); e
[ 1/2*sqrt(2) + 3/2 -1/4*sqrt(2) + 3/4]
[      -sqrt(2) + 3  1/2*sqrt(2) + 3/2]
```

```
sage: e.numerical_approx(53)
[ 2.20710678118655 0.396446609406726]
[ 1.58578643762690 2.20710678118655]
```

```
sage: e.numerical_approx(20)
[ 2.2071 0.39645]
[ 1.5858 2.2071]
```

```
sage: (e-I).numerical_approx(20)
[2.2071 - 1.0000*I      0.39645]
[      1.5858 2.2071 - 1.0000*I]
```

```
sage: M=matrix(QQ,4,[i/(i+1) for i in range(12)]);M
[ 0  1/2  2/3]
[ 3/4  4/5  5/6]
[ 6/7  7/8  8/9]
[ 9/10 10/11 11/12]
```

```
sage: M.numerical_approx()
[0.000000000000000 0.500000000000000 0.666666666666667]
[0.750000000000000 0.800000000000000 0.833333333333333]
[0.857142857142857 0.875000000000000 0.888888888888889]
[0.900000000000000 0.909090909090909 0.916666666666667]
```

```
sage: matrix(SR, 2, 2, range(4)).n()
[0.000000000000000 1.000000000000000]
[ 2.000000000000000 3.000000000000000]
```

```
sage: numerical_approx(M)
[0.000000000000000 0.500000000000000 0.666666666666667]
[0.750000000000000 0.800000000000000 0.833333333333333]
```

```
[0.857142857142857 0.875000000000000 0.888888888888889]
[0.900000000000000 0.909090909090909 0.916666666666667]
```

QR (*full=True*)

Returns a factorization of `self` as a unitary matrix and an upper-triangular matrix.

INPUT:

- `full` - default: `True` - if `True` then the returned matrices have dimensions as described below. If `False` the `R` matrix has no zero rows and the columns of `Q` are a basis for the column space of `self`.

OUTPUT:

If `self` is an $m \times n$ matrix and `full=True` then this method returns a pair of matrices: `Q` is an $m \times m$ unitary matrix (meaning its inverse is its conjugate-transpose) and `R` is an $m \times n$ upper-triangular matrix with non-negative entries on the diagonal. For a matrix of full rank this factorization is unique (due to the restriction to positive entries on the diagonal).

If `full=False` then `Q` has m rows and the columns form an orthonormal basis for the column space of `self`. So, in particular, the conjugate-transpose of `Q` times `Q` will be an identity matrix. The matrix `R` will still be upper-triangular but will also have full rank, in particular it will lack the zero rows present in a full factorization of a rank-deficient matrix.

The results obtained when `full=True` are cached, hence `Q` and `R` are immutable matrices in this case.

Note: This is an exact computation, so limited to exact rings. Also the base ring needs to have a fraction field implemented in Sage and this field must contain square roots. One example is the field of algebraic numbers, `QQbar`, as used in the examples below. If you need numerical results, convert the base ring to the field of complex double numbers, `CDF`, which will use a faster routine that is careful about numerical subtleties.

ALGORITHM:

“Modified Gram-Schmidt,” Algorithm 8.1 of [TREFETHEN-BAU].

EXAMPLES:

For a nonsingular matrix, the QR decomposition is unique.

```
sage: A = matrix(QQbar, [[-2, 0, -4, -1, -1],
...                       [-2, 1, -6, -3, -1],
...                       [1, 1, 7, 4, 5],
...                       [3, 0, 8, 3, 3],
...                       [-1, 1, -6, -6, 5]])
sage: Q, R = A.QR()
sage: Q
[ -0.4588314677411235? -0.1260506983326509?  0.3812120831224489? -0.394573711338418?
[ -0.4588314677411235?  0.4726901187474409? -0.05198346588033394?  0.717294125164660?
[  0.2294157338705618?  0.6617661662464172?  0.6619227988762521? -0.180872093737548?
[  0.6882472016116853?  0.1890760474989764? -0.2044682991293135?  0.096630296654307?
[ -0.2294157338705618?  0.5357154679137663? -0.609939332995919? -0.536422031427112?
sage: R
[  4.358898943540674? -0.4588314677411235?  13.07669683062202?  6.194224814505168?  2.982
[                0  1.670171752907625?  0.5987408170800917? -1.292019657909672?  6.207
[                0                0  5.444401659866974?  5.468660610611130? -0.682
[                0                0                0  1.027626039419836? -3.61
[                0                0                0                0  0.02
sage: Q.conjugate_transpose()*Q
[1.000000000000000?  0.?e-18  0.?e-17  0.?e-15  0.?e
[  0.?e-18  1.000000000000000?  0.?e-16  0.?e-15  0.?e
```

```

[          0.?e-17          0.?e-16 1.000000000000000?          0.?e-15          0.?e-15
[          0.?e-15          0.?e-15          0.?e-15 1.000000000000000?          0.?e-15
[          0.?e-12          0.?e-12          0.?e-12          0.?e-12          1.000000000000000?
sage: Q*R == A
True

```

An example with complex numbers in `QQbar`, the field of algebraic numbers.

```

sage: A = matrix(QQbar, [[-8, 4*I + 1, -I + 2, 2*I + 1],
...                       [1, -2*I - 1, -I + 3, -I + 1],
...                       [I + 7, 2*I + 1, -2*I + 7, -I + 1],
...                       [I + 2, 0, I + 12, -1]])
sage: Q, R = A.QR()
sage: Q
[          -0.7302967433402215?          0.2070566455055649? + 0.5383472783144687?
[          0.0912870929175277? -0.2070566455055649? - 0.3778783780476559?
[ 0.6390096504226938? + 0.0912870929175277?*I 0.1708217325420910? + 0.6677576817554466?
[ 0.1825741858350554? + 0.0912870929175277?*I -0.03623491296347385? + 0.0724698259269477?
sage: R
[          10.95445115010333?          0.?e-18 - 1.917028951268082?*I
[          0          4.829596256417300? + 0.?e-17*I
[          0          0
[          0          0
sage: Q.conjugate_transpose()*Q
[1.000000000000000? + 0.?e-19*I          0.?e-18 + 0.?e-17*I          0.?e-17 + 0.?e-17*I
[          0.?e-18 + 0.?e-17*I 1.000000000000000? + 0.?e-17*I          0.?e-17 + 0.?e-17*I
[          0.?e-17 + 0.?e-17*I          0.?e-17 + 0.?e-17*I 1.000000000000000? + 0.?e-16*I
[          0.?e-16 + 0.?e-16*I          0.?e-16 + 0.?e-16*I          0.?e-16 + 0.?e-16*I
sage: Q*R - A
[          0.?e-17 0.?e-17 + 0.?e-17*I 0.?e-16 + 0.?e-16*I 0.?e-16 + 0.?e-16*I]
[          0.?e-18 0.?e-17 + 0.?e-17*I 0.?e-16 + 0.?e-16*I 0.?e-15 + 0.?e-15*I]
[0.?e-17 + 0.?e-18*I 0.?e-17 + 0.?e-17*I 0.?e-16 + 0.?e-16*I 0.?e-16 + 0.?e-16*I]
[0.?e-18 + 0.?e-18*I 0.?e-18 + 0.?e-17*I 0.?e-16 + 0.?e-16*I 0.?e-15 + 0.?e-16*I]

```

A rank-deficient rectangular matrix, with both values of the `full` keyword.

```

sage: A = matrix(QQbar, [[2, -3, 3],
...                       [-1, 1, -1],
...                       [-1, 3, -3],
...                       [-5, 1, -1]])
sage: Q, R = A.QR()
sage: Q
[ 0.3592106040535498? -0.5693261797050169? 0.7239227659930268? 0.1509015305256380?]
[ -0.1796053020267749? 0.1445907757980996? 0 0.9730546968377341?]
[ -0.1796053020267749? 0.7048800320157352? 0.672213996993525? -0.1378927778941174?]
[ -0.8980265101338745? -0.3976246334447737? 0.1551263069985058? -0.10667177157846818?]
sage: R
[ 5.567764362830022? -2.694079530401624? 2.694079530401624?]
[ 0 3.569584777515583? -3.569584777515583?]
[ 0 0 0]
[ 0 0 0]
sage: Q.conjugate_transpose()*Q
[ 1          0.?e-18          0.?e-18          0.?e-18]
[ 0.?e-18          1          0.?e-18          0.?e-18]
[ 0.?e-18          0.?e-18 1.000000000000000?          0.?e-18]
[ 0.?e-18          0.?e-18          0.?e-18 1.000000000000000?]
sage: Q, R = A.QR(full=False)
sage: Q

```



```
[ 0.3592106040535498? -0.5693261797050169?]
[-0.1796053020267749?  0.1445907757980996?]
[-0.1796053020267749?  0.7048800320157352?]
[-0.8980265101338745? -0.3976246334447737?]
sage: R
[ 5.567764362830022? -2.694079530401624?  2.694079530401624?]
[          0  3.569584777515583? -3.569584777515583?]
sage: Q.conjugate_transpose()*Q
[      1 0.?e-18]
[0.?e-18      1]
```

Another rank-deficient rectangular matrix, with complex entries, as a reduced decomposition.

```
sage: A = matrix(QQbar, [[-3*I - 3, I - 3, -12*I + 1, -2],
...                      [-I - 1, -2, 5*I - 1, -I - 2],
...                      [-4*I - 4, I - 5, -7*I, -I - 4]])
sage: Q, R = A.QR(full=False)
sage: Q
[-0.4160251471689219? - 0.4160251471689219?*I  0.5370861555295747? + 0.1790287185098583?*I
[-0.1386750490563073? - 0.1386750490563073?*I  -0.7519206177414046? - 0.2506402059138015?*I
[-0.5547001962252291? - 0.5547001962252291?*I  -0.2148344622118299? - 0.07161148740394329?*I
sage: R
[          7.211102550927979?  3.328201177351375? - 5.269651864139676?*I  7.9
[          0          1.074172311059150? -1.6
sage: Q.conjugate_transpose()*Q
[1.000000000000000? + 0.?e-18*I          0.?e-18 + 0.?e-18*I]
[          0.?e-17 + 0.?e-17*I 1.000000000000000? + 0.?e-17*I]
sage: Q*R-A
[0.?e-18 + 0.?e-18*I 0.?e-18 + 0.?e-18*I 0.?e-17 + 0.?e-17*I 0.?e-18 + 0.?e-18*I]
[0.?e-18 + 0.?e-18*I 0.?e-18 + 0.?e-18*I 0.?e-18 + 0.?e-17*I 0.?e-18 + 0.?e-18*I]
[0.?e-18 + 0.?e-18*I 0.?e-17 + 0.?e-18*I 0.?e-17 + 0.?e-17*I 0.?e-18 + 0.?e-18*I]
```

Results of full decompositions are cached and thus returned immutable.

```
sage: A = random_matrix(QQbar, 2, 2)
sage: Q, R = A.QR()
sage: Q.is_mutable()
False
sage: R.is_mutable()
False
```

Trivial cases return trivial results of the correct size, and we check Q itself in one case.

```
sage: A = zero_matrix(QQbar, 0, 10)
sage: Q, R = A.QR()
sage: Q.nrows(), Q.ncols()
(0, 0)
sage: R.nrows(), R.ncols()
(0, 10)
sage: A = zero_matrix(QQbar, 3, 0)
sage: Q, R = A.QR()
sage: Q.nrows(), Q.ncols()
(3, 3)
sage: R.nrows(), R.ncols()
(3, 0)
sage: Q
[1 0 0]
[0 1 0]
[0 0 1]
```

TESTS:

Inexact rings are caught and CDF suggested.

```
sage: A = matrix(RealField(100), 2, range(4))
```

```
sage: A.QR()
```

```
Traceback (most recent call last):
```

```
...
```

```
NotImplementedError: QR decomposition is implemented over exact rings, try CDF for numerical
```

Without a fraction field, we cannot hope to run the algorithm.

```
sage: A = matrix(Integers(6), 2, range(4))
```

```
sage: A.QR()
```

```
Traceback (most recent call last):
```

```
...
```

```
ValueError: QR decomposition needs a fraction field of Ring of integers modulo 6
```

The biggest obstacle is making unit vectors, thus requiring square roots, though some small cases pass through.

```
sage: A = matrix(ZZ, 3, range(9))
```

```
sage: A.QR()
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: QR decomposition unable to compute square roots in Rational Field
```

```
sage: A = matrix(ZZ, 2, range(4))
```

```
sage: Q, R = A.QR()
```

```
sage: Q
```

```
[0 1]
```

```
[1 0]
```

```
sage: R
```

```
[2 3]
```

```
[0 1]
```

REFERENCES:

AUTHOR:

- Rob Beezer (2011-02-17)

T

Returns the transpose of a matrix.

EXAMPLE:

```
sage: A = matrix(QQ, 5, range(25))
```

```
sage: A.T
```

```
[ 0  5 10 15 20]
```

```
[ 1  6 11 16 21]
```

```
[ 2  7 12 17 22]
```

```
[ 3  8 13 18 23]
```

```
[ 4  9 14 19 24]
```

adjoint()

Returns the adjoint matrix of self (matrix of cofactors).

OUTPUT:

- N - the adjoint matrix, such that $N * M = M * N = M.parent(M.det())$

ALGORITHM:

Use PARI whenever the method `self._adjoint` is included to do so in an inheriting class. Otherwise, use a generic division-free algorithm to compute the characteristic polynomial and hence the adjoint.

The result is cached.

EXAMPLES:

```
sage: M = Matrix(ZZ, 2, 2, [5, 2, 3, 4]) ; M
[5 2]
[3 4]
sage: N = M.adjoint() ; N
[ 4 -2]
[-3  5]
sage: M * N
[14  0]
[ 0 14]
sage: N * M
[14  0]
[ 0 14]
sage: M = Matrix(QQ, 2, 2, [5/3, 2/56, 33/13, 41/10]) ; M
[ 5/3  1/28]
[33/13 41/10]
sage: N = M.adjoint() ; N
[ 41/10 -1/28]
[-33/13  5/3]
sage: M * N
[7363/1092      0]
[      0 7363/1092]
```

AUTHORS:

- Unknown: No author specified in the file from 2009-06-25
- Sebastian Pancratz (2009-06-25): Reflecting the change that `_adjoint` is now implemented in this class

`as_bipartite_graph()`

Construct a bipartite graph `B` representing the matrix uniquely.

Vertices are labeled 1 to `nrows` on the left and `nrows + 1` to `nrows + ncols` on the right, representing rows and columns correspondingly. Each row is connected to each column with an edge weighted by the value of the corresponding matrix entry.

This graph is a helper for calculating automorphisms of a matrix under row and column permutations. See `automorphisms_of_rows_and_columns()`.

OUTPUT:

- A bipartite graph.

EXAMPLES:

```
sage: M = matrix(QQ, [[1/3, 7], [6, 1/4], [8, -5]])
sage: M
[1/3  7]
[ 6 1/4]
[ 8 -5]

sage: B = M.as_bipartite_graph()
sage: B
Bipartite graph on 5 vertices
sage: B.edges()
[(1, 4, 1/3), (1, 5, 7), (2, 4, 6), (2, 5, 1/4), (3, 4, 8), (3, 5, -5)]
```

```

sage: len(B.left) == M.nrows()
True
sage: len(B.right) == M.ncols()
True

```

`as_sum_of_permutations()`

Returns the current matrix as a sum of permutation matrices

According to the Birkhoff-von Neumann Theorem, any bistochastic matrix can be written as a positive sum of permutation matrices, which also means that the polytope of bistochastic matrices is integer.

As a non-bistochastic matrix can obviously not be written as a sum of permutations, this theorem is an equivalence.

This function, given a bistochastic matrix, returns the corresponding decomposition.

•`bistochastic_as_sum_of_permutations` – for more information on this method.

EXAMPLE:

We create a bistochastic matrix from a convex sum of permutations, then try to deduce the decomposition from the matrix

```

sage: L = []
sage: L.append((9,Permutation([4, 1, 3, 5, 2])))
sage: L.append((6,Permutation([5, 3, 4, 1, 2])))
sage: L.append((3,Permutation([3, 1, 4, 2, 5])))
sage: L.append((2,Permutation([1, 4, 2, 3, 5])))
sage: M = sum([c * p.to_matrix() for (c,p) in L])
sage: decomp = sage.combinat.permutation.bistochastic_as_sum_of_permutations(M)
sage: print decomp
2*B[[1, 4, 2, 3, 5]] + 3*B[[3, 1, 4, 2, 5]] + 9*B[[4, 1, 3, 5, 2]] + 6*B[[5, 3, 4, 1, 2]]

```

An exception is raised when the matrix is not bistochastic:

```

sage: M = Matrix([[2,3],[2,2]])
sage: decomp = sage.combinat.permutation.bistochastic_as_sum_of_permutations(M)
Traceback (most recent call last):
...
ValueError: The matrix is not bistochastic

```

`automorphisms_of_rows_and_columns()`

Return the automorphisms of `self` under permutations of rows and columns as a list of pairs of `PermutationGroupElement` objects.

EXAMPLES:

```

sage: M = matrix(ZZ, [[1,0],[1,0],[0,1]])
sage: M
[1 0]
[1 0]
[0 1]
sage: A = M.automorphisms_of_rows_and_columns()
sage: A
[((), ()), ((1,2), ())]
sage: M = matrix(ZZ, [[1,1,1,1],[1,1,1,1]])
sage: A = M.automorphisms_of_rows_and_columns()
sage: len(A)
48

```

One can now apply these automorphisms to `M` to show that it leaves it invariant:

```
sage: all(M.with_permuted_rows_and_columns(*i) == M for i in A)
True
```

characteristic_polynomial (*args, **kws)

Synonym for self.charpoly(...).

EXAMPLES:

```
sage: a = matrix(QQ, 2, 2, [1, 2, 3, 4]); a
[1 2]
[3 4]
sage: a.characteristic_polynomial('T')
T^2 - 5*T - 2
```

charpoly (var='x', algorithm=None)

Returns the characteristic polynomial of self, as a polynomial over the base ring.

ALGORITHM:

In the generic case of matrices over a ring (commutative and with unity), there is a division-free algorithm, which can be accessed using "df", with complexity $O(n^4)$. Alternatively, by specifying "hessenberg", this method computes the Hessenberg form of the matrix and then reads off the characteristic polynomial. Moreover, for matrices over number fields, this method can use PARI's charpoly implementation instead.

The method's logic is as follows: If no algorithm is specified, first check if the base ring is a number field (and then use PARI), otherwise check if the base ring is the ring of integers modulo n (in which case compute the characteristic polynomial of a lift of the matrix to the integers, and then coerce back to the base), next check if the base ring is an exact field (and then use the Hessenberg form), or otherwise, use the generic division-free algorithm. If an algorithm is specified explicitly, if `algorithm == "hessenberg"`, use the Hessenberg form, or otherwise use the generic division-free algorithm.

The result is cached.

INPUT:

- **var** - a variable name (default: 'x')
- **algorithm - string:**
 - "df" - Generic $O(n^4)$ division-free algorithm
 - "hessenberg" - Use the Hessenberg form of the matrix

EXAMPLES:

First a matrix over \mathbf{Z} :

```
sage: A = MatrixSpace(ZZ, 2) ( [1, 2, 3, 4] )
sage: f = A.charpoly('x')
sage: f
x^2 - 5*x - 2
sage: f.parent()
Univariate Polynomial Ring in x over Integer Ring
sage: f(A)
[0 0]
[0 0]
```

An example over \mathbf{Q} :

```
sage: A = MatrixSpace(QQ, 3) (range(9))
sage: A.charpoly('x')
x^3 - 12*x^2 - 18*x
```

```
sage: A.trace()
12
sage: A.determinant()
0
```

We compute the characteristic polynomial of a matrix over the polynomial ring $\mathbb{Z}[a]$:

```
sage: R.<a> = PolynomialRing(ZZ)
sage: M = MatrixSpace(R, 2) ([a, 1, a, a+1]); M
[  a      1]
[  a a + 1]
sage: f = M.charpoly('x'); f
x^2 + (-2*a - 1)*x + a^2
sage: f.parent()
Univariate Polynomial Ring in x over Univariate Polynomial Ring in a over Integer Ring
sage: M.trace()
2*a + 1
sage: M.determinant()
a^2
```

We compute the characteristic polynomial of a matrix over the multi-variate polynomial ring $\mathbb{Z}[x, y]$:

```
sage: R.<x,y> = PolynomialRing(ZZ, 2)
sage: A = MatrixSpace(R, 2) ([x, y, x^2, y^2])
sage: f = A.charpoly('x'); f
x^2 + (-y^2 - x)*x - x^2*y + x*y^2
```

It's a little difficult to distinguish the variables. To fix this, we temporarily view the indeterminate as Z :

```
sage: with localvars(f.parent(), 'Z'): print f
Z^2 + (-y^2 - x)*Z - x^2*y + x*y^2
```

We could also compute f in terms of Z from the start:

```
sage: A.charpoly('Z')
Z^2 + (-y^2 - x)*Z - x^2*y + x*y^2
```

Here is an example over a number field:

```
sage: x = QQ['x'].gen()
sage: K.<a> = NumberField(x^2 - 2)
sage: m = matrix(K, [[a-1, 2], [a, a+1]])
sage: m.charpoly('Z')
Z^2 - 2*a*Z - 2*a + 1
sage: m.charpoly('a')(m) == 0
True
```

Here is an example over a general commutative ring, that is to say, as of version 4.0.2, SAGE does not even positively determine that S in the following example is an integral domain. But the computation of the characteristic polynomial succeeds as follows:

```
sage: R.<a,b> = QQ[]
sage: S.<x,y> = R.quo((b^3))
sage: A = matrix(S, [[x*y^2, 2*x], [2, x^10*y]])
sage: A
[ x*y^2      2*x]
[      2 x^10*y]
sage: A.charpoly('T')
T^2 + (-x^10*y - x*y^2)*T - 4*x
```

TESTS:

```
sage: P.<a,b,c> = PolynomialRing(Rationals())
sage: u = MatrixSpace(P,3) ([[0,0,a],[1,0,b],[0,1,c]])
sage: Q.<x> = PolynomialRing(P)
sage: u.charpoly('x')
x^3 - c*x^2 - b*x - a
```

A test case from [trac ticket #6442](#). Prior to [trac ticket #12292](#), the call to `A.det()` would attempt to use the cached charpoly, and crash if an empty dictionary was cached. We don't cache dictionaries anymore, but this test should still pass:

```
sage: z = Zp(p=5)
sage: A = matrix(z, [ [3 + O(5^1), 4 + O(5^1), 4 + O(5^1)],
...                  [2*5^2 + O(5^3), 2 + O(5^1), 1 + O(5^1)],
...                  [5 + O(5^2), 1 + O(5^1), 1 + O(5^1)] ])
sage: A.charpoly(algorithm='hessenberg')
Traceback (most recent call last):
...
ValueError: negative valuation
sage: A.det()
3 + O(5)
```

The cached polynomial should be independent of the `var` argument ([trac ticket #12292](#)). We check (indirectly) that the second call uses the cached value by noting that its result is not cached:

```
sage: M = MatrixSpace(RR, 2)
sage: A = M(range(0, 2^2))
sage: type(A)
<type 'sage.matrix.matrix_generic_dense.Matrix_generic_dense'>
sage: A.charpoly('x')
x^2 - 3.000000000000000*x - 2.000000000000000
sage: A.charpoly('y')
y^2 - 3.000000000000000*y - 2.000000000000000
sage: A._cache['charpoly']
x^2 - 3.000000000000000*x - 2.000000000000000
```

AUTHORS:

- Unknown: No author specified in the file from 2009-06-25
- Sebastian Pancratz (2009-06-25): Include the division-free algorithm

cholesky()

Returns the Cholesky decomposition of a symmetric or Hermitian matrix.

INPUT:

A square matrix that is real, symmetric and positive definite. Or a square matrix that is complex, Hermitian and positive definite. Generally, the base ring for the entries of the matrix needs to be a subfield of the algebraic numbers ($\overline{\mathbb{Q}\mathbb{Q}}$). Examples include the rational numbers ($\mathbb{Q}\mathbb{Q}$), some number fields, and real algebraic numbers and the algebraic numbers themselves.

OUTPUT:

For a matrix A the routine returns a lower triangular matrix L such that,

$$A = LL^*$$

where L^* is the conjugate-transpose in the complex case, and just the transpose in the real case. If the matrix fails to be positive definite (perhaps because it is not symmetric or Hermitian), then a `ValueError` results.

ALGORITHM:

Whether or not the matrix is positive definite is checked first in every case. This is accomplished with an indefinite factorization (see `indefinite_factorization()`) which caches its result. This algorithm is of an order $n^3/3$. If the matrix is positive definite, this computation always succeeds, using just field operations. The transition to a Cholesky decomposition “only” requires computing square roots of the positive (real) entries of the diagonal matrix produced in the indefinite factorization. Hence, there is no real penalty in the positive definite check (here, or prior to calling this routine), but a field extension with square roots may not be implemented in all reasonable cases.

EXAMPLES:

This simple example has a result with entries that remain in the field of rational numbers.

```
sage: A = matrix(QQ, [[ 4, -2,  4,  2],
...                  [-2, 10, -2, -7],
...                  [ 4, -2,  8,  4],
...                  [ 2, -7,  4,  7]])
sage: A.is_symmetric()
True
sage: L = A.cholesky()
sage: L
[ 2  0  0  0]
[-1  3  0  0]
[ 2  0  2  0]
[ 1 -2  1  1]
sage: L.parent()
Full MatrixSpace of 4 by 4 dense matrices over Rational Field
sage: L*L.transpose() == A
True
```

This seemingly simple example requires first moving to the rational numbers for field operations, and then square roots necessitate that the result has entries in the field of algebraic numbers.

```
sage: A = matrix(ZZ, [[ 78, -30, -37, -2],
...                  [-30, 102, 179, -18],
...                  [-37, 179, 326, -38],
...                  [-2, -18, -38, 15]])
sage: A.is_symmetric()
True
sage: L = A.cholesky()
sage: L
[ 8.83176086632785? 0 0]
[-3.396831102433787? 9.51112708681461? 0]
[-4.189425026335004? 17.32383862241232? 2.886751345948129? 0]
[-0.2264554068289192? -1.973397116652010? -1.649572197684645? 2.886751345948129?]
sage: L.parent()
Full MatrixSpace of 4 by 4 dense matrices over Algebraic Field
sage: L*L.transpose() == A
True
```

Some subfields of the complex numbers, such as this number field of complex numbers with rational real and imaginary parts, allow for this computation.

```
sage: C.<I> = QuadraticField(-1)
sage: A = matrix(C, [[ 23, 17*I + 3, 24*I + 25, 21*I],
...                  [-17*I + 3, 38, -69*I + 89, 7*I + 15],
...                  [-24*I + 25, 69*I + 89, 976, 24*I + 6],
...                  [-21*I, -7*I + 15, -24*I + 6, 28]])
sage: A.is_hermitian()
True
```



```

sage: L = A.cholesky()
sage: L
[
      4.79...?      0      0      0
[  0.62...? - 3.54...?*I      5.00...?      0      0
[  5.21...? - 5.00...?*I      13.58...? + 10.72...?*I      24.98...?      0
[      -4.37...?*I      -0.10...? - 0.85...?*I      -0.21...? + 0.37...?*I      2.81...?]
sage: L.parent()
Full MatrixSpace of 4 by 4 dense matrices over Algebraic Field
sage: (L*L.conjugate_transpose() - A.change_ring(QQbar)).norm() < 10^-10
True

```

The field of algebraic numbers is an ideal setting for this computation.

```

sage: A = matrix(QQbar, [[
      2,      4 + 2*I,      6 - 4*I],
...      [-2*I + 4,      11, 10 - 12*I],
...      [ 4*I + 6, 10 + 12*I,      37]])
sage: A.is_hermitian()
True
sage: L = A.cholesky()
sage: L
[
      1.414213562373095?      0      0
[2.828427124746190? - 1.414213562373095?*I      1      0
[4.242640687119285? + 2.828427124746190?*I      -2*I + 2      1.732050807568878?]
sage: L.parent()
Full MatrixSpace of 3 by 3 dense matrices over Algebraic Field
sage: (L*L.conjugate_transpose() - A.change_ring(QQbar)).norm() < 10^-10
True

```

Results are cached, hence immutable. Use the `copy` function if you need to make a change.

```

sage: A = matrix(QQ, [[ 4, -2,  4,  2],
...      [-2, 10, -2, -7],
...      [ 4, -2,  8,  4],
...      [ 2, -7,  4,  7]])
sage: L = A.cholesky()
sage: L.is_immutable()
True

```

```

sage: from copy import copy
sage: LC = copy(L)
sage: LC[0,0] = 1000
sage: LC
[1000      0      0      0]
[  -1      3      0      0]
[   2      0      2      0]
[   1     -2      1      1]

```

There are a variety of situations which will prevent the computation of a Cholesky decomposition.

The base ring must be exact. For numerical work, create a matrix with a base ring of RDF or CDF and use the `cholesky()` method for matrices of that type.

```

sage: F = RealField(100)
sage: A = matrix(F, [[1.0, 3.0], [3.0, -6.0]])
sage: A.cholesky()
Traceback (most recent call last):
...
TypeError: base ring of the matrix must be exact, not Real Field with 100 bits of precision

```

The base ring may not have a fraction field.

```
sage: A = matrix(Integers(6), [[2, 0], [0, 4]])
sage: A.cholesky()
Traceback (most recent call last):
...
ValueError: unable to check positive definiteness because
Unable to create the fraction field of Ring of integers modulo 6
```

The base field may not have elements that are comparable to zero.

```
sage: F.<a> = FiniteField(5^4)
sage: A = matrix(F, [[2+a^3, 3], [3, 3]])
sage: A.cholesky()
Traceback (most recent call last):
...
ValueError: unable to check positive definiteness because
cannot convert computations from Finite Field in a of size 5^4 into real numbers
```

The algebraic closure of the fraction field of the base ring may not be implemented.

```
sage: F = Integers(7)
sage: A = matrix(F, [[4, 0], [0, 3]])
sage: A.cholesky()
Traceback (most recent call last):
...
TypeError: base field needs an algebraic closure with square roots,
not Ring of integers modulo 7
```

The matrix may not be positive definite.

```
sage: C.<I> = QuadraticField(-1)
sage: B = matrix(C, [[2, 4 - 2*I, 2 + 2*I],
...                  [4 + 2*I, 8, 10*I],
...                  [2 - 2*I, -10*I, -3]])
sage: B.is_positive_definite()
False
sage: B.cholesky()
Traceback (most recent call last):
...
ValueError: matrix is not positive definite,
so cannot compute Cholesky decomposition
```

The matrix could be positive semi-definite, and thus lack a Cholesky decomposition.

```
sage: A = matrix(QQ, [[21, 15, 12, -3],
...                  [15, 12, 9, 12],
...                  [12, 9, 7, 3],
...                  [-3, 12, 3, 8]])
sage: A.is_positive_definite()
False
sage: [A[:i,:i].determinant() for i in range(1,A.nrows()+1)]
[21, 27, 0, 0]
sage: A.cholesky()
Traceback (most recent call last):
...
ValueError: matrix is not positive definite,
so cannot compute Cholesky decomposition
```

In certain cases, the algorithm can find an analogue of the Cholesky decomposition over finite fields:

```

sage: F.<a> = FiniteField(5^3)
sage: A = matrix(F, [[ 4,          2*a^2 + 3,          4*a + 1],
...                  [ 2*a^2 + 3,          2*a + 2, 4*a^2 + 4*a + 4],
...                  [ 4*a + 1, 4*a^2 + 4*a + 4,          a^2 + 4*a]])
sage: A.is_symmetric()
True
sage: L = A.cholesky()
sage: L*L.transpose() == A
True

sage: F = FiniteField(7)
sage: A = matrix(F, [[4, 0], [0, 3]])
sage: A.cholesky()
[ 2 0]
[ 0 2*z2 + 6]

```

TESTS:

This verifies that [trac ticket #11274](#) is resolved.

```

sage: E = matrix(QQ, [[2, 1], [1, 1]])
sage: E.is_symmetric()
True
sage: E.eigenvalues()
[0.38...?, 2.61...?]
sage: E.det()
1
sage: E.cholesky()
[ 1.414213562373095? 0]
[0.7071067811865475? 0.7071067811865475?]

```

AUTHOR:

•Rob Beezer (2012-05-27)

column_module()

Return the free module over the base ring spanned by the columns of this matrix.

EXAMPLES:

```

sage: t = matrix(QQ, 3, range(9)); t
[0 1 2]
[3 4 5]
[6 7 8]
sage: t.column_module()
Vector space of degree 3 and dimension 2 over Rational Field
Basis matrix:
[ 1 0 -1]
[ 0 1 2]

```

column_space()

Return the vector space over the base ring spanned by the columns of this matrix.

EXAMPLES:

```

sage: M = MatrixSpace(QQ, 3, 3)
sage: A = M([1, 9, -7, 4/5, 4, 3, 6, 4, 3])
sage: A.column_space()
Vector space of degree 3 and dimension 3 over Rational Field
Basis matrix:
[1 0 0]

```

```
[0 1 0]
[0 0 1]
sage: W = MatrixSpace(CC, 2, 2)
sage: B = W([1, 2+3*I, 4+5*I, 9]); B
[
  1.0000000000000000 2.0000000000000000 + 3.0000000000000000*I]
[4.0000000000000000 + 5.0000000000000000*I          9.0000000000000000]
sage: B.column_space()
Vector space of degree 2 and dimension 2 over Complex Field with 53 bits of precision
Basis matrix:
[ 1.0000000000000000 0.0000000000000000]
[0.0000000000000000 1.0000000000000000]
```

conjugate()

Return the conjugate of self, i.e. the matrix whose entries are the conjugates of the entries of self.

EXAMPLES:

```
sage: A = matrix(CDF, [[1+I, 1], [0, 2*I]])
sage: A.conjugate()
[1.0 - 1.0*I      1.0]
[      0.0      -2.0*I]
```

A matrix over a not-totally-real number field:

```
sage: K.<j> = NumberField(x^2+5)
sage: M = matrix(K, [[1+j, 1], [0, 2*j]])
sage: M.conjugate()
[-j + 1      1]
[      0    -2*j]
```

There is a shortcut for the conjugate:

```
sage: M.C
[-j + 1      1]
[      0    -2*j]
```

There is also a shortcut for the conjugate transpose, or “Hermitian transpose”:

```
sage: M.H
[-j + 1      0]
[      1    -2*j]
```

Conjugates work (trivially) for matrices over rings that embed canonically into the real numbers:

```
sage: M = random_matrix(ZZ, 2)
sage: M == M.conjugate()
True
sage: M = random_matrix(QQ, 3)
sage: M == M.conjugate()
True
sage: M = random_matrix(RR, 2)
sage: M == M.conjugate()
True
```

conjugate_transpose()

Returns the transpose of self after each entry has been converted to its complex conjugate.

Note: This function is sometimes known as the “adjoint” of a matrix, though there is substantial variation and some confusion with the use of that term.

OUTPUT:

A matrix formed by taking the complex conjugate of every entry of `self` and then transposing the resulting matrix.

Complex conjugation is implemented for many subfields of the complex numbers. See the examples below, or more at `conjugate()`.

EXAMPLES:

```
sage: M = matrix(SR, 2, 2, [[2-I, 3+4*I], [9-6*I, 5*I]])
```

```
sage: M.base_ring()
```

```
Symbolic Ring
```

```
sage: M.conjugate_transpose()
```

```
[  I + 2  6*I + 9]
[-4*I + 3  -5*I]
```

```
sage: P = matrix(CC, 3, 2, [0.95-0.63*I, 0.84+0.13*I, 0.94+0.23*I, 0.23+0.59*I, 0.52-0.41*I,
```

```
sage: P.base_ring()
```

```
Complex Field with 53 bits of precision
```

```
sage: P.conjugate_transpose()
```

```
[ 0.950... + 0.630...*I  0.940... - 0.230...*I  0.520... + 0.410...*I]
[ 0.840... - 0.130...*I  0.230... - 0.590...*I -0.500... - 0.900...*I]
```

There is also a shortcut for the conjugate transpose, or “Hermitian transpose”:

```
sage: M.H
```

```
[  I + 2  6*I + 9]
[-4*I + 3  -5*I]
```

Matrices over base rings that can be embedded in the real numbers will behave as expected.

```
sage: P = random_matrix(QQ, 3, 4)
```

```
sage: P.conjugate_transpose() == P.transpose()
```

```
True
```

The conjugate of a matrix is formed by taking conjugates of all the entries. Some specialized subfields of the complex numbers are implemented in Sage and complex conjugation can be applied. (Matrices over quadratic number fields are another class of examples.)

```
sage: C = CyclotomicField(5)
```

```
sage: a = C.gen(); a
zeta5
```

```
sage: CC(a)
```

```
0.309016994374947 + 0.951056516295154*I
```

```
sage: M = matrix(C, 1, 2, [a^2, a+a^3])
```

```
sage: M.conjugate_transpose()
```

```
[          zeta5^3]
[-zeta5^3 - zeta5 - 1]
```

Conjugation does not make sense over rings not containing complex numbers.

```
sage: N = matrix(GF(5), 2, [0,1,2,3])
```

```
sage: N.conjugate_transpose()
```

```
Traceback (most recent call last):
```

```
...
```

```
AttributeError: 'sage.rings.finite_rings.integer_mod.IntegerMod_int' object has no attribute
```

AUTHOR:

Rob Beezer (2010-12-13)

cyclic_subspace (*v*, *var*=None, *basis*='echelon')

Create a cyclic subspace for a vector, and optionally, a minimal polynomial for the iterated powers.

These subspaces are also known as Krylov subspaces. They are spanned by the vectors

$$\{v, Av, A^2v, A^3v, \dots\}$$

INPUT:

- *self* - a square matrix with entries from a field.
- *v* - a vector with a degree equal to the size of the matrix and entries compatible with the entries of the matrix.
- *var* - default: None - if specified as a string or a generator of a polynomial ring, then this will be used to construct a polynomial reflecting a relation of linear dependence on the powers $A^i v$ and this will cause the polynomial to be returned along with the subspace. A generator must create polynomials with coefficients from the same field as the matrix entries.
- *basis* - default: echelon - the basis for the subspace is “echelonized” by default, but the keyword ‘iterates’ will return a subspace with a user basis equal to the largest linearly independent set $\{v, Av, A^2v, A^3v, \dots, A^{k-1}v\}$.

OUTPUT:

Suppose k is the smallest power such that $\{v, Av, A^2v, A^3v, \dots, A^k v\}$ is linearly dependent. Then the subspace returned will have dimension k and be spanned by the powers 0 through $k - 1$.

If a polynomial is requested through the use of the *var* keyword, then a pair is returned, with the polynomial first and the subspace second. The polynomial is the unique monic polynomial whose coefficients provide a relation of linear dependence on the first k powers.

For less convenient, but more flexible output, see the helper method “_cyclic_subspace” in this module.

EXAMPLES:

```
sage: A = matrix(QQ, [[5,4,2,1],[0,1,-1,-1],[-1,-1,3,0],[1,1,-1,2]])
sage: v = vector(QQ, [0,1,0,0])
sage: E = A.cyclic_subspace(v); E
Vector space of degree 4 and dimension 3 over Rational Field
Basis matrix:
[ 1  0  0  0]
[ 0  1  0  0]
[ 0  0  1 -1]
sage: F = A.cyclic_subspace(v, basis='iterates'); F
Vector space of degree 4 and dimension 3 over Rational Field
User basis matrix:
[ 0  1  0  0]
[ 4  1 -1  1]
[23  1 -8  8]
sage: E == F
True
sage: p, S = A.cyclic_subspace(v, var='T'); p
T^3 - 9*T^2 + 24*T - 16
sage: gen = polygen(QQ, 'z')
sage: p, S = A.cyclic_subspace(v, var=gen); p
z^3 - 9*z^2 + 24*z - 16
sage: p.degree() == E.dimension()
True
```

The polynomial has coefficients that yield a non-trivial relation of linear dependence on the iterates. Or, equivalently, evaluating the polynomial with the matrix will create a matrix that annihilates the vector.

```

sage: A = matrix(QQ, [[15, 37/3, -16, -104/3, -29, -7/3, 35, 2/3, -29/3, -1/3],
...                  [ 2, 9, -1, -6, -6, 0, 7, 0, -2, 0],
...                  [24, 74/3, -29, -208/3, -58, -14/3, 70, 4/3, -58/3, -2/3],
...                  [-6, -19, 3, 21, 19, 0, -21, 0, 6, 0],
...                  [2, 6, -1, -6, -3, 0, 7, 0, -2, 0],
...                  [-96, -296/3, 128, 832/3, 232, 65/3, -279, -16/3, 232/3, 8/3],
...                  [0, 0, 0, 0, 0, 0, 3, 0, 0, 0],
...                  [20, 26/3, -30, -199/3, -42, -14/3, 70, 13/3, -55/3, -2/3],
...                  [18, 57, -9, -54, -57, 0, 63, 0, -15, 0],
...                  [0, 0, 0, 0, 0, 0, 0, 0, 0, 3]])
sage: u = zero_vector(QQ, 10); u[0] = 1
sage: p, S = A.cyclic_subspace(u, var='t', basis='iterates')
sage: S
Vector space of degree 10 and dimension 3 over Rational Field
User basis matrix:
[  1   0   0   0   0   0   0   0   0   0]
[ 15   2  24  -6   2 -96   0  20  18   0]
[ 79  12 140 -36  12 -560   0 116 108   0]
sage: p
t^3 - 9*t^2 + 27*t - 27
sage: k = p.degree()
sage: coeffs = p.list()
sage: iterates = S.basis() + [A^k*u]
sage: sum(coeffs[i]*iterates[i] for i in range(k+1))
(0, 0, 0, 0, 0, 0, 0, 0, 0, 0)
sage: u in p(A).right_kernel()
True

```

TESTS:

A small case.

```

sage: A = matrix(QQ, 5, range(25))
sage: u = zero_vector(QQ, 5)
sage: A.cyclic_subspace(u)
Vector space of degree 5 and dimension 0 over Rational Field
Basis matrix:
[]

```

Various problem inputs. Notice the vector must have entries that coerce into the base ring of the matrix, and a polynomial ring generator must have a base ring that agrees with the base ring of the matrix.

```

sage: A = matrix(QQ, 4, range(16))
sage: v = vector(QQ, 4, range(4))

sage: A.cyclic_subspace('junk')
Traceback (most recent call last):
...
TypeError: first input should be a vector, not junk

sage: A.cyclic_subspace(v, var=sin(x))
Traceback (most recent call last):
...
TypeError: polynomial variable must be a string or polynomial ring generator, not sin(x)

sage: t = polygen(GF(7), 't')
sage: A.cyclic_subspace(v, var=t)
Traceback (most recent call last):
...

```

```
TypeError: polynomial generator must be over the same ring as the matrix entries

sage: A.cyclic_subspace(v, basis='garbage')
Traceback (most recent call last):
...
ValueError: basis format must be 'echelon' or 'iterates', not garbage

sage: B = matrix(QQ, 4, 5, range(20))
sage: B.cyclic_subspace(v)
Traceback (most recent call last):
...
TypeError: matrix must be square, not 4 x 5

sage: C = matrix(QQ, 5, 5, range(25))
sage: C.cyclic_subspace(v)
Traceback (most recent call last):
...
TypeError: vector must have degree equal to the size of the matrix, not 4

sage: D = matrix(RDF, 4, 4, range(16))
sage: D.cyclic_subspace(v)
Traceback (most recent call last):
...
TypeError: matrix entries must be from an exact ring, not Real Double Field

sage: E = matrix(Integers(6), 4, 4, range(16))
sage: E.cyclic_subspace(v)
Traceback (most recent call last):
...
TypeError: matrix entries must be from an exact field, not Ring of integers modulo 6

sage: F.<a> = GF(2^4)
sage: G = matrix(QQ, 4, range(16))
sage: w = vector(F, 4, [1, a, a^2, a^3])
sage: G.cyclic_subspace(w)
Traceback (most recent call last):
...
TypeError: unable to make vector entries compatible with matrix entries
```

AUTHOR:

•Rob Beezer (2011-05-20)

decomposition (*algorithm='spin', is_diagonalizable=False, dual=False*)

Returns the decomposition of the free module on which this matrix A acts from the right (i.e., the action is x goes to $x A$), along with whether this matrix acts irreducibly on each factor. The factors are guaranteed to be sorted in the same way as the corresponding factors of the characteristic polynomial.

Let A be the matrix acting from the on the vector space V of column vectors. Assume that A is square. This function computes maximal subspaces W_1, \dots, W_n corresponding to Galois conjugacy classes of eigenvalues of A . More precisely, let $f(X)$ be the characteristic polynomial of A . This function computes the subspace $W_i = \ker(g_i(A)^n)$, where $g_i(X)$ is an irreducible factor of $f(X)$ and $g_i(X)$ exactly divides $f(X)$. If the optional parameter `is_diagonalizable` is `True`, then we let $W_i = \ker(g(A))$, since then we know that $\ker(g(A)) = \ker(g(A)^n)$.

INPUT:

•self - a matrix

- algorithm** - 'spin' (default): algorithm involves iterating the action of self on a vector. 'kernel': naively just compute $\ker(f_i(A))$ for each factor f_i .
- dual** - bool (default: False): If True, also returns the corresponding decomposition of V under the action of the transpose of A. The factors are guaranteed to correspond.
- is_diagonalizable** - if the matrix is known to be diagonalizable, set this to True, which might speed up the algorithm in some cases.

Note: If the base ring is not a field, the kernel algorithm is used.

OUTPUT:

- Sequence - list of pairs (V,t), where V is a vector spaces and t is a bool, and t is True exactly when the charpoly of self on V is irreducible.
- (optional) list - list of pairs (W,t), where W is a vector space and t is a bool, and t is True exactly when the charpoly of the transpose of self on W is irreducible.

EXAMPLES:

```
sage: A = matrix(ZZ, 4, [3,4,5,6,7,3,8,10,14,5,6,7,2,2,10,9])
sage: B = matrix(QQ, 6, range(36))
sage: B*11
[ 0 11 22 33 44 55]
[ 66 77 88 99 110 121]
[132 143 154 165 176 187]
[198 209 220 231 242 253]
[264 275 286 297 308 319]
[330 341 352 363 374 385]
sage: A.decomposition()
[
(Ambient free module of rank 4 over the principal ideal domain Integer Ring, True)
]
sage: B.decomposition()
[
(Vector space of degree 6 and dimension 2 over Rational Field
Basis matrix:
[ 1  0 -1 -2 -3 -4]
[ 0  1  2  3  4  5], True),
(Vector space of degree 6 and dimension 4 over Rational Field
Basis matrix:
[ 1  0  0  0 -5  4]
[ 0  1  0  0 -4  3]
[ 0  0  1  0 -3  2]
[ 0  0  0  1 -2  1], False)
]
```

decomposition_of_subspace (*M*, *check_restrict=True*, ***kws*)

Suppose the right action of self on M leaves M invariant. Return the decomposition of M as a list of pairs (W, is_irred) where is_irred is True if the charpoly of self acting on the factor W is irreducible.

Additional inputs besides M are passed onto the decomposition command.

INPUT:

- M* – A subspace of the free module self acts on.
- check_restrict** – A boolean (default: True); Call restrict with or without check.
- kws* – Keywords that will be forwarded to `decomposition()`.

We do a decomposition over $\mathbb{Z}\mathbb{Z}$:

TESTS:

denominator ()

Return the least common multiple of the denominators of the elements of self.

If there is no denominator function for the base field, or no LCM function for the denominators, raise a `TypeError`.

EXAMPLES:

```
sage: A = MatrixSpace(QQ,2)(['1/2', '1/3', '1/5', '1/7'])
sage: A.denominator()
210
```

A trivial example:

```
sage: A = matrix(QQ, 0, 2)
sage: A.denominator()
1
```

Denominators are not defined for real numbers:

```
sage: A = MatrixSpace(RealField(), 2) ([1, 2, 3, 4])
sage: A.denominator()
Traceback (most recent call last):
...
TypeError: denominator not defined for elements of the base ring
```

We can even compute the denominator of matrix over the fraction field of $\mathbb{Z}[x]$.

```
sage: K.<x> = Frac(ZZ['x'])
sage: A = MatrixSpace(K, 2) ([1/x, 2/(x+1), 1, 5/(x^3)])
sage: A.denominator()
x^4 + x^3
```

Here's an example involving a cyclotomic field:

```
sage: K.<z> = CyclotomicField(3)
sage: M = MatrixSpace(K, 3, sparse=True)
sage: A = M([(1+z)/3, (2+z)/3, z/3, 1, 1+z, -2, 1, 5, -1+z])
sage: print A
[1/3*z + 1/3 1/3*z + 2/3      1/3*z]
[          1          z + 1      -2]
[          1          5      z - 1]
sage: print A.denominator()
3
```

density()

Return the density of the matrix.

By density we understand the ratio of the number of nonzero positions and the `self.nrows() * self.ncols()`, i.e. the number of possible nonzero positions.

EXAMPLE:

First, note that the density parameter does not ensure the density of a matrix, it is only an upper bound.

```
sage: A = random_matrix(GF(127), 200, 200, density=0.3)
sage: A.density()
5211/20000
```

```
sage: A = matrix(QQ, 3, 3, [0, 1, 2, 3, 0, 0, 6, 7, 8])
sage: A.density()
2/3
```

```
sage: a = matrix([], [], [], [])
sage: a.density()
0
```

derivative(*args)

Derivative with respect to variables supplied in args.

Multiple variables and iteration counts may be supplied; see documentation for the `global derivative()` function for more details.

EXAMPLES:

```

sage: v = vector([1,x,x^2])
sage: v.derivative(x)
(0, 1, 2*x)
sage: type(v.derivative(x)) == type(v)
True
sage: v = vector([1,x,x^2], sparse=True)
sage: v.derivative(x)
(0, 1, 2*x)
sage: type(v.derivative(x)) == type(v)
True
sage: v.derivative(x,x)
(0, 0, 2)

```

det (*args, **kws)
 Synonym for self.determinant(...).

EXAMPLES:

```

sage: A = MatrixSpace(Integers(8),3) ([1,7,3, 1,1,1, 3,4,5])
sage: A.det ()
6

```

determinant (algorithm=None)
 Returns the determinant of self.

ALGORITHM:

For small matrices (n less than 4), this is computed using the naive formula. In the specific case of matrices over the integers modulo a non-prime, the determinant of a lift is computed over the integers. In general, the characteristic polynomial is computed either using the Hessenberg form (specified by "hessenberg") or the generic division-free algorithm (specified by "df"). When the base ring is an exact field, the default choice is "hessenberg", otherwise it is "df". Note that for matrices over most rings, more sophisticated algorithms can be used. (Type A.determinant? to see what is done for a specific matrix A.)

INPUT:

•algorithm - string:

- "df" - Generic $O(n^4)$ division-free algorithm
- "hessenberg" - Use the Hessenberg form of the matrix

EXAMPLES:

```

sage: A = MatrixSpace(Integers(8),3) ([1,7,3, 1,1,1, 3,4,5])
sage: A.determinant ()
6
sage: A.determinant () is A.determinant ()
True
sage: A[0,0] = 10
sage: A.determinant ()
7

```

We compute the determinant of the arbitrary 3x3 matrix:

```

sage: R = PolynomialRing(QQ,9,'x')
sage: A = matrix(R,3,R.gens())
sage: A
[x0 x1 x2]
[x3 x4 x5]
[x6 x7 x8]

```

```
sage: A.determinant()
-x^2*x^4*x^6 + x^1*x^5*x^6 + x^2*x^3*x^7 - x^0*x^5*x^7 - x^1*x^3*x^8 + x^0*x^4*x^8
```

We create a matrix over $\mathbb{Z}[x, y]$ and compute its determinant.

```
sage: R.<x,y> = PolynomialRing(IntegerRing(), 2)
sage: A = MatrixSpace(R, 2) ([x, y, x**2, y**2])
sage: A.determinant()
-x^2*y + x*y^2
```

TESTS:

```
sage: A = matrix(5, 5, [next_prime(i^2) for i in range(25)])
sage: B = MatrixSpace(ZZ['x'], 5, 5)(A)
sage: A.det() - B.det()
0
```

We verify that [trac ticket #5569](#) is resolved (otherwise the following would hang for hours):

```
sage: d = random_matrix(GF(next_prime(10^20)), 50).det()
sage: d = random_matrix(Integers(10^50), 50).det()
```

We verify that [trac 7704](#) is resolved:

```
sage: matrix(ZZ, {(0,0):1, (1,1):2, (2,2):3, (3,3):4}).det()
24
sage: matrix(QQ, {(0,0):1, (1,1):2, (2,2):3, (3,3):4}).det()
24
```

We verify that [trac 10063](#) is resolved:

```
sage: A = GF(2) ['x,y,z']
sage: A.inject_variables()
Defining x, y, z
sage: R = A.quotient(x^2 + 1).quotient(y^2 + 1).quotient(z^2 + 1)
sage: R.inject_variables()
Defining xbarbarbar, ybarbarbar, zbarbarbar
sage: M = matrix([[1,1,1,1], [xbarbarbar, ybarbarbar, 1, 1], [0, 1, zbarbarbar, 1], [xbarbarbar, zbarbarbar, 1, 1]])
sage: M.determinant()
xbarbarbar*ybarbarbar*zbarbarbar + xbarbarbar*ybarbarbar + xbarbarbar*zbarbarbar + ybarbarbar
```

Check that the determinant is computed from a cached charpoly properly:

```
sage: A = matrix(RR, [ [1, 0, 1/2],
...                   [0, 1, 0 ],
...                   [0, 0, -2 ] ])
sage: B = copy(A)
sage: _ = A.charpoly()
sage: A.determinant() == B.determinant()
True
```

AUTHORS:

- Unknown: No author specified in the file from 2009-06-25
- Sebastian Pancratz (2009-06-25): Use the division-free algorithm for charpoly
- Thierry Monteil (2010-10-05): Bugfix for [trac ticket #10063](#), so that the determinant is computed even for rings for which the `is_field` method is not implemented.

diagonal()

Return the diagonal entries of `self`.

OUTPUT:

A list containing the entries of the matrix that have equal row and column indices, in order of the indices. Behavior is not limited to square matrices.

EXAMPLES:

```
sage: A = matrix([[2, 5], [3, 7]]); A
[2 5]
[3 7]
sage: A.diagonal()
[2, 7]
```

Two rectangular matrices.

```
sage: B = matrix(3, 7, range(21)); B
[ 0  1  2  3  4  5  6]
[ 7  8  9 10 11 12 13]
[14 15 16 17 18 19 20]
sage: B.diagonal()
[0, 8, 16]
```

```
sage: C = matrix(3, 2, range(6)); C
[0 1]
[2 3]
[4 5]
sage: C.diagonal()
[0, 3]
```

Empty matrices behave properly.

```
sage: E = matrix(0, 5, []); E
[]
sage: E.diagonal()
[]
```

echelon_form(*algorithm*='default', *cutoff*=0, ***kws*)

Return the echelon form of self.

Note: This row reduction does not use division if the matrix is not over a field (e.g., if the matrix is over the integers). If you want to calculate the echelon form using division, then use `rref()`, which assumes that the matrix entries are in a field (specifically, the field of fractions of the base ring of the matrix).

INPUT:

- *algorithm* – string. Which algorithm to use. Choices are
 - 'default': Let Sage choose an algorithm (default).
 - 'classical': Gauss elimination.
 - 'strassen': use a Strassen divide and conquer algorithm (if available)
- *cutoff* – integer. Only used if the Strassen algorithm is selected.
- *transformation* – boolean. Whether to also return the transformation matrix. Some matrix backends do not provide this information, in which case this option is ignored.

OUTPUT:

The reduced row echelon form of `self`, as an immutable matrix. Note that `self` is *not* changed by this command. Use `echelonize()` to change `self` in place.

If the optional parameter `transformation=True` is specified, the output consists of a pair (E, T) of matrices where E is the echelon form of `self` and T is the transformation matrix.

EXAMPLES:

```
sage: MS = MatrixSpace(GF(19), 2, 3)
sage: C = MS.matrix([1, 2, 3, 4, 5, 6])
sage: C.rank()
2
sage: C.nullity()
0
sage: C.echelon_form()
[ 1  0 18]
[ 0  1  2]
```

The matrix library used for \mathbb{Z}/p -matrices does not return the transformation matrix, so the `transformation` option is ignored:

```
sage: C.echelon_form(transformation=True)
[ 1  0 18]
[ 0  1  2]

sage: D = matrix(ZZ, 2, 3, [1, 2, 3, 4, 5, 6])
sage: D.echelon_form(transformation=True)
(
[1 2 3]  [ 1  0]
[0 3 6], [ 4 -1]
)
sage: E, T = D.echelon_form(transformation=True)
sage: T*D == E
True
```

echelonize (*algorithm*='default', *cutoff*=0, ***kws*)

Transform `self` into a matrix in echelon form over the same base ring as `self`.

Note: This row reduction does not use division if the matrix is not over a field (e.g., if the matrix is over the integers). If you want to calculate the echelon form using division, then use `rref()`, which assumes that the matrix entries are in a field (specifically, the field of fractions of the base ring of the matrix).

INPUT:

- *algorithm* – string. Which algorithm to use. Choices are
 - 'default': Let Sage choose an algorithm (default).
 - 'classical': Gauss elimination.
 - 'strassen': use a Strassen divide and conquer algorithm (if available)
- *cutoff* – integer. Only used if the Strassen algorithm is selected.
- *transformation* – boolean. Whether to also return the transformation matrix. Some matrix backends do not provide this information, in which case this option is ignored.

OUTPUT:

The matrix `self` is put into echelon form. Nothing is returned unless the keyword option `transformation=True` is specified, in which case the transformation matrix is returned.

EXAMPLES:

```
sage: a = matrix(QQ, 3, range(9)); a
[0 1 2]
[3 4 5]
[6 7 8]
sage: a.echelonize()
sage: a
[ 1  0 -1]
[ 0  1  2]
[ 0  0  0]
```

An immutable matrix cannot be transformed into echelon form. Use `self.echelon_form()` instead:

```
sage: a = matrix(QQ, 3, range(9)); a.set_immutable()
sage: a.echelonize()
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead
(i.e., use copy(M) to change a copy of M).
sage: a.echelon_form()
[ 1  0 -1]
[ 0  1  2]
[ 0  0  0]
```

Echelon form over the integers is what is also classically often known as Hermite normal form:

```
sage: a = matrix(ZZ, 3, range(9))
sage: a.echelonize(); a
[ 3  0 -3]
[ 0  1  2]
[ 0  0  0]
```

We compute an echelon form both over a domain and fraction field:

```
sage: R.<x,y> = QQ[]
sage: a = matrix(R, 2, [x,y,x,y])
sage: a.echelon_form()                                     # not very useful? -- why two copies of the same row?
[x y]
[x y]

sage: b = a.change_ring(R.fraction_field())
sage: b.echelon_form()                                     # potentially useful
[ 1 y/x]
[ 0  0]
```

Echelon form is not defined over arbitrary rings:

```
sage: a = matrix(Integers(9), 3, range(9))
sage: a.echelon_form()
Traceback (most recent call last):
...
NotImplementedError: Echelon form not implemented over 'Ring of integers modulo 9'.
```

Involving a sparse matrix:

```
sage: m = matrix(3, [1, 1, 1, 1, 0, 2, 1, 2, 0], sparse=True); m
[1 1 1]
[1 0 2]
[1 2 0]
sage: m.echelon_form()
[ 1  0  2]
```



```

[ 0  1 -1]
[ 0  0  0]
sage: m.echelonize(); m
[ 1  0  2]
[ 0  1 -1]
[ 0  0  0]

```

The transformation matrix is optionally returned:

```

sage: m_original = m
sage: transformation_matrix = m.echelonize(transformation=True)
sage: m == transformation_matrix * m_original
True

```

eigenmatrix_left()

Return matrices D and P, where D is a diagonal matrix of eigenvalues and P is the corresponding matrix where the rows are corresponding eigenvectors (or zero vectors) so that $P \cdot \text{self} = D \cdot P$.

EXAMPLES:

```

sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: D, P = A.eigenmatrix_left()
sage: D
[
          0          0          0]
[
          0 -1.348469228349535?  0]
[
          0          0 13.34846922834954?]
sage: P
[
          1          -2          1]
[
          1 0.3101020514433644? -0.3797958971132713?]
[
          1 1.289897948556636?  1.579795897113272?]
sage: P*A == D*P
True

```

Because P is invertible, A is diagonalizable.

```

sage: A == (~P) * D * P
True

```

The matrix P may contain zero rows corresponding to eigenvalues for which the algebraic multiplicity is greater than the geometric multiplicity. In these cases, the matrix is not diagonalizable.

```

sage: A = jordan_block(2, 3); A
[2 1 0]
[0 2 1]
[0 0 2]
sage: A = jordan_block(2, 3)
sage: D, P = A.eigenmatrix_left()
sage: D
[2 0 0]
[0 2 0]
[0 0 2]
sage: P
[0 0 1]
[0 0 0]
[0 0 0]
sage: P*A == D*P
True

```

TESTS:

For matrices with floating point entries, some platforms will return eigenvectors that are negatives of those returned by the majority of platforms. This test accounts for that possibility. Running this test independently, without adjusting the eigenvectors could indicate this situation on your hardware.

```
sage: A = matrix(QQ, 3, 3, range(9))
sage: em = A.change_ring(RDF).eigenmatrix_left()
sage: evals = em[0]; evals.dense_matrix().zero_at(2e-15)
[ 13.3484692283      0.0      0.0]
[      0.0 -1.34846922835      0.0]
[      0.0      0.0      0.0]
sage: evectors = em[1];
sage: for i in range(3):
...     scale = evectors[i,0].sign()
...     evectors.rescale_row(i, scale)
sage: evectors
[ 0.440242867...  0.567868371...  0.695493875...]
[ 0.897878732...  0.278434036... -0.341010658...]
[ 0.408248290... -0.816496580...  0.408248290...]
```

eigenmatrix_right()

Return matrices D and P, where D is a diagonal matrix of eigenvalues and P is the corresponding matrix where the columns are corresponding eigenvectors (or zero vectors) so that $\text{self} \cdot P = P \cdot D$.

EXAMPLES:

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: D, P = A.eigenmatrix_right()
sage: D
[      0      0      0]
[      0 -1.348469228349535?      0]
[      0      0 13.34846922834954?]
sage: P
[      1      1      1]
[ -2  0.1303061543300932?  3.069693845669907?]
[      1 -0.7393876913398137?  5.139387691339814?]
sage: A*P == P*D
True
```

Because P is invertible, A is diagonalizable.

```
sage: A == P*D*(~P)
True
```

The matrix P may contain zero columns corresponding to eigenvalues for which the algebraic multiplicity is greater than the geometric multiplicity. In these cases, the matrix is not diagonalizable.

```
sage: A = jordan_block(2, 3); A
[2 1 0]
[0 2 1]
[0 0 2]
sage: A = jordan_block(2, 3)
sage: D, P = A.eigenmatrix_right()
sage: D
[2 0 0]
[0 2 0]
[0 0 2]
```

```
sage: P
[1 0 0]
[0 0 0]
[0 0 0]
sage: A*P == P*D
True
```

TESTS:

For matrices with floating point entries, some platforms will return eigenvectors that are negatives of those returned by the majority of platforms. This test accounts for that possibility. Running this test independently, without adjusting the eigenvectors could indicate this situation on your hardware.

```
sage: B = matrix(QQ, 3, 3, range(9))
sage: em = B.change_ring(RDF).eigenmatrix_right()
sage: evals = em[0]; evals.dense_matrix().zero_at(2e-15)
[ 13.3484692283      0.0      0.0]
[      0.0 -1.34846922835      0.0]
[      0.0      0.0      0.0]
sage: evectors = em[1];
sage: for i in range(3):
...     scale = evectors[0,i].sign()
...     evectors.rescale_col(i, scale)
sage: evectors
[ 0.164763817...  0.799699663...  0.408248290...]
[ 0.505774475...  0.104205787... -0.816496580...]
[ 0.846785134... -0.591288087...  0.408248290...]
```

eigenspaces_left (*format='all', var='a', algebraic_multiplicity=False*)

Compute the left eigenspaces of a matrix.

Note that `eigenspaces_left()` and `left_eigenspaces()` are identical methods. Here “left” refers to the eigenvectors being placed to the left of the matrix.

INPUT:

- `self` - a square matrix over an exact field. For inexact matrices consult the numerical or symbolic matrix classes.
- `format` - default: None
 - ‘all’ - attempts to create every eigenspace. This will always be possible for matrices with rational entries.
 - ‘galois’ - for each irreducible factor of the characteristic polynomial, a single eigenspace will be output for a single root/eigenvalue for the irreducible factor.
 - None - Uses the ‘all’ format if the base ring is contained in an algebraically closed field which is implemented. Otherwise, uses the ‘galois’ format.
- `var` - default: ‘a’ - variable name used to represent elements of the root field of each irreducible factor of the characteristic polynomial. If `var='a'`, then the root fields will be in terms of `a0, a1, a2, ...,`, where the numbering runs across all the irreducible factors of the characteristic polynomial, even for linear factors.
- `algebraic_multiplicity` - default: False - whether or not to include the algebraic multiplicity of each eigenvalue in the output. See the discussion below.

OUTPUT:

If `algebraic_multiplicity=False`, return a list of pairs (e, V) where e is an eigenvalue of the matrix, and V is

the corresponding left eigenspace. For Galois conjugates of eigenvalues, there may be just one representative eigenspace, depending on the `format` keyword.

If `algebraic_multiplicity=True`, return a list of triples (e, V, n) where e and V are as above and n is the algebraic multiplicity of the eigenvalue.

Warning: Uses a somewhat naive algorithm (simply factors the characteristic polynomial and computes kernels directly over the extension field).

EXAMPLES:

We compute the left eigenspaces of a 3×3 rational matrix. First, we request *all* of the eigenvalues, so the results are in the field of algebraic numbers, \overline{QQ} . Then we request just one eigenspace per irreducible factor of the characteristic polynomial with the *galois* keyword.

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: es = A.eigenspaces_left(format='all'); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]),
(-1.348469228349535?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 0.3101020514433644? -0.3797958971132713?]),
(13.34846922834954?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 1.289897948556636? 1.579795897113272?])
]

sage: es = A.eigenspaces_left(format='galois'); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
1 1/15*a1 + 2/5 2/15*a1 - 1/5])
]
sage: es = A.eigenspaces_left(format='galois', algebraic_multiplicity=True); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1], 1),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
1 1/15*a1 + 2/5 2/15*a1 - 1/5], 1)
]
sage: e, v, n = es[0]; v = v.basis()[0]
sage: delta = e*v - v*A
sage: abs(abs(delta)) < 1e-10
True
```

The same computation, but with implicit base change to a field.

```
sage: A = matrix(ZZ, 3, range(9)); A
[0 1 2]
```

```
[3 4 5]
[6 7 8]
sage: A.eigenspaces_left(format='galois')
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2 1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
1 1/15*a1 + 2/5 2/15*a1 - 1/5])
]
```

We compute the left eigenspaces of the matrix of the Hecke operator T_2 on level 43 modular symbols, both with all eigenvalues (the default) and with one subspace per factor.

```
sage: A = ModularSymbols(43).T(2).matrix(); A
[ 3  0  0  0  0  0 -1]
[ 0 -2  1  0  0  0  0]
[ 0 -1  1  1  0 -1  0]
[ 0 -1  0 -1  2 -1  1]
[ 0 -1  0  1  1 -1  1]
[ 0  0 -2  0  2 -2  1]
[ 0  0 -1  0  1  0 -1]
sage: A.base_ring()
Rational Field
sage: f = A.charpoly(); f
x^7 + x^6 - 12*x^5 - 16*x^4 + 36*x^3 + 52*x^2 - 32*x - 48
sage: factor(f)
(x - 3) * (x + 2)^2 * (x^2 - 2)^2
sage: A.eigenspaces_left(algebraic_multiplicity=True)
[
(3, Vector space of degree 7 and dimension 1 over Rational Field
User basis matrix:
[ 1  0 1/7  0 -1/7  0 -2/7], 1),
(-2, Vector space of degree 7 and dimension 2 over Rational Field
User basis matrix:
[ 0  1  0  1 -1  1 -1]
[ 0  0  1  0 -1  2 -1], 2),
(-1.414213562373095?, Vector space of degree 7 and dimension 2 over Algebraic Field
User basis matrix:
[
0 1 -1 1 -1 1 -1
0 0 1 0 -1 2 -1]
(1.414213562373095?, Vector space of degree 7 and dimension 2 over Algebraic Field
User basis matrix:
[
0 1 -1 1 -1 1 -1
0 0 1 0 -1 2 -1]
]
sage: A.eigenspaces_left(format='galois', algebraic_multiplicity=True)
[
(3, Vector space of degree 7 and dimension 1 over Rational Field
User basis matrix:
[ 1  0 1/7  0 -1/7  0 -2/7], 1),
(-2, Vector space of degree 7 and dimension 2 over Rational Field
User basis matrix:
[ 0  1  0  1 -1  1 -1]
[ 0  0  1  0 -1  2 -1], 2),
(a2, Vector space of degree 7 and dimension 2 over Number Field in a2 with defining polynomial
User basis matrix:
```

```
[
  0      1      0      -1 -a2 - 1      1      -1]
[
  0      0      1      0      -1      0 -a2 + 1], 2)
]
```

Next we compute the left eigenspaces over the finite field of order 11.

```
sage: A = ModularSymbols(43, base_ring=GF(11), sign=1).T(2).matrix(); A
[ 3  9  0  0]
[ 0  9  0  1]
[ 0 10  9  2]
[ 0  9  0  2]
sage: A.base_ring()
Finite Field of size 11
sage: A.charpoly()
x^4 + 10*x^3 + 3*x^2 + 2*x + 1
sage: A.eigenspaces_left(format='galois', var = 'beta')
[
  (9, Vector space of degree 4 and dimension 1 over Finite Field of size 11
  User basis matrix:
  [0 0 1 5]),
  (3, Vector space of degree 4 and dimension 1 over Finite Field of size 11
  User basis matrix:
  [1 6 0 6]),
  (beta2, Vector space of degree 4 and dimension 1 over Univariate Quotient Polynomial Ring in
  User basis matrix:
  [
      0      1      0 5*beta2 + 10])
]
```

This method is only applicable to exact matrices. The “eigenmatrix” routines for matrices with double-precision floating-point entries (RDF, CDF) are the best alternative. (Since some platforms return eigenvectors that are the negatives of those given here, this one example is not tested here.) There are also “eigenmatrix” routines for matrices with symbolic entries.

```
sage: A = matrix(QQ, 3, 3, range(9))
sage: A.change_ring(RR).eigenspaces_left()
Traceback (most recent call last):
...
NotImplementedError: eigenspaces cannot be computed reliably for inexact rings such as Real
consult numerical or symbolic matrix classes for other options

sage: em = A.change_ring(RDF).eigenmatrix_left()
sage: eigenvalues = em[0]; eigenvalues.dense_matrix().zero_at(2e-15)
[ 13.3484692283      0.0      0.0]
[      0.0 -1.34846922835      0.0]
[      0.0      0.0      0.0]
sage: eigenvectors = em[1]; eigenvectors # not tested
[ 0.440242867... 0.567868371... 0.695493875...]
[ 0.897878732... 0.278434036... -0.341010658...]
[ 0.408248290... -0.816496580... 0.408248290...]

sage: x, y = var('x y')
sage: S = matrix([[x, y], [y, 3*x^2]])
sage: em = S.eigenmatrix_left()
sage: eigenvalues = em[0]; eigenvalues
[3/2*x^2 + 1/2*x - 1/2*sqrt(9*x^4 - 6*x^3 + x^2 + 4*y^2)
 [
  0 3/2*x^2 + 1/2*x + 1/2*sqrt(9*x^4 -
sage: eigenvectors = em[1]; eigenvectors
[
  1 1/2*(3*x^2 - x - sqrt(9*x^4 - 6*x^3 +
  1 1/2*(3*x^2 - x + sqrt(9*x^4 - 6*x^3 +
```

A request for 'all' the eigenvalues, when it is not possible, will raise an error. Using the 'galois' format option is more likely to be successful.

```
sage: F.<b> = FiniteField(11^2)
sage: A = matrix(F, [[b + 1, b + 1], [10*b + 4, 5*b + 4]])
sage: A.eigenspaces_left(format='all')
Traceback (most recent call last):
...
NotImplementedError: unable to construct eigenspaces for eigenvalues outside the base field,
try the keyword option: format='galois'

sage: A.eigenspaces_left(format='galois')
[
(a0, Vector space of degree 2 and dimension 1 over Univariate Quotient Polynomial Ring in a0
User basis matrix:
[
1 6*b*a0 + 3*b + 1])
]
```

TESTS:

We make sure that [trac ticket #13308](#) is fixed.

```
sage: M = ModularSymbols(Gammal(23), sign=1)
sage: m = M.cuspidal_subspace().hecke_matrix(2)
sage: [j*m==i[0]*j for i in m.eigenspaces_left(format='all') for j in i[1].basis()] # long t
[True, True, True, True, True, True, True, True, True, True, True, True]

sage: B = matrix(QQ, 2, 3, range(6))
sage: B.eigenspaces_left()
Traceback (most recent call last):
...
TypeError: matrix must be square, not 2 x 3

sage: B = matrix(QQ, 4, 4, range(16))
sage: B.eigenspaces_left(format='junk')
Traceback (most recent call last):
...
ValueError: format keyword must be None, 'all' or 'galois', not junk

sage: B.eigenspaces_left(algebraic_multiplicity='garbage')
Traceback (most recent call last):
...
ValueError: algebraic_multiplicity keyword must be True or False
```

eigenspaces_right (*format='all', var='a', algebraic_multiplicity=False*)

Compute the right eigenspaces of a matrix.

Note that `eigenspaces_right()` and `right_eigenspaces()` are identical methods. Here “right” refers to the eigenvectors being placed to the right of the matrix.

INPUT:

- `self` - a square matrix over an exact field. For inexact matrices consult the numerical or symbolic matrix classes.
- `format` - default: None
 - ‘all’ - attempts to create every eigenspace. This will always be possible for matrices with rational entries.
 - ‘galois’ - for each irreducible factor of the characteristic polynomial, a single eigenspace will

be output for a single root/eigenvalue for the irreducible factor.

–None - Uses the ‘all’ format if the base ring is contained in an algebraically closed field which is implemented. Otherwise, uses the ‘galois’ format.

- var - default: ‘a’ - variable name used to represent elements of the root field of each irreducible factor of the characteristic polynomial. If var=‘a’, then the root fields will be in terms of a0, a1, a2, ..., where the numbering runs across all the irreducible factors of the characteristic polynomial, even for linear factors.

- algebraic_multiplicity - default: False - whether or not to include the algebraic multiplicity of each eigenvalue in the output. See the discussion below.

OUTPUT:

If algebraic_multiplicity=False, return a list of pairs (e, V) where e is an eigenvalue of the matrix, and V is the corresponding left eigenspace. For Galois conjugates of eigenvalues, there may be just one representative eigenspace, depending on the format keyword.

If algebraic_multiplicity=True, return a list of triples (e, V, n) where e and V are as above and n is the algebraic multiplicity of the eigenvalue.

Warning: Uses a somewhat naive algorithm (simply factors the characteristic polynomial and computes kernels directly over the extension field).

EXAMPLES:

Right eigenspaces are computed from the left eigenspaces of the transpose of the matrix. As such, there is a greater collection of illustrative examples at the [eigenspaces_left\(\)](#).

We compute the right eigenspaces of a 3×3 rational matrix.

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.eigenspaces_right()
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2 1]),
(-1.348469228349535?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 0.1303061543300932? -0.7393876913398137?]),
(13.34846922834954?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 3.069693845669907? 5.139387691339814?])
]
sage: es = A.eigenspaces_right(format='galois'); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2 1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
1 1/5*a1 + 2/5 2/5*a1 - 1/5])
]
sage: es = A.eigenspaces_right(format='galois', algebraic_multiplicity=True); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
```



```

[ 1 -2  1], 1),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
      1 1/5*a1 + 2/5 2/5*a1 - 1/5], 1)
]
sage: e, v, n = es[0]; v = v.basis()[0]
sage: delta = v*e - A*v
sage: abs(abs(delta)) < 1e-10
True

```

The same computation, but with implicit base change to a field:

```

sage: A = matrix(ZZ, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.eigenspaces_right(format='galois')
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
      1 1/5*a1 + 2/5 2/5*a1 - 1/5])
]

```

This method is only applicable to exact matrices. The “eigenmatrix” routines for matrices with double-precision floating-point entries (RDF, CDF) are the best alternative. (Since some platforms return eigenvectors that are the negatives of those given here, this one example is not tested here.) There are also “eigenmatrix” routines for matrices with symbolic entries.

```

sage: B = matrix(RR, 3, 3, range(9))
sage: B.eigenspaces_right()
Traceback (most recent call last):
...
NotImplementedError: eigenspaces cannot be computed reliably for inexact rings such as Real
consult numerical or symbolic matrix classes for other options

```

```

sage: em = B.change_ring(RDF).eigenmatrix_right()
sage: eigenvalues = em[0]; eigenvalues.dense_matrix().zero_at(1e-15)
[ 13.3484692283      0.0      0.0]
[      0.0 -1.34846922835      0.0]
[      0.0      0.0      0.0]
sage: eigenvectors = em[1]; eigenvectors # not tested
[ 0.164763817...  0.799699663...  0.408248290...]
[ 0.505774475...  0.104205787... -0.816496580...]
[ 0.846785134... -0.591288087...  0.408248290...]

sage: x, y = var('x y')
sage: S = matrix([x, y], [y, 3*x^2])
sage: em = S.eigenmatrix_right()
sage: eigenvalues = em[0]; eigenvalues
[3/2*x^2 + 1/2*x - 1/2*sqrt(9*x^4 - 6*x^3 + x^2 + 4*y^2)
[
0 3/2*x^2 + 1/2*x + 1/2*sqrt(9*x^4 -

sage: eigenvectors = em[1]; eigenvectors
[
1
[1/2*(3*x^2 - x - sqrt(9*x^4 - 6*x^3 + x^2 + 4*y^2))/y 1/2*(3*x^2 - x + sqrt(9*x^4 - 6*x^3 +

```

TESTS:

```
sage: B = matrix(QQ, 2, 3, range(6))
sage: B.eigenspaces_right()
Traceback (most recent call last):
...
TypeError: matrix must be square, not 2 x 3

sage: B = matrix(QQ, 4, 4, range(16))
sage: B.eigenspaces_right(format='junk')
Traceback (most recent call last):
...
ValueError: format keyword must be None, 'all' or 'galois', not junk

sage: B.eigenspaces_right(algebraic_multiplicity='garbage')
Traceback (most recent call last):
...
ValueError: algebraic_multiplicity keyword must be True or False
```

eigenvalues (*extend=True*)

Return a sequence of the eigenvalues of a matrix, with multiplicity. If the eigenvalues are roots of polynomials in QQ, then QQbar elements are returned that represent each separate root.

If the option `extend` is set to `False`, only eigenvalues in the base ring are considered.

EXAMPLES:

```
sage: a = matrix(ZZ, 4, range(16)); a
[ 0  1  2  3]
[ 4  5  6  7]
[ 8  9 10 11]
[12 13 14 15]
sage: sorted(a.eigenvalues(), reverse=True)
[32.46424919657298?, 0, 0, -2.464249196572981?]

sage: a=matrix([(1, 9, -1, -1), (-2, 0, -10, 2), (-1, 0, 15, -2), (0, 1, 0, -1)])
sage: a.eigenvalues()
[-0.9386318578049146?, 15.50655435353258?, 0.2160387521361705? - 4.713151979747493?*I, 0.2160387521361705? + 4.713151979747493?*I]
```

A symmetric matrix `a+a.transpose()` should have real eigenvalues

```
sage: b=a+a.transpose()
sage: ev = b.eigenvalues(); ev
[-8.35066086057957?, -1.107247901349379?, 5.718651326708515?, 33.73925743522043?]
```

The eigenvalues are elements of QQbar, so they really represent exact roots of polynomials, not just approximations.

```
sage: e = ev[0]; e
-8.35066086057957?
sage: p = e.minpoly(); p
x^4 - 30*x^3 - 171*x^2 + 1460*x + 1784
sage: p(e) == 0
True
```

To perform computations on the eigenvalue as an element of a number field, you can always convert back to a number field element.

```
sage: e.as_number_field_element()
(Number Field in a with defining polynomial y^4 - 2*y^3 - 507*y^2 - 3972*y - 4264,
a + 7,
```

Ring morphism:

```
From: Number Field in a with defining polynomial y^4 - 2*y^3 - 507*y^2 - 3972*y - 4264
To:   Algebraic Real Field
Defn: a |--> -15.35066086057957?)
```

Notice the effect of the extend option.

```
sage: M=matrix(QQ, [[0,-1,0],[1,0,0],[0,0,2]])
sage: M.eigenvalues()
[2, -1*I, 1*I]
sage: M.eigenvalues(extend=False)
[2]
```

The method also works for matrices over finite fields:

```
sage: M = matrix(GF(3), [[0,1,1],[1,2,0],[2,0,1]])
sage: ev = M.eigenvalues(); ev
[2*z3, 2*z3 + 2, 2*z3 + 1]
```

Similarly as in the case of $\overline{\mathbb{Q}\mathbb{Q}}$, the eigenvalues belong to some algebraic closure but they can be converted to elements of a finite field:

```
sage: e = ev[0]
sage: e.parent()
Algebraic closure of Finite Field of size 3
sage: e.as_finite_field_element()
(Finite Field in z3 of size 3^3, 2*z3, Ring morphism:
  From: Finite Field in z3 of size 3^3
  To:   Algebraic closure of Finite Field of size 3
  Defn: z3 |--> z3)
```

eigenvectors_left (*extend=True*)

Compute the left eigenvectors of a matrix.

For each distinct eigenvalue, returns a list of the form (e,V,n) where e is the eigenvalue, V is a list of eigenvectors forming a basis for the corresponding left eigenspace, and n is the algebraic multiplicity of the eigenvalue.

If the option extend is set to False, then only the eigenvalues that live in the base ring are considered.

EXAMPLES: We compute the left eigenvectors of a 3×3 rational matrix.

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: es = A.eigenvectors_left(); es
[(0, [
(1, -2, 1)
], 1),
(-1.348469228349535?, [(1, 0.3101020514433644?, -0.3797958971132713?)], 1),
(13.34846922834954?, [(1, 1.289897948556636?, 1.579795897113272?)], 1)]
sage: eval, [evect], mult = es[0]
sage: delta = eval*evect - evect*A
sage: abs(abs(delta)) < 1e-10
True
```

Notice the difference between considering ring extensions or not.

```

sage: M=matrix(QQ,[[0,-1,0],[1,0,0],[0,0,2]])
sage: M.eigenvectors_left()
[(2, [(0, 0, 1)
], 1), (-1*I, [(1, -1*I, 0)], 1), (1*I, [(1, 1*I, 0)], 1)]
sage: M.eigenvectors_left(extend=False)
[(2, [(0, 0, 1)
], 1)]

```

eigenvectors_right (*extend=True*)

Compute the right eigenvectors of a matrix.

For each distinct eigenvalue, returns a list of the form (e,V,n) where e is the eigenvalue, V is a list of eigenvectors forming a basis for the corresponding right eigenspace, and n is the algebraic multiplicity of the eigenvalue. If `extend = True` (the default), this will return eigenspaces over the algebraic closure of the base field where this is implemented; otherwise it will restrict to eigenvalues in the base field.

EXAMPLES: We compute the right eigenvectors of a 3×3 rational matrix.

```

sage: A = matrix(QQ,3,3,range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: es = A.eigenvectors_right(); es
[(0, [(1, -2, 1)
], 1),
(-1.348469228349535?, [(1, 0.1303061543300932?, -0.7393876913398137?)], 1),
(13.34846922834954?, [(1, 3.069693845669907?, 5.139387691339814?)], 1)]
sage: A.eigenvectors_right(extend=False)
[(0, [(1, -2, 1)
], 1)]
sage: eval, [evec], mult = es[0]
sage: delta = eval*evec - A*evec
sage: abs(abs(delta)) < 1e-10
True

```

elementary_divisors ()

If `self` is a matrix over a principal ideal domain R , return elements d_i for $1 \leq i \leq k = \min(r, s)$ where r and s are the number of rows and columns of `self`, such that the cokernel of `self` is isomorphic to

$$R/(d_1) \oplus R/(d_2) \oplus R/(d_k)$$

with $d_i \mid d_{i+1}$ for all i . These are the diagonal entries of the Smith form of `self` (see `smith_form()`).

EXAMPLES:

```

sage: OE = EquationOrder(x^2 - x + 2, 'w')
sage: w = OE.ring_generators()[0]
sage: m = Matrix([ [1, w], [w, 7] ])
sage: m.elementary_divisors()
[1, -w + 9]

```

See Also:

`smith_form()`

elementwise_product (*right*)

Returns the elementwise product of two matrices of the same size (also known as the Hadamard product).

INPUT:

- `right` - the right operand of the product. A matrix of the same size as `self` such that multiplication of elements of the base rings of `self` and `right` is defined, once Sage's coercion model is applied. If the matrices have different sizes, or if multiplication of individual entries cannot be achieved, a `TypeError` will result.

OUTPUT:

A matrix of the same size as `self` and `right`. The entry in location (i, j) of the output is the product of the two entries in location (i, j) of `self` and `right` (in that order).

The parent of the result is determined by Sage's coercion model. If the base rings are identical, then the result is dense or sparse according to this property for the left operand. If the base rings must be adjusted for one, or both, matrices then the result will be sparse only if both operands are sparse. No subdivisions are present in the result.

If the type of the result is not to your liking, or the ring could be "tighter," adjust the operands with `change_ring()`. Adjust sparse versus dense inputs with the methods `sparse_matrix()` and `dense_matrix()`.

EXAMPLES:

```
sage: A = matrix(ZZ, 2, range(6))
sage: B = matrix(QQ, 2, [5, 1/3, 2/7, 11/2, -3/2, 8])
sage: C = A.elementwise_product(B)
sage: C
[ 0 1/3 4/7]
[33/2 -6 40]
sage: C.parent()
Full MatrixSpace of 2 by 3 dense matrices over Rational Field
```

Notice the base ring of the results in the next two examples.

```
sage: D = matrix(ZZ[x], 2, [1+x^2, 2, 3, 4-x])
sage: E = matrix(QQ, 2, [1, 2, 3, 4])
sage: F = D.elementwise_product(E)
sage: F
[ x^2 + 1      4]
[      9 -4*x + 16]
sage: F.parent()
Full MatrixSpace of 2 by 2 dense matrices over Univariate Polynomial Ring in x over Rational
```

```
sage: G = matrix(GF(3), 2, [0, 1, 2, 2])
sage: H = matrix(ZZ, 2, [1, 2, 3, 4])
sage: J = G.elementwise_product(H)
sage: J
[0 2]
[0 2]
sage: J.parent()
Full MatrixSpace of 2 by 2 dense matrices over Finite Field of size 3
```

Non-commutative rings behave as expected. These are the usual quaternions.

```
sage: R.<i,j,k> = QuaternionAlgebra(-1, -1)
sage: A = matrix(R, 2, [1,i,j,k])
sage: B = matrix(R, 2, [i,i,i,i])
sage: A.elementwise_product(B)
[ i -1]
[-k  j]
sage: B.elementwise_product(A)
```

```
[ i -1]
[ k -j]
```

Input that is not a matrix will raise an error.

```
sage: A = random_matrix(ZZ, 5, 10, x=20)
sage: A.elementwise_product(vector(ZZ, [1, 2, 3, 4]))
Traceback (most recent call last):
...
TypeError: operand must be a matrix, not an element of Ambient free module of rank 4 over th
```

Matrices of different sizes for operands will raise an error.

```
sage: A = random_matrix(ZZ, 5, 10, x=20)
sage: B = random_matrix(ZZ, 10, 5, x=40)
sage: A.elementwise_product(B)
Traceback (most recent call last):
...
TypeError: incompatible sizes for matrices from: Full MatrixSpace of 5 by 10 dense matrices
```

Some pairs of rings do not have a common parent where multiplication makes sense. This will raise an error.

```
sage: A = matrix(QQ, 3, range(6))
sage: B = matrix(GF(3), 3, [2]*6)
sage: A.elementwise_product(B)
Traceback (most recent call last):
...
TypeError: no common canonical parent for objects with parents: 'Full MatrixSpace of 3 by 2
```

We illustrate various combinations of sparse and dense matrices. Notice how if base rings are unequal, both operands must be sparse to get a sparse result.

```
sage: A = matrix(ZZ, 5, range(30), sparse=False)
sage: B = matrix(ZZ, 5, range(30), sparse=True)
sage: C = matrix(QQ, 5, range(30), sparse=True)
sage: A.elementwise_product(C).is_dense()
True
sage: B.elementwise_product(C).is_sparse()
True
sage: A.elementwise_product(B).is_dense()
True
sage: B.elementwise_product(A).is_dense()
True
```

TESTS:

Implementation for dense and sparse matrices are different, this will provide a trivial test that they are working identically.

```
sage: A = random_matrix(ZZ, 10, x=1000, sparse=False)
sage: B = random_matrix(ZZ, 10, x=1000, sparse=False)
sage: C = A.sparse_matrix()
sage: D = B.sparse_matrix()
sage: E = A.elementwise_product(B)
sage: F = C.elementwise_product(D)
sage: E.is_dense() and F.is_sparse() and (E == F)
True
```

If the ring has zero divisors, the routines for setting entries of a sparse matrix should intercept zero results

and not create an entry.

```
sage: R = Integers(6)
sage: A = matrix(R, 2, [3, 2, 0, 0], sparse=True)
sage: B = matrix(R, 2, [2, 3, 1, 0], sparse=True)
sage: C = A.elementwise_product(B)
sage: len(C.nonzero_positions()) == 0
True
```

AUTHOR:

•Rob Beezer (2009-07-13)

exp()

Calculate the exponential of this matrix X, which is the matrix

$$e^X = \sum_{k=0}^{\infty} \frac{X^k}{k!}.$$

This function depends on maxima's matrix exponentiation function, which does not deal well with floating point numbers. If the matrix has floating point numbers, they will be rounded automatically to rational numbers during the computation. If you want approximations to the exponential that are calculated numerically, you may get better results by first converting your matrix to RDF or CDF, as shown in the last example.

EXAMPLES:

```
sage: a=matrix([[1,2],[3,4]])
sage: a.exp()
[-1/22*(e^sqrt(33)*(sqrt(33) - 11) - sqrt(33) - 11)*e^(-1/2*sqrt(33) + 5/2) 2/3
[
1/11*(sqrt(33)*e^sqrt(33) - sqrt(33))*e^(-1/2*sqrt(33) + 5/2) 1/22*(e^sqrt(33)

sage: type(a.exp())
<type 'sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense'>

sage: a=matrix([[1/2,2/3],[3/4,4/5]])
sage: a.exp()
[-1/418*(e^(1/10*sqrt(209))*(3*sqrt(209) - 209) - 3*sqrt(209) - 209)*e^(-1/20*sqrt(209) + 13
[
15/418*(sqrt(209)*e^(1/10*sqrt(209)) - sqrt(209))*e^(-1/20*sqrt(209) + 13

sage: a=matrix(RR, [[1,pi.n()], [1e2,1e-2]])
sage: a.exp()
[ 1/11882424341266*(e^(3/1275529100*sqrt(227345670387496707609)))*(11*sqrt(227345670387496707
[
10000/53470909535697*(sqrt(227345670387496707609)*e^(3

sage: a.change_ring(RDF).exp()
[42748127.3153 7368259.24416]
[234538976.138 40426191.4516]
```

extended_echelon_form(*subdivide=False*, ***kwds*)

Returns the echelon form of *self* augmented with an identity matrix.

INPUT:

- subdivide* - default: *False* - determines if the returned matrix is subdivided. See the description of the (output) below for details.
- kwds* - additional keywords that can be passed to the method that computes the echelon form.

OUTPUT:

If *A* is an $m \times n$ matrix, add the m columns of an $m \times m$ identity matrix to the right of *self*. Then row-reduce this $m \times (n + m)$ matrix. This matrix is returned as an immutable matrix.

If `subdivide` is `True` then the returned matrix has a single division among the columns and a single division among the rows. The column subdivision has n columns to the left and m columns to the right. The row division separates the non-zero rows from the zero rows, when restricted to the first n columns.

For a nonsingular matrix the final m columns of the extended echelon form are the inverse of `self`. For a matrix of any size, the final m columns provide a matrix that transforms `self` to echelon form when it multiplies `self` from the left. When the base ring is a field, the uniqueness of reduced row-echelon form implies that this transformation matrix can be taken as the coefficients giving a canonical set of linear combinations of the rows of `self` that yield reduced row-echelon form.

When subdivided as described above, and again over a field, the parts of the subdivision in the upper-left corner and lower-right corner satisfy several interesting relationships with the row space, column space, left kernel and right kernel of `self`. See the examples below.

Note: This method returns an echelon form. If the base ring is not a field, no attempt is made to move to the fraction field. See an example below where the base ring is changed manually.

EXAMPLES:

The four relationships at the end of this example hold in general.

```
sage: A = matrix(QQ, [[2, -1, 7, -1, 0, -3],
...                  [-1, 1, -5, 3, 4, 4],
...                  [2, -1, 7, 0, 2, -2],
...                  [2, 0, 4, 3, 6, 1],
...                  [2, -1, 7, 0, 2, -2]])
sage: E = A.extended_echelon_form(subdivide=True); E
[ 1  0  2  0  0 -1 | 0 -1  0  1 -1]
[ 0  1 -3  0 -2  0 | 0 -2  0  2 -3]
[ 0  0  0  1  2  1 | 0 2/3  0 -1/3 2/3]
[-----+-----]
[ 0  0  0  0  0  0 | 1 2/3  0 -1/3 -1/3]
[ 0  0  0  0  0  0 | 0  0  1  0 -1]
sage: J = E.matrix_from_columns(range(6,11)); J
[ 0 -1  0  1 -1]
[ 0 -2  0  2 -3]
[ 0 2/3  0 -1/3 2/3]
[ 1 2/3  0 -1/3 -1/3]
[ 0  0  1  0 -1]
sage: J*A == A.rref()
True
sage: C = E.subdivision(0,0); C
[ 1  0  2  0  0 -1]
[ 0  1 -3  0 -2  0]
[ 0  0  0  1  2  1]
sage: L = E.subdivision(1,1); L
[ 1 2/3  0 -1/3 -1/3]
[ 0  0  1  0 -1]
sage: A.right_kernel() == C.right_kernel()
True
sage: A.row_space() == C.row_space()
True
sage: A.column_space() == L.right_kernel()
True
sage: A.left_kernel() == L.row_space()
True
```

For a nonsingular matrix, the right half of the extended echelon form is the inverse matrix.


```

sage: B = matrix(QQ, [[1,3,4], [1,4,4], [0,-2,-1]])
sage: E = B.extended_echelon_form()
sage: J = E.matrix_from_columns(range(3,6)); J
[-4  5  4]
[-1  1  0]
[ 2 -2 -1]
sage: J == B.inverse()
True

```

The result is in echelon form, so if the base ring is not a field, the leading entry of each row may not be 1. But you can easily change to the fraction field if necessary.

```

sage: A = matrix(ZZ, [[16, 20, 4, 5, -4, 13, 5],
...                  [10, 13, 3, -3, 7, 11, 6],
...                  [-12, -15, -3, -3, 2, -10, -4],
...                  [10, 13, 3, 3, -1, 9, 4],
...                  [4, 5, 1, 8, -10, 1, -1]])
sage: E = A.extended_echelon_form(subdivide=True); E
[ 2  0 -2  2 -9 -3 -4 | 0  4 -3 -9  4]
[ 0  1  1  2  0  1  1 | 0  1  2  1  1]
[ 0  0  0  3 -4 -1 -1 | 0  3  1 -3  3]
[-----+-----]
[ 0  0  0  0  0  0  0 | 1  6  3 -6  5]
[ 0  0  0  0  0  0  0 | 0  7  2 -7  6]
sage: J = E.matrix_from_columns(range(7,12)); J
[ 0  4 -3 -9  4]
[ 0  1  2  1  1]
[ 0  3  1 -3  3]
[ 1  6  3 -6  5]
[ 0  7  2 -7  6]
sage: J*A == A.echelon_form()
True
sage: B = A.change_ring(QQ)
sage: B.extended_echelon_form(subdivide=True)
[  1      0      -1      0  -19/6  -7/6  -5/3 |  0      0  -89/42  -5/2   1/7]
[  0      1      1      0   8/3   5/3   5/3 |  0      0  34/21    2   -1/7]
[  0      0      0      1  -4/3  -1/3  -1/3 |  0      0   1/21    0   1/7]
[-----+-----]
[  0      0      0      0      0      0      0 |  1      0   9/7    0  -1/7]
[  0      0      0      0      0      0      0 |  0      1   2/7   -1   6/7]

```

Subdivided, or not, the result is immutable, so make a copy if you want to make changes.

```

sage: A = matrix(FiniteField(7), [[2,0,3], [5,5,3], [5,6,5]])
sage: E = A.extended_echelon_form()
sage: E.is_mutable()
False
sage: F = A.extended_echelon_form(subdivide=True)
sage: F
[1 0 0 | 0 4 6]
[0 1 0 | 4 2 2]
[0 0 1 | 5 2 3]
[-----]
sage: F.is_mutable()
False
sage: G = copy(F)
sage: G.subdivide([],[]); G
[1 0 0 0 4 6]
[0 1 0 4 2 2]

```

```
[0 0 1 5 2 3]
```

If you want to determine exactly which algorithm is used to compute the echelon form, you can add additional keywords to pass on to the `echelon_form()` routine employed on the augmented matrix. Sending the flag `include_zero_rows` is a bit silly, since the extended echelon form will never have any zero rows.

```
sage: A = matrix(ZZ, [[1,2], [5,0], [5,9]])
sage: E = A.extended_echelon_form(algorithm='padic', include_zero_rows=False)
sage: E
[ 1  0 36  1 -8]
[ 0  1  5  0 -1]
[ 0  0 45  1 -10]
```

TESTS:

The `subdivide` keyword is checked.

```
sage: A = matrix(QQ, 2, range(4))
sage: A.extended_echelon_form(subdivide='junk')
Traceback (most recent call last):
...
TypeError: subdivide must be True or False, not junk
```

AUTHOR:

•Rob Beezer (2011-02-02)

fcp (*var*='x')

Return the factorization of the characteristic polynomial of self.

INPUT:

•*var* - (default: 'x') name of variable of charpoly

EXAMPLES:

```
sage: M = MatrixSpace(QQ, 3, 3)
sage: A = M([1, 9, -7, 4/5, 4, 3, 6, 4, 3])
sage: A.fcp()
x^3 - 8*x^2 + 209/5*x - 286
sage: A = M([3, 0, -2, 0, -2, 0, 0, 0, 0])
sage: A.fcp('T')
(T - 3) * T * (T + 2)
```

find (*f*, *indices*=False)

Find elements in this matrix satisfying the constraints in the function *f*. The function is evaluated on each element of the matrix .

INPUT:

- f* - a function that is evaluated on each element of this matrix.
- indices* - whether or not to return the indices and elements of this matrix that satisfy the function.

OUTPUT: If *indices* is not specified, return a matrix with 1 where *f* is satisfied and 0 where it is not. If *indices* is specified, return a dictionary containing the elements of this matrix satisfying *f*.

EXAMPLES:

```
sage: M = matrix(4, 3, [1, -1/2, -1, 1, -1, -1/2, -1, 0, 0, 2, 0, 1])
sage: M.find(lambda entry: entry==0)
[0 0 0]
```

```
[0 0 0]
[0 1 1]
[0 1 0]
```

```
sage: M.find(lambda u:u<0)
```

```
[0 1 1]
[0 1 1]
[1 0 0]
[0 0 0]
```

```
sage: M = matrix(4,3,[1, -1/2, -1, 1, -1, -1/2, -1, 0, 0, 2, 0, 1])
```

```
sage: len(M.find(lambda u:u<1 and u>-1,indices=True))
```

```
5
```

```
sage: M.find(lambda u:u!=1/2)
```

```
[1 1 1]
[1 1 1]
[1 1 1]
[1 1 1]
```

```
sage: M.find(lambda u:u>1.2)
```

```
[0 0 0]
[0 0 0]
[0 0 0]
[1 0 0]
```

```
sage: sorted(M.find(lambda u:u!=0,indices=True).keys()) == M.nonzero_positions()
True
```

get_subdivisions()

Returns the current subdivision of self.

EXAMPLES:

```
sage: M = matrix(5, 5, range(25))
```

```
sage: M.subdivisions()
```

```
([], [])
```

```
sage: M.subdivide(2,3)
```

```
sage: M.subdivisions()
```

```
([2], [3])
```

```
sage: N = M.parent()(1)
```

```
sage: N.subdivide(M.subdivisions()); N
```

```
[1 0 0|0 0]
```

```
[0 1 0|0 0]
```

```
[-----+---]
```

```
[0 0 1|0 0]
```

```
[0 0 0|1 0]
```

```
[0 0 0|0 1]
```

gram_schmidt (*orthonormal=False*)

Performs Gram-Schmidt orthogonalization on the rows of the matrix, returning a new matrix and a matrix accomplishing the transformation.

INPUT:

- **self** - a matrix whose rows are to be orthogonalized.
- **orthonormal** - default: `False` - if `True` the returned orthogonal vectors are unit vectors. This keyword is ignored if the matrix is over RDF or CDF and the results are always orthonormal.

OUTPUT:

A pair of matrices, G and M such that if A represents `self`, where the parenthetical properties occur when `orthonormal = True`:

- $A = M \cdot G$
- The rows of G are an orthogonal (resp. orthonormal) set of vectors.
- G times the conjugate-transpose of G is a diagonal (resp. identity) matrix.
- The row space of G equals the row space of A .
- M is a full-rank matrix with zeros above the diagonal.

For exact rings, any zero vectors produced (when the original vectors are linearly dependent) are not output, thus the orthonormal set is linearly independent, and thus a basis for the row space of the original matrix.

Any notion of a Gram-Schmidt procedure requires that the base ring of the matrix has a fraction field implemented. In order to arrive at an orthonormal set, it must be possible to construct square roots of the elements of the base field. In Sage, your best option is the field of algebraic numbers, `QQbar`, which properly contains the rationals and number fields.

If you have an approximate numerical matrix, then this routine requires that your base field be the real and complex double-precision floating point numbers, `RDF` and `CDF`. In this case, the matrix is treated as having full rank, as no attempt is made to recognize linear dependence with approximate calculations.

EXAMPLES:

Inexact Rings, Numerical Matrices:

First, the inexact rings, `CDF` and `RDF`.

```
sage: A = matrix(CDF, [[ 0.6454 + 0.7491*I, -0.8662 + 0.1489*I,  0.7656 - 0.00344*I],
...                    [-0.2913 + 0.8057*I,  0.8321 + 0.8170*I, -0.6744 + 0.9248*I],
...                    [ 0.2554 + 0.3517*I, -0.4454 - 0.1715*I,  0.8325 - 0.6282*I]])
sage: G, M = A.gram_schmidt()
sage: G.round(6) # random signs
[-0.422243 - 0.490087*I  0.566698 - 0.097416*I -0.500882 + 0.002251*I]
[-0.057002 - 0.495035*I -0.35059 - 0.625323*I  0.255514 - 0.415284*I]
[ 0.394105 - 0.421778*I -0.392266 - 0.039345*I -0.352905 + 0.62195*I]
sage: M.round(6) # random
[
    -1.528503          0.0
[  0.459974 - 0.40061*I          -1.741233          0.0]
[-0.934304 + 0.148868*I  0.54833 + 0.073202*I          -0.550725]
sage: (A - M*G).zero_at(10^-12)
[0.0 0.0 0.0]
[0.0 0.0 0.0]
[0.0 0.0 0.0]
sage: (G*G.conjugate_transpose()).zero_at(10^-12)
[1.0 0.0 0.0]
[0.0 1.0 0.0]
[0.0 0.0 1.0]
```

A rectangular matrix. Note that the `orthonormal` keyword is ignored in these cases.

```
sage: A = matrix(RDF, [[-0.978325, -0.751994, 0.925305, -0.200512, 0.420458],
...                    [-0.474877, -0.983403, 0.089836, 0.132218, 0.672965]])
sage: G, M = A.gram_schmidt(orthonormal=False)
sage: G.round(6).zero_at(10^-6)
[-0.607223 -0.466745  0.574315 -0.124453  0.260968]
[ 0.123203 -0.617909 -0.530578  0.289773  0.487368]
```

```

sage: M.round(6).zero_at(10^-6)
[1.611147      0.0]
[0.958116 0.867778]
sage: (A-M*G).zero_at(10^-12)
[0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0]
sage: (G*G.transpose()).round(6).zero_at(10^-6)
[1.0 0.0]
[0.0 1.0]

```

Even though a set of vectors may be linearly dependent, no effort is made to decide when a zero vector is really the result of a relation of linear dependence. So in this regard, input matrices are treated as being of full rank. Try one of the base rings that provide exact results if you need exact results.

```

sage: entries = [[1,1,2], [2,1,3], [3,1,4]]
sage: A = matrix(QQ, entries)
sage: A.rank()
2
sage: B = matrix(RDF, entries)
sage: G, M = B.gram_schmidt()
sage: G.round(6) # random signs
[-0.408248 -0.408248 -0.816497]
[ 0.707107 -0.707107      -0.0]
[ -0.57735  -0.57735   0.57735]
sage: M.round(10) # random
[-2.4494897428      0.0      0.0]
[-3.6742346142  0.7071067812      0.0]
[-4.8989794856  1.4142135624      0.0]
sage: (A - M*G).zero_at(1e-14)
[0.0 0.0 0.0]
[0.0 0.0 0.0]
[0.0 0.0 0.0]
sage: (G*G.transpose()).zero_at(1e-14)
[1.0 0.0 0.0]
[0.0 1.0 0.0]
[0.0 0.0 1.0]

```

Exact Rings, Orthonormalization:

To scale a vector to unit length requires taking a square root, which often takes us outside the base ring. For the integers and the rationals, the field of algebraic numbers ($\overline{\mathbb{Q}\mathbb{Q}}$) is big enough to contain what we need, but the price is that the computations are very slow, hence mostly of value for small cases or instruction. Now we need to use the `orthonormal` keyword.

```

sage: A = matrix(QQbar, [[6, -8, 1],
...                      [4, 1, 3],
...                      [6, 3, 3],
...                      [7, 1, -5],
...                      [7, -3, 5]])
sage: G, M = A.gram_schmidt(orthonormal=True)
sage: G
[ 0.5970223141259934? -0.7960297521679913? 0.09950371902099891?]
[ 0.6063218341690895? 0.5289635311888953? 0.5937772444966257?]
[ 0.5252981913594170? 0.2941669871612735? -0.798453250866314?]
sage: M
[ 10.04987562112089?      0      0]
[ 1.890570661398980? 4.735582601355131?      0]
[ 1.492555785314984? 7.006153332071100? 1.638930357041381?]
[ 2.885607851608969? 1.804330147889395? 7.963520581008761?]

```

```
[ 7.064764050490923?  5.626248468100069? -1.197679876299471?]
sage: M=G-A
[0 0 0]
[0 0 0]
[0 0 0]
[0 0 0]
[0 0 0]
sage: (G*G.transpose()-identity_matrix(3)).norm() < 10^-10
True
sage: G.row_space() == A.row_space()
True
```

After [trac ticket #14047](#), the matrix can also be over the algebraic reals AA:

```
sage: A = matrix(AA, [[6, -8, 1],
...                  [4, 1, 3],
...                  [6, 3, 3],
...                  [7, 1, -5],
...                  [7, -3, 5]])
sage: G, M = A.gram_schmidt(orthonormal=True)
sage: G
[ 0.5970223141259934? -0.7960297521679913? 0.09950371902099891?]
[ 0.6063218341690895? 0.5289635311888953? 0.5937772444966257?]
[ 0.5252981913594170? 0.2941669871612735? -0.798453250866314?]
sage: M
[ 10.04987562112089? 0 0]
[ 1.890570661398980? 4.735582601355131? 0]
[ 1.492555785314984? 7.006153332071100? 1.638930357041381?]
[ 2.885607851608969? 1.804330147889395? 7.963520581008761?]
[ 7.064764050490923? 5.626248468100069? -1.197679876299471?]
```

Starting with complex numbers with rational real and imaginary parts. Note the use of the conjugate-transpose when checking the orthonormality.

```
sage: A = matrix(QQbar, [[ -2, -I - 1, 4*I + 2, -1],
...                     [-4*I, -2*I + 17, 0, 9*I + 1],
...                     [ 1, -2*I - 6, -I + 11, -5*I + 1]])
sage: G, M = A.gram_schmidt(orthonormal=True)
sage: (M*A.conjugate().transpose()).norm() < 10^-10
True
sage: id3 = G*G.conjugate().transpose()
sage: (id3 - identity_matrix(3)).norm() < 10^-10
True
sage: G.row_space() == A.row_space() # long time
True
```

A square matrix with small rank. The zero vectors produced as a result of linear dependence get eliminated, so the rows of G are a basis for the row space of A.

```
sage: A = matrix(QQbar, [[2, -6, 3, 8],
...                     [1, -3, 2, 5],
...                     [0, 0, 2, 4],
...                     [2, -6, 3, 8]])
sage: A.change_ring(QQ).rank()
2
sage: G, M = A.gram_schmidt(orthonormal=True)
sage: G
[ 0.1881441736767195? -0.5644325210301583? 0.2822162605150792? 0.7525766947068779?]
[-0.2502818123591464? 0.750845437077439? 0.3688363550555841? 0.4873908977520218?]
```

```

sage: M
[10.630145812734649? 0]
[ 6.208757731331742? 0.6718090752798139?]
[ 3.574739299857670? 2.687236301119256?]
[10.630145812734649? 0]
sage: M*G-A
[0 0 0 0]
[0 0 0 0]
[0 0 0 0]
[0 0 0 0]
sage: (G*G.transpose()-identity_matrix(2)).norm() < 10^-10
True
sage: G.row_space() == A.row_space()
True

```

Exact Rings, Orthogonalization:

If we forego scaling orthogonal vectors to unit vectors, we can apply Gram-Schmidt to a much greater variety of rings. Use the `orthonormal=False` keyword (or assume it as the default). Note that now the orthogonality check creates a diagonal matrix whose diagonal entries are the squares of the lengths of the vectors.

First, in the rationals, without involving `QQbar`.

```

sage: A = matrix(QQ, [[-1, 3, 2, 2],
...                  [-1, 0, -1, 0],
...                  [-1, -2, -3, -1],
...                  [ 1, 1, 2, 0]])
sage: A.rank()
3
sage: G, M = A.gram_schmidt()
sage: G
[ -1      3      2      2]
[-19/18   1/6   -8/9   1/9]
[ 2/35  -4/35  -2/35   9/35]
sage: M
[ 1      0      0]
[-1/18   1      0]
[-13/18  59/35   1]
[ 1/3  -48/35  -2]
sage: M*G-A
[0 0 0 0]
[0 0 0 0]
[0 0 0 0]
[0 0 0 0]
sage: G*G.transpose()
[ 18      0      0]
[ 0 35/18      0]
[ 0      0  3/35]
sage: G.row_space() == A.row_space()
True

```

A complex subfield of the complex numbers.

```

sage: C.<z> = CyclotomicField(5)
sage: A = matrix(C, [[ -z^3 - 2*z, -z^3 - 1, 2*z^3 - 2*z^2 + 2*z,
...                  [ z^3 - 2*z^2 + 1, -z^3 + 2*z^2 - z - 1, -1,
...                  [-1/2*z^3 - 2*z^2 + z + 1, -z^3 + z - 2, -2*z^3 + 1/2*z^2, 2
sage: G, M = A.gram_schmidt(orthonormal=False)
sage: G

```

```
[
      -z^3 - 2*z
[          155/139*z^3 - 161/139*z^2 + 31/139*z + 13/139          -175/139*z
[-10359/19841*z^3 - 36739/39682*z^2 + 24961/39682*z - 11879/39682 -28209/39682*z^3 - 3671/1
sage: M
[
      1
[          14/139*z^3 + 47/139*z^2 + 145/139*z + 95/139
[          -7/278*z^3 + 199/278*z^2 + 183/139*z + 175/278 -3785/39682*z^3 + 3346/19841*z
sage: M*G - A
[0 0 0 0]
[0 0 0 0]
[0 0 0 0]
sage: G*G.conjugate().transpose()
[          15*z^3 + 15*z^2 + 28
[                                0
[                                463/139*z^3 + 463/139*z^2
[                                0
sage: G.row_space() == A.row_space()
True
```

A slightly edited legacy example.

```
sage: A = matrix(ZZ, 3, [-1, 2, 5, -11, 1, 1, 1, -1, -3]); A
[ -1   2   5]
[-11   1   1]
[  1  -1  -3]
sage: G, mu = A.gram_schmidt()
sage: G
[      -1      2      5]
[ -52/5   -1/5   -2]
[ 2/187  36/187 -14/187]
sage: mu
[      1      0      0]
[  3/5      1      0]
[ -3/5 -7/187      1]
sage: G.row(0) * G.row(1)
0
sage: G.row(0) * G.row(2)
0
sage: G.row(1) * G.row(2)
0
```

The relation between μ and A is as follows.

```
sage: mu * G == A
True
```

hadamard bound()

Return an int n such that the absolute value of the determinant of this matrix is at most 10^n .

This is got using both the row norms and the column norms.

This function only makes sense when the base field can be coerced to the real double field RDF or the MPFR Real Field with 53-bits precision.

EXAMPLES:

```
sage: a = matrix(ZZ, 3, [1,2,5,7,-3,4,2,1,123])
sage: a.hadamard_bound()
4
sage: a.det()
-2014
```



```
sage: 10^4
10000
```

In this example the Hadamard bound has to be computed (automatically) using MPFR instead of doubles, since doubles overflow:

```
sage: a = matrix(ZZ, 2, [2^10000, 3^10000, 2^50, 3^19292])
sage: a.hadamard_bound()
12215
sage: len(str(a.det()))
12215
```

hermite_form(include_zero_rows=True, transformation=False)

Return the Hermite form of self, if it is defined.

INPUT:

- include_zero_rows – bool (default: True); if False the zero rows in the output matrix are deleted.
- transformation – bool (default: False) a matrix U such that $U \cdot \text{self} == H$.

OUTPUT:

- matrix H
- (optional) transformation matrix U such that $U \cdot \text{self} == H$, possibly with zero rows deleted...

EXAMPLES:

```
sage: M = FunctionField(GF(7), 'x').maximal_order()
sage: K.<x> = FunctionField(GF(7)); M = K.maximal_order()
sage: A = matrix(M, 2, 3, [x, 1, 2*x, x, 1+x, 2])
sage: A.hermite_form()
[      x      1      2*x]
[      0      x 5*x + 2]
sage: A.hermite_form(transformation=True)
(
[      x      1      2*x] [1 0]
[      0      x 5*x + 2], [6 1]
)
sage: A = matrix(M, 2, 3, [x, 1, 2*x, 2*x, 2, 4*x])
sage: A.hermite_form(transformation=True, include_zero_rows=False)
([ x  1 2*x], [1 0])
sage: H, U = A.hermite_form(transformation=True, include_zero_rows=True); H, U
(
[ x  1 2*x] [1 0]
[ 0  0  0], [5 1]
)
sage: U*A == H
True
sage: H, U = A.hermite_form(transformation=True, include_zero_rows=False)
sage: U*A
[ x  1 2*x]
sage: U*A == H
True
```

hessenberg_form()

Return Hessenberg form of self.

If the base ring is merely an integral domain (and not a field), the Hessenberg form will (in general) only be defined over the fraction field of the base ring.

EXAMPLES:

```
sage: A = matrix(ZZ,4,[2, 1, 1, -2, 2, 2, -1, -1, -1,1,2,3,4,5,6,7])
sage: h = A.hessenberg_form(); h
[  2  -7/2 -19/5   -2]
[  2   1/2 -17/5   -1]
[  0  25/4  15/2   5/2]
[  0    0  58/5    3]
sage: parent(h)
Full MatrixSpace of 4 by 4 dense matrices over Rational Field
sage: A.hessenbergize()
Traceback (most recent call last):
...
TypeError: Hessenbergize only possible for matrices over a field
```

hessenbergize()

Transform self to Hessenberg form.

The hessenberg form of a matrix A is a matrix that is similar to A , so has the same characteristic polynomial as A , and is upper triangular except possible for entries right below the diagonal.

ALGORITHM: See Henri Cohen's first book.

EXAMPLES:

```
sage: A = matrix(QQ,3, [2, 1, 1, -2, 2, 2, -1, -1, -1])
sage: A.hessenbergize(); A
[  2  3/2   1]
[ -2   3   2]
[  0  -3  -2]

sage: A = matrix(QQ,4, [2, 1, 1, -2, 2, 2, -1, -1, -1,1,2,3,4,5,6,7])
sage: A.hessenbergize(); A
[  2  -7/2 -19/5   -2]
[  2   1/2 -17/5   -1]
[  0  25/4  15/2   5/2]
[  0    0  58/5    3]
```

You can't Hessenbergize an immutable matrix:

```
sage: A = matrix(QQ, 3, [1..9])
sage: A.set_immutable()
sage: A.hessenbergize()
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead (i.e., use copy(M) to change a
```

image()

Return the image of the homomorphism on rows defined by this matrix.

EXAMPLES:

```
sage: MS1 = MatrixSpace(ZZ,4)
sage: MS2 = MatrixSpace(QQ,6)
sage: A = MS1.matrix([3,4,5,6,7,3,8,10,14,5,6,7,2,2,10,9])
sage: B = MS2.random_element()

sage: image(A)
Free module of degree 4 and rank 4 over Integer Ring
Echelon basis matrix:
[  1   0   0  426]
[  0   1   0  518]
```

```
[ 0 0 1 293]
[ 0 0 0 687]
```

```
sage: image(B) == B.row_module()
True
```

indefinite_factorization (*algorithm='symmetric', check=True*)

Decomposes a symmetric or Hermitian matrix into a lower triangular matrix and a diagonal matrix.

INPUT:

- **self** - a square matrix over a ring. The base ring must have an implemented fraction field.
- **algorithm** - default: 'symmetric'. Either 'symmetric' or 'hermitian', according to if the input matrix is symmetric or hermitian.
- **check** - default: True - if True then performs the check that the matrix is consistent with the algorithm keyword.

OUTPUT:

A lower triangular matrix L with each diagonal element equal to 1 and a vector of entries that form a diagonal matrix D . The vector of diagonal entries can be easily used to form the matrix, as demonstrated below in the examples.

For a symmetric matrix, A , these will be related by

$$A = LDL^T$$

If A is Hermitian matrix, then the transpose of L should be replaced by the conjugate-transpose of L .

If any leading principal submatrix (a square submatrix in the upper-left corner) is singular then this method will fail with a `ValueError`.

ALGORITHM:

The algorithm employed only uses field operations, but the computation of each diagonal entry has the potential for division by zero. The number of operations is of order $n^3/3$, which is half the count for an LU decomposition. This makes it an appropriate candidate for solving systems with symmetric (or Hermitian) coefficient matrices.

EXAMPLES:

There is no requirement that a matrix be positive definite, as indicated by the negative entries in the resulting diagonal matrix. The default is that the input matrix is symmetric.

```
sage: A = matrix(QQ, [[ 3, -6, 9, 6, -9],
...                  [-6, 11, -16, -11, 17],
...                  [ 9, -16, 28, 16, -40],
...                  [ 6, -11, 16, 9, -19],
...                  [-9, 17, -40, -19, 68]])
sage: A.is_symmetric()
True
sage: L, d = A.indefinite_factorization()
sage: D = diagonal_matrix(d)
sage: L
[ 1 0 0 0 0]
[-2 1 0 0 0]
[ 3 -2 1 0 0]
[ 2 -1 0 1 0]
[-3 1 -3 1 1]
sage: D
```

```

[ 3  0  0  0  0]
[ 0 -1  0  0  0]
[ 0  0  5  0  0]
[ 0  0  0 -2  0]
[ 0  0  0  0 -1]
sage: A == L*D*L.transpose()
True

```

Optionally, Hermitian matrices can be factored and the result has a similar property (but not identical). Here, the field is all complex numbers with rational real and imaginary parts. As theory predicts, the diagonal entries will be real numbers.

```

sage: C.<I> = QuadraticField(-1)
sage: B = matrix(C, [[ 2, 4 - 2*I, 2 + 2*I],
...                  [4 + 2*I, 8, 10*I],
...                  [2 - 2*I, -10*I, -3]])
sage: B.is_hermitian()
True
sage: L, d = B.indefinite_factorization(algorithm='hermitian')
sage: D = diagonal_matrix(d)
sage: L
[ 1 0 0]
[ I + 2 1 0]
[ -I + 1 2*I + 1 1]
sage: D
[ 2 0 0]
[ 0 -2 0]
[ 0 0 3]
sage: B == L*D*L.conjugate_transpose()
True

```

If a leading principal submatrix has zero determinant, this algorithm will fail. This will never happen with a positive definite matrix.

```

sage: A = matrix(QQ, [[21, 15, 12, -2],
...                  [15, 12, 9, 6],
...                  [12, 9, 7, 3],
...                  [-2, 6, 3, 8]])
sage: A.is_symmetric()
True
sage: A[0:3,0:3].det() == 0
True
sage: A.indefinite_factorization()
Traceback (most recent call last):
...
ValueError: 3x3 leading principal submatrix is singular,
so cannot create indefinite factorization

```

This algorithm only depends on field operations, so outside of the singular submatrix situation, any matrix may be factored. This provides a reasonable alternative to the Cholesky decomposition.

```

sage: F.<a> = FiniteField(5^3)
sage: A = matrix(F,
...              [[ a^2 + 2*a, 4*a^2 + 3*a + 4, 3*a^2 + a, 2*a^2 + 2*a + 1],
...              [4*a^2 + 3*a + 4, 4*a^2 + 2, 3*a, 2*a^2 + 4*a + 2],
...              [ 3*a^2 + a, 3*a, 3*a^2 + 2, 3*a^2 + 2*a + 3],
...              [2*a^2 + 2*a + 1, 2*a^2 + 4*a + 2, 3*a^2 + 2*a + 3, 3*a^2 + 2*a + 4]])
sage: A.is_symmetric()
True

```

```

sage: L, d = A.indefinite_factorization()
sage: D = diagonal_matrix(d)
sage: L
[
      1      0      0      0
[4*a^2 + 4*a + 3      1      0      0]
[      3      4*a^2 + a + 2      1      0]
[      4*a^2 + 4      2*a^2 + 3*a + 3      2*a^2 + 3*a + 1      1]
sage: D
[
      a^2 + 2*a      0      0      0
[      0      2*a^2 + 2*a + 4      0      0]
[      0      0      3*a^2 + 4*a + 3      0]
[      0      0      0      a^2 + 3*a]
sage: A == L*D*L.transpose()
True

```

AUTHOR:

•Rob Beezer (2012-05-24)

integer_kernel (ring='ZZ')

Return the kernel of this matrix over the given ring (which should be either the base ring, or a PID whose fraction field is the base ring).

Assume that the base field of this matrix has a numerator and denominator functions for its elements, e.g., it is the rational numbers or a fraction field. This function computes a basis over the integers for the kernel of self.

If the matrix is not coercible into `QQ`, then the PID itself should be given as a second argument, as in the third example below.

EXAMPLES:

```

sage: A = MatrixSpace(QQ, 4) (range(16))
sage: A.integer_kernel()
Free module of degree 4 and rank 2 over Integer Ring
Echelon basis matrix:
[ 1  0 -3  2]
[ 0  1 -2  1]

```

The integer kernel even makes sense for matrices with fractional entries:

```

sage: A = MatrixSpace(QQ, 2) ([1/2, 0, 0, 0])
sage: A.integer_kernel()
Free module of degree 2 and rank 1 over Integer Ring
Echelon basis matrix:
[0 1]

```

An example over a bigger ring:

```

sage: L.<w> = NumberField(x^2 - x + 2)
sage: OL = L.ring_of_integers()
sage: A = matrix(L, 2, [1, w/2])
sage: A.integer_kernel(OL)
Free module of degree 2 and rank 1 over Maximal Order in Number Field in w with defining pol
Echelon basis matrix:
[ -1 -w + 1]

```

inverse ()

Returns the inverse of self, without changing self.

Note that one can use the Python inverse operator to obtain the inverse as well.

EXAMPLES:

```
sage: m = matrix([[1,2],[3,4]])
sage: m^(-1)
[  -2    1]
[ 3/2 -1/2]
sage: m.inverse()
[  -2    1]
[ 3/2 -1/2]
sage: ~m
[  -2    1]
[ 3/2 -1/2]

sage: m = matrix([[1,2],[3,4]], sparse=True)
sage: m^(-1)
[  -2    1]
[ 3/2 -1/2]
sage: m.inverse()
[  -2    1]
[ 3/2 -1/2]
sage: ~m
[  -2    1]
[ 3/2 -1/2]
sage: m.I
[  -2    1]
[ 3/2 -1/2]
```

TESTS:

```
sage: matrix().inverse()
[]
```

is_bistochastic (*normalized=True*)

Returns True if this matrix is bistochastic.

A matrix is said to be bistochastic if both the sums of the entries of each row and the sum of the entries of each column are equal to 1.

INPUT:

- *normalized* – if set to True (default), checks that the sums are equal to 1. When set to False, checks that the row sums and column sums are all equal to some constant possibly different from 1.

EXAMPLES:

The identity matrix is clearly bistochastic:

```
sage: Matrix(5,5,1).is_bistochastic()
True
```

The same matrix, multiplied by 2, is not bistochastic anymore, though it verifies the constraints of `normalized == False`:

```
sage: (2 * Matrix(5,5,1)).is_bistochastic()
False
sage: (2 * Matrix(5,5,1)).is_bistochastic(normalized = False)
True
```

is_diagonalizable (*base_field=None*)

Determines if the matrix is similar to a diagonal matrix.

INPUT:

- `base_field` - a new field to use for entries of the matrix.

OUTPUT:

If `self` is the matrix A , then it is diagonalizable if there is an invertible matrix S and a diagonal matrix D such that

$$S^{-1}AS = D$$

This routine returns `True` if `self` is diagonalizable. The diagonal entries of the matrix D are the eigenvalues of A . It may be necessary to “increase” the base field to contain all of the eigenvalues. Over the rationals, the field of algebraic numbers, `sage.rings.qqbar` is a good choice.

To obtain the matrices S and D use the `jordan_form()` method with the `transformation=True` keyword.

ALGORITHM:

For each eigenvalue, this routine checks that the algebraic multiplicity (number of occurrences as a root of the characteristic polynomial) is equal to the geometric multiplicity (dimension of the eigenspace), which is sufficient to ensure a basis of eigenvectors for the columns of S .

EXAMPLES:

A matrix that is diagonalizable over the rationals, as evidenced by its Jordan form.

```
sage: A = matrix(QQ, [[-7, 16, 12, 0, 6],
...                  [-9, 15, 0, 12, -27],
...                  [ 9, -8, 11, -12, 51],
...                  [ 3, -4, 0, -1, 9],
...                  [-1, 0, -4, 4, -12]])
sage: A.jordan_form(subdivide=False)
[ 2  0  0  0  0]
[ 0  3  0  0  0]
[ 0  0  3  0  0]
[ 0  0  0 -1  0]
[ 0  0  0  0 -1]
sage: A.is_diagonalizable()
True
```

A matrix that is not diagonalizable over the rationals, as evidenced by its Jordan form.

```
sage: A = matrix(QQ, [[-3, -14, 2, -1, 15],
...                  [4, 6, -2, 3, -8],
...                  [-2, -14, 0, 0, 10],
...                  [3, 13, -2, 0, -11],
...                  [-1, 6, 1, -3, 1]])
sage: A.jordan_form(subdivide=False)
[-1  1  0  0  0]
[ 0 -1  0  0  0]
[ 0  0  2  1  0]
[ 0  0  0  2  1]
[ 0  0  0  0  2]
sage: A.is_diagonalizable()
False
```

If any eigenvalue of a matrix is outside the base ring, then this routine raises an error. However, the ring can be “expanded” to contain the eigenvalues.

```
sage: A = matrix(QQ, [[1, 0, 1, 1, -1],
...                  [0, 1, 0, 4, 8],
...                  [2, 1, 3, 5, 1],
```

```

...          [2, -1, 1, 0, -2],
...          [0, -1, -1, -5, -8]])

sage: [e in QQ for e in A.eigenvalues()]
[False, False, False, False, False]
sage: A.is_diagonalizable()
Traceback (most recent call last):
...
RuntimeError: an eigenvalue of the matrix is not contained in Rational Field

sage: [e in QQbar for e in A.eigenvalues()]
[True, True, True, True, True]
sage: A.is_diagonalizable(base_field=QQbar)
True

```

Other exact fields may be employed, though it will not always be possible to expand their base fields to contain all the eigenvalues.

```

sage: F.<b> = FiniteField(5^2)
sage: A = matrix(F, [[ 4, 3*b + 2, 3*b + 1, 3*b + 4],
...                  [2*b + 1, 4*b, 0, 2],
...                  [ 4*b, b + 2, 2*b + 3, 3],
...                  [ 2*b, 3*b, 4*b + 4, 3*b + 3]])
sage: A.jordan_form()
[ 4      1 | 0      0]
[ 0      4 | 0      0]
[-----+-----]
[ 0      0 | 2*b + 1  1]
[ 0      0 | 0 2*b + 1]
sage: A.is_diagonalizable()
False

sage: F.<c> = QuadraticField(-7)
sage: A = matrix(F, [[ c + 3, 2*c - 2, -2*c + 2, c - 1],
...                  [2*c + 10, 13*c + 15, -13*c - 17, 11*c + 31],
...                  [2*c + 10, 14*c + 10, -14*c - 12, 12*c + 30],
...                  [ 0, 2*c - 2, -2*c + 2, 2*c + 2]])
sage: A.jordan_form(subdivide=False)
[ 4      0      0      0]
[ 0     -2      0      0]
[ 0      0 c + 3      0]
[ 0      0      0 c + 3]
sage: A.is_diagonalizable()
True

```

A trivial matrix is diagonalizable, trivially.

```

sage: A = matrix(QQ, 0, 0)
sage: A.is_diagonalizable()
True

```

A matrix must be square to be diagonalizable.

```

sage: A = matrix(QQ, 3, 4)
sage: A.is_diagonalizable()
False

```

The matrix must have entries from a field, and it must be an exact field.


```

sage: A = matrix(ZZ, 4, range(16))
sage: A.is_diagonalizable()
Traceback (most recent call last):
...
ValueError: matrix entries must be from a field, not Integer Ring

sage: A = matrix(RDF, 4, range(16))
sage: A.is_diagonalizable()
Traceback (most recent call last):
...
ValueError: base field must be exact, not Real Double Field

```

AUTHOR:

•Rob Beezer (2011-04-01)

is_normal()

Returns True if the matrix commutes with its conjugate-transpose.

OUTPUT:

True if the matrix is square and commutes with its conjugate-transpose, and False otherwise.

Normal matrices are precisely those that can be diagonalized by a unitary matrix.

This routine is for matrices over exact rings and so may not work properly for matrices over RR or CC. For matrices with approximate entries, the rings of double-precision floating-point numbers, RDF and CDF, are a better choice since the `sage.matrix.matrix_double_dense.Matrix_double_dense.is_normal()` method has a tolerance parameter. This provides control over allowing for minor discrepancies between entries when checking equality.

The result is cached.

EXAMPLES:

Hermitian matrices are normal.

```

sage: A = matrix(QQ, 5, range(25)) + I*matrix(QQ, 5, range(0, 50, 2))
sage: B = A*A.conjugate_transpose()
sage: B.is_hermitian()
True
sage: B.is_normal()
True

```

Circulant matrices are normal.

```

sage: G = graphs.CirculantGraph(20, [3, 7])
sage: D = digraphs.Circuit(20)
sage: A = 3*D.adjacency_matrix() - 5*G.adjacency_matrix()
sage: A.is_normal()
True

```

Skew-symmetric matrices are normal.

```

sage: A = matrix(QQ, 5, range(25))
sage: B = A - A.transpose()
sage: B.is_skew_symmetric()
True
sage: B.is_normal()
True

```

A small matrix that does not fit into any of the usual categories of normal matrices.

```
sage: A = matrix(ZZ, [[1, -1],
...                  [1, 1]])
sage: A.is_normal()
True
sage: not A.is_hermitian() and not A.is_skew_symmetric()
True
```

Sage has several fields besides the entire complex numbers where conjugation is non-trivial.

```
sage: F.<b> = QuadraticField(-7)
sage: C = matrix(F, [[-2*b - 3, 7*b - 6, -b + 3],
...                  [-2*b - 3, -3*b + 2, -2*b],
...                  [ b + 1, 0, -2]])
sage: C = C*C.conjugate_transpose()
sage: C.is_normal()
True
```

A matrix that is nearly normal, but for a non-real diagonal entry.

```
sage: A = matrix(QQbar, [[ 2, 2-I, 1+4*I],
...                     [ 2+I, 3+I, 2-6*I],
...                     [1-4*I, 2+6*I, 5]])
sage: A.is_normal()
False
sage: A[1,1] = 132
sage: A.is_normal()
True
```

Rectangular matrices are never normal.

```
sage: A = matrix(QQbar, 3, 4)
sage: A.is_normal()
False
```

A square, empty matrix is trivially normal.

```
sage: A = matrix(QQ, 0, 0)
sage: A.is_normal()
True
```

AUTHOR:

•Rob Beezer (2011-03-31)

is_one()

Return True if this matrix is the identity matrix.

EXAMPLES:

```
sage: m = matrix(QQ, 2, range(4))
sage: m.is_one()
False
sage: m = matrix(QQ, 2, [5, 0, 0, 5])
sage: m.is_one()
False
sage: m = matrix(QQ, 2, [1, 0, 0, 1])
sage: m.is_one()
True
sage: m = matrix(QQ, 2, [1, 1, 1, 1])
sage: m.is_one()
False
```

is_permutation_of (*N*, *check=False*)

Return True if there exists a permutation of rows and columns sending *self* to *N* and False otherwise.

INPUT:

- *N* – a matrix.

- **check – boolean (default: False).** If **False** return Boolean indicating whether there exists a permutation of rows and columns sending *self* to *N* and False otherwise. If **True** return a tuple of a Boolean and a permutation mapping *self* to *N* if such a permutation exists, and (False, None) if it does not.

OUTPUT:

A Boolean or a tuple of a Boolean and a permutation.

EXAMPLES:

```
sage: M = matrix(ZZ, [[1, 2, 3], [3, 5, 3], [2, 6, 4]])
sage: M
[1 2 3]
[3 5 3]
[2 6 4]
sage: N = matrix(ZZ, [[1, 2, 3], [2, 6, 4], [3, 5, 3]])
sage: N
[1 2 3]
[2 6 4]
[3 5 3]
sage: M.is_permutation_of(N)
True
```

Some examples that are not permutations of each other:

```
sage: N = matrix(ZZ, [[1, 2, 3], [4, 5, 6], [7, 8, 9]])
sage: N
[1 2 3]
[4 5 6]
[7 8 9]
sage: M.is_permutation_of(N)
False
sage: N = matrix(ZZ, [[1, 2], [3, 4]])
sage: N
[1 2]
[3 4]
sage: M.is_permutation_of(N)
False
```

And for when *check* is **True**:

```
sage: N = matrix(ZZ, [[3, 5, 3], [2, 6, 4], [1, 2, 3]])
sage: N
[3 5 3]
[2 6 4]
[1 2 3]
sage: r = M.is_permutation_of(N, check=True)
sage: r
(True, ((1, 2, 3), ()))
sage: p = r[1]
sage: M.with_permuted_rows_and_columns(*p) == N
True
```

is_positive_definite ()

Determines if a real or symmetric matrix is positive definite.

A square matrix A is positive definite if it is symmetric with real entries or Hermitian with complex entries, and for every non-zero vector \vec{x}

$$\vec{x}^* A \vec{x} > 0$$

Here \vec{x}^* is the conjugate-transpose, which can be simplified to just the transpose in the real case.

ALGORITHM:

A matrix is positive definite if and only if the diagonal entries from the indefinite factorization are all positive (see `indefinite_factorization()`). So this algorithm is of order $n^3/3$ and may be applied to matrices with elements of any ring that has a fraction field contained within the reals or complexes.

INPUT:

Any square matrix.

OUTPUT:

This routine will return `True` if the matrix is square, symmetric or Hermitian, and meets the condition above for the quadratic form.

The base ring for the elements of the matrix needs to have a fraction field implemented and the computations that result from the indefinite factorization must be convertible to real numbers that are comparable to zero.

EXAMPLES:

A real symmetric matrix that is positive definite, as evidenced by the positive entries for the diagonal matrix of the indefinite factorization and the positive determinants of the leading principal submatrices.

```
sage: A = matrix(QQ, [[ 4, -2,  4,  2],
...                  [-2, 10, -2, -7],
...                  [ 4, -2,  8,  4],
...                  [ 2, -7,  4,  7]])
sage: A.is_positive_definite()
True
sage: _, d = A.indefinite_factorization(algorithm='symmetric')
sage: d
(4, 9, 4, 1)
sage: [A[:i,:i].determinant() for i in range(1,A.nrows()+1)]
[4, 36, 144, 144]
```

A real symmetric matrix which is not positive definite, along with a vector that makes the quadratic form negative.

```
sage: A = matrix(QQ, [[ 3, -6,  9,  6, -9],
...                  [-6, 11, -16, -11, 17],
...                  [ 9, -16, 28, 16, -40],
...                  [ 6, -11, 16,  9, -19],
...                  [-9, 17, -40, -19, 68]])
sage: A.is_positive_definite()
False
sage: _, d = A.indefinite_factorization(algorithm='symmetric')
sage: d
(3, -1, 5, -2, -1)
sage: [A[:i,:i].determinant() for i in range(1,A.nrows()+1)]
[3, -3, -15, 30, -30]
sage: u = vector(QQ, [2, 2, 0, 1, 0])
sage: u.row()*A*u
(-3)
```

A real symmetric matrix with a singular leading principal submatrix, that is therefore not positive definite. The vector u makes the quadratic form zero.

```
sage: A = matrix(QQ, [[21, 15, 12, -2],
...                  [15, 12, 9, 6],
...                  [12, 9, 7, 3],
...                  [-2, 6, 3, 8]])
sage: A.is_positive_definite()
False
sage: [A[:,i].determinant() for i in range(1,A.nrows()+1)]
[21, 27, 0, -75]
sage: u = vector(QQ, [1,1,-3,0])
sage: u.row()*A*u
(0)
```

An Hermitian matrix that is positive definite.

```
sage: C.<I> = NumberField(x^2 + 1)
sage: A = matrix(C, [[23, 17*I + 3, 24*I + 25, 21*I],
...                  [-17*I + 3, 38, -69*I + 89, 7*I + 15],
...                  [-24*I + 25, 69*I + 89, 976, 24*I + 6],
...                  [-21*I, -7*I + 15, -24*I + 6, 28]])
sage: A.is_positive_definite()
True
sage: _, d = A.indefinite_factorization(algorithm='hermitian')
sage: d
(23, 576/23, 89885/144, 142130/17977)
sage: [A[:,i].determinant() for i in range(1,A.nrows()+1)]
[23, 576, 359540, 2842600]
```

An Hermitian matrix that is not positive definite. The vector u makes the quadratic form negative.

```
sage: C.<I> = QuadraticField(-1)
sage: B = matrix(C, [[2, 4 - 2*I, 2 + 2*I],
...                  [4 + 2*I, 8, 10*I],
...                  [2 - 2*I, -10*I, -3]])
sage: B.is_positive_definite()
False
sage: _, d = B.indefinite_factorization(algorithm='hermitian')
sage: d
(2, -2, 3)
sage: [B[:,i].determinant() for i in range(1,B.nrows()+1)]
[2, -4, -12]
sage: u = vector(C, [-5 + 10*I, 4 - 3*I, 0])
sage: u.row().conjugate()*B*u
(-50)
```

A positive definite matrix over an algebraically closed field.

```
sage: A = matrix(QQbar, [[2, 4 + 2*I, 6 - 4*I],
...                      [-2*I + 4, 11, 10 - 12*I],
...                      [4*I + 6, 10 + 12*I, 37]])
sage: A.is_positive_definite()
True
sage: [A[:,i].determinant() for i in range(1,A.nrows()+1)]
[2, 2, 6]
```

TESTS:

If the base ring lacks a `conjugate` method, it will be assumed to not be Hermitian and thus symmetric. If the base ring does not make sense as a subfield of the reals, then this routine will fail since comparison

to zero is meaningless.

```
sage: F.<a> = FiniteField(5^3)
sage: a.conjugate()
Traceback (most recent call last):
...
AttributeError: 'sage.rings.finite_rings.element_givaro.FiniteField_givaroElement'
object has no attribute 'conjugate'
sage: A = matrix(F,
...      [[ a^2 + 2*a, 4*a^2 + 3*a + 4,          3*a^2 + a, 2*a^2 + 2*a + 1],
...      [4*a^2 + 3*a + 4,          4*a^2 + 2,          3*a, 2*a^2 + 4*a + 2],
...      [          3*a^2 + a,          3*a,          3*a^2 + 2, 3*a^2 + 2*a + 3],
...      [2*a^2 + 2*a + 1, 2*a^2 + 4*a + 2, 3*a^2 + 2*a + 3, 3*a^2 + 2*a + 4]])
sage: A.is_positive_definite()
Traceback (most recent call last):
...
TypeError: cannot convert computations from
Finite Field in a of size 5^3 into real numbers
```

AUTHOR:

- Rob Beezer (2012-05-24)

is_scalar (*a=None*)

Return True if this matrix is a scalar matrix.

INPUT

- base_ring element *a*, which is chosen as `self[0][0]` if *a* = None

OUTPUT

- whether `self` is a scalar matrix (in fact the scalar matrix *aI* if *a* is input)

EXAMPLES:

```
sage: m = matrix(QQ, 2, range(4))
sage: m.is_scalar(5)
False
sage: m = matrix(QQ, 2, [5, 0, 0, 5])
sage: m.is_scalar(5)
True
sage: m = matrix(QQ, 2, [1, 0, 0, 1])
sage: m.is_scalar(1)
True
sage: m = matrix(QQ, 2, [1, 1, 1, 1])
sage: m.is_scalar(1)
False
```

is_similar (*other, transformation=False*)

Returns True if `self` and `other` are similar, i.e. related by a change-of-basis matrix.

INPUT:

- other* - a matrix, which should be square, and of the same size as `self`, where the entries of the matrix have a fraction field equal to that of `self`. Inexact rings are not supported.
- transformation* - default: False - if True, the output will include the change-of-basis matrix. See below for an exact description.

OUTPUT:

Two matrices, *A* and *B* are similar if there is an invertible matrix *S* such that $A = S^{-1}BS$. *S* can be

interpreted as a change-of-basis matrix if A and B are viewed as matrix representations of the same linear transformation.

When `transformation=False` this method will return `True` if such a matrix S exists, otherwise it will return `False`. When `transformation=True` the method returns a pair. The first part of the pair is `True` or `False` depending on if the matrices are similar and the second part is the change-of-basis matrix, or `None` should it not exist.

When the transformation matrix is requested, it will satisfy `self = S.inverse()*other*S`.

If the base rings for any of the matrices is the integers, the rationals, or the field of algebraic numbers (`QQbar`), then the matrices are converted to have `QQbar` as their base ring prior to checking the equality of the base rings.

It is possible for this routine to fail over most fields, even when the matrices are similar. However, since the field of algebraic numbers is algebraically closed, the routine will always produce a result for matrices with rational entries.

EXAMPLES:

The two matrices in this example were constructed to be similar. The computations happen in the field of algebraic numbers, but we are able to convert the change-of-basis matrix back to the rationals (which may not always be possible).

```
sage: A = matrix(ZZ, [[-5, 2, -11],
...                  [-6, 7, -42],
...                  [0, 1, -6]])
sage: B = matrix(ZZ, [[ 1, 12, 3],
...                  [-1, -6, -1],
...                  [ 0, 6, 1]])
sage: A.is_similar(B)
True
sage: _, T = A.is_similar(B, transformation=True)
sage: T
[ 1.000000000000000? + 0.?e-13*I      0.?e-13 + 0.?e-13*I      0.?e-13 + 0.?e-13*I]
[-0.666666666666667? + 0.?e-13*I  0.166666666666667? + 0.?e-14*I -0.833333333333334? + 0.?e-13*I]
[ 0.666666666666667? + 0.?e-13*I      0.?e-13 + 0.?e-13*I      -0.333333333333334? + 0.?e-13*I]
sage: T.change_ring(QQ)
[ 1      0      0]
[-2/3  1/6 -5/6]
[ 2/3      0 -1/3]
sage: A == T.inverse()*B*T
True
```

Other exact fields are supported.

```
sage: F.<a> = FiniteField(7^2)
sage: A = matrix(F, [[2*a + 5, 6*a + 6, a + 3],
...                  [ a + 3, 2*a + 2, 4*a + 2],
...                  [2*a + 6, 5*a + 5, 3*a]])
sage: B = matrix(F, [[5*a + 5, 6*a + 4, a + 1],
...                  [ a + 5, 4*a + 3, 3*a + 3],
...                  [3*a + 5, a + 4, 5*a + 6]])
sage: A.is_similar(B)
True
sage: B.is_similar(A)
True
sage: _, T = A.is_similar(B, transformation=True)
sage: T
[ 1      0      0]
[6*a + 1 4*a + 3 4*a + 2]
```

```
[6*a + 3 3*a + 5 3*a + 6]
sage: A == T.inverse()*B*T
True
```

Two matrices with different sets of eigenvalues, so they cannot possibly be similar.

```
sage: A = matrix(QQ, [[ 2,  3, -3, -6],
...                  [ 0,  1, -2, -8],
...                  [-3, -3,  4,  3],
...                  [-1, -2,  2,  6]])
sage: B = matrix(QQ, [[ 1,  1,  2,  4],
...                  [-1,  2, -3, -7],
...                  [-2,  3, -4, -7],
...                  [ 0, -1,  0,  0]])
sage: A.eigenvalues() == B.eigenvalues()
False
sage: A.is_similar(B, transformation=True)
(False, None)
```

Similarity is an equivalence relation, so this routine computes a representative of the equivalence class for each matrix, the Jordan form, as provided by `jordan_form()`. The matrices below have identical eigenvalues (as evidenced by equal characteristic polynomials), but slightly different Jordan forms, and hence are not similar.

```
sage: A = matrix(QQ, [[ 19, -7, -29],
...                  [-16, 11,  30],
...                  [ 15, -7, -25]])
sage: B = matrix(QQ, [[-38, -63,  42],
...                  [ 14,  25, -14],
...                  [-14, -21,  18]])
sage: A.charpoly() == B.charpoly()
True
sage: A.jordan_form()
[-3| 0  0]
[--+----]
[ 0| 4  1]
[ 0| 0  4]
sage: B.jordan_form()
[-3| 0| 0]
[--+---]
[ 0| 4| 0]
[--+---]
[ 0| 0| 4]
sage: A.is_similar(B)
False
```

Obtaining the Jordan form requires computing the eigenvalues of the matrix, which may not lie in the field used for entries of the matrix. So the routine first checks the characteristic polynomials - if they are unequal, then the matrices cannot be similar. However, when the characteristic polynomials are equal, we must examine the Jordan form. In this case, the method may fail, EVEN when the matrices are similar. This is not the case for matrices over the integers, rationals or algebraic numbers, since the computations are done in the algebraically closed field of algebraic numbers.

Here is an example where the similarity is obvious, but the routine fails to compute a result.

```
sage: F.<a> = FiniteField(7^2)
sage: C = matrix(F, [[ a + 2, 5*a + 4],
...                  [6*a + 6, 6*a + 4]])
sage: S = matrix(ZZ, [[0, 1],
```



```

...          [1, 0]])
sage: D = S.inverse()*C*S
sage: C.is_similar(D)
Traceback (most recent call last):
...
ValueError: unable to compute Jordan canonical form for a matrix
sage: C.jordan_form()
Traceback (most recent call last):
...
RuntimeError: Some eigenvalue does not exist in Finite Field in a of size 7^2.

```

Inexact rings and fields are also not supported.

```

sage: A = matrix(CDF, 2, 2, range(4))
sage: B = copy(A)
sage: A.is_similar(B)
Traceback (most recent call last):
...
ValueError: unable to compute Jordan canonical form for a matrix

```

Rectangular matrices and mismatched sizes return quickly.

```

sage: A = matrix(3, 2, range(6))
sage: B = copy(A)
sage: A.is_similar(B)
False
sage: A = matrix(2, 2, range(4))
sage: B = matrix(3, 3, range(9))
sage: A.is_similar(B, transformation=True)
(False, None)

```

If the fraction fields of the entries are unequal, it is an error, except in the case when the rationals gets promoted to the algebraic numbers.

```

sage: A = matrix(ZZ, 2, 2, range(4))
sage: B = matrix(GF(2), 2, 2, range(4))
sage: A.is_similar(B, transformation=True)
Traceback (most recent call last):
...
TypeError: matrices need to have entries with identical fraction fields, not Algebraic Field
sage: A = matrix(ZZ, 2, 2, range(4))
sage: B = matrix(QQbar, 2, 2, range(4))
sage: A.is_similar(B)
True

```

Inputs are checked.

```

sage: A = matrix(ZZ, 2, 2, range(4))
sage: A.is_similar('garbage')
Traceback (most recent call last):
...
TypeError: similarity requires a matrix as an argument, not garbage
sage: B = copy(A)
sage: A.is_similar(B, transformation='junk')
Traceback (most recent call last):
...
ValueError: transformation keyword must be True or False, not junk

```

is_unitary()

Returns `True` if the columns of the matrix are an orthonormal basis.

For a matrix with real entries this determines if a matrix is “orthogonal” and for a matrix with complex entries this determines if the matrix is “unitary.”

OUTPUT:

`True` if the matrix is square and its conjugate-transpose is its inverse, and `False` otherwise. In other words, a matrix is orthogonal or unitary if the product of its conjugate-transpose times the matrix is the identity matrix.

For numerical matrices a specialized routine available over RDF and CDF is a good choice.

EXAMPLES:

```
sage: A = matrix(QQbar, [(1/sqrt(5))*(1+i), (1/sqrt(55))*(3+2*I), (1/sqrt(22))*(2+2*I)],
...                      [(1/sqrt(5))*(1-i), (1/sqrt(55))*(2+2*I), (1/sqrt(22))*(-3+I)],
...                      [(1/sqrt(5))*I, (1/sqrt(55))*(3-5*I), (1/sqrt(22))*(-2)])
sage: A.is_unitary()
True
```

A permutation matrix is always orthogonal.

```
sage: sigma = Permutation([1, 3, 4, 5, 2])
sage: P = sigma.to_matrix(); P
[1 0 0 0 0]
[0 0 0 0 1]
[0 1 0 0 0]
[0 0 1 0 0]
[0 0 0 1 0]
sage: P.is_unitary()
True
sage: P.change_ring(GF(3)).is_unitary()
True
sage: P.change_ring(GF(3)).is_unitary()
True
```

A square matrix far from unitary.

```
sage: A = matrix(QQ, 4, range(16))
sage: A.is_unitary()
False
```

Rectangular matrices are never unitary.

```
sage: A = matrix(QQbar, 3, 4)
sage: A.is_unitary()
False
```

jordan_form(*base_ring=None, sparse=False, subdivide=True, transformation=False, eigenvalues=None, check_input=True*)

Compute the Jordan normal form of this square matrix A , if it exists.

This computation is performed in a naive way using the ranks of powers of $A - xI$, where x is an eigenvalue of the matrix A . If desired, a transformation matrix P can be returned, which is such that the Jordan canonical form is given by $P^{-1}AP$.

INPUT:

- `base_ring` - Ring in which to compute the Jordan form.
- `sparse` - (default `False`) If `sparse=True`, return a sparse matrix.

- `subdivide` - (default `True`) If `subdivide=True`, the subdivisions for the Jordan blocks in the matrix are shown.
- `transformation` - (default `False`) If `transformation=True`, computes also the transformation matrix.
- `eigenvalues` - (default `None`) A complete set of roots, with multiplicity, of the characteristic polynomial of A , encoded as a list of pairs, each having the form (r, m) with r a root and m its multiplicity. If this is `None`, then Sage computes this list itself, but this is only possible over base rings in whose quotient fields polynomial factorization is implemented. Over all other rings, providing this list manually is the only way to compute Jordan normal forms.
- `check_input` - (default `True`) A Boolean specifying whether the list `eigenvalues` (if provided) has to be checked for correctness. Set this to `False` for a speedup if the eigenvalues are known to be correct.

NOTES:

Currently, the Jordan normal form is not computed over inexact rings in any but the trivial cases when the matrix is either 0×0 or 1×1 .

In the case of exact rings, this method does not compute any generalized form of the Jordan normal form, but is only able to compute the result if the characteristic polynomial of the matrix splits over the specific base ring.

Note that the base ring must be a field or a ring with an implemented fraction field.

EXAMPLES:

```
sage: a = matrix(ZZ, 4, [1, 0, 0, 0, 0, 1, 0, 0, 1, \
-1, 1, 0, 1, -1, 1, 2]); a
[ 1  0  0  0]
[ 0  1  0  0]
[ 1 -1  1  0]
[ 1 -1  1  2]
sage: a.jordan_form()
[2|0 0|0]
[-+---+-]
[0|1 1|0]
[0|0 1|0]
[-+---+-]
[0|0 0|1]
sage: a.jordan_form(subdivide=False)
[2 0 0 0]
[0 1 1 0]
[0 0 1 0]
[0 0 0 1]
sage: b = matrix(ZZ, 3, range(9)); b
[0 1 2]
[3 4 5]
[6 7 8]
sage: b.jordan_form()
Traceback (most recent call last):
...
RuntimeError: Some eigenvalue does not exist in Rational Field.
sage: b.jordan_form(RealField(15))
Traceback (most recent call last):
...
ValueError: Jordan normal form not implemented over inexact rings.
```

Here we need to specify a field, since the eigenvalues are not defined in the smallest ring containing the matrix entries ([trac ticket #14508](#)):

```
sage: c = matrix([[0,1,0],[0,0,1],[1,0,0]]);
sage: c.jordan_form(CyclotomicField(3))
[      1|      0|      0]
[-----+-----+-----]
[      0|     zeta3|      0]
[-----+-----+-----]
[      0|      0|-zeta3 - 1]
```

If you need the transformation matrix as well as the Jordan form of `self`, then pass the option `transformation=True`. For example:

```
sage: m = matrix([[5,4,2,1],[0,1,-1,-1],[-1,-1,3,0],[1,1,-1,2]]); m
[ 5  4  2  1]
[ 0  1 -1 -1]
[-1 -1  3  0]
[ 1  1 -1  2]
sage: jf, p = m.jordan_form(transformation=True)
sage: jf
[2|0|0 0]
[-+-+---]
[0|1|0 0]
[-+-+---]
[0|0|4 1]
[0|0|0 4]
sage: ~p * m * p
[2 0 0 0]
[0 1 0 0]
[0 0 4 1]
[0 0 0 4]
```

Note that for matrices over inexact rings, we do not attempt to compute the Jordan normal form, since it is not numerically stable:

```
sage: b = matrix(ZZ, 3, 3, range(9))
sage: jf, p = b.jordan_form(RealField(15), transformation=True)
Traceback (most recent call last):
...
ValueError: Jordan normal form not implemented over inexact rings.
```

TESTS:

```
sage: c = matrix(ZZ, 3, [1]*9); c
[1 1 1]
[1 1 1]
[1 1 1]
sage: c.jordan_form(subdivide=False)
[3 0 0]
[0 0 0]
[0 0 0]

sage: evals = [(i,i) for i in range(1,6)]
sage: n = sum(range(1,6))
sage: jf = block_diagonal_matrix([jordan_block(ev,size) for ev,size in evals])
sage: p = random_matrix(ZZ,n,n)
sage: while p.rank() != n: p = random_matrix(ZZ,n,n)
sage: m = p * jf * ~p
sage: mjf, mp = m.jordan_form(transformation=True)
```

```

sage: mjf == jf
True
sage: m = diagonal_matrix([1,1,0,0])
sage: jf,P = m.jordan_form(transformation=True)
sage: jf == ~P*m*P
True

```

We verify that the bug from trac ticket #6942 is fixed:

```

sage: M = Matrix(GF(2), [[1,0,1,0,0,0,1], [1,0,0,1,1,1,0], [1,1,0,1,1,1,1], [1,1,1,0,1,1,1], [1,1,1,1,0,1,1], [1,1,1,1,1,0,1], [1,1,1,1,1,1,0]])
sage: J, T = M.jordan_form(transformation=True)
sage: J
[1 1|0 0|0 0|0]
[0 1|0 0|0 0|0]
[---+---+---+---]
[0 0|1 1|0 0|0]
[0 0|0 1|0 0|0]
[---+---+---+---]
[0 0|0 0|1 1|0]
[0 0|0 0|0 1|0]
[---+---+---+---]
[0 0|0 0|0 0|1]
sage: M * T == T * J
True
sage: T.rank()
7
sage: M.rank()
7

```

We verify that the bug from trac ticket #6932 is fixed:

```

sage: M=Matrix(1,1,[1])
sage: M.jordan_form(transformation=True)
([1], [1])

```

We now go through three 10×10 matrices to exhibit cases where there are multiple blocks of the same size:

```

sage: A = matrix(QQ, [[15, 37/3, -16, -104/3, -29, -7/3, 0, 2/3, -29/3, -1/3], [2, 9, -1, -6, -6, 0, 0, 0, -2, 0],
[ 15, 37/3, -16, -104/3, -29, -7/3, 0, 2/3, -29/3, -1/3]
[ 2, 9, -1, -6, -6, 0, 0, 0, -2, 0]
[ 24, 74/3, -41, -208/3, -58, -23/3, 0, 4/3, -58/3, -2/3]
[ -6, -19, 3, 21, 19, 0, 0, 0, 6, 0]
[ 2, 6, 3, -6, -3, 1, 0, 0, -2, 0]
[ -96, -296/3, 176, 832/3, 232, 101/3, 0, -16/3, 232/3, 8/3]
[ -4, -2/3, 21, 16/3, 4, 14/3, 3, -1/3, 4/3, -25/3]
[ 20, 26/3, -66, -199/3, -42, -41/3, 0, 13/3, -55/3, -2/3]
[ 18, 57, -9, -54, -57, 0, 0, 0, -15, 0]
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 3]
sage: J, T = A.jordan_form(transformation=True); J
[3 1 0|0 0 0|0 0 0|0]
[0 3 1|0 0 0|0 0 0|0]
[0 0 3|0 0 0|0 0 0|0]
[-----+-----+-----+---]
[0 0 0|3 1 0|0 0 0|0]
[0 0 0|0 3 1|0 0 0|0]
[0 0 0|0 0 3|0 0 0|0]
[-----+-----+-----+---]
[0 0 0|0 0 0|3 1 0|0]

```

```

[0 0 0|0 0 0|0 3 1|0]
[0 0 0|0 0 0|0 0 3|0]
[-----+-----+-----+--]
[0 0 0|0 0 0|0 0 0|3]
sage: T * J * T**(-1) == A
True
sage: T.rank()
10

sage: A = matrix(QQ, [[15, 37/3, -16, -14/3, -29, -7/3, 0, 2/3, 1/3, 44/3], [2, 9, -1, 0, -6,
[ 15 37/3 -16 -14/3 -29 -7/3 0 2/3 1/3 44/3]
[ 2 9 -1 0 -6 0 0 0 0 3]
[ 24 74/3 -41 -28/3 -58 -23/3 0 4/3 2/3 88/3]
[ -6 -19 3 3 19 0 0 0 0 -9]
[ 2 6 3 0 -3 1 0 0 0 3]
[ -96 -296/3 176 112/3 232 101/3 0 -16/3 -8/3 -352/3]
[ -4 -2/3 21 16/3 4 14/3 3 -1/3 4/3 -25/3]
[ 20 26/3 -66 -28/3 -42 -41/3 0 13/3 2/3 82/3]
[ 18 57 -9 0 -57 0 0 0 3 28]
[ 0 0 0 0 0 0 0 0 0 3]
sage: J, T = A.jordan_form(transformation=True); J
[3 1 0|0 0 0|0 0|0 0]
[0 3 1|0 0 0|0 0|0 0]
[0 0 3|0 0 0|0 0|0 0]
[-----+-----+-----+---]
[0 0 0|3 1 0|0 0|0 0]
[0 0 0|0 3 1|0 0|0 0]
[0 0 0|0 0 3|0 0|0 0]
[-----+-----+-----+---]
[0 0 0|0 0 0|3 1|0 0]
[0 0 0|0 0 0|0 3|0 0]
[-----+-----+-----+---]
[0 0 0|0 0 0|0 0|3 1]
[0 0 0|0 0 0|0 0|0 3]
sage: T * J * T**(-1) == A
True
sage: T.rank()
10

sage: A = matrix(QQ, [[15, 37/3, -16, -104/3, -29, -7/3, 35, 2/3, -29/3, -1/3], [2, 9, -1, -6,
[ 15 37/3 -16 -104/3 -29 -7/3 35 2/3 -29/3 -1/3]
[ 2 9 -1 -6 -6 0 7 0 -2 0]
[ 24 74/3 -29 -208/3 -58 -14/3 70 4/3 -58/3 -2/3]
[ -6 -19 3 21 19 0 -21 0 6 0]
[ 2 6 -1 -6 -3 0 7 0 -2 0]
[ -96 -296/3 128 832/3 232 65/3 -279 -16/3 232/3 8/3]
[ 0 0 0 0 0 0 3 0 0 0]
[ 20 26/3 -30 -199/3 -42 -14/3 70 13/3 -55/3 -2/3]
[ 18 57 -9 -54 -57 0 63 0 -15 0]
[ 0 0 0 0 0 0 0 0 0 3]
sage: J, T = A.jordan_form(transformation=True); J
[3 1 0|0 0|0 0|0 0|0]
[0 3 1|0 0|0 0|0 0|0]
[0 0 3|0 0|0 0|0 0|0]
[-----+-----+-----+---]
[0 0 0|3 1|0 0|0 0|0]
[0 0 0|0 3|0 0|0 0|0]
[-----+-----+-----+---]

```

```

[0 0 0|0 0|3 1|0 0|0]
[0 0 0|0 0|0 3|0 0|0]
[-----+-----+-----+-----]
[0 0 0|0 0|0 0|3 1|0]
[0 0 0|0 0|0 0|0 3|0]
[-----+-----+-----+-----]
[0 0 0|0 0|0 0|0 0|3]
sage: T * J * T**(-1) == A
True
sage: T.rank()
10

```

Verify that we smoothly move to QQ from ZZ ([trac ticket #12693](#)), i.e. we work in the vector space over the field:

```

sage: M = matrix(((2,2,2),(0,0,0),(-2,-2,-2)))
sage: J, P = M.jordan_form(transformation=True)
sage: J; P
[0 1|0]
[0 0|0]
[---+--]
[0 0|0]
[ 2  1  0]
[ 0  0  1]
[-2  0 -1]
sage: J - ~P * M * P
[0 0 0]
[0 0 0]
[0 0 0]
sage: parent(M)
Full MatrixSpace of 3 by 3 dense matrices over Integer Ring
sage: parent(J) == parent(P) == MatrixSpace(QQ, 3)
True
sage: M.jordan_form(transformation=True) == (M/1).jordan_form(transformation=True)
True

```

By providing eigenvalues ourselves, we can compute the Jordan form even lacking a polynomial factorization algorithm.

```

sage: Qx = PolynomialRing(QQ, 'x11, x12, x13, x21, x22, x23, x31, x32, x33')
sage: x11, x12, x13, x21, x22, x23, x31, x32, x33 = Qx.gens()
sage: M = matrix(Qx, [[0, 0, x31], [0, 0, x21], [0, 0, 0]]) # This is a nilpotent matrix.
sage: M.jordan_form(eigenvalues=[(0, 3)])
[0 1|0]
[0 0|0]
[---+--]
[0 0|0]
sage: M.jordan_form(eigenvalues=[(0, 2)])
Traceback (most recent call last):
...
ValueError: The provided list of eigenvalues is not correct.
sage: M.jordan_form(transformation=True, eigenvalues=[(0, 3)])
(
[0 1|0]
[0 0|0]  [x31  0  1]
[---+--]  [x21  0  0]
[0 0|0], [  0  1  0]
)

```

TESTS:

The base ring for the matrix needs to have a fraction field and it needs to be implemented.

```
sage: A = matrix(Integers(6), 2, 2, range(4))
sage: A.jordan_form()
Traceback (most recent call last):
...
ValueError: Matrix entries must be from a field, not Ring of integers modulo 6
```

kernel (*args, **kws)

Returns the left kernel of this matrix, as a vector space or free module. This is the set of vectors x such that $x \cdot \text{self} = 0$.

Note: For the right kernel, use `right_kernel()`. The method `kernel()` is exactly equal to `left_kernel()`.

INPUT:

- `algorithm` - default: 'default' - a keyword that selects the algorithm employed. Allowable values are:
 - 'default' - allows the algorithm to be chosen automatically
 - 'generic' - naive algorithm usable for matrices over any field
 - 'pari' - PARI library code for matrices over number fields or the integers
 - 'padic' - padic algorithm from IML library for matrices over the rationals and integers
 - 'pluq' - PLUQ matrix factorization for matrices mod 2
- `basis` - default: 'echelon' - a keyword that describes the format of the basis used to construct the left kernel. Allowable values are:
 - 'echelon': the basis matrix is in echelon form
 - 'pivot': each basis vector is computed from the reduced row-echelon form of `self` by placing a single one in a non-pivot column and zeros in the remaining non-pivot columns. Only available for matrices over fields.
 - 'LLL': an LLL-reduced basis. Only available for matrices over the integers.

OUTPUT:

A vector space or free module whose degree equals the number of rows in `self` and contains all the vectors x such that $x \cdot \text{self} = 0$.

If `self` has 0 rows, the kernel has dimension 0, while if `self` has 0 columns the kernel is the entire ambient vector space.

The result is cached. Requesting the left kernel a second time, but with a different basis format will return the cached result with the format from the first computation.

Note: For much more detailed documentation of the various options see `right_kernel()`, since this method just computes the right kernel of the transpose of `self`.

EXAMPLES:

Over the rationals with a basis matrix in echelon form.


```

sage: A = matrix(QQ, [[1, 2, 4, -7, 4],
...                  [1, 1, 0, 2, -1],
...                  [1, 0, 3, -3, 1],
...                  [0, -1, -1, 3, -2],
...                  [0, 0, -1, 2, -1]])
sage: A.left_kernel()
Vector space of degree 5 and dimension 2 over Rational Field
Basis matrix:
[ 1  0 -1  2 -1]
[ 0  1 -1  1 -4]

```

Over a finite field, with a basis matrix in “pivot” format.

```

sage: A = matrix(FiniteField(7), [[5, 0, 5, 2, 4],
...                               [1, 3, 2, 3, 6],
...                               [1, 1, 6, 5, 3],
...                               [2, 5, 6, 0, 0]])
sage: A.kernel(basis='pivot')
Vector space of degree 4 and dimension 2 over Finite Field of size 7
User basis matrix:
[5 2 1 0]
[6 3 0 1]

```

The left kernel of a zero matrix is the entire ambient vector space whose degree equals the number of rows of self (i.e. everything).

```

sage: A = MatrixSpace(QQ, 3, 4)(0)
sage: A.kernel()
Vector space of degree 3 and dimension 3 over Rational Field
Basis matrix:
[1 0 0]
[0 1 0]
[0 0 1]

```

We test matrices with no rows or columns.

```

sage: A = matrix(QQ, 2, 0)
sage: A.left_kernel()
Vector space of degree 2 and dimension 2 over Rational Field
Basis matrix:
[1 0]
[0 1]
sage: A = matrix(QQ, 0, 2)
sage: A.left_kernel()
Vector space of degree 0 and dimension 0 over Rational Field
Basis matrix:
[]

```

The results are cached. Note that requesting a new format for the basis is ignored and the cached copy is returned. Work with a copy if you need a new left kernel, or perhaps investigate the `right_kernel_matrix()` method on the transpose, which does not cache its results and is more flexible.

```

sage: A = matrix(QQ, [[1,1],[2,2]])
sage: K1 = A.left_kernel()
sage: K1
Vector space of degree 2 and dimension 1 over Rational Field
Basis matrix:
[ 1 -1/2]

```

```

sage: K2 = A.left_kernel()
sage: K1 is K2
True
sage: K3 = A.left_kernel(basis='pivot')
sage: K3
Vector space of degree 2 and dimension 1 over Rational Field
Basis matrix:
[ 1 -1/2]
sage: B = copy(A)
sage: K3 = B.left_kernel(basis='pivot')
sage: K3
Vector space of degree 2 and dimension 1 over Rational Field
User basis matrix:
[-2  1]
sage: K3 is K1
False
sage: K3 == K1
True

```

kernel_on(*V*, *poly=None*, *check=True*)

Return the kernel of self restricted to the invariant subspace *V*. The result is a vector subspace of *V*, which is also a subspace of the ambient space.

INPUT:

- *V* - vector subspace
- *check* - (optional) default: True; whether to check that *V* is invariant under the action of self.
- *poly* - (optional) default: None; if not None, compute instead the kernel of *poly*(self) on *V*.

OUTPUT:

- a subspace

Warning: This function does *not* check that *V* is in fact invariant under self if *check* is False. With *check* False this function is much faster.

EXAMPLES:

```

sage: t = matrix(QQ, 4, [39, -10, 0, -12, 0, 2, 0, -1, 0, 1, -2, 0, 0, 2, 0, -2]); t
[ 39 -10  0 -12]
[  0  2  0 -1]
[  0  1 -2  0]
[  0  2  0 -2]
sage: t.fcp()
(x - 39) * (x + 2) * (x^2 - 2)
sage: s = (t-39)*(t^2-2)
sage: V = s.kernel(); V
Vector space of degree 4 and dimension 3 over Rational Field
Basis matrix:
[1 0 0 0]
[0 1 0 0]
[0 0 0 1]
sage: s.restrict(V)
[0 0 0]
[0 0 0]
[0 0 0]
sage: s.kernel_on(V)
Vector space of degree 4 and dimension 3 over Rational Field

```

```

Basis matrix:
[1 0 0 0]
[0 1 0 0]
[0 0 0 1]
sage: k = t-39
sage: k.restrict(V)
[ 0 -10 -12]
[ 0 -37  -1]
[ 0   2 -41]
sage: ker = k.kernel_on(V); ker
Vector space of degree 4 and dimension 1 over Rational Field
Basis matrix:
[ 1 -2/7  0 -2/7]
sage: ker.0 * k
(0, 0, 0, 0)

```

Test that trac ticket #9425 is fixed.

```

sage: V = span([[1/7,0,0],[0,1,0]], ZZ); V
Free module of degree 3 and rank 2 over Integer Ring
Echelon basis matrix:
[1/7  0  0]
[ 0  1  0]
sage: T = matrix(ZZ,3,[1,0,0,0,0,0,0,0]); T
[1 0 0]
[0 0 0]
[0 0 0]
sage: W = T.kernel_on(V); W.basis()
[
(0, 1, 0)
]
sage: W.is_submodule(V)
True

```

left_eigenmatrix()

Return matrices D and P, where D is a diagonal matrix of eigenvalues and P is the corresponding matrix where the rows are corresponding eigenvectors (or zero vectors) so that $P \cdot \text{self} = D \cdot P$.

EXAMPLES:

```

sage: A = matrix(QQ,3,3,range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: D, P = A.eigenmatrix_left()
sage: D
[
0 0 0
0 -1.348469228349535? 0
0 0 13.34846922834954?
]
sage: P
[
1 -2 1
1 0.3101020514433644? -0.3797958971132713?
1 1.289897948556636? 1.579795897113272?
]
sage: P*A == D*P
True

```

Because P is invertible, A is diagonalizable.

```
sage: A == (~P)*D*P
True
```

The matrix P may contain zero rows corresponding to eigenvalues for which the algebraic multiplicity is greater than the geometric multiplicity. In these cases, the matrix is not diagonalizable.

```
sage: A = jordan_block(2,3); A
[2 1 0]
[0 2 1]
[0 0 2]
sage: A = jordan_block(2,3)
sage: D, P = A.eigenmatrix_left()
sage: D
[2 0 0]
[0 2 0]
[0 0 2]
sage: P
[0 0 1]
[0 0 0]
[0 0 0]
sage: P*A == D*P
True
```

TESTS:

For matrices with floating point entries, some platforms will return eigenvectors that are negatives of those returned by the majority of platforms. This test accounts for that possibility. Running this test independently, without adjusting the eigenvectors could indicate this situation on your hardware.

```
sage: A = matrix(QQ, 3, 3, range(9))
sage: em = A.change_ring(RDF).eigenmatrix_left()
sage: evals = em[0]; evals.dense_matrix().zero_at(2e-15)
[ 13.3484692283      0.0      0.0]
[      0.0 -1.34846922835      0.0]
[      0.0      0.0      0.0]
sage: evectors = em[1];
sage: for i in range(3):
...     scale = evectors[i,0].sign()
...     evectors.rescale_row(i, scale)
sage: evectors
[ 0.440242867...  0.567868371...  0.695493875...]
[ 0.897878732...  0.278434036... -0.341010658...]
[ 0.408248290... -0.816496580...  0.408248290...]
```

left_eigenspaces (*format='all', var='a', algebraic_multiplicity=False*)

Compute the left eigenspaces of a matrix.

Note that `eigenspaces_left()` and `left_eigenspaces()` are identical methods. Here “left” refers to the eigenvectors being placed to the left of the matrix.

INPUT:

- `self` - a square matrix over an exact field. For inexact matrices consult the numerical or symbolic matrix classes.
- `format` - default: None
 - ‘all’ - attempts to create every eigenspace. This will always be possible for matrices with rational entries.

- 'galois' - for each irreducible factor of the characteristic polynomial, a single eigenspace will be output for a single root/eigenvalue for the irreducible factor.
- None - Uses the 'all' format if the base ring is contained in an algebraically closed field which is implemented. Otherwise, uses the 'galois' format.
- var - default: 'a' - variable name used to represent elements of the root field of each irreducible factor of the characteristic polynomial. If var='a', then the root fields will be in terms of a0, a1, a2, ..., where the numbering runs across all the irreducible factors of the characteristic polynomial, even for linear factors.
- algebraic_multiplicity - default: False - whether or not to include the algebraic multiplicity of each eigenvalue in the output. See the discussion below.

OUTPUT:

If algebraic_multiplicity=False, return a list of pairs (e, V) where e is an eigenvalue of the matrix, and V is the corresponding left eigenspace. For Galois conjugates of eigenvalues, there may be just one representative eigenspace, depending on the format keyword.

If algebraic_multiplicity=True, return a list of triples (e, V, n) where e and V are as above and n is the algebraic multiplicity of the eigenvalue.

Warning: Uses a somewhat naive algorithm (simply factors the characteristic polynomial and computes kernels directly over the extension field).

EXAMPLES:

We compute the left eigenspaces of a 3×3 rational matrix. First, we request *all* of the eigenvalues, so the results are in the field of algebraic numbers, \overline{QQ} . Then we request just one eigenspace per irreducible factor of the characteristic polynomial with the *galois* keyword.

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: es = A.eigenspaces_left(format='all'); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2 1]),
(-1.348469228349535?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 0.3101020514433644? -0.3797958971132713?]),
(13.34846922834954?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 1.289897948556636? 1.579795897113272?])
]

sage: es = A.eigenspaces_left(format='galois'); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2 1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
1 1/15*a1 + 2/5 2/15*a1 - 1/5])
]
sage: es = A.eigenspaces_left(format='galois', algebraic_multiplicity=True); es
[
```

```

(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1], 1),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
      1 1/15*a1 + 2/5 2/15*a1 - 1/5], 1)
]
sage: e, v, n = es[0]; v = v.basis()[0]
sage: delta = e*v - v*A
sage: abs(abs(delta)) < 1e-10
True

```

The same computation, but with implicit base change to a field.

```

sage: A = matrix(ZZ, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.eigenspaces_left(format='galois')
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[
      1 1/15*a1 + 2/5 2/15*a1 - 1/5])
]

```

We compute the left eigenspaces of the matrix of the Hecke operator T_2 on level 43 modular symbols, both with all eigenvalues (the default) and with one subspace per factor.

```

sage: A = ModularSymbols(43).T(2).matrix(); A
[ 3  0  0  0  0  0 -1]
[ 0 -2  1  0  0  0  0]
[ 0 -1  1  1  0 -1  0]
[ 0 -1  0 -1  2 -1  1]
[ 0 -1  0  1  1 -1  1]
[ 0  0 -2  0  2 -2  1]
[ 0  0 -1  0  1  0 -1]
sage: A.base_ring()
Rational Field
sage: f = A.charpoly(); f
x^7 + x^6 - 12*x^5 - 16*x^4 + 36*x^3 + 52*x^2 - 32*x - 48
sage: factor(f)
(x - 3) * (x + 2)^2 * (x^2 - 2)^2
sage: A.eigenspaces_left(algebraic_multiplicity=True)
[
(3, Vector space of degree 7 and dimension 1 over Rational Field
User basis matrix:
[  1    0  1/7    0 -1/7    0 -2/7], 1),
(-2, Vector space of degree 7 and dimension 2 over Rational Field
User basis matrix:
[ 0  1  0  1 -1  1 -1]
[ 0  0  1  0 -1  2 -1], 2),
(-1.414213562373095?, Vector space of degree 7 and dimension 2 over Algebraic Field
User basis matrix:
[
      0      1      0      -1 0.414213562
      0      0      1      0      0
(1.414213562373095?, Vector space of degree 7 and dimension 2 over Algebraic Field

```

```

User basis matrix:
[      0      1      0      -1  -2.414
 [      0      0      1      0
]
sage: A.eigenspaces_left(format='galois', algebraic_multiplicity=True)
[
(3, Vector space of degree 7 and dimension 1 over Rational Field
User basis matrix:
[ 1 0 1/7 0 -1/7 0 -2/7], 1),
(-2, Vector space of degree 7 and dimension 2 over Rational Field
User basis matrix:
[ 0 1 0 1 -1 1 -1]
[ 0 0 1 0 -1 2 -1], 2),
(a2, Vector space of degree 7 and dimension 2 over Number Field in a2 with defining polynomial
User basis matrix:
[ 0 1 0 -1 -a2 - 1 1 -1]
[ 0 0 1 0 -1 -1 0 -a2 + 1], 2)
]

```

Next we compute the left eigenspaces over the finite field of order 11.

```

sage: A = ModularSymbols(43, base_ring=GF(11), sign=1).T(2).matrix(); A
[ 3 9 0 0]
[ 0 9 0 1]
[ 0 10 9 2]
[ 0 9 0 2]
sage: A.base_ring()
Finite Field of size 11
sage: A.charpoly()
x^4 + 10*x^3 + 3*x^2 + 2*x + 1
sage: A.eigenspaces_left(format='galois', var = 'beta')
[
(9, Vector space of degree 4 and dimension 1 over Finite Field of size 11
User basis matrix:
[0 0 1 5]),
(3, Vector space of degree 4 and dimension 1 over Finite Field of size 11
User basis matrix:
[1 6 0 6]),
(beta2, Vector space of degree 4 and dimension 1 over Univariate Quotient Polynomial Ring in
User basis matrix:
[ 0 1 0 5*beta2 + 10])
]

```

This method is only applicable to exact matrices. The “eigenmatrix” routines for matrices with double-precision floating-point entries (RDF, CDF) are the best alternative. (Since some platforms return eigenvectors that are the negatives of those given here, this one example is not tested here.) There are also “eigenmatrix” routines for matrices with symbolic entries.

```

sage: A = matrix(QQ, 3, 3, range(9))
sage: A.change_ring(RR).eigenspaces_left()
Traceback (most recent call last):
...
NotImplementedError: eigenspaces cannot be computed reliably for inexact rings such as Real
consult numerical or symbolic matrix classes for other options

sage: em = A.change_ring(RDF).eigenmatrix_left()
sage: eigenvalues = em[0]; eigenvalues.dense_matrix().zero_at(2e-15)
[ 13.3484692283 0.0 0.0]
[ 0.0 -1.34846922835 0.0]

```

```

[
    0.0          0.0          0.0]
sage: eigenvectors = em[1]; eigenvectors # not tested
[ 0.440242867...  0.567868371...  0.695493875...]
[ 0.897878732...  0.278434036... -0.341010658...]
[ 0.408248290... -0.816496580...  0.408248290...]

sage: x, y = var('x y')
sage: S = matrix([[x, y], [y, 3*x^2]])
sage: em = S.eigenmatrix_left()
sage: eigenvalues = em[0]; eigenvalues
[3/2*x^2 + 1/2*x - 1/2*sqrt(9*x^4 - 6*x^3 + x^2 + 4*y^2)
 0 3/2*x^2 + 1/2*x + 1/2*sqrt(9*x^4 -
sage: eigenvectors = em[1]; eigenvectors
[
 1 1/2*(3*x^2 - x - sqrt(9*x^4 - 6*x^3 +
 1 1/2*(3*x^2 - x + sqrt(9*x^4 - 6*x^3 +

```

A request for 'all' the eigenvalues, when it is not possible, will raise an error. Using the 'galois' format option is more likely to be successful.

```

sage: F.<b> = FiniteField(11^2)
sage: A = matrix(F, [[b + 1, b + 1], [10*b + 4, 5*b + 4]])
sage: A.eigenspaces_left(format='all')
Traceback (most recent call last):
...
NotImplementedError: unable to construct eigenspaces for eigenvalues outside the base field,
try the keyword option: format='galois'

sage: A.eigenspaces_left(format='galois')
[
(a0, Vector space of degree 2 and dimension 1 over Univariate Quotient Polynomial Ring in a0
User basis matrix:
[
    1 6*b*a0 + 3*b + 1])
]

```

TESTS:

We make sure that [trac ticket #13308](#) is fixed.

```

sage: M = ModularSymbols(Gamma1(23), sign=1)
sage: m = M.cuspidal_subspace().hecke_matrix(2)
sage: [j*m==i[0]*j for i in m.eigenspaces_left(format='all') for j in i[1].basis()] # long t
[True, True, True, True, True, True, True, True, True, True, True, True]

sage: B = matrix(QQ, 2, 3, range(6))
sage: B.eigenspaces_left()
Traceback (most recent call last):
...
TypeError: matrix must be square, not 2 x 3

sage: B = matrix(QQ, 4, 4, range(16))
sage: B.eigenspaces_left(format='junk')
Traceback (most recent call last):
...
ValueError: format keyword must be None, 'all' or 'galois', not junk

sage: B.eigenspaces_left(algebraic_multiplicity='garbage')
Traceback (most recent call last):
...
ValueError: algebraic_multiplicity keyword must be True or False

```


left_eigenvectors (*extend=True*)

Compute the left eigenvectors of a matrix.

For each distinct eigenvalue, returns a list of the form (e,V,n) where e is the eigenvalue, V is a list of eigenvectors forming a basis for the corresponding left eigenspace, and n is the algebraic multiplicity of the eigenvalue.

If the option `extend` is set to `False`, then only the eigenvalues that live in the base ring are considered.

EXAMPLES: We compute the left eigenvectors of a 3×3 rational matrix.

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: es = A.eigenvectors_left(); es
[(0, [
(1, -2, 1)
], 1),
(-1.348469228349535?, [(1, 0.3101020514433644?, -0.3797958971132713?)], 1),
(13.34846922834954?, [(1, 1.289897948556636?, 1.579795897113272?)], 1)]
sage: eval, [evec], mult = es[0]
sage: delta = eval*evec - evec*A
sage: abs(abs(delta)) < 1e-10
True
```

Notice the difference between considering ring extensions or not.

```
sage: M=matrix(QQ, [[0,-1,0],[1,0,0],[0,0,2]])
sage: M.eigenvectors_left()
[(2, [
(0, 0, 1)
], 1), (-1*I, [(1, -1*I, 0)], 1), (1*I, [(1, 1*I, 0)], 1)]
sage: M.eigenvectors_left(extend=False)
[(2, [
(0, 0, 1)
], 1)]
```

left_kernel (**args, **kws*)

Returns the left kernel of this matrix, as a vector space or free module. This is the set of vectors x such that $x \cdot \text{self} = 0$.

Note: For the right kernel, use `right_kernel()`. The method `kernel()` is exactly equal to `left_kernel()`.

INPUT:

- `algorithm` - default: 'default' - a keyword that selects the algorithm employed. Allowable values are:

- 'default' - allows the algorithm to be chosen automatically
- 'generic' - naive algorithm usable for matrices over any field
- 'pari' - PARI library code for matrices over number fields or the integers
- 'padic' - padic algorithm from IML library for matrices over the rationals and integers
- 'pluq' - PLUQ matrix factorization for matrices mod 2

•**basis** - default: 'echelon' - a keyword that describes the format of the basis used to construct the left kernel. Allowable values are:

- 'echelon': the basis matrix is in echelon form
- 'pivot': each basis vector is computed from the reduced row-echelon form of `self` by placing a single one in a non-pivot column and zeros in the remaining non-pivot columns. Only available for matrices over fields.
- 'LLL': an LLL-reduced basis. Only available for matrices over the integers.

OUTPUT:

A vector space or free module whose degree equals the number of rows in `self` and contains all the vectors `x` such that `x*self = 0`.

If `self` has 0 rows, the kernel has dimension 0, while if `self` has 0 columns the kernel is the entire ambient vector space.

The result is cached. Requesting the left kernel a second time, but with a different basis format will return the cached result with the format from the first computation.

Note: For much more detailed documentation of the various options see `right_kernel()`, since this method just computes the right kernel of the transpose of `self`.

EXAMPLES:

Over the rationals with a basis matrix in echelon form.

```
sage: A = matrix(QQ, [[1, 2, 4, -7, 4],
...                  [1, 1, 0, 2, -1],
...                  [1, 0, 3, -3, 1],
...                  [0, -1, -1, 3, -2],
...                  [0, 0, -1, 2, -1]])
sage: A.left_kernel()
Vector space of degree 5 and dimension 2 over Rational Field
Basis matrix:
[ 1  0 -1  2 -1]
[ 0  1 -1  1 -4]
```

Over a finite field, with a basis matrix in "pivot" format.

```
sage: A = matrix(FiniteField(7), [[5, 0, 5, 2, 4],
...                               [1, 3, 2, 3, 6],
...                               [1, 1, 6, 5, 3],
...                               [2, 5, 6, 0, 0]])
sage: A.kernel(basis='pivot')
Vector space of degree 4 and dimension 2 over Finite Field of size 7
User basis matrix:
[5 2 1 0]
[6 3 0 1]
```

The left kernel of a zero matrix is the entire ambient vector space whose degree equals the number of rows of `self` (i.e. everything).

```
sage: A = MatrixSpace(QQ, 3, 4)(0)
sage: A.kernel()
Vector space of degree 3 and dimension 3 over Rational Field
Basis matrix:
[1 0 0]
```

```
[0 1 0]
[0 0 1]
```

We test matrices with no rows or columns.

```
sage: A = matrix(QQ, 2, 0)
sage: A.left_kernel()
Vector space of degree 2 and dimension 2 over Rational Field
Basis matrix:
[1 0]
[0 1]
sage: A = matrix(QQ, 0, 2)
sage: A.left_kernel()
Vector space of degree 0 and dimension 0 over Rational Field
Basis matrix:
[]
```

The results are cached. Note that requesting a new format for the basis is ignored and the cached copy is returned. Work with a copy if you need a new left kernel, or perhaps investigate the `right_kernel_matrix()` method on the transpose, which does not cache its results and is more flexible.

```
sage: A = matrix(QQ, [[1,1],[2,2]])
sage: K1 = A.left_kernel()
sage: K1
Vector space of degree 2 and dimension 1 over Rational Field
Basis matrix:
[ 1 -1/2]
sage: K2 = A.left_kernel()
sage: K1 is K2
True
sage: K3 = A.left_kernel(basis='pivot')
sage: K3
Vector space of degree 2 and dimension 1 over Rational Field
Basis matrix:
[ 1 -1/2]
sage: B = copy(A)
sage: K3 = B.left_kernel(basis='pivot')
sage: K3
Vector space of degree 2 and dimension 1 over Rational Field
User basis matrix:
[-2 1]
sage: K3 is K1
False
sage: K3 == K1
True
```

`left_nullity()`

Return the (left) nullity of this matrix, which is the dimension of the (left) kernel of this matrix acting from the right on row vectors.

EXAMPLES:

```
sage: M = Matrix(QQ, [[1,0,0,1],[0,1,1,0],[1,1,1,0]])
sage: M.nullity()
0
sage: M.left_nullity()
0
```

```

sage: A = M.transpose()
sage: A.nullity()
1
sage: A.left_nullity()
1

sage: M = M.change_ring(ZZ)
sage: M.nullity()
0
sage: A = M.transpose()
sage: A.nullity()
1

```

matrix_window (row=0, col=0, nrows=-1, ncols=-1, check=1)

Return the requested matrix window.

EXAMPLES:

```

sage: A = matrix(QQ, 3, range(9))
sage: A.matrix_window(1,1, 2, 1)
Matrix window of size 2 x 1 at (1,1):
[0 1 2]
[3 4 5]
[6 7 8]

```

We test the optional check flag.

```

sage: matrix([1]).matrix_window(0,1,1,1, check=False)
Matrix window of size 1 x 1 at (0,1):
[1]
sage: matrix([1]).matrix_window(0,1,1,1)
Traceback (most recent call last):
...
IndexError: matrix window index out of range

```

Another test of bounds checking:

```

sage: matrix([1]).matrix_window(1,1,1,1)
Traceback (most recent call last):
...
IndexError: matrix window index out of range

```

maxspin (v)

Computes the largest integer n such that the list of vectors $S = [v, v*A, \dots, v*A^n]$ are linearly independent, and returns that list.

INPUT:

- self - Matrix
- v - Vector

OUTPUT:

- list - list of Vectors

ALGORITHM: The current implementation just adds vectors to a vector space until the dimension doesn't grow. This could be optimized by directly using matrices and doing an efficient Echelon form. Also, when the base is \mathbb{Q} , maybe we could simultaneously keep track of what is going on in the reduction modulo p , which might make things much faster.

EXAMPLES:

```

sage: t = matrix(QQ, 3, range(9)); t
[0 1 2]
[3 4 5]
[6 7 8]
sage: v = (QQ^3).0
sage: t.maxspin(v)
[(1, 0, 0), (0, 1, 2), (15, 18, 21)]
sage: k = t.kernel(); k
Vector space of degree 3 and dimension 1 over Rational Field
Basis matrix:
[ 1 -2  1]
sage: t.maxspin(k.0)
[(1, -2, 1)]

```

minimal_polynomial (*var='x', **kws*)

This is a synonym for `self.minpoly`

EXAMPLES:

```

sage: a = matrix(QQ, 4, range(16))
sage: a.minimal_polynomial('z')
z^3 - 30*z^2 - 80*z
sage: a.minpoly()
x^3 - 30*x^2 - 80*x

```

minors (*k*)

Return the list of all $k \times k$ minors of self.

Let A be an $m \times n$ matrix and k an integer with $0 \leq k, k \leq m$ and $k \leq n$. A $k \times k$ minor of A is the determinant of a $k \times k$ matrix obtained from A by deleting $m - k$ rows and $n - k$ columns. There are no $k \times k$ minors of A if k is larger than either m or n .

The returned list is sorted in lexicographical row major ordering, e.g., if A is a 3×3 matrix then the minors returned are with these rows/columns: $[[0, 1] \times [0, 1], [0, 1] \times [0, 2], [0, 1] \times [1, 2], [0, 2] \times [0, 1], [0, 2] \times [0, 2], [0, 2] \times [1, 2], [1, 2] \times [0, 1], [1, 2] \times [0, 2], [1, 2] \times [1, 2]]$.

INPUT:

- k – integer

EXAMPLES:

```

sage: A = Matrix(ZZ, 2, 3, [1, 2, 3, 4, 5, 6]); A
[1 2 3]
[4 5 6]
sage: A.minors(2)
[-3, -6, -3]
sage: A.minors(1)
[1, 2, 3, 4, 5, 6]
sage: A.minors(0)
[1]
sage: A.minors(5)
[]

sage: k = GF(37)
sage: P.<x0,x1,x2> = PolynomialRing(k)
sage: A = Matrix(P, 2, 3, [x0*x1, x0, x1, x2, x2 + 16, x2 + 5*x1 ])
sage: A.minors(2)
[x0*x1*x2 + 16*x0*x1 - x0*x2, 5*x0*x1^2 + x0*x1*x2 - x1*x2, 5*x0*x1 + x0*x2 - x1*x2 - 16*x1]

```

minpoly (*var='x', **kws*)

Return the minimal polynomial of self.

This uses a simplistic - and potentially very very slow - algorithm that involves computing kernels to determine the powers of the factors of the charpoly that divide the minpoly.

EXAMPLES:

```
sage: A = matrix(GF(9, 'c'), 4, [1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 5, 0, 0, 0, 0, 5])
sage: factor(A.minpoly())
(x + 1) * (x + 2)^2
sage: A.minpoly() (A) == 0
True
sage: factor(A.charpoly())
(x + 1)^2 * (x + 2)^2
```

The default variable name is x , but you can specify another name:

```
sage: factor(A.minpoly('y'))
(y + 1) * (y + 2)^2
```

We can take the minimal polynomial of symbolic matrices:

```
sage: t = var('t')
sage: m = matrix(2, [1, 2, 4, t])
sage: m.minimal_polynomial()
x^2 + (-t - 1)*x + t - 8
```

n (*prec=None, digits=None, algorithm=None*)

Return a numerical approximation of self as either a real or complex number with at least the requested number of bits or digits of precision.

INPUT:

- *prec* - an integer: the number of bits of precision
- *digits* - an integer: digits of precision

OUTPUT: A matrix coerced to a real or complex field with *prec* bits of precision.

EXAMPLES:

```
sage: d = matrix([[3, 0], [0, sqrt(2)]]) ;
sage: b = matrix([[1, -1], [2, 2]]) ; e = b * d * b.inverse(); e
[ 1/2*sqrt(2) + 3/2 -1/4*sqrt(2) + 3/4]
[      -sqrt(2) + 3  1/2*sqrt(2) + 3/2]
```

```
sage: e.numerical_approx(53)
[ 2.20710678118655 0.396446609406726]
[ 1.58578643762690 2.20710678118655]
```

```
sage: e.numerical_approx(20)
[ 2.2071 0.39645]
[ 1.5858 2.2071]
```

```
sage: (e-I).numerical_approx(20)
[2.2071 - 1.0000*I      0.39645]
[      1.5858 2.2071 - 1.0000*I]
```

```
sage: M=matrix(QQ, 4, [i/(i+1) for i in range(12)]); M
[ 0 1/2 2/3]
[ 3/4 4/5 5/6]
```

```
[ 6/7  7/8  8/9]
[ 9/10 10/11 11/12]
```

```
sage: M.numerical_approx()
[0.000000000000000 0.500000000000000 0.666666666666667]
[0.750000000000000 0.800000000000000 0.833333333333333]
[0.857142857142857 0.875000000000000 0.888888888888889]
[0.900000000000000 0.909090909090909 0.916666666666667]
```

```
sage: matrix(SR, 2, 2, range(4)).n()
[0.000000000000000 1.000000000000000]
[ 2.000000000000000 3.000000000000000]
```

```
sage: numerical_approx(M)
[0.000000000000000 0.500000000000000 0.666666666666667]
[0.750000000000000 0.800000000000000 0.833333333333333]
[0.857142857142857 0.875000000000000 0.888888888888889]
[0.900000000000000 0.909090909090909 0.916666666666667]
```

norm($p=2$)

Return the p -norm of this matrix, where p can be 1, 2, inf, or the Frobenius norm.

INPUT:

- `self` - a matrix whose entries are coercible into CDF
- `p` - one of the following options:
 - 1 - the largest column-sum norm
 - 2 (default) - the Euclidean norm
 - Infinity - the largest row-sum norm
 - 'frob' - the Frobenius (sum of squares) norm

OUTPUT: RDF number

See Also:

- `sage.misc.functional.norm()`

EXAMPLES:

```
sage: A = matrix(ZZ, [[1,2,4,3],[-1,0,3,-10]])
sage: A.norm(1)
13.0
sage: A.norm(Infinity)
14.0
sage: B = random_matrix(QQ, 20, 21)
sage: B.norm(Infinity) == (B.transpose()).norm(1)
True

sage: Id = identity_matrix(12)
sage: Id.norm(2)
1.0
sage: A = matrix(RR, 2, 2, [13,-4,-4,7])
sage: A.norm()
15.0
```

Norms of numerical matrices over high-precision reals are computed by this routine. Faster routines for double precision entries from *RDF* or *CDF* are provided by the `Matrix_double_dense` class.

```
sage: A = matrix(CC, 2, 3, [3*I, 4, 1-I, 1, 2, 0])
sage: A.norm('frob')
5.65685424949
sage: A.norm(2)
5.47068444321
sage: A.norm(1)
6.0
sage: A.norm(Infinity)
8.41421356237
sage: a = matrix([[[]], [], [], []])
sage: a.norm()
0.0
sage: a.norm(Infinity) == a.norm(1)
True
```

nullity()

Return the (left) nullity of this matrix, which is the dimension of the (left) kernel of this matrix acting from the right on row vectors.

EXAMPLES:

```
sage: M = Matrix(QQ, [[1, 0, 0, 1], [0, 1, 1, 0], [1, 1, 1, 0]])
sage: M.nullity()
0
sage: M.left_nullity()
0

sage: A = M.transpose()
sage: A.nullity()
1
sage: A.left_nullity()
1

sage: M = M.change_ring(ZZ)
sage: M.nullity()
0
sage: A = M.transpose()
sage: A.nullity()
1
```

numerical_approx (*prec=None, digits=None, algorithm=None*)

Return a numerical approximation of `self` as either a real or complex number with at least the requested number of bits or digits of precision.

INPUT:

- `prec` - an integer: the number of bits of precision
- `digits` - an integer: digits of precision

OUTPUT: A matrix coerced to a real or complex field with `prec` bits of precision.

EXAMPLES:

```
sage: d = matrix([[3, 0], [0, sqrt(2)]]) ;
sage: b = matrix([[1, -1], [2, 2]]) ; e = b * d * b.inverse(); e
[ 1/2*sqrt(2) + 3/2 -1/4*sqrt(2) + 3/4]
[      -sqrt(2) + 3  1/2*sqrt(2) + 3/2]
```



```

sage: e.numerical_approx(53)
[ 2.20710678118655 0.396446609406726]
[ 1.58578643762690 2.20710678118655]

sage: e.numerical_approx(20)
[ 2.2071 0.39645]
[ 1.5858 2.2071]

sage: (e-I).numerical_approx(20)
[2.2071 - 1.0000*I      0.39645]
[      1.5858 2.2071 - 1.0000*I]

sage: M=matrix(QQ,4,[i/(i+1) for i in range(12)]);M
[ 0 1/2 2/3]
[ 3/4 4/5 5/6]
[ 6/7 7/8 8/9]
[ 9/10 10/11 11/12]

sage: M.numerical_approx()
[0.000000000000000 0.500000000000000 0.666666666666667]
[0.750000000000000 0.800000000000000 0.833333333333333]
[0.857142857142857 0.875000000000000 0.888888888888889]
[0.900000000000000 0.909090909090909 0.916666666666667]

sage: matrix(SR, 2, 2, range(4)).n()
[0.000000000000000 1.000000000000000]
[2.000000000000000 3.000000000000000]

sage: numerical_approx(M)
[0.000000000000000 0.500000000000000 0.666666666666667]
[0.750000000000000 0.800000000000000 0.833333333333333]
[0.857142857142857 0.875000000000000 0.888888888888889]
[0.900000000000000 0.909090909090909 0.916666666666667]

```

permanent()

Calculate and return the permanent of the $m \times n$ matrix `self` using Ryser's algorithm.

Let $A = (a_{i,j})$ be an $m \times n$ matrix over any commutative ring, with $m \leq n$. The permanent of A is

$$\text{per}(A) = \sum_{\pi} a_{1,\pi(1)} a_{2,\pi(2)} \cdots a_{m,\pi(m)}$$

where the summation extends over all one-to-one functions π from $\{1, \dots, m\}$ to $\{1, \dots, n\}$.

The product $a_{1,\pi(1)} a_{2,\pi(2)} \cdots a_{m,\pi(m)}$ is called diagonal product. So the permanent of an $m \times n$ matrix A is the sum of all the diagonal products of A .

INPUT:

- A – matrix of size $m \times n$ with $m \leq n$

OUTPUT:

permanent of the matrix A

ALGORITHM:

Modification of theorem 7.1.1. from Brualdi and Ryser: Combinatorial Matrix Theory. Instead of deleting columns from A , we choose columns from A and calculate the product of the row sums of the selected submatrix.

EXAMPLES:

```

sage: M = MatrixSpace(ZZ, 4, 4)
sage: A = M([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
sage: A.permanent()
24

sage: M = MatrixSpace(QQ, 3, 6)
sage: A = M([1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1])
sage: A.permanent()
36

sage: M = MatrixSpace(RR, 3, 6)
sage: A = M([1.0, 1.0, 1.0, 1.0, 0, 0, 0, 1.0, 1.0, 1.0, 1.0, 0, 0, 0, 1.0, 1.0, 1.0, 1.0])
sage: A.permanent()
36.000000000000000

```

See Sloane's sequence OEIS A079908(3) = 36, "The Dancing School Problems"

```

sage: oeis(79908)                                     # optional -- internet
A079908: Solution to the Dancing School Problem with 3 girls and n+3 boys: f(3,n).
sage: _(3)                                           # optional -- internet
36

```

```

sage: M = MatrixSpace(ZZ, 4, 5)
sage: A = M([1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0])
sage: A.permanent()
32

```

See Minc: Permanents, Example 2.1, p. 5.

```

sage: M = MatrixSpace(QQ, 2, 2)
sage: A = M([1/5, 2/7, 3/2, 4/5])
sage: A.permanent()
103/175

sage: R.<a> = PolynomialRing(ZZ)
sage: A = MatrixSpace(R, 2) ([a, 1], [a, a+1])
sage: A.permanent()
a^2 + 2*a

sage: R.<x,y> = PolynomialRing(ZZ, 2)
sage: A = MatrixSpace(R, 2) ([x, y, x^2, y^2])
sage: A.permanent()
x^2*y + x*y^2

```

AUTHORS:

- Jaap Spies (2006-02-16)
- Jaap Spies (2006-02-21): added definition of permanent

permanental_minor(k)

Return the permanental k -minor of an $m \times n$ matrix.

This is the sum of the permanents of all possible k by k submatrices of A .

See Brualdi and Ryser: Combinatorial Matrix Theory, p. 203. Note the typo $p_0(A) = 0$ in that reference! For applications see Theorem 7.2.1 and Theorem 7.2.4.

Note that the permanental m -minor equals $\text{per}(A)$ if $m = n$.

For a $(0,1)$ -matrix A the permanental k -minor counts the number of different selections of k 1's of A with

no two of the 1's on the same row and no two of the 1's on the same column.

INPUT:

•self – matrix of size $m \times n$ with $m \leq n$

OUTPUT:

The permanental k -minor of the matrix self.

EXAMPLES:

```
sage: M = MatrixSpace(ZZ, 4, 4)
sage: A = M([1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 10, 10, 1, 0, 1, 1])
sage: A.permanental_minor(2)
114

sage: M = MatrixSpace(ZZ, 3, 6)
sage: A = M([1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1])
sage: A.permanental_minor(0)
1
sage: A.permanental_minor(1)
12
sage: A.permanental_minor(2)
40
sage: A.permanental_minor(3)
36
```

Note that if $k = m = n$, the permanental k -minor equals $\text{per}(A)$:

```
sage: A.permanent()
36

sage: A.permanental_minor(5)
0
```

For C the “complement” of A :

```
sage: M = MatrixSpace(ZZ, 3, 6)
sage: C = M([0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0])
sage: m, n = 3, 6
sage: sum([(-1)^k * C.permanental_minor(k) * factorial(n-k) / factorial(n-m) for k in range(m+1)])
36
```

See Theorem 7.2.1 of Brualdi and Ryser: Combinatorial Matrix Theory: $\text{per}(A)$

AUTHORS:

•Jaap Spies (2006-02-19)

permutation_normal_form(*check=False*)

Take the set of matrices that are self permuted by any row and column permutation, and return the maximal one of the set where matrices are ordered lexicographically going along each row.

INPUT:

•**check** – (default: **False**) If **True** return a tuple of the maximal matrix and the permutations taking taking self to the maximal matrix. If **False**, return only the maximal matrix.

OUTPUT:

The maximal matrix.

EXAMPLES:

```
sage: M = matrix(ZZ, [[0, 0, 1], [1, 0, 2], [0, 0, 0]])
sage: M
[0 0 1]
[1 0 2]
[0 0 0]
```

```
sage: M.permutation_normal_form()
[2 1 0]
[1 0 0]
[0 0 0]
```

```
sage: M = matrix(ZZ, [[-1, 3], [-1, 5], [2, 4]])
sage: M
[-1  3]
[-1  5]
[ 2  4]
```

```
sage: M.permutation_normal_form(check=True)
(
  [ 5 -1]
  [ 4  2]
  [ 3 -1],
  ((1,2,3), (1,2))
)
```

TESTS:

```
sage: M = matrix(ZZ, [[3, 4, 5], [3, 4, 5], [3, 5, 4], [2, 0, 1]])
sage: M.permutation_normal_form()
[5 4 3]
[5 4 3]
[4 5 3]
[1 0 2]
```

pfaffian (*algorithm=None, check=True*)

Return the Pfaffian of `self`, assuming that `self` is an alternating matrix.

INPUT:

- `algorithm` – string, the algorithm to use; currently the following algorithms have been implemented:
 - ‘definition’ - using the definition given by perfect matchings
- `check` (default: `True`) – Boolean determining whether to check `self` for alternatingness and squareness. This has to be set to `False` if `self` is defined over a non-discrete ring.

The Pfaffian of an alternating matrix is defined as follows:

Let A be an alternating $k \times k$ matrix over a commutative ring. (Here, “alternating” means that $A^T = -A$ and that the diagonal entries of A are zero.) If k is odd, then the Pfaffian of the matrix A is defined to be 0. Let us now define it when k is even. In this case, set $n = k/2$ (this is an integer). For every i and j , we denote the (i, j) -th entry of A by $a_{i,j}$. Let M denote the set of all perfect matchings of the set $\{1, 2, \dots, 2n\}$ (see `sage.combinat.perfect_matching.PerfectMatchings`). For every matching $m \in M$, define the sign $\text{sign}(m)$ of m by writing m as $\{\{i_1, j_1\}, \{i_2, j_2\}, \dots, \{i_n, j_n\}\}$ with $i_k < j_k$ for all k , and setting $\text{sign}(m)$ to be the sign of the permutation $(i_1, j_1, i_2, j_2, \dots, i_n, j_n)$ (written here in one-line notation). For every matching $m \in M$, define the weight $w(m)$ of m by writing m as $\{\{i_1, j_1\}, \{i_2, j_2\}, \dots, \{i_n, j_n\}\}$ with $i_k < j_k$ for all k , and setting $w(m) = a_{i_1, j_1} a_{i_2, j_2} \cdots a_{i_n, j_n}$. Now,

the Pfaffian of the matrix A is defined to be the sum

$$\sum_{m \in M} \text{sign}(m) w(m).$$

The Pfaffian of A is commonly denoted by $\text{Pf}(A)$. It is well-known that $(\text{Pf}(A))^2 = \det A$ for every alternating matrix A , and that $\text{Pf}(U^T A U) = \det U \cdot \text{Pf}(A)$ for any $n \times n$ matrix U and any alternating $n \times n$ matrix A .

See [Kn95], [DW95] and [Rote2001], just to name three sources, for further properties of Pfaffians.

ALGORITHM:

The current implementation uses the definition given above. It checks alternatingness of the matrix `self` only if `check` is `True` (this is important because even if `self` is alternating, a non-discrete base ring might prevent Sage from being able to check this).

REFERENCES:

Todo

Implement faster algorithms, including a division-free one. Does [Rote2001], section 3.3 give one?

Check the implementation of the matchings used here for performance?

EXAMPLES:

A 3×3 alternating matrix has Pfaffian 0 independently of its entries:

```
sage: MSp = MatrixSpace(Integers(27), 3)
sage: A = MSp([0, 2, -3, -2, 0, 8, 3, -8, 0])
sage: A.pfaffian()
0
sage: parent(A.pfaffian())
Ring of integers modulo 27
```

The Pfaffian of a 2×2 alternating matrix is just its northeast entry:

```
sage: MSp = MatrixSpace(QQ, 2)
sage: A = MSp([0, 4, -4, 0])
sage: A.pfaffian()
4
sage: parent(A.pfaffian())
Rational Field
```

The Pfaffian of a 0×0 alternating matrix is 1:

```
sage: MSp = MatrixSpace(ZZ, 0)
sage: A = MSp([])
sage: A.pfaffian()
1
sage: parent(A.pfaffian())
Integer Ring
```

Let us compute the Pfaffian of a generic 4×4 alternating matrix:

```
sage: R = PolynomialRing(QQ, 'x12,x13,x14,x23,x24,x34')
sage: x12, x13, x14, x23, x24, x34 = R.gens()
sage: A = matrix(R, [[ 0, x12, x13, x14],
....:                [-x12, 0, x23, x24],
....:                [-x13, -x23, 0, x34],
....:                [-x14, -x24, -x34, 0]])
```

```
sage: A.pfaffian()
x14*x23 - x13*x24 + x12*x34
sage: parent(A.pfaffian())
Multivariate Polynomial Ring in x12, x13, x14, x23, x24, x34 over Rational Field
```

The Pfaffian of an alternating matrix squares to its determinant:

```
sage: A = [[0] * 6 for i in range(6)]
sage: for i in range(6):
....:     for j in range(i):
....:         u = floor(random() * 10)
....:         A[i][j] = u
....:         A[j][i] = -u
....:         A[i][i] = 0
sage: AA = Matrix(ZZ, A)
sage: AA.pfaffian() ** 2 == AA.det()
True
```

AUTHORS:

- Darij Grinberg (1 Oct 2013): first (slow) implementation.

pivot_rows()

Return the pivot row positions for this matrix, which are a topmost subset of the rows that span the row space and are linearly independent.

OUTPUT: a tuple of integers

EXAMPLES:

```
sage: A = matrix(QQ, 3, 3, [0, 0, 0, 1, 2, 3, 2, 4, 6]); A
[0 0 0]
[1 2 3]
[2 4 6]
sage: A.pivot_rows()
(1,)
sage: A.pivot_rows() # testing cached value
(1,)
```

plot(*args, **kws)

A plot of this matrix.

Each (ith, jth) matrix element is given a different color value depending on its relative size compared to the other elements in the matrix.

The tick marks drawn on the frame axes denote the (ith, jth) element of the matrix.

This method just calls `matrix_plot`. `*args` and `**kws` are passed to `matrix_plot`.

EXAMPLES:

A matrix over $\mathbb{Z}\mathbb{Z}$ colored with different grey levels:

```
sage: A = matrix([[1, 3, 5, 1], [2, 4, 5, 6], [1, 3, 5, 7]])
sage: A.plot()
```

Here we make a random matrix over $\mathbb{R}\mathbb{R}$ and use `cmap='hsv'` to color the matrix elements different RGB colors (see documentation for `matrix_plot` for more information on `cmaps`):

```
sage: A = random_matrix(RDF, 50)
sage: plot(A, cmap='hsv')
```

Another random plot, but over GF(389):

```
sage: A = random_matrix(GF(389), 10)
sage: A.plot(cmap='Oranges')
```

prod_of_row_sums (*cols*)

Calculate the product of all row sums of a submatrix of A for a list of selected columns *cols*.

EXAMPLES:

```
sage: a = matrix(QQ, 2, 2, [1, 2, 3, 2]); a
[1 2]
[3 2]
sage: a.prod_of_row_sums([0, 1])
15
```

Another example:

```
sage: a = matrix(QQ, 2, 3, [1, 2, 3, 2, 5, 6]); a
[1 2 3]
[2 5 6]
sage: a.prod_of_row_sums([1, 2])
55
```

AUTHORS:

•Jaap Spies (2006-02-18)

randomize (*density=1, nonzero=False, *args, **kws*)

Randomize density proportion of the entries of this matrix, leaving the rest unchanged.

Note: We actually choose at random *density* proportion of entries of the matrix and set them to random elements. It's possible that the same position can be chosen multiple times, especially for a very small matrix.

INPUT:

- density* - float (default: 1); rough measure of the proportion of nonzero entries in the random matrix
- nonzero* - Bool (default: False); whether the new entries have to be non-zero
- *args, **kws* - Remaining parameters may be passed to the `random_element` function of the base ring

EXAMPLES:

We construct the zero matrix over a polynomial ring.

```
sage: a = matrix(QQ['x'], 3); a
[0 0 0]
[0 0 0]
[0 0 0]
```

We then randomize roughly half the entries:

```
sage: a.randomize(0.5)
sage: a
[      1/2*x^2 - x - 12 1/2*x^2 - 1/95*x - 1/2      0]
[-5/2*x^2 + 2/3*x - 1/4      0      0]
[      -x^2 + 2/3*x      0      0]
```

Now we randomize all the entries of the resulting matrix:

```
sage: a.randomize()
sage: a
[ 1/3*x^2 - x + 1      -x^2 + 1      x^2 - x]
[ -1/14*x^2 - x - 1/4  -4*x - 1/5  -1/4*x^2 - 1/2*x + 4]
[ 1/9*x^2 + 5/2*x - 3   -x^2 + 3/2*x + 1   -2/7*x^2 - x - 1/2]
```

We create the zero matrix over the integers:

```
sage: a = matrix(ZZ, 2); a
[0 0]
[0 0]
```

Then we randomize it; the x and y parameters, which determine the size of the random elements, are passed onto the `ZZ` `random_element` method.

```
sage: a.randomize(x=-2^64, y=2^64)
sage: a
[-12401200298100116246  1709403521783430739]
[ -4417091203680573707  17094769731745295000]
```

rational_form(*format='right', subdivide=True*)

Returns the rational canonical form, also known as Frobenius form.

INPUT:

- `self` - a square matrix with entries from an exact field.
- `format` - default: 'right' - one of 'right', 'bottom', 'left', 'top' or 'invariants'. The first four will cause a matrix to be returned with companion matrices dictated by the keyword. The value 'invariants' will cause a list of lists to be returned, where each list contains coefficients of a polynomial associated with a companion matrix.
- `subdivide` - default: 'True' - if 'True' and a matrix is returned, then it contains subdivisions delineating the companion matrices along the diagonal.

OUTPUT:

The rational form of a matrix is a similar matrix composed of submatrices ("blocks") placed on the main diagonal. Each block is a companion matrix. Associated with each companion matrix is a polynomial. In rational form, the polynomial of one block will divide the polynomial of the next block (and thus, the polynomials of all subsequent blocks).

Rational form, also known as Frobenius form, is a canonical form. In other words, two matrices are similar if and only if their rational canonical forms are equal. The algorithm used does not provide the similarity transformation matrix (also known as the change-of-basis matrix).

Companion matrices may be written in one of four styles, and any such style may be selected with the `format` keyword. See the companion matrix constructor, `sage.matrix.constructor.companion_matrix()`, for more information about companion matrices.

If the 'invariants' value is used for the `format` keyword, then the return value is a list of lists, where each list is the coefficients of the polynomial associated with one of the companion matrices on the diagonal. These coefficients include the leading one of the monic polynomial and are ready to be coerced into any polynomial ring over the same field (see examples of this below). This return value is intended to be the most compact representation and the easiest to use for testing equality of rational forms.

Because the minimal and characteristic polynomials of a companion matrix are the associated polynomial, it is easy to see that the product of the polynomials of the blocks will be the characteristic polynomial and the final polynomial will be the minimal polynomial of the entire matrix.

ALGORITHM:

We begin with ZigZag form, which is due to Arne Storjohann and is documented at `zigzag_form()`. Then we eliminate “corner” entries enroute to rational form via an additional algorithm of Storjohann’s [STORJOHANN-EMAIL].

EXAMPLES:

The lists of coefficients returned with the `invariants` keyword are designed to easily convert to the polynomials associated with the companion matrices. This is illustrated by the construction below of the `polys` list. Then we can test the divisibility condition on the list of polynomials. Also the minimal and characteristic polynomials are easy to determine from this list.

```
sage: A = matrix(QQ, [[ 11, 14, -15, -4, -38, -29, 1, 23, 14, -63, 17, 24, 36, 32],
...                  [ 18, 6, -17, -11, -31, -43, 12, 26, 0, -69, 11, 13, 17, 24],
...                  [ 11, 16, -22, -8, -48, -34, 0, 31, 16, -82, 26, 31, 39, 37],
...                  [-8, -18, 22, 10, 46, 33, 3, -27, -12, 70, -19, -20, -42, -31],
...                  [-13, -21, 16, 10, 52, 43, 4, -28, -25, 89, -37, -20, -53, -62],
...                  [-2, -6, 0, 0, 6, 10, 1, 1, -7, 14, -11, -3, -10, -18],
...                  [-9, -19, -3, 4, 23, 30, 8, -3, -27, 55, -40, -5, -40, -69],
...                  [ 4, -8, -1, -1, 5, -4, 9, 5, -11, 4, -14, -2, -13, -17],
...                  [ 1, -2, 16, -1, 19, -2, -1, -17, 2, 19, 5, -25, -7, 14],
...                  [ 7, 7, -13, -4, -26, -21, 3, 18, 5, -40, 7, 15, 20, 14],
...                  [-6, -7, -12, 4, -1, 18, 3, 8, -11, 15, -18, 17, -15, -41],
...                  [ 5, 11, -11, -3, -26, -19, -1, 14, 10, -42, 14, 17, 25, 23],
...                  [-16, -15, 3, 10, 29, 45, -1, -13, -19, 71, -35, -2, -35, -65],
...                  [ 4, 2, 3, -2, -2, -10, 1, 0, 3, -11, 6, -4, 6, 17]])
sage: A.rational_form()
[ 0  -4| 0  0  0  0| 0  0  0  0  0  0  0  0  0]
[ 1   4| 0  0  0  0| 0  0  0  0  0  0  0  0  0]
[-----+-----+-----]
[ 0  0| 0  0  0 12| 0  0  0  0  0  0  0  0  0]
[ 0  0| 1  0  0 -4| 0  0  0  0  0  0  0  0  0]
[ 0  0| 0  1  0 -9| 0  0  0  0  0  0  0  0  0]
[ 0  0| 0  0  1  6| 0  0  0  0  0  0  0  0  0]
[-----+-----+-----]
[ 0  0| 0  0  0  0| 0  0  0  0  0  0  0 -216]
[ 0  0| 0  0  0  0| 1  0  0  0  0  0  0 108]
[ 0  0| 0  0  0  0| 0  1  0  0  0  0  0 306]
[ 0  0| 0  0  0  0| 0  0  1  0  0  0  0 -271]
[ 0  0| 0  0  0  0| 0  0  0  1  0  0  0 -41]
[ 0  0| 0  0  0  0| 0  0  0  0  1  0  0 134]
[ 0  0| 0  0  0  0| 0  0  0  0  0  1  0 -64]
[ 0  0| 0  0  0  0| 0  0  0  0  0  0  1 13]

sage: R = PolynomialRing(QQ, 'x')
sage: invariants = A.rational_form(format='invariants')
sage: invariants
[[4, -4, 1], [-12, 4, 9, -6, 1], [216, -108, -306, 271, 41, -134, 64, -13, 1]]
sage: polys = [R(p) for p in invariants]
sage: [p.factor() for p in polys]
[(x - 2)^2, (x - 3) * (x + 1) * (x - 2)^2, (x + 1)^2 * (x - 3)^3 * (x - 2)^3]
sage: all(polys[i].divides(polys[i+1]) for i in range(len(polys)-1))
True
sage: polys[-1] == A.minimal_polynomial(var='x')
True
sage: prod(polys) == A.characteristic_polynomial(var='x')
True
```

Rational form is a canonical form. Any two matrices are similar if and only if their rational forms are

equal. By starting with Jordan canonical forms, the matrices C and D below were built as similar matrices, while E was built to be just slightly different. All three matrices have equal characteristic polynomials though E's minimal polynomial differs.

```
sage: C = matrix(QQ, [[2, 31, -10, -9, -125, 13, 62, -12],
...                  [0, 48, -16, -16, -188, 20, 92, -16],
...                  [0, 9, -1, 2, -33, 5, 18, 0],
...                  [0, 15, -5, 0, -59, 7, 30, -4],
...                  [0, -21, 7, 2, 84, -10, -42, 5],
...                  [0, -42, 14, 8, 167, -17, -84, 13],
...                  [0, -50, 17, 10, 199, -23, -98, 14],
...                  [0, 15, -5, -2, -59, 7, 30, -2]])
sage: C.minimal_polynomial().factor()
(x - 2)^2
sage: C.characteristic_polynomial().factor()
(x - 2)^8
sage: C.rational_form()
[ 0 -4| 0 0| 0 0| 0 0]
[ 1 4| 0 0| 0 0| 0 0]
[-----+-----+-----+-----]
[ 0 0| 0 -4| 0 0| 0 0]
[ 0 0| 1 4| 0 0| 0 0]
[-----+-----+-----+-----]
[ 0 0| 0 0| 0 -4| 0 0]
[ 0 0| 0 0| 1 4| 0 0]
[-----+-----+-----+-----]
[ 0 0| 0 0| 0 0| 0 -4]
[ 0 0| 0 0| 0 0| 1 4]

sage: D = matrix(QQ, [[-4, 3, 7, 2, -4, 5, 7, -3],
...                  [-6, 5, 7, 2, -4, 5, 7, -3],
...                  [21, -12, 89, 25, 8, 27, 98, -95],
...                  [-9, 5, -44, -11, -3, -13, -48, 47],
...                  [23, -13, 74, 21, 12, 22, 85, -84],
...                  [31, -18, 135, 38, 12, 47, 155, -147],
...                  [-33, 19, -138, -39, -13, -45, -156, 151],
...                  [-7, 4, -29, -8, -3, -10, -34, 34]])
sage: D.minimal_polynomial().factor()
(x - 2)^2
sage: D.characteristic_polynomial().factor()
(x - 2)^8
sage: D.rational_form()
[ 0 -4| 0 0| 0 0| 0 0]
[ 1 4| 0 0| 0 0| 0 0]
[-----+-----+-----+-----]
[ 0 0| 0 -4| 0 0| 0 0]
[ 0 0| 1 4| 0 0| 0 0]
[-----+-----+-----+-----]
[ 0 0| 0 0| 0 -4| 0 0]
[ 0 0| 0 0| 1 4| 0 0]
[-----+-----+-----+-----]
[ 0 0| 0 0| 0 0| 0 -4]
[ 0 0| 0 0| 0 0| 1 4]

sage: E = matrix(QQ, [[0, -8, 4, -6, -2, 5, -3, 11],
...                  [-2, -4, 2, -4, -2, 4, -2, 6],
...                  [5, 14, -7, 12, 3, -8, 6, -27],
...                  [-3, -8, 7, -5, 0, 2, -6, 17],
...                  [0, 5, 0, 2, 4, -4, 1, 2],
```

```

...          [-3, -7, 5, -6, -1, 5, -4, 14],
...          [ 6, 18, -10, 14, 4, -10, 10, -28],
...          [-2, -6, 4, -5, -1, 3, -3, 13]])
sage: E.minimal_polynomial().factor()
(x - 2)^3
sage: E.characteristic_polynomial().factor()
(x - 2)^8
sage: E.rational_form()
[ 2| 0 0| 0 0| 0 0 0]
[---+-----+-----+-----]
[ 0| 0 -4| 0 0| 0 0 0]
[ 0| 1 4| 0 0| 0 0 0]
[---+-----+-----+-----]
[ 0| 0 0| 0 -4| 0 0 0]
[ 0| 0 0| 1 4| 0 0 0]
[---+-----+-----+-----]
[ 0| 0 0| 0 0| 0 0 8]
[ 0| 0 0| 0 0| 1 0 -12]
[ 0| 0 0| 0 0| 0 1 6]

```

The principal feature of rational canonical form is that it can be computed over any field using only field operations. Other forms, such as Jordan canonical form, are complicated by the need to determine the eigenvalues of the matrix, which can lie outside the field. The following matrix has all of its eigenvalues outside the rationals - some are irrational ($\pm\sqrt{2}$) and the rest are complex ($-1 \pm 2i$).

```

sage: A = matrix(QQ,
...   [[-154, -3, -54, 44, 48, -244, -19, 67, -326, 85, 355, 581],
...   [ 504, 25, 156, -145, -171, 793, 99, -213, 1036, -247, -1152, -1865],
...   [ 294, -1, 112, -89, -90, 469, 36, -128, 634, -160, -695, -1126],
...   [-49, -32, 25, 7, 37, -64, -58, 12, -42, -14, 72, 106],
...   [-261, -123, 65, 47, 169, -358, -254, 70, -309, -29, 454, 673],
...   [-448, -123, -10, 109, 227, -668, -262, 163, -721, 95, 896, 1410],
...   [ 38, 7, 8, -14, -17, 66, 6, -23, 73, -29, -78, -143],
...   [-96, 10, -55, 37, 24, -168, 17, 56, -231, 88, 237, 412],
...   [ 310, 67, 31, -81, -143, 473, 143, -122, 538, -98, -641, -1029],
...   [ 139, -35, 99, -49, -18, 236, -41, -70, 370, -118, -377, -619],
...   [ 243, 9, 81, -72, -81, 386, 43, -105, 508, -124, -564, -911],
...   [-155, -3, -55, 45, 50, -245, -27, 65, -328, 77, 365, 583]])
sage: A.characteristic_polynomial().factor()
(x^2 - 2)^2 * (x^2 + 2*x + 5)^4
sage: A.eigenvalues(extend=False)
[]
sage: A.rational_form()
[ 0 -5| 0 0 0 0| 0 0 0 0 0 0]
[ 1 -2| 0 0 0 0| 0 0 0 0 0 0]
[-----+-----+-----]
[ 0 0| 0 0 0 10| 0 0 0 0 0 0]
[ 0 0| 1 0 0 4| 0 0 0 0 0 0]
[ 0 0| 0 1 0 -3| 0 0 0 0 0 0]
[ 0 0| 0 0 1 -2| 0 0 0 0 0 0]
[-----+-----+-----]
[ 0 0| 0 0 0 0| 0 0 0 0 0 50]
[ 0 0| 0 0 0 0| 1 0 0 0 0 40]
[ 0 0| 0 0 0 0| 0 1 0 0 0 3]
[ 0 0| 0 0 0 0| 0 0 1 0 0 -12]
[ 0 0| 0 0 0 0| 0 0 0 1 0 -12]
[ 0 0| 0 0 0 0| 0 0 0 0 1 -4]
sage: F.<x> = QQ[]

```

```

sage: polys = A.rational_form(format='invariants')
sage: [F(p).factor() for p in polys]
[x^2 + 2*x + 5, (x^2 - 2) * (x^2 + 2*x + 5), (x^2 - 2) * (x^2 + 2*x + 5)^2]

```

Rational form may be computed over any field. The matrix below is an example where the eigenvalues lie outside the field.

```

sage: F.<a> = FiniteField(7^2)
sage: A = matrix(F,
...   [[5*a + 3, 4*a + 1, 6*a + 2, 2*a + 5, 6, 4*a + 5, 4*a + 5, 5, a + 6,
...   [6*a + 3, 2*a + 4, 0, 6, 5*a + 5, 2*a, 5*a + 1, 1, 5*a + 2,
...   [3*a + 1, 6*a + 6, a + 6, 2, 0, 3*a + 6, 5*a + 4, 5*a + 6, 5*a + 2,
...   [3*a, 6*a, 3*a, 4*a, 4*a + 4, 3*a + 6, 6*a, 4, 3*a + 4, 6*a
...   [4*a + 5, a + 1, 4*a + 3, 6*a + 5, 5*a + 2, 5*a + 2, 6*a, 4*a + 6, 6*a + 4, 5*a
...   [3*a, 6*a, 4*a + 1, 6*a + 2, 2*a + 5, 4*a + 6, 2, a + 5, 2*a + 4, 2*a
...   [4*a + 5, 3*a + 3, 6, 4*a + 1, 4*a + 3, 6*a + 3, 6, 3*a + 3, 3, a
...   [6*a + 6, a + 4, 2*a + 6, 3*a + 5, 4*a + 3, 2, a, 3*a + 4, 5*a, 2*a
...   [3*a + 5, 6*a + 2, 4*a, a + 5, 0, 5*a, 6*a + 5, 2*a + 1, 3*a + 1, 3*a
...   [3*a + 2, a + 3, 3*a + 6, a, 3*a + 5, 5*a + 1, 3*a + 2, a + 3, a + 2, 6*a
...   [6*a + 6, 5*a + 1, 4*a, 2, 5*a + 5, 3*a + 5, 3*a + 1, 2*a, 2*a, 2*a
sage: A.rational_form()
[ a + 2| 0 0 0| 0 0 0 0 0 0 0]
[-----+-----]
[ 0| 0 0 a + 6| 0 0 0 0 0 0 0]
[ 0| 1 0 6*a + 4| 0 0 0 0 0 0 0]
[ 0| 0 1 6*a + 4| 0 0 0 0 0 0 0]
[-----+-----]
[ 0| 0 0 0| 0 0 0 0 0 0 2*a]
[ 0| 0 0 0| 1 0 0 0 0 0 6*a + 3]
[ 0| 0 0 0| 0 1 0 0 0 0 6*a + 1]
[ 0| 0 0 0| 0 0 1 0 0 0 a + 2]
[ 0| 0 0 0| 0 0 0 1 0 0 a + 6]
[ 0| 0 0 0| 0 0 0 0 1 0 2*a + 1]
[ 0| 0 0 0| 0 0 0 0 0 1 2*a + 1]
sage: invariants = A.rational_form(format='invariants')
sage: invariants
[[6*a + 5, 1], [6*a + 1, a + 3, a + 3, 1], [5*a, a + 4, a + 6, 6*a + 5, 6*a + 1, 5*a + 6, 5*a
sage: R.<x> = F[]
sage: polys = [R(p) for p in invariants]
sage: [p.factor() for p in polys]
[x + 6*a + 5, (x + 6*a + 5) * (x^2 + (2*a + 5)*x + 5*a), (x + 6*a + 5) * (x^2 + (2*a + 5)*x
sage: polys[-1] == A.minimal_polynomial()
True
sage: prod(polys) == A.characteristic_polynomial()
True
sage: A.eigenvalues()
Traceback (most recent call last):
...
NotImplementedError: algebraic closures of finite fields are only implemented for prime field

```

Companion matrices may be selected as any one of four different types. See the documentation for the companion matrix constructor, `sage.matrix.constructor.companion_matrix()`, for more information.

```

sage: A = matrix(QQ, [[35, -18, -2, -45],
...   [22, -22, 12, -16],
...   [5, -12, 12, 4],
...   [16, -6, -4, -23]])
sage: A.rational_form(format='right')

```

```

[ 2| 0  0  0]
[--+-----]
[ 0| 0  0 10]
[ 0| 1  0 -1]
[ 0| 0  1  0]
sage: A.rational_form(format='bottom')
[ 2| 0  0  0]
[--+-----]
[ 0| 0  1  0]
[ 0| 0  0  1]
[ 0|10 -1  0]
sage: A.rational_form(format='left')
[ 2| 0  0  0]
[--+-----]
[ 0| 0  1  0]
[ 0|-1  0  1]
[ 0|10  0  0]
sage: A.rational_form(format='top')
[ 2| 0  0  0]
[--+-----]
[ 0| 0 -1 10]
[ 0| 1  0  0]
[ 0| 0  1  0]

```

TESTS:

```

sage: A = matrix(QQ, 2, 3, range(6))
sage: A.rational_form()
Traceback (most recent call last):
...
TypeError: matrix must be square, not 2 x 3

sage: A = matrix(Integers(6), 2, 2, range(4))
sage: A.rational_form()
Traceback (most recent call last):
...
TypeError: matrix entries must come from an exact field, not Ring of integers modulo 6

sage: A = matrix(RDF, 2, 2, range(4))
sage: A.rational_form()
Traceback (most recent call last):
...
TypeError: matrix entries must come from an exact field, not Real Double Field

sage: A = matrix(QQ, 2, range(4))
sage: A.rational_form(format='junk')
Traceback (most recent call last):
...
ValueError: 'format' keyword must be 'right', 'bottom', 'left', 'top' or 'invariants', not junk

sage: A = matrix(QQ, 2, range(4))
sage: A.rational_form(subdivide='garbage')
Traceback (most recent call last):
...
ValueError: 'subdivide' keyword must be True or False, not garbage

```

Citations

AUTHOR:

- Rob Beezer (2011-06-09)

restrict (*V*, *check=True*)

Returns the matrix that defines the action of self on the chosen basis for the invariant subspace *V*. If *V* is an ambient, returns self (not a copy of self).

INPUT:

- V* - vector subspace
- check* - (optional) default: True; if False may not check that *V* is invariant (hence can be faster).

OUTPUT: a matrix

Warning: This function returns an $n \times n$ matrix, where *V* has dimension *n*. It does *not* check that *V* is in fact invariant under self, unless *check* is True.

EXAMPLES:

```
sage: V = VectorSpace(QQ, 3)
sage: M = MatrixSpace(QQ, 3)
sage: A = M([1,2,0, 3,4,0, 0,0,0])
sage: W = V.subspace([[1,0,0], [0,1,0]])
sage: A.restrict(W)
[1 2]
[3 4]
sage: A.restrict(W, check=True)
[1 2]
[3 4]
```

We illustrate the warning about invariance not being checked by default, by giving a non-invariant subspace. With the default *check=False* this function returns the ‘restriction’ matrix, which is meaningless as *check=True* reveals.

```
sage: W2 = V.subspace([[1,0,0], [0,1,1]])
sage: A.restrict(W2, check=False)
[1 2]
[3 4]
sage: A.restrict(W2, check=True)
Traceback (most recent call last):
...
ArithmeticError: subspace is not invariant under matrix
```

restrict_codomain (*V*)

Suppose that self defines a linear map from some domain to a codomain that contains *V* and that the image of self is contained in *V*. This function returns a new matrix *A* that represents this linear map but as a map to *V*, in the sense that if *x* is in the domain, then *xA* is the linear combination of the elements of the basis of *V* that equals *v**self.

INPUT:

- V* - vector space (space of degree self.ncols()) that contains the image of self.

See Also:

```
restrict(), restrict_domain()
```

EXAMPLES:

```

sage: A = matrix(QQ, 3, [1..9])
sage: V = (QQ^3).span([[1,2,3], [7,8,9]]); V
Vector space of degree 3 and dimension 2 over Rational Field
Basis matrix:
[ 1  0 -1]
[ 0  1  2]
sage: z = vector(QQ, [1,2,5])
sage: B = A.restrict_codomain(V); B
[1 2]
[4 5]
[7 8]
sage: z*B
(44, 52)
sage: z*A
(44, 52, 60)
sage: 44*V.0 + 52*V.1
(44, 52, 60)

```

restrict_domain(V)

Compute the matrix relative to the basis for V on the domain obtained by restricting self to V, but not changing the codomain of the matrix. This is the matrix whose rows are the images of the basis for V.

INPUT:

- V - vector space (subspace of ambient space on which self acts)

See Also:

`restrict()`

EXAMPLES:

```

sage: V = QQ^3
sage: A = matrix(QQ, 3, [1,2,0, 3,4,0, 0,0,0])
sage: W = V.subspace([[1,0,0], [1,2,3]])
sage: A.restrict_domain(W)
[1 2 0]
[3 4 0]
sage: W2 = V.subspace_with_basis([[1,0,0], [1,2,3]])
sage: A.restrict_domain(W2)
[ 1  2  0]
[ 7 10  0]

```

right_eigenmatrix()

Return matrices D and P, where D is a diagonal matrix of eigenvalues and P is the corresponding matrix where the columns are corresponding eigenvectors (or zero vectors) so that $\text{self} \cdot P = P \cdot D$.

EXAMPLES:

```

sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: D, P = A.eigenmatrix_right()
sage: D
[
0 0 0
0 -1.348469228349535? 0
0 0 13.34846922834954?]
sage: P
[
1 1 1
-2 0.1303061543300932? 3.069693845669907?]

```

```
[
1 -0.7393876913398137? 5.139387691339814?]
sage: A*P == P*D
True
```

Because P is invertible, A is diagonalizable.

```
sage: A == P*D*(~P)
True
```

The matrix P may contain zero columns corresponding to eigenvalues for which the algebraic multiplicity is greater than the geometric multiplicity. In these cases, the matrix is not diagonalizable.

```
sage: A = jordan_block(2, 3); A
[2 1 0]
[0 2 1]
[0 0 2]
sage: A = jordan_block(2, 3)
sage: D, P = A.eigenmatrix_right()
sage: D
[2 0 0]
[0 2 0]
[0 0 2]
sage: P
[1 0 0]
[0 0 0]
[0 0 0]
sage: A*P == P*D
True
```

TESTS:

For matrices with floating point entries, some platforms will return eigenvectors that are negatives of those returned by the majority of platforms. This test accounts for that possibility. Running this test independently, without adjusting the eigenvectors could indicate this situation on your hardware.

```
sage: B = matrix(QQ, 3, 3, range(9))
sage: em = B.change_ring(RDF).eigenmatrix_right()
sage: evals = em[0]; evals.dense_matrix().zero_at(2e-15)
[ 13.3484692283 0.0 0.0]
[ 0.0 -1.34846922835 0.0]
[ 0.0 0.0 0.0]
sage: evectors = em[1];
sage: for i in range(3):
...     scale = evectors[0,i].sign()
...     evectors.rescale_col(i, scale)
sage: evectors
[ 0.164763817... 0.799699663... 0.408248290...]
[ 0.505774475... 0.104205787... -0.816496580...]
[ 0.846785134... -0.591288087... 0.408248290...]
```

right_eigenspaces (*format='all', var='a', algebraic_multiplicity=False*)

Compute the right eigenspaces of a matrix.

Note that `eigenspaces_right()` and `right_eigenspaces()` are identical methods. Here “right” refers to the eigenvectors being placed to the right of the matrix.

INPUT:

- `self` - a square matrix over an exact field. For inexact matrices consult the numerical or symbolic matrix classes.

- `format` - default: `None`
 - `'all'` - attempts to create every eigenspace. This will always be possible for matrices with rational entries.
 - `'galois'` - for each irreducible factor of the characteristic polynomial, a single eigenspace will be output for a single root/eigenvalue for the irreducible factor.
 - `None` - Uses the `'all'` format if the base ring is contained in an algebraically closed field which is implemented. Otherwise, uses the `'galois'` format.
- `var` - default: `'a'` - variable name used to represent elements of the root field of each irreducible factor of the characteristic polynomial. If `var='a'`, then the root fields will be in terms of `a0`, `a1`, `a2`, ..., where the numbering runs across all the irreducible factors of the characteristic polynomial, even for linear factors.
- `algebraic_multiplicity` - default: `False` - whether or not to include the algebraic multiplicity of each eigenvalue in the output. See the discussion below.

OUTPUT:

If `algebraic_multiplicity=False`, return a list of pairs (e, V) where e is an eigenvalue of the matrix, and V is the corresponding left eigenspace. For Galois conjugates of eigenvalues, there may be just one representative eigenspace, depending on the `format` keyword.

If `algebraic_multiplicity=True`, return a list of triples (e, V, n) where e and V are as above and n is the algebraic multiplicity of the eigenvalue.

Warning: Uses a somewhat naive algorithm (simply factors the characteristic polynomial and computes kernels directly over the extension field).

EXAMPLES:

Right eigenspaces are computed from the left eigenspaces of the transpose of the matrix. As such, there is a greater collection of illustrative examples at the `eigenspaces_left()`.

We compute the right eigenspaces of a 3×3 rational matrix.

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.eigenspaces_right()
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]),
(-1.348469228349535?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 0.1303061543300932? -0.7393876913398137?]),
(13.34846922834954?, Vector space of degree 3 and dimension 1 over Algebraic Field
User basis matrix:
[
1 3.069693845669907? 5.139387691339814?])
]
sage: es = A.eigenspaces_right(format='galois'); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
```

```

[          1 1/5*a1 + 2/5 2/5*a1 - 1/5])
]
sage: es = A.eigenspaces_right(format='galois', algebraic_multiplicity=True); es
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1], 1),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[          1 1/5*a1 + 2/5 2/5*a1 - 1/5], 1)
]
sage: e, v, n = es[0]; v = v.basis()[0]
sage: delta = v*e - A*v
sage: abs(abs(delta)) < 1e-10
True

```

The same computation, but with implicit base change to a field:

```

sage: A = matrix(ZZ, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.eigenspaces_right(format='galois')
[
(0, Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]),
(a1, Vector space of degree 3 and dimension 1 over Number Field in a1 with defining polynomial
User basis matrix:
[          1 1/5*a1 + 2/5 2/5*a1 - 1/5])
]

```

This method is only applicable to exact matrices. The “eigenmatrix” routines for matrices with double-precision floating-point entries (RDF, CDF) are the best alternative. (Since some platforms return eigenvectors that are the negatives of those given here, this one example is not tested here.) There are also “eigenmatrix” routines for matrices with symbolic entries.

```

sage: B = matrix(RR, 3, 3, range(9))
sage: B.eigenspaces_right()
Traceback (most recent call last):
...
NotImplementedError: eigenspaces cannot be computed reliably for inexact rings such as Real
consult numerical or symbolic matrix classes for other options

```

```

sage: em = B.change_ring(RDF).eigenmatrix_right()
sage: eigenvalues = em[0]; eigenvalues.dense_matrix().zero_at(1e-15)
[ 13.3484692283      0.0      0.0]
[      0.0 -1.34846922835      0.0]
[      0.0      0.0      0.0]
sage: eigenvectors = em[1]; eigenvectors # not tested
[ 0.164763817...  0.799699663...  0.408248290...]
[ 0.505774475...  0.104205787... -0.816496580...]
[ 0.846785134... -0.591288087...  0.408248290...]

sage: x, y = var('x y')
sage: S = matrix([[x, y], [y, 3*x^2]])
sage: em = S.eigenmatrix_right()
sage: eigenvalues = em[0]; eigenvalues
[3/2*x^2 + 1/2*x - 1/2*sqrt(9*x^4 - 6*x^3 + x^2 + 4*y^2)

```

```
[
0 3/2*x^2 + 1/2*x + 1/2*sqrt(9*x^4 -
sage: eigenvectors = em[1]; eigenvectors
[
1
[1/2*(3*x^2 - x - sqrt(9*x^4 - 6*x^3 + x^2 + 4*y^2))/y 1/2*(3*x^2 - x + sqrt(9*x^4 - 6*x^3 +
```

TESTS:

```
sage: B = matrix(QQ, 2, 3, range(6))
sage: B.eigenspaces_right()
Traceback (most recent call last):
...
TypeError: matrix must be square, not 2 x 3

sage: B = matrix(QQ, 4, 4, range(16))
sage: B.eigenspaces_right(format='junk')
Traceback (most recent call last):
...
ValueError: format keyword must be None, 'all' or 'galois', not junk

sage: B.eigenspaces_right(algebraic_multiplicity='garbage')
Traceback (most recent call last):
...
ValueError: algebraic_multiplicity keyword must be True or False
```

right_eigenvectors (*extend=True*)

Compute the right eigenvectors of a matrix.

For each distinct eigenvalue, returns a list of the form (e,V,n) where e is the eigenvalue, V is a list of eigenvectors forming a basis for the corresponding right eigenspace, and n is the algebraic multiplicity of the eigenvalue. If *extend* = True (the default), this will return eigenspaces over the algebraic closure of the base field where this is implemented; otherwise it will restrict to eigenvalues in the base field.

EXAMPLES: We compute the right eigenvectors of a 3×3 rational matrix.

```
sage: A = matrix(QQ, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: es = A.eigenvectors_right(); es
[(0, [
(1, -2, 1)
], 1),
(-1.348469228349535?, [(1, 0.1303061543300932?, -0.7393876913398137?)], 1),
(13.34846922834954?, [(1, 3.069693845669907?, 5.139387691339814?)], 1)]
sage: A.eigenvectors_right(extend=False)
[(0, [
(1, -2, 1)
], 1)]
sage: eval, [evec], mult = es[0]
sage: delta = eval*evec - A*evec
sage: abs(abs(delta)) < 1e-10
True
```

right_kernel (**args, **kws*)

Returns the right kernel of this matrix, as a vector space or free module. This is the set of vectors x such that $\text{self} \cdot x = 0$.

Note: For the left kernel, use `left_kernel()`. The method `kernel()` is exactly equal to `left_kernel()`.

INPUT:

- `algorithm` - default: 'default' - a keyword that selects the algorithm employed. Allowable values are:
 - 'default' - allows the algorithm to be chosen automatically
 - 'generic' - naive algorithm usable for matrices over any field
 - 'pari' - PARI library code for matrices over number fields or the integers
 - 'padic' - padic algorithm from IML library for matrices over the rationals and integers
 - 'pluq' - PLUQ matrix factorization for matrices mod 2
- `basis` - default: 'echelon' - a keyword that describes the format of the basis used to construct the left kernel. Allowable values are:
 - 'echelon': the basis matrix is in echelon form
 - 'pivot': each basis vector is computed from the reduced row-echelon form of `self` by placing a single one in a non-pivot column and zeros in the remaining non-pivot columns. Only available for matrices over fields.
 - 'LLL': an LLL-reduced basis. Only available for matrices over the integers.

OUTPUT:

A vector space or free module whose degree equals the number of columns in `self` and contains all the vectors `x` such that `self*x = 0`.

If `self` has 0 columns, the kernel has dimension 0, while if `self` has 0 rows the kernel is the entire ambient vector space.

The result is cached. Requesting the right kernel a second time, but with a different basis format, will return the cached result with the format from the first computation.

Note: For more detailed documentation on the selection of algorithms used and a more flexible method for computing a basis matrix for a right kernel (rather than computing a vector space), see `right_kernel_matrix()`, which powers the computations for this method.

EXAMPLES:

```
sage: A = matrix(QQ, [[0, 0, 1, 2, 2, -5, 3],
...                  [-1, 5, 2, 2, 1, -7, 5],
...                  [0, 0, -2, -3, -3, 8, -5],
...                  [-1, 5, 0, -1, -2, 1, 0]])
sage: K = A.right_kernel(); K
Vector space of degree 7 and dimension 4 over Rational Field
Basis matrix:
[ 1  0  0  0 -1 -1 -1]
[ 0  1  0  0  5  5  5]
[ 0  0  1  0 -1 -2 -3]
[ 0  0  0  1  0  1  1]
sage: A*K.basis_matrix().transpose() == zero_matrix(QQ, 4, 4)
True
```

The default is basis vectors that form a matrix in echelon form. A “pivot basis” instead has a basis matrix where the columns of an identity matrix are in the locations of the non-pivot columns of the original matrix. This alternate format is available whenever the base ring is a field.

```
sage: A = matrix(QQ, [[0, 0, 1, 2, 2, -5, 3],
...                  [-1, 5, 2, 2, 1, -7, 5],
...                  [0, 0, -2, -3, -3, 8, -5],
...                  [-1, 5, 0, -1, -2, 1, 0]])
sage: A.rref()
[ 1 -5  0  0  1  1 -1]
[ 0  0  1  0  0 -1  1]
[ 0  0  0  1  1 -2  1]
[ 0  0  0  0  0  0  0]
sage: A.nonpivots()
(1, 4, 5, 6)
sage: K = A.right_kernel(basis='pivot'); K
Vector space of degree 7 and dimension 4 over Rational Field
User basis matrix:
[ 5  1  0  0  0  0  0]
[-1  0  0 -1  1  0  0]
[-1  0  1  2  0  1  0]
[ 1  0 -1 -1  0  0  1]
sage: A*K.basis_matrix().transpose() == zero_matrix(QQ, 4, 4)
True
```

Matrices may have any field as a base ring. Number fields are computed by PARI library code, matrices over $GF(2)$ are computed by the M4RI library, and matrices over the rationals are computed by the IML library. For any of these specialized cases, general-purpose code can be called instead with the keyword setting `algorithm='generic'`.

Over an arbitrary field, with two basis formats. Same vector space, different bases.

```
sage: F.<a> = FiniteField(5^2)
sage: A = matrix(F, 3, 4, [[ 1,  a,  1+a,  a^3+a^5],
...                       [ a, a^4,  a+a^4, a^4+a^8],
...                       [a^2, a^6, a^2+a^6, a^5+a^10]])
sage: K = A.right_kernel(); K
Vector space of degree 4 and dimension 2 over Finite Field in a of size 5^2
Basis matrix:
[      1      0 3*a + 4 2*a + 2]
[      0      1 2*a 3*a + 3]
sage: A*K.basis_matrix().transpose() == zero_matrix(F, 3, 2)
True
sage: B = copy(A)
sage: P = B.right_kernel(basis = 'pivot'); P
Vector space of degree 4 and dimension 2 over Finite Field in a of size 5^2
User basis matrix:
[      4      4      1      0]
[ a + 2 3*a + 3      0      1]
sage: B*P.basis_matrix().transpose() == zero_matrix(F, 3, 2)
True
sage: K == P
True
```

Over number fields, PARI is used by default, but general-purpose code can be requested. Same vector space, same bases, different code.:

```
sage: Q = QuadraticField(-7)
sage: a = Q.gen(0)
sage: A = matrix(Q, [[ 2, 5-a, 15-a, 16+4*a],
```

```

...          [2+a,  a, -7 + 5*a, -3+3*a]])
sage: K = A.right_kernel(algorithm='default'); K
Vector space of degree 4 and dimension 2 over Number Field in a with defining polynomial x^2
Basis matrix:
[      1      0      7/88*a + 3/88 -3/176*a - 39/176]
[      0      1 -1/88*a - 13/88 13/176*a - 7/176]
sage: A*K.basis_matrix().transpose() == zero_matrix(Q, 2, 2)
True
sage: B = copy(A)
sage: G = A.right_kernel(algorithm='generic'); G
Vector space of degree 4 and dimension 2 over Number Field in a with defining polynomial x^2
Basis matrix:
[      1      0      7/88*a + 3/88 -3/176*a - 39/176]
[      0      1 -1/88*a - 13/88 13/176*a - 7/176]
sage: B*G.basis_matrix().transpose() == zero_matrix(Q, 2, 2)
True
sage: K == G
True

```

For matrices over the integers, several options are possible. The basis can be an LLL-reduced basis or an echelon basis. The pivot basis is not available. A heuristic will decide whether to use a p-adic algorithm from the IML library or an algorithm from the PARI library. Note how specifying the algorithm can mildly influence the LLL basis.

```

sage: A = matrix(ZZ, [[0, -1, -1, 2, 9, 4, -4],
...                  [-1, 1, 0, -2, -7, -1, 6],
...                  [2, 0, 1, 0, 1, -5, -2],
...                  [-1, -1, -1, 3, 10, 10, -9],
...                  [-1, 2, 0, -3, -7, 1, 6]])
sage: A.right_kernel(basis='echelon')
Free module of degree 7 and rank 2 over Integer Ring
Echelon basis matrix:
[ 1  5 -8  3 -1 -1 -1]
[ 0 11 -19 5 -2 -3 -3]
sage: B = copy(A)
sage: B.right_kernel(basis='LLL')
Free module of degree 7 and rank 2 over Integer Ring
User basis matrix:
[ 2 -1  3  1  0  1  1]
[-5 -3  2 -5  1 -1 -1]
sage: C = copy(A)
sage: C.right_kernel(basis='pivot')
Traceback (most recent call last):
...
ValueError: pivot basis only available over a field, not over Integer Ring
sage: D = copy(A)
sage: D.right_kernel(algorithm='pari')
Free module of degree 7 and rank 2 over Integer Ring
Echelon basis matrix:
[ 1  5 -8  3 -1 -1 -1]
[ 0 11 -19 5 -2 -3 -3]
sage: E = copy(A)
sage: E.right_kernel(algorithm='padic', basis='LLL')
Free module of degree 7 and rank 2 over Integer Ring
User basis matrix:
[-2  1 -3 -1  0 -1 -1]
[ 5  3 -2  5 -1  1  1]

```

Besides the integers, rings may be as general as principal ideal domains. Results are then free modules.

```
sage: R.<y> = QQ[]
sage: A = matrix(R, [[ 1, y, 1+y^2],
...                  [y^3, y^2, 2*y^3]])
sage: K = A.right_kernel(algorithm='default', basis='echelon'); K
Free module of degree 3 and rank 1 over Univariate Polynomial Ring in y over Rational Field
Echelon basis matrix:
[-1 -y  1]
sage: A*K.basis_matrix().transpose() == zero_matrix(ZZ, 2, 1)
True
```

It is possible to compute a kernel for a matrix over an integral domain which is not a PID, but usually this will fail.

```
sage: D.<x> = ZZ[]
sage: A = matrix(D, 2, 2, [[x^2 - x, -x + 5],
...                       [x^2 - 8, -x + 2]])
sage: A.right_kernel()
Traceback (most recent call last):
...
ArithmeticError: Ideal Ideal (x^2 - x, x^2 - 8) of Univariate Polynomial Ring in x over Inte
```

Matrices over non-commutative rings are not a good idea either. These are the “usual” quaternions.

```
sage: Q.<i,j,k> = QuaternionAlgebra(-1,-1)
sage: A = matrix(Q, 2, [i,j,-1,k])
sage: A.right_kernel()
Traceback (most recent call last):
...
NotImplementedError: Cannot compute a matrix kernel over Quaternion Algebra (-1, -1) with ba
```

Sparse matrices, over the rationals and the integers, use the same routines as the dense versions.

```
sage: A = matrix(ZZ, [[0, -1, 1, 1, 2],
...                  [1, -2, 0, 1, 3],
...                  [-1, 2, 0, -1, -3]],
...               sparse=True)
sage: A.right_kernel()
Free module of degree 5 and rank 3 over Integer Ring
Echelon basis matrix:
[ 1  0  0  2 -1]
[ 0  1  0 -1  1]
[ 0  0  1 -3  1]
sage: B = A.change_ring(QQ)
sage: B.is_sparse()
True
sage: B.right_kernel()
Vector space of degree 5 and dimension 3 over Rational Field
Basis matrix:
[ 1  0  0  2 -1]
[ 0  1  0 -1  1]
[ 0  0  1 -3  1]
```

With no columns, the kernel can only have dimension zero. With no rows, every possible vector is in the kernel.

```
sage: A = matrix(QQ, 2, 0)
sage: A.right_kernel()
Vector space of degree 0 and dimension 0 over Rational Field
Basis matrix:
```

```
[ ]
sage: A = matrix(QQ, 0, 2)
sage: A.right_kernel()
Vector space of degree 2 and dimension 2 over Rational Field
Basis matrix:
[1 0]
[0 1]
```

Every vector is in the kernel of a zero matrix, the dimension is the number of columns.

```
sage: A = zero_matrix(QQ, 10, 20)
sage: A.right_kernel()
Vector space of degree 20 and dimension 20 over Rational Field
Basis matrix:
20 x 20 dense matrix over Rational Field
```

Results are cached as the right kernel of the matrix. Subsequent requests for the right kernel will return the cached result, without regard for new values of the algorithm or format keyword. Work with a copy if you need a new right kernel, or perhaps investigate the `right_kernel_matrix()` method, which does not cache its results and is more flexible.

```
sage: A = matrix(QQ, 3, range(9))
sage: K1 = A.right_kernel(basis='echelon')
sage: K1
Vector space of degree 3 and dimension 1 over Rational Field
Basis matrix:
[ 1 -2  1]
sage: K2 = A.right_kernel(basis='pivot')
sage: K2
Vector space of degree 3 and dimension 1 over Rational Field
Basis matrix:
[ 1 -2  1]
sage: K1 is K2
True
sage: B = copy(A)
sage: K3 = B.kernel(basis='pivot')
sage: K3
Vector space of degree 3 and dimension 1 over Rational Field
User basis matrix:
[ 1 -2  1]
sage: K3 is K1
False
sage: K3 == K1
True
```

right_kernel_matrix(*args, **kws)

Returns a matrix whose rows form a basis for the right kernel of `self`.

INPUT:

- `algorithm` - default: 'default' - a keyword that selects the algorithm employed. Allowable values are:
 - 'default' - allows the algorithm to be chosen automatically
 - 'generic' - naive algorithm usable for matrices over any field
 - 'pari' - PARI library code for matrices over number fields or the integers
 - 'padic' - padic algorithm from IML library for matrices over the rationals and integers

- ‘pluq’ - PLUQ matrix factorization for matrices mod 2
- basis** - default: ‘echelon’ - a keyword that describes the format of the basis returned. Allowable values are:
 - ‘echelon’: the basis matrix is in echelon form
 - ‘pivot’: each basis vector is computed from the reduced row-echelon form of `self` by placing a single one in a non-pivot column and zeros in the remaining non-pivot columns. Only available for matrices over fields.
 - ‘LLL’: an LLL-reduced basis. Only available for matrices over the integers.
 - ‘computed’: no work is done to transform the basis, it is returned exactly as provided by whichever routine actually computed the basis. Request this for the least possible computation possible, but with no guarantees about the format of the basis.

OUTPUT:

A matrix `X` whose rows are an independent set spanning the right kernel of `self`. So `self*X.transpose()` is a zero matrix.

The output varies depending on the choice of `algorithm` and the format chosen by `basis`.

The results of this routine are not cached, so you can call it again with different options to get possibly different output (like the basis format). Conversely, repeated calls on the same matrix will always start from scratch.

Note: If you want to get the most basic description of a kernel, with a minimum of overhead, then ask for the right kernel matrix with the basis format requested as ‘computed’. You are then free to work with the output for whatever purpose. For a left kernel, call this method on the transpose of your matrix.

For greater convenience, plus cached results, request an actual vector space or free module with `right_kernel()` or `left_kernel()`.

EXAMPLES:

Over the Rational Numbers:

Kernels are computed by the IML library in `_right_kernel_matrix()`. Setting the `algorithm` keyword to ‘default’, ‘padic’ or unspecified will yield the same result, as there is no optional behavior. The ‘computed’ format of the basis vectors are exactly the negatives of the vectors in the ‘pivot’ format.

```
sage: A = matrix(QQ, [[1, 0, 1, -3, 1],
...                  [-5, 1, 0, 7, -3],
...                  [0, -1, -4, 6, -2],
...                  [4, -1, 0, -6, 2]])
sage: C = A.right_kernel_matrix(algorithm='default', basis='computed'); C
[-1  2 -2 -1  0]
[ 1  2  0  0 -1]
sage: A*C.transpose() == zero_matrix(QQ, 4, 2)
True
sage: P = A.right_kernel_matrix(algorithm='padic', basis='pivot'); P
[ 1 -2  2  1  0]
[-1 -2  0  0  1]
sage: A*P.transpose() == zero_matrix(QQ, 4, 2)
True
sage: C == -P
True
sage: E = A.right_kernel_matrix(algorithm='default', basis='echelon'); E
[  1  0  1  1/2 -1/2]
```

```
[ 0 1 -1/2 -1/4 -1/4]
sage: A*E.transpose() == zero_matrix(QQ, 4, 2)
True
```

Since the rationals are a field, we can call the general code available for any field by using the ‘generic’ keyword.

```
sage: A = matrix(QQ, [[1, 0, 1, -3, 1],
...                  [-5, 1, 0, 7, -3],
...                  [0, -1, -4, 6, -2],
...                  [4, -1, 0, -6, 2]])
sage: G = A.right_kernel_matrix(algorithm='generic', basis='echelon'); G
[ 1 0 1 1/2 -1/2]
[ 0 1 -1/2 -1/4 -1/4]
sage: A*G.transpose() == zero_matrix(QQ, 4, 2)
True
```

We verify that the rational matrix code is called for both dense and sparse rational matrices, with equal result.

```
sage: A = matrix(QQ, [[1, 0, 1, -3, 1],
...                  [-5, 1, 0, 7, -3],
...                  [0, -1, -4, 6, -2],
...                  [4, -1, 0, -6, 2]])
...                  sparse=False)
sage: B = copy(A).sparse_matrix()
sage: set_verbosity(1)
sage: D = A.right_kernel(); D
verbose ...
verbose 1 (<module>) computing right kernel matrix over the rationals for 4x5 matrix
...
verbose 1 (<module>) done computing right kernel matrix over the rationals for 4x5 matrix
...
Vector space of degree 5 and dimension 2 over Rational Field
Basis matrix:
[ 1 0 1 1/2 -1/2]
[ 0 1 -1/2 -1/4 -1/4]
sage: S = B.right_kernel(); S
verbose ...
verbose 1 (<module>) computing right kernel matrix over the rationals for 4x5 matrix
...
verbose 1 (<module>) done computing right kernel matrix over the rationals for 4x5 matrix
...
Vector space of degree 5 and dimension 2 over Rational Field
Basis matrix:
[ 1 0 1 1/2 -1/2]
[ 0 1 -1/2 -1/4 -1/4]
sage: set_verbosity(0)
sage: D == S
True
```

Over Number Fields:

Kernels are by default computed by PARI, (except for exceptions like the rationals themselves). The raw results from PARI are a pivot basis, so the *basis* keywords ‘computed’ and ‘pivot’ will return the same results.

```
sage: Q = QuadraticField(-7)
sage: a = Q.gen(0)
sage: A = matrix(Q, [[2, 5-a, 15-a, 16+4*a],
```

```

...          [2+a, a, -7 + 5*a, -3+3*a]])
sage: C = A.right_kernel_matrix(algorithm='default', basis='computed'); C
[  -a      -3      1      0]
[ -2 -a - 1      0      1]
sage: A*C.transpose() == zero_matrix(Q, 2, 2)
True
sage: P = A.right_kernel_matrix(algorithm='pari', basis='pivot'); P
[  -a      -3      1      0]
[ -2 -a - 1      0      1]
sage: A*P.transpose() == zero_matrix(Q, 2, 2)
True
sage: E = A.right_kernel_matrix(algorithm='default', basis='echelon'); E
[          1          0      7/88*a + 3/88 -3/176*a - 39/176]
[          0          1 -1/88*a - 13/88 13/176*a - 7/176]
sage: A*E.transpose() == zero_matrix(Q, 2, 2)
True

```

We can bypass using PARI for number fields and use Sage's general code for matrices over any field. The basis vectors as computed are in pivot format.

```

sage: Q = QuadraticField(-7)
sage: a = Q.gen(0)
sage: A = matrix(Q, [[2, 5-a, 15-a, 16+4*a], [2+a, a, -7 + 5*a, -3+3*a]])
sage: G = A.right_kernel_matrix(algorithm='generic', basis='computed'); G
[  -a      -3      1      0]
[ -2 -a - 1      0      1]
sage: A*G.transpose() == zero_matrix(Q, 2, 2)
True

```

We check that number fields are handled by the right routine as part of typical right kernel computation.

```

sage: Q = QuadraticField(-7)
sage: a = Q.gen(0)
sage: A = matrix(Q, [[2, 5-a, 15-a, 16+4*a], [2+a, a, -7 + 5*a, -3+3*a]])
sage: set_verbosity(1)
sage: A.right_kernel(algorithm='default')
verbose ...
verbose 1 (<module>) computing right kernel matrix over a number field for 2x4 matrix
verbose 1 (<module>) done computing right kernel matrix over a number field for 2x4 matrix
...
Vector space of degree 4 and dimension 2 over Number Field in a with defining polynomial x^2
Basis matrix:
[          1          0      7/88*a + 3/88 -3/176*a - 39/176]
[          0          1 -1/88*a - 13/88 13/176*a - 7/176]
sage: set_verbosity(0)

```

Over the Finite Field of Order 2:

Kernels are computed by the M4RI library using PLUQ matrix decomposition in the `_right_kernel_matrix()` method. There are no options for the algorithm used.

```

sage: A = matrix(GF(2), [[0, 1, 1, 0, 0, 0],
...                      [1, 0, 0, 0, 1, 1],
...                      [1, 0, 0, 0, 1, 1]])
sage: E = A.right_kernel_matrix(algorithm='default', format='echelon'); E
[1 0 0 0 0 1]
[0 1 1 0 0 0]
[0 0 0 1 0 0]
[0 0 0 0 1 1]
sage: A*E.transpose() == zero_matrix(GF(2), 3, 4)

```

True

Since GF(2) is a field we can route this computation to the generic code and obtain the ‘pivot’ form of the basis. The algorithm keywords, ‘plug’, ‘default’ and unspecified, all have the same effect as there is no optional behavior.

```
sage: A = matrix(GF(2), [[0, 1, 1, 0, 0, 0],
...                      [1, 0, 0, 0, 1, 1],
...                      [1, 0, 0, 0, 1, 1]])
sage: P = A.right_kernel_matrix(algorithm='generic', basis='pivot'); P
[0 1 1 0 0 0]
[0 0 0 1 0 0]
[1 0 0 0 1 0]
[1 0 0 0 0 1]
sage: A*P.transpose() == zero_matrix(GF(2), 3, 4)
True
sage: DP = A.right_kernel_matrix(algorithm='default', basis='pivot'); DP
[0 1 1 0 0 0]
[0 0 0 1 0 0]
[1 0 0 0 1 0]
[1 0 0 0 0 1]
sage: A*DP.transpose() == zero_matrix(GF(2), 3, 4)
True
sage: A.right_kernel_matrix(algorithm='plug', basis='echelon')
[1 0 0 0 0 1]
[0 1 1 0 0 0]
[0 0 0 1 0 0]
[0 0 0 0 1 1]
```

We test that the mod 2 code is called for matrices over GF(2).

```
sage: A = matrix(GF(2), [[0, 1, 1, 0, 0, 0],
...                      [1, 0, 0, 0, 1, 1],
...                      [1, 0, 0, 0, 1, 1]])
sage: set_verbose(1)
sage: A.right_kernel(algorithm='default')
verbose ...
verbose 1 (<module>) computing right kernel matrix over integers mod 2 for 3x6 matrix
verbose 1 (<module>) done computing right kernel matrix over integers mod 2 for 3x6 matrix
...
Vector space of degree 6 and dimension 4 over Finite Field of size 2
Basis matrix:
[1 0 0 0 0 1]
[0 1 1 0 0 0]
[0 0 0 1 0 0]
[0 0 0 0 1 1]
sage: set_verbose(0)
```

Over Arbitrary Fields:

For kernels over fields not listed above, totally general code will compute a set of basis vectors in the pivot format. These could be returned as a basis in echelon form.

```
sage: F.<a> = FiniteField(5^2)
sage: A = matrix(F, 3, 4, [[ 1, a, 1+a, a^3+a^5],
...                      [ a, a^4, a+a^4, a^4+a^8],
...                      [a^2, a^6, a^2+a^6, a^5+a^10]])
sage: P = A.right_kernel_matrix(algorithm='default', basis='pivot'); P
[      4      4      1      0]
```

```

[ a + 2 3*a + 3      0      1]
sage: A*P.transpose() == zero_matrix(F, 3, 2)
True
sage: E = A.right_kernel_matrix(algorithm='default', basis='echelon'); E
[      1      0 3*a + 4 2*a + 2]
[      0      1      2*a 3*a + 3]
sage: A*E.transpose() == zero_matrix(F, 3, 2)
True

```

This general code can be requested for matrices over any field with the `algorithm` keyword 'generic'. Normally, matrices over the rationals would be handled by specific routines from the IML library. The default format is an echelon basis, but a pivot basis may be requested, which is identical to the computed basis.

```

sage: A = matrix(QQ, 3, 4, [[1,3,-2,4],
...                        [2,0,2,2],
...                        [-1,1,-2,0]])
sage: G = A.right_kernel_matrix(algorithm='generic'); G
[  1      0 -1/2 -1/2]
[  0      1  1/2 -1/2]
sage: A*G.transpose() == zero_matrix(QQ, 3, 2)
True
sage: C = A.right_kernel_matrix(algorithm='generic', basis='computed'); C
[-1  1  1  0]
[-1 -1  0  1]
sage: A*C.transpose() == zero_matrix(QQ, 3, 2)
True

```

We test that the generic code is called for matrices over fields, lacking any more specific routine.

```

sage: F.<a> = FiniteField(5^2)
sage: A = matrix(F, 3, 4, [[ 1,  a,  1+a,  a^3+a^5],
...                        [ a, a^4,  a+a^4, a^4+a^8],
...                        [a^2, a^6, a^2+a^6, a^5+a^10]])
sage: set_verbose(1)
sage: A.right_kernel(algorithm='default')
verbose ...
verbose 1 (<module>) computing right kernel matrix over an arbitrary field for 3x4 matrix
...
verbose 1 (<module>) done computing right kernel matrix over an arbitrary field for 3x4 matrix
...
Vector space of degree 4 and dimension 2 over Finite Field in a of size 5^2
Basis matrix:
[      1      0 3*a + 4 2*a + 2]
[      0      1      2*a 3*a + 3]
sage: set_verbose(0)

```

Over the Integers:

Either the IML or PARI libraries are used to provide a set of basis vectors. The `algorithm` keyword can be used to select either, or when set to 'default' a heuristic will choose between the two. Results can be returned in the 'compute' format, straight out of the libraries. Unique to the integers, the basis vectors can be returned as an LLL basis. Note the similarities and differences in the results. The 'pivot' format is not available, since the integers are not a field.

```

sage: A = matrix(ZZ, [[8, 0, 7, 1, 3, 4, 6],
...                  [4, 0, 3, 4, 2, 7, 7],
...                  [1, 4, 6, 1, 2, 8, 5],
...                  [0, 3, 1, 2, 3, 6, 2]])

```

```

sage: X = A.right_kernel_matrix(algorithm='default', basis='echelon'); X
[ 1 12  3 14 -3 -10  1]
[ 0 35  0 25 -1 -31 17]
[ 0  0  7 12 -3 -1 -8]
sage: A*X.transpose() == zero_matrix(ZZ, 4, 3)
True

```

```

sage: X = A.right_kernel_matrix(algorithm='padic', basis='LLL'); X
[-3 -1  5  7  2 -3 -2]
[ 3  1  2  5 -5  2 -6]
[-4 -13  2 -7  5  7 -3]
sage: A*X.transpose() == zero_matrix(ZZ, 4, 3)
True

```

```

sage: X = A.right_kernel_matrix(algorithm='pari', basis='computed'); X
[ 3  1 -5 -7 -2  3  2]
[ 3  1  2  5 -5  2 -6]
[ 4 13 -2  7 -5 -7  3]
sage: A*X.transpose() == zero_matrix(ZZ, 4, 3)
True

```

```

sage: X = A.right_kernel_matrix(algorithm='padic', basis='computed'); X
[ 265 345 -178 17 -297  0  0]
[-242 -314 163 -14 271 -1  0]
[-36 -47 25 -1 40  0 -1]
sage: A*X.transpose() == zero_matrix(ZZ, 4, 3)
True

```

We test that the code for integer matrices is called for matrices defined over the integers, both dense and sparse, with equal result.

```

sage: A = matrix(ZZ, [[8, 0, 7, 1, 3, 4, 6],
...                  [4, 0, 3, 4, 2, 7, 7],
...                  [1, 4, 6, 1, 2, 8, 5],
...                  [0, 3, 1, 2, 3, 6, 2]],
...               sparse=False)
sage: B = copy(A).sparse_matrix()
sage: set_verbose(1)
sage: D = A.right_kernel(); D
verbose ...
verbose 1 (<module>) computing right kernel matrix over the integers for 4x7 matrix
verbose 1 (<module>) done computing right kernel matrix over the integers for 4x7 matrix
...
Free module of degree 7 and rank 3 over Integer Ring
Echelon basis matrix:
[ 1 12  3 14 -3 -10  1]
[ 0 35  0 25 -1 -31 17]
[ 0  0  7 12 -3 -1 -8]
sage: S = B.right_kernel(); S
verbose ...
verbose 1 (<module>) computing right kernel matrix over the integers for 4x7 matrix
verbose 1 (<module>) done computing right kernel matrix over the integers for 4x7 matrix
...
Free module of degree 7 and rank 3 over Integer Ring
Echelon basis matrix:
[ 1 12  3 14 -3 -10  1]
[ 0 35  0 25 -1 -31 17]
[ 0  0  7 12 -3 -1 -8]

```

```
sage: set_verbose(0)
sage: D == S
True
```

Over Principal Ideal Domains:

Kernels can be computed using Smith normal form. Only the default algorithm is available, and the ‘pivot’ basis format is not available.

```
sage: R.<y> = QQ[]
sage: A = matrix(R, [[ 1, y, 1+y^2],
...                  [y^3, y^2, 2*y^3]])
sage: E = A.right_kernel_matrix(algorithm='default', basis='echelon'); E
[-1 -y  1]
sage: A*E.transpose() == zero_matrix(ZZ, 2, 1)
True
```

It can be computationally expensive to determine if an integral domain is a principal ideal domain. The Smith normal form routine can fail for non-PIDs, as in this example.

```
sage: D.<x> = ZZ[]
sage: A = matrix(D, 2, 2, [[x^2 - x, -x + 5],
...                       [x^2 - 8, -x + 2]])
sage: A.right_kernel_matrix()
Traceback (most recent call last):
...
ArithmeticError: Ideal Ideal (x^2 - x, x^2 - 8) of Univariate Polynomial Ring in x over Inte
```

We test that the domain code is called for domains that lack any extra structure.

```
sage: R.<y> = QQ[]
sage: A = matrix(R, [[ 1, y, 1+y^2],
...                  [y^3, y^2, 2*y^3]])
sage: set_verbose(1)
sage: A.right_kernel(algorithm='default', basis='echelon')
verbose ...
verbose 1 (<module>) computing right kernel matrix over a domain for 2x3 matrix
verbose 1 (<module>) done computing right kernel matrix over a domain for 2x3 matrix
...
Free module of degree 3 and rank 1 over Univariate Polynomial Ring in y over Rational Field
Echelon basis matrix:
[-1 -y  1]
sage: set_verbose(0)
```

Trivial Cases:

We test two trivial cases. Any possible values for the keywords (algorithm, basis) will return identical results.

```
sage: A = matrix(ZZ, 0, 2)
sage: A.right_kernel_matrix()
[1 0]
[0 1]
sage: A = matrix(FiniteField(7), 2, 0)
sage: A.right_kernel_matrix().parent()
Full MatrixSpace of 0 by 0 dense matrices over Finite Field of size 7
```

TESTS:

The “usual” quaternions are a non-commutative ring and computations of kernels over these rings are not

implemented.

```
sage: Q.<i,j,k> = QuaternionAlgebra(-1,-1)
sage: A = matrix(Q, 2, [i,j,-1,k])
sage: A.right_kernel_matrix()
Traceback (most recent call last):
...
NotImplementedError: Cannot compute a matrix kernel over Quaternion Algebra (-1, -1) with ba
```

We test error messages for improper choices of the ‘algorithm’ keyword.

```
sage: matrix(ZZ, 2, 2).right_kernel_matrix(algorithm='junk')
Traceback (most recent call last):
...
ValueError: matrix kernel algorithm 'junk' not recognized
sage: matrix(GF(2), 2, 2).right_kernel_matrix(algorithm='padic')
Traceback (most recent call last):
...
ValueError: 'padic' matrix kernel algorithm only available over the rationals and the integers
sage: matrix(QQ, 2, 2).right_kernel_matrix(algorithm='pari')
Traceback (most recent call last):
...
ValueError: 'pari' matrix kernel algorithm only available over non-trivial number fields and
sage: matrix(Integers(6), 2, 2).right_kernel_matrix(algorithm='generic')
Traceback (most recent call last):
...
ValueError: 'generic' matrix kernel algorithm only available over a field, not over Ring of
sage: matrix(QQ, 2, 2).right_kernel_matrix(algorithm='pluq')
Traceback (most recent call last):
...
ValueError: 'pluq' matrix kernel algorithm only available over integers mod 2, not over Rati
```

We test error messages for improper basis format requests.

```
sage: matrix(ZZ, 2, 2).right_kernel_matrix(basis='junk')
Traceback (most recent call last):
...
ValueError: matrix kernel basis format 'junk' not recognized
sage: matrix(ZZ, 2, 2).right_kernel_matrix(basis='pivot')
Traceback (most recent call last):
...
ValueError: pivot basis only available over a field, not over Integer Ring
sage: matrix(QQ, 2, 2).right_kernel_matrix(basis='LLL')
Traceback (most recent call last):
...
ValueError: LLL-reduced basis only available over the integers, not over Rational Field
```

Finally, error messages for the ‘proof’ keyword.

```
sage: matrix(ZZ, 2, 2).right_kernel_matrix(proof='junk')
Traceback (most recent call last):
...
ValueError: 'proof' must be one of True, False or None, not junk
sage: matrix(QQ, 2, 2).right_kernel_matrix(proof=True)
Traceback (most recent call last):
...
ValueError: 'proof' flag only valid for matrices over the integers
```

AUTHOR:

•Rob Beezer (2011-02-05)

right_nullity()

Return the right nullity of this matrix, which is the dimension of the right kernel.

EXAMPLES:

```
sage: A = MatrixSpace(QQ, 3, 2) (range(6))
```

```
sage: A.right_nullity()
```

```
0
```

```
sage: A = matrix(ZZ, 3, range(9))
```

```
sage: A.right_nullity()
```

```
1
```

rook_vector (*check=False*)

Return the rook vector of the matrix `self`.

Let A be an m by n $(0,1)$ -matrix with $m \leq n$. We identify A with a chessboard where rooks can be placed on the fields (i, j) with $a_{i,j} = 1$. The number $r_k = p_k(A)$ (the permanental k -minor) counts the number of ways to place k rooks on this board so that no rook can attack another.

The rook vector of the matrix A is the list consisting of r_0, r_1, \dots, r_m .

The rook polynomial is defined by $r(x) = \sum_{k=0}^m r_k x^k$.

INPUT:

- `self` – an m by n $(0,1)$ -matrix with $m \leq n$
- `check` – Boolean (default: `False`) determining whether to check that `self` is a $(0,1)$ -matrix.

OUTPUT:

The rook vector of the matrix `self`.

EXAMPLES:

```
sage: M = MatrixSpace(ZZ, 3, 6)
```

```
sage: A = M([1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1])
```

```
sage: A.rook_vector()
```

```
[1, 12, 40, 36]
```

```
sage: R.<x> = PolynomialRing(ZZ)
```

```
sage: rv = A.rook_vector()
```

```
sage: rook_polynomial = sum([rv[k] * x^k for k in range(len(rv))])
```

```
sage: rook_polynomial
```

```
36*x^3 + 40*x^2 + 12*x + 1
```

AUTHORS:

- Jaap Spies (2006-02-24)

row_module (*base_ring=None*)

Return the free module over the base ring spanned by the rows of `self`.

EXAMPLES:

```
sage: A = MatrixSpace(IntegerRing(), 2) ([1, 2, 3, 4])
```

```
sage: A.row_module()
```

Free module of degree 2 and rank 2 over Integer Ring

Echelon basis matrix:

```
[1 0]
```

```
[0 2]
```

row_space (*base_ring=None*)

Return the row space of this matrix. (Synonym for `self.row_module()`.)

EXAMPLES:

```
sage: t = matrix(QQ, 3, range(9)); t
```

```
[0 1 2]
```

```
[3 4 5]
```

```
[6 7 8]
```

```
sage: t.row_space()
```

Vector space of degree 3 and dimension 2 over Rational Field

Basis matrix:

```
[ 1  0 -1]
```

```
[ 0  1  2]
```

```
sage: m = Matrix(Integers(5), 2, 2, [2, 2, 2, 2]);
```

```
sage: m.row_space()
```

Vector space of degree 2 and dimension 1 over Ring of integers modulo 5

Basis matrix:

```
[1 1]
```

rref (**args, **kws*)

Return the reduced row echelon form of the matrix, considered as a matrix over a field.

If the matrix is over a ring, then an equivalent matrix is constructed over the fraction field, and then row reduced.

All arguments are passed on to `:meth:echelon_form`.

Note: Because the matrix is viewed as a matrix over a field, every leading coefficient of the returned matrix will be one and will be the only nonzero entry in its column.

EXAMPLES:

```
sage: A=matrix(3, range(9)); A
```

```
[0 1 2]
```

```
[3 4 5]
```

```
[6 7 8]
```

```
sage: A.rref()
```

```
[ 1  0 -1]
```

```
[ 0  1  2]
```

```
[ 0  0  0]
```

Note that there is a difference between `rref()` and `echelon_form()` when the matrix is not over a field (in this case, the integers instead of the rational numbers):

```
sage: A.base_ring()
```

Integer Ring

```
sage: A.echelon_form()
```

```
[ 3  0 -3]
```

```
[ 0  1  2]
```

```
[ 0  0  0]
```

```
sage: B=random_matrix(QQ, 3, num_bound=10); B
```

```
[ -4  -3   6]
```

```
[  5  -5  9/2]
```

```
[3/2  -4  -7]
```

```
sage: B.rref()
```

```
[1 0 0]
```

```
[0 1 0]
[0 0 1]
```

In this case, since B is a matrix over a field (the rational numbers), `rref()` and `echelon_form()` are exactly the same:

```
sage: B.echelon_form()
[1 0 0]
[0 1 0]
[0 0 1]
sage: B.echelon_form() is B.rref()
True
```

Since `echelon_form()` is not implemented for every ring, sometimes behavior varies, as here:

```
sage: R.<x>=ZZ[]
sage: C = matrix(3, [2,x,x^2,x+1,3-x,-1,3,2,1])
sage: C.rref()
[1 0 0]
[0 1 0]
[0 0 1]
sage: C.base_ring()
Univariate Polynomial Ring in x over Integer Ring
sage: C.echelon_form()
Traceback (most recent call last):
...
NotImplementedError: Ideal Ideal (2, x + 1) of Univariate Polynomial Ring in x over Integer
Echelon form not implemented over 'Univariate Polynomial Ring in x over Integer Ring'.
sage: C = matrix(3, [2,x,x^2,x+1,3-x,-1,3,2,1/2])
sage: C.echelon_form()
[
                                2
                                x
                                1
                                15*x^2 - 3/2*x
                                0 5/2*x^3 - 15/4*x^2 - 9/4*x
                                0
sage: C.rref()
[1 0 0]
[0 1 0]
[0 0 1]
sage: C = matrix(3, [2,x,x^2,x+1,3-x,-1/x,3,2,1/2])
sage: C.echelon_form()
[1 0 0]
[0 1 0]
[0 0 1]
```

set_block(*row*, *col*, *block*)

Sets the sub-matrix of self, with upper left corner given by *row*, *col* to *block*.

EXAMPLES:

```
sage: A = matrix(QQ, 3, 3, range(9))/2
sage: B = matrix(ZZ, 2, 1, [100,200])
sage: A.set_block(0, 1, B)
sage: A
[ 0 100 1]
[3/2 200 5/2]
[ 3 7/2 4]
```

We test that an exception is raised when the block is out of bounds:

```
sage: matrix([1]).set_block(0,1,matrix([1]))
Traceback (most recent call last):
```

```
...
IndexError: matrix window index out of range
```

smith_form()

If self is a matrix over a principal ideal domain R, return matrices D, U, V over R such that $D = U * \text{self} * V$, U and V have unit determinant, and D is diagonal with diagonal entries the ordered elementary divisors of self, ordered so that $D_i \mid D_{i+1}$. Note that U and V are not uniquely defined in general, and D is defined only up to units.

INPUT:

- self - a matrix over an integral domain. If the base ring is not a PID, the routine might work, or else it will fail having found an example of a non-principal ideal. Note that we do not call any methods to check whether or not the base ring is a PID, since this might be quite expensive (e.g. for rings of integers of number fields of large degree).

ALGORITHM: Lifted wholesale from http://en.wikipedia.org/wiki/Smith_normal_form

See Also:

`elementary_divisors()`

AUTHORS:

- David Loeffler (2008-12-05)

EXAMPLES:

An example over the ring of integers of a number field (of class number 1):

```
sage: OE = NumberField(x^2 - x + 2, 'w').ring_of_integers()
sage: w = OE.ring_generators()[0]
sage: m = Matrix([ [1, w], [w, 7] ])
sage: d, u, v = m.smith_form()
sage: (d, u, v)
(
  [      1      0]  [ 1  0]  [ 1 -w]
  [      0 -w + 9], [-w  1], [ 0  1]
)
sage: u * m * v == d
True
sage: u.base_ring() == v.base_ring() == d.base_ring() == OE
True
sage: u.det().is_unit() and v.det().is_unit()
True
```

An example over the polynomial ring $\mathbb{Q}\mathbb{Q}[x]$:

```
sage: R.<x> = QQ[]; m=x*matrix(R,2,2,1) - matrix(R, 2,2,[3,-4,1,-1]); m.smith_form()
(
  [      1      0]  [ 0  -1]  [ 1 x + 1]
  [      0 x^2 - 2*x + 1], [ 1 x - 3], [ 0  1]
)

```

An example over a field:

```
sage: m = matrix( GF(17), 3, 3, [11,5,1,3,6,8,1,16,0] ); d,u,v = m.smith_form()
sage: d
[1 0 0]
[0 1 0]
[0 0 0]
```

```
sage: u*m*v == d
True
```

Some examples over non-PID's work anyway:

```
sage: R = EquationOrder(x^2 + 5, 's') # class number 2
sage: s = R.ring_generators()[0]
sage: A = matrix(R, 2, 2, [s-1, -s, -s, 2*s+1])
sage: D, U, V = A.smith_form()
sage: D, U, V
(
[      1      0] [      4 s + 4] [      1 -5*s + 6]
[      0 -s - 6], [      s s - 1], [      0      1]
)
sage: D == U*A*V
True
```

Others don't, but they fail quite constructively:

```
sage: matrix(R, 2, 2, [s-1, -s-2, -2*s, -s-2]).smith_form()
Traceback (most recent call last):
...
ArithmeticError: Ideal Fractional ideal (2, s + 1) not principal
```

Empty matrices are handled safely:

```
sage: m = MatrixSpace(OE, 2, 0)(0); d,u,v=m.smith_form(); u*m*v == d
True
sage: m = MatrixSpace(OE, 0, 2)(0); d,u,v=m.smith_form(); u*m*v == d
True
sage: m = MatrixSpace(OE, 0, 0)(0); d,u,v=m.smith_form(); u*m*v == d
True
```

Some pathological cases that crashed earlier versions:

```
sage: m = Matrix(OE, [[2*w, 2*w-1, -w+1], [2*w+2, -2*w-1, w-1], [-2*w-1, -2*w-2, 2*w-1]]); d, u, v = m.smith_form()
True
sage: m = matrix(OE, 3, 3, [-5*w-1, -2*w-2, 4*w-10, 8*w, -w, w-1, -1, 1, -8]); d,u,v = m.smith_form()
True
```

solve_left (*B*, *check=True*)

If *self* is a matrix *A*, then this function returns a vector or matrix *X* such that $XA = B$. If *B* is a vector then *X* is a vector and if *B* is a matrix, then *X* is a matrix.

INPUT:

- *B* - a matrix
- *check* - bool (default: True) - if False and *self* is nonsquare, may not raise an error message even if there is no solution. This is faster but more dangerous.

EXAMPLES:

```
sage: A = matrix(QQ, 4, 2, [0, -1, 1, 0, -2, 2, 1, 0])
sage: B = matrix(QQ, 2, 2, [1, 0, 1, -1])
sage: X = A.solve_left(B)
sage: X*A == B
True
```

TESTS:

```
sage: A = matrix(QQ,4,2, [0, -1, 1, 0, -2, 2, 1, 0])
sage: B = vector(QQ,2, [2,1])
sage: X = A.solve_left(B)
sage: X*A == B
True
sage: X
(-1, 2, 0, 0)
sage: A = Matrix(Zmod(128), 2, 3, [5, 29, 33, 64, 0, 7])
sage: B = vector(Zmod(128), [31,39,56])
sage: X = A.solve_left(B); X
(19, 83)
sage: X * A == B
True
```

solve_right (*B*, *check=True*)

If *self* is a matrix *A*, then this function returns a vector or matrix *X* such that $AX = B$. If *B* is a vector then *X* is a vector and if *B* is a matrix, then *X* is a matrix.

Note: In Sage one can also write `A \ B` for `A.solve_right(B)`, i.e., Sage implements the “the MATLAB/Octave backslash operator”.

INPUT:

- *B* - a matrix or vector
- *check* - bool (default: True) - if False and *self* is nonsquare, may not raise an error message even if there is no solution. This is faster but more dangerous.

OUTPUT: a matrix or vector

See Also:

`solve_left()`

EXAMPLES:

```
sage: A = matrix(QQ, 3, [1,2,3,-1,2,5,2,3,1])
sage: b = vector(QQ,[1,2,3])
sage: x = A \ b; x
(-13/12, 23/12, -7/12)
sage: A * x
(1, 2, 3)
```

We solve with *A* nonsquare:

```
sage: A = matrix(QQ,2,4, [0, -1, 1, 0, -2, 2, 1, 0]); B = matrix(QQ,2,2, [1, 0, 1, -1])
sage: X = A.solve_right(B); X
[-3/2  1/2]
[ -1    0]
[  0    0]
[  0    0]
sage: A*X == B
True
```

Another nonsingular example:

```
sage: A = matrix(QQ,2,3, [1,2,3,2,4,6]); v = vector([-1/2,-1])
sage: x = A \ v; x
(-1/2, 0, 0)
```

```
sage: A*x == v
True
```

Same example but over \mathbb{Z} :

```
sage: A = matrix(ZZ,2,3, [1,2,3,2,4,6]); v = vector([-1,-2])
sage: A \ v
(-1, 0, 0)
```

An example in which there is no solution:

```
sage: A = matrix(QQ,2,3, [1,2,3,2,4,6]); v = vector([1,1])
sage: A \ v
Traceback (most recent call last):
...
ValueError: matrix equation has no solutions
```

A `ValueError` is raised if the input is invalid:

```
sage: A = matrix(QQ,4,2, [0, -1, 1, 0, -2, 2, 1, 0])
sage: B = matrix(QQ,2,2, [1, 0, 1, -1])
sage: X = A.solve_right(B)
Traceback (most recent call last):
...
ValueError: number of rows of self must equal number of rows of B
```

We solve with A singular:

```
sage: A = matrix(QQ,2,3, [1,2,3,2,4,6]); B = matrix(QQ,2,2, [6, -6, 12, -12])
sage: X = A.solve_right(B); X
[ 6 -6]
[ 0  0]
[ 0  0]
sage: A*X == B
True
```

We illustrate left associativity, etc., of the backslash operator.

```
sage: A = matrix(QQ, 2, [1,2,3,4])
sage: A \ A
[1 0]
[0 1]
sage: A \ A \ A
[1 2]
[3 4]
sage: A.parent()(1) \ A
[1 2]
[3 4]
sage: A \ (A \ A)
[ -2  1]
[ 3/2 -1/2]
sage: X = A \ (A - 2); X
[ 5 -2]
[-3  2]
sage: A * X
[-1  2]
[ 3  2]
```

Solving over a polynomial ring:

```

sage: x = polygen(QQ, 'x')
sage: A = matrix(2, [x, 2*x, -5*x^2+1, 3])
sage: v = vector([3, 4*x - 2])
sage: X = A \ v
sage: X
((-8*x^2 + 4*x + 9)/(10*x^3 + x), (19*x^2 - 2*x - 3)/(10*x^3 + x))
sage: A * X == v
True

```

Solving some systems over $\mathbf{Z}/n\mathbf{Z}$:

```

sage: A = Matrix(Zmod(6), 3, 2, [1, 2, 3, 4, 5, 6])
sage: B = vector(Zmod(6), [1, 1, 1])
sage: A.solve_right(B)
(5, 1)
sage: B = vector(Zmod(6), [5, 1, 1])
sage: A.solve_right(B)
Traceback (most recent call last):
...
ValueError: matrix equation has no solutions
sage: A = Matrix(Zmod(128), 2, 3, [23, 11, 22, 4, 1, 0])
sage: B = Matrix(Zmod(128), 2, 1, [1, 0])
sage: A.solve_right(B)
[ 1]
[124]
[ 1]

```

Solving a system over the p-adics:

```

sage: k = Qp(5, 4)
sage: a = matrix(k, 3, [1, 7, 3, 2, 5, 4, 1, 1, 2]); a
[ 1 + O(5^4) 2 + 5 + O(5^4) 3 + O(5^4)]
[ 2 + O(5^4) 5 + O(5^5) 4 + O(5^4)]
[ 1 + O(5^4) 1 + O(5^4) 2 + O(5^4)]
sage: v = vector(k, 3, [1, 2, 3])
sage: x = a \ v; x
(4 + 5 + 5^2 + 3*5^3 + O(5^4), 2 + 5 + 3*5^2 + 5^3 + O(5^4), 1 + 5 + O(5^4))
sage: a * x == v
True

```

Solving a system of linear equation symbolically using symbolic matrices:

```

sage: var('a,b,c,d,x,y')
(a, b, c, d, x, y)
sage: A=matrix(SR, 2, [a,b,c,d]); A
[a b]
[c d]
sage: result=vector(SR, [3,5]); result
(3, 5)
sage: soln=A.solve_right(result)
sage: soln
(-b*(3*c/a - 5)/(a*(b*c/a - d)) + 3/a, (3*c/a - 5)/(b*c/a - d))
sage: (a*x+b*y).subs(x=soln[0],y=soln[1]).simplify_full()
3
sage: (c*x+d*y).subs(x=soln[0],y=soln[1]).simplify_full()
5
sage: (A*soln).apply_map(lambda x: x.simplify_full())
(3, 5)

```


subdivide (*row_lines=None, col_lines=None*)

Divides self into logical submatrices which can then be queried and extracted. If a subdivision already exists, this method forgets the previous subdivision and flushes the cache.

INPUT:

- *row_lines* - None, an integer, or a list of integers
- *col_lines* - None, an integer, or a list of integers

OUTPUT: changes self

Note: One may also pass a tuple into the first argument which will be interpreted as (*row_lines*, *col_lines*)

EXAMPLES:

```
sage: M = matrix(5, 5, prime_range(100))
sage: M.subdivide(2,3); M
[ 2  3  5| 7 11]
[13 17 19|23 29]
[-----+-----]
[31 37 41|43 47]
[53 59 61|67 71]
[73 79 83|89 97]
sage: M.subdivision(0,0)
[ 2  3  5]
[13 17 19]
sage: M.subdivision(1,0)
[31 37 41]
[53 59 61]
[73 79 83]
sage: M.subdivision_entry(1,0,0,0)
31
sage: M.subdivisions()
([2], [3])
sage: M.subdivide(None, [1,3]); M
[ 2| 3  5| 7 11]
[13|17 19|23 29]
[31|37 41|43 47]
[53|59 61|67 71]
[73|79 83|89 97]
```

Degenerate cases work too.

```
sage: M.subdivide([2,5], [0,1,3]); M
[| 2| 3  5| 7 11]
[|13|17 19|23 29]
[+--+-----+-----]
[|31|37 41|43 47]
[|53|59 61|67 71]
[|73|79 83|89 97]
[+--+-----+-----]
sage: M.subdivision(0,0)
[]
sage: M.subdivision(0,1)
[ 2]
[13]
sage: M.subdivide([2,2,3], [0,0,1,1]); M
[|| 2|| 3  5  7 11]
[||13||17 19 23 29]
```

```
[++---+-----]
[++---+-----]
[|31|37 41 43 47]
[++---+-----]
[|53|59 61 67 71]
[|73|79 83 89 97]
sage: M.subdivision(0,0)
[]
sage: M.subdivision(2,4)
[37 41 43 47]
```

AUTHORS:

•Robert Bradshaw (2007-06-14)

subdivision (*i, j*)

Returns an immutable copy of the (*i,j*)th submatrix of self, according to a previously set subdivision.

Before a subdivision is set, the only valid arguments are (0,0) which returns self.

EXAMPLE:

```
sage: M = matrix(3, 4, range(12))
sage: M.subdivide(1,2); M
[ 0  1| 2  3]
[-----]
[ 4  5| 6  7]
[ 8  9|10 11]
sage: M.subdivision(0,0)
[0 1]
sage: M.subdivision(0,1)
[2 3]
sage: M.subdivision(1,0)
[4 5]
[8 9]
```

It handles size-zero subdivisions as well.

```
sage: M = matrix(3, 4, range(12))
sage: M.subdivide([0], [0,2,2,4]); M
[+-----+-----+]
[| 0  1|| 2  3|]
[| 4  5|| 6  7|]
[| 8  9||10 11|]
sage: M.subdivision(0,0)
[]
sage: M.subdivision(1,1)
[0 1]
[4 5]
[8 9]
sage: M.subdivision(1,2)
[]
sage: M.subdivision(1,0)
[]
sage: M.subdivision(0,1)
[]
```

subdivision_entry (*i, j, x, y*)

Returns the *x,y* entry of the *i,j* submatrix of self.

EXAMPLES:

```

sage: M = matrix(5, 5, range(25))
sage: M.subdivide(3,3); M
[ 0  1  2| 3  4]
[ 5  6  7| 8  9]
[10 11 12|13 14]
[-----+-----]
[15 16 17|18 19]
[20 21 22|23 24]
sage: M.subdivision_entry(0,0,1,2)
7
sage: M.subdivision(0,0)[1,2]
7
sage: M.subdivision_entry(0,1,0,0)
3
sage: M.subdivision_entry(1,0,0,0)
15
sage: M.subdivision_entry(1,1,1,1)
24

```

Even though this entry exists in the matrix, the index is invalid for the submatrix.

```

sage: M.subdivision_entry(0,0,4,0)
Traceback (most recent call last):
...
IndexError: Submatrix 0,0 has no entry 4,0

```

subdivisions()

Returns the current subdivision of self.

EXAMPLES:

```

sage: M = matrix(5, 5, range(25))
sage: M.subdivisions()
([], [])
sage: M.subdivide(2,3)
sage: M.subdivisions()
([2], [3])
sage: N = M.parent()(1)
sage: N.subdivide(M.subdivisions()); N
[1 0 0|0 0]
[0 1 0|0 0]
[-----+-----]
[0 0 1|0 0]
[0 0 0|1 0]
[0 0 0|0 1]

```

subs (in_dict=None, **kws)

EXAMPLES:

```

sage: var('a,b,d,e')
(a, b, d, e)
sage: m = matrix([[a,b], [d,e]])
sage: m.substitute(a=1)
[1 b]
[d e]
sage: m.subs(a=b, b=d)
[b d]
[d e]

```

symplectic_form()

Find a symplectic form for self if self is an anti-symmetric, alternating matrix defined over a field.

Returns a pair (F, C) such that the rows of C form a symplectic basis for self and $F = C * self * C.transpose()$.

Raises a ValueError if not over a field, or self is not anti-symmetric, or self is not alternating.

Anti-symmetric means that $M = -M^t$. Alternating means that the diagonal of M is identically zero.

A symplectic basis is a basis of the form $e_1, \dots, e_j, f_1, \dots, f_j, z_1, \dots, z_k$ such that

- $z_i M v^t = 0$ for all vectors v
- $e_i M e_j^t = 0$ for all i, j
- $f_i M f_j^t = 0$ for all i, j
- $e_i M f_i^t = 1$ for all i
- $e_i M f_j^t = 0$ for all i not equal j .

See the example for a pictorial description of such a basis.

EXAMPLES:

```
sage: E = matrix(QQ, 8, 8, [0, -1/2, -2, 1/2, 2, 0, -2, 1, 1/2, 0, -1, -3, 0, 2, 5/2, -3, 2,
[ 0 -1/2  -2  1/2   2   0  -2   1]
[ 1/2  0   -1  -3    0   2  5/2  -3]
[ 2    1    0  3/2  -1    0  -1  -2]
[-1/2  3 -3/2    0   1  3/2 -1/2 -1/2]
[ -2    0    1  -1    0    0    1  -1]
[  0   -2    0 -3/2    0    0  1/2  -2]
[  2 -5/2    1  1/2   -1 -1/2    0  -1]
[ -1    3    2  1/2    1    2    1   0]
sage: F, C = E.symplectic_form(); F
[ 0  0  0  0  1  0  0  0]
[ 0  0  0  0  0  1  0  0]
[ 0  0  0  0  0  0  1  0]
[ 0  0  0  0  0  0  0  1]
[-1  0  0  0  0  0  0  0]
[ 0 -1  0  0  0  0  0  0]
[ 0  0 -1  0  0  0  0  0]
[ 0  0  0 -1  0  0  0  0]
sage: F == C * E * C.transpose()
True
```

tensor_product (A, subdivide=True)

Returns the tensor product of two matrices.

INPUT:

- A - a matrix
- subdivide - default: True - whether or not to return natural subdivisions with the matrix

OUTPUT:

Replace each element of self by a copy of A, but first create a scalar multiple of A by the element it replaces. So if self is an $m \times n$ matrix and A is a $p \times q$ matrix, then the tensor product is an $mp \times nq$ matrix. By default, the matrix will be subdivided into submatrices of size $p \times q$.

EXAMPLES:

```

sage: M1=Matrix(QQ, [[-1,0],[-1/2,-1]])
sage: M2=Matrix(ZZ, [[1,-1,2],[-2,4,8]])
sage: M1.tensor_product(M2)
[ -1   1   -2|  0   0   0]
[  2  -4  -8|  0   0   0]
[-----+-----]
[-1/2  1/2  -1| -1   1  -2]
[  1   -2  -4|  2  -4  -8]
sage: M2.tensor_product(M1)
[ -1   0|  1   0| -2   0]
[-1/2  -1| 1/2   1| -1  -2]
[-----+-----]
[  2   0| -4   0| -8   0]
[  1   2| -2  -4| -4  -8]

```

Subdivisions can be optionally suppressed.

```

sage: M1.tensor_product(M2, subdivide=False)
[ -1   1   -2   0   0   0]
[  2  -4  -8   0   0   0]
[-1/2  1/2  -1  -1   1  -2]
[  1   -2  -4   2  -4  -8]

```

Different base rings are handled sensibly.

```

sage: A = matrix(ZZ, 2, 3, range(6))
sage: B = matrix(FiniteField(23), 3, 4, range(12))
sage: C = matrix(FiniteField(29), 4, 5, range(20))
sage: D = A.tensor_product(B)
sage: D.parent()

```

Full MatrixSpace of 6 by 12 dense matrices over Finite Field of size 23

```
sage: E = C.tensor_product(B)
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: unsupported operand parent(s) for '*': 'Finite Field of size 29' and 'Full MatrixSpace of 6 by 12 dense matrices over Finite Field of size 23'
```

The input is checked to be sure it is a matrix.

```
sage: A = matrix(QQ, 2, range(4))
```

```
sage: A.tensor_product('junk')
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: tensor product requires a second matrix, not junk
```

trace()

Return the trace of self, which is the sum of the diagonal entries of self.

INPUT:

- self - a square matrix

OUTPUT: element of the base ring of self

EXAMPLES:

```
sage: a = matrix(3, range(9)); a
```

```
[0 1 2]
```

```
[3 4 5]
```

```
[6 7 8]
```

```
sage: a.trace()
```

```
12
```

```

sage: a = matrix({(1,1):10, (2,1):-3, (2,2):4/3}); a
[ 0  0  0]
[ 0 10  0]
[ 0 -3 4/3]
sage: a.trace()
34/3

```

trace_of_product (*other*)

Returns the trace of $\text{self} * \text{other}$ without computing the entire product.

EXAMPLES:

```

sage: M = random_matrix(ZZ, 10, 20)
sage: N = random_matrix(ZZ, 20, 10)
sage: M.trace_of_product(N)
-1629
sage: (M*N).trace()
-1629

```

visualize_structure (*filename=None, maxsize=512*)

Write a PNG image to ‘filename’ which visualizes self by putting black pixels in those positions which have nonzero entries.

White pixels are put at positions with zero entries. If ‘maxsize’ is given, then the maximal dimension in either x or y direction is set to ‘maxsize’ depending on which is bigger. If the image is scaled, the darkness of the pixel reflects how many of the represented entries are nonzero. So if e.g. one image pixel actually represents a 2x2 submatrix, the dot is darker the more of the four values are nonzero.

INPUT:

- *filename* - either a path or None in which case a filename in the current directory is chosen automatically (default:None)

maxsize - maximal dimension in either x or y direction of the resulting image. If None or a maxsize larger than $\max(\text{self.nrows()}, \text{self.ncols}())$ is given the image will have the same pixelsize as the matrix dimensions (default: 512)

EXAMPLE:

```

sage: M = random_matrix(CC, 4)
sage: M.visualize_structure()

```

weak_popov_form (*ascend=True*)

This function computes a weak Popov form of a matrix over a rational function field $k(x)$, for k a field.

INPUT:

- *ascend* - if True, rows of output matrix W are sorted so degree (= the maximum of the degrees of the elements in the row) increases monotonically, and otherwise degrees decrease.

OUTPUT:

A 3-tuple (W, N, d) consisting of:

1. W - a matrix over $k(x)$ giving a weak the Popov form of self
2. N - a matrix over $k[x]$ representing row operations used to transform *self* to W
3. d - degree of respective columns of W ; the degree of a column is the maximum of the degree of its elements

N is invertible over $k(x)$. These matrices satisfy the relation $N * \text{self} = W$.

EXAMPLES:

The routine expects matrices over the rational function field, but other examples below show how one can provide matrices over the ring of polynomials (whose quotient field is the rational function field).

```
sage: R.<t> = GF(3)['t']
sage: K = FractionField(R)
sage: M = matrix([[ (t-1)^2/t], [(t-1) ]])
sage: M.weak_popov_form()
(
[
0] [
t 2*t + 1]
[(2*t + 1)/t], [
1
2], [-Infinity, 0]
)
```

If *self* is an $n \times 1$ matrix with at least one non-zero entry, *W* has a single non-zero entry and that entry is a scalar multiple of the greatest-common-divisor of the entries of *self*.

```
sage: M1 = matrix([[t*(t-1)*(t+1)], [t*(t-2)*(t+2)], [t]])
sage: output1 = M1.weak_popov_form()
sage: output1
(
[0] [
1
0 2*t^2 + 1]
[0] [
0
1 2*t^2 + 1]
[t], [
0
0
1], [-Infinity, -Infinity, 1]
)
```

We check that the output is the same for a matrix *M* if its entries are rational functions instead of polynomials. We also check that the type of the output follows the documentation. See #9063

```
sage: M2 = M1.change_ring(K)
sage: output2 = M2.weak_popov_form()
sage: output1 == output2
True
sage: output1[0].base_ring() is K
True
sage: output2[0].base_ring() is K
True
sage: output1[1].base_ring() is R
True
sage: output2[1].base_ring() is R
True
```

The following is the first half of example 5 in [H] *except* that we have transposed *self*; [H] uses column operations and we use row.

```
sage: R.<t> = QQ['t']
sage: M = matrix([[t^3 - t, t^2 - 2], [0, t]]).transpose()
sage: M.weak_popov_form()
(
[
t
-t^2] [
1 -t]
[t^2 - 2
t], [
0
1], [2, 2]
)
```

The next example demonstrates what happens when *self* is a zero matrix.

```
sage: R.<t> = GF(5)['t']
sage: K = FractionField(R)
sage: M = matrix([[K(0), K(0)], [K(0), K(0)]])
sage: M.weak_popov_form()
(
[0 0] [1 0]
```

```
[0 0], [0 1], [-Infinity, -Infinity]
)
```

In the following example, *self* has more rows than columns.

```
sage: R.<t> = QQ['t']
sage: M = matrix([[t,t,t],[0,0,t]], ascend=False)
sage: M.weak_popov_form()
(
[t t t]  [1 0]
[0 0 t], [0 1], [1, 1]
)
```

The next example shows that *M* must be a matrix with coefficients in a rational function field $k(t)$.

```
sage: M = matrix([[1,0],[1,1]])
sage: M.weak_popov_form()
Traceback (most recent call last):
...
TypeError: the coefficients of M must lie in a univariate
polynomial ring
```

NOTES:

- For consistency with LLL and other algorithms in sage, we have opted for row operations; however, references such as [H] transpose and use column operations.
- There are multiple weak Popov forms of a matrix, so one may want to extend this code to produce a Popov form (see section 1.2 of [V]). The latter is canonical, but more work to produce.

REFERENCES:

wiedemann (*i*, *t*=0)

Application of Wiedemann's algorithm to the *i*-th standard basis vector.

INPUT:

- i* - an integer
- t* - an integer (default: 0) if *t* is nonzero, use only the first *t* linear recurrence relations.

IMPLEMENTATION: This is a toy implementation.

EXAMPLES:

```
sage: t = matrix(QQ, 3, range(9)); t
[0 1 2]
[3 4 5]
[6 7 8]
sage: t.wiedemann(0)
x^2 - 12*x - 18
sage: t.charpoly()
x^3 - 12*x^2 - 18*x
```

zigzag_form (*subdivide*=True, *transformation*=False)

Find a matrix in ZigZag form that is similar to *self*.

INPUT:

- self* - a square matrix with entries from an exact field.
- transformation* - default: False - if True return a change-of-basis matrix relating the matrix and its ZigZag form.

- `subdivide` - default: `True` - if `True` the ZigZag form matrix is subdivided according to the companion matrices described in the output section below.

OUTPUT:

A matrix in ZigZag form has blocks on the main diagonal that are companion matrices. The first companion matrix has ones just below the main diagonal. The last column has the negatives of coefficients of a monic polynomial, but not the leading one. Low degree monomials have their coefficients in the earlier rows. The second companion matrix is like the first only transposed. The third is like the first. The fourth is like the second. And so on.

These blocks on the main diagonal define blocks just off the diagonal. To the right of the first companion matrix, and above the second companion matrix is a block that is totally zero, except the entry of the first row and first column may be a one. Below the second block and to the left of the third block is a block that is totally zero, except the entry of the first row and first column may be one. This alternating pattern continues. It may now be apparent how this form gets its name. Any other entry of the matrix is zero. So this form is reminiscent of rational canonical form and is a good precursor to that form.

If `transformation` is `True`, then the output is a pair of matrices. The first is the form `Z` and the second is an invertible matrix `U` such that `U.inverse()*self*U` equals `Z`. In other words, the representation of `self` with respect to the columns of `U` will be `Z`.

If `subdivide` is `True` then the matrix returned as the form is partitioned according to the companion matrices and these may be manipulated by several different matrix methods.

For output that may be more useful as input to other routines, see the helper method `_zigzag_form()`.

Note: An effort has been made to optimize computation of the form, but no such work has been done for the computation of the transformation matrix, so for fastest results do not request the transformation matrix.

ALGORITHM:

ZigZag form, and its computation, are due to Arne Storjohann and are described in [STORJOHANN-THESIS] and [STORJOHANN-ISACC98], where the former is more representative of the code here.

EXAMPLES:

Two examples that illustrate ZigZag form well. Notice that this is *not* a canonical form. The two matrices below are similar, since they have equal Jordan canonical forms, yet their ZigZag forms are quite different. In other words, while the computation of the form is deterministic, the final result, when viewed as a property of a linear transformation, is dependent on the basis used for the matrix representation.

```
sage: A = matrix(QQ, [[-68, 69, -27, -11, -65, 9, -181, -32],
...                  [-52, 52, -27, -8, -52, -16, -133, -14],
...                  [ 92, -97, 47, 14, 90, 32, 241, 18],
...                  [139, -144, 60, 18, 148, -10, 362, 77],
...                  [ 40, -41, 12, 6, 45, -24, 105, 42],
...                  [-46, 48, -20, -7, -47, 0, -122, -22],
...                  [-26, 27, -13, -4, -29, -6, -66, -14],
...                  [-33, 34, -13, -5, -35, 7, -87, -23]])
sage: Z, U = A.zigzag_form(transformation=True)
sage: Z
[ 0  0  0  40| 1  0| 0  0]
[ 1  0  0  52| 0  0| 0  0]
[ 0  1  0  18| 0  0| 0  0]
[ 0  0  1 -11| 0  0| 0  0]
[-----+-----+-----]
[ 0  0  0  0| 0  1| 0  0]
[ 0  0  0  0|-25 10| 0  0]
```

```

[-----+-----+-----]
[  0   0   0   0|  1   0|  0  -4]
[  0   0   0   0|  0   0|  1  -4]
sage: U.inverse()*A*U == Z
True

sage: B = matrix(QQ, [[ 16,  69, -13,  2, -52, 143,  90, -3],
...                  [ 26,  54,  6, -5, -28,  73,  73, -48],
...                  [-16, -79, 12, -10,  64, -142, -115, 41],
...                  [ 27, -7, 21, -33,  39, -20, -42, 43],
...                  [  8, -75, 34, -32,  86, -156, -130, 42],
...                  [  2, -17,  7, -8,  20, -33, -31, 16],
...                  [-24, -80,  7, -3,  56, -136, -112, 42],
...                  [-6, -19,  0, -1, 13, -28, -27, 15]])
sage: Z, U = B.zigzag_form(transformation=True)
sage: Z
[  0   0   0   0   0 1000|  0|  0]
[  1   0   0   0   0  900|  0|  0]
[  0   1   0   0   0 -30|  0|  0]
[  0   0   1   0   0 -153|  0|  0]
[  0   0   0   1   0   3|  0|  0]
[  0   0   0   0   1   9|  0|  0]
[-----+-----+-----]
[  0   0   0   0   0   0| -2|  0]
[-----+-----+-----]
[  0   0   0   0   0   0|  1| -2]
sage: U.inverse()*B*U == Z
True

sage: A.jordan_form() == B.jordan_form()
True

```

Two more examples, illustrating the two extremes of the zig-zag nature of this form. The first has a one in each of the off-diagonal blocks, the second has all zeros in each off-diagonal block. Notice again that the two matrices are similar, since their Jordan canonical forms are equal.

```

sage: C = matrix(QQ, [[2,  31, -10, -9, -125, 13,  62, -12],
...                  [0,  48, -16, -16, -188, 20,  92, -16],
...                  [0,  9, -1,  2, -33,  5,  18,  0],
...                  [0, 15, -5,  0, -59,  7,  30, -4],
...                  [0, -21,  7,  2,  84, -10, -42,  5],
...                  [0, -42, 14,  8, 167, -17, -84, 13],
...                  [0, -50, 17, 10, 199, -23, -98, 14],
...                  [0, 15, -5, -2, -59,  7,  30, -2]])
sage: Z, U = C.zigzag_form(transformation=True)
sage: Z
[2|1|0|0|0|0|0|0|0]
[-+-+--+--+--+--+]
[0|2|0|0|0|0|0|0|0]
[-+-+--+--+--+--+]
[0|1|2|1|0|0|0|0|0]
[-+-+--+--+--+--+]
[0|0|0|2|0|0|0|0|0]
[-+-+--+--+--+--+]
[0|0|0|1|2|1|0|0|0]
[-+-+--+--+--+--+]
[0|0|0|0|0|2|0|0|0]
[-+-+--+--+--+--+]

```

```

[0|0|0|0|0|1|2|1]
[-+-+--+--+--+--+]
[0|0|0|0|0|0|0|2]
sage: U.inverse()*C*U == Z
True

sage: D = matrix(QQ, [[ -4,  3,  7,  2, -4,  5,  7, -3],
...                    [ -6,  5,  7,  2, -4,  5,  7, -3],
...                    [ 21, -12, 89, 25,  8, 27, 98, -95],
...                    [ -9,  5, -44, -11, -3, -13, -48, 47],
...                    [ 23, -13, 74, 21, 12, 22, 85, -84],
...                    [ 31, -18, 135, 38, 12, 47, 155, -147],
...                    [-33, 19, -138, -39, -13, -45, -156, 151],
...                    [ -7,  4, -29, -8, -3, -10, -34, 34]])
sage: Z, U = D.zigzag_form(transformation=True)
sage: Z
[ 0 -4| 0  0| 0  0| 0  0]
[ 1  4| 0  0| 0  0| 0  0]
[-----+-----+-----+-----]
[ 0  0| 0  1| 0  0| 0  0]
[ 0  0|-4  4| 0  0| 0  0]
[-----+-----+-----+-----]
[ 0  0| 0  0| 0 -4| 0  0]
[ 0  0| 0  0| 1  4| 0  0]
[-----+-----+-----+-----]
[ 0  0| 0  0| 0  0| 0  1]
[ 0  0| 0  0| 0 -4| 4]
sage: U.inverse()*D*U == Z
True

sage: C.jordan_form() == D.jordan_form()
True

```

ZigZag form is achieved entirely with the operations of the field, so while the eigenvalues may lie outside the field, this does not impede the computation of the form.

```

sage: F.<a> = GF(5^4)
sage: A = matrix(F, [[ a,      0,  0, a + 3],
...                  [ 0, a^2 + 1, 0,  0],
...                  [ 0,      0, a^3,  0],
...                  [a^2 + 4,  0,  0, a + 2]])
sage: A.zigzag_form()
[          0 a^3 + 2*a^2 + 2*a + 2|          0|          0]
[          1          2*a + 2|          0|          0]
[-----+-----+-----+-----]
[          0          0|          a^3|          0]
[-----+-----+-----+-----]
[          0          0|          0|          a^2 + 1]
sage: A.eigenvalues()
Traceback (most recent call last):
...
NotImplementedError: algebraic closures of finite fields are only implemented for prime fields

```

Subdivisions are optional.

```

sage: F.<a> = GF(5^4)
sage: A = matrix(F, [[ a,      0,  0, a + 3],
...                  [ 0, a^2 + 1, 0,  0],
...                  [ 0,      0, a^3,  0],

```

```

...          [a^2 + 4 ,      0,      0, a + 2]])
sage: A.zigzag_form(subdivide=False)
[      0 a^3 + 2*a^2 + 2*a + 2      0      0]
[      1      2*a + 2      0      0]
[      0      0      a^3      0]
[      0      0      0      a^2 + 1]

```

TESTS:

```

sage: A = matrix(QQ, 2, 3, range(6))
sage: A.zigzag_form()
Traceback (most recent call last):
...
TypeError: matrix must be square, not 2 x 3

sage: A = matrix(Integers(6), 2, 2, range(4))
sage: A.zigzag_form()
Traceback (most recent call last):
...
TypeError: matrix entries must come from an exact field, not Ring of integers modulo 6

sage: A = matrix(RDF, 2, 2, range(4))
sage: A.zigzag_form()
Traceback (most recent call last):
...
TypeError: matrix entries must come from an exact field, not Real Double Field

sage: A = matrix(QQ, 2, range(4))
sage: A.zigzag_form(transformation='junk')
Traceback (most recent call last):
...
ValueError: 'transformation' keyword must be True or False, not junk

sage: A = matrix(QQ, 2, range(4))
sage: A.zigzag_form(subdivide='garbage')
Traceback (most recent call last):
...
ValueError: 'subdivide' keyword must be True or False, not garbage

```

Citations**AUTHOR:**

•Rob Beezer (2011-06-09)

`sage.matrix.matrix2.cmp_pivots(x, y)`

Compare two sequences of pivot columns.

- If x is shorter than y , return -1 , i.e., $x < y$, “not as good”.
- If x is longer than y , $x > y$, “better”.
- If the length is the same then x is better, i.e., $x > y$ if the entries of x are correspondingly \geq those of y with one being greater.

`sage.matrix.matrix2.decomp_seq(v)`

This function is used internally by the decomposition matrix method. It takes a list of tuples and produces a sequence that is correctly sorted and prints with carriage returns.

EXAMPLES:

```
sage: from sage.matrix.matrix2 import decomp_seq
sage: V = [(QQ^3, 2), (QQ^2, 1)]
sage: decomp_seq(V)
[
  (Vector space of dimension 2 over Rational Field, 1),
  (Vector space of dimension 3 over Rational Field, 2)
]
```


GENERIC ASYMPTOTICALLY FAST STRASSEN ALGORITHMS

Generic Asymptotically Fast Strassen Algorithms

Sage implements asymptotically fast echelon form and matrix multiplication algorithms.

class `sage.matrix.strassen.int_range` (*indices=None, range=None*)
Represent a list of integers as a list of integer intervals.

Note: Repetitions are not considered.

Useful class for dealing with pivots in the strassen echelon, could have much more general application

INPUT:

It can be one of the following:

- *indices* - integer, start of the unique interval
- *range* - integer, length of the unique interval

OR

- *indices* - list of integers, the integers to wrap into intervals

OR

- *indices* - None (default), shortcut for an empty list

OUTPUT:

An instance of `int_range`, i.e. a list of pairs (*start*, *length*).

EXAMPLES:

From a pair of integers:

```
sage: from sage.matrix.strassen import int_range
sage: int_range(2, 4)
[(2, 4)]
```

Default:

```
sage: int_range()
[]
```

From a list of integers:

```
sage: int_range([1,2,3,4])
[(1, 4)]
sage: int_range([1,2,3,4,6,7,8])
[(1, 4), (6, 3)]
sage: int_range([1,2,3,4,100,101,102])
[(1, 4), (100, 3)]
sage: int_range([1,1000,2,101,3,4,100,102])
[(1, 4), (100, 3), (1000, 1)]
```

Repetitions are not considered:

```
sage: int_range([1,2,3])
[(1, 3)]
sage: int_range([1,1,1,1,2,2,2,3])
[(1, 3)]
```

AUTHORS:

- Robert Bradshaw

intervals()

Return the list of intervals.

OUTPUT:

A list of pairs of integers.

EXAMPLES:

```
sage: from sage.matrix.strassen import int_range
sage: I = int_range([4,5,6,20,21,22,23])
sage: I.intervals()
[(4, 3), (20, 4)]
sage: type(I.intervals())
<type 'list'>
```

to_list()

Return the (sorted) list of integers represented by this object.

OUTPUT:

A list of integers.

EXAMPLES:

```
sage: from sage.matrix.strassen import int_range
sage: I = int_range([6,20,21,4,5,22,23])
sage: I.to_list()
[4, 5, 6, 20, 21, 22, 23]

sage: I = int_range(34, 9)
sage: I.to_list()
[34, 35, 36, 37, 38, 39, 40, 41, 42]
```

Repetitions are not considered:

```
sage: I = int_range([1,1,1,1,2,2,2,3])
sage: I.to_list()
[1, 2, 3]
```

`sage.matrix.strassen.strassen_echelon(A, cutoff)`

Compute echelon form, in place. Internal function, call with `M.echelonize(algorithm="strassen")` Based on

INPUT:

- **A** - matrix window
- **cutoff** - size at which algorithm reverts to naive Gaussian elimination and multiplication must be at least 1.

EXAMPLE:

```
sage: A = matrix(QQ, 7, [5, 0, 0, 0, 0, 0, -1, 0, 0, 1, 0, 0, 0, 0, 0, -1, 3, 1, 0, -1, 0, 0, -1]
sage: B = A.__copy__(); B._echelon_strassen(1); B
[ 1  0  0  0  0  0  0]
[ 0  1  0 -1  0  1  0]
[ 0  0  1  0  0  0  0]
[ 0  0  0  0  1  0  0]
[ 0  0  0  0  0  0  1]
[ 0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0]
sage: C = A.__copy__(); C._echelon_strassen(2); C == B
True
sage: C = A.__copy__(); C._echelon_strassen(4); C == B
True

sage: n = 32; A = matrix(Integers(389), n, range(n^2))
sage: B = A.__copy__(); B._echelon_in_place_classical()
sage: C = A.__copy__(); C._echelon_strassen(2)
sage: B == C
True
```

```
sage: A = matrix(Integers(7), 4, 4, [1,2,0,3,0,0,1,0,0,1,0,0,0,0,0,1])
sage: B = A.__copy__(); B._echelon_in_place_classical()
sage: C = A.__copy__(); C._echelon_strassen(2)
sage: B == C
True
```

```
sage: A = matrix(Integers(7), 4, 4, [1,0,5,0,2,0,3,6,5,1,2,6,4,6,1,1])
sage: B = A.__copy__(); B._echelon_in_place_classical()
sage: C = A.__copy__(); C._echelon_strassen(2)      #indirect doctest
sage: B == C
True
```

•Robert Bradshaw

```
sage.matrix.strassen.strassen_window_multiply(C,A,B,cutoff)
```

Multiplies the submatrices specified by A and B, places result in C. Assumes that A and B have compatible dimensions to be multiplied, and that C is the correct size to receive the product, and that they are all defined over the same ring.

Uses strassen multiplication at high levels and then uses MatrixWindow methods at low levels. EXAMPLES: The following matrix dimensions are chosen especially to exercise the eight possible parity combinations that could occur while subdividing the matrix in the strassen recursion. The base case in both cases will be a (4x5) matrix times a (5x6) matrix.

```
sage: A = MatrixSpace(Integers(2^65), 64, 83).random_element()
sage: B = MatrixSpace(Integers(2^65), 83, 101).random_element()
sage: A._multiply_classical(B) == A._multiply_strassen(B, 3) #indirect doctest
True
```

AUTHORS:

- David Harvey
- Simon King (2011-07): Improve memory efficiency; trac ticket #11610

sage.matrix.strassen.**test**(*n, m, R, c=2*)

INPUT:

- n* - integer
- m* - integer
- R* - ring
- c* - integer (optional, default:2)

EXAMPLES:

```
sage: from sage.matrix.strassen import test
sage: for n in range(5): print n, test(2*n,n,Frac(QQ['x']),2)
0 True
1 True
2 True
3 True
4 True
```

MINIMAL POLYNOMIALS OF LINEAR RECURRENCE SEQUENCES

AUTHORS:

- William Stein

`sage.matrix.berlekamp_massey.berlekamp_massey(a)`

Use the Berlekamp-Massey algorithm to find the minimal polynomial of a linearly recurrence sequence a .

The minimal polynomial of a linear recurrence $\{a_r\}$ is by definition the unique monic polynomial g , such that if $\{a_r\}$ satisfies a linear recurrence $a_{j+k} + b_{j-1}a_{j-1+k} + \cdots + b_0a_k = 0$ (for all $k \geq 0$), then g divides the polynomial $x^j + \sum_{i=0}^{j-1} b_i x^i$.

INPUT:

- a - a list of even length of elements of a field (or domain)

OUTPUT:

- Polynomial - the minimal polynomial of the sequence (as a polynomial over the field in which the entries of a live)

EXAMPLES:

sage: `berlekamp_massey([1,2,1,2,1,2])`

$x^2 - 1$

sage: `berlekamp_massey(GF(7)(1),19,1,19)`

$x^2 + 6$

sage: `berlekamp_massey([2,2,1,2,1,191,393,132])`

$x^4 - 36727/11711x^3 + 34213/5019x^2 + 7024942/35133x - 335813/1673$

sage: `berlekamp_massey(prime_range(2,38))`

$x^6 - 14/9x^5 - 7/9x^4 + 157/54x^3 - 25/27x^2 - 73/18x + 37/9$

BASE CLASS FOR DENSE MATRICES

Base class for dense matrices

TESTS:

```
sage: R.<a,b> = QQ[]
sage: m = matrix(R,2,[0,a,b,b^2])
sage: TestSuite(m).run()
```

class sage.matrix.matrix_dense.**Matrix_dense**

Bases: sage.matrix.matrix.Matrix

The initialization routine of the `Matrix` base class ensures that it sets the attributes `self._parent`, `self._base_ring`, `self._nrows`, `self._ncols`. It sets the latter ones by accessing the relevant information on `parent`, which is often slower than what a more specific subclass can do.

Subclasses of `Matrix` can safely skip calling `Matrix.__init__` provided they take care of initializing these attributes themselves.

The private attributes `self._is_immutable` and `self._cache` are implicitly initialized to valid values upon memory allocation.

EXAMPLES:

```
sage: import sage.matrix.matrix0
sage: A = sage.matrix.matrix0.Matrix(MatrixSpace(QQ,2))
sage: type(A)
<type 'sage.matrix.matrix0.Matrix'>
```

antitranspose()

Returns the antitranspose of `self`, without changing `self`.

EXAMPLES:

```
sage: A = matrix(2,3,range(6)); A
[0 1 2]
[3 4 5]
sage: A.antitranspose()
[5 2]
[4 1]
[3 0]
```

```
sage: A.subdivide(1,2); A
[0 1|2]
[---+-]
[3 4|5]
sage: A.antitranspose()
```

```
[5|2]
[-+-]
[4|1]
[3|0]
```

apply_map (*phi*, *R=None*, *sparse=False*)

Apply the given map *phi* (an arbitrary Python function or callable object) to this dense matrix. If *R* is not given, automatically determine the base ring of the resulting matrix.

INPUT: *sparse* – True to make the output a sparse matrix; default False

- *phi* - arbitrary Python function or callable object
- *R* - (optional) ring

OUTPUT: a matrix over *R*

EXAMPLES:

```
sage: m = matrix(ZZ, 3, range(9))
sage: k.<a> = GF(9)
sage: f = lambda x: k(x)
sage: n = m.apply_map(f); n
[0 1 2]
[0 1 2]
[0 1 2]
sage: n.parent()
Full MatrixSpace of 3 by 3 dense matrices over Finite Field in a of size 3^2
```

In this example, we explicitly specify the codomain.

```
sage: s = GF(3)
sage: f = lambda x: s(x)
sage: n = m.apply_map(f, k); n
[0 1 2]
[0 1 2]
[0 1 2]
sage: n.parent()
Full MatrixSpace of 3 by 3 dense matrices over Finite Field in a of size 3^2
```

If *self* is subdivided, the result will be as well:

```
sage: m = matrix(2, 2, srange(4))
sage: m.subdivide(None, 1); m
[0|1]
[2|3]
sage: m.apply_map(lambda x: x*x)
[0|1]
[4|9]
```

If the map sends most of the matrix to zero, then it may be useful to get the result as a sparse matrix.

```
sage: m = matrix(ZZ, 3, 3, range(1, 10))
sage: n = m.apply_map(lambda x: 1//x, sparse=True); n
[1 0 0]
[0 0 0]
[0 0 0]
sage: n.parent()
Full MatrixSpace of 3 by 3 sparse matrices over Integer Ring
```

TESTS:

```
sage: m = matrix([])
sage: m.apply_map(lambda x: x*x) == m
True

sage: m.apply_map(lambda x: x*x, sparse=True).parent()
Full MatrixSpace of 0 by 0 sparse matrices over Integer Ring
```

apply_morphism(*phi*)

Apply the morphism *phi* to the coefficients of this dense matrix.

The resulting matrix is over the codomain of *phi*.

INPUT:

- *phi* - a morphism, so *phi* is callable and *phi*.domain() and *phi*.codomain() are defined. The codomain must be a ring.

OUTPUT: a matrix over the codomain of *phi*

EXAMPLES:

```
sage: m = matrix(ZZ, 3, range(9))
sage: phi = ZZ.hom(GF(5))
sage: m.apply_morphism(phi)
[0 1 2]
[3 4 0]
[1 2 3]
sage: parent(m.apply_morphism(phi))
Full MatrixSpace of 3 by 3 dense matrices over Finite Field of size 5
```

We apply a morphism to a matrix over a polynomial ring:

```
sage: R.<x,y> = QQ[]
sage: m = matrix(2, [x,x^2 + y, 2/3*y^2-x, x]); m
[      x      x^2 + y]
[2/3*y^2 - x      x]
sage: phi = R.hom([y,x])
sage: m.apply_morphism(phi)
[      y      y^2 + x]
[2/3*x^2 - y      y]
```

transpose()

Returns the transpose of self, without changing self.

EXAMPLES: We create a matrix, compute its transpose, and note that the original matrix is not changed.

```
sage: M = MatrixSpace(QQ, 2)
sage: A = M([1,2,3,4])
sage: B = A.transpose()
sage: print B
[1 3]
[2 4]
sage: print A
[1 2]
[3 4]
```

.T is a convenient shortcut for the transpose:

```
sage: A.T
[1 3]
[2 4]
```

```
sage: A.subdivide(None, 1); A
[1|2]
[3|4]
sage: A.transpose()
[1 3]
[---]
[2 4]
```


BASE CLASS FOR SPARSE MATRICES

Base class for sparse matrices

```
class sage.matrix.matrix_sparse.Matrix_sparse
    Bases: sage.matrix.matrix.Matrix
```

The initialization routine of the `Matrix` base class ensures that it sets the attributes `self._parent`, `self._base_ring`, `self._nrows`, `self._ncols`. It sets the latter ones by accessing the relevant information on `parent`, which is often slower than what a more specific subclass can do.

Subclasses of `Matrix` can safely skip calling `Matrix.__init__` provided they take care of initializing these attributes themselves.

The private attributes `self._is_immutable` and `self._cache` are implicitly initialized to valid values upon memory allocation.

EXAMPLES:

```
sage: import sage.matrix.matrix0
sage: A = sage.matrix.matrix0.Matrix(MatrixSpace(QQ, 2))
sage: type(A)
<type 'sage.matrix.matrix0.Matrix'>
```

antitranspose()

apply_map (*phi*, *R=None*, *sparse=True*)

Apply the given map *phi* (an arbitrary Python function or callable object) to this matrix. If *R* is not given, automatically determine the base ring of the resulting matrix.

INPUT: *sparse* – False to make the output a dense matrix; default True

- *phi* - arbitrary Python function or callable object
- *R* - (optional) ring

OUTPUT: a matrix over *R*

EXAMPLES:

```
sage: m = matrix(ZZ, 10000, {(1,2): 17}, sparse=True)
sage: k.<a> = GF(9)
sage: f = lambda x: k(x)
sage: n = m.apply_map(f)
sage: n.parent()
Full MatrixSpace of 10000 by 10000 sparse matrices over Finite Field in a of size 3^2
sage: n[1,2]
2
```

An example where the codomain is explicitly specified.

```
sage: n = m.apply_map(lambda x:x%3, GF(3))
sage: n.parent()
Full MatrixSpace of 10000 by 10000 sparse matrices over Finite Field of size 3
sage: n[1,2]
2
```

If we didn't specify the codomain, the resulting matrix in the above case ends up over ZZ again:

```
sage: n = m.apply_map(lambda x:x%3)
sage: n.parent()
Full MatrixSpace of 10000 by 10000 sparse matrices over Integer Ring
sage: n[1,2]
2
```

If self is subdivided, the result will be as well:

```
sage: m = matrix(2, 2, [0, 0, 3, 0])
sage: m.subdivide(None, 1); m
[0|0]
[3|0]
sage: m.apply_map(lambda x: x*x)
[0|0]
[9|0]
```

If the map sends zero to a non-zero value, then it may be useful to get the result as a dense matrix.

```
sage: m = matrix(ZZ, 3, 3, [0] * 7 + [1,2], sparse=True); m
[0 0 0]
[0 0 0]
[0 1 2]
sage: parent(m)
Full MatrixSpace of 3 by 3 sparse matrices over Integer Ring
sage: n = m.apply_map(lambda x: x+polygen(QQ), sparse=False); n
[  x      x      x]
[  x      x      x]
[  x x + 1 x + 2]
sage: parent(n)
Full MatrixSpace of 3 by 3 dense matrices over Univariate Polynomial Ring in x over Rational
```

TESTS:

```
sage: m = matrix([], sparse=True)
sage: m.apply_map(lambda x: x*x) == m
True

sage: m.apply_map(lambda x: x*x, sparse=False).parent()
Full MatrixSpace of 0 by 0 dense matrices over Integer Ring
```

Check that we don't unnecessarily apply phi to 0 in the sparse case:

```
sage: m = matrix(QQ, 2, 2, range(1, 5), sparse=True)
sage: m.apply_map(lambda x: 1/x)
[ 1 1/2]
[1/3 1/4]
```

Test subdivisions when phi maps 0 to non-zero:

```
sage: m = matrix(2, 2, [0, 0, 3, 0])
sage: m.subdivide(None, 1); m
[0|0]
```

```
[3|0]
sage: m.apply_map(lambda x: x+1)
[1|1]
[4|1]
```

apply_morphism (*phi*)

Apply the morphism *phi* to the coefficients of this sparse matrix.

The resulting matrix is over the codomain of *phi*.

INPUT:

- *phi* - a morphism, so *phi* is callable and *phi*.domain() and *phi*.codomain() are defined. The codomain must be a ring.

OUTPUT: a matrix over the codomain of *phi*

EXAMPLES:

```
sage: m = matrix(ZZ, 3, range(9), sparse=True)
sage: phi = ZZ.hom(GF(5))
sage: m.apply_morphism(phi)
[0 1 2]
[3 4 0]
[1 2 3]
sage: m.apply_morphism(phi).parent()
Full MatrixSpace of 3 by 3 sparse matrices over Finite Field of size 5
```

augment (*right*, *subdivide=False*)

Return the augmented matrix of the form:

```
[self | right].
```

EXAMPLES:

```
sage: M = MatrixSpace(QQ, 2, 2, sparse=True)
sage: A = M([1,2, 3,4])
sage: A
[1 2]
[3 4]
sage: N = MatrixSpace(QQ, 2, 1, sparse=True)
sage: B = N([9,8])
sage: B
[9]
[8]
sage: A.augment(B)
[1 2 9]
[3 4 8]
sage: B.augment(A)
[9 1 2]
[8 3 4]
```

A vector may be augmented to a matrix.

```
sage: A = matrix(QQ, 3, 4, range(12), sparse=True)
sage: v = vector(QQ, 3, range(3), sparse=True)
sage: A.augment(v)
[ 0 1 2 3 0]
[ 4 5 6 7 1]
[ 8 9 10 11 2]
```

The `subdivide` option will add a natural subdivision between `self` and `right`. For more details about how subdivisions are managed when augmenting, see `sage.matrix.matrix1.Matrix.augment()`.

```
sage: A = matrix(QQ, 3, 5, range(15), sparse=True)
sage: B = matrix(QQ, 3, 3, range(9), sparse=True)
sage: A.augment(B, subdivide=True)
[ 0  1  2  3  4 | 0  1  2]
[ 5  6  7  8  9 | 3  4  5]
[10 11 12 13 14 | 6  7  8]
```

TESTS:

Verify that Trac #12689 is fixed:

```
sage: A = identity_matrix(QQ, 2, sparse=True)
sage: B = identity_matrix(ZZ, 2, sparse=True)
sage: A.augment(B)
[1 0 1 0]
[0 1 0 1]
```

change_ring (*ring*)

Return the matrix obtained by coercing the entries of this matrix into the given ring.

Always returns a copy (unless `self` is immutable, in which case returns `self`).

EXAMPLES:

```
sage: A = matrix(QQ['x,y'], 2, [0,-1,2*x,-2], sparse=True); A
[ 0 -1]
[2*x -2]
sage: A.change_ring(QQ['x,y,z'])
[ 0 -1]
[2*x -2]
```

Subdivisions are preserved when changing rings:

```
sage: A.subdivide([2], []); A
[ 0 -1]
[2*x -2]
[-----]
sage: A.change_ring(RR['x,y'])
[ 0 -1.000000000000000]
[2.000000000000000*x -2.000000000000000]
[-----]
```

charpoly (*var='x', **kws*)

Return the characteristic polynomial of this matrix.

Note - the generic sparse charpoly implementation in Sage is to just compute the charpoly of the corresponding dense matrix, so this could use a lot of memory. In particular, for this matrix, the charpoly will be computed using a dense algorithm.

EXAMPLES:

```
sage: A = matrix(ZZ, 4, range(16), sparse=True)
sage: A.charpoly()
x^4 - 30*x^3 - 80*x^2
sage: A.charpoly('y')
y^4 - 30*y^3 - 80*y^2
sage: A.charpoly()
x^4 - 30*x^3 - 80*x^2
```

density()

Return the density of the matrix.

By density we understand the ratio of the number of nonzero positions and the `self.nrows() * self.ncols()`, i.e. the number of possible nonzero positions.

EXAMPLES:

```
sage: a = matrix([[[]],[[]],[[]],[[]], sparse=True); a.density()
0
sage: a = matrix(5000,5000,{(1,2): 1}); a.density()
1/25000000
```

determinant (kws)**

Return the determinant of this matrix.

Note: the generic sparse determinant implementation in Sage is to just compute the determinant of the corresponding dense matrix, so this could use a lot of memory. In particular, for this matrix, the determinant will be computed using a dense algorithm.

EXAMPLES:

```
sage: A = matrix(ZZ, 4, range(16), sparse=True)
sage: B = A + identity_matrix(ZZ, 4, sparse=True)
sage: B.det()
-49
```

matrix_from_rows_and_columns (rows, columns)

Return the matrix constructed from self from the given rows and columns.

EXAMPLES:

```
sage: M = MatrixSpace(Integers(8),3,3, sparse=True)
sage: A = M(range(9)); A
[0 1 2]
[3 4 5]
[6 7 0]
sage: A.matrix_from_rows_and_columns([1], [0,2])
[3 5]
sage: A.matrix_from_rows_and_columns([1,2], [1,2])
[4 5]
[7 0]
```

Note that row and column indices can be reordered or repeated:

```
sage: A.matrix_from_rows_and_columns([2,1], [2,1])
[0 7]
[5 4]
```

For example here we take from row 1 columns 2 then 0 twice, and do this 3 times.

```
sage: A.matrix_from_rows_and_columns([1,1,1], [2,0,0])
[5 3 3]
[5 3 3]
[5 3 3]
```

We can efficiently extract large submatrices:

```
sage: A = random_matrix(ZZ, 100000, density=.00005, sparse=True) # long time (4s on sage.ma
sage: B = A[50000:,:50000] # long time
sage: len(B.nonzero_positions()) # long time
```

```
17550          # 32-bit
125449          # 64-bit
```

We must pass in a list of indices:

```
sage: A=random_matrix(ZZ,100,density=.02,sparse=True)
sage: A.matrix_from_rows_and_columns(1,[2,3])
Traceback (most recent call last):
...
TypeError: rows must be a list of integers
sage: A.matrix_from_rows_and_columns([1,2],3)
Traceback (most recent call last):
...
TypeError: columns must be a list of integers
```

AUTHORS:

- Jaap Spies (2006-02-18)
- Didier Deshommes: some Pyrex speedups implemented
- Jason Grout: sparse matrix optimizations

stack (*bottom*, *subdivide=False*)

Stack self on top of bottom.

```
[ self ] [ bottom ]
```

EXAMPLES:

```
sage: M = Matrix(QQ, 2, 3, range(6), sparse=True)
sage: N = Matrix(QQ, 1, 3, [10,11,12], sparse=True)
sage: M.stack(N)
[ 0  1  2]
[ 3  4  5]
[10 11 12]
```

A vector may be stacked below a matrix.

```
sage: A = matrix(QQ, 2, 5, range(10), sparse=True)
sage: v = vector(QQ, 5, range(5), sparse=True)
sage: A.stack(v)
[0 1 2 3 4]
[5 6 7 8 9]
[0 1 2 3 4]
```

The subdivide option will add a natural subdivision between self and other. For more details about how subdivisions are managed when stacking, see `sage.matrix.matrix1.Matrix.stack()`.

```
sage: A = matrix(ZZ, 3, 4, range(12), sparse=True)
sage: B = matrix(ZZ, 2, 4, range(8), sparse=True)
sage: A.stack(B, subdivide=True)
[ 0  1  2  3]
[ 4  5  6  7]
[ 8  9 10 11]
[-----]
[ 0  1  2  3]
[ 4  5  6  7]
```

transpose ()

Returns the transpose of self, without changing self.

EXAMPLES: We create a matrix, compute its transpose, and note that the original matrix is not changed.

```
sage: M = MatrixSpace(QQ, 2, sparse=True)
```

```
sage: A = M([1,2,3,4])
```

```
sage: B = A.transpose()
```

```
sage: print B
```

```
[1 3]
```

```
[2 4]
```

```
sage: print A
```

```
[1 2]
```

```
[3 4]
```

.T is a convenient shortcut for the transpose:

```
sage: A.T
```

```
[1 3]
```

```
[2 4]
```


DENSE MATRICES OVER A GENERAL RING

Dense Matrices over a general ring

```
class sage.matrix.matrix_generic_dense.Matrix_generic_dense  
    Bases: sage.matrix.matrix_dense.Matrix_dense
```

The `Matrix_generic_dense` class derives from `Matrix`, and defines functionality for dense matrices over any base ring. Matrices are represented by a list of elements in the base ring, and element access operations are implemented in this class.

EXAMPLES:

```
sage: A = random_matrix(Integers(25) ['x'], 2); A  
[      x^2 + 12*x + 2   4*x^2 + 13*x + 8]  
[ 22*x^2 + 2*x + 17 19*x^2 + 22*x + 14]  
sage: type(A)  
<type 'sage.matrix.matrix_generic_dense.Matrix_generic_dense'>  
sage: TestSuite(A).run()
```


SPARSE MATRICES OVER A GENERAL RING

Sparse Matrices over a general ring

EXAMPLES:

```
sage: R.<x> = PolynomialRing(QQ)
sage: M = MatrixSpace(QQ['x'], 2, 3, sparse=True); M
Full MatrixSpace of 2 by 3 sparse matrices over Univariate Polynomial Ring in x over Rational Field
sage: a = M(range(6)); a
[0 1 2]
[3 4 5]
sage: b = M([x^n for n in range(6)]); b
[ 1   x x^2]
[x^3 x^4 x^5]
sage: a * b.transpose()
[      2*x^2 + x      2*x^5 + x^4]
[ 5*x^2 + 4*x + 3 5*x^5 + 4*x^4 + 3*x^3]
sage: pari(a)*pari(b.transpose())
[2*x^2 + x, 2*x^5 + x^4; 5*x^2 + 4*x + 3, 5*x^5 + 4*x^4 + 3*x^3]
sage: c = copy(b); c
[ 1   x x^2]
[x^3 x^4 x^5]
sage: c[0,0] = 5; c
[ 5   x x^2]
[x^3 x^4 x^5]
sage: b[0,0]
1
sage: c.dict()
{(0, 1): x, (1, 2): x^5, (0, 0): 5, (1, 0): x^3, (0, 2): x^2, (1, 1): x^4}
sage: c.list()
[5, x, x^2, x^3, x^4, x^5]
sage: c.rows()
[(5, x, x^2), (x^3, x^4, x^5)]
sage: TestSuite(c).run()
sage: d = c.change_ring(CC['x']); d
[5.000000000000000      x      x^2]
[      x^3      x^4      x^5]
sage: latex(c)
\left(\begin{array}{rrr}
5 & x & x^2 \\
x^3 & x^4 & x^5
\end{array}\right)
```

```
sage: c.sparse_rows()
[(5, x, x^2), (x^3, x^4, x^5)]
sage: d = c.dense_matrix(); d
[ 5   x x^2]
[x^3 x^4 x^5]
sage: parent(d)
Full MatrixSpace of 2 by 3 dense matrices over Univariate Polynomial Ring in x over Rational Field
sage: c.sparse_matrix() is c
True
sage: c.is_sparse()
True
```

class `sage.matrix.matrix_generic_sparse.Matrix_generic_sparse`
Bases: `sage.matrix.matrix_sparse.Matrix_sparse`

The `Matrix_generic_sparse` class derives from `Matrix`, and defines functionality for sparse matrices over any base ring. A generic sparse matrix is represented using a dictionary with keys pairs (i, j) and values in the base ring.

The values of the dictionary must never be zero.

`sage.matrix.matrix_generic_sparse.Matrix_sparse_from_rows(X)`
INPUT:

- X - nonempty list of `SparseVector` rows

OUTPUT: `Sparse_matrix` with those rows.

EXAMPLES:

```
sage: V = VectorSpace(QQ, 20, sparse=True)
sage: v = V(0)
sage: v[9] = 4
sage: from sage.matrix.matrix_generic_sparse import Matrix_sparse_from_rows
sage: Matrix_sparse_from_rows([v])
[0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
sage: Matrix_sparse_from_rows([v, v, v, V(0)])
[0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 4 0 0 0 0 0 0 0 0 0 0]
[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]
```

DENSE MATRICES OVER $\mathbf{Z}/N\mathbf{Z}$ FOR N SMALL

Dense matrices over $\mathbf{Z}/n\mathbf{Z}$ for n small

AUTHORS:

- William Stein
- Robert Bradshaw

This is a compiled implementation of dense matrices over $\mathbf{Z}/n\mathbf{Z}$ for n small.

EXAMPLES:

```
sage: a = matrix(Integers(37), 3, range(9), sparse=False); a
[0 1 2]
[3 4 5]
[6 7 8]
sage: a.rank()
2
sage: type(a)
<type 'sage.matrix.matrix_modn_dense_float.Matrix_modn_dense_float'>
sage: a[0,0] = 5
sage: a.rank()
3
sage: parent(a)
Full MatrixSpace of 3 by 3 dense matrices over Ring of integers modulo 37
```

```
sage: a^2
[ 3 23 31]
[20 17 29]
[25 16  0]
sage: a+a
[10  2  4]
[ 6  8 10]
[12 14 16]
```

```
sage: b = a.new_matrix(2,3,range(6)); b
[0 1 2]
[3 4 5]
sage: a*b
```

```
Traceback (most recent call last):
```

```
...
```

```
TypeError: unsupported operand parent(s) for '*': 'Full MatrixSpace of 3 by 3 dense matrices over Ring of integers modulo 37', 'Full MatrixSpace of 2 by 3 dense matrices over Ring of integers modulo 37'
```

```
sage: b*a
[15 18 21]
[20 17 29]
```

```
sage: TestSuite(a).run()
sage: TestSuite(b).run()
```

```
sage: a.echelonize(); a
[1 0 0]
[0 1 0]
[0 0 1]
sage: b.echelonize(); b
[ 1  0 36]
[ 0  1  2]
```

We create a matrix group:

```
sage: M = MatrixSpace(GF(3), 3, 3)
sage: G = MatrixGroup([M([0, 1, 0], [0, 0, 1], [1, 0, 0]), M([0, 1, 0], [1, 0, 0], [0, 0, 1])])
sage: G
Matrix group over Finite Field of size 3 with 2 generators (
[0 1 0] [0 1 0]
[0 0 1] [1 0 0]
[1 0 0], [0 0 1]
)
sage: G.gap()
Group([ [ [ 0*Z(3), Z(3)^0, 0*Z(3) ], [ 0*Z(3), 0*Z(3), Z(3)^0 ], [ Z(3)^0, 0*Z(3), 0*Z(3) ] ],
        [ [ 0*Z(3), Z(3)^0, 0*Z(3) ], [ Z(3)^0, 0*Z(3), 0*Z(3) ], [ 0*Z(3), 0*Z(3), Z(3)^0 ] ] )
```

TESTS:

```
sage: M = MatrixSpace(GF(5), 2, 2)
sage: A = M([1, 0, 0, 1])
sage: A - int(-1)
[2 0]
[0 2]
sage: B = M([4, 0, 0, 1])
sage: B - int(-1)
[0 0]
[0 2]
sage: Matrix(GF(5), 0, 0, sparse=False).inverse()
[]
```

```
class sage.matrix.matrix_modn_dense.Matrix_modn_dense
    Bases: sage.matrix.matrix_dense.Matrix_dense
```

TESTS:

```
sage: matrix(GF(7), 2, 2, [-1, int(-2), GF(7)(-3), 1/4])
[6 5]
[4 2]
```

charpoly (*var*='x', *algorithm*='generic')

Returns the characteristic polynomial of self.

INPUT:

- var - a variable name

- algorithm- 'generic' (default)

EXAMPLES:

```
sage: A = Mat(GF(7), 3, 3) (range(3)*3)
sage: A.charpoly()
x^3 + 4*x^2

sage: A = Mat(Integers(6), 3, 3) (range(9))
sage: A.charpoly()
x^3
```

determinant()

Return the determinant of this matrix.

EXAMPLES:

```
sage: m = matrix(GF(101), 5, range(25))
sage: m.det()
0

sage: m = matrix(Integers(4), 2, [2, 2, 2, 2])
sage: m.det()
0
```

TESTS:

```
sage: m = random_matrix(GF(3), 3, 4)
sage: m.determinant()
Traceback (most recent call last):
...
ValueError: self must be a square matrix
```

echelonize (algorithm='gauss', **kwds)

Puts self in row echelon form.

INPUT:

- self - a mutable matrix
- algorithm- 'gauss' - uses a custom slower $O(n^3)$ Gauss elimination implemented in Sage.
- **kwds - these are all ignored

OUTPUT:

- self is put in reduced row echelon form.
- the rank of self is computed and cached
- the pivot columns of self are computed and cached.
- the fact that self is now in echelon form is recorded and cached so future calls to echelonize return immediately.

EXAMPLES:

```
sage: a = matrix(GF(97), 3, 4, range(12))
sage: a.echelonize(); a
[ 1  0 96 95]
[ 0  1  2  3]
[ 0  0  0  0]
sage: a.pivots()
(0, 1)
```

hessenbergize()

Transforms self in place to its Hessenberg form.

lift()

Return the lift of this matrix to the integers.

EXAMPLES:

```
sage: a = matrix(GF(7), 2, 3, [1..6])
sage: a.lift()
[1 2 3]
[4 5 6]
sage: a.lift().parent()
Full MatrixSpace of 2 by 3 dense matrices over Integer Ring
```

Subdivisions are preserved when lifting:

```
sage: a.subdivide([], [1,1]); a
[1|2 3]
[4|5 6]
sage: a.lift()
[1|2 3]
[4|5 6]
```

matrix_window(row=0, col=0, nrows=-1, ncols=-1, check=1)

Return the requested matrix window.

EXAMPLES:

```
sage: a = matrix(GF(7), 3, range(9)); a
[0 1 2]
[3 4 5]
[6 0 1]
sage: type(a)
<type 'sage.matrix.matrix_modn_dense_float.Matrix_modn_dense_float'>
```

We test the optional check flag.

```
sage: matrix(GF(7), [1]).matrix_window(0,1,1,1)
Traceback (most recent call last):
...
IndexError: matrix window index out of range
sage: matrix(GF(7), [1]).matrix_window(0,1,1,1, check=False)
Matrix window of size 1 x 1 at (0,1):
[1]
```

minpoly(var='x', algorithm='generic', proof=None)

Returns the minimal polynomial of self.

INPUT:

- var - a variable name
- algorithm - 'generic' (default)
- proof - (default: True); whether to provably return the true minimal polynomial; if False, we only guarantee to return a divisor of the minimal polynomial. There are also certainly cases where the computed results is frequently not exactly equal to the minimal polynomial (but is instead merely a divisor of it).

WARNING: If proof=True, minpoly is insanely slow compared to proof=False.

EXAMPLES:


```
sage: R.<x>=GF(3)[]
sage: A = matrix(GF(3), 2, [0, 0, 1, 2])
sage: A.minpoly()
x^2 + x
```

Minpoly with `proof=False` is *dramatically* (“millions” of times!) faster than minpoly with `proof=True`. This matters since `proof=True` is the default, unless you first type “`proof.linear_algebra(False)`”:

```
sage: A.minpoly(proof=False) in [x, x+1, x^2+x]
True
```

randomize (*density=1, nonzero=False*)

Randomize `density` proportion of the entries of this matrix, leaving the rest unchanged.

INPUT:

- `density` - Integer; proportion (roughly) to be considered for changes
- `nonzero` - Bool (default: `False`); whether the new entries are forced to be non-zero

OUTPUT:

- None, the matrix is modified in-space

EXAMPLES:

```
sage: A = matrix(GF(5), 5, 5, 0)
sage: A.randomize(0.5); A
[0 0 0 2 0]
[0 3 0 0 2]
[4 0 0 0 0]
[4 0 0 0 0]
[0 1 0 0 0]
sage: A.randomize(); A
[3 3 2 1 2]
[4 3 3 2 2]
[0 3 3 3 3]
[3 3 2 2 4]
[2 2 2 1 4]
```

rank ()

Return the rank of this matrix.

EXAMPLES:

```
sage: m = matrix(GF(7), 5, range(25))
sage: m.rank()
2
```

Rank is not implemented over the integers modulo a composite yet.:

```
sage: m = matrix(Integers(4), 2, [2, 2, 2, 2])
sage: m.rank()
Traceback (most recent call last):
...
NotImplementedError: Echelon form not implemented over 'Ring of integers modulo 4'.
```

```
sage.matrix.matrix_modn_dense.is_Matrix_modn_dense(self)
```


SPARSE MATRICES OVER $\mathbb{Z}/N\mathbb{Z}$ FOR N SMALL

Sparse matrices over $\mathbb{Z}/n\mathbb{Z}$ for n small

This is a compiled implementation of sparse matrices over $\mathbb{Z}/n\mathbb{Z}$ for n small.

TODO: - move vectors into a Cython vector class - add `_add_` and `_mul_` methods.

EXAMPLES:

```
sage: a = matrix(Integers(37), 3, 3, range(9), sparse=True); a
[0 1 2]
[3 4 5]
[6 7 8]
sage: type(a)
<type 'sage.matrix.matrix_modn_sparse.Matrix_modn_sparse'>
sage: parent(a)
Full MatrixSpace of 3 by 3 sparse matrices over Ring of integers modulo 37
sage: a^2
[15 18 21]
[ 5 17 29]
[32 16  0]
sage: a+a
[ 0  2  4]
[ 6  8 10]
[12 14 16]
sage: b = a.new_matrix(2, 3, range(6)); b
[0 1 2]
[3 4 5]
sage: a*b
Traceback (most recent call last):
...
TypeError: unsupported operand parent(s) for '*': 'Full MatrixSpace of 3 by 3 sparse matrices over R
sage: b*a
[15 18 21]
[ 5 17 29]

sage: TestSuite(a).run()
sage: TestSuite(b).run()

sage: a.echelonize(); a
[ 1  0 36]
[ 0  1  2]
```

```

[ 0  0  0]
sage: b.echelonize(); b
[ 1  0 36]
[ 0  1  2]
sage: a.pivots()
(0, 1)
sage: b.pivots()
(0, 1)
sage: a.rank()
2
sage: b.rank()
2
sage: a[2,2] = 5
sage: a.rank()
3

```

TESTS: sage: matrix(Integers(37),0,0,sparse=True).inverse() []

class sage.matrix.matrix_modn_sparse.**Matrix_modn_sparse**
 Bases: sage.matrix.matrix_sparse.Matrix_sparse

Create a sparse matrix modulo n.

INPUT:

- parent - a matrix space
- entries
 - a Python list of triples (i,j,x), where $0 \leq i < \text{rows}$, $0 \leq j < \text{ncols}$, and x is coercible to an int. The i,j entry of self is set to x. The x's can be 0.
 - Alternatively, entries can be a list of *all* the entries of the sparse matrix (so they would be mostly 0).
- copy - ignored
- coerce - ignored

density()

Return the density of self, i.e., the ratio of the number of nonzero entries of self to the total size of self.

EXAMPLES:

```

sage: A = matrix(QQ,3,3,[0,1,2,3,0,0,6,7,8],sparse=True)
sage: A.density()
2/3

```

Notice that the density parameter does not ensure the density of a matrix; it is only an upper bound.

```

sage: A = random_matrix(GF(127),200,200,density=0.3,sparse=True)
sage: A.density()
2073/8000

```

lift()

Return lift of this matrix to a sparse matrix over the integers.

EXAMPLES: sage: a = matrix(GF(7),2,3,[1..6], sparse=True) sage: a.lift() [1 2 3] [4 5 6] sage: a.lift().parent() Full MatrixSpace of 2 by 3 sparse matrices over Integer Ring

Subdivisions are preserved when lifting:

```

sage: a.subdivide([], [1,1]); a
[1|2 3]
[4|5 6]

```

```
sage: a.lift()
[1||2 3]
[4||5 6]
```

matrix_from_columns (*cols*)

Return the matrix constructed from self using columns with indices in the columns list.

EXAMPLES:

```
sage: M = MatrixSpace(GF(127), 3, 3, sparse=True)
sage: A = M(range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.matrix_from_columns([2, 1])
[2 1]
[5 4]
[8 7]
```

matrix_from_rows (*rows*)

Return the matrix constructed from self using rows with indices in the rows list.

INPUT:

- rows - list or tuple of row indices

EXAMPLE:

```
sage: M = MatrixSpace(GF(127), 3, 3, sparse=True)
sage: A = M(range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.matrix_from_rows([2, 1])
[6 7 8]
[3 4 5]
```

P**rank** (*gauss=False*)

Compute the rank of self.

INPUT:

- gauss - if True LinBox' Gaussian elimination is used. If False 'Symbolic Reordering' as implemented in LinBox is used. If 'native' the native Sage implementation is used. (default: False)

EXAMPLE:

```
sage: A = random_matrix(GF(127), 200, 200, density=0.01, sparse=True)
sage: r1 = A.rank(gauss=False)
sage: r2 = A.rank(gauss=True)
sage: r3 = A.rank(gauss='native')
sage: r1 == r2 == r3
True
sage: r1
155
```

ALGORITHM: Uses LinBox or native implementation.

REFERENCES:

- Jean-Guillaume Dumas and Gilles Villars. ‘Computing the Rank of Large Sparse Matrices over Finite Fields’. Proc. CASC’2002, The Fifth International Workshop on Computer Algebra in Scientific Computing, Big Yalta, Crimea, Ukraine, 22-27 sept. 2002, Springer-Verlag, <http://perso.ens-lyon.fr/gilles.villard/BIBLIOGRAPHIE/POSTSCRIPT/rankjgd.ps>

Note: For very sparse matrices Gaussian elimination is faster because it barely has anything to do. If the fill in needs to be considered, ‘Symbolic Reordering’ is usually much faster.

swap_rows (*r1*, *r2*)

transpose ()

Return the transpose of self.

EXAMPLE:

```
sage: A = matrix(GF(127), 3, 3, [0, 1, 0, 2, 0, 0, 3, 0, 0], sparse=True)
```

```
sage: A
```

```
[0 1 0]
```

```
[2 0 0]
```

```
[3 0 0]
```

```
sage: A.transpose()
```

```
[0 2 3]
```

```
[1 0 0]
```

```
[0 0 0]
```

.T is a convenient shortcut for the transpose:

```
sage: A.T
```

```
[0 2 3]
```

```
[1 0 0]
```

```
[0 0 0]
```

visualize_structure (*filename=None*, *maxsize=512*)

Write a PNG image to ‘filename’ which visualizes self by putting black pixels in those positions which have nonzero entries.

White pixels are put at positions with zero entries. If ‘maxsize’ is given, then the maximal dimension in either x or y direction is set to ‘maxsize’ depending on which is bigger. If the image is scaled, the darkness of the pixel reflects how many of the represented entries are nonzero. So if e.g. one image pixel actually represents a 2x2 submatrix, the dot is darker the more of the four values are nonzero.

INPUT:

- filename - either a path or None in which case a filename in the current directory is chosen automatically (default:None)
- maxsize - maximal dimension in either x or y direction of the resulting image. If None or a maxsize larger than max(self.nrows(),self.ncols()) is given the image will have the same pixelsize as the matrix dimensions (default: 512)

EXAMPLES:

```
sage: M = Matrix(GF(7), [[0, 0, 0, 1, 0, 0, 0, 0], [0, 1, 0, 0, 0, 0, 1, 0]], sparse=True); M
```

```
[0 0 0 1 0 0 0 0]
```

```
[0 1 0 0 0 0 1 0]
```

```
sage: M.visualize_structure()
```

SYMBOLIC MATRICES

Symbolic matrices

Matrices with symbolic entries. The underlying representation is a pointer to a Maxima object.

EXAMPLES:

```
sage: matrix(SR, 2, 2, range(4))
[0 1]
[2 3]
sage: matrix(SR, 2, 2, var('t'))
[t 0]
[0 t]
```

Arithmetic:

```
sage: -matrix(SR, 2, range(4))
[ 0 -1]
[-2 -3]
sage: m = matrix(SR, 2, [1..4]); sqrt(2)*m
[ sqrt(2) 2*sqrt(2)]
[3*sqrt(2) 4*sqrt(2)]
sage: m = matrix(SR, 4, [1..4^2])
sage: m * m
[ 90 100 110 120]
[202 228 254 280]
[314 356 398 440]
[426 484 542 600]
```

```
sage: m = matrix(SR, 3, [1, 2, 3]); m
[1]
[2]
[3]
sage: m.transpose() * m
[14]
```

Computing inverses:

```
sage: M = matrix(SR, 2, var('a,b,c,d'))
sage: ~M
[1/a - b*c/(a^2*(b*c/a - d))      b/(a*(b*c/a - d))]
[      c/(a*(b*c/a - d))      -1/(b*c/a - d)]
sage: (~M*M).simplify_rational()
[1 0]
[0 1]
```

```
sage: M = matrix(SR, 3, 3, range(9)) - var('t')
sage: (~M * M).simplify_rational()
[1 0 0]
[0 1 0]
[0 0 1]
```

```
sage: matrix(SR, 1, 1, 1).inverse()
[1]
sage: matrix(SR, 0, 0).inverse()
[]
sage: matrix(SR, 3, 0).inverse()
Traceback (most recent call last):
...
ArithmeticError: self must be a square matrix
```

Transposition:

```
sage: m = matrix(SR, 2, [sqrt(2), -1, pi, e^2])
sage: m.transpose()
[sqrt(2)      pi]
[      -1      e^2]
```

.T is a convenient shortcut for the transpose:

```
sage: m.T
[sqrt(2)      pi]
[      -1      e^2]
```

Test pickling:

```
sage: m = matrix(SR, 2, [sqrt(2), 3, pi, e]); m
[sqrt(2)      3]
[      pi      e]
sage: TestSuite(m).run()
```

Comparison:

```
sage: m = matrix(SR, 2, [sqrt(2), 3, pi, e])
sage: cmp(m,m)
0
sage: cmp(m,3) != 0
True
sage: m = matrix(SR,2,[1..4]); n = m^2
sage: (exp(m+n) - exp(m)*exp(n)).simplify_rational() == 0      # indirect test
True
```

Determinant:

```
sage: M = matrix(SR, 2, 2, [x,2,3,4])
sage: M.determinant()
4*x - 6
sage: M = matrix(SR, 3,3,range(9))
sage: M.det()
0
sage: t = var('t')
sage: M = matrix(SR, 2, 2, [cos(t), sin(t), -sin(t), cos(t)])
sage: M.det()
cos(t)^2 + sin(t)^2
```



```
sage: M = matrix([[sqrt(x), 0, 0, 0], [0, 1, 0, 0], [0, 0, 1, 0], [0, 0, 0, 1]])
sage: det(M)
sqrt(x)
```

Permanents:

```
sage: M = matrix(SR, 2, 2, [x, 2, 3, 4])
sage: M.permanent()
4*x + 6
```

Rank:

```
sage: M = matrix(SR, 5, 5, range(25))
sage: M.rank()
2
sage: M = matrix(SR, 5, 5, range(25)) - var('t')
sage: M.rank()
5
```

```
.. warning::
```

```
:meth:`rank` may return the wrong answer if it cannot determine that a
matrix element that is equivalent to zero is indeed so.
```

Copying symbolic matrices:

```
sage: m = matrix(SR, 2, [sqrt(2), 3, pi, e])
sage: n = copy(m)
sage: n[0,0] = sin(1)
sage: m
[sqrt(2)      3]
[      pi      e]
sage: n
[sin(1)      3]
[      pi      e]
```

Conversion to Maxima:

```
sage: m = matrix(SR, 2, [sqrt(2), 3, pi, e])
sage: m._maxima_()
matrix([sqrt(2), 3], [%pi, %e])
```

```
class sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense
    Bases: sage.matrix.matrix_generic_dense.Matrix_generic_dense
```

See `Matrix_generic_dense` for documentation.

TESTS:

We check that the problem related to Trac #9049 is not an issue any more:

```
sage: S.<t>=PolynomialRing(QQ)
sage: F.<q>=QQ.extension(t^4+1)
sage: R.<x,y>=PolynomialRing(F)
sage: M = MatrixSpace(R, 1, 2)
sage: from sage.matrix.matrix_generic_dense import Matrix_generic_dense
sage: Matrix_generic_dense(M, (x, y), True, True)
[x y]
```

arguments()

Returns a tuple of the arguments that self can take.

EXAMPLES:

```
sage: var('x,y,z')
(x, y, z)
sage: M = MatrixSpace(SR, 2, 2)
sage: M(x).arguments()
(x,)
sage: M(x+sin(x)).arguments()
(x,)
```

charpoly (*var='x', algorithm=None*)

Compute the characteristic polynomial of self, using maxima.

EXAMPLES:

```
sage: M = matrix(SR, 2, 2, var('a,b,c,d'))
sage: M.charpoly('t')
t^2 + (-a - d)*t - b*c + a*d
sage: matrix(SR, 5, [1..5^2]).charpoly()
x^5 - 65*x^4 - 250*x^3
```

TESTS:

The cached polynomial should be independent of the `var` argument ([trac ticket #12292](#)). We check (indirectly) that the second call uses the cached value by noting that its result is not cached:

```
sage: M = MatrixSpace(SR, 2)
sage: A = M(range(0, 2^2))
sage: type(A)
<type 'sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense'>
sage: A.charpoly('x')
x^2 - 3*x - 2
sage: A.charpoly('y')
y^2 - 3*y - 2
sage: A._cache['charpoly']
x^2 - 3*x - 2
```

Ensure the variable name of the polynomial does not conflict with variables used within the matrix ([trac ticket #14403](#)):

```
sage: Matrix(SR, [[sqrt(x), x], [1, x]]).charpoly().list()
[x^(3/2) - x, -x - sqrt(x), 1]
```

eigenvalues()

Compute the eigenvalues by solving the characteristic polynomial in maxima

EXAMPLES:

```
sage: a=matrix(SR, [[1,2],[3,4]])
sage: a.eigenvalues()
[-1/2*sqrt(33) + 5/2, 1/2*sqrt(33) + 5/2]
```

eigenvectors_left()

Compute the left eigenvectors of a matrix.

For each distinct eigenvalue, returns a list of the form (e, V, n) where e is the eigenvalue, V is a list of eigenvectors forming a basis for the corresponding left eigenspace, and n is the algebraic multiplicity of the eigenvalue.

EXAMPLES:

```

sage: A = matrix(SR, 3, 3, range(9)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: es = A.eigenvectors_left(); es
[(-3*sqrt(6) + 6, [(1, -1/5*sqrt(6) + 4/5, -2/5*sqrt(3)*sqrt(2) + 3/5)], 1), (3*sqrt(6) + 6,
sage: eval, [evec], mult = es[0]
sage: delta = eval*evec - evec*A
sage: abs(abs(delta)) < 1e-10
abs(sqrt(1/25*(3*(2*sqrt(3)*sqrt(2) - 3)*(sqrt(6) - 2) + 16*sqrt(3)*sqrt(2) + 5*sqrt(6) - 54
sage: abs(abs(delta)).n() < 1e-10
True

sage: A = matrix(SR, 2, 2, var('a,b,c,d'))
sage: A.eigenvectors_left()
[(1/2*a + 1/2*d - 1/2*sqrt(a^2 + 4*b*c - 2*a*d + d^2), [(1, -1/2*(a - d + sqrt(a^2 + 4*b*c -
sage: es = A.eigenvectors_left()); es
[(1/2*a + 1/2*d - 1/2*sqrt(a^2 + 4*b*c - 2*a*d + d^2), [(1, -1/2*(a - d + sqrt(a^2 + 4*b*c -
sage: eval, [evec], mult = es[0]
sage: delta = eval*evec - evec*A
sage: delta.apply_map(lambda x: x.full_simplify())
(0, 0)

```

This routine calls Maxima and can struggle with even small matrices with a few variables, such as a 3×3 matrix with three variables. However, if the entries are integers or rationals it can produce exact values in a reasonable time. These examples create 0-1 matrices from the adjacency matrices of graphs and illustrate how the format and type of the results differ when the base ring changes. First for matrices over the rational numbers, then the same matrix but viewed as a symbolic matrix.

```

sage: G=graphs.CycleGraph(5)
sage: am = G.adjacency_matrix()
sage: spectrum = am.eigenvectors_left()
sage: qqbar_value = spectrum[2][0]
sage: type(qqbar_value)
<class 'sage.rings.qqbar.AlgebraicNumber'>
sage: qqbar_value
0.618033988749895?

sage: am = G.adjacency_matrix().change_ring(SR)
sage: spectrum = am.eigenvectors_left()
sage: symbolic_value = spectrum[2][0]
sage: type(symbolic_value)
<type 'sage.symbolic.expression.Expression'>
sage: symbolic_value
1/2*sqrt(5) - 1/2

sage: qqbar_value == symbolic_value
True

```

A slightly larger matrix with a “nice” spectrum.

```

sage: G=graphs.CycleGraph(6)
sage: am = G.adjacency_matrix().change_ring(SR)
sage: am.eigenvectors_left()
[(-1, [(1, 0, -1, 1, 0, -1), (0, 1, -1, 0, 1, -1)], 2), (1, [(1, 0, -1, -1, 0, 1), (0, 1, 1,

```

eigenvectors_right()

Compute the right eigenvectors of a matrix.

For each distinct eigenvalue, returns a list of the form (e,V,n) where e is the eigenvalue, V is a list of eigenvectors forming a basis for the corresponding right eigenspace, and n is the algebraic multiplicity of the eigenvalue.

EXAMPLES:

```
sage: A = matrix(SR, 2, 2, range(4)); A
[0 1]
[2 3]
sage: right = A.eigenvectors_right(); right
[(-1/2*sqrt(17) + 3/2, [(1, -1/2*sqrt(17) + 3/2)], 1), (1/2*sqrt(17) + 3/2, [(1, 1/2*sqrt(17) + 3/2)], 1)]
```

The right eigenvectors are nothing but the left eigenvectors of the transpose matrix:

```
sage: left = A.transpose().eigenvectors_left(); left
[(-1/2*sqrt(17) + 3/2, [(1, -1/2*sqrt(17) + 3/2)], 1), (1/2*sqrt(17) + 3/2, [(1, 1/2*sqrt(17) + 3/2)], 1)]
sage: right[0][1] == left[0][1]
True
```

exp()

Return the matrix exponential of this matrix X , which is the matrix

$$e^X = \sum_{k=0}^{\infty} \frac{X^k}{k!}.$$

This function depends on maxima's matrix exponentiation function, which does not deal well with floating point numbers. If the matrix has floating point numbers, they will be rounded automatically to rational numbers during the computation.

EXAMPLES:

```
sage: m = matrix(SR, 2, [0, x, x, 0]); m
[0 x]
[x 0]
sage: m.exp()
[1/2*(e^(2*x) + 1)*e^(-x) 1/2*(e^(2*x) - 1)*e^(-x)]
[1/2*(e^(2*x) - 1)*e^(-x) 1/2*(e^(2*x) + 1)*e^(-x)]
sage: exp(m)
[1/2*(e^(2*x) + 1)*e^(-x) 1/2*(e^(2*x) - 1)*e^(-x)]
[1/2*(e^(2*x) - 1)*e^(-x) 1/2*(e^(2*x) + 1)*e^(-x)]
```

Exp works on 0x0 and 1x1 matrices:

```
sage: m = matrix(SR, 0, []); m
[]
sage: m.exp()
[]
sage: m = matrix(SR, 1, [2]); m
[2]
sage: m.exp()
[e^2]
```

Commuting matrices m, n have the property that $e^{m+n} = e^m e^n$ (but non-commuting matrices need not):

```
sage: m = matrix(SR, 2, [1..4]); n = m^2
sage: m*n
[ 37  54]
[ 81 118]
sage: n*m
[ 37  54]
[ 81 118]
```

```
sage: a = exp(m+n) - exp(m)*exp(n)
sage: a.simplify_rational() == 0
True
```

The input matrix must be square:

```
sage: m = matrix(SR, 2, 3, [1..6]); exp(m)
Traceback (most recent call last):
...
ValueError: exp only defined on square matrices
```

In this example we take the symbolic answer and make it numerical at the end:

```
sage: exp(matrix(SR, [[1.2, 5.6], [3, 4]])).change_ring(RDF)
[346.557487298 661.734590934]
[354.500673715 677.424782765]
```

Another example involving the reversed identity matrix, which we clumsily create:

```
sage: m = identity_matrix(SR, 4); m = matrix(list(reversed(m.rows())) * x
sage: exp(m)
[1/2*(e^(2*x) + 1)*e^(-x) 0 0 1/2*(e^(2*x) - 1)
[ 0 1/2*(e^(2*x) + 1)*e^(-x) 1/2*(e^(2*x) - 1)*e^(-x)
[ 0 1/2*(e^(2*x) - 1)*e^(-x) 1/2*(e^(2*x) + 1)*e^(-x)
[1/2*(e^(2*x) - 1)*e^(-x) 0 0 1/2*(e^(2*x) + 1)
```

expand()

Operates point-wise on each element.

EXAMPLES:

```
sage: M = matrix(2, 2, range(4)) - var('x')
sage: M*M
[ x^2 + 2 -2*x + 3]
[ -4*x + 6 (x - 3)^2 + 2]
sage: (M*M).expand()
[ x^2 + 2 -2*x + 3]
[ -4*x + 6 x^2 - 6*x + 11]
```

factor()

Operates point-wise on each element.

EXAMPLES:

```
sage: M = matrix(SR, 2, 2, x^2 - 2*x + 1); M
[x^2 - 2*x + 1 0]
[ 0 x^2 - 2*x + 1]
sage: M.factor()
[(x - 1)^2 0]
[ 0 (x - 1)^2]
```

fcp (var='x')

Return the factorization of the characteristic polynomial of self.

INPUT:

- var - (default: 'x') name of variable of charpoly

EXAMPLES:

```

sage: a = matrix(SR, [[1,2],[3,4]])
sage: a.fcp()
x^2 - 5*x - 2
sage: [i for i in a.fcp()]
[(x^2 - 5*x - 2, 1)]
sage: a = matrix(SR, [[1,0],[0,2]])
sage: a.fcp()
(x - 2) * (x - 1)
sage: [i for i in a.fcp()]
[(x - 2, 1), (x - 1, 1)]
sage: a = matrix(SR, 5, [1..5^2])
sage: a.fcp()
(x^2 - 65*x - 250) * x^3
sage: list(a.fcp())
[(x^2 - 65*x - 250, 1), (x, 3)]

```

number_of_arguments()

Returns the number of arguments that self can take.

EXAMPLES:

```

sage: var('a,b,c,x,y')
(a, b, c, x, y)
sage: m = matrix([[a, (x+y)/(x+y)], [x^2, y^2+2]]); m
[
  a      1]
[
  x^2 y^2 + 2]
sage: m.number_of_arguments()
3

```

simplify()

Simplifies self.

EXAMPLES:

```

sage: var('x,y,z')
(x, y, z)
sage: m = matrix([[z, (x+y)/(x+y)], [x^2, y^2+2]]); m
[
  z      1]
[
  x^2 y^2 + 2]
sage: m.simplify()
[
  z      1]
[
  x^2 y^2 + 2]

```

simplify_rational()

EXAMPLES:

```

sage: M = matrix(SR, 3, 3, range(9)) - var('t')
sage: (~M*M)[0,0]
t*(3*(2/t + (6/t + 7)/((t - 3/t - 4)*t))*(2/t + (6/t + 5)/((t - 3/t - 4)*t))/((t - (6/t + 7)*(6/t + 5)/(t - 3/t - 4) - 12/t - 8) + 1/t + 3/((t - 3/t - 4)*t^2)) - 6*(2/t + (6/t + 5)/((t - 3/t - 4)*t))/((t - (6/t + 7)*(6/t + 5)/(t - 3/t - 4) - 12/t - 8) - 3*(6/t + 7)*(2/t + (6/t + 5)/((t - 3/t - 4)*t))/((t - (6/t + 7)*(6/t + 5)/(t - 3/t - 4) - 12/t - 8)*(t - 3/t - 4)) - 3/((t - 3/t - 4)*t)
sage: expand((~M*M)[0,0])
1
sage: (~M * M).simplify_rational()
[1 0 0]
[0 1 0]
[0 0 1]

```

simplify_trig()

EXAMPLES:

```
sage: theta = var('theta')
sage: M = matrix(SR, 2, 2, [cos(theta), sin(theta), -sin(theta), cos(theta)])
sage: ~M
[1/cos(theta) - sin(theta)^2/((sin(theta)^2/cos(theta) + cos(theta))*cos(theta)^2)
 [
      sin(theta)/((sin(theta)^2/cos(theta) + cos(theta))*cos(theta))
sage: (~M).simplify_trig()
[ cos(theta) -sin(theta)]
[ sin(theta)  cos(theta)]
```

variables()

Returns the variables of self.

EXAMPLES:

```
sage: var('a,b,c,x,y')
(a, b, c, x, y)
sage: m = matrix([[x, x+2], [x^2, x^2+2]]); m
[
  x   x + 2]
[
  x^2 x^2 + 2]
sage: m.variables()
(x,)
sage: m = matrix([[a, b+c], [x^2, y^2+2]]); m
[
  a   b + c]
[
  x^2 y^2 + 2]
sage: m.variables()
(a, b, c, x, y)
```


DENSE MATRICES OVER THE INTEGER RING

Dense matrices over the integer ring

AUTHORS:

- William Stein
- Robert Bradshaw

EXAMPLES:

```
sage: a = matrix(ZZ, 3, 3, range(9)); a
[0 1 2]
[3 4 5]
[6 7 8]
sage: a.det()
0
sage: a[0,0] = 10; a.det()
-30
sage: a.charpoly()
x^3 - 22*x^2 + 102*x + 30
sage: b = -3*a
sage: a == b
False
sage: b < a
True
```

TESTS:

```
sage: a = matrix(ZZ, 2, range(4), sparse=False)
sage: TestSuite(a).run()
sage: Matrix(ZZ, 0, 0).inverse()
[]
```

```
class sage.matrix.matrix_integer_dense.Matrix_integer_dense
    Bases: sage.matrix.matrix_dense.Matrix_dense
    Matrix over the integers.
```

On a 32-bit machine, they can have at most $2^{32} - 1$ rows or columns. On a 64-bit machine, matrices can have at most $2^{64} - 1$ rows or columns.

EXAMPLES:

```

sage: a = MatrixSpace(ZZ, 3) (2); a
[2 0 0]
[0 2 0]
[0 0 2]
sage: a = matrix(ZZ, 1, 3, [1, 2, -3]); a
[ 1  2 -3]
sage: a = MatrixSpace(ZZ, 2, 4) (2); a
Traceback (most recent call last):
...
TypeError: nonzero scalar matrix must be square

```

BKZ (*delta=None, algorithm='fpLLL', fp=None, block_size=10, prune=0, use_givens=False, precision=0, max_loops=0, max_time=0, auto_abort=False*)
 Block Korkin-Zolotarev reduction.

INPUT:

- *delta* – (default: 0.99) LLL parameter
- *algorithm* – (default: "fpLLL") "fpLLL" or "NTL"
- *fp* – floating point number implementation
 - None – NTL's exact reduction or fpLLL's wrapper (default)
 - 'fp' – double precision: NTL's FP or fpLLL's double
 - 'qd' – NTL's QP or fpLLL's long doubles
 - 'qd1' – quad doubles: Uses quad_float precision to compute Gram-Schmidt, but uses double precision in the search phase of the block reduction algorithm. This seems adequate for most purposes, and is faster than 'qd', which uses quad_float precision uniformly throughout (NTL only).
 - 'xd' – extended exponent: NTL's XD or fpLLL's dpe
 - 'rr' – arbitrary precision: NTL's RR or fpLLL's MPFR
- *block_size* – (default: 10) Specifies the size of the blocks in the reduction. High values yield shorter vectors, but the running time increases double exponentially with *block_size*. *block_size* should be between 2 and the number of rows of *self*.

NLT SPECIFIC INPUTS:

- *prune* – (default: 0) The optional parameter *prune* can be set to any positive number to invoke the Volume Heuristic from [SH95]. This can significantly reduce the running time, and hence allow much bigger block size, but the quality of the reduction is of course not as good in general. Higher values of *prune* mean better quality, and slower running time. When *prune* is 0, pruning is disabled. Recommended usage: for *block_size*==30, set $10 \leq \text{prune} \leq 15$.
- *use_givens* – Use Given's orthogonalization. This is a bit slower, but generally much more stable, and is really the preferred orthogonalization strategy. For a nice description of this, see Chapter 5 of [GL96].

fpLLL SPECIFIC INPUTS:

- *precision* – (default: 0 for automatic choice) bit precision to use if *fp*='rr' is set
- *max_loops* – (default: 0 for no restriction) maximum number of full loops
- *max_time* – (default: 0 for no restriction) stop after time seconds (up to loop completion)
- *auto_abort* – (default: False) heuristic, stop when the average slope of $\log(\|b_i^*\|)$ does not decrease fast enough

EXAMPLES:

```
sage: A = Matrix(ZZ, 3, 3, range(1, 10))
sage: A.BKZ()
[ 0  0  0]
[ 2  1  0]
[-1  1  3]
```

```
sage: A = Matrix(ZZ, 3, 3, range(1, 10))
sage: A.BKZ(use_givens=True)
[ 0  0  0]
[ 2  1  0]
[-1  1  3]
```

```
sage: A = Matrix(ZZ, 3, 3, range(1, 10))
sage: A.BKZ(fp="fp")
[ 0  0  0]
[ 2  1  0]
[-1  1  3]
```

ALGORITHM:

Calls either NTL or fpLLL.

REFERENCES:

LLL (*delta=None, eta=None, algorithm='fpLLL:wrapper', fp=None, prec=0, early_red=False, use_givens=False, use_siegel=False*)
Return LLL reduced or approximated LLL reduced lattice R for this matrix interpreted as a lattice.

A lattice (b_1, b_2, \dots, b_d) is (δ, η) -LLL-reduced if the two following conditions hold:

- For any $i > j$, we have $|\mu_{i,j}| \leq \eta$.
- For any $i < d$, we have $\delta |b_i^*|^2 \leq |b_{i+1} + \mu_{i+1,i} b_i^*|^2$,

where $\mu_{i,j} = \langle b_i, b_j^* \rangle / \langle b_j^*, b_j^* \rangle$ and b_i^* is the i -th vector of the Gram-Schmidt orthogonalisation of (b_1, b_2, \dots, b_d) .

The default reduction parameters are $\delta = 3/4$ and $\eta = 0.501$. The parameters δ and η must satisfy: $0.25 < \delta \leq 1.0$ and $0.5 \leq \eta < \sqrt{\delta}$. Polynomial time complexity is only guaranteed for $\delta < 1$.

The lattice is returned as a matrix. Also the rank (and the determinant) of `self` are cached if those are computed during the reduction. Note that in general this only happens when `self.rank() == self.ncols()` and the exact algorithm is used.

INPUT:

- `delta` – (default: 0.99) δ parameter as described above
- `eta` – (default: 0.501) η parameter as described above, ignored by NTL
- `algorithm` – string one of the algorithms listed below (default: "fpLLL:wrapper").
- `fp` – floating point number implementation:
 - `None` – NTL's exact reduction or fpLLL's wrapper
 - `'fp'` – double precision: NTL's FP or fpLLL's double
 - `'qd'` – NTL's QP or fpLLL's long doubles
 - `'xd'` – extended exponent: NTL's XD or fpLLL's dpe
 - `'rr'` – arbitrary precision: NTL's RR or fpLLL's MPFR

- `prec` – (default: auto choose) precision, ignored by NTL
- `early_red` – (default: `False`) perform early reduction, ignored by NTL
- `use_givens` – (default: `False`) use Givens orthogonalization only applicable to approximate reductions and NTL; this is more stable but slower
- `use_siegel` – (default: `False`) use Siegel’s condition instead of Lovasz’s condition, ignored by NTL

Also, if the verbose level is at least 2, some more verbose output is printed during the computation.

AVAILABLE ALGORITHMS:

- `NTL:LLL` - NTL’s LLL + choice of `fp`.
- `fpLLL:heuristic` - `fpLLL`’s heuristic + choice of `fp`.
- `fpLLL:fast` - `fpLLL`’s fast + choice of `fp`.
- `fpLLL:proved` - `fpLLL`’s proved + choice of `fp`.
- `fpLLL:wrapper` - `fpLLL`’s automatic choice (default).

OUTPUT:

A matrix over the integers.

EXAMPLES:

```
sage: A = Matrix(ZZ, 3, 3, range(1, 10))
sage: A.LLL()
[ 0  0  0]
[ 2  1  0]
[-1  1  3]
```

We compute the extended GCD of a list of integers using LLL, this example is from the Magma handbook:

```
sage: Q = [ 67015143, 248934363018, 109210, 25590011055, 74631449,
....:      10230248, 709487, 68965012139, 972065, 864972271 ]
sage: n = len(Q)
sage: S = 100
sage: X = Matrix(ZZ, n, n + 1)
sage: for i in xrange(n):
...     X[i, i + 1] = 1
sage: for i in xrange(n):
...     X[i, 0] = S*Q[i]
sage: L = X.LLL()
sage: M = L.row(n-1).list()[1:]
sage: M
[-3, -1, 13, -1, -4, 2, 3, 4, 5, -1]
sage: add([Q[i]*M[i] for i in range(n)])
-1
```

TESTS:

```
sage: matrix(ZZ, 0, 0).LLL()
[]
sage: matrix(ZZ, 3, 0).LLL()
[]
sage: matrix(ZZ, 0, 3).LLL()
[]

sage: M = matrix(ZZ, [[1, 2, 3], [31, 41, 51], [101, 201, 301]])
sage: A = M.LLL()
```

```

sage: A
[ 0  0  0]
[-1  0  1]
[ 1  1  1]
sage: B = M.LLL(algorithm='NTL:LLL')
sage: C = M.LLL(algorithm='NTL:LLL', fp=None)
sage: D = M.LLL(algorithm='NTL:LLL', fp='fp')
sage: F = M.LLL(algorithm='NTL:LLL', fp='xd')
sage: G = M.LLL(algorithm='NTL:LLL', fp='rr')
sage: A == B == C == D == F == G
True
sage: H = M.LLL(algorithm='NTL:LLL', fp='qd')
Traceback (most recent call last):
...
TypeError: algorithm NTL:LLL_QD not supported

```

Note: See `ntl.mat_ZZ` or `sage.libs.fplll.fplll` for details on the used algorithms.

LLL_gram()

LLL reduction of the lattice whose gram matrix is `self`.

INPUT:

- `M` – gram matrix of a definite quadratic form

OUTPUT:

`U` - unimodular transformation matrix such that $U.T * M * U$ is LLL-reduced.

ALGORITHM: Use PARI

EXAMPLES:

```

sage: M = Matrix(ZZ, 2, 2, [5, 3, 3, 2]) ; M
[5 3]
[3 2]
sage: U = M.LLL_gram(); U
[-1  1]
[ 1 -2]
sage: U.transpose() * M * U
[1 0]
[0 1]

```

Semidefinite and indefinite forms no longer raise a `ValueError`:

```

sage: Matrix(ZZ, 2, 2, [2, 6, 6, 3]).LLL_gram()
[-3 -1]
[ 1  0]
sage: Matrix(ZZ, 2, 2, [1, 0, 0, -1]).LLL_gram()
[ 0 -1]
[ 1  0]

```

antitranspose()

Returns the antitranspose of `self`, without changing `self`.

EXAMPLES:

```

sage: A = matrix(2, 3, range(6))
sage: type(A)
<type 'sage.matrix.matrix_integer_dense.Matrix_integer_dense'>

```

```
sage: A.antitranspose()
[5 2]
[4 1]
[3 0]
sage: A
[0 1 2]
[3 4 5]

sage: A.subdivide(1,2); A
[0 1|2]
[---+--]
[3 4|5]
sage: A.antitranspose()
[5|2]
[-+-]
[4|1]
[3|0]
```

augment (*right*, *subdivide=False*)

Returns a new matrix formed by appending the matrix (or vector) *right* on the right side of *self*.

INPUT:

- *right* - a matrix, vector or free module element, whose dimensions are compatible with *self*.
- *subdivide* - default: `False` - request the resulting matrix to have a new subdivision, separating *self* from *right*.

OUTPUT:

A new matrix formed by appending *right* onto the right side of *self*. If *right* is a vector (or free module element) then in this context it is appropriate to consider it as a column vector. (The code first converts a vector to a 1-column matrix.)

EXAMPLES:

```
sage: A = matrix(ZZ, 4, 5, range(20))
sage: B = matrix(ZZ, 4, 3, range(12))
sage: A.augment(B)
[ 0  1  2  3  4  0  1  2]
[ 5  6  7  8  9  3  4  5]
[10 11 12 13 14  6  7  8]
[15 16 17 18 19  9 10 11]
```

A vector may be augmented to a matrix.

```
sage: A = matrix(ZZ, 3, 5, range(15))
sage: v = vector(ZZ, 3, range(3))
sage: A.augment(v)
[ 0  1  2  3  4  0]
[ 5  6  7  8  9  1]
[10 11 12 13 14  2]
```

The `subdivide` option will add a natural subdivision between *self* and *right*. For more details about how subdivisions are managed when augmenting, see `sage.matrix.matrix1.Matrix.augment()`.

```
sage: A = matrix(ZZ, 3, 5, range(15))
sage: B = matrix(ZZ, 3, 3, range(9))
sage: A.augment(B, subdivide=True)
[ 0  1  2  3  4| 0  1  2]
```

```
[ 5  6  7  8  9| 3  4  5]
[10 11 12 13 14| 6  7  8]
```

Errors are raised if the sizes are incompatible.

```
sage: A = matrix(ZZ, [[1, 2],[3, 4]])
sage: B = matrix(ZZ, [[10, 20], [30, 40], [50, 60]])
sage: A.augment(B)
Traceback (most recent call last):
...
TypeError: number of rows must be the same, not 2 != 3
```

charpoly (*var='x', algorithm='linbox'*)

INPUT:

- var - a variable name
- algorithm - 'linbox' (default) 'generic'

Note: Linbox charpoly disabled on 64-bit machines, since it hangs in many cases.

EXAMPLES:

```
sage: A = matrix(ZZ, 6, range(36))
sage: f = A.charpoly(); f
x^6 - 105*x^5 - 630*x^4
sage: f(A) == 0
True
sage: n=20; A = Mat(ZZ,n)(range(n^2))
sage: A.charpoly()
x^20 - 3990*x^19 - 266000*x^18
sage: A.minpoly()
x^3 - 3990*x^2 - 266000*x
```

TESTS:

The cached polynomial should be independent of the `var` argument ([trac ticket #12292](#)). We check (indirectly) that the second call uses the cached value by noting that its result is not cached:

```
sage: M = MatrixSpace(ZZ, 2)
sage: A = M(range(0, 2^2))
sage: type(A)
<type 'sage.matrix.matrix_integer_dense.Matrix_integer_dense'>
sage: A.charpoly('x')
x^2 - 3*x - 2
sage: A.charpoly('y')
y^2 - 3*y - 2
sage: A._cache['charpoly_linbox']
x^2 - 3*x - 2
```

decomposition (***kws*)

Returns the decomposition of the free module on which this matrix `A` acts from the right (i.e., the action is `x` goes to `x A`), along with whether this matrix acts irreducibly on each factor. The factors are guaranteed to be sorted in the same way as the corresponding factors of the characteristic polynomial, and are saturated as `ZZ` modules.

INPUT:

- self - a matrix over the integers

- **kwds - these are passed onto to the decomposition over QQ command.

EXAMPLES:

```
sage: t = ModularSymbols(11, sign=1).hecke_matrix(2)
sage: w = t.change_ring(ZZ)
sage: w.list()
[3, -1, 0, -2]
```

determinant (*algorithm='default', proof=None, stabilize=2*)

Return the determinant of this matrix.

INPUT:

- algorithm
 - 'default' - automatically determine which algorithm to use depending on the matrix.
 - 'padic' - uses a p-adic / multimodular algorithm that relies on code in IML and linbox
 - 'linbox' - calls linbox det (you *must* set proof=False to use this!)
 - 'ntl' - calls NTL's det function
 - 'pari' - uses PARI
- proof - bool or None; if None use proof.linear_algebra(); only relevant for the padic algorithm.

Note: It would be *VERY VERY* hard for det to fail even with proof=False.

- stabilize - if proof is False, require det to be the same for this many CRT primes in a row. Ignored if proof is True.

ALGORITHM: The p-adic algorithm works by first finding a random vector v , then solving $A \cdot x = v$ and taking the denominator d . This gives a divisor of the determinant. Then we compute $\det(A)/d$ using a multimodular algorithm and the Hadamard bound, skipping primes that divide d .

TIMINGS: This is perhaps the fastest implementation of determinants in the world. E.g., for a 500x500 random matrix with 32-bit entries on a core2 duo 2.6Ghz running OS X, Sage takes 4.12 seconds, whereas Magma takes 62.87 seconds (both with proof False). With proof=True on the same problem Sage takes 5.73 seconds. For another example, a 200x200 random matrix with 1-digit entries takes 4.18 seconds in pari, 0.18 in Sage with proof True, 0.11 in Sage with proof False, and 0.21 seconds in Magma with proof True and 0.18 in Magma with proof False.

EXAMPLES:

```
sage: A = matrix(ZZ, 8, 8, [3..66])
sage: A.determinant()
0

sage: A = random_matrix(ZZ, 20, 20)
sage: D1 = A.determinant()
sage: A._clear_cache()
sage: D2 = A.determinant(algorithm='ntl')
sage: D1 == D2
True
```

We have a special-case algorithm for 4 x 4 determinants:

```
sage: A = matrix(ZZ, 4, [1, 2, 3, 4, 4, 3, 2, 1, 0, 5, 0, 1, 9, 1, 2, 3])
sage: A.determinant()
270
```


Next we try the Linbox det. Note that we must have `proof=False`.

```
sage: A = matrix(ZZ, 5, [1, 2, 3, 4, 5, 4, 6, 3, 2, 1, 7, 9, 7, 5, 2, 1, 4, 6, 7, 8, 3, 2, 4, 6, 7])
sage: A.determinant(algorithm='linbox')
Traceback (most recent call last):
...
RuntimeError: you must pass the proof=False option to the determinant command to use LinBox
sage: A.determinant(algorithm='linbox', proof=False)
-21
sage: A._clear_cache()
sage: A.determinant()
-21
```

A bigger example:

```
sage: A = random_matrix(ZZ, 30)
sage: d = A.determinant()
sage: A._clear_cache()
sage: A.determinant(algorithm='linbox', proof=False) == d
True
```

TESTS:

This shows that we can compute determinants for all sizes up to 80. The check that the determinant of a squared matrix is a square is a sanity check that the result is probably correct:

```
sage: for s in [1..80]: # long time (6s on sage.math, 2013)
...     M = random_matrix(ZZ, s)
...     d = (M*M).determinant()
...     assert d.is_square()
```

echelon_form(algorithm='default', proof=None, include_zero_rows=True, transformation=False, D=None)

Return the echelon form of this matrix over the integers, also known as the hermit normal form (HNF).

INPUT:

- **algorithm** – String. The algorithm to use. Valid options are:
 - 'default' – Let Sage pick an algorithm (default). Up to 10 rows or columns: pari with flag 0; Up to 75 rows or columns: pari with flag 1; Larger: use padic algorithm.
 - 'padic' – an asymptotically fast p-adic modular algorithm, If your matrix has large coefficients and is small, you may also want to try this.
 - 'pari' – use PARI with flag 1
 - 'pari0' – use PARI with flag 0
 - 'pari4' – use PARI with flag 4 (use heuristic LLL)
 - 'ntl' – use NTL (only works for square matrices of full rank!)
- **proof** – (default: True); if proof=False certain determinants are computed using a randomized hybrid p-adic multimodular strategy until it stabilizes twice (instead of up to the Hadamard bound). It is *incredibly* unlikely that one would ever get an incorrect result with proof=False.
- **include_zero_rows** – (default: True) if False, don't include zero rows
- **transformation** – if given, also compute transformation matrix; only valid for padic algorithm
- **D** – (default: None) if given and the algorithm is 'ntl', then D must be a multiple of the determinant and this function will use that fact.

OUTPUT:

The Hermite normal form (=echelon form over \mathbb{Z}) of self.

EXAMPLES:

```
sage: A = MatrixSpace(ZZ, 2) ([1, 2, 3, 4])
```

```
sage: A.echelon_form()
```

```
[1 0]
```

```
[0 2]
```

```
sage: A = MatrixSpace(ZZ, 5) (range(25))
```

```
sage: A.echelon_form()
```

```
[ 5  0 -5 -10 -15]
```

```
[ 0  1  2  3  4]
```

```
[ 0  0  0  0  0]
```

```
[ 0  0  0  0  0]
```

```
[ 0  0  0  0  0]
```

Getting a transformation matrix in the nonsquare case:

```
sage: A = matrix(ZZ, 5, 3, [1..15])
```

```
sage: H, U = A.hermite_form(transformation=True, include_zero_rows=False)
```

```
sage: H
```

```
[1 2 3]
```

```
[0 3 6]
```

```
sage: U
```

```
[ 0  0  0  4 -3]
```

```
[ 0  0  0 13 -10]
```

```
sage: U*A == H
```

```
True
```

TESTS: Make sure the zero matrices are handled correctly:

```
sage: m = matrix(ZZ, 3, 3, [0]*9)
```

```
sage: m.echelon_form()
```

```
[0 0 0]
```

```
[0 0 0]
```

```
[0 0 0]
```

```
sage: m = matrix(ZZ, 3, 1, [0]*3)
```

```
sage: m.echelon_form()
```

```
[0]
```

```
[0]
```

```
[0]
```

```
sage: m = matrix(ZZ, 1, 3, [0]*3)
```

```
sage: m.echelon_form()
```

```
[0 0 0]
```

The ultimate border case!

```
sage: m = matrix(ZZ, 0, 0, [])
```

```
sage: m.echelon_form()
```

```
[]
```

Note: If 'ntl' is chosen for a non square matrix this function raises a ValueError.

Special cases: 0 or 1 rows:

```
sage: a = matrix(ZZ, 1, 2, [0, -1])
```

```
sage: a.hermite_form()
```

```
[0 1]
```

```

sage: a.pivots()
(1,)
sage: a = matrix(ZZ, 1, 2, [0, 0])
sage: a.hermite_form()
[0 0]
sage: a.pivots()
()
sage: a = matrix(ZZ, 1, 3); a
[0 0 0]
sage: a.echelon_form(include_zero_rows=False)
[]
sage: a.echelon_form(include_zero_rows=True)
[0 0 0]

```

Illustrate using various algorithms.:

```

sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='pari')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='pari0')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='pari4')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='padic')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='default')
[1 2 3]
[0 3 6]
[0 0 0]

```

The 'ntl' algorithm doesn't work on matrices that do not have full rank.:

```

sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='ntl')
Traceback (most recent call last):
...
ValueError: ntl only computes HNF for square matrices of full rank.
sage: matrix(ZZ, 3, [0] + [2..9]).hermite_form(algorithm='ntl')
[1 0 0]
[0 1 0]
[0 0 3]

```

TESTS:

This example illustrated trac 2398:

```

sage: a = matrix([(0, 0, 3), (0, -2, 2), (0, 1, 2), (0, -2, 5)])
sage: a.hermite_form()
[0 1 2]
[0 0 3]
[0 0 0]
[0 0 0]

```

Check that #12280 is fixed:

```
sage: m = matrix([(-2, 1, 9, 2, -8, 1, -3, -1, -4, -1),
...               (5, -2, 0, 1, 0, 4, -1, 1, -2, 0),
...               (-11, 3, 1, 0, -3, -2, -1, -11, 2, -2),
...               (-1, 1, -1, -2, 1, -1, -1, -1, -1, 7),
...               (-2, -1, -1, 1, 1, -2, 1, 0, 2, -4)]).stack(
...               200 * identity_matrix(ZZ, 10))
sage: matrix(ZZ,m).hermite_form(algorithm='pari', include_zero_rows=False)
[ 1  0  2  0 13  5  1 166 72 69]
[ 0  1  1  0 20  4 15 195 65 190]
[ 0  0  4  0 24  5 23 22 51 123]
[ 0  0  0  1 23  7 20 105 60 151]
[ 0  0  0  0 40  4  0 80 36 68]
[ 0  0  0  0  0 10  0 100 190 170]
[ 0  0  0  0  0  0 25  0 100 150]
[ 0  0  0  0  0  0  0 200  0  0]
[ 0  0  0  0  0  0  0  0 200  0]
[ 0  0  0  0  0  0  0  0  0 200]
sage: matrix(ZZ,m).hermite_form(algorithm='padic', include_zero_rows=False)
[ 1  0  2  0 13  5  1 166 72 69]
[ 0  1  1  0 20  4 15 195 65 190]
[ 0  0  4  0 24  5 23 22 51 123]
[ 0  0  0  1 23  7 20 105 60 151]
[ 0  0  0  0 40  4  0 80 36 68]
[ 0  0  0  0  0 10  0 100 190 170]
[ 0  0  0  0  0  0 25  0 100 150]
[ 0  0  0  0  0  0  0 200  0  0]
[ 0  0  0  0  0  0  0  0 200  0]
[ 0  0  0  0  0  0  0  0  0 200]
```

elementary_divisors (*algorithm*='pari')

Return the elementary divisors of self, in order.

Warning: This is MUCH faster than the smith_form function.

The elementary divisors are the invariants of the finite abelian group that is the cokernel of *left* multiplication of this matrix. They are ordered in reverse by divisibility.

INPUT:

- self - matrix
- algorithm - (default: 'pari')
 - 'pari': works robustly, but is slower.
 - 'linbox' - use linbox (currently off, broken)

OUTPUT: list of integers

Note: These are the invariants of the cokernel of *left* multiplication:

```
sage: M = Matrix([[3,0,1],[0,1,0]])
sage: M
[3 0 1]
[0 1 0]
sage: M.elementary_divisors()
[1, 1]
```

```
sage: M.transpose().elementary_divisors()
[1, 1, 0]
```

EXAMPLES:

```
sage: matrix(3, range(9)).elementary_divisors()
[1, 3, 0]
sage: matrix(3, range(9)).elementary_divisors(algorithm='pari')
[1, 3, 0]
sage: C = MatrixSpace(ZZ, 4) ([3, 4, 5, 6, 7, 3, 8, 10, 14, 5, 6, 7, 2, 2, 10, 9])
sage: C.elementary_divisors()
[1, 1, 1, 687]

sage: M = matrix(ZZ, 3, [1, 5, 7, 3, 6, 9, 0, 1, 2])
sage: M.elementary_divisors()
[1, 1, 6]
```

This returns a copy, which is safe to change:

```
sage: edivs = M.elementary_divisors()
sage: edivs.pop()
6
sage: M.elementary_divisors()
[1, 1, 6]
```

See Also:

[smith_form\(\)](#)

frobenius (*flag=0*, *var='x'*)

Return the Frobenius form (rational canonical form) of this matrix.

INPUT:

- *flag* – 0 (default), 1 or 2 as follows:
 - 0 – (default) return the Frobenius form of this matrix.
 - 1 – return only the elementary divisor polynomials, as polynomials in *var*.
 - 2 – return a two-components vector [F,B] where F is the Frobenius form and B is the basis change so that $M = B^{-1}FB$.
- *var* – a string (default: 'x')

ALGORITHM: uses PARI's `matfrobenius()`

EXAMPLES:

```
sage: A = MatrixSpace(ZZ, 3) (range(9))
sage: A.frobenius(0)
[ 0  0  0]
[ 1  0 18]
[ 0  1 12]
sage: A.frobenius(1)
[x^3 - 12*x^2 - 18*x]
sage: A.frobenius(1, var='y')
[y^3 - 12*y^2 - 18*y]
sage: F, B = A.frobenius(2)
sage: A == B^(-1)*F*B
True
sage: a=matrix([])
```

```
sage: a.frobenius(2)
([], [])
sage: a.frobenius(0)
[]
sage: a.frobenius(1)
[]
sage: B = random_matrix(ZZ, 2, 3)
sage: B.frobenius()
Traceback (most recent call last):
...
ArithmeticError: frobenius matrix of non-square matrix not defined.
```

AUTHORS:

- Martin Albrect (2006-04-02)

TODO: - move this to work for more general matrices than just over \mathbb{Z} . This will require fixing how PARI polynomials are coerced to Sage polynomials.

gcd()

Return the gcd of all entries of self; very fast.

EXAMPLES:

```
sage: a = matrix(ZZ, 2, [6, 15, -6, 150])
sage: a.gcd()
3
```

height()

Return the height of this matrix, i.e., the max absolute value of the entries of the matrix.

OUTPUT: A nonnegative integer.

EXAMPLE:

```
sage: a = Mat(ZZ, 3)(range(9))
sage: a.height()
8
sage: a = Mat(ZZ, 2, 3)([-17, 3, -389, 15, -1, 0]); a
[ -17   3 -389]
[  15  -1   0]
sage: a.height()
389
```

hermite_form(*algorithm*='default', *proof*=None, *include_zero_rows*=True, *transformation*=False, *D*=None)

Return the echelon form of this matrix over the integers, also known as the hermit normal form (HNF).

INPUT:

- algorithm* – String. The algorithm to use. Valid options are:
 - 'default' – Let Sage pick an algorithm (default). Up to 10 rows or columns: pari with flag 0; Up to 75 rows or columns: pari with flag 1; Larger: use padic algorithm.
 - 'padic' – an asymptotically fast p-adic modular algorithm, If your matrix has large coefficients and is small, you may also want to try this.
 - 'pari' – use PARI with flag 1
 - 'pari0' – use PARI with flag 0
 - 'pari4' – use PARI with flag 4 (use heuristic LLL)

- 'ntl' - use NTL (only works for square matrices of full rank!)

- `proof` - (default: True); if `proof=False` certain determinants are computed using a randomized hybrid p-adic multimodular strategy until it stabilizes twice (instead of up to the Hadamard bound). It is *incredibly* unlikely that one would ever get an incorrect result with `proof=False`.
- `include_zero_rows` - (default: True) if False, don't include zero rows
- `transformation` - if given, also compute transformation matrix; only valid for padic algorithm
- `D` - (default: None) if given and the algorithm is 'ntl', then D must be a multiple of the determinant and this function will use that fact.

OUTPUT:

The Hermite normal form (=echelon form over \mathbb{Z}) of self.

EXAMPLES:

```
sage: A = MatrixSpace(ZZ, 2) ([1, 2, 3, 4])
```

```
sage: A.echelon_form()
```

```
[1 0]
```

```
[0 2]
```

```
sage: A = MatrixSpace(ZZ, 5) (range(25))
```

```
sage: A.echelon_form()
```

```
[ 5  0 -5 -10 -15]
```

```
[ 0  1  2  3  4]
```

```
[ 0  0  0  0  0]
```

```
[ 0  0  0  0  0]
```

```
[ 0  0  0  0  0]
```

Getting a transformation matrix in the nonsquare case:

```
sage: A = matrix(ZZ, 5, 3, [1..15])
```

```
sage: H, U = A.hermite_form(transformation=True, include_zero_rows=False)
```

```
sage: H
```

```
[1 2 3]
```

```
[0 3 6]
```

```
sage: U
```

```
[ 0  0  0  4 -3]
```

```
[ 0  0  0 13 -10]
```

```
sage: U*A == H
```

```
True
```

TESTS: Make sure the zero matrices are handled correctly:

```
sage: m = matrix(ZZ, 3, 3, [0]*9)
```

```
sage: m.echelon_form()
```

```
[0 0 0]
```

```
[0 0 0]
```

```
[0 0 0]
```

```
sage: m = matrix(ZZ, 3, 1, [0]*3)
```

```
sage: m.echelon_form()
```

```
[0]
```

```
[0]
```

```
[0]
```

```
sage: m = matrix(ZZ, 1, 3, [0]*3)
```

```
sage: m.echelon_form()
```

```
[0 0 0]
```

The ultimate border case!

```
sage: m = matrix(ZZ, 0, 0, [])
sage: m.echelon_form()
[]
```

Note: If 'ntl' is chosen for a non square matrix this function raises a ValueError.

Special cases: 0 or 1 rows:

```
sage: a = matrix(ZZ, 1, 2, [0, -1])
sage: a.hermite_form()
[0 1]
sage: a.pivots()
(1,)
sage: a = matrix(ZZ, 1, 2, [0, 0])
sage: a.hermite_form()
[0 0]
sage: a.pivots()
()
sage: a = matrix(ZZ, 1, 3); a
[0 0 0]
sage: a.echelon_form(include_zero_rows=False)
[]
sage: a.echelon_form(include_zero_rows=True)
[0 0 0]
```

Illustrate using various algorithms.:

```
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='pari')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='pari0')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='pari4')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='padic')
[1 2 3]
[0 3 6]
[0 0 0]
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='default')
[1 2 3]
[0 3 6]
[0 0 0]
```

The 'ntl' algorithm doesn't work on matrices that do not have full rank.:

```
sage: matrix(ZZ, 3, [1..9]).hermite_form(algorithm='ntl')
Traceback (most recent call last):
...
ValueError: ntl only computes HNF for square matrices of full rank.
sage: matrix(ZZ, 3, [0] + [2..9]).hermite_form(algorithm='ntl')
[1 0 0]
[0 1 0]
[0 0 3]
```


TESTS:

This example illustrated trac 2398:

```
sage: a = matrix([(0, 0, 3), (0, -2, 2), (0, 1, 2), (0, -2, 5)])
sage: a.hermite_form()
[0 1 2]
[0 0 3]
[0 0 0]
[0 0 0]
```

Check that #12280 is fixed:

```
sage: m = matrix([(-2, 1, 9, 2, -8, 1, -3, -1, -4, -1),
...               (5, -2, 0, 1, 0, 4, -1, 1, -2, 0),
...               (-11, 3, 1, 0, -3, -2, -1, -11, 2, -2),
...               (-1, 1, -1, -2, 1, -1, -1, -1, -1, 7),
...               (-2, -1, -1, 1, 1, -2, 1, 0, 2, -4)]).stack(
...               200 * identity_matrix(ZZ, 10))
sage: matrix(ZZ,m).hermite_form(algorithm='pari', include_zero_rows=False)
[ 1  0  2  0 13  5  1 166 72 69]
[ 0  1  1  0 20  4 15 195 65 190]
[ 0  0  4  0 24  5 23  22 51 123]
[ 0  0  0  1 23  7 20 105 60 151]
[ 0  0  0  0 40  4  0  80 36 68]
[ 0  0  0  0  0 10  0 100 190 170]
[ 0  0  0  0  0  0 25  0 100 150]
[ 0  0  0  0  0  0  0 200  0  0]
[ 0  0  0  0  0  0  0  0 200  0]
[ 0  0  0  0  0  0  0  0  0 200]
sage: matrix(ZZ,m).hermite_form(algorithm='padic', include_zero_rows=False)
[ 1  0  2  0 13  5  1 166 72 69]
[ 0  1  1  0 20  4 15 195 65 190]
[ 0  0  4  0 24  5 23  22 51 123]
[ 0  0  0  1 23  7 20 105 60 151]
[ 0  0  0  0 40  4  0  80 36 68]
[ 0  0  0  0  0 10  0 100 190 170]
[ 0  0  0  0  0  0 25  0 100 150]
[ 0  0  0  0  0  0  0 200  0  0]
[ 0  0  0  0  0  0  0  0 200  0]
[ 0  0  0  0  0  0  0  0  0 200]
```

index_in_saturation (*proof=None*)

Return the index of self in its saturation.

INPUT:

- *proof* - (default: use `proof.linear_algebra()`); if False, the determinant calculations are done with `proof=False`.

OUTPUT:

- *positive integer* - the index of the row span of this matrix in its saturation

ALGORITHM: Use Hermite normal form twice to find an invertible matrix whose inverse transforms a matrix with the same row span as self to its saturation, then compute the determinant of that matrix.

EXAMPLES:

```
sage: A = matrix(ZZ, 2,3, [1..6]); A
[1 2 3]
[4 5 6]
```

```
sage: A.index_in_saturation()
3
sage: A.saturation()
[1 2 3]
[1 1 1]
```

insert_row (*index*, *row*)

Create a new matrix from self with.

INPUT:

- *index* - integer
- *row* - a vector

EXAMPLES:

```
sage: X = matrix(ZZ, 3, range(9)); X
[0 1 2]
[3 4 5]
[6 7 8]
sage: X.insert_row(1, [1, 5, -10])
[ 0 1 2]
[ 1 5 -10]
[ 3 4 5]
[ 6 7 8]
sage: X.insert_row(0, [1, 5, -10])
[ 1 5 -10]
[ 0 1 2]
[ 3 4 5]
[ 6 7 8]
sage: X.insert_row(3, [1, 5, -10])
[ 0 1 2]
[ 3 4 5]
[ 6 7 8]
[ 1 5 -10]
```

is_LLL_reduced (*delta*=None, *eta*=None)

Return True if this lattice is (δ, η) -LLL reduced. See `self.LLL` for a definition of LLL reduction.

INPUT:

- *delta* – (default: 0.99) parameter δ as described above
- *eta* – (default: 0.501) parameter η as described above

EXAMPLES:

```
sage: A = random_matrix(ZZ, 10, 10)
sage: L = A.LLL()
sage: A.is_LLL_reduced()
False
sage: L.is_LLL_reduced()
True
```

minpoly (*var*='x', *algorithm*='linbox')

INPUT:

- *var* - a variable name
- *algorithm* - 'linbox' (default) 'generic'

Note: Linbox charpoly disabled on 64-bit machines, since it hangs in many cases.

EXAMPLES:

```
sage: A = matrix(ZZ, 6, range(36))
sage: A.minpoly()
x^3 - 105*x^2 - 630*x
sage: n=6; A = Mat(ZZ, n) ([k^2 for k in range(n^2)])
sage: A.minpoly()
x^4 - 2695*x^3 - 257964*x^2 + 1693440*x
```

pivots()

Return the pivot column positions of this matrix.

OUTPUT: a tuple of Python integers: the position of the first nonzero entry in each row of the echelon form.

EXAMPLES:

```
sage: n = 3; A = matrix(ZZ, n, range(n^2)); A
[0 1 2]
[3 4 5]
[6 7 8]
sage: A.pivots()
(0, 1)
sage: A.echelon_form()
[ 3  0 -3]
[ 0  1  2]
[ 0  0  0]
```

prod_of_row_sums(cols)

Return the product of the sums of the entries in the submatrix of `self` with given columns.

INPUT:

- `cols` – a list (or set) of integers representing columns of `self`

OUTPUT: an integer

EXAMPLES:

```
sage: a = matrix(ZZ, 2, 3, [1..6]); a
[1 2 3]
[4 5 6]
sage: a.prod_of_row_sums([0, 2])
40
sage: (1+3) * (4+6)
40
sage: a.prod_of_row_sums(set([0, 2]))
40
```

randomize(density=1, x=None, y=None, distribution=None, nonzero=False)

Randomize density proportion of the entries of this matrix, leaving the rest unchanged.

The parameters are the same as the ones for the integer ring's `random_element` function.

If `x` and `y` are given, randomized entries of this matrix have to be between `x` and `y` and have density 1.

INPUT:

- `self` – a mutable matrix over `ZZ`

- `density` - a float between 0 and 1
- `x`, `y` - if not `None`, these are passed to the `ZZ.random_element` function as the upper and lower endpoints in the uniform distribution
- `distribution` - would also be passed into `ZZ.random_element` if given
- `nonzero` - bool (default: `False`); whether the new entries are guaranteed to be zero

OUTPUT:

- `None`, the matrix is modified in-place

EXAMPLES:

```
sage: A = matrix(ZZ, 2, 3, [1..6]); A
[1 2 3]
[4 5 6]
sage: A.randomize()
sage: A
[-8  2  0]
[ 0  1 -1]
sage: A.randomize(x=-30,y=30)
sage: A
[  5 -19  24]
[ 24  23  -9]
```

rank()

Return the rank of this matrix.

OUTPUT:

- nonnegative integer - the rank

Note: The rank is cached.

ALGORITHM: First check if the matrix has maxim possible rank by working modulo one random prime. If not call `LinBox`'s rank function.

EXAMPLES:

```
sage: a = matrix(ZZ, 2, 3, [1..6]); a
[1 2 3]
[4 5 6]
sage: a.rank()
2
sage: a = matrix(ZZ, 3, 3, [1..9]); a
[1 2 3]
[4 5 6]
[7 8 9]
sage: a.rank()
2
```

Here's a bigger example - the rank is of course still 2:

```
sage: a = matrix(ZZ, 100, [1..100^2]); a.rank()
2
```

rational_reconstruction(N)

Use rational reconstruction to lift self to a matrix over the rational numbers (if possible), where we view self as a matrix modulo `N`.

INPUT:

- `N` - an integer

OUTPUT:

- `matrix` - over `QQ` or raise a `ValueError`

EXAMPLES: We create a random 4x4 matrix over `ZZ`.

```
sage: A = matrix(ZZ, 4, [4, -4, 7, 1, -1, 1, -1, -12, -1, -1, 1, -1, -3, 1, 5, -1])
```

There isn't a unique rational reconstruction of it:

```
sage: A.rational_reconstruction(11)
Traceback (most recent call last):
...
ValueError: Rational reconstruction of 4 (mod 11) does not exist.
```

We throw in a denominator and reduce the matrix modulo 389 - it does rationally reconstruct.

```
sage: B = (A/3 % 389).change_ring(ZZ)
sage: B.rational_reconstruction(389) == A/3
True
```

TEST:

Check that ticket #9345 is fixed:

```
sage: A = random_matrix(ZZ, 3, 3)
sage: A.rational_reconstruction(0)
Traceback (most recent call last):
...
ZeroDivisionError: The modulus cannot be zero
```

saturation ($p=0$, $proof=None$, $max_dets=5$)

Return a saturation matrix of self, which is a matrix whose rows span the saturation of the row span of self. This is not unique.

The saturation of a \mathbf{Z} module M embedded in \mathbf{Z}^n is the a module S that contains M with finite index such that \mathbf{Z}^n/S is torsion free. This function takes the row span M of self, and finds another matrix of full rank with row span the saturation of M .

INPUT:

- `p` - (default: 0); if nonzero given, saturate only at the prime p , i.e., return a matrix whose row span is a \mathbf{Z} -module S that contains self and such that the index of S in its saturation is coprime to p . If p is `None`, return full saturation of self.
- `proof` - (default: `use proof.linear_algebra()`); if `False`, the determinant calculations are done with `proof=False`.
- `max_dets` - (default: 5); technical parameter - max number of determinant to compute when bounding prime divisor of self in its saturation.

OUTPUT:

- `matrix` - a matrix over `ZZ`

Note: The result is *not* cached.

ALGORITHM: 1. Replace input by a matrix of full rank got from a subset of the rows. 2. Divide out any common factors from rows. 3. Check `max_dets` random dets of submatrices to see if their GCD (with

p) is 1 - if so matrix is saturated and we're done. 4. Finally, use that if A is a matrix of full rank, then $\text{hnf}(\text{transpose}(A))^{-1} * A$ is a saturation of A .

EXAMPLES:

```
sage: A = matrix(ZZ, 3, 5, [-51, -1509, -71, -109, -593, -19, -341, 4, 86, 98, 0, -246, -11,
sage: A.echelon_form()
[ 1      5      2262      20364      56576]
[  0      6      35653      320873      891313]
[  0      0      42993      386937      1074825]
sage: S = A.saturation(); S
[ -51 -1509   -71  -109  -593]
[ -19  -341    4    86    98]
[  35   994   43    51   347]
```

Notice that the saturation spans a different module than A .

```
sage: S.echelon_form()
[ 1  2  0  8 32]
[ 0  3  0 -2 -6]
[ 0  0  1  9 25]
sage: V = A.row_space(); W = S.row_space()
sage: V.is_submodule(W)
True
sage: V.index_in(W)
85986
sage: V.index_in_saturation()
85986
```

We illustrate each option:

```
sage: S = A.saturation(p=2)
sage: S = A.saturation(proof=False)
sage: S = A.saturation(max_dets=2)
```

smith_form()

Returns matrices S , U , and V such that $S = U * \text{self} * V$, and S is in Smith normal form. Thus S is diagonal with diagonal entries the ordered elementary divisors of S .

Warning: The `elementary_divisors` function, which returns the diagonal entries of S , is VASTLY faster than this function.

The elementary divisors are the invariants of the finite abelian group that is the cokernel of this matrix. They are ordered in reverse by divisibility.

EXAMPLES:

```
sage: A = MatrixSpace(IntegerRing(), 3)(range(9))
sage: D, U, V = A.smith_form()
sage: D
[1 0 0]
[0 3 0]
[0 0 0]
sage: U
[ 0  1  0]
[ 0 -1  1]
[-1  2 -1]
sage: V
[-1  4  1]
[ 1 -3 -2]
```

```
[ 0  0  1]
sage: U*A*V
[1 0 0]
[0 3 0]
[0 0 0]
```

It also makes sense for nonsquare matrices:

```
sage: A = Matrix(ZZ, 3, 2, range(6))
sage: D, U, V = A.smith_form()
sage: D
[1 0]
[0 2]
[0 0]
sage: U
[ 0  1  0]
[ 0 -1  1]
[-1  2 -1]
sage: V
[-1  3]
[ 1 -2]
sage: U * A * V
[1 0]
[0 2]
[0 0]
```

Empty matrices are handled sensibly (see trac #3068):

```
sage: m = MatrixSpace(ZZ, 2, 0)(0); d,u,v = m.smith_form(); u*m*v == d
True
sage: m = MatrixSpace(ZZ, 0, 2)(0); d,u,v = m.smith_form(); u*m*v == d
True
sage: m = MatrixSpace(ZZ, 0, 0)(0); d,u,v = m.smith_form(); u*m*v == d
True
```

See Also:

`elementary_divisors()`

stack (*bottom*, *subdivide=False*)

Return the matrix self on top of bottom: [self] [bottom]

EXAMPLES:

```
sage: M = Matrix(ZZ, 2, 3, range(6))
sage: N = Matrix(ZZ, 1, 3, [10,11,12])
sage: M.stack(N)
[ 0  1  2]
[ 3  4  5]
[10 11 12]
```

A vector may be stacked below a matrix.

```
sage: A = matrix(ZZ, 2, 4, range(8))
sage: v = vector(ZZ, 4, range(4))
sage: A.stack(v)
[0 1 2 3]
[4 5 6 7]
[0 1 2 3]
```

The `subdivide` option will add a natural subdivision between `self` and

bottom. For more details about how subdivisions are managed when stacking, see `sage.matrix.matrix1.Matrix.stack()`.

```
sage: A = matrix(ZZ, 3, 4, range(12))
```

```
sage: B = matrix(ZZ, 2, 4, range(8))
```

```
sage: A.stack(B, subdivide=True)
```

```
[ 0  1  2  3]
```

```
[ 4  5  6  7]
```

```
[ 8  9 10 11]
```

```
[-----]
```

```
[ 0  1  2  3]
```

```
[ 4  5  6  7]
```

TESTS:

Stacking a dense matrix atop a sparse one should work:

```
sage: M = Matrix(ZZ, 2, 3, range(6))
```

```
sage: M.is_sparse()
```

```
False
```

```
sage: N = diagonal_matrix([10,11,12], sparse=True)
```

```
sage: N.is_sparse()
```

```
True
```

```
sage: P = M.stack(N); P
```

```
[ 0  1  2]
```

```
[ 3  4  5]
```

```
[10  0  0]
```

```
[ 0 11  0]
```

```
[ 0  0 12]
```

```
sage: P.is_sparse()
```

```
False
```

`symplectic_form()`

Find a symplectic basis for self if self is an anti-symmetric, alternating matrix.

Returns a pair (F, C) such that the rows of C form a symplectic basis for self and $F = C * \text{self} * C.\text{transpose}()$.

Raises a `ValueError` if self is not anti-symmetric, or self is not alternating.

Anti-symmetric means that $M = -M^t$. Alternating means that the diagonal of M is identically zero.

A symplectic basis is a basis of the form $e_1, \dots, e_j, f_1, \dots, f_j, z_1, \dots, z_k$ such that

- $z_i M v^t = 0$ for all vectors v
- $e_i M e_j^t = 0$ for all i, j
- $f_i M f_j^t = 0$ for all i, j
- $e_i M f_i^t = 1$ for all i
- $e_i M f_j^t = 0$ for all i not equal j .

The ordering for the factors $d_i | d_{i+1}$ and for the placement of zeroes was chosen to agree with the output of `smith_form`.

See the example for a pictorial description of such a basis.

EXAMPLES:

```
sage: E = matrix(ZZ, 5, 5, [0, 14, 0, -8, -2, -14, 0, -3, -11, 4, 0, 3, 0, 0, 0, 8, 11, 0, 0, 0])
```

```
[ 0  14  0 -8 -2]
```

```
[-14  0 -3 -11  4]
```



```

[ 0  3  0  0  0]
[ 8 11  0  0  8]
[ 2 -4  0 -8  0]
sage: F, C = E.symplectic_form()
sage: F
[ 0  0  1  0  0]
[ 0  0  0  2  0]
[-1  0  0  0  0]
[ 0 -2  0  0  0]
[ 0  0  0  0  0]
sage: F == C * E * C.transpose()
True
sage: E.smith_form()[0]
[1 0 0 0 0]
[0 1 0 0 0]
[0 0 2 0 0]
[0 0 0 2 0]
[0 0 0 0 0]

```

transpose()

Returns the transpose of self, without changing self.

EXAMPLES:

We create a matrix, compute its transpose, and note that the original matrix is not changed.

```

sage: A = matrix(ZZ, 2, 3, xrange(6))
sage: type(A)
<type 'sage.matrix.matrix_integer_dense.Matrix_integer_dense'>
sage: B = A.transpose()
sage: print B
[0 3]
[1 4]
[2 5]
sage: print A
[0 1 2]
[3 4 5]

```

.T is a convenient shortcut for the transpose:

```

sage: A.T
[0 3]
[1 4]
[2 5]

sage: A.subdivide(None, 1); A
[0|1 2]
[3|4 5]
sage: A.transpose()
[0 3]
[---]
[1 4]
[2 5]

```

```
sage.matrix.matrix_integer_dense.tune_multiplication(k, nmin=10, nmax=200, bit-
min=2, bitmax=64)
```

Compare various multiplication algorithms.

INPUT:

- k - integer; affects numbers of trials
- nmin - integer; smallest matrix to use
- nmax - integer; largest matrix to use
- bitmin - integer; smallest bitsize
- bitmax - integer; largest bitsize

OUTPUT:

- prints what doing then who wins - multimodular or classical

EXAMPLES:

```
sage: from sage.matrix.matrix_integer_dense import tune_multiplication
sage: tune_multiplication(2, nmin=10, nmax=60, bitmin=2, bitmax=8)
10 2 0.2
...
```

DENSE MATRICES OVER THE RATIONAL FIELD

Dense matrices over the rational field

EXAMPLES:

We create a 3x3 matrix with rational entries and do some operations with it.

```
sage: a = matrix(QQ, 3, 3, [1, 2/3, -4/5, 1, 1, 1, 8, 2, -3/19]); a
[ 1  2/3 -4/5]
[ 1  1    1]
[ 8  2 -3/19]
sage: a.det()
2303/285
sage: a.charpoly()
x^3 - 35/19*x^2 + 1259/285*x - 2303/285
sage: b = a^(-1); b
[ -615/2303  -426/2303   418/2303]
[ 2325/2303  1779/2303  -513/2303]
[-1710/2303   950/2303    95/2303]
sage: b.det()
285/2303
sage: a == b
False
sage: a < b
False
sage: b < a
True
sage: a > b
True
sage: a*b
[1 0 0]
[0 1 0]
[0 0 1]
```

TESTS:

```
sage: a = matrix(QQ, 2, range(4), sparse=False)
sage: TestSuite(a).run()
```

```
class sage.matrix.matrix_rational_dense.MatrixWindow
    Bases: object
    x.__init__(...) initializes x; see help(type(x)) for signature
```

class sage.matrix.matrix_rational_dense.**Matrix_rational_dense**

Bases: sage.matrix.matrix_dense.Matrix_dense

antitranspose()

Returns the antitranspose of self, without changing self.

EXAMPLES:

```
sage: A = matrix(QQ, 2, 3, range(6))
sage: type(A)
<type 'sage.matrix.matrix_rational_dense.Matrix_rational_dense'>
sage: A.antitranspose()
[5 2]
[4 1]
[3 0]
sage: A
[0 1 2]
[3 4 5]

sage: A.subdivide(1, 2); A
[0 1|2]
[---+--]
[3 4|5]
sage: A.antitranspose()
[5|2]
[-+-]
[4|1]
[3|0]
```

change_ring(R)

Create the matrix over R with entries the entries of self coerced into R.

EXAMPLES:

```
sage: a = matrix(QQ, 2, [1/2, -1, 2, 3])
sage: a.change_ring(GF(3))
[2 2]
[2 0]
sage: a.change_ring(ZZ)
Traceback (most recent call last):
...
TypeError: matrix has denominators so can't change to ZZ.
sage: b = a.change_ring(QQ['x']); b
[1/2 -1]
[ 2 3]
sage: b.parent()
Full MatrixSpace of 2 by 2 dense matrices over Univariate Polynomial Ring in x over Rational
```

TESTS:

Make sure that subdivisions are preserved when changing rings:

```
sage: a = matrix(QQ, 3, range(9))
sage: a.subdivide(2, 1); a
[0|1 2]
[3|4 5]
[-+---]
[6|7 8]
sage: a.change_ring(ZZ).change_ring(QQ)
[0|1 2]
[3|4 5]
```

```

[+----]
[6|7 8]
sage: a.change_ring(GF(3))
[0|1 2]
[0|1 2]
[+----]
[0|1 2]

```

charpoly (*var*='x', *algorithm*='linbox')

Return the characteristic polynomial of this matrix.

INPUT:

- *var* - 'x' (string)
- *algorithm* - 'linbox' (default) or 'generic'

OUTPUT: a polynomial over the rational numbers.

EXAMPLES:

```

sage: a = matrix(QQ, 3, [4/3, 2/5, 1/5, 4, -3/2, 0, 0, -2/3, 3/4])
sage: f = a.charpoly(); f
x^3 - 7/12*x^2 - 149/40*x + 97/30
sage: f(a)
[0 0 0]
[0 0 0]
[0 0 0]

```

TESTS:

The cached polynomial should be independent of the *var* argument ([trac ticket #12292](#)). We check (indirectly) that the second call uses the cached value by noting that its result is not cached:

```

sage: M = MatrixSpace(QQ, 2)
sage: A = M(range(0, 2^2))
sage: type(A)
<type 'sage.matrix.matrix_rational_dense.Matrix_rational_dense'>
sage: A.charpoly('x')
x^2 - 3*x - 2
sage: A.charpoly('y')
y^2 - 3*y - 2
sage: A._cache['charpoly_linbox']
x^2 - 3*x - 2

```

column (*i*, *from_list*=False)

Return the *i*-th column of this matrix as a dense vector.

INPUT:

- *i* - integer
- *from_list* - ignored

EXAMPLES:

```

sage: matrix(QQ, 2, [1/5, -2/3, 3/4, 4/9]).column(1)
(-2/3, 4/9)
sage: matrix(QQ, 2, [1/5, -2/3, 3/4, 4/9]).column(1, from_list=True)
(-2/3, 4/9)
sage: matrix(QQ, 2, [1/5, -2/3, 3/4, 4/9]).column(-1)
(-2/3, 4/9)

```

```
sage: matrix(QQ, 2, [1/5, -2/3, 3/4, 4/9]).column(-2)
(1/5, 3/4)
```

decomposition (*is_diagonalizable=False*, *dual=False*, *algorithm='default'*, *height_guess=None*, *proof=None*)

Returns the decomposition of the free module on which this matrix A acts from the right (i.e., the action is x goes to $x A$), along with whether this matrix acts irreducibly on each factor. The factors are guaranteed to be sorted in the same way as the corresponding factors of the characteristic polynomial.

Let A be the matrix acting from the on the vector space V of column vectors. Assume that A is square. This function computes maximal subspaces W_1, \dots, W_n corresponding to Galois conjugacy classes of eigenvalues of A . More precisely, let $f(X)$ be the characteristic polynomial of A . This function computes the subspace $W_i = \ker(g_i(A)^n)$, where $g_i(X)$ is an irreducible factor of $f(X)$ and $g_i(X)$ exactly divides $f(X)$. If the optional parameter *is_diagonalizable* is *True*, then we let $W_i = \ker(g(A))$, since then we know that $\ker(g(A)) = \ker(g(A)^n)$.

If *dual* is *True*, also returns the corresponding decomposition of V under the action of the transpose of A . The factors are guaranteed to correspond.

INPUT:

- *is_diagonalizable* - ignored
- *dual* - whether to also return decompositions for the dual
- *algorithm*
 - ‘default’: use default algorithm for computing Echelon forms
 - ‘multimodular’: much better if the answers factors have small height
- *height_guess* - positive integer; only used by the multimodular algorithm
- *proof* - bool or None (default: None, see `proof.linear_algebra` or `sage.structure.proof`); only used by the multimodular algorithm. Note that the Sage global default is *proof=True*.

Note: IMPORTANT: If you expect that the subspaces in the answer are spanned by vectors with small height coordinates, use *algorithm='multimodular'* and *height_guess=1*; this is potentially much faster than the default. If you know for a fact the answer will be very small, use *algorithm='multimodular'*, *height_guess=bound on height*, *proof=False*.

You can get very very fast decomposition with *proof=False*.

EXAMPLES:

```
sage: a = matrix(QQ, 3, [1..9])
sage: a.decomposition()
[
  (Vector space of degree 3 and dimension 1 over Rational Field
  Basis matrix:
  [ 1 -2  1], True),
  (Vector space of degree 3 and dimension 2 over Rational Field
  Basis matrix:
  [ 1  0 -1]
  [ 0  1  2], True)
]
```

denominator ()

Return the denominator of this matrix.

OUTPUT: a Sage Integer

EXAMPLES:

```
sage: b = matrix(QQ, 2, range(6)); b[0,0] = -5007/293; b
[-5007/293      1      2]
[      3      4      5]
sage: b.denominator()
293
```

determinant (*algorithm='default', proof=None*)

Return the determinant of this matrix.

INPUT:

- *proof* - bool or None; if None use `proof.linear_algebra()`; only relevant for the p-adic algorithm.
- *algorithm*:
 - “default” – use PARI for up to 7 rows, then use integer
 - “pari” – use PARI
 - “integer” – clear denominators and call `det` on integer matrix

Note: It would be *VERY VERY* hard for `det` to fail even with `proof=False`.

ALGORITHM: Clear denominators and call the integer determinant function.

EXAMPLES:

```
sage: m = matrix(QQ, 3, [1, 2/3, 4/5, 2, 2, 2, 5, 3, 2/5])
sage: m.determinant()
-34/15
sage: m.charpoly()
x^3 - 17/5*x^2 - 122/15*x + 34/15
```

echelon_form (*algorithm='default', height_guess=None, proof=None, **kws*)

INPUT:

- *algorithm*
 - “default” (default): use heuristic choice
 - “p-adic”: an algorithm based on the IML p-adic solver.
 - “multimodular”: uses a multimodular algorithm the uses `linbox` modulo many primes.
 - “classical”: just clear each column using Gauss elimination
- *height_guess, **kws* - all passed to the multimodular algorithm; ignored by the p-adic algorithm.
- *proof* - bool or None (default: None, see `proof.linear_algebra` or `sage.structure.proof`). Passed to the multimodular algorithm. Note that the Sage global default is `proof=True`.

OUTPUT: the reduced row echelon form of self.

EXAMPLES:

```
sage: a = matrix(QQ, 4, range(16)); a[0,0] = 1/19; a[0,1] = 1/5; a
[1/19 1/5 2 3]
[ 4 5 6 7]
[ 8 9 10 11]
[ 12 13 14 15]
```

```

sage: a.echelon_form()
[ 1 0 0 0 -76/157]
[ 0 1 0 0 -5/157]
[ 0 0 1 238/157]
[ 0 0 0 0 0]
sage: a.echelon_form(algorithm='multimodular')
[ 1 0 0 0 -76/157]
[ 0 1 0 0 -5/157]
[ 0 0 1 238/157]
[ 0 0 0 0 0]

```

The result is an immutable matrix, so if you want to modify the result then you need to make a copy. This checks that Trac #10543 is fixed.

```

sage: A = matrix(QQ, 2, range(6))
sage: E = A.echelon_form()
sage: E.is_mutable()
False
sage: F = copy(E)
sage: F[0,0] = 50
sage: F
[50 0 -1]
[ 0 1 2]

```

echelonize (*algorithm='default', height_guess=None, proof=None, **kws*)

INPUT:

- *algorithm*
- 'default' (default): use heuristic choice
- 'padic': an algorithm based on the IML p-adic solver.
- 'multimodular': uses a multimodular algorithm the uses linbox modulo many primes.
- 'classical': just clear each column using Gauss elimination
- *height_guess, **kws* - all passed to the multimodular algorithm; ignored by the p-adic algorithm.
- *proof* - bool or None (default: None, see `proof.linear_algebra` or `sage.structure.proof`). Passed to the multimodular algorithm. Note that the Sage global default is `proof=True`.

OUTPUT:

- *matrix* - the reduced row echelon form of self.

EXAMPLES:

```

sage: a = matrix(QQ, 4, range(16)); a[0,0] = 1/19; a[0,1] = 1/5; a
[1/19 1/5 2 3]
[ 4 5 6 7]
[ 8 9 10 11]
[ 12 13 14 15]
sage: a.echelonize(); a
[ 1 0 0 0 -76/157]
[ 0 1 0 0 -5/157]
[ 0 0 1 238/157]
[ 0 0 0 0 0]

```



```

sage: a = matrix(QQ, 4, range(16)); a[0,0] = 1/19; a[0,1] = 1/5
sage: a.echelonize(algorithm='multimodular'); a
[   1   0   0 -76/157]
[   0   1   0 -5/157]
[   0   0   1 238/157]
[   0   0   0   0]

```

height()

Return the height of this matrix, which is the maximum of the absolute values of all numerators and denominators of entries in this matrix.

OUTPUT: an Integer

EXAMPLES:

```

sage: b = matrix(QQ, 2, range(6)); b[0,0] = -5007/293; b
[-5007/293   1   2]
[   3   4   5]
sage: b.height()
5007

```

minpoly(var='x', algorithm='linbox')

Return the minimal polynomial of this matrix.

INPUT:

- var - 'x' (string)
- algorithm - 'linbox' (default) or 'generic'

OUTPUT: a polynomial over the rational numbers.

EXAMPLES:

```

sage: a = matrix(QQ, 3, [4/3, 2/5, 1/5, 4, -3/2, 0, 0, -2/3, 3/4])
sage: f = a.minpoly(); f
x^3 - 7/12*x^2 - 149/40*x + 97/30
sage: a = Mat(ZZ, 4) (range(16))
sage: f = a.minpoly(); f.factor()
x * (x^2 - 30*x - 80)
sage: f(a) == 0
True

sage: a = matrix(QQ, 4, [1..4^2])
sage: factor(a.minpoly())
x * (x^2 - 34*x - 80)
sage: factor(a.minpoly('y'))
y * (y^2 - 34*y - 80)
sage: factor(a.charpoly())
x^2 * (x^2 - 34*x - 80)
sage: b = matrix(QQ, 4, [-1, 2, 2, 0, 0, 4, 2, 2, 0, 0, -1, -2, 0, -4, 0, 4])
sage: a = matrix(QQ, 4, [1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 5, 0, 0, 0, 0, 5])
sage: c = b^(-1)*a*b
sage: factor(c.minpoly())
(x - 5) * (x - 1)^2
sage: factor(c.charpoly())
(x - 5)^2 * (x - 1)^2

```

prod_of_row_sums(cols)

randomize (*density=1, num_bound=2, den_bound=2, distribution=None, nonzero=False*)

Randomize density proportion of the entries of this matrix, leaving the rest unchanged.

If *x* and *y* are given, randomized entries of this matrix have numerators and denominators bounded by *x* and *y* and have density 1.

INPUT:

- *density* - number between 0 and 1 (default: 1)
- *num_bound* - numerator bound (default: 2)
- *den_bound* - denominator bound (default: 2)
- *distribution* - None or '1/n' (default: None); if '1/n' then *num_bound*, *den_bound* are ignored and numbers are chosen using the GMP function `mpq_randomize_entry_recip_uniform`
- *nonzero* - Bool (default: False); whether the new entries are forced to be non-zero

OUTPUT:

- None, the matrix is modified in-space

EXAMPLES:

```
sage: a = matrix(QQ,2,4); a.randomize(); a
[ 0 -1  2 -2]
[ 1 -1  2  1]
sage: a = matrix(QQ,2,4); a.randomize(density=0.5); a
[ -1 -2  0  0]
[  0  0 1/2  0]
sage: a = matrix(QQ,2,4); a.randomize(num_bound=100, den_bound=100); a
[ 14/27  21/25  43/42 -48/67]
[-19/55  64/67 -11/51   76]
sage: a = matrix(QQ,2,4); a.randomize(distribution='1/n'); a
[      3      1/9      1/2      1/4]
[      1     1/39      2 -1955/2]
```

rank ()

Return the rank of this matrix.

EXAMPLES:: `sage: matrix(QQ,3,[1..9]).rank()` 2 `sage: matrix(QQ,100,[1..100^2]).rank()` 2

row (*i, from_list=False*)

Return the *i*-th row of this matrix as a dense vector.

INPUT:

- *i* - integer
- *from_list* - ignored

EXAMPLES:

```
sage: matrix(QQ,2,[1/5,-2/3,3/4,4/9]).row(1)
(3/4, 4/9)
sage: matrix(QQ,2,[1/5,-2/3,3/4,4/9]).row(1,from_list=True)
(3/4, 4/9)
sage: matrix(QQ,2,[1/5,-2/3,3/4,4/9]).row(-2)
(1/5, -2/3)
```

set_row_to_multiple_of_row (*i, j, s*)

Set row *i* equal to *s* times row *j*.

EXAMPLES:

```
sage: a = matrix(QQ, 2, 3, range(6)); a
[0 1 2]
[3 4 5]
sage: a.set_row_to_multiple_of_row(1, 0, -3)
sage: a
[ 0  1  2]
[ 0 -3 -6]
```

transpose()

Returns the transpose of self, without changing self.

EXAMPLES:

We create a matrix, compute its transpose, and note that the original matrix is not changed.

```
sage: A = matrix(QQ, 2, 3, xrange(6))
sage: type(A)
<type 'sage.matrix.matrix_rational_dense.Matrix_rational_dense'>
sage: B = A.transpose()
sage: print B
[0 3]
[1 4]
[2 5]
sage: print A
[0 1 2]
[3 4 5]
```

.T is a convenient shortcut for the transpose:

```
sage: print A.T
[0 3]
[1 4]
[2 5]

sage: A.subdivide(None, 1); A
[0|1 2]
[3|4 5]
sage: A.transpose()
[0 3]
[---]
[1 4]
[2 5]
```


DENSE MATRICES USING A NUMPY BACKEND.

Dense matrices using a NumPy backend.

This serves as a base class for dense matrices over Real Double Field and Complex Double Field.

AUTHORS:

- Jason Grout, Sep 2008: switch to NumPy backend, factored out the `Matrix_double_dense` class
- Josh Kantor
- William Stein: many bug fixes and touch ups.

EXAMPLES:

```
sage: b=Mat(RDF,2,3).basis()
sage: b[0]
[1.0 0.0 0.0]
[0.0 0.0 0.0]
```

We deal with the case of zero rows or zero columns:

```
sage: m = MatrixSpace(RDF,0,3)
sage: m.zero_matrix()
[]
```

TESTS:

```
sage: a = matrix(RDF,2,range(4), sparse=False)
sage: TestSuite(a).run()
sage: a = matrix(CDF,2,range(4), sparse=False)
sage: TestSuite(a).run()
```

```
class sage.matrix.matrix_double_dense.Matrix_double_dense
    Bases: sage.matrix.matrix_dense.Matrix_dense
```

Base class for matrices over the Real Double Field and the Complex Double Field. These are supposed to be fast matrix operations using C doubles. Most operations are implemented using numpy which will call the underlying BLAS on the system.

This class cannot be instantiated on its own. The numpy matrix creation depends on several variables that are set in the subclasses.

EXAMPLES:

```
sage: m = Matrix(RDF, [[1,2],[3,4]])
sage: m**2
[ 7.0 10.0]
[15.0 22.0]
sage: n= m^(-1); n
[-2.0  1.0]
[ 1.5 -0.5]
```

LU()

Returns a decomposition of the (row-permuted) matrix as a product of a lower-triangular matrix (“L”) and an upper-triangular matrix (“U”).

OUTPUT:

For an $m \times n$ matrix A this method returns a triple of immutable matrices P , L , U such that

- $P \cdot A = L \cdot U$
- P is a square permutation matrix, of size $m \times m$, so is all zeroes, but with exactly a single one in each row and each column.
- L is lower-triangular, square of size $m \times m$, with every diagonal entry equal to one.
- U is upper-triangular with size $m \times n$, i.e. entries below the “diagonal” are all zero.

The computed decomposition is cached and returned on subsequent calls, thus requiring the results to be immutable.

Effectively, P permutes the rows of A . Then L can be viewed as a sequence of row operations on this matrix, where each operation is adding a multiple of a row to a subsequent row. There is no scaling (thus 1’s on the diagonal of L) and no row-swapping (P does that). As a result U is close to being the result of Gaussian-elimination. However, round-off errors can make it hard to determine the zero entries of U .

Note: Sometimes this decomposition is written as $A=P \cdot L \cdot U$, where P represents the inverse permutation and is the matrix inverse of the P returned by this method. The computation of this matrix inverse can be accomplished quickly with just a transpose as the matrix is orthogonal/unitary.

EXAMPLES:

```
sage: m = matrix(RDF, 4, range(16))
sage: P,L,U = m.LU()
sage: P*m
[12.0 13.0 14.0 15.0]
[ 0.0  1.0  2.0  3.0]
[ 8.0  9.0 10.0 11.0]
[ 4.0  5.0  6.0  7.0]
sage: L*U
[12.0 13.0 14.0 15.0]
[ 0.0  1.0  2.0  3.0]
[ 8.0  9.0 10.0 11.0]
[ 4.0  5.0  6.0  7.0]
```

Trac 10839 made this routine available for rectangular matrices.

```
sage: A = matrix(RDF, 5, 6, range(30)); A
[ 0.0  1.0  2.0  3.0  4.0  5.0]
[ 6.0  7.0  8.0  9.0 10.0 11.0]
[12.0 13.0 14.0 15.0 16.0 17.0]
[18.0 19.0 20.0 21.0 22.0 23.0]
[24.0 25.0 26.0 27.0 28.0 29.0]
```

```

sage: P, L, U = A.LU()
sage: P
[0.0 0.0 0.0 0.0 1.0]
[1.0 0.0 0.0 0.0 0.0]
[0.0 0.0 1.0 0.0 0.0]
[0.0 0.0 0.0 1.0 0.0]
[0.0 1.0 0.0 0.0 0.0]
sage: L.zero_at(0) # Use zero_at(0) to get rid of signed zeros
[ 1.0  0.0  0.0  0.0  0.0]
[ 0.0  1.0  0.0  0.0  0.0]
[ 0.5  0.5  1.0  0.0  0.0]
[0.75 0.25  0.0  1.0  0.0]
[0.25 0.75  0.0  0.0  1.0]
sage: U.zero_at(0) # Use zero_at(0) to get rid of signed zeros
[24.0 25.0 26.0 27.0 28.0 29.0]
[ 0.0  1.0  2.0  3.0  4.0  5.0]
[ 0.0  0.0  0.0  0.0  0.0  0.0]
[ 0.0  0.0  0.0  0.0  0.0  0.0]
[ 0.0  0.0  0.0  0.0  0.0  0.0]
sage: P*A-L*U
[0.0 0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0 0.0]
sage: P.transpose()*L*U
[ 0.0  1.0  2.0  3.0  4.0  5.0]
[ 6.0  7.0  8.0  9.0 10.0 11.0]
[12.0 13.0 14.0 15.0 16.0 17.0]
[18.0 19.0 20.0 21.0 22.0 23.0]
[24.0 25.0 26.0 27.0 28.0 29.0]

```

Trivial cases return matrices of the right size and characteristics.

```

sage: A = matrix(RDF, 5, 0, entries=0)
sage: P, L, U = A.LU()
sage: P.parent()
Full MatrixSpace of 5 by 5 dense matrices over Real Double Field
sage: L.parent()
Full MatrixSpace of 5 by 5 dense matrices over Real Double Field
sage: U.parent()
Full MatrixSpace of 5 by 0 dense matrices over Real Double Field
sage: P*A-L*U
[]

```

The results are immutable since they are cached.

```

sage: P, L, U = matrix(RDF, 2, 2, range(4)).LU()
sage: L[0,0] = 0
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead (i.e., use copy(M) to change a
sage: P[0,0] = 0
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead (i.e., use copy(M) to change a
sage: U[0,0] = 0
Traceback (most recent call last):
...

```

ValueError: matrix is immutable; please change a copy instead (i.e., use `copy(M)` to change a

LU_valid()

Returns True if the LU form of this matrix has already been computed.

EXAMPLES:

```
sage: A = random_matrix(RDF, 3) ; A.LU_valid()
False
sage: P, L, U = A.LU()
sage: A.LU_valid()
True
```

QR()

Returns a factorization into a unitary matrix and an upper-triangular matrix.

INPUT:

Any matrix over RDF or CDF.

OUTPUT:

Q, R – a pair of matrices such that if A is the original matrix, then

$$A = QR, \quad Q^*Q = I$$

where R is upper-triangular. Q^* is the conjugate-transpose in the complex case, and just the transpose in the real case. So Q is a unitary matrix (or rather, orthogonal, in the real case), or equivalently Q has orthogonal columns. For a matrix of full rank this factorization is unique up to adjustments via multiples of rows and columns by multiples with scalars having modulus 1. So in the full-rank case, R is unique if the diagonal entries are required to be positive real numbers.

The resulting decomposition is cached.

ALGORITHM:

Calls “linalg.qr” from SciPy, which is in turn an interface to LAPACK routines.

EXAMPLES:

Over the reals, the inverse of Q is its transpose, since including a conjugate has no effect. In the real case, we say Q is orthogonal.

```
sage: A = matrix(RDF, [[-2, 0, -4, -1, -1],
...                    [-2, 1, -6, -3, -1],
...                    [1, 1, 7, 4, 5],
...                    [3, 0, 8, 3, 3],
...                    [-1, 1, -6, -6, 5]])
sage: Q, R = A.QR()
```

At this point, Q is only well-defined up to the signs of its columns, and similarly for R and its rows, so we normalize them:

```
sage: Qnorm = Q._normalize_columns()
sage: Rnorm = R._normalize_rows()
sage: Qnorm.round(6).zero_at(10^-6)
[ 0.458831  0.126051  0.381212  0.394574  0.68744]
[ 0.458831 -0.47269 -0.051983 -0.717294  0.220963]
[-0.229416 -0.661766  0.661923  0.180872 -0.196411]
[-0.688247 -0.189076 -0.204468 -0.09663  0.662889]
[ 0.229416 -0.535715 -0.609939  0.536422 -0.024551]
sage: Rnorm.round(6).zero_at(10^-6)
```



```
[ 4.358899 -0.458831 13.076697  6.194225  2.982405]
[      0.0  1.670172  0.598741 -1.29202  6.207997]
[      0.0      0.0  5.444402  5.468661 -0.682716]
[      0.0      0.0      0.0  1.027626 -3.6193]
[      0.0      0.0      0.0      0.0  0.024551]
```

```
sage: (Q*Q.transpose()).zero_at(10^-14)
```

```
[1.0 0.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0 0.0]
[0.0 0.0 1.0 0.0 0.0]
[0.0 0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 0.0 1.0]
```

```
sage: (Q*R - A).zero_at(10^-14)
```

```
[0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0 0.0]
```

Now over the complex numbers, demonstrating that the SciPy libraries are (properly) using the Hermitian inner product, so that Q is a unitary matrix (its inverse is the conjugate-transpose).

```
sage: A = matrix(CDF, [[-8, 4*I + 1, -I + 2, 2*I + 1],
...                    [1, -2*I - 1, -I + 3, -I + 1],
...                    [I + 7, 2*I + 1, -2*I + 7, -I + 1],
...                    [I + 2, 0, I + 12, -1]])
```

```
sage: Q, R = A.QR()
```

```
sage: Q._normalize_columns().round(6).zero_at(10^-6)
```

```
[      0.730297  0.207057 + 0.538347*I  0.246305 - 0.076446*I  0.238162 - 0.10366*I
[      -0.091287 -0.207057 - 0.377878*I  0.378656 - 0.195222*I  0.701244 - 0.364371*I
[    -0.63901 - 0.091287*I  0.170822 + 0.667758*I -0.034115 + 0.040902*I  0.314017 - 0.082519*I
[-0.182574 - 0.091287*I  -0.036235 + 0.07247*I  0.863228 + 0.063228*I -0.449969 - 0.011612*I]
```

```
sage: R._normalize_rows().round(6).zero_at(10^-6)
```

```
[      10.954451      -1.917029*I  5.385938 - 2.19089*I  -0.273861 - 2.19089*I
[      0.0      4.829596 -0.869638 - 5.864879*I  0.993872 - 0.305409*I
[      0.0      0.0      12.001608 -0.270953 + 0.442063*I
[      0.0      0.0      0.0      0.0      1.942964]
```

```
sage: (Q.conjugate().transpose()*Q).zero_at(10^-15)
```

```
[1.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0]
[0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 1.0]
```

```
sage: (Q*R - A).zero_at(10^-14)
```

```
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
```

An example of a rectangular matrix that is also rank-deficient. If you run this example yourself, you may see a very small, nonzero entries in the third row, in the third column, even though the exact version of the matrix has rank 2. The final two columns of Q span the left kernel of A (as evidenced by the two zero rows of R). Different platforms will compute different bases for this left kernel, so we do not exhibit the actual matrix.

```
sage: Arat = matrix(QQ, [[2, -3, 3],
...                     [-1, 1, -1],
...                     [-1, 3, -3],
...                     [-5, 1, -1]])
sage: Arat.rank()
```

```
2
sage: A = Arat.change_ring(CDF)
sage: Q, R = A.QR()
sage: R._normalize_rows().round(6).zero_at(10^-6)
[ 5.567764 -2.69408  2.69408]
[      0.0  3.569585 -3.569585]
[      0.0      0.0      0.0]
[      0.0      0.0      0.0]
sage: (Q.conjugate_transpose()*Q).zero_at(10^-14)
[1.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0]
[0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 1.0]
```

Results are cached, meaning they are immutable matrices. Make a copy if you need to manipulate a result.

```
sage: A = random_matrix(CDF, 2, 2)
sage: Q, R = A.QR()
sage: Q.is_mutable()
False
sage: R.is_mutable()
False
sage: Q[0,0] = 0
Traceback (most recent call last):
...
ValueError: matrix is immutable; please change a copy instead (i.e., use copy(M) to change a
sage: Qcopy = copy(Q)
sage: Qcopy[0,0] = 679
sage: Qcopy[0,0]
679.0
```

TESTS:

Trivial cases return trivial results of the correct size, and we check Q itself in one case, verifying a fix for [trac ticket #10795](#).

```
sage: A = zero_matrix(RDF, 0, 10)
sage: Q, R = A.QR()
sage: Q.nrows(), Q.ncols()
(0, 0)
sage: R.nrows(), R.ncols()
(0, 10)
sage: A = zero_matrix(RDF, 3, 0)
sage: Q, R = A.QR()
sage: Q.nrows(), Q.ncols()
(3, 3)
sage: R.nrows(), R.ncols()
(3, 0)
sage: Q
[1.0 0.0 0.0]
[0.0 1.0 0.0]
[0.0 0.0 1.0]
```

SVD()

Return the singular value decomposition of this matrix.

The U and V matrices are not unique and may be returned with different values in the future or on different systems. The S matrix is unique and contains the singular values in descending order.

The computed decomposition is cached and returned on subsequent calls.

INPUT:

- A – a matrix

OUTPUT:

- U, S, V – immutable matrices such that $A = U * S * V.conj().transpose()$ where U and V are orthogonal and S is zero off of the diagonal.

Note that if self is m-by-n, then the dimensions of the matrices that this returns are (m,m), (m,n), and (n, n).

Note: If all you need is the singular values of the matrix, see the more convenient `singular_values()`.

EXAMPLES:

```
sage: m = matrix(RDF, 4, range(1, 17))
sage: U, S, V = m.SVD()
sage: U*S*V.transpose()
[ 1.0  2.0  3.0  4.0]
[ 5.0  6.0  7.0  8.0]
[ 9.0 10.0 11.0 12.0]
[13.0 14.0 15.0 16.0]
```

A non-square example:

```
sage: m = matrix(RDF, 2, range(1, 7)); m
[1.0 2.0 3.0]
[4.0 5.0 6.0]
sage: U, S, V = m.SVD()
sage: U*S*V.transpose()
[1.0 2.0 3.0]
[4.0 5.0 6.0]
```

S contains the singular values:

```
sage: S.round(4)
[ 9.508  0.0  0.0]
[ 0.0 0.7729  0.0]
sage: [round(sqrt(abs(x)), 4) for x in (S*S.transpose()).eigenvalues()]
[9.508, 0.7729]
```

U and V are orthogonal matrices:

```
sage: U # random, SVD is not unique
[-0.386317703119 -0.922365780077]
[-0.922365780077  0.386317703119]
[-0.274721127897 -0.961523947641]
[-0.961523947641  0.274721127897]
sage: (U*U.transpose()).zero_at(1e-15)
[1.0 0.0]
[0.0 1.0]
sage: V # random, SVD is not unique
[-0.428667133549  0.805963908589  0.408248290464]
[-0.566306918848  0.112382414097 -0.816496580928]
[-0.703946704147 -0.581199080396  0.408248290464]
sage: (V*V.transpose()).zero_at(1e-15)
[1.0 0.0 0.0]
[0.0 1.0 0.0]
[0.0 0.0 1.0]
```

TESTS:

```
sage: m = matrix(RDF, 3, 2, range(1, 7)); m
[1.0 2.0]
[3.0 4.0]
[5.0 6.0]
sage: U, S, V = m.SVD()
sage: U*S*V.transpose()
[1.0 2.0]
[3.0 4.0]
[5.0 6.0]

sage: m = matrix(RDF, 3, 0, []); m
[]
sage: m.SVD()
([], [], [])
sage: m = matrix(RDF, 0, 3, []); m
[]
sage: m.SVD()
([], [], [])
sage: def shape(x): return (x.nrows(), x.ncols())
sage: m = matrix(RDF, 2, 3, range(6))
sage: map(shape, m.SVD())
[(2, 2), (2, 3), (3, 3)]
sage: for x in m.SVD(): x.is_immutable()
True
True
True
```

cholesky()

Returns the Cholesky factorization of a matrix that is real symmetric, or complex Hermitian.

INPUT:

Any square matrix with entries from RDF that is symmetric, or with entries from CDF that is Hermitian. The matrix must be positive definite for the Cholesky decomposition to exist.

OUTPUT:

For a matrix A the routine returns a lower triangular matrix L such that,

$$A = LL^*$$

where L^* is the conjugate-transpose in the complex case, and just the transpose in the real case. If the matrix fails to be positive definite (perhaps because it is not symmetric or Hermitian), then this function raises a `ValueError`.

IMPLEMENTATION:

The existence of a Cholesky decomposition and the positive definite property are equivalent. So this method and the `is_positive_definite()` method compute and cache both the Cholesky decomposition and the positive-definiteness. So the `is_positive_definite()` method or catching a `ValueError` from the `cholesky()` method are equally expensive computationally and if the decomposition exists, it is cached as a side-effect of either routine.

EXAMPLES:

A real matrix that is symmetric and positive definite.

```
sage: M = matrix(RDF, [[ 1, 1, 1, 1, 1],
...                   [ 1, 5, 31, 121, 341],
...                   [ 1, 31, 341, 1555, 4681],
```

```

...           [ 1,121, 1555, 7381, 22621],
...           [ 1,341, 4681, 22621, 69905]])
sage: M.is_symmetric()
True
sage: L = M.cholesky()
sage: L.round(6).zero_at(10^-10)
[ 1.0  0.0  0.0  0.0  0.0]
[ 1.0  2.0  0.0  0.0  0.0]
[ 1.0 15.0 10.723805 0.0 0.0]
[ 1.0 60.0 60.985814 7.792973 0.0]
[ 1.0 170.0 198.623524 39.366567 1.7231]
sage: (L*L.transpose()).round(6).zero_at(10^-10)
[ 1.0  1.0  1.0  1.0  1.0]
[ 1.0  5.0  31.0 121.0 341.0]
[ 1.0 31.0 341.0 1555.0 4681.0]
[ 1.0 121.0 1555.0 7381.0 22621.0]
[ 1.0 341.0 4681.0 22621.0 69905.0]

```

A complex matrix that is Hermitian and positive definite.

```

sage: A = matrix(CDF, [[ 23, 17*I + 3, 24*I + 25, 21*I],
...                    [-17*I + 3, 38, -69*I + 89, 7*I + 15],
...                    [-24*I + 25, 69*I + 89, 976, 24*I + 6],
...                    [-21*I, -7*I + 15, -24*I + 6, 28]])
sage: A.is_hermitian()
True
sage: L = A.cholesky()
sage: L.round(6).zero_at(10^-10)
[ 4.795832 0.0 0.0 0.0]
[ 0.625543 - 3.544745*I 5.004346 0.0 0.0]
[ 5.21286 - 5.004346*I 13.588189 + 10.721116*I 24.984023 0.0]
[ -4.378803*I -0.104257 - 0.851434*I -0.21486 + 0.371348*I 2.811799]
sage: (L*L.conjugate_transpose()).round(6).zero_at(10^-10)
[ 23.0 3.0 + 17.0*I 25.0 + 24.0*I 21.0*I]
[ 3.0 - 17.0*I 38.0 89.0 - 69.0*I 15.0 + 7.0*I]
[ 25.0 - 24.0*I 89.0 + 69.0*I 976.0 6.0 + 24.0*I]
[ -21.0*I 15.0 - 7.0*I 6.0 - 24.0*I 28.0]

```

This routine will recognize when the input matrix is not positive definite. The negative eigenvalues are an equivalent indicator. (Eigenvalues of a Hermitian matrix must be real, so there is no loss in ignoring the imprecise imaginary parts).

```

sage: A = matrix(RDF, [[ 3, -6, 9, 6, -9],
...                    [-6, 11, -16, -11, 17],
...                    [ 9, -16, 28, 16, -40],
...                    [ 6, -11, 16, 9, -19],
...                    [-9, 17, -40, -19, 68]])
sage: A.is_symmetric()
True
sage: A.eigenvalues()
[108.07..., 13.02..., -0.02..., -0.70..., -1.37...]
sage: A.cholesky()
Traceback (most recent call last):
...
ValueError: matrix is not positive definite

sage: B = matrix(CDF, [[ 2, 4 - 2*I, 2 + 2*I],
...                    [4 + 2*I, 8, 10*I],
...                    [2 - 2*I, -10*I, -3]])

```

```

sage: B.is_hermitian()
True
sage: [ev.real() for ev in B.eigenvalues()]
[15.88..., 0.08..., -8.97...]
sage: B.cholesky()
Traceback (most recent call last):
...
ValueError: matrix is not positive definite

```

TESTS:

A trivial case.

```

sage: A = matrix(RDF, 0, [])
sage: A.cholesky()
[]

```

The Cholesky factorization is only defined for square matrices.

```

sage: A = matrix(RDF, 4, 5, range(20))
sage: A.cholesky()
Traceback (most recent call last):
...
ValueError: Cholesky decomposition requires a square matrix, not a 4 x 5 matrix

```

condition ($p='frob'$)

Returns the condition number of a square nonsingular matrix.

Roughly speaking, this is a measure of how sensitive the matrix is to round-off errors in numerical computations. The minimum possible value is 1.0, and larger numbers indicate greater sensitivity.

INPUT:

- p - default: 'frob' - controls which norm is used to compute the condition number, allowable values are 'frob' (for the Frobenius norm), integers -2, -1, 1, 2, positive and negative infinity. See output discussion for specifics.

OUTPUT:

The condition number of a matrix is the product of a norm of the matrix times the norm of the inverse of the matrix. This requires that the matrix be square and invertible (nonsingular, full rank).

Returned value is a double precision floating point value in RDF, or `Infinity`. Row and column sums described below are sums of the absolute values of the entries, where the absolute value of the complex number $a + bi$ is $\sqrt{a^2 + b^2}$. Singular values are the “diagonal” entries of the “S” matrix in the singular value decomposition.

- $p = 'frob'$: the default norm employed in computing the condition number, the Frobenius norm, which for a matrix $A = (a_{ij})$ computes

$$\left(\sum_{i,j} |a_{i,j}|^2 \right)^{1/2}$$

- $p = 'sv'$: the quotient of the maximal and minimal singular value.
- $p = \text{Infinity}$ or $p = \infty$: the maximum row sum.
- $p = -\text{Infinity}$ or $p = -\infty$: the minimum column sum.
- $p = 1$: the maximum column sum.

- $p = -1$: the minimum column sum.
- $p = 2$: the 2-norm, equal to the maximum singular value.
- $p = -2$: the minimum singular value.

ALGORITHM:

Computation is performed by the `cond()` function of the SciPy/NumPy library.

EXAMPLES:

First over the reals.

```
sage: A = matrix(RDF, 4, [(1/4)*x^3 for x in range(16)]); A
[  0.0  0.25  2.0  6.75]
[ 16.0 31.25 54.0 85.75]
[128.0 182.25 250.0 332.75]
[432.0 549.25 686.0 843.75]
sage: A.condition()
9923.88955...
sage: A.condition(p='frob')
9923.88955...
sage: A.condition(p=Infinity)
22738.5
sage: A.condition(p=-Infinity)
17.5
sage: A.condition(p=1)
12139.21...
sage: A.condition(p=-1)
550.0
sage: A.condition(p=2)
9897.8088...
sage: A.condition(p=-2)
0.000101032462...
```

And over the complex numbers.

```
sage: B = matrix(CDF, 3, [x + x^2*I for x in range(9)]); B
[  0.0  1.0 + 1.0*I  2.0 + 4.0*I]
[ 3.0 + 9.0*I  4.0 + 16.0*I  5.0 + 25.0*I]
[ 6.0 + 36.0*I  7.0 + 49.0*I  8.0 + 64.0*I]
sage: B.condition()
203.851798...
sage: B.condition(p='frob')
203.851798...
sage: B.condition(p=Infinity)
369.55630...
sage: B.condition(p=-Infinity)
5.46112969...
sage: B.condition(p=1)
289.251481...
sage: B.condition(p=-1)
20.4566639...
sage: B.condition(p=2)
202.653543...
sage: B.condition(p=-2)
0.00493453005...
```

Hilbert matrices are famously ill-conditioned, while an identity matrix can hit the minimum with the right norm.

```
sage: A = matrix(RDF, 10, [1/(i+j+1) for i in range(10) for j in range(10)])
sage: A.condition()
1.633...e+13
sage: id = identity_matrix(CDF, 10)
sage: id.condition(p=1)
1.0
```

Return values are in *RDF*.

```
sage: A = matrix(CDF, 2, range(1,5))
sage: A.condition() in RDF
True
```

Rectangular and singular matrices raise errors if *p* is not 'sv'.

```
sage: A = matrix(RDF, 2, 3, range(6))
sage: A.condition()
Traceback (most recent call last):
...
TypeError: matrix must be square if p is not 'sv', not 2 x 3

sage: A.condition('sv')
7.34...
```

```
sage: A = matrix(QQ, 5, range(25))
sage: A.is_singular()
True
sage: B = A.change_ring(CDF)
sage: B.condition()
Traceback (most recent call last):
...
LinAlgError: Singular matrix
```

Improper values of *p* are caught.

```
sage: A = matrix(CDF, 2, range(1,5))
sage: A.condition(p='bogus')
Traceback (most recent call last):
...
ValueError: condition number 'p' must be +/- infinity, 'frob', 'sv' or an integer, not bogus
sage: A.condition(p=632)
Traceback (most recent call last):
...
ValueError: condition number integer values of 'p' must be -2, -1, 1 or 2, not 632
```

TESTS:

Some condition numbers, first by the definition which also exercises `norm()`, then by this method.

```
sage: A = matrix(CDF, [[1,2,4],[5,3,9],[7,8,6]])
sage: c = A.norm(2)*A.inverse().norm(2)
sage: d = A.condition(2)
sage: abs(c-d) < 1.0e-12
True
sage: c = A.norm(1)*A.inverse().norm(1)
sage: d = A.condition(1)
sage: abs(c-d) < 1.0e-12
True
```

`determinant()`

Return the determinant of self.

ALGORITHM:

Use numpy

EXAMPLES:

```
sage: m = matrix(RDF, 2, range(4)); m.det()
-2.0
sage: m = matrix(RDF, 0, []); m.det()
1.0
sage: m = matrix(RDF, 2, range(6)); m.det()
Traceback (most recent call last):
...
ValueError: self must be a square matrix
```

eigenvalues (*algorithm*='default', *tol*=None)

Returns a list of eigenvalues.

INPUT:

- *self* - a square matrix
- *algorithm* - default: 'default'
 - 'default' - applicable to any matrix with double-precision floating point entries. Uses the `eigvals()` method from SciPy.
 - 'symmetric' - converts the matrix into a real matrix (i.e. with entries from RDF), then applies the algorithm for Hermitian matrices. This algorithm can be significantly faster than the 'default' algorithm.
 - 'hermitian' - uses the `eigh()` method from SciPy, which applies only to real symmetric or complex Hermitian matrices. Since Hermitian is defined as a matrix equaling its conjugate-transpose, for a matrix with real entries this property is equivalent to being symmetric. This algorithm can be significantly faster than the 'default' algorithm.
- *tol* - default: None - if set to a value other than None this is interpreted as a small real number used to aid in grouping eigenvalues that are numerically similar. See the output description for more information.

Warning: When using the 'symmetric' or 'hermitian' algorithms, no check is made on the input matrix, and only the entries below, and on, the main diagonal are employed in the computation. Methods such as `is_symmetric()` and `is_hermitian()` could be used to verify this beforehand.

OUTPUT:

Default output for a square matrix of size n is a list of n eigenvalues from the complex double field, CDF. If the 'symmetric' or 'hermitian' algorithms are chosen, the returned eigenvalues are from the real double field, RDF.

If a tolerance is specified, an attempt is made to group eigenvalues that are numerically similar. The return is then a list of pairs, where each pair is an eigenvalue followed by its multiplicity. The eigenvalue reported is the mean of the eigenvalues computed, and these eigenvalues are contained in an interval (or disk) whose radius is less than $5 * tol$ for $n < 10,000$ in the worst case.

More precisely, for an $n \times n$ matrix, the diameter of the interval containing similar eigenvalues could be as large as sum of the reciprocals of the first n integers times tol .

Warning: Use caution when using the `tol` parameter to group eigenvalues. See the examples below to see how this can go wrong.

EXAMPLES:

```
sage: m = matrix(RDF, 2, 2, [1,2,3,4])
sage: ev = m.eigenvalues(); ev
[-0.372281323..., 5.37228132...]
sage: ev[0].parent()
Complex Double Field

sage: m = matrix(RDF, 2, 2, [0,1,-1,0])
sage: m.eigenvalues(algorithm='default')
[1.0*I, -1.0*I]

sage: m = matrix(CDF, 2, 2, [I,1,-I,0])
sage: m.eigenvalues()
[-0.624810533... + 1.30024259...*I, 0.624810533... - 0.30024259...*I]
```

The adjacency matrix of a graph will be symmetric, and the eigenvalues will be real.

```
sage: A = graphs.PetersenGraph().adjacency_matrix()
sage: A = A.change_ring(RDF)
sage: ev = A.eigenvalues(algorithm='symmetric'); ev
[-2.0, -2.0, -2.0, -2.0, 1.0, 1.0, 1.0, 1.0, 1.0, 3.0]
sage: ev[0].parent()
Real Double Field
```

The matrix A is “random”, but the construction of B provides a positive-definite Hermitian matrix. Note that the eigenvalues of a Hermitian matrix are real, and the eigenvalues of a positive-definite matrix will be positive.

```
sage: A = matrix([[ 4*I + 5, 8*I + 1, 7*I + 5, 3*I + 5],
...               [ 7*I - 2, -4*I + 7, -2*I + 4, 8*I + 8],
...               [-2*I + 1, 6*I + 6, 5*I + 5, -I - 4],
...               [ 5*I + 1, 6*I + 2, I - 4, -I + 3]])
sage: B = (A*A.conjugate_transpose()).change_ring(CDF)
sage: ev = B.eigenvalues(algorithm='hermitian'); ev
[2.68144025..., 49.5167998..., 274.086188..., 390.71557...]
sage: ev[0].parent()
Real Double Field
```

A tolerance can be given to aid in grouping eigenvalues that are similar numerically. However, if the parameter is too small it might split too finely. Too large, and it can go wrong very badly. Use with care.

```
sage: G = graphs.PetersenGraph()
sage: G.spectrum()
[3, 1, 1, 1, 1, 1, -2, -2, -2, -2]

sage: A = G.adjacency_matrix().change_ring(RDF)
sage: A.eigenvalues(algorithm='symmetric', tol=1.0e-5)
[(-2.0, 4), (1.0, 5), (3.0, 1)]

sage: A.eigenvalues(algorithm='symmetric', tol=2.5)
[(-2.0, 4), (1.33333333333, 6)]
```

An (extreme) example of properly grouping similar eigenvalues.

```
sage: G = graphs.HigmanSimsGraph()
sage: A = G.adjacency_matrix().change_ring(RDF)
sage: A.eigenvalues(algorithm='symmetric', tol=1.0e-5)
[(-8.0, 22), (2.0, 77), (22.0, 1)]
```

TESTS:

Testing bad input.

```
sage: A = matrix(CDF, 2, range(4))
sage: A.eigenvalues(algorithm='junk')
Traceback (most recent call last):
...
ValueError: algorithm must be 'default', 'symmetric', or 'hermitian', not junk
```

```
sage: A = matrix(CDF, 2, 3, range(6))
sage: A.eigenvalues()
Traceback (most recent call last):
...
ValueError: matrix must be square, not 2 x 3
```

```
sage: A = matrix(CDF, 2, [1, 2, 3, 4*I])
sage: A.eigenvalues(algorithm='symmetric')
Traceback (most recent call last):
...
TypeError: cannot apply symmetric algorithm to matrix with complex entries
```

```
sage: A = matrix(CDF, 2, 2, range(4))
sage: A.eigenvalues(tol='junk')
Traceback (most recent call last):
...
TypeError: tolerance parameter must be a real number, not junk
```

```
sage: A = matrix(CDF, 2, 2, range(4))
sage: A.eigenvalues(tol=-0.01)
Traceback (most recent call last):
...
ValueError: tolerance parameter must be positive, not -0.01
```

A very small matrix.

```
sage: matrix(CDF, 0, 0).eigenvalues()
[]
```

eigenvectors_left()

Compute the left eigenvectors of a matrix of double precision real or complex numbers (i.e. RDF or CDF).

OUTPUT: Returns a list of triples, each of the form $(e, [v], 1)$, where e is the eigenvalue, and v is an associated left eigenvector. If the matrix is of size n , then there are n triples. Values are computed with the SciPy library.

The format of this output is designed to match the format for exact results. However, since matrices here have numerical entries, the resulting eigenvalues will also be numerical. No attempt is made to determine if two eigenvalues are equal, or if eigenvalues might actually be zero. So the algebraic multiplicity of each eigenvalue is reported as 1. Decisions about equal eigenvalues or zero eigenvalues should be addressed in the calling routine.

The SciPy routines used for these computations produce eigenvectors normalized to have length 1, but on different hardware they may vary by a sign. So for doctests we have normalized output by forcing their

eigenvectors to have their first non-zero entry equal to one.

EXAMPLES:

```
sage: m = matrix(RDF, [[-5, 3, 2, 8], [10, 2, 4, -2], [-1, -10, -10, -17], [-2, 7, 6, 13]])
sage: m
[ -5.0   3.0   2.0   8.0]
[ 10.0   2.0   4.0  -2.0]
[ -1.0 -10.0 -10.0 -17.0]
[ -2.0   7.0   6.0  13.0]
sage: spectrum = m.left_eigenvectors()
sage: for i in range(len(spectrum)):
...     spectrum[i][1][0]=matrix(RDF, spectrum[i][1]).echelon_form()[0]
sage: spectrum[0]
(2.0, [(1.0, 1.0, 1.0, 1.0)], 1)
sage: spectrum[1]
(1.0, [(1.0, 0.8, 0.8, 0.6)], 1)
sage: spectrum[2]
(-2.0, [(1.0, 0.4, 0.6, 0.2)], 1)
sage: spectrum[3]
(-1.0, [(1.0, 1.0, 2.0, 2.0)], 1)
```

eigenvectors_right()

Compute the right eigenvectors of a matrix of double precision real or complex numbers (i.e. RDF or CDF).

OUTPUT:

Returns a list of triples, each of the form $(e, [v], 1)$, where e is the eigenvalue, and v is an associated right eigenvector. If the matrix is of size n , then there are n triples. Values are computed with the SciPy library.

The format of this output is designed to match the format for exact results. However, since matrices here have numerical entries, the resulting eigenvalues will also be numerical. No attempt is made to determine if two eigenvalues are equal, or if eigenvalues might actually be zero. So the algebraic multiplicity of each eigenvalue is reported as 1. Decisions about equal eigenvalues or zero eigenvalues should be addressed in the calling routine.

The SciPy routines used for these computations produce eigenvectors normalized to have length 1, but on different hardware they may vary by a sign. So for doctests we have normalized output by forcing their eigenvectors to have their first non-zero entry equal to one.

EXAMPLES:

```
sage: m = matrix(RDF, [[-9, -14, 19, -74], [-1, 2, 4, -11], [-4, -12, 6, -32], [0, -2, -1, 1]])
sage: m
[ -9.0 -14.0  19.0 -74.0]
[ -1.0   2.0   4.0 -11.0]
[ -4.0 -12.0   6.0 -32.0]
[  0.0  -2.0  -1.0   1.0]
sage: spectrum = m.right_eigenvectors()
sage: for i in range(len(spectrum)):
...     spectrum[i][1][0]=matrix(RDF, spectrum[i][1]).echelon_form()[0]
sage: spectrum[0]
(2.0, [(1.0, -2.0, 3.0, 1.0)], 1)
sage: spectrum[1]
(1.0, [(1.0, -0.6666666666667, 1.3333333333333, 0.3333333333333)], 1)
sage: spectrum[2]
(-2.0, [(1.0, -0.2, 1.0, 0.2)], 1)
sage: spectrum[3]
(-1.0, [(1.0, -0.5, 2.0, 0.5)], 1)
```

exp (*algorithm*='pade', *order*=None)

Calculate the exponential of this matrix X, which is the matrix

$$e^X = \sum_{k=0}^{\infty} \frac{X^k}{k!}.$$

INPUT:

- *algorithm* – 'pade', 'eig', or 'taylor'; the algorithm used to compute the exponential.
- *order* – for the Taylor series algorithm, this specifies the order of the Taylor series used. This is ignored for the other algorithms. The current default (from scipy) is 20.

EXAMPLES:

```
sage: A=matrix(RDF, 2, [1,2,3,4]); A
[1.0 2.0]
[3.0 4.0]
sage: A.exp()
[51.9689561987  74.736564567]
[112.104846851 164.073803049]
sage: A.exp(algorithm='eig')
[51.9689561987  74.736564567]
[112.104846851 164.073803049]
sage: A.exp(algorithm='taylor', order=5)
[19.9583333333 28.0833333333]
[ 42.125 62.0833333333]
sage: A.exp(algorithm='taylor')
[51.9689035511 74.7364878369]
[112.104731755 164.073635306]

sage: A=matrix(CDF, 2, [1,2+I,3*I,4]); A
[ 1.0 2.0 + 1.0*I]
[ 3.0*I 4.0]
sage: A.exp()
[-19.6146029538 + 12.5177438468*I  3.79496364496 + 28.8837993066*I]
[-32.3835809809 + 21.8842359579*I  2.26963300409 + 44.9013248277*I]
sage: A.exp(algorithm='eig')
[-19.6146029538 + 12.5177438468*I  3.79496364496 + 28.8837993066*I]
[-32.3835809809 + 21.8842359579*I  2.26963300409 + 44.9013248277*I]
sage: A.exp(algorithm='taylor', order=5)
[ -6.29166666667 + 14.25*I 14.0833333333 + 15.7916666667*I]
[ -10.5 + 26.375*I 20.0833333333 + 24.75*I]
sage: A.exp(algorithm='taylor')
[-19.6146006163 + 12.5177432169*I  3.79496442472 + 28.8837964828*I]
[-32.3835771246 + 21.8842351994*I  2.26963458304 + 44.9013203415*I]
```

is_hermitian (*tol*=1e-12, *algorithm*='orthonormal')

Returns True if the matrix is equal to its conjugate-transpose.

INPUT:

- *tol* - default: 1e-12 - the largest value of the absolute value of the difference between two matrix entries for which they will still be considered equal.
- *algorithm* - default: 'orthonormal' - set to 'orthonormal' for a stable procedure and set to 'naive' for a fast procedure.

OUTPUT:

True if the matrix is square and equal to the transpose with every entry conjugated, and False otherwise.

Note that if conjugation has no effect on elements of the base ring (such as for integers), then the `is_symmetric()` method is equivalent and faster.

The tolerance parameter is used to allow for numerical values to be equal if there is a slight difference due to round-off and other imprecisions.

The result is cached, on a per-tolerance and per-algorithm basis.

ALGORITHMS:

The naive algorithm simply compares corresponding entries on either side of the diagonal (and on the diagonal itself) to see if they are conjugates, with equality controlled by the tolerance parameter.

The orthonormal algorithm first computes a Schur decomposition (via the `schur()` method) and checks that the result is a diagonal matrix with real entries.

So the naive algorithm can finish quickly for a matrix that is not Hermitian, while the orthonormal algorithm will always compute a Schur decomposition before going through a similar check of the matrix entry-by-entry.

EXAMPLES:

```
sage: A = matrix(CDF, [[ 1 + I, 1 - 6*I, -1 - I],
...                  [-3 - I, -4*I, -2],
...                  [-1 + I, -2 - 8*I, 2 + I]])
sage: A.is_hermitian(algorithm='orthonormal')
False
sage: A.is_hermitian(algorithm='naive')
False
sage: B = A*A.conjugate_transpose()
sage: B.is_hermitian(algorithm='orthonormal')
True
sage: B.is_hermitian(algorithm='naive')
True
```

A matrix that is nearly Hermitian, but for one non-real diagonal entry.

```
sage: A = matrix(CDF, [[ 2, 2-I, 1+4*I],
...                  [ 2+I, 3+I, 2-6*I],
...                  [1-4*I, 2+6*I, 5]])
sage: A.is_hermitian(algorithm='orthonormal')
False
sage: A[1,1] = 132
sage: A.is_hermitian(algorithm='orthonormal')
True
```

We get a unitary matrix from the SVD routine and use this numerical matrix to create a matrix that should be Hermitian (indeed it should be the identity matrix), but with some imprecision. We use this to illustrate that if the tolerance is set too small, then we can be too strict about the equality of entries and may achieve the wrong result (depending on the system):

```
sage: A = matrix(CDF, [[ 1 + I, 1 - 6*I, -1 - I],
...                  [-3 - I, -4*I, -2],
...                  [-1 + I, -2 - 8*I, 2 + I]])
sage: U, _, _ = A.SVD()
sage: B=U*U.conjugate_transpose()
sage: B.is_hermitian(algorithm='naive')
True
sage: B.is_hermitian(algorithm='naive', tol=1.0e-17) # random
False
sage: B.is_hermitian(algorithm='naive', tol=1.0e-15)
True
```

A square, empty matrix is trivially Hermitian.

```
sage: A = matrix(RDF, 0, 0)
sage: A.is_hermitian()
True
```

Rectangular matrices are never Hermitian, no matter which algorithm is requested.

```
sage: A = matrix(CDF, 3, 4)
sage: A.is_hermitian()
False
```

TESTS:

The tolerance must be strictly positive.

```
sage: A = matrix(RDF, 2, range(4))
sage: A.is_hermitian(tol = -3.1)
Traceback (most recent call last):
...
ValueError: tolerance must be positive, not -3.1
```

The algorithm keyword gets checked.

```
sage: A = matrix(RDF, 2, range(4))
sage: A.is_hermitian(algorithm='junk')
Traceback (most recent call last):
...
ValueError: algorithm must be 'naive' or 'orthonormal', not junk
```

AUTHOR:

- Rob Beezer (2011-03-30)

is_normal (*tol=1e-12, algorithm='orthonormal'*)

Returns True if the matrix commutes with its conjugate-transpose.

INPUT:

- tol - default: 1e-12 - the largest value of the absolute value of the difference between two matrix entries for which they will still be considered equal.
- algorithm - default: 'orthonormal' - set to 'orthonormal' for a stable procedure and set to 'naive' for a fast procedure.

OUTPUT:

True if the matrix is square and commutes with its conjugate-transpose, and False otherwise.

Normal matrices are precisely those that can be diagonalized by a unitary matrix.

The tolerance parameter is used to allow for numerical values to be equal if there is a slight difference due to round-off and other imprecisions.

The result is cached, on a per-tolerance and per-algorithm basis.

ALGORITHMS:

The naive algorithm simply compares entries of the two possible products of the matrix with its conjugate-transpose, with equality controlled by the tolerance parameter.

The orthonormal algorithm first computes a Schur decomposition (via the `schur()` method) and checks that the result is a diagonal matrix. An orthonormal diagonalization is equivalent to being normal.

So the naive algorithm can finish fairly quickly for a matrix that is not normal, once the products have been computed. However, the orthonormal algorithm will compute a Schur decomposition before going through a similar check of a matrix entry-by-entry.

EXAMPLES:

First over the complexes. B is Hermitian, hence normal.

```
sage: A = matrix(CDF, [[ 1 + I, 1 - 6*I, -1 - I],
...                  [-3 - I, -4*I, -2],
...                  [-1 + I, -2 - 8*I, 2 + I]])
sage: B = A*A.conjugate_transpose()
sage: B.is_hermitian()
True
sage: B.is_normal(algorithm='orthonormal')
True
sage: B.is_normal(algorithm='naive')
True
sage: B[0,0] = I
sage: B.is_normal(algorithm='orthonormal')
False
sage: B.is_normal(algorithm='naive')
False
```

Now over the reals. Circulant matrices are normal.

```
sage: G = graphs.CirculantGraph(20, [3, 7])
sage: D = digraphs.Circuit(20)
sage: A = 3*D.adjacency_matrix() - 5*G.adjacency_matrix()
sage: A = A.change_ring(RDF)
sage: A.is_normal()
True
sage: A.is_normal(algorithm = 'naive')
True
sage: A[19,0] = 4.0
sage: A.is_normal()
False
sage: A.is_normal(algorithm = 'naive')
False
```

Skew-Hermitian matrices are normal.

```
sage: A = matrix(CDF, [[ 1 + I, 1 - 6*I, -1 - I],
...                  [-3 - I, -4*I, -2],
...                  [-1 + I, -2 - 8*I, 2 + I]])
sage: B = A - A.conjugate_transpose()
sage: B.is_hermitian()
False
sage: B.is_normal()
True
sage: B.is_normal(algorithm='naive')
True
```

A small matrix that does not fit into any of the usual categories of normal matrices.

```
sage: A = matrix(RDF, [[1, -1],
...                  [1, 1]])
sage: A.is_normal()
True
sage: not A.is_hermitian() and not A.is_skew_symmetric()
True
```


Sage has several fields besides the entire complex numbers where conjugation is non-trivial.

```
sage: F.<b> = QuadraticField(-7)
sage: C = matrix(F, [[-2*b - 3, 7*b - 6, -b + 3],
...                  [-2*b - 3, -3*b + 2, -2*b],
...                  [ b + 1, 0, -2]])
sage: C = C*C.conjugate_transpose()
sage: C.is_normal()
True
```

A square, empty matrix is trivially normal.

```
sage: A = matrix(CDF, 0, 0)
sage: A.is_normal()
True
```

Rectangular matrices are never normal, no matter which algorithm is requested.

```
sage: A = matrix(CDF, 3, 4)
sage: A.is_normal()
False
```

TESTS:

The tolerance must be strictly positive.

```
sage: A = matrix(RDF, 2, range(4))
sage: A.is_normal(tol = -3.1)
Traceback (most recent call last):
...
ValueError: tolerance must be positive, not -3.1
```

The algorithm keyword gets checked.

```
sage: A = matrix(RDF, 2, range(4))
sage: A.is_normal(algorithm='junk')
Traceback (most recent call last):
...
ValueError: algorithm must be 'naive' or 'orthonormal', not junk
```

AUTHOR:

•Rob Beezer (2011-03-31)

is_positive_definite()

Determines if a matrix is positive definite.

A matrix A is positive definite if it is square, is Hermitian (which reduces to symmetric in the real case), and for every nonzero vector \vec{x} ,

$$\vec{x}^* A \vec{x} > 0$$

where \vec{x}^* is the conjugate-transpose in the complex case and just the transpose in the real case. Equivalently, a positive definite matrix has only positive eigenvalues and only positive determinants of leading principal submatrices.

INPUT:

Any matrix over RDF or CDF.

OUTPUT:

True if and only if the matrix is square, Hermitian, and meets the condition above on the quadratic form. The result is cached.

IMPLEMENTATION:

The existence of a Cholesky decomposition and the positive definite property are equivalent. So this method and the `cholesky()` method compute and cache both the Cholesky decomposition and the positive-definiteness. So the `is_positive_definite()` method or catching a `ValueError` from the `cholesky()` method are equally expensive computationally and if the decomposition exists, it is cached as a side-effect of either routine.

EXAMPLES:

A matrix over RDF that is positive definite.

```
sage: M = matrix(RDF, [[ 1, 1, 1, 1, 1],
...                    [ 1, 5, 31, 121, 341],
...                    [ 1, 31, 341, 1555, 4681],
...                    [ 1, 121, 1555, 7381, 22621],
...                    [ 1, 341, 4681, 22621, 69905]])
sage: M.is_symmetric()
True
sage: M.eigenvalues()
[77547.66..., 82.44..., 2.41..., 0.46..., 0.011...]
sage: [round(M[:i,:i].determinant()) for i in range(1, M.nrows()+1)]
[1, 4, 460, 27936, 82944]
sage: M.is_positive_definite()
True
```

A matrix over CDF that is positive definite.

```
sage: C = matrix(CDF, [[ 23, 17*I + 3, 24*I + 25, 21*I],
...                    [-17*I + 3, 38, -69*I + 89, 7*I + 15],
...                    [-24*I + 25, 69*I + 89, 976, 24*I + 6],
...                    [-21*I, -7*I + 15, -24*I + 6, 28]])
sage: C.is_hermitian()
True
sage: [x.real() for x in C.eigenvalues()]
[991.46..., 55.96..., 3.69..., 13.87...]
sage: [round(C[:i,:i].determinant().real()) for i in range(1, C.nrows()+1)]
[23, 576, 359540, 2842600]
sage: C.is_positive_definite()
True
```

A matrix over RDF that is not positive definite.

```
sage: A = matrix(RDF, [[ 3, -6, 9, 6, -9],
...                    [-6, 11, -16, -11, 17],
...                    [ 9, -16, 28, 16, -40],
...                    [ 6, -11, 16, 9, -19],
...                    [-9, 17, -40, -19, 68]])
sage: A.is_symmetric()
True
sage: A.eigenvalues()
[108.07..., 13.02..., -0.02..., -0.70..., -1.37...]
sage: [round(A[:i,:i].determinant()) for i in range(1, A.nrows()+1)]
[3, -3, -15, 30, -30]
sage: A.is_positive_definite()
False
```

A matrix over CDF that is not positive definite.

```
sage: B = matrix(CDF, [[ 2, 4 - 2*I, 2 + 2*I],
...                    [4 + 2*I, 8, 10*I],
```

```

... [2 - 2*I, -10*I, -3]])
sage: B.is_hermitian()
True
sage: [ev.real() for ev in B.eigenvalues()]
[15.88..., 0.08..., -8.97...]
sage: [round(B[:i,:i].determinant().real()) for i in range(1, B.nrows()+1)]
[2, -4, -12]
sage: B.is_positive_definite()
False

```

A large random matrix that is guaranteed by theory to be positive definite.

```

sage: R = random_matrix(CDF, 200)
sage: H = R.conjugate_transpose()*R
sage: H.is_positive_definite()
True

```

TESTS:

A trivially small case.

```

sage: S = matrix(CDF, [])
sage: S.nrows(), S.ncols()
(0, 0)
sage: S.is_positive_definite()
True

```

A rectangular matrix will never be positive definite.

```

sage: R = matrix(RDF, 2, 3, range(6))
sage: R.is_positive_definite()
False

```

A non-Hermitian matrix will never be positive definite.

```

sage: T = matrix(CDF, 8, 8, range(64))
sage: T.is_positive_definite()
False

```

AUTHOR:

•Rob Beezer (2012-05-28)

is_symmetric (*tol=1e-12*)

Return whether this matrix is symmetric, to the given tolerance.

EXAMPLES:

```

sage: m = matrix(RDF, 2, 2, range(4)); m
[0.0 1.0]
[2.0 3.0]
sage: m.is_symmetric()
False
sage: m[1,0] = 1.1; m
[0.0 1.0]
[1.1 3.0]
sage: m.is_symmetric()
False

```

The tolerance inequality is strict: sage: m.is_symmetric(tol=0.1) False sage: m.is_symmetric(tol=0.11)
True

is_unitary (*tol=1e-12, algorithm='orthonormal'*)

Returns `True` if the columns of the matrix are an orthonormal basis.

For a matrix with real entries this determines if a matrix is “orthogonal” and for a matrix with complex entries this determines if the matrix is “unitary.”

INPUT:

- `tol` - default: `1e-12` - the largest value of the absolute value of the difference between two matrix entries for which they will still be considered equal.
- `algorithm` - default: `'orthonormal'` - set to `'orthonormal'` for a stable procedure and set to `'naive'` for a fast procedure.

OUTPUT:

`True` if the matrix is square and its conjugate-transpose is its inverse, and `False` otherwise. In other words, a matrix is orthogonal or unitary if the product of its conjugate-transpose times the matrix is the identity matrix.

The tolerance parameter is used to allow for numerical values to be equal if there is a slight difference due to round-off and other imprecisions.

The result is cached, on a per-tolerance and per-algorithm basis.

ALGORITHMS:

The naive algorithm simply computes the product of the conjugate-transpose with the matrix and compares the entries to the identity matrix, with equality controlled by the tolerance parameter.

The orthonormal algorithm first computes a Schur decomposition (via the `schur()` method) and checks that the result is a diagonal matrix with entries of modulus 1, which is equivalent to being unitary.

So the naive algorithm might finish fairly quickly for a matrix that is not unitary, once the product has been computed. However, the orthonormal algorithm will compute a Schur decomposition before going through a similar check of a matrix entry-by-entry.

EXAMPLES:

A matrix that is far from unitary.

```
sage: A = matrix(RDF, 4, range(16))
sage: A.conjugate().transpose()*A
[224.0 248.0 272.0 296.0]
[248.0 276.0 304.0 332.0]
[272.0 304.0 336.0 368.0]
[296.0 332.0 368.0 404.0]
sage: A.is_unitary()
False
sage: A.is_unitary(algorithm='naive')
False
sage: A.is_unitary(algorithm='orthonormal')
False
```

The QR decomposition will produce a unitary matrix as `Q` and the SVD decomposition will create two unitary matrices, `U` and `V`.

```
sage: A = matrix(CDF, [[ 1 - I, -3*I, -2 + I, 1, -2 + 3*I],
... [ 1 - I, -2 + I, 1 + 4*I, 0, 2 + I],
... [ -1, -5 + I, -2 + I, 1 + I, -5 - 4*I],
... [-2 + 4*I, 2 - I, 8 - 4*I, 1 - 8*I, 3 - 2*I]])
sage: Q, R = A.QR()
sage: Q.is_unitary()
True
```

```

sage: U, S, V = A.SVD()
sage: U.is_unitary(algorithm='naive')
True
sage: U.is_unitary(algorithm='orthonormal')
True
sage: V.is_unitary(algorithm='naive') # not tested - known bug (trac #11248)
True

```

If we make the tolerance too strict we can get misleading results.

```

sage: A = matrix(RDF, 10, 10, [1/(i+j+1) for i in range(10) for j in range(10)])
sage: Q, R = A.QR()
sage: Q.is_unitary(algorithm='naive', tol=1e-16)
False
sage: Q.is_unitary(algorithm='orthonormal', tol=1e-17)
False

```

Rectangular matrices are not unitary/orthogonal, even if their columns form an orthonormal set.

```

sage: A = matrix(CDF, [[1,0], [0,0], [0,1]])
sage: A.is_unitary()
False

```

The smallest cases. The Schur decomposition used by the orthonormal algorithm will fail on a matrix of size zero.

```

sage: P = matrix(CDF, 0, 0)
sage: P.is_unitary(algorithm='naive')
True

```

```

sage: P = matrix(CDF, 1, 1, [1])
sage: P.is_unitary(algorithm='orthonormal')
True

```

```

sage: P = matrix(CDF, 0, 0,)
sage: P.is_unitary(algorithm='orthonormal')
Traceback (most recent call last):
...

```

ValueError: failed to create intent(cache|hide)|optional array-- must have defined dimension

TESTS:

```

sage: P = matrix(CDF, 2, 2)
sage: P.is_unitary(tol='junk')
Traceback (most recent call last):
...

```

TypeError: tolerance must be a real number, not junk

```

sage: P.is_unitary(tol=-0.3)
Traceback (most recent call last):
...

```

ValueError: tolerance must be positive, not -0.3

```

sage: P.is_unitary(algorithm='junk')
Traceback (most recent call last):
...

```

ValueError: algorithm must be 'naive' or 'orthonormal', not junk

AUTHOR:

•Rob Beezer (2011-05-04)

`left_eigenvectors()`

Compute the left eigenvectors of a matrix of double precision real or complex numbers (i.e. RDF or CDF).

OUTPUT: Returns a list of triples, each of the form $(e, [v], 1)$, where e is the eigenvalue, and v is an associated left eigenvector. If the matrix is of size n , then there are n triples. Values are computed with the SciPy library.

The format of this output is designed to match the format for exact results. However, since matrices here have numerical entries, the resulting eigenvalues will also be numerical. No attempt is made to determine if two eigenvalues are equal, or if eigenvalues might actually be zero. So the algebraic multiplicity of each eigenvalue is reported as 1. Decisions about equal eigenvalues or zero eigenvalues should be addressed in the calling routine.

The SciPy routines used for these computations produce eigenvectors normalized to have length 1, but on different hardware they may vary by a sign. So for doctests we have normalized output by forcing their eigenvectors to have their first non-zero entry equal to one.

EXAMPLES:

```
sage: m = matrix(RDF, [[-5, 3, 2, 8], [10, 2, 4, -2], [-1, -10, -10, -17], [-2, 7, 6, 13]])
sage: m
[ -5.0   3.0   2.0   8.0]
[ 10.0   2.0   4.0  -2.0]
[ -1.0 -10.0 -10.0 -17.0]
[ -2.0   7.0   6.0  13.0]
sage: spectrum = m.left_eigenvectors()
sage: for i in range(len(spectrum)):
...     spectrum[i][1][0]=matrix(RDF, spectrum[i][1]).echelon_form()[0]
sage: spectrum[0]
(2.0, [(1.0, 1.0, 1.0, 1.0)], 1)
sage: spectrum[1]
(1.0, [(1.0, 0.8, 0.8, 0.6)], 1)
sage: spectrum[2]
(-2.0, [(1.0, 0.4, 0.6, 0.2)], 1)
sage: spectrum[3]
(-1.0, [(1.0, 1.0, 2.0, 2.0)], 1)
```

`log_determinant()`

Compute the log of the absolute value of the determinant using LU decomposition.

Note: This is useful if the usual determinant overflows.

EXAMPLES:

```
sage: m = matrix(RDF, 2, 2, range(4)); m
[0.0 1.0]
[2.0 3.0]
sage: RDF(log(abs(m.determinant())))
0.69314718056
sage: m.log_determinant()
0.69314718056
sage: m = matrix(RDF, 0, 0, []); m
[]
sage: m.log_determinant()
0.0
sage: m = matrix(CDF, 2, 2, range(4)); m
[0.0 1.0]
[2.0 3.0]
```

```

sage: RDF(log(abs(m.determinant()))))
0.69314718056
sage: m.log_determinant()
0.69314718056
sage: m = matrix(CDF, 0, 0, []); m
[]
sage: m.log_determinant()
0.0

```

norm(p=2)

Returns the norm of the matrix.

INPUT:

- p - default: 2 - controls which norm is computed, allowable values are ‘frob’ (for the Frobenius norm), integers -2, -1, 1, 2, positive and negative infinity. See output discussion for specifics.

OUTPUT:

Returned value is a double precision floating point value in RDF. Row and column sums described below are sums of the absolute values of the entries, where the absolute value of the complex number $a + bi$ is $\sqrt{a^2 + b^2}$. Singular values are the “diagonal” entries of the “S” matrix in the singular value decomposition.

- p = ‘frob’: the Frobenius norm, which for a matrix $A = (a_{ij})$ computes

$$\left(\sum_{i,j} |a_{i,j}|^2 \right)^{1/2}$$

- p = Infinity or p = oo: the maximum row sum.
- p = -Infinity or p = -oo: the minimum column sum.
- p = 1: the maximum column sum.
- p = -1: the minimum column sum.
- p = 2: the induced 2-norm, equal to the maximum singular value.
- p = -2: the minimum singular value.

ALGORITHM:

Computation is performed by the `norm()` function of the SciPy/NumPy library.

EXAMPLES:

First over the reals.

```

sage: A = matrix(RDF, 3, range(-3, 6)); A
[-3.0 -2.0 -1.0]
[ 0.0  1.0  2.0]
[ 3.0  4.0  5.0]
sage: A.norm()
7.99575670...
sage: A.norm(p='frob')
8.30662386...
sage: A.norm(p=Infinity)
12.0
sage: A.norm(p=-Infinity)
3.0
sage: A.norm(p=1)
8.0

```

```
sage: A.norm(p=-1)
6.0
sage: A.norm(p=2)
7.99575670...
sage: A.norm(p=-2) < 10^-15
True
```

And over the complex numbers.

```
sage: B = matrix(CDF, 2, [[1+I, 2+3*I], [3+4*I, 3*I]]); B
[1.0 + 1.0*I 2.0 + 3.0*I]
[3.0 + 4.0*I      3.0*I]
sage: B.norm()
6.66189877...
sage: B.norm(p='frob')
7.0
sage: B.norm(p=Infinity)
8.0
sage: B.norm(p=-Infinity)
5.01976483...
sage: B.norm(p=1)
6.60555127...
sage: B.norm(p=-1)
6.41421356...
sage: B.norm(p=2)
6.66189877...
sage: B.norm(p=-2)
2.14921023...
```

Since it is invariant under unitary multiplication, the Frobenius norm is equal to the square root of the sum of squares of the singular values.

```
sage: A = matrix(RDF, 5, range(1,26))
sage: f = A.norm(p='frob')
sage: U, S, V = A.SVD()
sage: s = sqrt(sum([S[i,i]^2 for i in range(5)]))
sage: abs(f-s) < 1.0e-12
True
```

Return values are in *RDF*.

```
sage: A = matrix(CDF, 2, range(4))
sage: A.norm() in RDF
True
```

Improper values of *p* are caught.

```
sage: A.norm(p='bogus')
Traceback (most recent call last):
...
ValueError: matrix norm 'p' must be +/- infinity, 'frob' or an integer, not bogus
sage: A.norm(p=632)
Traceback (most recent call last):
...
ValueError: matrix norm integer values of 'p' must be -2, -1, 1 or 2, not 632
```

numpy (*dtype=None*)

This method returns a copy of the matrix as a numpy array. It uses the numpy C/api so is very fast.

INPUT:

- `dtype` - The desired data-type for the array. If not given, then the type will be determined as the minimum type required to hold the objects in the sequence.

EXAMPLES:

```
sage: m = matrix(RDF, [[1, 2], [3, 4]])
sage: n = m.numpy()
sage: import numpy
sage: numpy.linalg.eig(n)
(array([-0.37228132,  5.37228132]), array([[ -0.82456484, -0.41597356],
      [ 0.56576746, -0.90937671]]))
sage: m = matrix(RDF, 2, range(6)); m
[0.0 1.0 2.0]
[3.0 4.0 5.0]
sage: m.numpy()
array([[ 0.,  1.,  2.],
       [ 3.,  4.,  5.]])
```

Alternatively, `numpy` automatically calls this function (via the magic `__array__()` method) to convert Sage matrices to `numpy` arrays:

```
sage: import numpy
sage: m = matrix(RDF, 2, range(6)); m
[0.0 1.0 2.0]
[3.0 4.0 5.0]
sage: numpy.array(m)
array([[ 0.,  1.,  2.],
       [ 3.,  4.,  5.]])
sage: numpy.array(m).dtype
dtype('float64')
sage: m = matrix(CDF, 2, range(6)); m
[0.0 1.0 2.0]
[3.0 4.0 5.0]
sage: numpy.array(m)
array([[ 0.+0.j,  1.+0.j,  2.+0.j],
       [ 3.+0.j,  4.+0.j,  5.+0.j]])
sage: numpy.array(m).dtype
dtype('complex128')
```

TESTS:

```
sage: m = matrix(RDF, 0, 5, []); m
[]
sage: m.numpy()
array([], shape=(0, 5), dtype=float64)
sage: m = matrix(RDF, 5, 0, []); m
[]
sage: m.numpy()
array([], shape=(5, 0), dtype=float64)
```

`right_eigenvectors()`

Compute the right eigenvectors of a matrix of double precision real or complex numbers (i.e. `RDF` or `CDF`).

OUTPUT:

Returns a list of triples, each of the form $(e, [v], 1)$, where e is the eigenvalue, and v is an associated right eigenvector. If the matrix is of size n , then there are n triples. Values are computed with the SciPy library.

The format of this output is designed to match the format for exact results. However, since matrices here

have numerical entries, the resulting eigenvalues will also be numerical. No attempt is made to determine if two eigenvalues are equal, or if eigenvalues might actually be zero. So the algebraic multiplicity of each eigenvalue is reported as 1. Decisions about equal eigenvalues or zero eigenvalues should be addressed in the calling routine.

The SciPy routines used for these computations produce eigenvectors normalized to have length 1, but on different hardware they may vary by a sign. So for doctests we have normalized output by forcing their eigenvectors to have their first non-zero entry equal to one.

EXAMPLES:

```
sage: m = matrix(RDF, [[-9, -14, 19, -74], [-1, 2, 4, -11], [-4, -12, 6, -32], [0, -2, -1, 1]])
sage: m
[ -9.0 -14.0  19.0 -74.0]
[ -1.0   2.0   4.0 -11.0]
[ -4.0 -12.0   6.0 -32.0]
[  0.0  -2.0  -1.0   1.0]
sage: spectrum = m.right_eigenvectors()
sage: for i in range(len(spectrum)):
...     spectrum[i][1][0]=matrix(RDF, spectrum[i][1]).echelon_form()[0]
sage: spectrum[0]
(2.0, [(1.0, -2.0, 3.0, 1.0)], 1)
sage: spectrum[1]
(1.0, [(1.0, -0.6666666666667, 1.3333333333333, 0.3333333333333)], 1)
sage: spectrum[2]
(-2.0, [(1.0, -0.2, 1.0, 0.2)], 1)
sage: spectrum[3]
(-1.0, [(1.0, -0.5, 2.0, 0.5)], 1)
```

round (*ndigits=0*)

Returns a copy of the matrix where all entries have been rounded to a given precision in decimal digits (default 0 digits).

INPUT:

- *ndigits* - The precision in number of decimal digits

OUTPUT:

A modified copy of the matrix

EXAMPLES:

```
sage: M=matrix(CDF, [[10.234r + 34.2343jr, 34e10r]])
sage: M
[10.234 + 34.2343*I      3.4e+11]
sage: M.round(2)
[10.23 + 34.23*I      3.4e+11]
sage: M.round()
[10.0 + 34.0*I      3.4e+11]
```

schur (*base_ring=None*)

Returns the Schur decomposition of the matrix.

INPUT:

- *base_ring* - optional, defaults to the base ring of *self*. Use this to request the base ring of the returned matrices, which will affect the format of the results.

OUTPUT:

A pair of immutable matrices. The first is a unitary matrix Q . The second, T , is upper-triangular when returned over the complex numbers, while it is almost upper-triangular over the reals. In the latter case,

there can be some 2×2 blocks on the diagonal which represent a pair of conjugate complex eigenvalues of `self`.

If `self` is the matrix A , then

$$A = QT(\overline{Q})^t$$

where the latter matrix is the conjugate-transpose of Q , which is also the inverse of Q , since Q is unitary.

Note that in the case of a normal matrix (Hermitian, symmetric, and others), the upper-triangular matrix is a diagonal matrix with eigenvalues of `self` on the diagonal, and the unitary matrix has columns that form an orthonormal basis composed of eigenvectors of `self`. This is known as “orthonormal diagonalization”.

Warning: The Schur decomposition is not unique, as there may be numerous choices for the vectors of the orthonormal basis, and consequently different possibilities for the upper-triangular matrix. However, the diagonal of the upper-triangular matrix will always contain the eigenvalues of the matrix (in the complex version), or 2×2 block matrices in the real version representing pairs of conjugate complex eigenvalues.

In particular, results may vary across systems and processors.

EXAMPLES:

First over the complexes. The similar matrix is always upper-triangular in this case.

```
sage: A = matrix(CDF, 4, 4, range(16)) + matrix(CDF, 4, 4, [x^3*I for x in range(0, 16)])
sage: Q, T = A.schur()
sage: (Q*Q.conjugate().transpose()).zero_at(1.0e-12)
[1.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0]
[0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 1.0]
sage: all([T.zero_at(1.0e-12)[i,j] == 0 for i in range(4) for j in range(i)])
True
sage: (Q*T*Q.conjugate().transpose()-A).zero_at(1.0e-11)
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
sage: eigenvalues = [T[i,i] for i in range(4)]; eigenvalues
[30.733... + 4648.541...*I, -0.184... - 159.057...*I, -0.523... + 11.158...*I, -0.025... - 0.000...*I]
sage: A.eigenvalues()
[30.733... + 4648.541...*I, -0.184... - 159.057...*I, -0.523... + 11.158...*I, -0.025... - 0.000...*I]
sage: abs(A.norm()-T.norm()) < 1e-10
True
```

We begin with a real matrix but ask for a decomposition over the complexes. The result will yield an upper-triangular matrix over the complex numbers for T .

```
sage: A = matrix(RDF, 4, 4, [x^3 for x in range(16)])
sage: Q, T = A.schur(base_ring=CDF)
sage: (Q*Q.conjugate().transpose()).zero_at(1.0e-12)
[1.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0]
[0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 1.0]
sage: T.parent()
Full MatrixSpace of 4 by 4 dense matrices over Complex Double Field
sage: all([T.zero_at(1.0e-12)[i,j] == 0 for i in range(4) for j in range(i)])
True
sage: (Q*T*Q.conjugate().transpose()-A).zero_at(1.0e-11)
```

```
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
```

Now totally over the reals. But with complex eigenvalues, the similar matrix may not be upper-triangular. But “at worst” there may be some 2×2 blocks on the diagonal which represent a pair of conjugate complex eigenvalues. These blocks will then just interrupt the zeros below the main diagonal. This example has a pair of these of the blocks.

```
sage: A = matrix(RDF, 4, 4, [[1, 0, -3, -1],
...                          [4, -16, -7, 0],
...                          [1, 21, 1, -2],
...                          [26, -1, -2, 1]])
sage: Q, T = A.schur()
sage: (Q*Q.conjugate().transpose()).zero_at(1.0e-12)
[1.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0]
[0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 1.0]
sage: all([T.zero_at(1.0e-12)[i,j] == 0 for i in range(4) for j in range(i)])
False
sage: all([T.zero_at(1.0e-12)[i,j] == 0 for i in range(4) for j in range(i-1)])
True
sage: (Q*T*Q.conjugate().transpose()-A).zero_at(1.0e-11)
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
sage: sorted(T[0:2,0:2].eigenvalues() + T[2:4,2:4].eigenvalues())
[-5.710... - 8.382...*I, -5.710... + 8.382...*I, -0.789... - 2.336...*I, -0.789... + 2.336...*I]
sage: sorted(A.eigenvalues())
[-5.710... - 8.382...*I, -5.710... + 8.382...*I, -0.789... - 2.336...*I, -0.789... + 2.336...*I]
sage: abs(A.norm()-T.norm()) < 1e-12
True
```

Starting with complex numbers and requesting a result over the reals will never happen.

```
sage: A = matrix(CDF, 2, 2, [[2+I, -1+3*I], [5-4*I, 2-7*I]])
sage: A.schur(base_ring=RDF)
Traceback (most recent call last):
...
TypeError: unable to convert input matrix over CDF to a matrix over RDF
```

If theory predicts your matrix is real, but it contains some very small imaginary parts, you can specify the cutoff for “small” imaginary parts, then request the output as real matrices, and let the routine do the rest.

```
sage: A = matrix(RDF, 2, 2, [1, 1, -1, 0]) + matrix(CDF, 2, 2, [1.0e-14*I]*4)
sage: B = A.zero_at(1.0e-12)
sage: B.parent()
Full MatrixSpace of 2 by 2 dense matrices over Complex Double Field
sage: Q, T = B.schur(RDF)
sage: Q.parent()
Full MatrixSpace of 2 by 2 dense matrices over Real Double Field
sage: T.parent()
Full MatrixSpace of 2 by 2 dense matrices over Real Double Field
sage: Q.round(6)
[ 0.707107  0.707107]
[-0.707107  0.707107]
```

```

sage: T.round(6)
[ 0.5  1.5]
[-0.5  0.5]
sage: (Q*T*Q.conjugate().transpose()-B).zero_at(1.0e-11)
[0.0 0.0]
[0.0 0.0]

```

A Hermitian matrix has real eigenvalues, so the similar matrix will be upper-triangular. Furthermore, a Hermitian matrix is diagonalizable with respect to an orthonormal basis, composed of eigenvectors of the matrix. Here that basis is the set of columns of the unitary matrix.

```

sage: A = matrix(CDF, [[ 52, -9*I - 8, 6*I - 187, -188*I + 2],
...                     [ 9*I - 8, 12, -58*I + 59, 30*I + 42],
...                     [-6*I - 187, 58*I + 59, 2677, 2264*I + 65],
...                     [188*I + 2, -30*I + 42, -2264*I + 65, 2080]])
sage: Q, T = A.schur()
sage: T = T.zero_at(1.0e-12).change_ring(RDF)
sage: T.round(6)
[4680.13301 0.0 0.0 0.0]
[ 0.0 102.715967 0.0 0.0]
[ 0.0 0.0 35.039344 0.0]
[ 0.0 0.0 0.0 3.11168]
sage: (Q*Q.conjugate().transpose()).zero_at(1.0e-12)
[1.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0]
[0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 1.0]
sage: (Q*T*Q.conjugate().transpose()-A).zero_at(1.0e-11)
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]

```

Similarly, a real symmetric matrix has only real eigenvalues, and there is an orthonormal basis composed of eigenvectors of the matrix.

```

sage: A = matrix(RDF, [[ 1, -2, 5, -3],
...                     [-2, 9, 1, 5],
...                     [ 5, 1, 3, 7],
...                     [-3, 5, 7, -8]])
sage: Q, T = A.schur()
sage: Q.round(4)
[-0.3027 -0.751 0.576 -0.1121]
[ 0.139 -0.3892 -0.2648 0.8713]
[ 0.4361 0.359 0.7599 0.3217]
[-0.836 0.3945 0.1438 0.3533]
sage: T = T.zero_at(10^-12)
sage: all(abs(e) < 10^-4 for e in (T - diagonal_matrix(RDF, [-13.5698, -0.8508, 7.7664, 11.6
True
sage: (Q*Q.transpose()).zero_at(1.0e-12)
[1.0 0.0 0.0 0.0]
[0.0 1.0 0.0 0.0]
[0.0 0.0 1.0 0.0]
[0.0 0.0 0.0 1.0]
sage: (Q*T*Q.transpose()-A).zero_at(1.0e-11)
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]
[0.0 0.0 0.0 0.0]

```

The results are cached, both as a real factorization and also as a complex factorization. This means the returned matrices are immutable.

```
sage: A = matrix(RDF, 2, 2, [[0, -1], [1, 0]])
sage: Qr, Tr = A.schur(base_ring=RDF)
sage: Qc, Tc = A.schur(base_ring=CDF)
sage: all([M.is_immutable() for M in [Qr, Tr, Qc, Tc]])
True
sage: Tr.round(6) != Tc.round(6)
True
```

TESTS:

The Schur factorization is only defined for square matrices.

```
sage: A = matrix(RDF, 4, 5, range(20))
sage: A.schur()
Traceback (most recent call last):
...
ValueError: Schur decomposition requires a square matrix, not a 4 x 5 matrix
```

A base ring request is checked.

```
sage: A = matrix(RDF, 3, range(9))
sage: A.schur(base_ring=QQ)
Traceback (most recent call last):
...
ValueError: base ring of Schur decomposition matrices must be RDF or CDF, not Rational Field
```

AUTHOR:

- Rob Beezer (2011-03-31)

singular_values (*eps=None*)

Returns a sorted list of the singular values of the matrix.

INPUT:

- eps* - default: *None* - the largest number which will be considered to be zero. May also be set to the string 'auto'. See the discussion below.

OUTPUT:

A sorted list of the singular values of the matrix, which are the diagonal entries of the “S” matrix in the SVD decomposition. As such, the values are real and are returned as elements of RDF. The list is sorted with larger values first, and since theory predicts these values are always positive, for a rank-deficient matrix the list should end in zeros (but in practice may not). The length of the list is the minimum of the row count and column count for the matrix.

The number of non-zero singular values will be the rank of the matrix. However, as a numerical matrix, it is impossible to control the difference between zero entries and very small non-zero entries. As an informed consumer it is up to you to use the output responsibly. We will do our best, and give you the tools to work with the output, but we cannot give you a guarantee.

With *eps* set to *None* you will get the raw singular values and can manage them as you see fit. You may also set *eps* to any positive floating point value you wish. If you set *eps* to 'auto' this routine will compute a reasonable cutoff value, based on the size of the matrix, the largest singular value and the smallest nonzero value representable by the 53-bit precision values used. See the discussion at page 268 of [WATKINS].

See the examples for a way to use the “verbose” facility to easily watch the zero cutoffs in action.

ALGORITHM:

The singular values come from the SVD decomposition computed by SciPy/NumPy.

EXAMPLES:

Singular values close to zero have trailing digits that may vary on different hardware. For exact matrices, the number of non-zero singular values will equal the rank of the matrix. So for some of the doctests we round the small singular values that ideally would be zero, to control the variability across hardware.

This matrix has a determinant of one. A chain of two or three theorems implies the product of the singular values must also be one.

```
sage: A = matrix(QQ, [[ 1, 0, 0, 0, 0, 1, 3],
...                  [-2, 1, 1, -2, 0, -4, 0],
...                  [ 1, 0, 1, -4, -6, -3, 7],
...                  [-2, 2, 1, 1, 7, 1, -1],
...                  [-1, 0, -1, 5, 8, 4, -6],
...                  [ 4, -2, -2, 1, -3, 0, 8],
...                  [-2, 1, 0, 2, 7, 3, -4]])
sage: A.determinant()
1
sage: B = A.change_ring(RDF)
sage: sv = B.singular_values(); sv
[20.5239806589, 8.48683702854, 5.86168134845, 2.44291658993,
 0.583197014472, 0.269332872866, 0.00255244880761]
sage: prod(sv)
1.0
```

An exact matrix that is obviously not of full rank, and then a computation of the singular values after conversion to an approximate matrix.

```
sage: A = matrix(QQ, [[1/3, 2/3, 11/3],
...                  [2/3, 1/3, 7/3],
...                  [2/3, 5/3, 27/3]])
sage: A.rank()
2
sage: B = A.change_ring(CDF)
sage: sv = B.singular_values()
sage: sv[0:2]
[10.1973039..., 0.487045871...]
sage: sv[2] < 1e-14
True
```

A matrix of rank 3 over the complex numbers.

```
sage: A = matrix(CDF, [[46*I - 28, -47*I - 50, 21*I + 51, -62*I - 782, 13*I + 22],
...                  [35*I - 20, -32*I - 46, 18*I + 43, -57*I - 670, 7*I + 3],
...                  [22*I - 13, -23*I - 23, 9*I + 24, -26*I - 347, 7*I + 13],
...                  [-44*I + 23, 41*I + 57, -19*I - 54, 60*I + 757, -11*I - 9],
...                  [30*I - 18, -30*I - 34, 14*I + 34, -42*I - 522, 8*I + 12]])
sage: sv = A.singular_values()
sage: sv[0:3]
[1440.733666, 18.4044034134, 6.83970779714]
sage: (10^-15 < sv[3]) and (sv[3] < 10^-13)
True
sage: (10^-16 < sv[4]) and (sv[4] < 10^-14)
True
```

A full-rank matrix that is ill-conditioned. We use this to illustrate ways of using the various possibilities for eps, including one that is ill-advised. Notice that the automatically computed cutoff gets this (difficult) example slightly wrong. This illustrates the impossibility of any automated process always getting this right. Use with caution and judgement.

```
sage: entries = [1/(i+j+1) for i in range(12) for j in range(12)]
sage: B = matrix(QQ, 12, 12, entries)
sage: B.rank()
12
sage: A = B.change_ring(RDF)
sage: A.condition() > 1.6e16 or A.condition()
True

sage: A.singular_values(eps=None) # abs tol 1e-16
[1.79537205956, 0.380275245955, 0.0447385487522, 0.00372231223789, 0.000233089089022, 1.1163...

sage: A.singular_values(eps='auto') # abs tol 1e-16
[1.79537205956, 0.380275245955, 0.0447385487522, 0.00372231223789, 0.000233089089022, 1.1163...

sage: A.singular_values(eps=1e-4)
[1.79537205956, 0.380275245955, 0.0447385487522, 0.00372231223789, 0.000233089089022, 0.0, 0...
```

With Sage’s “verbose” facility, you can compactly see the cutoff at work. In any application of this routine, or those that build upon it, it would be a good idea to conduct this exercise on samples. We also test here that all the values are returned in *RDF* since singular values are always real.

```
sage: A = matrix(CDF, 4, range(16))
sage: set_verbose(1)
sage: sv = A.singular_values(eps='auto'); sv
verbose 1 (<module>) singular values,
smallest-non-zero:cutoff:largest-zero,
2.2766...:6.2421...e-14:...
[35.139963659, 2.27661020871, 0.0, 0.0]
sage: set_verbose(0)

sage: all([s in RDF for s in sv])
True
```

TESTS:

Bogus values of the `eps` keyword will be caught.

```
sage: A.singular_values(eps='junk')
Traceback (most recent call last):
...
ValueError: could not convert string to float: junk
```

REFERENCES:

AUTHOR:

- Rob Beezer - (2011-02-18)

`solve_left(b)`

Solve the vector equation $x * A = b$ for a nonsingular A .

INPUT:

- `self` - a square matrix that is nonsingular (of full rank).
- `b` - a vector of the correct size. Elements of the vector must coerce into the base ring of the coefficient matrix. In particular, if `b` has entries from *CDF* then `self` must have *CDF* as its base ring.

OUTPUT:

The unique solution x to the matrix equation $x * A = b$, as a vector over the same base ring as `self`.

ALGORITHM:

Uses the `solve()` routine from the SciPy `scipy.linalg` module, after taking the transpose of the coefficient matrix.

EXAMPLES:

Over the reals.

```
sage: A = matrix(RDF, 3, 3, [1, 2, 5, 7.6, 2.3, 1, 1, 2, -1]); A
[ 1.0  2.0  5.0]
[ 7.6  2.3  1.0]
[ 1.0  2.0 -1.0]
sage: b = vector(RDF, [1, 2, 3])
sage: x = A.solve_left(b); x.zero_at(1e-18) # fix noisy zeroes
(0.666666666..., 0.0, 0.333333333...)
sage: x.parent()
Vector space of dimension 3 over Real Double Field
sage: x*A
(1.0, 2.0, 3.0)
```

Over the complex numbers.

```
sage: A = matrix(CDF, [[ 0, -1 + 2*I, 1 - 3*I, I],
... [2 + 4*I, -2 + 3*I, -1 + 2*I, -1 - I],
... [ 2 + I, 1 - I, -1, 5],
... [ 3*I, -1 - I, -1 + I, -3 + I]])
sage: b = vector(CDF, [2 - 3*I, 3, -2 + 3*I, 8])
sage: x = A.solve_left(b); x
(-1.55765124... - 0.644483985...*I, 0.183274021... + 0.286476868...*I, 0.270818505... + 0.24...
sage: x.parent()
Vector space of dimension 4 over Complex Double Field
sage: abs(x*A - b) < 1e-14
True
```

The vector of constants, `b`, can be given in a variety of forms, so long as it coerces to a vector over the same base ring as the coefficient matrix.

```
sage: A=matrix(CDF, 5, [1/(i+j+1) for i in range(5) for j in range(5)])
sage: A.solve_left([1]*5) # rel tol 1e-11
(5.0, -120.0, 630.0, -1120.0, 630.0)
```

TESTS:

A degenerate case.

```
sage: A = matrix(RDF, 0, 0, [])
sage: A.solve_left(vector(RDF, []))
()
```

The coefficient matrix must be square.

```
sage: A = matrix(RDF, 2, 3, range(6))
sage: b = vector(RDF, [1, 2, 3])
sage: A.solve_left(b)
Traceback (most recent call last):
...
ValueError: coefficient matrix of a system over RDF/CDF must be square, not 2 x 3
```

The coefficient matrix must be nonsingular.

```
sage: A = matrix(RDF, 5, range(25))
sage: b = vector(RDF, [1,2,3,4,5])
sage: A.solve_left(b)
Traceback (most recent call last):
...
LinAlgError: singular matrix
```

The vector of constants needs the correct degree.

```
sage: A = matrix(RDF, 5, range(25))
sage: b = vector(RDF, [1,2,3,4])
sage: A.solve_left(b)
Traceback (most recent call last):
...
TypeError: vector of constants over Real Double Field incompatible with matrix over Real Dou
```

The vector of constants needs to be compatible with the base ring of the coefficient matrix.

```
sage: F.<a> = FiniteField(27)
sage: b = vector(F, [a,a,a,a,a])
sage: A.solve_left(b)
Traceback (most recent call last):
...
TypeError: vector of constants over Finite Field in a of size 3^3 incompatible with matrix o
```

With a coefficient matrix over RDF, a vector of constants over CDF can be accommodated by converting the base ring of the coefficient matrix.

```
sage: A = matrix(RDF, 2, range(4))
sage: b = vector(CDF, [1+I,2])
sage: A.solve_left(b)
Traceback (most recent call last):
...
TypeError: vector of constants over Complex Double Field incompatible with matrix over Real

sage: B = A.change_ring(CDF)
sage: B.solve_left(b)
(0.5 - 1.5*I, 0.5 + 0.5*I)
```

solve_left_LU(*b*)

Solve the equation $Ax = b$ using LU decomposition.

Warning: This function is broken. See trac 4932.

INPUT:

- self – an invertible matrix
- b – a vector

Note: This method precomputes and stores the LU decomposition before solving. If many equations of the form $Ax=b$ need to be solved for a single matrix A , then this method should be used instead of `solve`. The first time this method is called it will compute the LU decomposition. If the matrix has not changed then subsequent calls will be very fast as the precomputed LU decomposition will be reused.

EXAMPLES:

```

sage: A = matrix(RDF, 3, 3, [1, 2, 5, 7.6, 2.3, 1, 1, 2, -1]); A
[ 1.0  2.0  5.0]
[ 7.6  2.3  1.0]
[ 1.0  2.0 -1.0]
sage: b = vector(RDF, [1, 2, 3])
sage: x = A.solve_left_LU(b); x
Traceback (most recent call last):
...
NotImplementedError: this function is not finished (see trac 4932)

```

TESTS:

We test two degenerate cases:

```

sage: A = matrix(RDF, 0, 3, [])
sage: A.solve_left_LU(vector(RDF, []))
(0.0, 0.0, 0.0)
sage: A = matrix(RDF, 3, 0, [])
sage: A.solve_left_LU(vector(RDF, 3, [1, 2, 3]))
()

```

solve_right(*b*)

Solve the vector equation $A \cdot x = b$ for a nonsingular A .

INPUT:

- *self* - a square matrix that is nonsingular (of full rank).
- *b* - a vector of the correct size. Elements of the vector must coerce into the base ring of the coefficient matrix. In particular, if *b* has entries from CDF then *self* must have CDF as its base ring.

OUTPUT:

The unique solution x to the matrix equation $A \cdot x = b$, as a vector over the same base ring as *self*.

ALGORITHM:

Uses the `solve()` routine from the SciPy `scipy.linalg` module.

EXAMPLES:

Over the reals.

```

sage: A = matrix(RDF, 3, 3, [1, 2, 5, 7.6, 2.3, 1, 1, 2, -1]); A
[ 1.0  2.0  5.0]
[ 7.6  2.3  1.0]
[ 1.0  2.0 -1.0]
sage: b = vector(RDF, [1, 2, 3])
sage: x = A.solve_right(b); x
(-0.113695090439, 1.39018087855, -0.333333333333)
sage: x.parent()
Vector space of dimension 3 over Real Double Field
sage: A*x
(1.0, 2.0, 3.0)

```

Over the complex numbers.

```

sage: A = matrix(CDF, [[ 0, -1 + 2*I, 1 - 3*I, I],
...                    [2 + 4*I, -2 + 3*I, -1 + 2*I, -1 - I],
...                    [ 2 + I, 1 - I, -1, 5],
...                    [ 3*I, -1 - I, -1 + I, -3 + I]])
sage: b = vector(CDF, [2 - 3*I, 3, -2 + 3*I, 8])
sage: x = A.solve_right(b); x

```

```
(1.96841637... - 1.07606761...*I, -0.614323843... + 1.68416370...*I, 0.0733985765... + 1.734...
sage: x.parent()
Vector space of dimension 4 over Complex Double Field
sage: abs(A*x - b) < 1e-14
True
```

The vector of constants, b , can be given in a variety of forms, so long as it coerces to a vector over the same base ring as the coefficient matrix.

```
sage: A=matrix(CDF, 5, [1/(i+j+1) for i in range(5) for j in range(5)])
sage: A.solve_right([1]*5) # rel tol 1e-11
(5.0, -120.0, 630.0, -1120.0, 630.0)
```

TESTS:

A degenerate case.

```
sage: A = matrix(RDF, 0, 0, [])
sage: A.solve_right(vector(RDF, []))
()
```

The coefficient matrix must be square.

```
sage: A = matrix(RDF, 2, 3, range(6))
sage: b = vector(RDF, [1,2,3])
sage: A.solve_right(b)
Traceback (most recent call last):
...
ValueError: coefficient matrix of a system over RDF/CDF must be square, not 2 x 3
```

The coefficient matrix must be nonsingular.

```
sage: A = matrix(RDF, 5, range(25))
sage: b = vector(RDF, [1,2,3,4,5])
sage: A.solve_right(b)
Traceback (most recent call last):
...
LinAlgError: singular matrix
```

The vector of constants needs the correct degree.

```
sage: A = matrix(RDF, 5, range(25))
sage: b = vector(RDF, [1,2,3,4])
sage: A.solve_right(b)
Traceback (most recent call last):
...
TypeError: vector of constants over Real Double Field incompatible with matrix over Real Dou
```

The vector of constants needs to be compatible with the base ring of the coefficient matrix.

```
sage: F.<a> = FiniteField(27)
sage: b = vector(F, [a,a,a,a,a])
sage: A.solve_right(b)
Traceback (most recent call last):
...
TypeError: vector of constants over Finite Field in a of size 3^3 incompatible with matrix o
```

With a coefficient matrix over RDF, a vector of constants over CDF can be accommodated by converting the base ring of the coefficient matrix.

```

sage: A = matrix(RDF, 2, range(4))
sage: b = vector(CDF, [1+I, 2])
sage: A.solve_right(b)
Traceback (most recent call last):
...
TypeError: vector of constants over Complex Double Field incompatible with matrix over Real Double Field

sage: B = A.change_ring(CDF)
sage: B.solve_right(b)
(-0.5 - 1.5*I, 1.0 + 1.0*I)

```

transpose()

Return the transpose of this matrix, without changing self.

EXAMPLES:

```

sage: m = matrix(RDF, 2, 3, range(6)); m
[0.0 1.0 2.0]
[3.0 4.0 5.0]
sage: m2 = m.transpose()
sage: m[0,0] = 2
sage: m2
[0.0 3.0]
[1.0 4.0]
[2.0 5.0]
#note that m2 hasn't changed

```

.T is a convenient shortcut for the transpose:

```

sage: m.T
[2.0 3.0]
[1.0 4.0]
[2.0 5.0]

```

```

sage: m = matrix(RDF, 0, 3); m
[]
sage: m.transpose()
[]
sage: m.transpose().parent()
Full MatrixSpace of 3 by 0 dense matrices over Real Double Field

```

zero_at(*eps*)

Returns a copy of the matrix where elements smaller than or equal to *eps* are replaced with zeroes. For complex matrices, the real and imaginary parts are considered individually.

This is useful for modifying output from algorithms which have large relative errors when producing zero elements, e.g. to create reliable doctests.

INPUT:

- *eps* - Cutoff value

OUTPUT:

A modified copy of the matrix.

EXAMPLES:

```

sage: a=matrix([[1, 1e-4r, 1+1e-100jr], [1e-8+3j, 0, 1e-58r]])
sage: a
[
      1.0      0.0001 1.0 + 1e-100*I
[ 1e-08 + 3.0*I      0.0      1e-58]

```

```
sage: a.zero_at(1e-50)
[      1.0      0.0001      1.0]
[1e-08 + 3.0*I      0.0      0.0]
sage: a.zero_at(1e-4)
[  1.0  0.0  1.0]
[3.0*I  0.0  0.0]
```

DENSE MATRICES OVER THE REAL DOUBLE FIELD USING NUMPY

Dense matrices over the Real Double Field using NumPy

EXAMPLES:

```
sage: b=Mat(RDF,2,3).basis()
sage: b[0]
[1.0 0.0 0.0]
[0.0 0.0 0.0]
```

We deal with the case of zero rows or zero columns:

```
sage: m = MatrixSpace(RDF,0,3)
sage: m.zero_matrix()
[]
```

TESTS:

```
sage: a = matrix(RDF,2,range(4), sparse=False)
sage: TestSuite(a).run()
sage: MatrixSpace(RDF,0,0).zero_matrix().inverse()
[]
```

AUTHORS:

- Jason Grout (2008-09): switch to NumPy backend, factored out the `Matrix_double_dense` class
- Josh Kantor
- William Stein: many bug fixes and touch ups.

```
class sage.matrix.matrix_real_double_dense.Matrix_real_double_dense
Bases: sage.matrix.matrix_double_dense.Matrix_double_dense
```

Class that implements matrices over the real double field. These are supposed to be fast matrix operations using C doubles. Most operations are implemented using numpy which will call the underlying BLAS on the system.

EXAMPLES:

```
sage: m = Matrix(RDF, [[1,2],[3,4]])
sage: m**2
[ 7.0 10.0]
[15.0 22.0]
sage: n= m^(-1); n
```

```
[-2.0  1.0]
[ 1.5 -0.5]
```

To compute eigenvalues the use the functions `left_eigenvectors` or `right_eigenvectors`

```
sage: p,e = m.right_eigenvectors()
```

the result of `eigen` is a pair (p,e) , where p is a list of eigenvalues and the e is a matrix whose columns are the eigenvectors.

To solve a linear system $Ax = b$ where $A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix}$ and $b = \begin{bmatrix} 5 \\ 6 \end{bmatrix}$.

```
sage: b = vector(RDF, [5, 6])
```

```
sage: m.solve_right(b)
(-4.0, 4.5)
```

See the commands `qr`, `lu`, and `svd` for QR, LU, and singular value decomposition.

DENSE MATRICES OVER THE COMPLEX DOUBLE FIELD USING NUMPY

Dense matrices over the Complex Double Field using NumPy

EXAMPLES:

```
sage: b=Mat(CDF,2,3).basis()
sage: b[0]
[1.0 0.0 0.0]
[0.0 0.0 0.0]
```

We deal with the case of zero rows or zero columns:

```
sage: m = MatrixSpace(CDF,0,3)
sage: m.zero_matrix()
[]
```

TESTS:

```
sage: a = matrix(CDF,2,[i+(4-i)*I for i in range(4)], sparse=False)
sage: TestSuite(a).run()
sage: Mat(CDF,0,0).zero_matrix().inverse()
[]
```

AUTHORS:

- Jason Grout (2008-09): switch to NumPy backend
- Josh Kantor
- William Stein: many bug fixes and touch ups.

```
class sage.matrix.matrix_complex_double_dense.Matrix_complex_double_dense
Bases: sage.matrix.matrix_double_dense.Matrix_double_dense
```

Class that implements matrices over the real double field. These are supposed to be fast matrix operations using C doubles. Most operations are implemented using numpy which will call the underlying BLAS on the system.

EXAMPLES:

```
sage: m = Matrix(CDF, [[1,2*I],[3+I,4]])
sage: m**2
[-1.0 + 6.0*I      10.0*I]
```

```
[15.0 + 5.0*I 14.0 + 6.0*I]
```

```
sage: n= m^(-1); n
```

```
[ 0.333333333333 + 0.333333333333*I 0.166666666667 - 0.166666666667*I]
[ -0.166666666667 - 0.333333333333*I 0.083333333333 + 0.083333333333*I]
```

To compute eigenvalues the use the functions `left_eigenvectors` or `right_eigenvectors`

```
sage: p,e = m.right_eigenvectors()
```

the result of `eigen` is a pair (p,e), where p is a list of eigenvalues and the e is a matrix whose columns are the eigenvectors.

To solve a linear system $Ax = b$ where $A = \begin{bmatrix} 1 & 2 \end{bmatrix}$ and $b = \begin{bmatrix} 5 & 6 \end{bmatrix}$

```
sage: b = vector(CDF, [5, 6])
```

```
sage: m.solve_right(b)
```

```
(2.66666666667 + 0.66666666667*I, -0.333333333333 - 1.16666666667*I)
```

See the commands `qr`, `lu`, and `svd` for QR, LU, and singular value decomposition.

DENSE MATRICES OVER MULTIVARIATE POLYNOMIALS OVER FIELDS

Dense matrices over multivariate polynomials over fields

This implementation inherits from `Matrix_generic_dense`, i.e. it is not optimized for speed only some methods were added.

AUTHOR:

- Martin Albrecht <malb@informatik.uni-bremen.de>

```
class sage.matrix.matrix_mpolynomial_dense.Matrix_mpolynomial_dense
    Bases: sage.matrix.matrix_generic_dense.Matrix_generic_dense
```

Dense matrix over a multivariate polynomial ring over a field.

determinant (*algorithm=None*)

Return the determinant of this matrix

INPUT:

- `algorithm` – ignored

EXAMPLES:

We compute the determinant of the arbitrary 3×3 matrix:

```
sage: R = PolynomialRing(QQ, 9, 'x')
sage: A = matrix(R, 3, R.gens())
sage: A
[x0 x1 x2]
[x3 x4 x5]
[x6 x7 x8]
sage: A.determinant()
-x2*x4*x6 + x1*x5*x6 + x2*x3*x7 - x0*x5*x7 - x1*x3*x8 + x0*x4*x8
```

We check if two implementations agree on the result:

```
sage: R.<x,y> = QQ[]
sage: C = random_matrix(R, 2, 2, terms=2)
sage: C
[-6/5*x*y - y^2 -6*y^2 - 1/4*y]
[ -1/3*x*y - 3      x*y - x]
sage: C.determinant()
```

$$-6/5*x^2*y^2 - 3*x*y^3 + 6/5*x^2*y + 11/12*x*y^2 - 18*y^2 - 3/4*y$$

```
sage: C.change_ring(R.change_ring(QQbar)).det()
(-6/5)*x^2*y^2 + (-3)*x*y^3 + 6/5*x^2*y + 11/12*x*y^2 + (-18)*y^2 + (-3/4)*y
```

Finally, we check whether the Singular interface is working:

```
sage: R.<x,y> = RR[]
sage: C = random_matrix(R, 2, 2, terms=2)
sage: C
[0.368965517352886*y^2 + 0.425700773972636*x -0.800362171389760*y^2 - 0.807635502485287]
[ 0.706173539423122*y^2 - 0.915986060298440 0.897165181570476*y + 0.107903328188376]
sage: C.determinant()
0.565194587390682*y^4 + 0.331023015369146*y^3 + 0.381923912175852*x*y - 0.122977163520282*y^2
```

ALGORITHM: Calls Singular, libSingular or native implementation.

TESTS:

```
sage: R = PolynomialRing(QQ, 9, 'x')
sage: matrix(R, 0, 0).det()
1

sage: R.<h,y> = QQ[]
sage: m = matrix([[y,y,y,y]] * 4) # larger than 3x3
sage: m.charpoly() # put charpoly in the cache
x^4 - 4*y*x^3
sage: m.det()
0
```

echelon_form(algorithm='row_reduction', **kws)

Return an echelon form of self using chosen algorithm.

By default only a usual row reduction with no divisions or column swaps is returned.

If Gauss-Bareiss algorithm is chosen, column swaps are recorded and can be retrieved via `swapped_columns()`.

INPUT:

- `algorithm` – string, which algorithm to use (default: 'row_reduction'). Valid options are:
 - 'row_reduction' (default) – reduce as far as possible, only divide by constant entries
 - 'frac' – reduced echelon form over fraction field
 - 'bareiss' – fraction free Gauss-Bareiss algorithm with column swaps

OUTPUT:

The row echelon form of A depending on the chosen algorithm, as an immutable matrix. Note that `self` is *not* changed by this command. Use `A.echelonize()` to change A in place.

EXAMPLES:

```
sage: P.<x,y> = PolynomialRing(GF(127), 2)
sage: A = matrix(P, 2, 2, [1, x, 1, y])
sage: A
[1 x]
[1 y]
sage: A.echelon_form()
[ 1 x]
[ 0 -x + y]
```

The reduced row echelon form over the fraction field is as follows:

```
sage: A.echelon_form('frac') # over fraction field
[1 0]
[0 1]
```

Alternatively, the Gauss-Bareiss algorithm may be chosen:

```
sage: E = A.echelon_form('bareiss'); E
[ 1      y]
[ 0 x - y]
```

After the application of the Gauss-Bareiss algorithm the swapped columns may inspected:

```
sage: E.swapped_columns(), E.pivots()
((0, 1), (0, 1))
sage: A.swapped_columns(), A.pivots()
(None, (0, 1))
```

Another approach is to row reduce as far as possible:

```
sage: A.echelon_form('row_reduction')
[ 1      x]
[ 0 -x + y]
```

echelonize (*algorithm*='row_reduction', ***kws*)

Transform self into a matrix in echelon form over the same base ring as *self*.

If Gauss-Bareiss algorithm is chosen, column swaps are recorded and can be retrieved via `swapped_columns()`.

INPUT:

- *algorithm* – string, which algorithm to use. Valid options are:
 - 'row_reduction' – reduce as far as possible, only divide by constant entries
 - 'bareiss' – fraction free Gauss-Bareiss algorithm with column swaps

EXAMPLES:

```
sage: P.<x,y> = PolynomialRing(QQ, 2)
sage: A = matrix(P, 2, 2, [1/2, x, 1, 3/4*y+1])
sage: A
[ 1/2      x]
[ 1 3/4*y + 1]
```

```
sage: B = copy(A)
sage: B.echelonize('bareiss'); B
[      1      3/4*y + 1]
[      0 x - 3/8*y - 1/2]
```

```
sage: B = copy(A)
sage: B.echelonize('row_reduction'); B
[      1      2*x]
[      0 -2*x + 3/4*y + 1]
```

```
sage: P.<x,y> = PolynomialRing(QQ, 2)
sage: A = matrix(P, 2, 3, [2, x, 0, 3, y, 1]); A
[2 x 0]
[3 y 1]
```

```
sage: E = A.echelon_form('bareiss'); E
```

```
[1 3 y]
[0 2 x]
sage: E.swapped_columns()
(2, 0, 1)
sage: A.pivots()
(0, 1, 2)
```

pivots()

Return the pivot column positions of this matrix as a list of integers.

This returns a list, of the position of the first nonzero entry in each row of the echelon form.

OUTPUT:

A list of Python ints.

EXAMPLES:

```
sage: matrix([PolynomialRing(GF(2), 2, 'x').gen()]).pivots()
(0,)
sage: K = QQ['x,y']
sage: x, y = K.gens()
sage: m = matrix(K, [(-x, 1, y, x - y), (-x*y, y, y^2 - 1, x*y - y^2 + x), (-x*y + x, y - 1,
sage: m.pivots()
(0, 2)
sage: m.rank()
2
```

swapped_columns()

Return which columns were swapped during the Gauss-Bareiss reduction

OUTPUT:

Return a tuple representing the column swaps during the last application of the Gauss-Bareiss algorithm (see `echelon_form()` for details).

The tuple as length equal to the rank of self and the value at the i -th position indicates the source column which was put as the i -th column.

If no Gauss-Bareiss reduction was performed yet, None is returned.

EXAMPLES:

```
sage: R.<x,y> = QQ[]
sage: C = random_matrix(R, 2, 2, terms=2)
sage: C.swapped_columns()
sage: E = C.echelon_form('bareiss')
sage: E.swapped_columns()
(0, 1)
```

OPERATION TABLES

This module implements general operation tables, which are very matrix-like.

```
class sage.matrix.operation_table.OperationTable(S, operation, names='letters', elements=None)
```

Bases: `sage.structure.sage_object.SageObject`

An object that represents a binary operation as a table.

Primarily this object is used to provide a `multiplication_table()` for objects in the category of magmas (monoids, groups, ...) and `addition_table()` for objects in the category of commutative additive magmas (additive monoids, groups, ...).

INPUT:

- **S** - a finite algebraic structure (or finite iterable)
- **operation** - a function of two variables that accepts pairs of elements from *S*. A natural source of such functions is the Python `operator` module, and in particular `operator.add()` and `operator.mul()`. This may also be a function defined with `lambda` or `def`.
- **names** - (default: 'letters') The type of names used, values are:
 - 'letters' - lowercase ASCII letters are used for a base 26 representation of the elements' positions in the list given by `column_keys()`, padded to a common width with leading 'a's.
 - 'digits' - base 10 representation of the elements' positions in the list given by `column_keys()`, padded to a common width with leading zeros.
 - 'elements' - the string representations of the elements themselves.
 - a list - a list of strings, where the length of the list equals the number of elements.
- **elements** - (default: None) A list of elements of *S*, in forms that can be coerced into the structure, eg. their string representations. This may be used to impose an alternate ordering on the elements of *S*, perhaps when this is used in the context of a particular structure. The default is to use whatever ordering the `S.list()` method returns. *elements* can also be a subset which is closed under the operation, useful perhaps when the set is infinite.

OUTPUT: An object with methods that abstracts multiplication tables, addition tables, Cayley tables, etc. It should be general enough to be useful for any finite algebraic structure whose elements can be combined with a binary operation. This is not necessarily meant to be constructed directly, but instead should be useful for constructing operation tables of various algebraic structures that have binary operations.

EXAMPLES:

In it's most basic use, the table needs a structure and an operation:

```

sage: from sage.matrix.operation_table import OperationTable
sage: G=SymmetricGroup(3)
sage: OperationTable(G, operation = operator.mul)
*  a b c d e f
+-----+
a| a b c d e f
b| b a d c f e
c| c e a f b d
d| d f b e a c
e| e c f a d b
f| f d e b c a

```

With two operations present, we can specify which operation we want:

```

sage: from sage.matrix.operation_table import OperationTable
sage: R=Integers(6)
sage: OperationTable(R, operation=operator.add)
+  a b c d e f
+-----+
a| a b c d e f
b| b c d e f a
c| c d e f a b
d| d e f a b c
e| e f a b c d
f| f a b c d e

```

The default symbol set for elements is lowercase ASCII letters, which take on a base 26 flavor for structures with more than 26 elements.

```

sage: from sage.matrix.operation_table import OperationTable
sage: G=DihedralGroup(14)
sage: OperationTable(G, operator.mul, names='letters')
*  aa ab ac ad ae af ag ah ai aj ak al am an ao ap aq ar as at au av aw ax ay az ba bb
+-----+
aa| aa ab ac ad ae af ag ah ai aj ak al am an ao ap aq ar as at au av aw ax ay az ba bb
ab| ab aa ad ac af ae ah ag aj ai al ak an am ap ao ar aq at as av au ax aw az ay bb ba
ac| ac ba aa ae ad ag af ai ah ak aj am al ao an aq ap as ar au at aw av ay ax bb ab az
ad| ad bb ab af ac ah ae aj ag al ai an ak ap am ar ao at aq av as ax au az aw ba aa ay
ae| ae az ba ag aa ai ad ak af am ah ao aj aq al as an au ap aw ar ay at bb av ab ac ax
af| af ay bb ah ab aj ac al ae an ag ap ai ar ak at am av ao ax aq az as ba au aa ad aw
ag| ag ax az ai ba ak aa am ad ao af aq ah as aj au al aw an ay ap bb ar ab at ac ae av
ah| ah aw ay aj bb al ab an ac ap ae ar ag at ai av ak ax am az ao ba aq aa as ad af au
ai| ai av ax ak az am ba ao aa aq ad as af au ah aw aj ay al bb an ab ap ac ar ae ag at
aj| aj au aw al ay an bb ap ab ar ac at ae av ag ax ai az ak ba am aa ao ad aq af ah as
ak| ak at av am ax ao az aq ba as aa au ad aw af ay ah bb aj ab al ac an ae ap ag ai ar
al| al as au an aw ap ay ar bb at ab av ac ax ae az ag ba ai aa ak ad am af ao ah aj aq
am| am ar at ao av aq ax as az au ba aw aa ay ad bb af ab ah ac aj ae al ag an ai ak ap
an| an aq as ap au ar aw at ay av bb ax ab az ac ba ae aa ag ad ai af ak ah am aj al ao
ao| ao ap ar aq at as av au ax aw az ay ba bb aa ab ad ac af ae ah ag aj ai al ak am an
ap| ap ao aq ar as at au av aw ax ay az bb ba ab aa ac ad ae af ag ah ai aj ak al an am
aq| aq an ap as ar au at aw av ay ax bb az ab ba ac aa ae ad ag af ai ah ak aj am ao al
ar| ar am ao at aq av as ax au az aw ba ay aa bb ad ab af ac ah ae aj ag al ai an ap ak
as| as al an au ap aw ar ay at bb av ab ax ac az ae ba ag aa ai ad ak af am ah ao aq aj
at| at ak am av ao ax aq az as ba au aa aw ad ay af bb ah ab aj ac al ae an ag ap ar ai
au| au aj al aw an ay ap bb ar ab at ac av ae ax ag az ai ba ak aa am ad ao af aq as ah
av| av ai ak ax am az ao ba aq aa as ad au af aw ah ay aj bb al ab an ac ap ae ar at ag
aw| aw ah aj ay al bb an ab ap ac ar ae at ag av ai ax ak az am ba ao aa aq ad as au af
ax| ax ag ai az ak ba am aa ao ad aq af as ah au aj aw al ay an bb ap ab ar ac at av ae
ay| ay af ah bb aj ab al ac an ae ap ag ar ai at ak av am ax ao az aq ba as aa au aw ad

```



```
az| az ae ag ba ai aa ak ad am af ao ah aq aj as al au an aw ap ay ar bb at ab av ax ac
ba| ba ac ae aa ag ad ai af ak ah am aj ao al aq an as ap au ar aw at ay av bb ax az ab
bb| bb ad af ab ah ac aj ae al ag an ai ap ak ar am at ao av aq ax as az au ba aw ay aa
```

Another symbol set is base 10 digits, padded with leading zeros to make a common width.

```
sage: from sage.matrix.operation_table import OperationTable
sage: G=AlternatingGroup(4)
sage: OperationTable(G, operator.mul, names='digits')
*  00 01 02 03 04 05 06 07 08 09 10 11
+-----+
00| 00 01 02 03 04 05 06 07 08 09 10 11
01| 01 02 00 05 03 04 07 08 06 11 09 10
02| 02 00 01 04 05 03 08 06 07 10 11 09
03| 03 06 09 00 07 10 01 04 11 02 05 08
04| 04 08 10 02 06 11 00 05 09 01 03 07
05| 05 07 11 01 08 09 02 03 10 00 04 06
06| 06 09 03 10 00 07 04 11 01 08 02 05
07| 07 11 05 09 01 08 03 10 02 06 00 04
08| 08 10 04 11 02 06 05 09 00 07 01 03
09| 09 03 06 07 10 00 11 01 04 05 08 02
10| 10 04 08 06 11 02 09 00 05 03 07 01
11| 11 05 07 08 09 01 10 02 03 04 06 00
```

If the group's elements are not too cumbersome, or the group is small, then the string representation of the elements can be used.

```
sage: from sage.matrix.operation_table import OperationTable
sage: G=AlternatingGroup(3)
sage: OperationTable(G, operator.mul, names='elements')
*      () (1,2,3) (1,3,2)
+-----+
() |      () (1,2,3) (1,3,2)
(1,2,3) | (1,2,3) (1,3,2)      ()
(1,3,2) | (1,3,2)      () (1,2,3)
```

You can give the elements any names you like, but they need to be ordered in the same order as returned by the `column_keys()` method.

```
sage: from sage.matrix.operation_table import OperationTable
sage: G=QuaternionGroup()
sage: T=OperationTable(G, operator.mul)
sage: T.column_keys()
(((), (1,2,3,4) (5,6,7,8)), ..., (1,8,3,6) (2,7,4,5))
sage: names=['1', 'I', '-1', '-I', 'J', '-K', '-J', 'K']
sage: T.change_names(names=names)
sage: sorted(T.translation().items())
[('-1', (1,3) (2,4) (5,7) (6,8)), ..., ('K', (1,8,3,6) (2,7,4,5))]
sage: T
*  1  I -1 -I  J -K -J  K
+-----+
1|  1  I -1 -I  J -K -J  K
I|  I -1 -I  1  K  J -K -J
-1| -1 -I  1  I -J  K  J -K
-I| -I  1  I -1 -K -J  K  J
J|  J -K -J  K -1 -I  1  I
-K| -K -J  K  J  I -1 -I  1
-J| -J  K  J -K  1  I -1 -I
K|  K  J -K -J -I  1  I -1
```

With the right functions and a list comprehension, custom names can be easier. A multiplication table for hex digits (without carries):

```
sage: from sage.matrix.operation_table import OperationTable
sage: R=Integers(16)
sage: names=[hex(Integer(a)) for a in R]
sage: OperationTable(R, operation=operator.mul, names=names)
*  0 1 2 3 4 5 6 7 8 9 a b c d e f
+-----+
0| 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1| 0 1 2 3 4 5 6 7 8 9 a b c d e f
2| 0 2 4 6 8 a c e 0 2 4 6 8 a c e
3| 0 3 6 9 c f 2 5 8 b e 1 4 7 a d
4| 0 4 8 c 0 4 8 c 0 4 8 c 0 4 8 c
5| 0 5 a f 4 9 e 3 8 d 2 7 c 1 6 b
6| 0 6 c 2 8 e 4 a 0 6 c 2 8 e 4 a
7| 0 7 e 5 c 3 a 1 8 f 6 d 4 b 2 9
8| 0 8 0 8 0 8 0 8 0 8 0 8 0 8 0 8
9| 0 9 2 b 4 d 6 f 8 1 a 3 c 5 e 7
a| 0 a 4 e 8 2 c 6 0 a 4 e 8 2 c 6
b| 0 b 6 1 c 7 2 d 8 3 e 9 4 f a 5
c| 0 c 8 4 0 c 8 4 0 c 8 4 0 c 8 4
d| 0 d a 7 4 1 e b 8 5 2 f c 9 6 3
e| 0 e c a 8 6 4 2 0 e c a 8 6 4 2
f| 0 f e d c b a 9 8 7 6 5 4 3 2 1
```

This should be flexible enough to create a variety of such tables.

```
sage: from sage.matrix.operation_table import OperationTable
sage: from operator import xor
sage: T=OperationTable(ZZ, xor, elements=range(8))
sage: T
.  a b c d e f g h
+-----+
a| a b c d e f g h
b| b a d c f e h g
c| c d a b g h e f
d| d c b a h g f e
e| e f g h a b c d
f| f e h g b a d c
g| g h e f c d a b
h| h g f e d c b a
sage: names=['000', '001', '010', '011', '100', '101', '110', '111']
sage: T.change_names(names)
sage: T.set_print_symbols('^', '\\land')
sage: T
^  000 001 010 011 100 101 110 111
+-----+
000| 000 001 010 011 100 101 110 111
001| 001 000 011 010 101 100 111 110
010| 010 011 000 001 110 111 100 101
011| 011 010 001 000 111 110 101 100
100| 100 101 110 111 000 001 010 011
101| 101 100 111 110 001 000 011 010
110| 110 111 100 101 010 011 000 001
111| 111 110 101 100 011 010 001 000

sage: T = OperationTable([False, True], operator.or_, names = 'elements')
sage: T
.  False  True
```

```

+-----+
False| False  True
True|  True  True

```

TESTS:

Empty structures behave acceptably, though the ASCII table looks a bit odd. The LaTeX version works much better.

```

sage: from sage.matrix.operation_table import OperationTable
sage: L=FiniteSemigroups().example(())
sage: L

```

An example of a finite semigroup: the left regular band generated by ()

```

sage: T=OperationTable(L, operation=operator.mul)
sage: T

```

```

*
+

```

```

sage: T._latex_()

```

```

'\setlength{\arraycolsep}{2ex}\n\begin{array}{r|*{0}{r}}\n\multicolumn{1}{c|}{\ast}\n\\hline

```

If the algebraic structure cannot be listed (like when it is infinite) then there is no way to create a table.

```

sage: from sage.matrix.operation_table import OperationTable
sage: OperationTable(ZZ, operator.mul)
Traceback (most recent call last):
...
ValueError: Integer Ring is infinite

```

The value of elements must be a subset of the algebraic structure, in forms that can be coerced into the structure. Here we demonstrate the proper use first:

```

sage: from sage.matrix.operation_table import OperationTable
sage: H=CyclicPermutationGroup(4)
sage: H.list()
[(), (1,2,3,4), (1,3)(2,4), (1,4,3,2)]
sage: elts = ['()', '(1,3)(2,4)']
sage: OperationTable(H, operator.mul, elements=elts)
*   a b
+----
a| a b
b| b a

```

This can be rewritten so as to pass the actual elements of the group H, using a simple for loop:

```

sage: L = H.list()      #list of elements of the group H
sage: elts = [L[i] for i in {0, 2}]
sage: elts
[(), (1,3)(2,4)]
sage: OperationTable(H, operator.mul, elements=elts)
*   a b
+----
a| a b
b| b a

```

Here are a couple of improper uses

```

sage: elts.append(5)
sage: OperationTable(H, operator.mul, elements=elts)
Traceback (most recent call last):
...

```

```
TypeError: unable to coerce 5 into Cyclic group of order 4 as a permutation group
sage: elts[2]='(1,3,2,4)'
sage: OperationTable(H, operator.mul, elements=elts)
Traceback (most recent call last):
...
TypeError: unable to coerce (1,3,2,4) into Cyclic group of order 4 as a permutation group
sage: elts[2]='(1,2,3,4)'
sage: OperationTable(H, operator.mul, elements=elts)
Traceback (most recent call last):
...
ValueError: (1,3) (2,4) * (1,2,3,4)=(1,4,3,2), and so the set is not closed
```

Unusable functions should be recognized as such:

```
sage: H=CyclicPermutationGroup(4)
sage: OperationTable(H, operator.add)
Traceback (most recent call last):
...
TypeError: elements () and () of Cyclic group of order 4 as a permutation group are incompatible
sage: from operator import xor
sage: OperationTable(H, xor)
Traceback (most recent call last):
...
TypeError: elements () and () of Cyclic group of order 4 as a permutation group are incompatible
```

TODO:

Provide color and grayscale graphical representations of tables. See commented-out stubs in source code.

AUTHOR:

- Rob Beezer (2010-03-15)

change_names (*names*)

For an existing operation table, change the names used for the elements.

INPUT:

- names - the type of names used, values are:
 - 'letters' - lowercase ASCII letters are used for a base 26 representation of the elements' positions in the list given by `list()`, padded to a common width with leading 'a's.
 - 'digits' - base 10 representation of the elements' positions in the list given by `list()`, padded to a common width with leading zeros.
 - 'elements' - the string representations of the elements themselves.
- a list - a list of strings, where the length of the list equals the number of elements.

OUTPUT: None. This method changes the table “in-place”, so any printed version will change and the output of the `dict()` will also change. So any items of interest about a particular table need to be copied/saved prior to calling this method.

EXAMPLES:

More examples can be found in the documentation for `OperationTable` since creating a new operation table uses the same routine.

```
sage: from sage.matrix.operation_table import OperationTable
sage: D=DihedralGroup(2)
sage: T=OperationTable(D, operator.mul)
sage: T
```

```

*   a b c d
+-----+
a| a b c d
b| b a d c
c| c d a b
d| d c b a
sage: T.translation()['c']
(1,2)
sage: T.change_names('digits')
sage: T
*   0 1 2 3
+-----+
0| 0 1 2 3
1| 1 0 3 2
2| 2 3 0 1
3| 3 2 1 0
sage: T.translation()['2']
(1,2)
sage: T.change_names('elements')
sage: T
*               ()      (3,4)      (1,2) (1,2) (3,4)
+-----+
() |               ()      (3,4)      (1,2) (1,2) (3,4)
(3,4) |           (3,4)      () (1,2) (3,4)      (1,2)
(1,2) |           (1,2) (1,2) (3,4)      ()      (3,4)
(1,2) (3,4) | (1,2) (3,4)      (1,2)      (3,4)      ()
sage: T.translation()['(1,2)']
(1,2)
sage: T.change_names(['w', 'x', 'y', 'z'])
sage: T
*   w x y z
+-----+
w| w x y z
x| x w z y
y| y z w x
z| z y x w
sage: T.translation()['y']
(1,2)

```

column_keys()

Returns a tuple of the elements used to build the table.

Note: `column_keys` and `row_keys` are identical. Both list the elements in the order used to label the table.

OUTPUT:

The elements of the algebraic structure used to build the table, as a list. But most importantly, elements are present in the list in the order which they appear in the table's column headings.

EXAMPLES:

```

sage: from sage.matrix.operation_table import OperationTable
sage: G=AlternatingGroup(3)
sage: T=OperationTable(G, operator.mul)
sage: T.column_keys()
((), (1,2,3), (1,3,2))

```

matrix_of_variables()

This method provides some backward compatibility for Cayley tables of groups, whose output was restricted to this single format.

EXAMPLES:

The output here is from the doctests for the old `cayley_table()` method for permutation groups.

```
sage: from sage.matrix.operation_table import OperationTable
sage: G=PermutationGroup(['(1,2,3)', '(2,3)'])
sage: T=OperationTable(G, operator.mul)
sage: T.matrix_of_variables()
[x0 x1 x2 x3 x4 x5]
[x1 x0 x3 x2 x5 x4]
[x2 x4 x0 x5 x1 x3]
[x3 x5 x1 x4 x0 x2]
[x4 x2 x5 x0 x3 x1]
[x5 x3 x4 x1 x2 x0]
sage: T.column_keys()[3]*T.column_keys()[3] == T.column_keys()[4]
True
```

row_keys()

Returns a tuple of the elements used to build the table.

Note: `column_keys` and `row_keys` are identical. Both list the elements in the order used to label the table.

OUTPUT:

The elements of the algebraic structure used to build the table, as a list. But most importantly, elements are present in the list in the order which they appear in the table's column headings.

EXAMPLES:

```
sage: from sage.matrix.operation_table import OperationTable
sage: G=AlternatingGroup(3)
sage: T=OperationTable(G, operator.mul)
sage: T.column_keys()
((), (1,2,3), (1,3,2))
```

set_print_symbols(ascii, latex)

Set the symbols used for text and LaTeX printing of operation tables.

INPUT:

- `ascii` - a single character for text table
- `latex` - a string to represent an operation in LaTeX math mode. Note the need for double-backslashes to escape properly.

EXAMPLE:

```
sage: from sage.matrix.operation_table import OperationTable
sage: G=AlternatingGroup(3)
sage: T=OperationTable(G, operator.mul)
sage: T.set_print_symbols('@', '\\times')
sage: T
@  a b c
+-----
a| a b c
b| b c a
```

```

c| c a b
sage: T._latex_()
'{{\setlength{\arraycolsep}{2ex}\n\begin{array}{r|*{3}{r}}\n\multicolumn{1}{c|}{\times}&

```

TESTS:

```

sage: from sage.matrix.operation_table import OperationTable
sage: G=AlternatingGroup(3)
sage: T=OperationTable(G, operator.mul)
sage: T.set_print_symbols('@', 5)
Traceback (most recent call last):
...
ValueError: LaTeX symbol must be a string, not 5
sage: T.set_print_symbols('@x@', '\\times')
Traceback (most recent call last):
...
ValueError: ASCII symbol should be a single character, not @x@
sage: T.set_print_symbols(5, '\\times')
Traceback (most recent call last):
...
ValueError: ASCII symbol should be a single character, not 5

```

table()

Returns the table as a list of lists, using integers to reference the elements.

OUTPUT: The rows of the table, as a list of rows, each row being a list of integer entries. The integers correspond to the order of the elements in the headings of the table and the order of the output of the `list()` method.

EXAMPLE:

```

sage: from sage.matrix.operation_table import OperationTable
sage: C=CyclicPermutationGroup(3)
sage: T=OperationTable(C, operator.mul)
sage: T.table()
[[0, 1, 2], [1, 2, 0], [2, 0, 1]]

```

translation()

Returns a dictionary associating names with elements.

OUTPUT: A dictionary whose keys are strings used as names for entries of the table and values that are the actual elements of the algebraic structure.

EXAMPLES:

```

sage: from sage.matrix.operation_table import OperationTable
sage: G=AlternatingGroup(3)
sage: T=OperationTable(G, operator.mul, names=['p','q','r'])
sage: sorted(T.translation().items())
[('p', ()), ('q', (1,2,3)), ('r', (1,3,2))]

```


BENCHMARKS FOR MATRICES

This file has many functions for computing timing benchmarks of various methods for random matrices with given bounds for the entries. The systems supported are Sage and Magma.

The basic command syntax is as follows:

```
sage: import sage.matrix.benchmark as b
sage: print "starting"; import sys; sys.stdout.flush(); b.report([b.det_ZZ], 'Test', systems=['sage'])
starting...
=====
Test
=====
...
=====
```

`sage.matrix.benchmark.MatrixVector_QQ` ($n=1000$, $h=100$, $system='sage'$, $times=1$)

Compute product of square n matrix by random vector with num and denom bounded by h the given number of times.

INPUT:

- n - matrix dimension (default: 300)
- h - numerator and denominator bound (default: bnd)
- $system$ - either 'sage' or 'magma' (default: 'sage')
- $times$ - number of experiments (default: 1)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.MatrixVector_QQ(500)
sage: tm = b.MatrixVector_QQ(500, system='magma') # optional - magma
```

`sage.matrix.benchmark.charpoly_GF` ($n=100$, $p=16411$, $system='sage'$)

Given a $n \times n$ matrix over GF with random entries, compute the charpoly.

INPUT:

- n - matrix dimension (default: 100)
- p - prime number (default: 16411)
- $system$ - either 'magma' or 'sage' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.charpoly_GF(100)
sage: tm = b.charpoly_GF(100, system='magma') # optional - magma
```

`sage.matrix.benchmark.charpoly_ZZ` ($n=100$, $min=0$, $max=9$, $system='sage'$)

Characteristic polynomial over ZZ: Given an $n \times n$ matrix over ZZ with random entries between min and max, compute the charpoly.

INPUT:

- n - matrix dimension (default: 100)
- min - minimal value for entries of matrix (default: 0)
- max - maximal value for entries of matrix (default: 9)
- $system$ - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.charpoly_ZZ(100)
sage: tm = b.charpoly_ZZ(100, system='magma') # optional - magma
```

`sage.matrix.benchmark.det_GF` ($n=400$, $p=16411$, $system='sage'$)

Dense determinant over GF(p). Given an $n \times n$ matrix A over GF with random entries compute $\det(A)$.

INPUT:

- n - matrix dimension (default: 300)
- p - prime number (default: 16411)
- $system$ - either 'magma' or 'sage' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.det_GF(1000)
sage: tm = b.det_GF(1000, system='magma') # optional - magma
```

`sage.matrix.benchmark.det_QQ` ($n=300$, $num_bound=10$, $den_bound=10$, $system='sage'$)

Dense rational determinant over QQ. Given an $n \times n$ matrix A over QQ with random entries with numerator bound and denominator bound, compute $\det(A)$.

INPUT:

- n - matrix dimension (default: 200)
- num_bound - numerator bound, inclusive (default: 10)
- den_bound - denominator bound, inclusive (default: 10)
- $system$ - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.det_QQ(200)
sage: ts = b.det_QQ(10, num_bound=100000, den_bound=10000)
sage: tm = b.det_QQ(200, system='magma') # optional - magma
```

`sage.matrix.benchmark.det_ZZ (n=200, min=1, max=100, system='sage')`

Dense integer determinant over ZZ. Given an $n \times n$ matrix A over ZZ with random entries between min and max, inclusive, compute $\det(A)$.

INPUT:

- n - matrix dimension (default: 200)
- min - minimal value for entries of matrix (default: 1)
- max - maximal value for entries of matrix (default: 100)
- system - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.det_ZZ(200)
sage: tm = b.det_ZZ(200, system='magma') # optional - magma
```

`sage.matrix.benchmark.det_hilbert_QQ (n=80, system='sage')`

Runs the benchmark for calculating the determinant of the hilbert matrix over rationals of dimension n.

INPUT:

- n - matrix dimension (default: 300)
- system - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.det_hilbert_QQ(50)
sage: tm = b.det_hilbert_QQ(50, system='magma') # optional - magma
```

`sage.matrix.benchmark.echelon_QQ (n=100, min=0, max=9, system='sage')`

Given an $n \times (2*n)$ matrix over QQ with random integer entries between min and max, compute the reduced row echelon form.

INPUT:

- n - matrix dimension (default: 300)
- min - minimal value for entries of matrix (default: -9)
- max - maximal value for entries of matrix (default: 9)
- system - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.echelon_QQ(100)
sage: tm = b.echelon_QQ(100, system='magma') # optional - magma
```

`sage.matrix.benchmark.hilbert_matrix (n)`

Returns the Hilbert matrix of size n over rationals.

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: b.hilbert_matrix(3)
[ 1 1/2 1/3]
[1/2 1/3 1/4]
[1/3 1/4 1/5]
```

`sage.matrix.benchmark.inverse_QQ (n=100, min=0, max=9, system='sage')`

Given a $n \times n$ matrix over $\mathbb{Q}\mathbb{Q}$ with random integer entries between \min and \max , compute the reduced row echelon form.

INPUT:

- n - matrix dimension (default: 300)
- \min - minimal value for entries of matrix (default: -9)
- \max - maximal value for entries of matrix (default: 9)
- system - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.inverse_QQ(100)
sage: tm = b.inverse_QQ(100, system='magma') # optional - magma
```

`sage.matrix.benchmark.invert_hilbert_QQ (n=40, system='sage')`

Runs the benchmark for calculating the inverse of the hilbert matrix over rationals of dimension n .

INPUT:

- n - matrix dimension (default: 300)
- system - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.invert_hilbert_QQ(30)
sage: tm = b.invert_hilbert_QQ(30, system='magma') # optional - magma
```

`sage.matrix.benchmark.matrix_add_GF (n=1000, p=16411, system='sage', times=100)`

Given two $n \times n$ matrix over $\text{GF}(p)$ with random entries, add them.

INPUT:

- n - matrix dimension (default: 300)
- p - prime number (default: 16411)
- system - either 'magma' or 'sage' (default: 'sage')
- times - number of experiments (default: 100)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.matrix_add_GF(500, p=19)
sage: tm = b.matrix_add_GF(500, p=19, system='magma') # optional - magma
```

`sage.matrix.benchmark.matrix_add_ZZ (n=200, min=-9, max=9, system='sage', times=50)`

Matrix addition over $\mathbb{Z}\mathbb{Z}$ Given an $n \times n$ matrix A and B over $\mathbb{Z}\mathbb{Z}$ with random entries between \min and \max , inclusive, compute $A + B$ times times.

INPUT:

- n - matrix dimension (default: 200)
- \min - minimal value for entries of matrix (default: -9)
- \max - maximal value for entries of matrix (default: 9)
- system - either 'sage' or 'magma' (default: 'sage')

- times - number of experiments (default: 50)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.matrix_add_ZZ(200)
sage: tm = b.matrix_add_ZZ(200, system='magma') # optional - magma
```

sage.matrix.benchmark.**matrix_add_ZZ_2**(n=200, bits=16, system='sage', times=50)

Matrix addition over ZZ. Given an n x n matrix A and B over ZZ with random bits-bit entries, compute A + B.

INPUT:

- n - matrix dimension (default: 200)
- bits - bitsize of entries
- system - either 'sage' or 'magma' (default: 'sage')
- times - number of experiments (default: 50)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.matrix_add_ZZ_2(200)
sage: tm = b.matrix_add_ZZ_2(200, system='magma') # optional - magma
```

sage.matrix.benchmark.**matrix_multiply_GF**(n=100, p=16411, system='sage', times=3)

Given an n x n matrix A over GF(p) with random entries, compute A * (A+1).

INPUT:

- n - matrix dimension (default: 100)
- p - prime number (default: 16411)
- system - either 'magma' or 'sage' (default: 'sage')
- times - number of experiments (default: 3)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.matrix_multiply_GF(100, p=19)
sage: tm = b.matrix_multiply_GF(100, p=19, system='magma') # optional - magma
```

sage.matrix.benchmark.**matrix_multiply_QQ**(n=100, bnd=2, system='sage', times=1)

Given an n x n matrix A over QQ with random entries whose numerators and denominators are bounded by bnd, compute A * (A+1).

INPUT:

- n - matrix dimension (default: 300)
- bnd - numerator and denominator bound (default: bnd)
- system - either 'sage' or 'magma' (default: 'sage')
- times - number of experiments (default: 1)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.matrix_multiply_QQ(100)
sage: tm = b.matrix_multiply_QQ(100, system='magma') # optional - magma
```

```
sage.matrix.benchmark.matrix_multiply_ZZ (n=300, min=-9, max=9, system='sage',
                                           times=1)
```

Matrix multiplication over ZZ: Given an $n \times n$ matrix A over ZZ with random entries between min and max, inclusive, compute $A * (A+1)$.

INPUT:

- `n` - matrix dimension (default: 300)
- `min` - minimal value for entries of matrix (default: -9)
- `max` - maximal value for entries of matrix (default: 9)
- `system` - either 'sage' or 'magma' (default: 'sage')
- `times` - number of experiments (default: 1)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.matrix_multiply_ZZ(200)
sage: tm = b.matrix_multiply_ZZ(200, system='magma') # optional - magma
```

```
sage.matrix.benchmark.nullspace_GF (n=300, p=16411, system='sage')
```

Given a $n+1 \times n$ matrix over GF(p) with random entries, compute the nullspace.

INPUT:

- `n` - matrix dimension (default: 300)
- `p` - prime number (default: 16411)
- `system` - either 'magma' or 'sage' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.nullspace_GF(300)
sage: tm = b.nullspace_GF(300, system='magma') # optional - magma
```

```
sage.matrix.benchmark.nullspace_RDF (n=300, min=0, max=10, system='sage')
```

Nullspace over RDF: Given a $n+1 \times n$ matrix over RDF with random entries between min and max, compute the nullspace.

INPUT:

- `n` - matrix dimension (default: 300)
- `min` - minimal value for entries of matrix (default: 0)
- `max` - maximal value for entries of matrix (default: 10)
- `system` - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.nullspace_RDF(100) # long time
sage: tm = b.nullspace_RDF(100, system='magma') # optional - magma
```

```
sage.matrix.benchmark.nullspace_RR (n=300, min=0, max=10, system='sage')
```

Nullspace over RR: Given a $n+1 \times n$ matrix over RR with random entries between min and max, compute the nullspace.

INPUT:

- `n` - matrix dimension (default: 300)
- `min` - minimal value for entries of matrix (default: 0)
- `max` - maximal value for entries of matrix (default: 10)
- `system` - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.nullspace_RR(100)
sage: tm = b.nullspace_RR(100, system='magma') # optional - magma
```

`sage.matrix.benchmark.nullspace_ZZ` ($n=200$, $min=0$, $max=4294967296$, $system='sage'$)

Nullspace over ZZ: Given a $n+1 \times n$ matrix over ZZ with random entries between min and max, compute the nullspace.

INPUT:

- `n` - matrix dimension (default: 200)
- `min` - minimal value for entries of matrix (default: 0)
- `max` - maximal value for entries of matrix (default: 2^{32})
- `system` - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.nullspace_ZZ(200)
sage: tm = b.nullspace_ZZ(200, system='magma') # optional - magma
```

`sage.matrix.benchmark.rank2_GF` ($n=500$, $p=16411$, $system='sage'$)

Rank over GF(p): Given a $(n+10) \times n$ matrix over GF(p) with random entries, compute the rank.

INPUT:

- `n` - matrix dimension (default: 300)
- `p` - prime number (default: 16411)
- `system` - either 'magma' or 'sage' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.rank2_GF(500)
sage: tm = b.rank2_GF(500, system='magma') # optional - magma
```

`sage.matrix.benchmark.rank2_ZZ` ($n=400$, $min=0$, $max=18446744073709551616L$, $system='sage'$)

Rank 2 over ZZ: Given a $(n+10) \times n$ matrix over ZZ with random entries between min and max, compute the rank.

INPUT:

- `n` - matrix dimension (default: 400)
- `min` - minimal value for entries of matrix (default: 0)
- `max` - maximal value for entries of matrix (default: 2^{64})
- `system` - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.rank2_ZZ(300)
sage: tm = b.rank2_ZZ(300, system='magma') # optional - magma
```

`sage.matrix.benchmark.rank_GF` ($n=500$, $p=16411$, $system='sage'$)

Rank over GF(p): Given a $n \times (n+10)$ matrix over GF(p) with random entries, compute the rank.

INPUT:

- n - matrix dimension (default: 300)
- p - prime number (default: 16411)
- $system$ - either 'magma' or 'sage' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.rank_GF(1000)
sage: tm = b.rank_GF(1000, system='magma') # optional - magma
```

`sage.matrix.benchmark.rank_ZZ` ($n=700$, $min=0$, $max=9$, $system='sage'$)

Rank over ZZ: Given a $n \times (n+10)$ matrix over ZZ with random entries between min and max, compute the rank.

INPUT:

- n - matrix dimension (default: 700)
- min - minimal value for entries of matrix (default: 0)
- max - maximal value for entries of matrix (default: 9)
- $system$ - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.rank_ZZ(300)
sage: tm = b.rank_ZZ(300, system='magma') # optional - magma
```

`sage.matrix.benchmark.report` (F , $title$, $systems=['sage', 'magma']$, $**kwds$)

Run benchmarks with default arguments for each function in the list F .

INPUT:

- F - a list of callables used for benchmarking
- $title$ - a string describing this report
- $systems$ - a list of systems (supported entries are 'sage' and 'magma')
- $**kwds$ - keyword arguments passed to all functions in F

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: print "starting"; import sys; sys.stdout.flush(); b.report([b.det_ZZ], 'Test', systems=['sage', 'magma'])
starting...
=====
Test
=====
...
=====
```



```
sage.matrix.benchmark.report_GF(p=16411, **kws)
```

Runs all the reports for finite field matrix operations, for prime $p=16411$.

INPUT:

- `p` - ignored
- `**kws` - passed through to `report()`

Note: right now, even though `p` is an input, it is being ignored! If you need to check the performance for other primes, you can call individual benchmark functions.

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: print "starting"; import sys; sys.stdout.flush(); b.report_GF(systems=['sage'])
starting...
=====
Dense benchmarks over GF with prime 16411
=====
...
=====
```

```
sage.matrix.benchmark.report_ZZ(**kws)
```

Reports all the benchmarks for integer matrices and few rational matrices.

INPUT:

- `**kws` - passed through to `report()`

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: print "starting"; import sys; sys.stdout.flush(); b.report_ZZ(systems=['sage']) # long ti
starting...
=====
Dense benchmarks over ZZ
=====
...
=====
```

```
sage.matrix.benchmark.smithform_ZZ(n=128, min=0, max=9, system='sage')
```

Smith Form over ZZ: Given a $n \times n$ matrix over ZZ with random entries between `min` and `max`, compute the Smith normal form.

INPUT:

- `n` - matrix dimension (default: 128)
- `min` - minimal value for entries of matrix (default: 0)
- `max` - maximal value for entries of matrix (default: 9)
- `system` - either 'sage' or 'magma' (default: 'sage')

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.smithform_ZZ(100)
sage: tm = b.smithform_ZZ(100, system='magma') # optional - magma
```

```
sage.matrix.benchmark.vecmat_ZZ(n=300, min=-9, max=9, system='sage', times=200)
```

Vector matrix multiplication over ZZ.

Given an $n \times n$ matrix A over $\mathbb{Z}\mathbb{Z}$ with random entries between \min and \max , inclusive, and v the first row of A , compute the product $v * A$.

INPUT:

- `n` - matrix dimension (default: 300)
- `min` - minimal value for entries of matrix (default: -9)
- `max` - maximal value for entries of matrix (default: 9)
- `system` - either 'sage' or 'magma' (default: 'sage')
- `times` - number of runs (default: 200)

EXAMPLES:

```
sage: import sage.matrix.benchmark as b
sage: ts = b.vecmat_ZZ(300) # long time
sage: tm = b.vecmat_ZZ(300, system='magma') # optional - magma
```

INDICES AND TABLES

- [Index](#)
- [Module Index](#)
- [Search Page](#)

BIBLIOGRAPHY

- [BEEZER] *A First Course in Linear Algebra*. Robert A. Beezer, accessed 15 July 2010.
- [TREFETHEN-BAU] Trefethen, Lloyd N., Bau, David, III “Numerical Linear Algebra” SIAM, Philadelphia, 1997.
- [Kn95] Donald E. Knuth, *Overlapping Pfaffians*, *Arxiv math/9503234v1*.
- [Rote2001] Gunter Rote, *Division-Free Algorithms for the Determinant and the Pfaffian: Algebraic and Combinatorial Approaches*, H. Alt (Ed.): Computational Discrete Mathematics, LNCS 2122, pp. 119–135, 2001. <http://page.mi.fu-berlin.de/rote/Papers/pdf/Division-free+algorithms.pdf>
- [DW95] Andreas W.M. Dress, Walter Wenzel, *A Simple Proof of an Identity Concerning Pfaffians of Skew Symmetric Matrices*, *Advances in Mathematics*, volume 112, Issue 1, April 1995, pp. 120-134. <http://www.sciencedirect.com/science/article/pii/S0001870885710298>
- [STORJOHANN-EMAIL] 1. Storjohann, Email Communication. 30 May 2011.
- [H] F. Hess, “Computing Riemann-Roch spaces in algebraic function fields and related topics,” *J. Symbolic Comput.* 33 (2002), no. 4, 425–445.
- [MS] T. Mulders, A. Storjohann, “On lattice reduction for polynomial matrices,” *J. Symbolic Comput.* 35 (2003), no. 4, 377–401
- [V] G. Villard, “Computing Popov and Hermite forms of polynomial matrices”, ISSAC ‘96: Proceedings of the 1996 international symposium on Symbolic and algebraic computation, 250–258.
- [STORJOHANN-THESIS] A. Storjohann, Algorithms for Matrix Canonical Forms. PhD Thesis. Department of Computer Science, Swiss Federal Institute of Technology – ETH, 2000.
- [STORJOHANN-ISACC98] A. Storjohann, An $O(n^3)$ algorithm for Frobenius normal form. Proceedings of the International Symposium on Symbolic and Algebraic Computation (ISSAC’98), ACM Press, 1998, pp. 101-104.
- [SH95] C. P. Schnorr and H. H. Hörner. *Attacking the Chor-Rivest Cryptosystem by Improved Lattice Reduction*. *Advances in Cryptology - EUROCRYPT ‘95*. LNCS Volume 921, 1995, pp 1-12.
- [GL96] G. Golub and C. van Loan. *Matrix Computations*. 3rd edition, Johns Hopkins Univ. Press, 1996.
- [WATKINS] Watkins, David S. *Fundamentals of Matrix Computations*, Third Edition. Wiley, Hoboken, New Jersey, 2010.

PYTHON MODULE INDEX

m

- `sage.matrix.benchmark`, 413
- `sage.matrix.berlekamp_massey`, 279
- `sage.matrix.constructor`, 15
- `sage.matrix.docs`, 59
- `sage.matrix.matrix`, 67
- `sage.matrix.matrix0`, 69
- `sage.matrix.matrix1`, 101
- `sage.matrix.matrix2`, 123
- `sage.matrix.matrix_complex_double_dense`, 397
- `sage.matrix.matrix_dense`, 281
- `sage.matrix.matrix_double_dense`, 353
- `sage.matrix.matrix_generic_dense`, 293
- `sage.matrix.matrix_generic_sparse`, 295
- `sage.matrix.matrix_integer_dense`, 317
- `sage.matrix.matrix_modn_dense`, 297
- `sage.matrix.matrix_modn_sparse`, 303
- `sage.matrix.matrix_mpolynomial_dense`, 399
- `sage.matrix.matrix_rational_dense`, 343
- `sage.matrix.matrix_real_double_dense`, 395
- `sage.matrix.matrix_space`, 3
- `sage.matrix.matrix_sparse`, 285
- `sage.matrix.matrix_symbolic_dense`, 307
- `sage.matrix.operation_table`, 403
- `sage.matrix.strassen`, 275

INDEX

A

`act_on_polynomial()` (sage.matrix.matrix0.Matrix method), 70
`add_multiple_of_column()` (sage.matrix.matrix0.Matrix method), 70
`add_multiple_of_row()` (sage.matrix.matrix0.Matrix method), 70
`adjoint()` (sage.matrix.matrix2.Matrix method), 134
`anticommutator()` (sage.matrix.matrix0.Matrix method), 71
`antitranspose()` (sage.matrix.matrix_dense.Matrix_dense method), 281
`antitranspose()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 321
`antitranspose()` (sage.matrix.matrix_rational_dense.Matrix_rational_dense method), 344
`antitranspose()` (sage.matrix.matrix_sparse.Matrix_sparse method), 285
`apply_map()` (sage.matrix.matrix_dense.Matrix_dense method), 282
`apply_map()` (sage.matrix.matrix_sparse.Matrix_sparse method), 285
`apply_morphism()` (sage.matrix.matrix_dense.Matrix_dense method), 283
`apply_morphism()` (sage.matrix.matrix_sparse.Matrix_sparse method), 287
`arguments()` (sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense method), 309
`as_bipartite_graph()` (sage.matrix.matrix2.Matrix method), 135
`as_sum_of_permutations()` (sage.matrix.matrix2.Matrix method), 136
`augment()` (sage.matrix.matrix1.Matrix method), 101
`augment()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 322
`augment()` (sage.matrix.matrix_sparse.Matrix_sparse method), 287
`automorphisms_of_rows_and_columns()` (sage.matrix.matrix2.Matrix method), 136

B

`base_extend()` (sage.matrix.matrix_space.MatrixSpace method), 4
`base_ring()` (sage.matrix.matrix0.Matrix method), 71
`basis()` (sage.matrix.matrix_space.MatrixSpace method), 4
`berlekamp_massey()` (in module sage.matrix.berlekamp_massey), 279
`BKZ()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 318
`block_diagonal_matrix()` (in module sage.matrix.constructor), 22
`block_matrix()` (in module sage.matrix.constructor), 22
`block_sum()` (sage.matrix.matrix1.Matrix method), 104

C

`C` (sage.matrix.matrix2.Matrix attribute), 123
`cached_method()` (sage.matrix.matrix_space.MatrixSpace method), 4
`change_names()` (sage.matrix.operation_table.OperationTable method), 408

`change_ring()` (`sage.matrix.matrix0.Matrix` method), 71
`change_ring()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 344
`change_ring()` (`sage.matrix.matrix_space.MatrixSpace` method), 5
`change_ring()` (`sage.matrix.matrix_sparse.Matrix_sparse` method), 288
`characteristic_polynomial()` (`sage.matrix.matrix2.Matrix` method), 137
`charpoly()` (`sage.matrix.matrix2.Matrix` method), 137
`charpoly()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 323
`charpoly()` (`sage.matrix.matrix_modn_dense.Matrix_modn_dense` method), 298
`charpoly()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 345
`charpoly()` (`sage.matrix.matrix_sparse.Matrix_sparse` method), 288
`charpoly()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 310
`charpoly_GF()` (in module `sage.matrix.benchmark`), 413
`charpoly_ZZ()` (in module `sage.matrix.benchmark`), 414
`cholesky()` (`sage.matrix.matrix2.Matrix` method), 139
`cholesky()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 360
`cmp_pivots()` (in module `sage.matrix.matrix2`), 272
`column()` (`sage.matrix.matrix1.Matrix` method), 104
`column()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 345
`column_keys()` (`sage.matrix.operation_table.OperationTable` method), 409
`column_matrix()` (in module `sage.matrix.constructor`), 25
`column_module()` (`sage.matrix.matrix2.Matrix` method), 143
`column_space()` (`sage.matrix.matrix2.Matrix` method), 143
`column_space()` (`sage.matrix.matrix_space.MatrixSpace` method), 5
`columns()` (`sage.matrix.matrix1.Matrix` method), 105
`commutator()` (`sage.matrix.matrix0.Matrix` method), 72
`companion_matrix()` (in module `sage.matrix.constructor`), 26
`condition()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 362
`conjugate()` (`sage.matrix.matrix2.Matrix` method), 144
`conjugate_transpose()` (`sage.matrix.matrix2.Matrix` method), 144
`construction()` (`sage.matrix.matrix_space.MatrixSpace` method), 5
`cyclic_subspace()` (`sage.matrix.matrix2.Matrix` method), 146

D

`decomp_seq()` (in module `sage.matrix.matrix2`), 272
`decomposition()` (`sage.matrix.matrix2.Matrix` method), 148
`decomposition()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 323
`decomposition()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 346
`decomposition_of_subspace()` (`sage.matrix.matrix2.Matrix` method), 149
`delete_columns()` (`sage.matrix.matrix1.Matrix` method), 105
`delete_rows()` (`sage.matrix.matrix1.Matrix` method), 106
`denominator()` (`sage.matrix.matrix2.Matrix` method), 150
`denominator()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 346
`dense_columns()` (`sage.matrix.matrix1.Matrix` method), 107
`dense_matrix()` (`sage.matrix.matrix1.Matrix` method), 108
`dense_rows()` (`sage.matrix.matrix1.Matrix` method), 108
`density()` (`sage.matrix.matrix2.Matrix` method), 151
`density()` (`sage.matrix.matrix_modn_sparse.Matrix_modn_sparse` method), 304
`density()` (`sage.matrix.matrix_sparse.Matrix_sparse` method), 288
`derivative()` (`sage.matrix.matrix2.Matrix` method), 151
`det()` (`sage.matrix.matrix2.Matrix` method), 152

`det_GF()` (in module `sage.matrix.benchmark`), 414
`det_hilbert_QQ()` (in module `sage.matrix.benchmark`), 415
`det_QQ()` (in module `sage.matrix.benchmark`), 414
`det_ZZ()` (in module `sage.matrix.benchmark`), 414
`determinant()` (`sage.matrix.matrix2.Matrix` method), 152
`determinant()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 364
`determinant()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 324
`determinant()` (`sage.matrix.matrix_modn_dense.Matrix_modn_dense` method), 299
`determinant()` (`sage.matrix.matrix_mpolynomial_dense.Matrix_mpolynomial_dense` method), 399
`determinant()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 347
`determinant()` (`sage.matrix.matrix_sparse.Matrix_sparse` method), 289
`diagonal()` (`sage.matrix.matrix2.Matrix` method), 153
`diagonal_matrix()` (in module `sage.matrix.constructor`), 28
`dict()` (`sage.matrix.matrix0.Matrix` method), 72
`dict_to_list()` (in module `sage.matrix.matrix_space`), 12
`dimension()` (`sage.matrix.matrix_space.MatrixSpace` method), 6
`dimensions()` (`sage.matrix.matrix0.Matrix` method), 72
`dims()` (`sage.matrix.matrix_space.MatrixSpace` method), 6

E

`echelon_form()` (`sage.matrix.matrix2.Matrix` method), 154
`echelon_form()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 325
`echelon_form()` (`sage.matrix.matrix_mpolynomial_dense.Matrix_mpolynomial_dense` method), 400
`echelon_form()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 347
`echelon_QQ()` (in module `sage.matrix.benchmark`), 415
`echelonize()` (`sage.matrix.matrix2.Matrix` method), 155
`echelonize()` (`sage.matrix.matrix_modn_dense.Matrix_modn_dense` method), 299
`echelonize()` (`sage.matrix.matrix_mpolynomial_dense.Matrix_mpolynomial_dense` method), 401
`echelonize()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 348
`eigenmatrix_left()` (`sage.matrix.matrix2.Matrix` method), 157
`eigenmatrix_right()` (`sage.matrix.matrix2.Matrix` method), 158
`eigenspaces_left()` (`sage.matrix.matrix2.Matrix` method), 159
`eigenspaces_right()` (`sage.matrix.matrix2.Matrix` method), 163
`eigenvalues()` (`sage.matrix.matrix2.Matrix` method), 166
`eigenvalues()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 365
`eigenvalues()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 310
`eigenvectors_left()` (`sage.matrix.matrix2.Matrix` method), 167
`eigenvectors_left()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 367
`eigenvectors_left()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 310
`eigenvectors_right()` (`sage.matrix.matrix2.Matrix` method), 168
`eigenvectors_right()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 368
`eigenvectors_right()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 311
`elementary_divisors()` (`sage.matrix.matrix2.Matrix` method), 168
`elementary_divisors()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 328
`elementary_matrix()` (in module `sage.matrix.constructor`), 31
`elementwise_product()` (`sage.matrix.matrix2.Matrix` method), 168
`exp()` (`sage.matrix.matrix2.Matrix` method), 171
`exp()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 368
`exp()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 312
`expand()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 313

`extended_echelon_form()` (`sage.matrix.matrix2.Matrix` method), 171

F

`factor()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 313

`fcf()` (`sage.matrix.matrix2.Matrix` method), 174

`fcf()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 313

`find()` (`sage.matrix.matrix2.Matrix` method), 174

`frobenius()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 329

`full_category_initialisation()` (`sage.matrix.matrix_space.MatrixSpace` method), 6

G

`gcd()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 330

`gen()` (`sage.matrix.matrix_space.MatrixSpace` method), 6

`get_action_impl()` (`sage.matrix.matrix_space.MatrixSpace` method), 7

`get_subdivisions()` (`sage.matrix.matrix2.Matrix` method), 175

`gram_schmidt()` (`sage.matrix.matrix2.Matrix` method), 175

H

`H` (`sage.matrix.matrix2.Matrix` attribute), 124

`hadamard_bound()` (`sage.matrix.matrix2.Matrix` method), 180

`height()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 330

`height()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 349

`hermite_form()` (`sage.matrix.matrix2.Matrix` method), 181

`hermite_form()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 330

`hessenberg_form()` (`sage.matrix.matrix2.Matrix` method), 181

`hessenbergize()` (`sage.matrix.matrix2.Matrix` method), 182

`hessenbergize()` (`sage.matrix.matrix_modn_dense.Matrix_modn_dense` method), 299

`hilbert_matrix()` (in module `sage.matrix.benchmark`), 415

I

`I` (`sage.matrix.matrix2.Matrix` attribute), 124

`identity_matrix()` (in module `sage.matrix.constructor`), 36

`identity_matrix()` (`sage.matrix.matrix_space.MatrixSpace` method), 7

`image()` (`sage.matrix.matrix2.Matrix` method), 182

`indefinite_factorization()` (`sage.matrix.matrix2.Matrix` method), 183

`index_in_saturation()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 333

`insert_row()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 334

`int_range` (class in `sage.matrix.strassen`), 275

`integer_kernel()` (`sage.matrix.matrix2.Matrix` method), 185

`intervals()` (`sage.matrix.strassen.int_range` method), 276

`inverse()` (`sage.matrix.matrix2.Matrix` method), 185

`inverse_QQ()` (in module `sage.matrix.benchmark`), 415

`invert_hilbert_QQ()` (in module `sage.matrix.benchmark`), 416

`is_alternating()` (`sage.matrix.matrix0.Matrix` method), 73

`is_bistochastic()` (`sage.matrix.matrix2.Matrix` method), 186

`is_dense()` (`sage.matrix.matrix0.Matrix` method), 73

`is_dense()` (`sage.matrix.matrix_space.MatrixSpace` method), 7

`is_diagonalizable()` (`sage.matrix.matrix2.Matrix` method), 186

`is_finite()` (`sage.matrix.matrix_space.MatrixSpace` method), 7

`is_hermitian()` (`sage.matrix.matrix0.Matrix` method), 73

[is_hermitian\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 369
[is_immutable\(\)](#) (sage.matrix.matrix0.Matrix method), 74
[is_invertible\(\)](#) (sage.matrix.matrix0.Matrix method), 74
[is_LLL_reduced\(\)](#) (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 334
[is_Matrix\(\)](#) (in module sage.matrix.matrix), 67
[is_Matrix_modn_dense\(\)](#) (in module sage.matrix.matrix_modn_dense), 301
[is_MatrixSpace\(\)](#) (in module sage.matrix.matrix_space), 13
[is_mutable\(\)](#) (sage.matrix.matrix0.Matrix method), 75
[is_normal\(\)](#) (sage.matrix.matrix2.Matrix method), 189
[is_normal\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 371
[is_one\(\)](#) (sage.matrix.matrix2.Matrix method), 190
[is_permutation_of\(\)](#) (sage.matrix.matrix2.Matrix method), 190
[is_positive_definite\(\)](#) (sage.matrix.matrix2.Matrix method), 191
[is_positive_definite\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 373
[is_scalar\(\)](#) (sage.matrix.matrix2.Matrix method), 194
[is_similar\(\)](#) (sage.matrix.matrix2.Matrix method), 194
[is_singular\(\)](#) (sage.matrix.matrix0.Matrix method), 75
[is_skew_symmetric\(\)](#) (sage.matrix.matrix0.Matrix method), 76
[is_skew_symmetrizable\(\)](#) (sage.matrix.matrix0.Matrix method), 77
[is_sparse\(\)](#) (sage.matrix.matrix0.Matrix method), 77
[is_sparse\(\)](#) (sage.matrix.matrix_space.MatrixSpace method), 7
[is_square\(\)](#) (sage.matrix.matrix0.Matrix method), 78
[is_symmetric\(\)](#) (sage.matrix.matrix0.Matrix method), 78
[is_symmetric\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 375
[is_symmetrizable\(\)](#) (sage.matrix.matrix0.Matrix method), 78
[is_unitary\(\)](#) (sage.matrix.matrix2.Matrix method), 197
[is_unitary\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 376
[iterates\(\)](#) (sage.matrix.matrix0.Matrix method), 79
[ith_to_zero_rotation_matrix\(\)](#) (in module sage.matrix.constructor), 36

J

[jordan_block\(\)](#) (in module sage.matrix.constructor), 38
[jordan_form\(\)](#) (sage.matrix.matrix2.Matrix method), 198

K

[kernel\(\)](#) (sage.matrix.matrix2.Matrix method), 204
[kernel_on\(\)](#) (sage.matrix.matrix2.Matrix method), 206

L

[left_eigenmatrix\(\)](#) (sage.matrix.matrix2.Matrix method), 207
[left_eigenspaces\(\)](#) (sage.matrix.matrix2.Matrix method), 208
[left_eigenvectors\(\)](#) (sage.matrix.matrix2.Matrix method), 212
[left_eigenvectors\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 378
[left_kernel\(\)](#) (sage.matrix.matrix2.Matrix method), 213
[left_nullity\(\)](#) (sage.matrix.matrix2.Matrix method), 215
[lift\(\)](#) (sage.matrix.matrix1.Matrix method), 109
[lift\(\)](#) (sage.matrix.matrix_modn_dense.Matrix_modn_dense method), 300
[lift\(\)](#) (sage.matrix.matrix_modn_sparse.Matrix_modn_sparse method), 304
[linear_combination_of_columns\(\)](#) (sage.matrix.matrix0.Matrix method), 79
[linear_combination_of_rows\(\)](#) (sage.matrix.matrix0.Matrix method), 80

`list()` (sage.matrix.matrix0.Matrix method), 81
`list_to_dict()` (in module sage.matrix.matrix_space), 13
`LLL()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 319
`LLL_gram()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 321
`log_determinant()` (sage.matrix.matrix_double_dense.Matrix_double_dense method), 378
`LU()` (sage.matrix.matrix2.Matrix method), 124
`LU()` (sage.matrix.matrix_double_dense.Matrix_double_dense method), 354
`LU_valid()` (sage.matrix.matrix_double_dense.Matrix_double_dense method), 356

M

`Matrix` (class in sage.matrix.matrix), 67
`Matrix` (class in sage.matrix.matrix0), 69
`Matrix` (class in sage.matrix.matrix1), 101
`Matrix` (class in sage.matrix.matrix2), 123
`matrix()` (sage.matrix.matrix_space.MatrixSpace method), 8
`matrix_add_GF()` (in module sage.matrix.benchmark), 416
`matrix_add_ZZ()` (in module sage.matrix.benchmark), 416
`matrix_add_ZZ_2()` (in module sage.matrix.benchmark), 417
`Matrix_complex_double_dense` (class in sage.matrix.matrix_complex_double_dense), 397
`Matrix_dense` (class in sage.matrix.matrix_dense), 281
`Matrix_double_dense` (class in sage.matrix.matrix_double_dense), 353
`matrix_from_columns()` (sage.matrix.matrix1.Matrix method), 109
`matrix_from_columns()` (sage.matrix.matrix_modn_sparse.Matrix_modn_sparse method), 305
`matrix_from_rows()` (sage.matrix.matrix1.Matrix method), 110
`matrix_from_rows()` (sage.matrix.matrix_modn_sparse.Matrix_modn_sparse method), 305
`matrix_from_rows_and_columns()` (sage.matrix.matrix1.Matrix method), 110
`matrix_from_rows_and_columns()` (sage.matrix.matrix_sparse.Matrix_sparse method), 289
`Matrix_generic_dense` (class in sage.matrix.matrix_generic_dense), 293
`Matrix_generic_sparse` (class in sage.matrix.matrix_generic_sparse), 296
`Matrix_integer_dense` (class in sage.matrix.matrix_integer_dense), 317
`matrix_method()` (in module sage.matrix.constructor), 38
`Matrix_modn_dense` (class in sage.matrix.matrix_modn_dense), 298
`Matrix_modn_sparse` (class in sage.matrix.matrix_modn_sparse), 304
`Matrix_mpolynomial_dense` (class in sage.matrix.matrix_mpolynomial_dense), 399
`matrix_multiply_GF()` (in module sage.matrix.benchmark), 417
`matrix_multiply_QQ()` (in module sage.matrix.benchmark), 417
`matrix_multiply_ZZ()` (in module sage.matrix.benchmark), 417
`matrix_of_variables()` (sage.matrix.operation_table.OperationTable method), 409
`matrix_over_field()` (sage.matrix.matrix1.Matrix method), 110
`Matrix_rational_dense` (class in sage.matrix.matrix_rational_dense), 343
`Matrix_real_double_dense` (class in sage.matrix.matrix_real_double_dense), 395
`matrix_space()` (sage.matrix.matrix1.Matrix method), 111
`matrix_space()` (sage.matrix.matrix_space.MatrixSpace method), 9
`Matrix_sparse` (class in sage.matrix.matrix_sparse), 285
`Matrix_sparse_from_rows()` (in module sage.matrix.matrix_generic_sparse), 296
`Matrix_symbolic_dense` (class in sage.matrix.matrix_symbolic_dense), 309
`matrix_window()` (sage.matrix.matrix2.Matrix method), 216
`matrix_window()` (sage.matrix.matrix_modn_dense.Matrix_modn_dense method), 300
`MatrixFactory` (class in sage.matrix.constructor), 15
`MatrixSpace` (class in sage.matrix.matrix_space), 3

MatrixVector_QQ() (in module sage.matrix.benchmark), 413
 MatrixWindow (class in sage.matrix.matrix_rational_dense), 343
 maxspin() (sage.matrix.matrix2.Matrix method), 216
 minimal_polynomial() (sage.matrix.matrix2.Matrix method), 217
 minors() (sage.matrix.matrix2.Matrix method), 217
 minpoly() (sage.matrix.matrix2.Matrix method), 217
 minpoly() (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 334
 minpoly() (sage.matrix.matrix_modn_dense.Matrix_modn_dense method), 300
 minpoly() (sage.matrix.matrix_rational_dense.Matrix_rational_dense method), 349
 mod() (sage.matrix.matrix0.Matrix method), 82
 multiplicative_order() (sage.matrix.matrix0.Matrix method), 82
 mutate() (sage.matrix.matrix0.Matrix method), 83

N

N() (sage.matrix.matrix2.Matrix method), 130
 n() (sage.matrix.matrix2.Matrix method), 218
 ncols() (sage.matrix.matrix0.Matrix method), 84
 ncols() (sage.matrix.matrix_space.MatrixSpace method), 10
 ncols_from_dict() (in module sage.matrix.constructor), 39
 new_matrix() (sage.matrix.matrix1.Matrix method), 111
 ngens() (sage.matrix.matrix_space.MatrixSpace method), 10
 nonpivots() (sage.matrix.matrix0.Matrix method), 84
 nonzero_positions() (sage.matrix.matrix0.Matrix method), 84
 nonzero_positions_in_column() (sage.matrix.matrix0.Matrix method), 85
 nonzero_positions_in_row() (sage.matrix.matrix0.Matrix method), 85
 norm() (sage.matrix.matrix2.Matrix method), 219
 norm() (sage.matrix.matrix_double_dense.Matrix_double_dense method), 379
 nrows() (sage.matrix.matrix0.Matrix method), 85
 nrows() (sage.matrix.matrix_space.MatrixSpace method), 10
 nrows_from_dict() (in module sage.matrix.constructor), 39
 nullity() (sage.matrix.matrix2.Matrix method), 220
 nullspace_GF() (in module sage.matrix.benchmark), 418
 nullspace_RDF() (in module sage.matrix.benchmark), 418
 nullspace_RR() (in module sage.matrix.benchmark), 418
 nullspace_ZZ() (in module sage.matrix.benchmark), 419
 number_of_arguments() (sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense method), 314
 numerical_approx() (sage.matrix.matrix2.Matrix method), 220
 numpy() (sage.matrix.matrix1.Matrix method), 112
 numpy() (sage.matrix.matrix_double_dense.Matrix_double_dense method), 380

O

one() (sage.matrix.matrix_space.MatrixSpace method), 10
 ones_matrix() (in module sage.matrix.constructor), 39
 OperationTable (class in sage.matrix.operation_table), 403

P

p (sage.matrix.matrix_modn_sparse.Matrix_modn_sparse attribute), 305
 permanent() (sage.matrix.matrix2.Matrix method), 221
 permanental_minor() (sage.matrix.matrix2.Matrix method), 222
 permutation_normal_form() (sage.matrix.matrix2.Matrix method), 223

`permute_columns()` (sage.matrix.matrix0.Matrix method), 86
`permute_rows()` (sage.matrix.matrix0.Matrix method), 86
`permute_rows_and_columns()` (sage.matrix.matrix0.Matrix method), 87
`pfaffian()` (sage.matrix.matrix2.Matrix method), 224
`pivot_rows()` (sage.matrix.matrix2.Matrix method), 226
`pivots()` (sage.matrix.matrix0.Matrix method), 88
`pivots()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 335
`pivots()` (sage.matrix.matrix_mpolynomial_dense.Matrix_mpolynomial_dense method), 402
`plot()` (sage.matrix.matrix2.Matrix method), 226
`prepare()` (in module sage.matrix.constructor), 41
`prepare_dict()` (in module sage.matrix.constructor), 41
`prod_of_row_sums()` (sage.matrix.matrix2.Matrix method), 227
`prod_of_row_sums()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 335
`prod_of_row_sums()` (sage.matrix.matrix_rational_dense.Matrix_rational_dense method), 349

Q

`QR()` (sage.matrix.matrix2.Matrix method), 131
`QR()` (sage.matrix.matrix_double_dense.Matrix_double_dense method), 356

R

`random_diagonalizable_matrix()` (in module sage.matrix.constructor), 41
`random_echelonizable_matrix()` (in module sage.matrix.constructor), 44
`random_element()` (sage.matrix.matrix_space.MatrixSpace method), 11
`random_matrix()` (in module sage.matrix.constructor), 46
`random_rref_matrix()` (in module sage.matrix.constructor), 51
`random_subspaces_matrix()` (in module sage.matrix.constructor), 53
`random_unimodular_matrix()` (in module sage.matrix.constructor), 55
`randomize()` (sage.matrix.matrix2.Matrix method), 227
`randomize()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 335
`randomize()` (sage.matrix.matrix_modn_dense.Matrix_modn_dense method), 301
`randomize()` (sage.matrix.matrix_rational_dense.Matrix_rational_dense method), 349
`rank()` (sage.matrix.matrix0.Matrix method), 88
`rank()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 336
`rank()` (sage.matrix.matrix_modn_dense.Matrix_modn_dense method), 301
`rank()` (sage.matrix.matrix_modn_sparse.Matrix_modn_sparse method), 305
`rank()` (sage.matrix.matrix_rational_dense.Matrix_rational_dense method), 350
`rank2_GF()` (in module sage.matrix.benchmark), 419
`rank2_ZZ()` (in module sage.matrix.benchmark), 419
`rank_GF()` (in module sage.matrix.benchmark), 420
`rank_ZZ()` (in module sage.matrix.benchmark), 420
`rational_form()` (sage.matrix.matrix2.Matrix method), 228
`rational_reconstruction()` (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 336
`report()` (in module sage.matrix.benchmark), 420
`report_GF()` (in module sage.matrix.benchmark), 420
`report_ZZ()` (in module sage.matrix.benchmark), 421
`rescale_col()` (sage.matrix.matrix0.Matrix method), 88
`rescale_row()` (sage.matrix.matrix0.Matrix method), 89
`restrict()` (sage.matrix.matrix2.Matrix method), 234
`restrict_codomain()` (sage.matrix.matrix2.Matrix method), 234
`restrict_domain()` (sage.matrix.matrix2.Matrix method), 235

[right_eigenmatrix\(\)](#) (sage.matrix.matrix2.Matrix method), 235
[right_eigenspaces\(\)](#) (sage.matrix.matrix2.Matrix method), 236
[right_eigenvectors\(\)](#) (sage.matrix.matrix2.Matrix method), 239
[right_eigenvectors\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 381
[right_kernel\(\)](#) (sage.matrix.matrix2.Matrix method), 239
[right_kernel_matrix\(\)](#) (sage.matrix.matrix2.Matrix method), 244
[right_nullity\(\)](#) (sage.matrix.matrix2.Matrix method), 253
[rook_vector\(\)](#) (sage.matrix.matrix2.Matrix method), 253
[round\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 382
[row\(\)](#) (sage.matrix.matrix1.Matrix method), 112
[row\(\)](#) (sage.matrix.matrix_rational_dense.Matrix_rational_dense method), 350
[row_keys\(\)](#) (sage.matrix.operation_table.OperationTable method), 410
[row_module\(\)](#) (sage.matrix.matrix2.Matrix method), 253
[row_space\(\)](#) (sage.matrix.matrix2.Matrix method), 253
[row_space\(\)](#) (sage.matrix.matrix_space.MatrixSpace method), 11
[rows\(\)](#) (sage.matrix.matrix1.Matrix method), 113
[rref\(\)](#) (sage.matrix.matrix2.Matrix method), 254

S

[sage.matrix.benchmark](#) (module), 413
[sage.matrix.berlekamp_massey](#) (module), 279
[sage.matrix.constructor](#) (module), 15
[sage.matrix.docs](#) (module), 59
[sage.matrix.matrix](#) (module), 67
[sage.matrix.matrix0](#) (module), 69
[sage.matrix.matrix1](#) (module), 101
[sage.matrix.matrix2](#) (module), 123
[sage.matrix.matrix_complex_double_dense](#) (module), 397
[sage.matrix.matrix_dense](#) (module), 281
[sage.matrix.matrix_double_dense](#) (module), 353
[sage.matrix.matrix_generic_dense](#) (module), 293
[sage.matrix.matrix_generic_sparse](#) (module), 295
[sage.matrix.matrix_integer_dense](#) (module), 317
[sage.matrix.matrix_modn_dense](#) (module), 297
[sage.matrix.matrix_modn_sparse](#) (module), 303
[sage.matrix.matrix_mpolynomial_dense](#) (module), 399
[sage.matrix.matrix_rational_dense](#) (module), 343
[sage.matrix.matrix_real_double_dense](#) (module), 395
[sage.matrix.matrix_space](#) (module), 3
[sage.matrix.matrix_sparse](#) (module), 285
[sage.matrix.matrix_symbolic_dense](#) (module), 307
[sage.matrix.operation_table](#) (module), 403
[sage.matrix.strassen](#) (module), 275
[saturation\(\)](#) (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 337
[schur\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 382
[set_block\(\)](#) (sage.matrix.matrix2.Matrix method), 255
[set_col_to_multiple_of_col\(\)](#) (sage.matrix.matrix0.Matrix method), 90
[set_column\(\)](#) (sage.matrix.matrix1.Matrix method), 114
[set_immutable\(\)](#) (sage.matrix.matrix0.Matrix method), 90
[set_max_cols\(\)](#) (in module sage.matrix.matrix0), 98

`set_max_rows()` (in module `sage.matrix.matrix0`), 98
`set_print_symbols()` (`sage.matrix.operation_table.OperationTable` method), 410
`set_row()` (`sage.matrix.matrix1.Matrix` method), 115
`set_row_to_multiple_of_row()` (`sage.matrix.matrix0.Matrix` method), 91
`set_row_to_multiple_of_row()` (`sage.matrix.matrix_rational_dense.Matrix_rational_dense` method), 350
`simplify()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 314
`simplify_rational()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 314
`simplify_trig()` (`sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense` method), 314
`singular_values()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 386
`smith_form()` (`sage.matrix.matrix2.Matrix` method), 256
`smith_form()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 338
`smithform_ZZ()` (in module `sage.matrix.benchmark`), 421
`solve_left()` (`sage.matrix.matrix2.Matrix` method), 257
`solve_left()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 388
`solve_left_LU()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 390
`solve_right()` (`sage.matrix.matrix2.Matrix` method), 258
`solve_right()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 391
`sparse_columns()` (`sage.matrix.matrix1.Matrix` method), 115
`sparse_matrix()` (`sage.matrix.matrix1.Matrix` method), 116
`sparse_rows()` (`sage.matrix.matrix1.Matrix` method), 117
`stack()` (`sage.matrix.matrix1.Matrix` method), 117
`stack()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 339
`stack()` (`sage.matrix.matrix_sparse.Matrix_sparse` method), 290
`str()` (`sage.matrix.matrix0.Matrix` method), 91
`strassen_echelon()` (in module `sage.matrix.strassen`), 276
`strassen_window_multiply()` (in module `sage.matrix.strassen`), 277
`subdivide()` (`sage.matrix.matrix2.Matrix` method), 260
`subdivision()` (`sage.matrix.matrix2.Matrix` method), 262
`subdivision_entry()` (`sage.matrix.matrix2.Matrix` method), 262
`subdivisions()` (`sage.matrix.matrix2.Matrix` method), 263
`submatrix()` (`sage.matrix.matrix1.Matrix` method), 120
`subs()` (`sage.matrix.matrix2.Matrix` method), 263
`SVD()` (`sage.matrix.matrix_double_dense.Matrix_double_dense` method), 358
`swap_columns()` (`sage.matrix.matrix0.Matrix` method), 92
`swap_rows()` (`sage.matrix.matrix0.Matrix` method), 92
`swap_rows()` (`sage.matrix.matrix_modn_sparse.Matrix_modn_sparse` method), 306
`swapped_columns()` (`sage.matrix.matrix_mpolynomial_dense.Matrix_mpolynomial_dense` method), 402
`symplectic_form()` (`sage.matrix.matrix2.Matrix` method), 263
`symplectic_form()` (`sage.matrix.matrix_integer_dense.Matrix_integer_dense` method), 340

T

`T` (`sage.matrix.matrix2.Matrix` attribute), 134
`table()` (`sage.matrix.operation_table.OperationTable` method), 411
`tensor_product()` (`sage.matrix.matrix2.Matrix` method), 264
`test()` (in module `sage.matrix.strassen`), 278
`test_trivial_matrices_inverse()` (in module `sage.matrix.matrix_space`), 13
`to_list()` (`sage.matrix.strassen.int_range` method), 276
`trace()` (`sage.matrix.matrix2.Matrix` method), 265
`trace_of_product()` (`sage.matrix.matrix2.Matrix` method), 266
`translation()` (`sage.matrix.operation_table.OperationTable` method), 411

[transpose\(\)](#) (sage.matrix.matrix_dense.Matrix_dense method), 283
[transpose\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 393
[transpose\(\)](#) (sage.matrix.matrix_integer_dense.Matrix_integer_dense method), 341
[transpose\(\)](#) (sage.matrix.matrix_modn_sparse.Matrix_modn_sparse method), 306
[transpose\(\)](#) (sage.matrix.matrix_rational_dense.Matrix_rational_dense method), 351
[transpose\(\)](#) (sage.matrix.matrix_sparse.Matrix_sparse method), 290
[tune_multiplication\(\)](#) (in module sage.matrix.matrix_integer_dense), 341

U

[unpickle\(\)](#) (in module sage.matrix.matrix0), 99

V

[variables\(\)](#) (sage.matrix.matrix_symbolic_dense.Matrix_symbolic_dense method), 315
[vecmat_ZZ\(\)](#) (in module sage.matrix.benchmark), 421
[vector_on_axis_rotation_matrix\(\)](#) (in module sage.matrix.constructor), 57
[visualize_structure\(\)](#) (sage.matrix.matrix2.Matrix method), 266
[visualize_structure\(\)](#) (sage.matrix.matrix_modn_sparse.Matrix_modn_sparse method), 306

W

[weak_popov_form\(\)](#) (sage.matrix.matrix2.Matrix method), 266
[wiedemann\(\)](#) (sage.matrix.matrix2.Matrix method), 268
[with_added_multiple_of_column\(\)](#) (sage.matrix.matrix0.Matrix method), 93
[with_added_multiple_of_row\(\)](#) (sage.matrix.matrix0.Matrix method), 93
[with_col_set_to_multiple_of_col\(\)](#) (sage.matrix.matrix0.Matrix method), 93
[with_permuted_columns\(\)](#) (sage.matrix.matrix0.Matrix method), 94
[with_permuted_rows\(\)](#) (sage.matrix.matrix0.Matrix method), 94
[with_permuted_rows_and_columns\(\)](#) (sage.matrix.matrix0.Matrix method), 95
[with_rescaled_col\(\)](#) (sage.matrix.matrix0.Matrix method), 96
[with_rescaled_row\(\)](#) (sage.matrix.matrix0.Matrix method), 96
[with_row_set_to_multiple_of_row\(\)](#) (sage.matrix.matrix0.Matrix method), 97
[with_swapped_columns\(\)](#) (sage.matrix.matrix0.Matrix method), 97
[with_swapped_rows\(\)](#) (sage.matrix.matrix0.Matrix method), 98

Z

[zero\(\)](#) (sage.matrix.matrix_space.MatrixSpace method), 11
[zero_at\(\)](#) (sage.matrix.matrix_double_dense.Matrix_double_dense method), 393
[zero_matrix\(\)](#) (in module sage.matrix.constructor), 58
[zero_matrix\(\)](#) (sage.matrix.matrix_space.MatrixSpace method), 12
[zigzag_form\(\)](#) (sage.matrix.matrix2.Matrix method), 268