Sage Reference Manual: Parallel Computing

Release 6.3

The Sage Development Team

August 11, 2014

CONTENTS

DECORATE INTERFACE FOR PARALLEL COMPUTATION

class `sage.parallel.decorate.`**`Fork`**(*timeout=0*, *verbose=False*)
>   A `fork` decorator class.

class `sage.parallel.decorate.`**`Parallel`**(*p_iter='fork'*, *ncpus=None*, ***kwds*)
>   Create a `parallel`-decorated function. This is the object created by `parallel()`.

class `sage.parallel.decorate.`**`ParallelFunction`**(*parallel*, *func*)
>   Bases: `object`
>
>   Class which parallelizes a function or class method. This is typically accessed indirectly through `Parallel.__call__()`.

`sage.parallel.decorate.`**`fork`**(*f=None*, *timeout=0*, *verbose=False*)
>   Decorate a function so that when called it runs in a forked subprocess. This means that it won't have any in-memory side effects on the parent Sage process. The pexpect interfaces are all reset.
>
>   INPUT:
>
>   > - `f` – a function
>   >
>   > - `timeout` – (default: 0) if positive, kill the subprocess after this many seconds (wall time)
>   >
>   > - `verbose` – (default: `False`) whether to print anything about what the decorator does (e.g., killing the subprocess)
>
>   > **Warning:** The forked subprocess will not have access to data created in pexpect interfaces. This behavior with respect to pexpect interfaces is very important to keep in mind when setting up certain computations. It's the one big limitation of this decorator.
>
>   EXAMPLES:
>
>   We create a function and run it with the `fork` decorator. Note that it does not have a side effect. Despite trying to change the global variable `a` below in `g`, the variable `a` does not get changed:
>
> ```
> sage: a = 5
> sage: @fork
> ... def g(n, m):
> ...        global a
> ...        a = 10
> ...        return factorial(n).ndigits() + m
> sage: g(5, m=5)
> 8
> ```

```
sage: a
5
```

We use `fork` to make sure that the function terminates after one second, no matter what:

```
sage: @fork(timeout=1, verbose=True)
... def g(n, m): return factorial(n).ndigits() + m
sage: g(5, m=5)
8
sage: g(10^7, m=5)
Killing subprocess ... with input ((10000000,), {'m': 5}) which took too long
'NO DATA (timed out)'
```

We illustrate that the state of the pexpect interface is not altered by forked functions (they get their own new pexpect interfaces!):

```
sage: gp.eval('a = 5')
'5'
sage: @fork()
... def g():
...         gp.eval('a = 10')
...         return gp.eval('a')
sage: g()
'10'
sage: gp.eval('a')
'5'
```

We illustrate that the forked function has its own pexpect interface:

```
sage: gp.eval('a = 15')
'15'
sage: @fork()
... def g(): return gp.eval('a')
sage: g()
'a'
```

We illustrate that segfaulting subprocesses are no trouble at all:

```
sage: cython('def f(): print <char*>0')
sage: @fork
... def g(): f()
sage: print "this works"; g()
this works...

--------------------------------------------------------------------
Unhandled SIG...
--------------------------------------------------------------------
'NO DATA'
```

sage.parallel.decorate.**normalize_input**(*a*)

Convert `a` to a pair `(args, kwds)` using some rules:

   •if already of that form, leave that way.

   •if `a` is a tuple make `(a, {})`

   •if `a` is a dict make `(tuple([]),a)`

   •otherwise make `((a,),{})`

INPUT:

- a – object

OUTPUT:

- `args` – tuple

- `kwds` – dictionary

EXAMPLES:
```
sage: sage.parallel.decorate.normalize_input( (2, {3:4}) )
((2, {3: 4}), {})
sage: sage.parallel.decorate.normalize_input( (2,3) )
((2, 3), {})
sage: sage.parallel.decorate.normalize_input( {3:4} )
((), {3: 4})
sage: sage.parallel.decorate.normalize_input( 5 )
((5,), {})
```

`sage.parallel.decorate.`**`parallel`**(*p_iter='fork'*, *ncpus=None*, ***kwds*)

This is a decorator that gives a function a parallel interface, allowing it to be called with a list of inputs, whose values will be computed in parallel.

> **Warning:** The parallel subprocesses will not have access to data created in pexpect interfaces. This behavior with respect to pexpect interfaces is very important to keep in mind when setting up certain computations. It's the one big limitation of this decorator.

INPUT:

- **`p_iter` – parallel iterator function or string:**

    - `'fork'` – (default) use a new forked subprocess for each input

    - `'multiprocessing'` – use multiprocessing library

    - `'reference'` – use a fake serial reference implementation

- `ncpus` – integer, maximal number of subprocesses to use at the same time

- `timeout` – number of seconds until each subprocess is killed (only supported by 'fork'; zero means not at all)

> **Warning:** If you use anything but `'fork'` above, then a whole new subprocess is spawned, so none of your local state (variables, certain functions, etc.) is available.

EXAMPLES:

We create a simple decoration for a simple function. The number of cpus (or cores, or hardware threads) is automatically detected:
```
sage: @parallel
... def f(n): return n*n
sage: f(10)
100
sage: sorted(list(f([1,2,3])))
[(((1,), {}), 1), (((2,), {}), 4), (((3,), {}), 9)]
```

We use exactly two cpus:
```
sage: @parallel(2)
... def f(n): return n*n
```

We create a decorator that uses three subprocesses, and times out individual processes after 10 seconds:

```
sage: @parallel(ncpus=3, timeout=10)
...  def fac(n): return factor(2^n-1)
sage: for X, Y in sorted(list(fac([101,119,151,197,209]))): print X,Y
((101,), {}) 7432339208719 * 341117531003194129
((119,), {}) 127 * 239 * 20231 * 131071 * 62983048367 * 131105292137
((151,), {}) 18121 * 55871 * 165799 * 2332951 * 7289088383388253664437433
((197,), {}) 7487 * 26828803997912886929710867041891989490486893845712448833
((209,), {}) 23 * 89 * 524287 * 94803416684681 * 1512348937147247 * 5346950541323960232319657

sage: @parallel('multiprocessing')
...  def f(N): return N^2
sage: v = list(f([1,2,4])); v.sort(); v
[(((1,), {}), 1), (((2,), {}), 4), (((4,), {}), 16)]
sage: @parallel('reference')
...  def f(N): return N^2
sage: v = list(f([1,2,4])); v.sort(); v
[(((1,), {}), 1), (((2,), {}), 4), (((4,), {}), 16)]
```

For functions that take multiple arguments, enclose the arguments in tuples when calling the parallel function:

```
sage: @parallel
...  def f(a,b): return a*b
sage: for X, Y in sorted(list(f([(2,3),(3,5),(5,7)]))): print X, Y
((2, 3), {}) 6
((3, 5), {}) 15
((5, 7), {}) 35
```

For functions that take a single tuple as an argument, enclose it in an additional tuple at call time, to distinguish it as the first argument, as opposed to a tuple of arguments:

```
sage: @parallel
...  def firstEntry(aTuple): return aTuple[0]
sage: for X, Y in sorted(list(firstEntry([((1,2,3,4),),((5,6,7,8),)]))): print X, Y
(((1, 2, 3, 4),), {}) 1
(((5, 6, 7, 8),), {}) 5
```

The parallel decorator also works with methods, classmethods, and staticmethods. Be sure to apply the parallel decorator after ("above") either the classmethod or staticmethod decorators:

```
sage: class Foo(object):
...        @parallel(2)
...        def square(self, n):
...            return n*n
...        @parallel(2)
...        @classmethod
...        def square_classmethod(cls, n):
...            return n*n
sage: a = Foo()
sage: a.square(3)
9
sage: sorted(a.square([2,3]))
[(((2,), {}), 4), (((3,), {}), 9)]
sage: Foo.square_classmethod(3)
9
sage: sorted(Foo.square_classmethod([2,3]))
[(((2,), {}), 4), (((3,), {}), 9)]
sage: Foo.square_classmethod(3)
9
```

Warning: Currently, parallel methods do not work with the multiprocessing implementation.

REFERENCE PARALLEL PRIMITIVES

These are reference implementations of basic parallel primitives. These are not actually parallel, but work the same way. They are good for testing.

sage.parallel.reference.**parallel_iter**(*f*, *inputs*)
Reference parallel iterator implementation.

INPUT:

•f – a Python function that can be pickled using the pickle_function command.

•inputs – a list of pickleable pairs (args, kwds), where args is a tuple and kwds is a dictionary.

OUTPUT:

•iterator over 2-tuples (inputs[i], f(inputs[i])), where the order may be completely random

EXAMPLES:
```
sage: def f(N,M=10): return N*M
sage: inputs = [((2,3),{}),  (tuple([]), {'N':3,'M':5}), ((2,),{})]
sage: set_random_seed(0)
sage: for a, val in sage.parallel.reference.parallel_iter(f, inputs):
...       print a, val
((2,), {}) 20
((), {'M': 5, 'N': 3}) 15
((2, 3), {}) 6
sage: for a, val in sage.parallel.reference.parallel_iter(f, inputs):
...       print a, val
((), {'M': 5, 'N': 3}) 15
((2,), {}) 20
((2, 3), {}) 6
```

PARALLEL ITERATOR BUILT USING THE `FORK()` SYSTEM CALL

class `sage.parallel.use_fork.`**`p_iter_fork`**(*ncpus*, *timeout=0*, *verbose=False*, *reset_interfaces=True*)

A parallel iterator implemented using `fork()`.

EXAMPLES:

```
sage: X = sage.parallel.use_fork.p_iter_fork(2,3, False); X
<sage.parallel.use_fork.p_iter_fork instance at ...>
sage: X.ncpus
2
sage: X.timeout
3.0
sage: X.verbose
False
```

PARALLEL ITERATOR BUILT USING PYTHON'S MULTIPROCESSING MODULE

sage.parallel.multiprocessing_sage.**parallel_iter**(*processes*, *f*, *inputs*)
> Return a parallel iterator.

> INPUT:

>> •`processes` – integer

>> •`f` – function

>> •`inputs` – an iterable of pairs (args, kwds)

> OUTPUT:

>> •iterator over values of `f` at `args,kwds` in some random order.

> EXAMPLES:

```
sage: def f(x): return x+x
sage: import sage.parallel.multiprocessing_sage
sage: v = list(sage.parallel.multiprocessing_sage.parallel_iter(2, f, [((2,), {}), ((3,),{})]))
sage: v.sort(); v
[(((2,), {}), 4), (((3,), {}), 6)]
```

sage.parallel.multiprocessing_sage.**pyprocessing**(*processes=0*)
> Return a parallel iterator using a given number of processes implemented using pyprocessing.

> INPUTS:

>> •`processes` – integer (default: 0); if 0, set to the number of processors on the computer.

> OUTPUT:

>> •a (partially evaluated) function

> EXAMPLES:

```
sage: from sage.parallel.multiprocessing_sage import pyprocessing
sage: p_iter = pyprocessing(4)
sage: P = parallel(p_iter=p_iter)
sage: def f(x): return x+x
sage: v = list(P(f)(range(10))); v.sort(); v
[(((0,), {}), 0), (((1,), {}), 2), (((2,), {}), 4), (((3,), {}), 6), (((4,), {}), 8), (((5,), {}
```

CPU DETECTION

sage.parallel.ncpus.**ncpus**()

Detects the number of effective CPUs in the system.

EXAMPLES:

```
sage: sage.parallel.ncpus.ncpus()   # random output -- depends on machine.
2
```

INDICES AND TABLES

- Index
- Module Index
- Search Page

PYTHON MODULE INDEX

p

INDEX

F

N

P

S